


The x86 Microprocessors

PEARSON



*Architecture, Programming and Interfacing
(8086 to Pentium)*

Lyla B. Das

THE x86 MICROPROCESSORS

Two roads diverged in a wood, and I —
I took the one less traveled by,
And that has made all the difference.

—ROBERT FROST

THE X86 MICROPROCESSORS

*Architecture, Programming and Interfacing
(8086 to Pentium)*



LYLA B DAS

ASSOCIATE PROFESSOR
DEPARTMENT OF ELECTRONICS AND COMMUNICATION ENGINEERING
NATIONAL INSTITUTE OF TECHNOLOGY CALICUT
KOZHIKODE
KERALA

PEARSON

Chennai • Delhi • Chandigarh

Publishing Manager: K. Srinivas
Senior Managing Editor: Thomas Mathew Rajesh
Acquisitions Editor: Sojan Jose
Assistant Acquisitions Editor: S. Shankari
Managing Editor, Production Editorial: Shadan Perween
Senior Production Editor: M.E. Sethurajan
Production Editor: M.R. Ramesh
Cover Design: Rahul Sharma
Interior Design: Syed Dilshad Ali

General Manager, Marketing: J. Saravanan
Marketing Manager: Vikram Singh
Senior Rights Manager: Sumita Roy
VP, Production: Subhasis Ganguli
Deputy General Manager, Design and Manufacturing: Madhur Bhatia
Assistant Manager, Manufacturing: Priti Singh
Composition: White Lotus Infotech Pvt. Ltd., Pondicherry
Printer: Chennai Micro Print, Chennai

The pin diagrams, timing diagrams and internal architecture of x86 processors are reproduced in this book from the manuals of Intel with the permission of Intel Corporation, California.

Copyright © 2010 Dorling Kindersley (India) Pvt. Ltd.

This book is sold subject to the condition that it shall not, by way of trade or otherwise, be lent, resold, hired out, or otherwise circulated without the publisher's prior written consent in any form of binding or cover other than that in which it is published and without a similar condition including this condition being imposed on the subsequent purchaser and without limiting the rights under copyright reserved above, no part of this publication may be reproduced, stored in or introduced into a retrieval system, or transmitted in any form or by any means (electronic, mechanical, photocopying, recording or otherwise), without the prior written permission of both the copyright owner and the publisher of this book.

ISBN: 978-81-317-3246-5

10 9 8 7 6 5 4 3 2 1

Published by Dorling Kindersley (India) Pvt. Ltd., licensees of Pearson Education in South Asia.

Head Office: 7th Floor, Knowledge Boulevard, A-8 (A), Sector 62, Noida 201309, UP, India.
Registered Office: 11 Community Centre, Panchsheel Park, New Delhi 110 017, India.

This book is dedicated
to all my students
of the past, present and future,
and to my children, Sagar and Thushar.

CONTENTS

	<i>Preface</i>	xv
	<i>About the Author</i>	xix
0	Basics of Computer Systems	1
	0.1 A Brief History of Microprocessors	1
	0.2 Basics of Computer Architecture	5
	0.3 Computer Languages	10
	0.4 RISC and CISC Architectures	12
	0.5 Number Systems	12
	0.6 Number Format Conversions	15
	0.7 Computer Arithmetic	22
	0.8 Units of Memory Capacity	30
	0.9 The 8085 Microprocessor	31
	<i>Key Points of this Chapter</i>	47
	<i>Questions</i>	47
	<i>Exercise</i>	48
1	The Architecture of 8086	51
	1.1 Internal Block Diagram of the 8086	51
	1.2 The Execution Unit	51
	1.3 Bus Interface Unit	58
	1.4 Addressing Modes	63
	<i>Key Points of this Chapter</i>	68
	<i>Questions</i>	69
	<i>Exercise</i>	69
2	Programming Concepts – I	71
	2.1 The Assembly Process	71
	2.2 Assemblers for x86	73
	2.3 Memory Models	75

2.4	Instruction Design	86
	<i>Key Points of this Chapter</i>	93
	<i>Questions</i>	93
	<i>Exercise</i>	94

3 Programming Concepts – II 95

3.1	Approaches to Programming	95
3.2	Data Transfer Instructions	98
3.3	Branch Instructions	107
3.4	Arithmetic Instructions	113
3.5	Logical Instructions	129
3.6	Shift and Rotate Instructions	131
	<i>Key Points of this Chapter</i>	136
	<i>Questions</i>	137
	<i>Exercise</i>	137

4 Programming Concepts – III 139

4.1	String Instructions	139
4.2	Procedures	145
4.3	Macros	155
4.4	Number Format Conversions	158
4.5	ASCII Operations	162
4.6	Conversions for Computations and Display/Entry	165
4.7	Signed Number Arithmetic	166
4.8	Programming Using High Level Language Constructs	172
	<i>Key Points of this Chapter</i>	176
	<i>Questions</i>	177
	<i>Exercise</i>	177

5 Programming Concepts – IV 179

5.1	Input/Output Programming	179
5.2	I/O Instructions	180
5.3	Modular Programming	183
5.4	Programming in C with Assembly Modules	189

<i>Key Points of this Chapter</i>	193
<i>Questions</i>	193
<i>Exercise</i>	193

6	The Hardware Structure of 8086	195
6.1	Pin Configuration	195
6.2	Clock	208
6.3	Other Processor Activities	211
6.4	Maximum Mode	214
6.5	Instruction Cycle	218
	<i>Key Points of this Chapter</i>	223
	<i>Questions</i>	223
	<i>Exercise</i>	224

7	Memory and I/O Decoding	225
7.1	Memory Device Pins	225
7.2	Memory Address Decoding	227
7.3	Memory Banks	235
7.4	I/O Address Decoding	239
	<i>Key Points of this Chapter</i>	246
	<i>Questions</i>	246
	<i>Exercise</i>	246

8	The Interrupt Structure of 8086	249
8.1	Interrupts of 8086	250
8.2	Dedicated Interrupt Types	253
8.3	Software Interrupts	255
8.4	Hardware Interrupts	256
8.5	Priority of Interrupts	259
8.6	Interrupt Type Allocation for Current PCs	259
8.7	BIOS 10H Functions	264
8.8	Addressing Video Memory Directly	270
8.9	Keyboard Interfacing	272
8.10	Hooking an Interrupt	277

<i>Key Points of this Chapter</i>	286
<i>Questions</i>	287
<i>Exercise</i>	287

9 Peripheral Interfacing – I 289

9.1	Trainer Kit	290
9.2	Programmable Peripheral Interface (PPI)-8255A	291
9.3	Modes of Operation	296
9.4	Mode 0	296
9.5	Mode 1	302
9.6	Mode 2 (Strobed Bidirectional Bus I/O)	310
9.7	Centronics Printer Interface	310
9.8	Interfacing an Analog to Digital Converter to the 8086	313
9.9	Interfacing to a Digital to Analog Converter	317
9.10	Interfacing Liquid Crystal Displays to the 8086	322
9.11	Interfacing a Stepper Motor to the 8086	327
9.12	Hex Keyboard Interfacing	337
9.13	Interfacing Led Displays	341
	<i>Key Points of this Chapter</i>	347
	<i>Questions</i>	348
	<i>Exercise</i>	348

10 Peripheral Interfacing – II 349

10.1	The Programmable Interval Timer 8253/8254	349
10.2	The Programmable Keyboard Display Interface – 8279	363
10.3	The Programmable Interrupt Controller (PIC) 8259	377
10.4	Cascade Mode	390
	<i>Key Points of this Chapter</i>	393
	<i>Questions</i>	393
	<i>Exercise</i>	394

11 Peripheral Interfacing – III 395

11.1	Serial Communication Principles	395
11.2	Simplex, Half Duplex and Full Duplex Communication	395
11.3	The Programmable Serial Communication Interface	403

11.4	Internal Reset on Power Up	414
11.5	Direct Memory Access	415
11.6	The DMA Controller – 8237	419
11.7	DMA and IBM-PC	426
11.8	PCI Based Computers	428
	<i>Key Points of this Chapter</i>	428
	<i>Questions</i>	429
	<i>Exercise</i>	429

12	Semiconductor Memory Devices	431
12.1	Semiconductor Memory	432
12.2	Dynamic RAM	435
12.3	Synchronous DRAM (SDRAM)	440
12.4	ROM (Read Only Memory)	443
12.5	Cache Memory	444
12.6	Mapping Techniques	447
12.7	Cache and the x86 Family	451
	<i>Key Points of this Chapter</i>	453
	<i>Questions</i>	453
	<i>Exercise</i>	454

13	Multiprocessor Configurations	455
13.1	Multiprocessor Systems	456
13.2	Multiprocessing Using 8086	457
13.3	The 8086 and 8089 in a Tightly Coupled Configuration	462
13.4	Loosely Coupled Configurations and Bus Arbitration	464
13.5	Bus Arbitration Using the 8289 Bus Arbiter IC	466
13.6	The Arithmetic Co-Processor 8087	471
	<i>Key Points of this Chapter</i>	483
	<i>Questions</i>	484
	<i>Exercise</i>	484

14	80186 – The Embedded Microprocessor	487
14.1	Additions in the Instruction Set	488
14.2	Instruction Set Enhancements	490

14.3	Block Diagram of the 80186	492
14.4	Programming the Timer Unit	495
14.5	Programming	497
	<i>Key Points of this Chapter</i>	502
	<i>Questions</i>	502
	<i>Exercise</i>	502

15 The 80286 and 80386 Processors 503

15.1	The 80286 Processor	503
15.2	The 80386	504
15.3	Internal Architecture	505
15.4	Programming Enhancements	506
15.5	Hardware Features of 80386	510
15.6	Virtual Memory	514
15.7	Memory Management Unit	515
15.8	Converting a Logical Address to a Physical Address	522
15.9	Calculating the size of the Logical Address Space	523
15.10	Protection	528
15.11	Multi Tasking	534
15.12	Interrupts of 80386	538
15.13	Privileged Instructions	539
15.14	Conclusion	541
	<i>Key Points of this Chapter</i>	541
	<i>Questions</i>	542
	<i>Exercise</i>	543

16 The Pentium Processor 545

16.1	The Enhanced Features of 80486	545
16.2	Data Alignment	548
16.3	The Pentium Processor	549
16.4	Pentium Pro	553
16.5	Pentium-II And Pentium-III	554
16.6	Pentium-IV	555
16.7	Latest Trends in Microprocessor Design	555
16.8	Multi-Core Technology and Intel	558
16.9	Mobile Processors	559

16.10	Legacy Support	559
	<i>Key Points of this Chapter</i>	559
	<i>Questions</i>	560
	<i>Exercise</i>	560

17 The x86 Based Personal Computer 561

17.1	The Modern PC	562
17.2	The Mother Board	562
17.3	Chipset	564
17.4	Transfer Speed	566
17.5	Expansion Buses	568
17.6	ATA	573
17.7	Memory – SIMM and DIMM	576
17.8	System BIOS	578
17.9	New Motherboards	580
17.10	Other I/O Devices	582
17.11	PS/2	582
17.12	Form Factors	582
17.13	Laptops	583
	<i>Key Points of this Chapter</i>	585
	<i>Questions</i>	586
	<i>Exercise</i>	586

<i>Appendix A</i>	587
<i>Appendix B</i>	599
<i>Appendix C</i>	607
<i>Appendix D</i>	617
<i>Appendix E</i>	625
<i>Bibliography</i>	635
<i>Index</i>	637

PREFACE

Preamble

In my twenty-five years of teaching experience (almost), I have taught subjects as varied as basic electronics, electronic circuits, digital signal processing, communications, information theory, digital image processing, computer architecture, computer programming and microprocessors, to name just a few. However, microprocessors, microcontrollers, assembly language programming and hardware interfacing caught my fancy and interest, at some point in time. This book is a result of my continued interest in these topics.

Although I have used a number of text books for teaching microprocessors, I have felt that something is missing in most of them – either they do not touch upon programming concepts well or their approach to the x86 series of microprocessors creates an impression that it is tough to understand and manage, and that assembly language programming is unfriendly and difficult to master. In this book, I have tried to eliminate these shortcomings by delineating the concepts in a step-by-step approach, aiming to keep simplicity of ideas, lucidity of explanations and clarity in presentation as my guiding principles.

This book proffers a detailed study of the x86 family of microprocessors. The x86 family comprises Intel's 8086, 80186, 80286, 80386, 80486 and 80586 (Pentium) processors. The approach is to study the x86 family architecture based on the architecture of the elementary processor, i.e. the 8086. The higher-order processors are discussed based on the enhancements, improvements and differences vis-à-vis the basic 8086. This is the best approach to learning the family architecture and it is followed by students worldwide. Most PCs across the world use the x86 architecture. Hence, it is an important subject that is taught and learnt at the academic and at the professional level.

Prerequisite

Microprocessors form a key subject of study at the bachelor's level degree program of engineering, where it is taught as a core subject for all circuit-related branches, i.e. electronics, electrical, computer science and information technology. A prerequisite for mastering this subject is a course on logic design, implying that students need to know the basic building blocks of a digital system. A course on computer organization and architecture would be helpful to the student, but it is not mandatory for understanding the subject. However, not all institutions deal with computer architecture in their study modules before teaching microprocessors. Hence, this book is aimed at being the first introduction to microprocessors.

Approach

The theme of the book is centered around the architecture of the x86 microprocessor and a detailed study of assembly language programming and interfacing to external chips. Throughout the book, the emphasis is on ensuring that the reader can grasp concepts and ideas easily. To this end, solved examples, worked-out problems, tested programs and explanatory diagrams have been included.

Organization of the book

Students who start learning microprocessors would have already learnt binary and other number systems. However, years of teaching have convinced me that a fresh look at these concepts would be in order, to understand assembly language programming. That is why an elaborate treatment has been meted out to

these concepts in Chapter 0. It is important to be clear about topics like sign extension, signed arithmetic and BCD arithmetic. Readers would do well to be conversant with the concepts presented in Chapter 0. The discussion on the elements of computer architecture in this chapter is meant for those who have not studied this topic earlier. Similarly, the chapter also gives an overview of the 8085 processor for the benefit of those who have to learn about 8085 as part of their curriculum requirements. A study of the 8085 processor is not necessary to understand the x86 family of processors.

Chapter 1 explains the basic architecture of the 8086 processor. This chapter, fundamental to understanding the topics covered in the book, is profuse with numerical problems that explain important concepts.

Chapters 2 to 5 are devoted, for the most part, to assembly programming. Chapter 2 introduces the MASM assembler. We discuss Version 6.14, which is useful for effective assembly language programming. The steps for using the assembler to run programs are discussed in this chapter. The DOS and debugging commands of Appendix B may also be useful for getting a good grasp of programming skills. Chapter 5 includes an introduction to C programming with embedded assembly modules. These four chapters cover most of the instructions of the 8086 processor the use of which has been highlighted in solved examples. Adequate end-of-chapter questions have been provided, to ensure proficiency in programming. Advanced concepts like modular programming and high-level language constructs of MASM have also been described.

Chapters 6 and 7 are devoted to hardware. Chapter 6 talks about the pins of the 8086 processor and how these pins are used in the minimum and maximum mode configurations. Timing diagrams are introduced here. Chapter 7 elaborates on the techniques of address decoding and is important for understanding the hardware interfacing chapters that follow.

Chapter 8 is an interesting chapter because it introduces the concept of ‘interrupts’, which is an important theme in the study of computers. How hardware can be manipulated using software interrupts is explained here. Text mode video and TSR programming are introduced with practical worked-out examples. These help one to use the knowledge of assembly language to understand the PC.

Chapters 9 to 11 deal with the interfacing of the 8086 processor to various peripherals. A number of interfacing chips are introduced here. A detailed study of a few of these chips would stand the student in good stead. For example, understanding the 8255 chip would help us to learn about other peripheral chips with ease. These chapters have been designed to meet the needs of undergraduate students who use these chips for their laboratory work.

Chapter 12 discusses memory from the user’s point of view. State-of-the-art memory trends like synchronous dynamic random access memory (SDRAM) and terms like double data rate (DDR) have been explained. Chapter 13 is devoted to the basic principles of multiprocessing and bus arbitration techniques. In addition, it discusses floating point arithmetic and the use of the arithmetic co-processor, including the programming aspects.

Chapters 14 and 15 deal with the higher versions of the x86 processor family – the 80186 processor is described as an ‘embedded processor’ – while the rest of them are depicted as processors used for PCs. The 80386 processor, which is the actual fore-runner of the Pentium processor, has been dealt with in great detail. Difficult topics like address translation, protection and multitasking are elucidated for the student’s benefit. Repeated reading of this topic will sort out many of the initial difficulties faced by the reader.

In Chapter 16, the Pentium processor family and its enhanced features have been listed. The idea of ‘multi core’ processing has also been touched upon.

Chapter 17 is about the modern PC, with all its intricacies and mysteries unraveled. The features, parts, boards, buses, connectors and BIOS have been discussed in simple terms to help an average student understand the applications of the x86 processor. I have heard many people using terms such as chipset, bus and adapter without knowing what exactly these expressions mean. This chapter aims to clarify the meaning of such terminology in unambiguous words. Photographs of typical motherboards and other parts of the PC have been included in this chapter. It also elaborates on how the PC has developed over the years, how the hardware and standards have changed and evolved and also the current trends in the field. No doubt contemporary technology is bound to change again, in the times to come.

The book comes with five appendices and relate, in order, to the Intel manual of the 8086 processor, the use of DOS and debug commands, the instruction set and instruction timing of the 8086 processor, DOS and BIOS interrupt list and the instruction set of the 8087 processor. In addition, there is an appendix on the installation of MASM 32 and MASM 6.14 made available at the book’s companion Web site, www.pearsoned.co.in/llylabdas. Supplements to the book in the form of slides in PowerPoint can be accessed at the companion Web site. The appendices and PowerPoint slides add value to the book by disseminating additional information on selected topics to the discerning learner.

I hope I have effectively addressed the learning of the x86 family of microprocessors. I suggest that all teachers of this subject emphasize the use of an assembler and practical programming to make the topics more interesting.

Contact

Your feedback and suggestions for the improvement of this book are welcome. While every attempt has been made to eliminate errors in this book, a few may still have managed to creep in. Kindly point them out to me – my email id is lbd@nitc.ac.in.

ACKNOWLEDGEMENTS

As I complete writing this book, I realize, with a sense of deep humility that I have a lot of people to thank. I first thank Sojan Jose, editor at Pearson Education, who discovered the author in me. It was only his enthusiasm, encouragement and support that gave me the courage to embark upon this venture. Ramesh, Thomas and other team members of this project at Pearson Education have impressed me with their professionalism and I thank them wholeheartedly for all the work they have put in to complete this book.

After I finished writing the first few chapters, I used the material to teach a course to the fourth-semester students (B070EC batch). I remember that many students participated actively in the teaching–learning process, which, in effect, gave me tips and suggestions on how this book should be. I thank them all and place on record my appreciation of their curiosity and determination to delve into the subject beyond mere superficiality.

Many of my colleagues were instrumental in helping me in this venture. The discussions I had with Jagandand, EED, and Saidalvi, CSed, were fruitful and helped to evolve some of the topics discussed in this book. I am grateful to Anand, senior mechanic at the Embedded Systems laboratory, who assisted me in carrying out the hardware work associated with interfacing, and his expertise has indeed made the work easier for me. I thank Dr Lillykutty Jacob,

former head of the department of electronics and communications engineering, who reduced the quantum of topics that I had to teach for one semester and enabled me to speed up the writing of this book. I am indebted to Kishore, who shared my duties at the microprocessor lab so that I could find time to complete the last few chapters as scheduled. I am obliged to my M.Tech students, Shiny and Shaeen, for their help in reading through some of the chapters, and to Sneha, Nitin, Venkat and Divya for their assistance in a few hardware experiments that were conducted in the course of writing this book.

It was Dr Mohamed Rafiquzzaman (Professor, California State Polytechnic University, Pomona, USA) who introduced me to the MASM32 assembler and the Olly debugger that have been discussed in Chapter 15. I am indebted to him. I also thank Dr Krishna Vedula who was the chief organizer of the IUCEE workshop, where I got the opportunity to meet Prof. Rafi.

This book contains a lot of diagrams, and I was lucky to find a few people who could draw really well. I thank Shelitha, Beljith and Ranjusha for their help with the illustrations.

The last chapter of this book contains information on how the modern PC works. I thank Sajth, of Dot Computers, Calicut, for clarifying my doubts and queries on this topic and also for lending me the motherboard and components, which I have used for generating the images presented in that chapter. I appreciate the contribution of Rajesh of Raja Studio, Kattangal, who photographed the chips, boards and other components that feature in the opening pages of each chapter.

I acknowledge the support of my friends Dr Elizabeth Elias, Dr Sally George, Dr Jeevamma Jacob, Dr Sathidevi, Dr E. Gopinathan, Dr Suresh Babu, Dr Sreelekha and Dr Deepthi whose companionship has always been a source of great encouragement to me.

My department colleagues have always been helpful and I think I am lucky to be a part of this group of motivated individuals. I am deeply indebted to my institution for nurturing me for the past twenty-five years (almost) and giving me the freedom to grow.

I am happy that my family has always been a source of solace for me.

Last, but not the least, I thank my students. All these years, I have been inspired by them. I have wanted to learn and know more, only on account of their 'demand'. I hope I will continue to have such students in the future as well.

LYLA B. DAS

ABOUT THE AUTHOR

Lyla B. Das is Associate Professor, Department of Electronics Engineering, National Institute of Technology Calicut (NITC), Kerala. She has a diverse mix of industrial, teaching and research experience spanning about 30 years. As a young graduate specializing in Electronics and Communications from the College of Engineering, Trivandrum, Lyla B. Das joined Keltron Controls as Deputy Engineer in 1981. Four years later, she joined NITC (then Regional Engineering College, Calicut) as lecturer and proceeded to complete her master's degree in digital communications from the same college. Over the years, she was successively elevated as Assistant Professor and then Associate Professor, a position which she currently holds.

Keen to actively seek and impart knowledge, Lyla B. Das currently teaches courses on microprocessors, microcontrollers, digital system design using VHDL, and system design using embedded processors at the undergraduate as well as postgraduate level. She has presented research papers in conferences of national and international stature and has worked on numerous projects based on microprocessors and microcontrollers, such as microprocessor-based voting machines and microcontroller-based rail track switching system. An avid reader of contemporary research material, she keeps herself abreast of the current trends in her chosen field and guides students in their M. Tech. research theses.

Lyla B. Das has worked on various projects funded by the ministry of human resource development (MHRD) in thrust areas of growth including the setting up of an embedded systems laboratory in 2005–2008. She has delivered expert lectures on image compression using wavelets, advanced microprocessors and microcontrollers, FPGA-based systems and embedded systems at several engineering colleges across Kerala. She has also participated in numerous tutorials and workshops conducted by the Indian Institute of Technology (IIT) and the Indian Institute of Science (IISc). She was a Fellow in the national conference on 'VLSI Design and Embedded Systems' held at IISc Bangalore (2003) and IIT Mumbai (2004). She is a life member of the System Society of India and a member of the Indian Society for Technical Education and the Computer Society of India.

0 BASICS OF COMPUTER SYSTEMS



In this chapter, you will learn

- The brief history of Intel's microprocessors and their use in personal computers.
- The general principles of computer architecture.
- The operation of the data, address and control buses of a computer.
- The distinction between RISC and CISC computing.
- The comparison between assembly and high level language programming.
- The binary, hexadecimal and BCD number systems.
- Number format conversions.
- Computer arithmetic using different number systems.
- The processing of negative numbers as done by a computer.
- The programming model and programming of the 8085 microprocessor.
- The pin out and hardware features of 8085.

0.1 | A Brief History of Microprocessors

What is a microprocessor? We usually answer that it is a processor on a single microchip, where the word 'micro' stands for small. A processor is a device which has computing i.e., data processing capability. Over the years, the computing power of microprocessor chips has increased steeply, along with the speed of processing, and this trend is continuing as newer applications are envisaged.

The idea of a single chip computer was around even as early as the 1950s, but at that time, electronics and integrated circuit technology were still in their infancy. By the end of the 1960s, things had changed and integrated circuits had already made inroads into the technology market. Around that time, there were many companies in the field trying to develop a single chip computer and one of them was Intel. Intel released its 4-bit all-purpose chip, the Intel 4004, in November 1971. It had a clock speed of 108 KHz and 2,300 transistors with ports for ROM, RAM, and I/O. This was just the beginning of newer innovations. In April 1972, Intel came up with the 8008, which was just an 8-bit version of the 4004. Since it had nothing very spectacular to offer, this microprocessor did not do much for Intel or the microprocessor field. Meanwhile, other companies started making their presence felt, by their own versions of microchips with computing capabilities.

The first major break came with the introduction of the 8080 (again, by Intel) in 1974. Intel put itself back on the map with the 8080, which used the same instruction set as the earlier 8008, but the 8080 is generally considered to be the first truly usable microprocessor. The 8080 had a

16-bit address bus and an 8-bit data bus, a 16-bit stack pointer to memory which replaced the 8-level internal stack of the 8008, and a 16-bit program counter. It also contained 256 I/O ports, so that I/O devices could be connected without taking away or interfering with the addressing space. It also possessed a signal pin that allowed the stack to occupy a separate bank of memory. These features are what made this a truly modern microprocessor.

Other Players in the Market

All this does not mean that there were no other companies in the microprocessor field – there were, and each of them had contributions to make. We are talking more about Intel's role, because it will take us to the advent and development of Intel's x86 microprocessor family the class of microprocessors about which we discuss in this book.

Motorola (now Freescale) is another company that came up with its own original design of microprocessors. It started with the 6800 chip, which was followed by more advanced versions of this basic chip. In terms of computational capability, Motorola's processors are as good (or in some ways, better) as any Intel processor. The market segments of both these companies are separate and distinct now.

In contrast, AMD is a company that did not go for original design but worked on improving the 8080 design under license agreements. There was also this very popular microprocessor Z-80 developed by Zilog Corporation. This 8-bit microprocessor was binary compatible with the 8080 and surprisingly, is still in widespread use today in many embedded applications. It was supposed to be superior to 8080. Though Zilog made improved versions of this microprocessor, the company did not take off as well as it should have because it could not handle the competition of bigger players in the market.

In 1976, Intel updated the 8080 design with the 8085, and it became very popular as a well-designed and simple microprocessor, though it was never used in a PC. However, it is still very popular and used in various applications just as Zilog's Z-80 processor is.

In 1978, Intel introduced the 16-bit 8086, a 16-bit processor which gave rise to the x86 architecture. In 1979, Intel also released the 8088, a microprocessor similar in every way to 8086, except that its external data bus was 8 bits wide, while being a 16-bit microprocessor internally.

0.1.1 | The x86 Family

The real breakthrough for Intel came when in 1981, IBM picked Intel's 8088 for its personal computer. With this, the market segment of Intel grew by leaps and bounds. Intel kept on updating and improving the x86 architecture with newer and newer innovations and since x86 had already established its place in the PC world, PCs also became more and more sophisticated. Table 0.1 gives the timeline of the release of the various x86 processors. The points to remember is that 8088 was the first x86 processor used in a PC and it was a 16-bit architecture with an external data bus of 8 bits. Note that 8086 was manufactured earlier with the same internal architecture as the 8088, but it had an external data bus of 16 bits which made it difficult to connect it to 8-bit peripherals and was costlier too. That is why it was not chosen for the PC.

The 80186 was never used in PCs, but found applications in the field of embedded computing. The 80286, was the next x86 processor used in PCs. It was also a 16-bit architecture, but also had an external data bus of 16 bits. It marked a step upwards in PC architecture as it brought into PCs the concept of virtual memory and protected mode operations. All later PCs have these features. The next major step came with the 80386, which was a 32-bit

Table 0.1 | Year of Release of Various x86 Processors

x86 processor	Year of release
8086	1978
8088	1979
80186	1982
80286	1982
89386	1985
80486	1989
Pentium	1993
Pentium Pro	1995
Pentium-2	1997
Pentium-3	1999
Pentium-M	2002
Pentium-4	2004
Pentium-D	2005
Core-2	2006
Core-2 Quad	2007




Figure 0.1 | The first x86 processor – the 8086

processor internally and externally. All x86 processors are still 32-bit internally, but Pentium has a 64-bit external data bus.

0.1.2 | The x86 Philosophy

What is meant by the term x86? It refers to a set of machine language instructions, the philosophy and usage of which was introduced in the 8086 processor. It was designed by Intel, but other companies like AMD also used the same vocabulary, though their internal design could be different. The same was followed by Intel for processors 80186, '286, '386, '486 and the different versions of Pentium. Over the years, the x86 instruction set has grown and expanded, but backward compatibility with earlier members has been maintained. What this means is that a program written for 8086 will run on a Pentium machine (but not the other way round). The current status of the x86 processors is that it is the de facto standard for PCs from desktops

to servers, and laptops to supercomputers. There are other competitors in the market (some offering superior processors), but their market share is relatively small.

Now that we have had a peep into the history of the x86 processor family, let us also have a glimpse of the history of PCs, because that is where x86 processors have the maximum number of applications.

0.1.3 | Personal Computers

We have seen that IBM manufactured its first personal computer in August 1981, using the 8088 (see Fig 0.2). The IBM Personal Computer was introduced as the IBM 5150. During its development, the IBM 5150 had been internally referred to as 'Project Chess' and was created by a team of 12 people headed by Don Estridge and Larry Potter. The first IBM PC had an Intel 8088 processor, 64 KB of RAM (expandable to 256 KB), a floppy disk drive (which could be used to boot the computer with a rebranded version of MS-DOS (PC-DOS) and a CGA or monochrome video card. The machine also had a version of Microsoft BASIC in ROM.

In March 1983, the IBM PC-XT was marketed, with the change being that it had a hard disk as well. XT stands for 'Extended Technology'. The next change came about in August 1984 when the 80286, a 16-bit processor, which also had virtual memory capability, was used by IBM in its PC. They called it the PC-AT with AT standing for 'Advanced Technology'. Then, one after the other, 80386, 80486 and Pentium came to be used in PCs. The processor developed after the '486 was to have been named 80586, and Intel wanted to trademark this number, but a court ruling prevented this. So the word Pentium was coined and Intel did not have to look back after that. Different versions of Pentium rule the PC market as we all know.

Incidentally, it was not IBM which first used the term PC. The word 'Personal Computer' was in common usage quite some time before IBM developed its first PC, and many other



Figure 0.2 | The first IBM PC 5150

companies also made PCs using various other processors. However, when the IBM PC became a great success, many manufacturers started making similar PCs and software companies also had the inclination to develop software for IBM PC compatible PCs. That is how the PC revolution came about – and we know now, that it is not IBM that makes most of the PCs in the world now – various other manufacturers do, but the x86 processor still forms the core of all these PCs. With more and more people using PCs, standardization of software and hardware came about slowly. This has made the world of computing more ordered and comfortable, and compatibility issues for hardware and software are much, much less than in the early days.

Currently we talk about the Pentium PC, which has many versions and different clock speeds and the trend is to indicate the clock speed to specify the ‘Pentium’ model that we want for our processor. For example, the model ‘Pentium-4 2.2’ means a Pentium-4 processor with a clock speed of **2.2GHz**. Indeed, technology has come a long way since the advent of the first use of the x86 processor at **4.77MHz** frequency!

Note In this book, whenever we use the word PC, we mean the IBM compatible PC. There is the other great PC, the Apple Mac, which while being low in numbers, is an elite class of computers. ‘Our goal is not to build the most computers. It’s to build the best’ – according to the CEO of Apple Inc. Apple has a worldwide market share of just 10% of the PC market, but has more than 90% of the market share of premium computers.

What this Book is About

In this book, we will study the x86 architecture as developed by Intel. We will go through the ideas of programming in assembly language and also study the interfacing of the x86 processors with various devices. Frequently, we will also see how a particular interface is used in the standard PC. Even though we started by defining a microprocessor as a single chip computer, we know that a PC cannot be made with just a single chip. When we use the phrase ‘single chip computer’, we only mean that all the computing and processing power is held within that single chip which we call the microprocessor – but to get this chip to communicate with the external world, and make it comfortable for us to use, requires a lot more of hardware and software support. We will make a trip to understand the concepts behind how the personal computer has become an integrated part of our lives.

0.2 | Basics of Computer Architecture

0.2.1 | The Block Diagram of a Computer

A computer, as its name indicates is a machine used for computing. Computing, which many years ago meant arithmetic calculations, has now given way to large amounts of ‘data processing’. As such, it is more reasonable to designate the computer now as a ‘data processing machine’. For performing its designated tasks, this machine requires many components, which can broadly be divided as hardware and software. Hardware is obviously, the physical constituents of a computer. Software is the collection of programs which directs the hardware to perform its tasks.

Let us first look at a computer in terms of its hardware. Fig 0.3 shows the architectural description of a computer system. It shows the major parts of the computer and also indicates how these parts are connected together, to form the computing machine. The major parts are the **CPU, memory and input/output devices.**

The heart of a computer is the ‘central processing unit’. It is this unit which gives ‘life’ to a computer. The CPU usually is a ‘microprocessor’, which means that it is usually a separate and self contained chip. The CPU processes the data given to it, according to the programs meant to operate on these data. The program consists of ‘instructions’. These instructions are decoded




Figure 0.3 | The block diagram of a computer

by the CPU, which generates control signals necessary to activate the arithmetic and logic units of the CPU. As such, the CPU contains the arithmetic logic unit and the control unit. All these activities are timed and synchronized by a pulse train of fixed frequency. This is the clock signal, and it also has the job of synchronizing the activity of the CPU with the activity on the bus.

0.2.2 | The System Bus

A bus is collection of signal wires which connect between the components of the computer systems – the figure shows that the CPU is connected to the memory as well as I/O through the system bus, but only one at a time – if the memory and I/O wants to use the bus at the same time, there is a conflict, as there is only one system bus. The system bus comprises of the address bus, data bus and the control bus.

The Data Bus The set of lines used to transfer data is called the data bus. It is a bidirectional bus, as data has to be sent from the CPU to memory and I/O, and has to be received as well by the CPU. The width of the data bus determines the data transfer rate, size of the internal registers of the CPU and the processing capability of the CPU. In short, it is a reflection of the complexity of the processor. As we see, the 8086 has a data bus width of 16 bits, while the 80486 has a 32-bit bus width. Thus the 80486 can process data of 32 bits at a time while the 8086 can only handle 16 bits.

The Address Bus The address bus width determines the maximum size of the physical memory that the CPU can access. With an address bus width of 20 bits, the 8086 can address 2^{20} different locations. It can use a memory size of 2^{20} bytes or 1 MB. For Pentium with an address bus width of 32 bits, the corresponding numbers are 2^{32} bytes i.e., 4 GB. When a particular memory location is to be accessed, the corresponding address is placed on the address bus by the CPU. I/O devices also have addresses. In both cases, it is the CPU which supplies the address, and as such, the address bus is unidirectional.

The Control Bus The control bus is a set of control signals which needs to be activated for activities like writing/reading to/from memory/I/O, or special activities of the CPU like interrupts and DMA. Thus, we see signals like Memory Read, I/O Read, Memory Write and Interrupt Acknowledge as part of the control bus. These control signals dictate the actions taking place on the system bus that involve communications with devices like memory or I/O. For example, the Memory Read signal will be asserted for reading from memory. It is sent to




Figure 0.4 | The system bus and its components

memory from the processor. A signal such as 'Interrupt' is received by the processor from an I/O device. Hence in the control bus, we have signals traveling in either direction. **Some control lines may be bidirectional too.**

Now that we have discussed a computer system in general, let us go a bit deeper into its individual constituents.

0.2.3 | The Processor

The processor or the microprocessor as we might call it, is the component **responsible for controlling all the activity in the system.** It performs the following three actions continuously.

- i) **Fetch an instruction from memory.**
- ii) **Decode the instruction.**
- iii) **Execute the instruction.**

When we write a program, it is stored in memory. Our code has to be brought to the processor for the required action to be performed. The **first step** obviously, is to **'fetch' it from memory.** The **next step** i.e., **decoding**, involves the interpretation of the code as to what action is to be performed. **After decoding, the action required is performed.** This is termed **'instruction execution'.** **The sequence of these three actions is called the 'execution cycle'.** To do all this, the processor has 'control circuitry' to fetch and decode instructions. The ALU part of the processor performs the required arithmetic/logic operations. **The sequence of fetch-decode-execute is done continuously and infinitely by the processor.** An important implication of this cycle is that **instruction execution is 'sequential' in nature** – it is only after the first instruction is dealt with, will the second one be taken up. However, there will be situations when the sequential nature of program execution is disturbed. This is **when a 'branch' instruction appears in the sequence, and a new sequence of instructions will be taken up starting from a new location.**




Figure 0.5 | The execution cycle




Figure 0.6 | System clock

0.2.4 | System Clock

All the activities of the processor and buses are synchronized by a clock, which is as shown in Fig 0.6 a square wave with a particular frequency. The reciprocal of the clock frequency is the cycle time T , also called the clock period. $T = 1/f$ where f is the clock frequency. An execution cycle may require many clock periods. This depends on the architectural features of the processor, as well as the complexity of the instruction to be executed. Since an execution cycle also involves fetching instructions and data from memory, it also depends on how many clock cycles are needed to access memory. Obviously, the time for execution depends on the clock speed as well. i.e., a clock speed of 3 GHz implies faster processing than a clock of 1 GHz. However, the technology used for the processor must be able to support the clock frequency used.

0.2.5 | Memory

The memory associated with a computer system includes the primary memory as well as secondary memory. However, for the time being, we will think of memory as constituting the primary or **main memory** only, which is usually **RAM** (Random Access Memory). **Memory is organized as bytes**, and the capacity of a memory chip is stated in terms of the number of bytes it can store. Thus, we can have chips of size 256 bytes, 1KB, 1MB and so on. If a computer has a total memory space of 20 MB it can use RAM chips of the available capacity to get that much of memory.

There are **two basic operations** associated with memory – **read and write**. **Reading causes a data stored in a memory location to be transferred to the CPU**, without erasing the content in memory. **Writing causes a new data to be placed in a memory location** (it overwrites the previous value). There is a certain amount of time required for these operations and this is termed as ‘access time’.




Figure 0.7 | Memory and associated control signals

Memory Read Cycle The steps involved in a typical read cycle are:

- i) Place on the address bus, the address of the memory location whose content is to be read. This action is performed by the processor.
- ii) Assert the **memory read** signal which is part of the control bus.
- iii) Wait until the content of the addressed location appears on the data bus.
- iv) Transfer the data on the data bus to the processor.
- v) De-activate the memory read signal. The memory read operation is over and the address on the address bus is not relevant anymore.

Memory Write Cycle As a continuation, let us also examine the steps in a typical write cycle.

- i) Place on the address bus, the address of the location to which data is to be written.
- ii) On the data bus, place the data to be written.
- iii) Assert the **memory write** signal which is part of the control bus.
- iv) Wait until the data is stored in the addressed location.
- v) De-activate the memory write signal. This ends the memory write operation.

At this stage, we should remember that these operations are synchronized with the system clock. An **8086 processor takes at least four clock cycles for reading/writing**. These four cycles constitute the 'memory read' and 'memory write' cycles for the processor. Other processors may require more/less clock cycles for the same operations.

0.2.6 | The I/O System

For a computer to communicate with the outside world there is the need for what are called peripherals. Some of these peripherals are purely input devices like the keyboard and mouse; some are purely output devices like the printer and video monitor and some like the modem transfer data in both directions. All this just means that such I/O devices are needed for us to use a computer. However, **it is difficult for a processor to deal directly with I/O devices, because of their incompatibility with the processor** – each peripheral is different and the operating conditions, voltages, speeds and standards are not understandable to the processor. **The processor does not have the necessary control signals to deal with different peripherals**. Hence, the normal practice is **for each peripheral to have a controller which acts as an interface between the peripheral and the processor**. This controller, which may be a special purpose chip, understands the characteristics of the particular device and provides the necessary control signals to the processor to communicate with the peripheral. Thus, we have specialized controllers for most




Figure 0.8 | The I/O system

peripherals – like the keyboard display interfacing chip, parallel port interfacing chip and serial communication chip. All these chips are programmable – they have registers for commands, data and status. By suitably programming these chips, we can get the processor to communicate correctly with any peripheral. Fig 0.8 shows the use of an I/O interface between an I/O device and a processor. The processor is not shown in the figure, but the system bus which comes from the processor is shown.

In the final analysis, we can think of a computer that we usually use, as a conglomeration of components which include memory and I/O devices of various types, applications and specifications.

0.3 | Computer Languages

0.3.1 | Machine Language, Assembly Language and High Level Language

The computer is just a dumb piece of equipment unless we are able to make it work for us. For that, we must be able to 'program' it, so that it will perform the tasks we assign it. Programming a computer entails the use of a language that the computer understands. The language native to computers is 'machine language' which consists of binary ones and zeros. The computer knows this language, and the series of ones and zeros fed to it are 'operation codes' for it, which tells it what action is to be performed. Thus there is one binary code for addition and another one for subtraction. These operation codes are called 'opcodes' and this language is called 'machine language'. Programming in machine language means writing the opcodes of the tasks we want to get done by the computer.

However, the problem with machine language, as is obvious, is that it is cumbersome and error-prone. Human beings are not good at remembering or using binary codes. Programming using machine language is not something that any one of us is likely to enjoy. To make it easier for us to communicate with computers, there is a language at a slightly higher level and that is called 'assembly language'. This is more intelligible to users than machine code. This language uses 'mnemonics' for specifying the operation the computer is to perform. These mnemonics are a direct translation of the machine code to a symbol. For example, the binary code for addition is replaced by the symbol 'ADD' – the binary code for multiplication is given the symbolic name 'MUL'. The exact mnemonic used depends on the processor type, but it will be related to the operation to be done.

How does this help? A user does not have to remember binary codes or enter binary code for programming. He only needs to remember the symbolic codes and the associated syntax. We say that assembly coding is at a higher level than machine coding. However, does the computer understand the mnemonics? No, which means that should be an interface between assembly language and machine language. This interface converts the symbolic codes fed in by the user into machine codes. The software which does this is called an ‘**assembler**’.

Since machine language is native to a processor, each processor will have its own machine language and thus it has its own assembly language also. Translating from assembly language to machine language and vice versa is a one to one process – one opcode translates to a unique machine code for a particular processor.

However, we human beings always look for easier ways to get things done. So there are ‘higher level languages’ which has the vocabulary and grammar similar to the language spoken by us. Such languages are very easy to use because the communication process is similar to English. We have heard of languages like C, FORTRAN, COBOL and many, many such ‘high level languages’. The features of such languages are that they are

- i) easy to understand and write,
- ii) are not processor specific.

Thus if we write a program in C, we can use it to run on any processor – as long as the ‘compiler’ for the language is available. The compiler is the software which ‘translates the high level language statements’ to statements in a lower level language. The lower level may be assembly or machine language. However, finally the processor needs the machine code.

The program that we write in assembly or high level language is called the source program or source code. A compiler or assembler converts this into an object code which is ‘executable’ in the sense that the processor understands the code and performs the tasks indicated.

0.3.2 | Comparison

Programming in machine language is too cumbersome and hence ruled out in the present world. However, assembly language programming is frequently done, so let us now make a comparison between assembly language programming and high level language programming.

Assembly code is specific to a processor – which means that the assembly code of 8086 does not make any sense to 8085 (though both are Intel made). Assembly programs need the programmer to know the architecture of the processor intimately. He should know the registers and flags and the way each instruction handles data. So doing assembly coding involves the study of the concerned processor. However, once this part is done, coding is very efficient, compact and executes very fast. Speed advantage of a hundred times or more, is fairly common. Assembly language programming also gives direct access to key machine features essential for implementing certain kinds of low level routines, such as an operating system kernel or microkernel, device drivers, and machine control.

High level language programming, on the other hand, is not processor-specific. It is easy to learn and master. However, high level languages are abstract. Typically a single high level instruction is translated into several (sometimes dozens or in rare cases even hundreds) executable machine language instructions. The object code generated by a compiler is usually not compact. However, the advantage of high level languages is that since it is easy to learn, semi-skilled designers can be employed for development activities, and so development and maintenance times are much less.

This book focuses on the x86 architecture and on assembly language programming. The aim is to impart good assembly language skills and a thorough knowledge of the x86 architecture.

0.4 | RISC and CISC Architectures

Two terms that are likely to be encountered frequently while reading about computer architecture are RISC and CISC. RISC stands for **Reduced Instruction Set Computer** and CISC means **Complex Instruction Set Computer**. Since a lot of controversy surrounds these two terms, let us try to find out what it is all about.

In the early days of microprocessor development, the trend was to **have complex instructions implemented fully using hardware**. For example, the multiply instruction is a complex instruction which needs a dedicated hardware multiplier. Because hardware is fast, execution is fast, but with lots of such complex instructions, the hardware budget is naturally high. This is the philosophy and the main feature of **CISC**.

RISC on the other hand, views this matter in a different way. On an average, **the number of complex instructions a computer uses is relatively less**. So, why not **realize a complex instruction using a set of simple instructions**? This is possible, and the advantage is that the hardware budget is much less. The **instruction set is also small**. However, software is to be written to realize complex instructions with simple instructions. This amounts to trading software for hardware.

There exists a long history of controversy regarding which is better. The x86 architecture was based on the CISC philosophy, right from the beginning. By the time RISC principles became popular and software development for RISC became established, the x86 CISC processors had already carved a niche for themselves in the processor market. So, even though the supporters of RISC were able to establish their point, most developers did not want to take the risk of switching over to an untested domain. However, most of the newer processors used the RISC philosophy for their architectures – examples are ARM, Power PC, Sun's Sparc processors and the like. Many of them found their applications in the embedded processing field.

The main features of RISC are that they have only simple instructions implemented in a single clock. However, there is an irony in that, many RISC processors have as many complex instructions as CISC processors. Probably this can be justified by explaining that such complex instructions have been implemented using microprogramming rather than a direct hardware realization. Microprogramming is a method of implementing the control unit of a computer by breaking down instructions into a sequence of small programming steps.

While the RISC versus CISC controversy is still raging, the distinction between what exactly is RISC and what is CISC is reducing to the extent of being almost indistinguishable except at the basic philosophical level. Intel which held on to CISC for many years, bowed down to the RISC architecture by designing its Pentium Pro with complex instructions which internally were broken down to simple RISC like instructions. So the comment on it is that Pentium Pro is a RISC processor than runs CISC instructions.

0.5 | Number Systems

Motivation In the study of microprocessors, we will have to use many different number systems, and conversions from one system to the other. Clarity of these ideas is very important for correct computation and the right interpretation of results. This is the motivation for a review on it, though most of you have had an introduction to it already.

We have become quite used to the number system which we call the decimal number system, which is a system with a base (radix) 10. We are so used to this system of numbers that our visualization of quantity is always based on this. Our mental faculties are tuned to perform all calculations in this number system. In contrast, computers are not comfortable with this

system – we know that they use the binary system of numbers and all computations are done in the binary format. Thus we have a problem when we use computers to perform computations for us. So let us start this discussion by first understanding the intricacies of each of the commonly used number systems. We will discuss the ones that we most often might have to use in the context of computers.

0.5.1 | The Decimal System

The base of this system is 10 (ten) – and it naturally follows that there are ten defined symbols in this system – the combinations of these ten symbols give us various values. The ten symbols here are 0, 1, 2, 3, 4, 5, 6, 7, 8 and 9, and they are called ‘digits’. The **position** of a digit in a number is what gives its value.

For example, how does the number 346 get its value?

$$\begin{aligned} 346 &= 3 \times 10^2 + 4 \times 10^1 + 6 \times 10^0 \\ &= 300 + 40 + 6 \end{aligned}$$

This means that, associated with each position, there is a weight. Here the weight is a power of 10. Thus

$$\begin{aligned} 56785 &= 5 \times 10^4 + 6 \times 10^3 + 7 \times 10^2 + 8 \times 10^1 + 5 \times 10^0 \\ &= 50000 + 6000 + 700 + 80 + 5 \end{aligned}$$

What about fractional numbers? The positions on the right side of the decimal number also have weights, but the powers (of 10) are negative.

$$\begin{aligned} 6.785 &= 6 \times 10^0 + 7 \times 10^{-1} + 8 \times 10^{-2} + 5 \times 10^{-3} \\ &= 6 \times 1 + 7 \times 0.1 + 8 \times 0.01 + 5 \times 0.001 \\ &= 6 + 0.7 + 0.08 + 0.005 \end{aligned}$$

These are things that we know very well. They have been reviewed here to show that the same concept applies to other number systems as well.

0.5.2 | The Binary Number System

The base of this system is 2, and so it has two symbols, 0 and 1, each of them being called a bit. So each position has a weight which is a power of 2. Take the number 110110. Let us find its value. Since there are 6 bits, there are six positions with weights as shown for the bits:

Power of 2:	2 ⁵	+2 ⁴	+2 ³	+2 ²	+2 ¹	+2 ⁰
Weight:	32	+16	+8	+4	+2	+1
Number:	1	1	0	1	1	0
Value:	1 × 32	+1 × 16	+0 × 8	+1 × 4	+1 × 2	+0 × 1

Adding the values in all the bit positions gives 32 + 16 + 0 + 4 + 2 + 0 = 54. This is the equivalent value in the decimal system. We cannot help putting back everything into the decimal system, because this is the number system with which we are most familiar and comfortable.

Note Sometimes binary numbers are suffixed with B to indicate that they are binary numbers e.g., 110110B, 1010110B. Sometimes the notation 110110₂ is also used.

Note Keep the calculator in the PC (Accessories of Windows) open in the scientific mode. This will help to verify all the calculations we are going to do from now on.

Next, let us try to understand the concept of fractional binary numbers.

Example 0.1

Find the decimal values of the binary number 1001.011 B

Solution

Power of 2:	2^3	2^2	2^1	2^0	2^{-1}	2^{-2}	2^{-3}		
Weight	8	4	2	1	0.5	0.25	0.125		
Number	1	0	0	1	0	1	1		
Value	8	+	0	+	0	+	0.25	+	0.125
	= 9.375								

Example 0.1 shows 1001.011 in binary (often written as 1001.011₂).

It also shows the power and weight or value of each digit position. Thus 1001.001 is equivalent in decimal to 9.375 ($8 + 1 + 0.25 + 0.125$).

Notice that this is the sum of $2^3 + 2^0 + 2^{-2} + 2^{-3}$, but 2^2 and 2^1 are not added, as the bit under these positions is 0. The fractional part is composed of 2^{-2} and 2^{-3} , but there is no digit under 2^{-1} , so 0.5 is not added.

0.5.3 | The Hexadecimal Number System

Next is the hexadecimal system of numbering which has 16 symbols namely 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, A, B, C, D, E, F. The base of the system is 16 and each symbol is called a 'hex digit'. Each position in the hexadecimal number system has a weight which is a power of 16. Let us find the value of 240FCH. The letter 'H' is suffixed to the number if it is needed to make clear that it is a hexadecimal number. A to F have the decimal values of 10 to 15.

Power of 16:	16^4	$+16^3$	$+16^2$	$+16^1$	$+16^0$
Weight:	65536	+4096	+256	+16	+1
Number:	2	4	0	F	C
Value:	2×65536	$+4 \times 4096$	$+0 \times 256$	$+15 \times 16$	$+12 \times 1$
i.e.,	131072	+16384	+0	+240	+12
	= 147708				

So we have calculated the equivalent decimal value of the given hex number by using the concept of positional weights.

Example 0.2

Find the decimal value of the hex number 25.1H

Solution

Power of 16:	16^1	$+16^0$	$+16^{-1}$
Weight:	16	+1	+0.0625
Number:	2	5	1
Value	2×16	$+5 \times 1$	$+1 \times 0.0625$
i.e.,	32	+5	+0.0625
	= 37.0625		

There is also an octal system whose base is 8. The equivalent calculations involved in this are left as an exercise for the interested student. In all the above, we have done the conversion to decimal form from other number systems. Now we will see how we will convert a decimal number to other systems of numbering.

- Note** i) In most computers, the default number system for writing numbers is decimal. When we mean decimal numbers, we simply write it as it is – like 35, 687 and 234 and so on. A number in hex form is suffixed with the letter H, for example, 56H, 8FH, 0AH and so on.
- ii) The numbers from 0 to 9 are the same in the decimal and the hexadecimal system. So, in the forthcoming chapters, you will see that no ‘H’ is added when writing numbers from 0 to 9, though there is nothing wrong in writing 7H, 8H, 01H and so on.

0.6 | Number Format Conversions

0.6.1 | Conversion from Decimal to Binary

The method is to divide the decimal number by 2, until the quotient is 0. See the technique illustrated below.

Example 0.3

Find the binary value of 13.

Solution

Divide 13 by 2 repeatedly and save the remainders

2)13	remainder = 1
2)6	remainder = 0
2)3	remainder = 1
2)1	remainder = 1
0	

Now write the remainders from bottom to top, as one line from left to right. We get 1101 as the converted binary number.

Thus, we have been able to convert from decimal to binary by repeated division by 2, the base of the binary number system. To verify, try converting this binary number back to decimal. It should be 13. Or simply use the scientific calculator to verify the conversion. (Make sure it is kept open on your PC’s desktop.)

Example 0.4

Convert the number 213 to binary form.

Solution

2)213	remainder = 1
2)106	remainder = 0
2)53	remainder = 1
2)26	remainder = 0
2)13	remainder = 1
2)6	remainder = 0
2)3	remainder = 1
2)1	remainder = 1
0	

Now write the remainders from bottom to top in one line, from left to right. The number is 11010101.

0.6.2 | Conversion from Decimal to Hexadecimal

Conversion from decimal to hexadecimal is accomplished by dividing by 16 and finding the remainders. Remainders ranging from 10 to 15 will be written using the hexadecimal symbols A to F. See how 225 is converted to a hexadecimal form.

16)225

16)14

0

remainder = 1

remainder = E

Result = E1

The method for this is obvious i.e., divide repeatedly the decimal number by 16, keep the remainders. Do this until the quotient is 0.

Example 0.5

Convert the decimal number 4152 to hexadecimal.

Solution

16)4152

16)259

16)16

16)1

0

remainder = 8

remainder = 3

remainder = 0

remainder = 1

Take the remainders from bottom to top and write it in a single line from left to right. The number is 1038H.

0.6.3 | Converting from Binary to Hexadecimal

If we take any hex digit, note that its decimal value ranges from 0 to 15. For example F is 15, A is 10 and so on. If a hex digit has to be converted to a binary number, the maximum number of bits required is 4.

See Table 0.2. Any hex digit can be written as a group of four bits. Taking an example, 4C57FH can be written in binary, by writing the equivalent binary of each of the digits

Hex	4.	C	5	7	F
Binary	0100	1100	0101	0111	1111

Table 0.2 | Hex, Binary and Decimal Representations

Decimal	Hex	Binary	Decimal	Hex	Binary
0	0	0000	8	8	1000
1	1	0001	9	9	1001
2	2	0010	10	A	1010
3	3	0011	11	B	1011
4	4	0100	12	C	1100
5	5	0101	13	D	1101
6	6	0110	14	E	1110
7	7	0111	15	F	1111

The binary value of 4C57FH is 01001100010101111111. Looking at both the representations tells us the biggest problem with binary numbers – they are long and cumbersome to handle. Putting them into a hex form makes the representation short and concise – we conclude that the binary representation is an expanded form of the hexadecimal representation where each hex digit is expanded to its 4-bit binary form.

If we have a long binary number, what we can do to convert it into hex form is to divide it into groups of 4 bits (starting from the right i.e., the LSB). Then write the hex representation of each 4-bit binary group. Try this technique with the following binary number: 11100101010100011101.

1110	0101	0101	0001	1101	B	i.e., binary
E	5	5	1	D	H	i.e., hex

Example 0.6

Convert the following binary number to hex form.
10111110000111110001.

Solution

It is the practice to write binary as groups of four with a space between the groups. This increases the readability of the binary number.

1011	1110	0001	1111	0001
A	E	1	F	1

The equivalent hex number is AE1F1H.

Example 0.7

Convert the following hexadecimal number to binary form.
3AF24H

Solution

Take each hexadecimal digit and write its equivalent four-bit binary value.

3	A	F	2	4
0011	1010	1111	0010	0100

From all this, we should realize that the hexadecimal notation is a contracted form of rebinary number representation. Computers do all their processing using binary numbers, but it is easier for us to represent that binary number in hex form. So when we ask the computer to add 34H and 5DH, it actually expands these into binary form and does the addition.

0.6.4 | BCD Numbers

BCD stands for ‘Binary Coded Decimal’ but there is more to it than being just a binary representation of a decimal number. Let us look into the details.

Decimal numbers are represented by 10 symbols from 0 to 9, each of them being called a digit. We know the binary code for each of these decimal numbers. Suppose we represent one

decimal digit as a byte, it is called ‘unpacked BCD’. Consider the representation of 9 – it is written as 00001001. Now if we want to write 98 in unpacked BCD, it is written as two bytes:

9 8
00001001 00001000.

Thus the binary code of each decimal digit is in one byte.

Packed BCD What, then, is ‘packed BCD’? When each digit is packed into 4 binary bits, it is packed BCD. Thus 98 is

9 8
1001 1000.

Each digit needs a nibble (four bits) to represent it. The packed BCD form of 675 is 0110 0111 0101. The important point to remember is that since there is no digit greater than 9, no BCD nibble can have a code greater than ‘1001’. Computers do process BCD numbers, but the user must be aware of the number representation that is being used.

Can we write BCD numbers in hex? Yes, because the hex representation is just a concise representation of binary numbers. The decimal number 675 when written as 675H represents the packed BCD, in hex form. There is no need to be confused about this, because the steps involved are:

- i) write the binary equivalent of each decimal number, as a nibble,
- ii) write the hex equivalent of each nibble.

675 is 0110 0111 0101
 6 7 5 H

Spend a few moments thinking, to make it clear. One important point to keep in mind is that when we represent BCD in hex form, no digit will ever take the value of A to F, since decimal digits are limited to 9. So there will never be a BCD number such as 8F5H or 56DH or A34H.

Example 0.8

Find the binary, hex and packed BCD representation of the decimal numbers 126 and 245. Also write the packed BCD in the hex format.

Solution

Number	Binary	Hex	BCD	BCD in hex form
126	0111 1110	7EH	0001 0010 0110	126H
245	1111 0101	F5H	0010 0100 0101	245H

Example 0.9

Find the packed BCD value of the decimal number 2347654, and represent the BCD in hex format.

Solution

To find the BCD, each digit is to be coded in 4-bit binary.

Hence 2347659 is

0010 0011 0100 0111 0110 0101 1001 i.e., 2347659H is the hex representation of the BCD number.

It is very important to keep this in mind, when we do programs using BCD arithmetic. Whenever you have doubts then, just refer back to this chapter.

0.6.5 | ASCII Code

This word pronounced as 'ask-ee' is the abbreviation of the words 'American Standard Code for Information Interchange'. This is the code used when entering data through the keyboard and displaying text on the video display. It is very important to know what it is and how this code is used.

ASCII is a **seven bit code, which is written as a byte**. It has representations for numbers, lower case and upper case English alphabets, special characters (like #, ^, &) and control characters. For example, there are ASCII codes for 'new line', carriage return and the space bar. A number of characters are related to printing. **When we type a character on the keyboard, it is the ASCII value of the key that is read in**. The computer must convert it from this form to binary form, for processing. The list of ASCII codes is shown in Table 0.3. Note that the ASCII value of numbers from 0 to 9 is 30H to 39H. The ASCII of upper case alphabets starts from 41H and that for lower case starts from 61H. This table will be needed as quick reference for various calculations we will do in the programming chapters.

0.6.6 | Representation of Negative Numbers

There are various ways of representing negative numbers – like signed magnitude, one's complement, two's complement and so on, but we will straightway discuss the representation used by computers for this. Computers use the 'two's complement' representation for negative numbers. The method is to complement each bit of the number and add a '1' to this. Let us see how it is done.

We will start with 4-bit numbers. Say we want to represent -6.

- i) Write the 4-bit binary value of 6: 0110.
- ii) Complement each bit: 1001.
- iii) Add '1' to this: 1010.

So -6 is '1010', for computers.

Let us try this for all the numbers from 0 to 7. See Table 0.4 which shows the positive and negative number representation of numbers possible to be represented in four bits. A number of observations can be made from Table 0.4.

- i) The range of numbers that can be represented by 4 bits is -8 to +7. For an n -bit number, this range works out to be (-2^{n-1}) to $(+2^{n-1}-1)$.
- ii) In this notation, the most significant bit (MSB) is considered to be the sign bit. The MSB for positive numbers is '0' and for negative numbers is 1.
- iii) There is a unique representation for 0.

Since we will deal mostly with bytes and words (16-bit) let's have a feel of 8-bit negative number representation.

Table 0.3 | The ASCII Code – Symbols versus Hex Value

Symbol	ASCII (Hex)	Symbol	ASCII (Hex)	Symbol	ASCII (Hex)	Symbol	ASCII (Hex)
NUL	0	DLE	10	(Space)	20	0	30
SOH	1	DC1	11	!	21	1	31
STX	2	DC2	12	"	22	2	32
ETX	3	DC3	13	#	23	3	33
EOT	4	DC4	14	\$	24	4	34
ENQ	5	NAK	15	%	25	5	35
ACK	6	SYN	16	&	26	6	36
BEL	7	ETB	17	.	27	7	37
BS	8	CAN	18	(	28	8	38
Tab	9	EM	19	)	29	9	39
LF	A	SUB	1A	*	2A	:	3A
VT	B	ESC	1B	+	2B	;	3B
FF	C	FS	1C	,	2C	<	3C
CR	D	GS	1D	-	2D	=	3D
SO	E	RS	1E	.	2E	>	3E
SI	F	US	1F	/	2F	?	3F
@	40	P	50	`	60	P	70
A	41	Q	51	a	61	q	71
B	42	R	52	b	62	r	72
C	43	S	53	c	63	s	73
D	44	T	54	d	64	t	74
E	45	U	55	e	65	u	75
F	46	V	56	f	66	v	76
G	47	W	57	g	67	w	77
H	48	X	58	h	68	x	78
I	49	Y	59	i	69	y	79
J	4A	Z	5A	j	6A	z	7A
K	4B	[	5B	k	6B	{	7B
L	4C	/	5C	l	6C		7C
M	4D	]	5D	m	6D	}	7D
N	4E	^	5E	n	6E	~	7E
O	4F	-	5F	o	6F		7F

Table 0.4 | Negative and Positive Number Representation in 4-bit Binary

Negative Numbers	Binary	Hex	Positive numbers	Binary	Hex
−8	1000	8			
−7	1001	9	+ 7	0111	7
−6	1010	A	+ 6	0110	6
−5	1011	B	+ 5	0101	5
−4	1100	C	+ 4	0100	4
−3	1101	D	+ 3	0011	3
−2	1110	E	+ 2	0010	2
−1	1111	F	+ 1	0001	1
−0	0000	0	+ 0	0000	0

Example 0.10

Find the two’s complement number corresponding to −6 when 6 is represented in 8 bits as 0000 0110.

Solution

The steps: 0000 0110

1111 1001;complement each bit

1111 1010;add ‘1’ to it

F A;in hex

Thus −6 is FAH in 8-bit form, while it is AH in 4-bit form (from Table 0.4)

Note H is the notation for ‘hexadecimal’.

One very important point we need to observe and keep in mind is that, when a 4-bit number is expanded into an 8-bit form, its sign bit has to be extended into the 8 bits. The sign bit in the 4-bit representation of −6 is ‘1’. When expanding the number to fill into 8 bits, the 1 is replicated 4 more times to fill the whole byte. Thus −6 which is AH in 4-bit form, becomes FAH in byte form, and will be FFFAH in 16-bit format, and FFFF FFFAH in 32-bit format. We need to understand this for negative numbers. For positive numbers, we do it without much thinking. So +6 is 0110, which expands to be 0000 0110 (byte) or 06H, and 0006H in 16-bit format and 0000 0006 H in the 32-bit format. Note that for positive numbers, the sign bit is 0; effectively we are doing sign extension here too. This concept of ‘sign extension’ is important and we will deal with it in greater detail later.

Conversion from Two’s Complement Form Given the two’s complement representation of a decimal number, how do we find the decimal number which it represents? The answer is – two’s complement it again.

Take FA

1111 1010...the number

0000 0101 +...invert each bit

1...add 1

0000 0110...its 2’s complement

This is 6. Thus FA is the two’s complement representation of −6.

Example 0.11

Find the decimal number whose two's complement representation is given.

- i) FFF2H
- ii) F9H

Solution

- i) FFF2H

Taking two's complement gives 000E
which is 1110. i.e., 14 – which means that –14 is the number represented by FFF2H.

- ii) F9H

Taking two's complement gives
0111 i.e., 7, which means that –7 is the number represented by F9H.

Question Looking at the result of various arithmetic operations on binary numbers, how do we know whether it is a positive or a negative number? What is your observation regarding signed numbers?

Answer We should know how many bits are used for the representation of a signed number in the system. Then, if the MSB is a '1', it is a negative number, if the MSB is a '0', it is a positive number.

0.7 | Computer Arithmetic

0.7.1 | Addition of Unsigned Numbers

When we say that a number is unsigned, it implies that the sign of the number is irrelevant, which actually means that we consider the numbers as having no sign bit – all the bits allotted for the data are used for the magnitude alone, in effect, it turns out that these refer to positive numbers. With 8 bits, numbers from 0 to 255 can be used.

Binary addition is something that you have already learnt. Here we are reviewing it to bring into focus some important points which we may have to be taken care of, in the study of micro-processor programming.

Binary addition is done by adding bits column wise. We will consider byte sized data.

Case 1

Binary		Decimal		Hexadecimal
0101 1001	+	89	+	59H
0110 1001		105		69H
<u>1100 0010</u>		<u>194</u>		<u>C2H</u>

Addition of the same numbers in the binary, decimal and hexadecimal formats is shown. Since the sum lies within a value of 255, there is no special problem in this case.

Case 2

0111 1000		120	+	78H	+
<u>1001 1001</u>		<u>153</u>		<u>99H</u>	
<u>10001 0001</u>		<u>273</u>		<u>111H</u>	

In this case, the sum is greater than the number of bits allotted for the operand, and the extra bit, beyond the 8 bits of the sum, is called a 'carry'. **Whenever a carry appears, it indicates the insufficiency of the space allocated for the result. In microprocessors, there is a flag that indicates this condition.**

0.7.2 | Addition of Packed BCD Numbers

Now let us add packed BCD numbers

Case 1

Consider the case of two packed BCD bytes that are to be added, say 45 and 22.

Packed BCD		Packed BCD in hex form		Decimal
0100 0101	+	45H	+	45
<u>0010 0010</u>		<u>22H</u>		<u>22</u>
<u>0110 0111</u>		<u>67H</u>		<u>67</u>

In this case, the upper nibble and lower nibble are within 0 to 9. So the addition proceeds just like normal decimal addition.

Case 2

Consider the case of two packed BCD bytes that are to be added, say 45 and 27. In BCD form, the correct answer should be 72. However, this is not obtained directly.

Packed BCD		Packed BCD in hex form		Decimal
0100 0101	+	45H	+	45
<u>0010 0111</u>		<u>27H</u>		<u>27</u>
<u>0110 1100</u>		<u>6CH</u>		<u>72</u>

When adding in binary form, the lower nibble of the sum is greater than 9. Since no BCD digit can have a value greater than 9, a correction needs to be applied here. The correction to get the sum back to BCD form is to add 6 (0110) to the lower nibble alone.

Correction

0110 1100	+
<u>0000 0110</u>	
<u>0111 0010</u>	

This gives the correct sum of 72.

Case 3

This is when the upper nibble of the sum is greater than 9. The correction is to add 6 to the upper nibble alone.

Add BCD 76 and 62. In binary form, the additions are

0111 0110	+	76H	+	
<u>0110 0010</u>		<u>62H</u>		
<u>1101 1000</u>	+	<u>D8H</u>	+	Now adding 6 to the upper nibble,
<u>0110 0000</u>		<u>60H</u>		
<u>1 0011 1000</u>		<u>138H</u>		

However, note that the data size exceeds 99, which is the maximum number that 8 bits can accommodate for a packed BCD number. Thus there is a 'carry' generated from the addition operation. However, if the carry is also included in the answer, the sum of 138 is correct. However, more than 8 bits are needed for the sum.

Case 4

When both the upper and lower nibbles of the sum are greater than 9, add 6 to both nibbles. Add BCD 89 and 72.

1000 1001	+	89H	+	
<u>0111 0010</u>		<u>72H</u>		
1111 1011		FBH		add 06 to both nibbles
<u>0110 0110</u>		<u>66H</u>		
1 0110 0001		1 61H		

The right answer of 161 is obtained. However, the sum needs more than one byte space.

Example 0.12

Perform the addition of the following numbers, after converting to decimal and hexadecimal forms.

- i) 39 and 99
- ii) 117 and 156

Solution

Decimal	+	Binary	+	Hexadecimal	+
i) 39	+	0010 0111	+	27H	+
<u>99</u>		<u>0110 0011</u>		<u>63H</u>	
138		1000 1010		8AH	
ii) 117	+	0111 0101	+	75H	+
<u>156</u>		<u>1001 1100</u>		<u>9CH</u>	
273		1 0001 0001		1 1 1H	

In the second addition, the data has exceeded the size which can be accommodated in 8 bits. Hence a carry will be generated. In microprocessors, there is a flag which indicates this condition.

0.7.3 | Addition of Negative Numbers

We know now that negative numbers are represented in two's complement notation. Let's consider adding two negative numbers.

Example 0.13

Add -43 and -56

Solution

Convert the two numbers into their two's complement form, as both are negative numbers.

-43	+	1101 0101	+
<u>-56</u>		<u>1100 1000</u>	
-99		1 1001 1101	

We are adding two 8-bit numbers. If the sum exceeds 8 bits, an extra bit is generated from the addition. Ignore this carry and look at the eight bits of the sum. (This is the rule for two's complement addition.)

It is 1001 1101. The MSB is found to be '1'. So we know that it is a negative number. To find the decimal number whose two's complement representation this is, take the two's complement of the sum. This comes to be 0110 0011 i.e., 99. Thus, we verify the correctness of our addition procedure.

Example 0.14

Add +90 and -26.

Solution

One number is positive and the other is negative.

$$\begin{array}{r}
 +90 \qquad 0101 \ 1010 \quad + \\
 -26 \qquad 1110 \ 0110 \\
 \hline
 64 \qquad 1 \ 0100 \ 0000
 \end{array}$$

Ignore the end around carry. The sum is 0100 0000. Since the MSB of the number is '0', we understand that the sum is positive. So convert it to decimal. The result is 64.

Example 0.15

Add -120 and +45

Solution

$$\begin{array}{r}
 -120 \qquad 1000 \ 1000 \quad + \\
 +45 \qquad 0010 \ 1101 \\
 \hline
 -75 \qquad 1011 \ 0101
 \end{array}$$

Look at the sum – the MSB of the sum is '1'. Hence, it is a negative number. The two's complement of this is 0100 1011 i.e., 75. Thus, the result of the calculation is -75.

Note In all the above calculations, we have used data of 8 bits. The result of the calculations was in the range of -128 to +127. Thus, the answers are correct. If the sum goes outside this (for eight-bit data), the answers will be wrong, and havoc will be created if one is not aware of that. Computers have 'flags' to let us know of this. This will be discussed in later sections.

0.7.4 | Subtraction

Unsigned Numbers

- i) Binary numbers
- ii) Hexadecimal numbers
- iii) BCD numbers

The procedure here is similar to addition i.e., bit by bit, column by column subtraction. Sometimes, borrows from the columns on the left are needed.

Example 0.16

Subtract 56 from 230. Do this subtraction after converting numbers to binary and hex.

Solution

230	–	1110 0110	–	E6H	–
<u>56</u>		<u>0011 1000</u>		<u>38H</u>	
174		1010 1110		AEH	

In the above subtraction, we are subtracting a smaller number from a bigger number. However, when subtracting column-wise, sometimes there is the issue of having to subtract a bigger number from a smaller number. We know the idea of ‘borrow’ from the left-hand column. However, for the borrowing with which we append the number, depends on the base of the number system. For the decimal system, we borrow 10, for binary 2 and for hex we borrow 16.

Check this hexadecimal subtraction:

E6H	–
<u>38H</u>	
AEH	

Starting from the rightmost column, we see that we cannot subtract 8 from 6. So, borrowing from E is needed. Borrowing from E leaves E to become D and 6 becomes 6 + 16 = 22 (in decimal). Subtracting 8 from 22 gives 14 which is E in hex. That is how we get E in rightmost column of the result. Then, going over to the left, subtract 3 from D (13 in decimal). This is 10 (in decimal) and A in hex. That is how the result of the subtraction is A.

This idea has been explained here in detail, so that we can use a similar idea in BCD subtraction.

0.7.5 | Packed BCD Subtraction

Let us use the same numbers for BCD subtraction as we did in Example 0.16. i.e., subtract 56 from 230. The BCD representation is shown below. Each decimal digit is packed in to 4-bit binary bits.

Decimal		Packed BCD	
230	–	0010 0011 0000	–
<u>56</u>		<u>0000 0101 0110</u>	
174		0001 0111 0100	

The point to remember here is that each group of 4 bits represents a ‘decimal number’, the base of which is ten. Thus, when we try to subtract a bigger number from a smaller number, we have to consider the ‘four bits together’ as a decimal number. Let us review the steps in the above subtraction.

First step

Thus, when we have to subtract 6 from 0 in the rightmost group of four bits, we need to borrow. Borrow from the group on the left a decimal 10, and add it to the ‘0000’ on the right. That makes it ‘1010’ (because of borrowing, the 0011 on the left is now ‘0010’). Then subtract 0110 from this. The result is 0100 as seen (within the group, binary subtraction is done).

0010 0010 1010	–
<u>0110</u>	
0100	

Second step

This is the second group. For subtracting 0101 from 0010, borrowing of decimal 10 is taken from the leftmost group. Thus 0010 is '1100', 12 in decimal. Subtracting '0101' (5) from it, gives '0111' (7) as shown.

$$\begin{array}{r} 0001\ 1100\ 1010\ - \\ \underline{0101\ 0110} \\ 0111\ 0100 \end{array}$$

Third step

The leftmost group is now 0001. Subtract 0000 from it. Thus, the final answer is 174 in packed BCD form.

$$\begin{array}{r} 0001\ 1101\ 1010\ - \\ \underline{0000\ 0101\ 0110} \\ 0001\ 0111\ 0100 \end{array}$$

All this shows that BCD subtraction also needs extra care as BCD addition. In computers, special instructions take care of this.

Example 0.17

Express the numbers 53 and 18 in packed BCD and subtract the latter from the former.

Solution

Decimal	Packed BCD
53 -	0101 0011 -
18	0001 1000

First step

Borrowing from the left side nibble to the nibble on the right side gives

$$\begin{array}{r} 0100\ 1101\ - \\ \underline{0001\ 1000} \\ 0101 \end{array}$$

Second step

$$\begin{array}{r} 0100\ 1101\ - \\ \underline{0001\ 1000} \\ 0011\ 0101 \end{array}$$

The result is 35, as it should be.

0.7.6 | Subtraction of Signed Numbers

Subtraction is the process of changing the sign of the second number and adding to the first. 65-34 is 65 + (-34).

So when we do subtraction, we actually add the two's complement form (i.e., the negative) of the second number to the first number. This is what computers actually do when they perform subtraction. In the discussion of subtraction in Section 0.7.4, this was not explicitly mentioned, because the idea then was to present certain other intricacies related to subtraction. Now let us

discuss subtraction for 8-bit signed numbers. Keep in mind that the range of signed numbers usable with 8 bits is -128 to $+127$.

Example 0.18

Perform subtraction of the following signed numbers:

- i) $+26$ from $+68$
- ii) $+26$ from -68

Solution

- i) $+26$ from $+68$

This comes to be a computation in the form of $68 + (-26)$. For this, the two's complement form of 26 should be added to 68.

Decimal		Binary	
68	is	0100 0100	+
-26	is	<u>1110 0110</u>	
		<u>1 0010 1010</u>	

Ignore the extra bit generated. Since the MSB is '0', the result is positive. The result is 0010 1010 i.e., 42.

- ii) $+26$ from -68 i.e., $-68 - (26) -68 + (-26)$

-68	is (in two's complement form)	1011 1100	+
-26		<u>1110 0110</u>	
-94		<u>1 1010 0010</u>	

Ignore the extra bit generated. Since the MSB of the 8-bit result 1010 0010 is '1', the difference of the two numbers is negative. Take the two's complement of this. 0101 1110 i.e., 94. So the result of the computation is -94 .

Example 0.19

Find the result of the following subtraction:

- i) -56 from $+23$
- ii) -56 from -23

Solution

- i) -56 from $+23$

The computation to be done is $+23 - (-56)$ i.e., $23 + 56$. This turns out to be the addition of the two positive numbers 23 and 56.

23	+	0001 0111	+
<u>56</u>		<u>0011 1000</u>	
<u>79</u>		<u>0100 1111</u>	i.e., 79

- ii) -56 from -23

The computation to be done is $-23 - (-56)$ i.e., $-23 + 56$.

-23	+	1110 1001	+
<u>56</u>		<u>0011 1000</u>	
<u>33</u>		<u>1 0010 0001</u>	

Ignore the extra bit generated. The MSB of the 8-bit result 0010 0001 is '0'. So the number is positive and is 33 in decimal, as it should be.

Overflow into the Signed Bit

Whenever we use 8-bit signed numbers in addition or subtraction, the result is found to be correct in sign and magnitude if it is within the range of -128 to $+127$. However, suppose this is violated? What happens then? A typical case is when two negative numbers are added. Try adding -100 and -55 . Both the operands are within the allowed range. See the addition.

$$\begin{array}{r}
 -100 \quad + \quad 1001 \ 1100 \quad + \\
 \underline{-55} \quad \quad \underline{1100 \ 1001} \\
 -155 \quad \quad \underline{1 \ 0110 \ 0101}
 \end{array}$$

Ignore the extra carry bit and look at the 8-bit result. The MSB of the result is '0' indicating that it is a positive number. However, we know that the answer is negative. What caused the error? Because the sum was too large (larger than -128) to fit into the 8 bits allotted to it, there was an 'overflow into the sign bit' causing the sign bit to be changed. (A similar issue occurs when we add two positive numbers and the sum is greater than $+127$). In computers there is a flag which tells us when there is an overflow into the sign bit causing it to be inverted. These matters will be discussed in detail when we do programming.

0.7.7 | Addition of Numbers of Different Lengths

We have discussed computer arithmetic in detail, because it is very important to be clear about it, so as to be able to understand how the microprocessor responds to different data types and arithmetic operations. Now let's try to understand how data of different data widths are dealt with.

Data can have different sizes depending on the processor. The 8086 can have data of 8 bits and 16 bits, while Pentium can handle 8, 16 and 32 bits internally. Sometimes it may be required to add/subtract data of different widths. In these cases, the important thing to do is to equalize the size of the data involved. Processors do not allow addition/subtraction of data of different widths. So a byte will have to be converted to a 16-bit word, if it has to be added to a 16-bit number. The way it is done depends on whether the data is signed or unsigned. For unsigned data, the byte is appended with zeros in the upper byte, and converted to a 16-bit word. For signed data, the byte should be 'sign extended' to make it a 16-bit word. Refer Section 0.8.6 once again to convince yourself of the necessity for this.

Example 0.20

Add the unsigned numbers 35H and 7890H.

Solution

In this, 35H is appended with zeros to make it 0035H.

$$\begin{array}{r}
 0035\text{H} \quad + \\
 \underline{7890\text{H}} \\
 78\text{C5H}
 \end{array}$$

Example 0.21

Add the following signed numbers:

- i) 45H and A87CH
- ii) A8H and 1045H
- iii) F5H and B45CH

Solution

- i) In this 45H should be made into a 16-bit number. Check the MSB of this byte. It is '0', meaning that it is a positive number. The sign bit when extended to 16 bits makes the number 0045H. Then the addition is

$$\begin{array}{r} 0045\text{H} \quad + \\ \underline{\text{A87CH}} \\ \text{A8C1H} \end{array}$$

- ii) In this the byte is A8H, which has an MSB of '1'. Thus, sign extension makes it FFA8H. Now the addition is

$$\begin{array}{r} \text{FFA8H} \quad + \\ \underline{1045\text{H}} \\ \text{1 0FEDH} \end{array}$$

The extra bit generated is ignored, like we have done in Section 0.7.3 on signed number computation. To be sure that this is correct, verification can be done as below.

A8H is -88
1045H is +4165

Adding the two, gives us 4077 whose hex representation is 0FEDH.

- iii) Add F5H and B45CH

In this, F5H is sign extended to be FFF5H

Adding

$$\begin{array}{r} \text{FFF5H} \quad + \\ \underline{\text{B45CH}} \quad \dots \text{note that this is a negative number} \\ \text{1 B451H} \end{array}$$

Ignoring the extra carry bit, the sum is B451H, a negative number. To verify, find the decimal equivalents of the numbers which are -11 and -19364, which when added, give -19375.

Note You may also verify that, without extending the negative sign, a wrong result is obtained. All the calculations we have done can be verified easily using the scientific calculator available on the PC. So, try to be adept in the use of that calculator.

0.8 | Units of Memory Capacity

A memory device is one in which data is stored. How much data a memory device can store depends on its capacity. The capacity of memory is specified as multiples of bytes since memory is byte organized, which means that one byte is stored in one location in memory. So, if there

are 100 locations in a memory device, 100 bytes are stored. We all have heard of memory capacity being mentioned in terms such as bytes, kilobytes and megabytes. Now let us quantify these terms. You will be hearing these terms throughout the use of this book.

A byte is 8 bits. A word is not really defined. It depends on the processor used. For the 8086, a word size is 16 bits. A 32-bit processor may claim to have a word size of 32 bits. Memory capacity is always specified in bytes.

$$2^8 = 256 \text{ bytes}$$

$$2^{10} = 1024 \text{ bytes} = 1 \text{ KiloByte or 1KB}$$

$$2^6 \times 2^{10} = 2^{16} = 64 \text{ KB} = 65,536 \text{ bytes}$$

$$2^{10} \times 2^{10} = 2^{20} = \text{one Mega Byte (1MB)} = 1024 \times 1024 = 1,048,576 \text{ bytes}$$

$$2^{10} \times 2^{20} = 2^{30} = \text{one Giga Byte (1 GB)} = 1024 \times 1024 \times 1024 = 1,073,741,824 \text{ bytes}$$

$$2^{10} \times 2^{30} = 2^{40} = \text{one Terra Byte (TB)} = 1024 \times 1024 \times 1024 \times 1024 = 1,099,511,627,776 \text{ bytes}$$

There are also higher units, which are not so common in usage as yet, but things will change soon, no doubt about it. Some of these units are:

$$\text{Peta Byte (PB)} = 2^{50} \text{ bytes}$$

$$\text{Exa Byte (EB)} = 2^{60} \text{ bytes}$$

$$\text{Zetta Byte (ZB)} = 2^{70} \text{ bytes}$$

0.9 | The 8085 Microprocessor

In this section, we will do a brief review of the 8085 microprocessor. Even though it is not a member of the x86 family, it has an important place in the history of Intel's microprocessors. It was a very popular microprocessor in the days before the advent of the x86 family. It was never used in a PC, but became popular for various uses where computation was involved. It also remained popular in the academic field where it was considered an easy model to teach and understand. Hence, a peep into the attributes of the 8085 would be in order. This section may also serve as a first step in the study of microprocessors, in general.

The 8085 was released by Intel in 1976. It is a 40-pin DIP (dual in line package, in which the pins are mounted as two parallel rows of a rectangular housing). This chip has eight data pins and 16 address pins. Thus data can be read in and out from the chip as eight bits at a time. The fact that it has 16 address lines imply that the amount of memory it can access is 2^{16} bytes – i.e., it can have access to memory of 64 KB at the maximum. Besides the data and address pins, it has a number of control pins as well. We will discuss the hardware aspects later and concentrate on the programming model first, because it is easier to understand a processor starting from such a model.

0.9.1 | The Programming Model of 8085

Figure 0.9 shows the programming model of the processor. The functional block diagram shows an arithmetic logic unit (ALU), and its associated blocks. The whole system is controlled by the timing and control unit which synchronizes the activities of the processor with a central clock. All arithmetic and logic computations take place in the ALU. Instructions from memory are brought to an instruction register (not shown) and decoded by the instruction decoder. Accordingly, the specified operation takes place in the ALU. The accumulator is an 8-bit register, called A, and its importance in 8085 is that it serves to hold one of the operands in most two-operand arithmetic and logical operations.




Figure 0.9 | Programming model of the 8085

Also, in many instructions, it is an implied register. There is a flag register which gives information about the result of arithmetic and logical computations. In such cases, 'flag bits' get altered as part of instruction execution.

The model also shows two sets of registers – one is a set of general-purpose registers and the other is a set of address registers (See Fig 0.10). All processors have registers which are used as temporary storage for operands. 'General-purpose registers' are used for data and they are 8-bit registers. The address registers are used for holding memory addresses and since addresses are 16-bit long, the address registers are 16 bits in size.

0.9.1.1 | General-Purpose Registers

They are also designated as scratchpad registers, and are 8 bits long. The 8085 is an 'accumulator' based machine, which means that for many operations (mostly arithmetic, logical and I/O), it is mandatory to have the operand/one of the operands in the accumulator. In such cases, this register is implied and its name is not written in the instruction. It is notated as 'A' when used in other instructions.

The other 8-bit registers are B, C, D, E, H and L. There is the provision to combine them in pairs to make them 16-bit registers. The allowed combinations are BC, DE and HL where, in the 16-bit form, the first register holds the upper byte of the 16-bit number. Thus, in the BC combination, B holds the upper byte and C the lower byte. Some instructions view the

A register and flag register together as a 16-bit register called the processor status word (PSW). In that case, the A register is the most significant byte and the flag register is the least significant byte. See Fig 0.11

Flag Register What is the importance of the flag register? It is an 8-bit register which is a combination of 1-bit flip flops are designated as 'flag bits'. The flags of 8085 are all status or conditional flags, in the sense they are set/reset as a result of the status or condition of the result of certain (not all) instructions. The conditions of the flags are important because it is based on this, that conditional branching is performed. See Fig 0.12. Now let us understand each bit of the flag register.

Z or Zero flag When the result of a computation is zero, this flag is set ($Z = 1$), otherwise, it is reset.

PROGRAM COUNTER (16)			
STACK POINTER (16)			
H	(8)	L	(8)
D	(8)	E	(8)
B	(8)	C	(8)
A	(8)	FLAGS	(8)

Figure 0.10 | Register structure of the 8085




Figure 0.11 | The A and the flag register combined to form the PSW




Figure 0.12 | The Flag register bits

C or Carry Flag If an operations causes a bit to overflow from the MSB, this flag is set ($C = 1$), otherwise it is reset.

S or Sign flag If the MSB of the result is '1', this flag is set ($S = 1$), otherwise it is reset.

P or Parity flag When the result has even parity, this is set ($P = 1$), otherwise, it is reset.

AC or Auxiliary Carry flag This flag is set ($AC = 1$) when there is a carry from D3 to D4 of the result. This flag is used only for BCD calculations, and there is no 'branching' operation associated with this flag.

0.9.1.2 | Address Registers

Figure 0.10 shows two 16-bit registers – the program counter and the stack pointer. They are address registers.

Program Counter (PC) This is a 16-bit register and it is used to 'sequence' the instructions being executed. To elaborate this point, think of the process of instruction execution. Instructions are stored in memory and when execution is to be done, they are fetched one byte at a time from memory to the processor. The PC always points to the 'next instruction' to be executed. When a branch instruction appears, it is apparent that the current sequence of instruction fetching is disturbed and a new sequence will be taken up. In that case, PC will be loaded with the 'branch' address.

Stack Pointer (SP) The stack pointer (SP) is a 16-bit register which points to a memory location in RAM which will hold temporary values in an area of RAM called the stack. The stack is explained in detail in Chapter 3.

0.9.2 | Assembly Language Programming

We will now attempt to write a few programs using the assembly language instructions of 8085. Programs that are written are saved in RAM from where instruction bytes are fetched one byte at a time for execution. An instruction may need one or more than one byte space for storage. Executing an instruction may need data from memory addresses as specified in the instruction. Program execution proceeds sequentially one after the other in the order that they are fetched. In this sequence, it may happen that a 'branch' instruction is encountered. Branching means 'jumping' to a location which is not in the sequence that is followed now. A branch instruction takes control to a different set of instructions written in a different area in memory. Most cases of branching depend on conditions, though unconditional branching is also possible.

The general format of writing assembly code is as shown below:

```
label:           instruction           ; comment
```

Labels are not needed for all lines, and comments are optional – they are used only to enhance the understandability of the instruction line. With these basic ideas, let us go on to the finer aspects of 8085 programming. For most assemblers, a semi-colon precedes the comment, and a colon is needed along with a label.

0.9.2.1 | Modes of Addressing

Now that we have the list of registers available in 8085, we can start programming the chip – but we also have to know the modes of addressing, before we can actually plunge into programming.

An addressing mode for an instruction implies the way an operand is presented for execution by the instruction. There are single operand as well as double operand instructions for the 8085. In the former case, the operand will be in a register or memory. In the double operand case, the operands will have to be specified to be in some memory address or register and another register. However, there is no mode in which both operands are in memory – this means that if there are two operands, at least one of them has to be in a register.

The general format of a two-operand instruction of any Intel processor can be illustrated by the ‘move’ instruction which is of the form,

MOV destination, source

In this, the content of the source is copied to the destination. The source remains unchanged. We will discuss the ideas related to addressing modes using the ‘MOV’ and ADD instructions.

Note: The code lines are not case sensitive.

The addressing modes available are:

1. Register addressing.
2. Direct addressing.
3. Indirect addressing.
4. Immediate addressing.

Register Addressing

In this mode, only registers are involved. Both the source and destination are registers.

MOV A, B	;copy the contents of B to A
MOV C, A	;copy the contents of A to C
ADD A, H	;add the contents of H and A – sum to be in destination i.e., A

Direct Addressing

In this, one of the operands is ‘directly’ addressed. This operand is in memory and its address is ‘directly’ mentioned in the instruction.

LDA 0987H	;load accumulator direct – the address is the 16-bit number 0987H. ;The content of this location is copied to the accumulator (register A)
STA 2345H	;store accumulator direct – the content of the accumulator is stored in the address 2345H
IN 56H	;copy the data in input port with address 56H to register A

Indirect Addressing

Here, one of the operands is ‘indirectly’ specified. It is in memory, but instead of writing its address directly in the instruction, the address is loaded into a 16-bit register and that register is specified in the instruction.

MOV C, M	;copy the content of the ‘memory specified indirectly’ to the C register. For this instruction, the ‘address’ is to be loaded into the HL register pair previous to this instruction
----------	--

STAX B	;‘store accumulator indirect’ in this, the content of the accumulator is stored in the address specified by the BC register pair.
--------	---

Immediate Addressing

Here, the source data is written in the instruction itself

MOV A, 78H	;copy the number 78H to A
ADD A, 67	;add the content of A and the number 67. The sum is to be in A.
LXI H, 3453H	;the immediate number 3453H is loaded in the HL register pair

0.9.3 | Instruction set of 8085

The instruction set can be divided into functional groups as:

- i) Data transfer instructions.
- ii) Arithmetic instructions.
- iii) Logical instructions.
- iv) Branch instructions.

We will discuss each group in more detail now, using a sample set of each group.

0.9.3.1 | Data Transfer Instructions

These instructions are responsible for transferring data between registers, register and memory, I/O and accumulator. These also include instructions for loading an 8-bit or 16-bit number into a register of the same size. For certain instructions, the accumulator (A register) is implied. The instructions are listed as:

- | | |
|-----------------------|---|
| i) MOV Rd, Rs | ;copy data from Rs(source) to destination register(Rd) |
| ii) MVI Rd.data8 | ;move 8-bit data to 8-bit register |
| iii) OUT port address | ;send data from accumulator to output port |
| iv) IN port address | ;take in data from input port to accumulator |
| v) LXI Rp, data16 | ;load 16-bit data to register pair |
| vi) MOV R, M | ;copy data to register, from indirectly specified memory |
| vii) MOV M, R | ;copy data from register, to indirectly specified memory |
| viii) LDA address | ;load data to accumulator from specified address |
| ix) STA address | ;store data from accumulator to specified address |
| x) LDAX Rp | ;load to accumulator, the data from memory address specified by the register pair (Rp) |
| xi) STAX Rp | ;store into the memory address specified by the register pair Rp, the data in the accumulator |

0.9.3.2 | Arithmetic Instructions

For all arithmetic instructions, the A register is implied to contain one operand (if there are two operands for the instruction). The arithmetic group includes instructions to add, subtract, increment, decrement and decimal adjust after addition of BCD. Note that there are no

instructions for multiplication and division. **It is important to remember that all conditional flags are affected by arithmetic instructions.**

- i) ADD R ;add the content of R to A – sum in A
- ii) ADI data8 ;add the immediate data to A – sum in A
- iii) ADC R ;add the content of R, A and the carry bit – sum in A
- iv) ADD M ;add the content of indirectly specified memory to A
- v) SUB R ;subtract from A the data in R – result in A
- vi) SUI data8 ;subtract from A the immediate data – result in A
- vii) SBB R ;subtract from A, the content of R and carry – result in A
- viii) INR R ;increment the content of the 8-bit register R by 1
- ix) DCR R ;decrement the content of the 8-bit the register R by 1
- x) INR M ;increment by, the content of the indirectly specified memory
- xi) DCR M ;increment by 1, the content of the indirectly specified memory
- xii) INX Rp ;increment by 1, the content of the register pair Rp
- xiii) DCX Rp ;decrement by 1 the content of the register pair Rp
- xiv) DAA ;decimal adjust accumulator – to be done after BCD addition

Now let us write a few programs using the instructions we have just learned.

Example 0.22

Bring data from the memory location 4567H, add 56H to it and store the sum in location 0567H.

Solution

```
LDA 4567H ;copy into the A,data from address 4567H
ADI A,56H ;add 56H to this – sum is in A
STA 0567H ;store the content of A in address 0567H.
```

The above is a very simple program in which we have used only the load, store and an immediate add instruction.

Example 0.23

Write a program to add an immediate data to a data in memory address 0987H. Store the sum in the next address in memory.

Solution

In this, an immediate data is loaded into register B. The other operand which is in memory is indirectly addressed. For that, the address of the memory location is loaded into the register pair HL, using immediate addressing. Then the data is moved to A, using indirect addressing. This is added to B and the sum is now in A. To store the sum in the next address, increment the address pointer in the HL register pair. See the program below.

```
MVI B,3CH ;load the immediate data in B
LXI H,0987H ;load address in the H-L register pair
```

MOV A,M	;copy data in the address pointed by HL
ADD A,B	;add the content of B to A – sum in A
INX H	;increment the content of HL
MOV M,A	;store A in the address pointed by HL

Example 0.24

There is a byte of data in location 6756H, which is to be overwritten by a number which is less than it by 5. This is done below.

Solution

LXI B,6756H	;load the address of data in the BC pair
LDAX B	;load in A the content pointed by BC
SUI 05	;subtract 05 from A-result is in A
STAX 6756H	;store A in the address pointed by BC

Using data transfer and arithmetic instructions, a number of computations can be done. However, a problem arises if we want to do repetitive operations on large chunks of data. The problem is that of having no ‘stopping condition’. Say, we want to do the same operation on different data ten times – for that, we need a counter which will tell us when the count of 10 is reached. When this happens, the repetitive action can be stopped. The capability of being able to take decisions based on conditions is what gives any computer its power, and the 8085 is no exception. It has instructions for branching which are as listed below. Branching effectively changes the sequential nature of instruction execution, by causing control to move to a new set of instructions written at a new address. Branching can be conditional or unconditional – the conditions used are the settings of the flags, as we will see now.

0.9.3.3 | Branch Instructions

- | | |
|--------------------|--|
| i) JMP address16 | ;jump unconditionally to the given address |
| ii) JZ address16 | ;jump on zero-jump to the given address if Z = 1 |
| iii) JNZ address16 | ;jump on non-zero-jump to the given address if Z = 0 |
| iv) JC address16 | ;jump on carry-jump to the given address if C = 1 |
| v) JNC address16 | ;jump on no carry-jump to the given address if C = 0 |
| vi) CALL address16 | ;call a subprogram written in the address given |
| vii) RET | ;return to the original program after executing the subprogram |

For 8085, all addresses are 16-bit in length.

Example 0.25

Now let us write a program for multiplying 24 by 10. Since there is no multiply instruction for 8085, multiplication is achieved by repeated addition. We do it by adding 24 to itself 10 times.

Solution

MOV A,0	;make the content of A = 0
MOV B,10	;load the multiplier (10) in B

```

AGAIN:    ADI A,24      ;add the multiplicand(24) to A
           DCR B         ;decrement B
           JNZ AGAIN     ;check if Z (zero flag) is non - zero - if
                           so, jump to AGAIN
           LDA 3400H     ;if Z = 0, store the sum in address 3400H

```

In the Example 0.25, B acts as a counter. Initially A = 0. Then 24 is added to it and each time this is done, the counter B is decremented by one. By the time, 10 additions are over, the content of B will be equal to 0. When the decrement operation causes the result to be zero i.e., B = 0, the Zero flag is set (Z = 1). This is the stopping condition for the addition loop. Once the addition is over, the result is saved in memory. Then the loop is exited and the content of A is copied to a location in memory.

Note Since A has only a data size of 8 bits, the maximum number it can hold is 255. This should be kept in mind.

Example 0.26

Two numbers are stored in memory. Verify if their sum is greater than 255. If yes, send the ASCII value of Y to an output port with address 78H. Otherwise send N.

Solution

If the sum is greater than 255 (FFH), a carry will be generated which will set the carry flag. So the verification consists of testing the carry flag. See the program flow.

```

LXI H,0897H      ;load HL with the first address
MOV A,M          ;load data from that address to A
MOV B,A          ;move that data from A to B
INX H            ;increment pointer register pair HL
MOV A,M          ;load the second data from memory into A
ADD A,B          ;add the content of A and B-sum in A
JC YES           ;if C = 1, sum >255; jump to YES
MOV A,'N'        ;if C = 0, it means sum <255, so A = 'N'
JMP XIT          ;jump unconditionally to XIT
YES: MOV A,'Y'    ;this is when sum>255.A = 'Y'
XIT: OUT 78H     ;send A to output port with address 78H

```

In Example 0.26, a number of points need elaboration. We are already familiar with getting data from memory using indirect addressing. Once the two 8-bit numbers are brought to registers A and B, they are added. If the sum >FFH, the carry flag gets set, which is tested by the JC instruction. If the condition is true, control branches to the label 'YES'. Here the ASCII value of Y is copied to A. ASCII is represented by single quotes of the character. When we write 'Y', actually the number corresponding to the ASCII of Y i.e., 59H is what is being copied. If the YES branch is not taken, the A register is to be loaded with the ASCII of N. Either way, there is an ASCII character in A. The last step is sending this number to an output port. For that, the output port address is mentioned in the instruction, but the register A is implicit. The output port may be a display device.

0.9.3.4 | Logical and Bit Manipulation Instructions

8085 has a set of logical and bit manipulation instructions as well. The list is given below. Flag bits are affected by these instructions:

- i) ANA R ;logically AND the contents of A and R – result in A
- ii) ANI data8 ;logically AND the contents of A and data given – result in A
- iii) ANA M ;logically AND the contents of A and memory – result in A
- iv) ORA R ;logically OR the contents of A and R – result in A
- v) ORI data8 ;logically OR the contents of A and data given – result in A
- vi) ORA M ;logically OR the contents of A and the memory – result in A
- vii) XRA R ;logically XOR the contents of A with that of R – result in A
- viii) XRI data8 ;logically XOR the contents of A and data given – result in A
- ix) XRA M ;logically XOR the contents of A and the memory – result in A
- x) CMP R ;compare the contents of A and R – only flags affected,
- xi) CMP M ;compare the content of A and memory – only flags affected
- xii) CPI data8 ;compare the contents of A and given data – only flags affected
- xiii) RLC ;rotate the bits in A once, to the left
- xiv) RAL ;rotate the bits in A and the carry bit, to the left, once
- xv) RRC ;rotate the bits in A once, to the right
- xvi) RAR ;rotate the bits in A and the carry bit, to the right, once

Most of the above instructions do not need much explanation – but the ‘compare’ instruction does. The format of this instruction is

CMP destination, source

This instruction compares the two operands, causes the conditional flags to be affected, **but neither the destination nor the source changes**. Comparison is done by a subtraction operation, and the flags are set/reset according to the result of this. However, only two flags really matter – they are the Zero flag and the Carry flag. Consider the instruction CMP destination, source. The condition of the flags will be as shown in Table 0.5. Thus to test for equality of two data, the zero flag can be used, and for comparing the magnitude of the numbers, the carry flag is tested. We have used the latter condition in Example 0.27.

Table 0.5 | Flag Settings After a Compare Instruction

If	C flag	Z flag
destination > source	0	0
destination < source	1	0
destination = source	0	1

Example 0.27

Compare the number in the memory location 4566H with the content of B. Send the bigger number to the output port with address 67H.

Solution

```
LDA 4566H      ;load into A, the number in memory
CMP A,B        ;compare it to the number in B
JNC YES        ;If C = 0, it means that A>B, jump to YES
MOV A,B        ;C = 1, means B>A, so copy bigger no to A
YES: OUT 67H    ;send the bigger number to the output port
```

With this, we come to an end of our discussion on 8085 programming. There are a few more instructions, but since only a brief introduction to 8085 is possible here, we will straight-way go into the hardware aspects of 8085.

0.9.4 | Hardware Aspects of the 8085

Fig 0.13 shows the pin diagram of the processor. Let us discuss the pins in detail. As we do that, the various features of the processor will become clear to us.




Figure 0.13 | Pin Diagram of the 8085

0.9.4.1 | Multiplexed Address/Data Bus

The 8085 has 16 address lines and 8 data lines, but if we look at the pin diagram we will not find $16 + 8 = 24$ pins dedicated for address and data. The reason is that the address and data pins are multiplexed. The lower address lines carry data also. See the pins titled AD_0 to AD_7 . They are multiplexed pins. They cater to address as well as data, but not at the same time. The bus cycle for accessing memory or I/O is arranged in such a way that first the address appears on the address/data pins AD_0 to AD_7 (lower byte of address) as well as on A_8 to A_{15} . Along with this event, a signal named ALE (Address Latch Enable) will go high. An 8-bit latch is to be connected (externally) to the higher order address lines and the ALE is to act as a clock for the latch. So, when ALE goes high, the higher order address information appears at the output of the latch and remains there. Fig 0.14 makes this clear. After this, the address information can be removed from the AD_0 to AD_7 pins, and then they function as data pins named D_0 to D_7 . Thus, using the latch 74LS373, address de-multiplexing has been achieved. However, why was multiplexing done, in the first place? In the early days of microprocessor development, packing problems required pin counts to be as small as possible.




Figure 0.14 | Address and data de-multiplexing

0.9.4.2 | Read and Write

Refer to the pin diagram in Fig 0.13 again. Pin nos. 31 and 32 are $\overline{WR}$ and $\overline{RD}$. They are active low signals. They are the control signals for reading from/to memory/I/O. Recollect the way these signals are used (Section 0.2.5). However, these control signals are used irrespective of whether reading/writing pertains to memory or I/O. The processor has another pin to differentiate between these two cases. See Pin no. 34. It has the designation $IO/\overline{M}$. This is a control signal sent from the processor, and goes high for I/O access and low for memory. External logic is used as shown in Fig 0.15 to get the four separate control signals – $\overline{MEMRD}$, $\overline{MEMWR}$, $\overline{IORD}$ and $\overline{IOWR}$. Fig 0.16 shows the processor with the separate address bus, data bus and read/write control signals.

0.9.4.3 | Power Supply and Clock

Now refer Fig 0.17 which show the functional block diagram of the processor, power and clock pins. Pins 1 and 2 are the pins between which a crystal is to be connected. There is an internal crystal oscillator for the chip, and the only thing that we do is to connect a crystal externally. This defines the clock frequency for the processor. Internally the clock frequency gets divided by 2. So, if we need an internal clock of 4 MHz, an 8 MHz crystal is to be connected. The maximum allowable frequency of operation for the 8085 is 5MHz. For that, a crystal of 10 MHz is to be connected between X_1 and X_2 . Pin No. 37 is CLK OUT. From this pin, the clock of the processor can be extracted and used for timing any unit in the system which is required to be synchronized to the processor. The clock can also be divided to get lower frequencies which may be used for devices like an ADC (analog to digital converter), which may be part of the 8085 based system.

0.9.4.4 | Reset and Ready Pins

There are two pins pertaining to reset. $\overline{RESETIN}$ is an input signal to the processor. When this is pulled low, the processor resets and the program counter is cleared. Program execution starts




Figure 0.15 | Generating separate control signals for memory and I/O




Figure 0.16 | External hardware used for getting separate address and data bus and control signals

at the address 0000, which is the address corresponding to the value of the PC when the system is reset. While this is happening, RESET OUT can be used to reset other devices in the system concurrently.

Ready This is a signal sent to the processor by a peripheral. Usually peripherals are slower than processors. The time for accessing any peripheral or memory location is specified in terms of a number of clock cycles. For slower peripherals, more clock cycles will be needed. When that is the case, the processor looks at the READY pin that is controlled by the peripheral. As long as the READY pin is low, the access cycle does not end and thus more time is obtained for reading or writing. It must be arranged that the READY pin goes high once the required delay is obtained.

0.9.4.5 | Interrupt

An important feature of any processor is the capability to be interrupted. However, before discussing the interrupt pins, let us try to understand what is meant by an interrupt.

Say a processor is performing a particular task, executing instructions corresponding to that task. If another more important task is required to be done, the current task will have to be stopped temporarily and the other more urgent one must be taken up. Once the urgent task is




Figure 0.17 | Functional block diagram of 8085

done with and finished, the old one can be resumed. This is the scenario regarding interrupts. This means that the processor will have to be ‘interrupted’. Who gives the interrupt request? It could be a peripheral requesting urgent service. In this case, the peripheral places its interrupt request on a hardware pin of the 8085. Another case is when the processor is interrupted by software. Here, an instruction is written to interrupt the execution of the current task.

In either case, the sequence of events that follow the interrupt are the same. The processor completes the current instruction, saves the program counter and processor status and takes up the new task by going to a new location in memory. This new memory location is called the ‘interrupt vector’ corresponding to the specific interrupt. What is there at the interrupt vector? It is at this address that the Interrupt Service Routine (ISR) is stored. The ISR is the program that the interrupting device wants to get executed. Thus, for each interrupting device, there should be an ISR at an ‘interrupt vector’.

Table 0.6 | Interrupts Pins and Their Vectors

Interrupt pins	Interrupt vector
TRAP	0024H
RST7.5	003CH
RST6.5	0034H
RST5.5	002CH

Interrupt Pins of the 8085 This processor has five hardware interrupt pins on which peripherals can place their interrupt requests. The first one is the INTR (Interrupt) pin. A peripheral should place a high on this pin for a minimum specified time, which causes the processor to get interrupted. The processor responds by lowering the *INTA* line which is an ‘Interrupt Acknowledge’ signal. The peripheral on getting this, tells the processor where the ISR for this task is placed – we say that the peripheral gives the ‘interrupt vector’ for its task. The processor goes to this address, executes that task, and after that goes back to the task that had been interrupted.

This is the case for an interrupt on the INTR pin. How about the other interrupt pins? They are TRAP, RST 7.5, RST 6.5, and RST 5.5. The specialty of these pins is they are ‘vectored’ interrupt pins. This means that when a peripheral places its request on any of these pins, the peripheral does not have to give the address of the ISR. For each of these interrupts, there is a specific pre-defined address to which control automatically goes – the hardware of the processor takes care of this. The interrupt vectors for these pins are as in Table 0.6. There is also another issue regarding interrupts. Suppose a number of peripherals interrupt simultaneously on the many interrupt pins of the 8085, obviously, only one of the requests can be serviced. Which one? There is a pre-defined priority for the interrupt pins – TRAP is the highest priority pin followed by RST 7.5, 6.5, 5.5 in that order and last, the INTR pin.

0.9.4.6 | Hold and HLDA

See pins 38 and 39, they correspond to the HOLD and HLDA (hold acknowledge) pins. These are the pins that co-ordinate DMA activities of the processor. Now, what is DMA? Normally, a data transfer between memory and an I/O device always takes place with the processor acting as an intermediary. If data from memory is to be sent to an output device, it first has to be sent to the processor’s accumulator and from there, it is transferred to the peripheral. This is okay if the data involved is small. However, for bulk data transfer, this becomes cumbersome. So, the processor allows data to be transferred directly between a peripheral and memory, while the processor stands aside after just initiating the process.

If, say, a peripheral wants DMA service, it sends a request to the processor. The DMA request is placed on the HOLD pin. Since DMA is a high priority issue, the processor stops whatever it is doing and sends the acknowledge signal HLDA to the peripheral. With this, the processor’s buses are tri stated, such that the processor effectively gets cut off from the system. Thus a direct path of data transfer is established between the peripheral and memory. Once this data transfer is over, the HOLD request is removed by the peripheral and the processor gets back its status as the master of the system.

0.9.4.7 | SID and SOD

These are two pins that most other microprocessors do not have. They are the Serial Input Data (SID) and Serial Output Data (SOD) pins. Using SID, data can be sent in serially, converted

to parallel form (internally) and saved in a register. Similarly parallel data can be converted to serial form and sent out as one bit at a time, using the SOD. Of course, there are special instructions to aid in the serial transmission process.

Conclusion

We have discussed all pins of the 8085 processor, except the status pins S_0 and S_1 . These are pins which carry some status information, which is not very important for small systems. With this, we conclude our review of the 8085, and will go on to the task of understanding the x86 microprocessors, starting from the 8086. That will be from Chapter 1 onwards.

KEY POINTS OF THIS CHAPTER

- The first PC which made a great impact in the world of computing was the IBM PC with Intel's 8088 microprocessor being its CPU. This PC was launched in August 1981.
- Gradually the x86 family of microprocessors became the *de facto* standard for PCs the world over.
- A computer system consists of a CPU, memory and I/O which communicate with one another through the system bus.
- The system bus comprises the data bus, address bus and the control bus.
- A processor's activities are restricted to fetching, decoding and executing instructions.
- For reading and writing from/to memory, a number of clock cycles of time are required. The time expended for this is called the memory access time.
- When comparing assembly language programming with high level language programming, we conclude that the former is faster in execution and more efficient and compact, but is more difficult to learn and master.
- RISC and CISC are two different philosophies in computer design, and even though a lot of controversy still rages around which is better, the two seem to have merged, more or less.
- Computers do all the computations in binary, but for entering data through the keyboard and for displaying it on the monitor, ASCII codes are used.
- Negative numbers are represented in two's complement form by all computers.
- The 8085 microprocessor is still a very popular chip even though it is not used in PCs.
- The 8085 has 16 address lines and 8 data lines.
- The 8085 is one of the processors of Intel which has the clock generator inside the chip.

QUESTIONS

1. Which is the first noteworthy microprocessor developed by Intel?
2. Name two 8-bit microprocessors of Intel which are not part of the x86 family.
3. Which microprocessor was used in the first IBM PC?
4. What is IBM's contribution to the PC revolution?
5. What is meant by the term 'x86 family' of microprocessors?
6. Name the three most important components of a computer system.
7. Have you heard of the term 'bus contention'? What does it mean in the context of a computer system?

8. If the data bus width of a processor is 64 bits, what would you say about its complexity and capability?
9. If the address bus of a processor is 64 bits, what is its address space?
10. What could be a 'multi processing' system?
11. What is the first step in the execution cycle of a processor?
12. How does the system clock frequency influence the speed of processing?
13. If a system uses a 1.5-GHz clock, what is its clock period?
14. What is meant by the word 'system bus'?
15. Why should a computer have an I/O controller?
16. What are the difficulties involved in learning and using assembly language programming?
17. Name one distinguishing feature each of RISC and CISC computers.
18. How are the hexadecimal and binary number systems related?
19. When two signed positive numbers are added and the sum exceeds 127, what is the problem that arises?
20. What is the range of signed numbers that can be represented in 12 bits?
21. Why is the 8085 called an 'accumulator based' processor?
22. Distinguish between the LDA and LDAX instruction.
23. List the interrupt lines of 8085.
24. What is meant by DMA?

EXERCISE

1. Write briefly the history of the x86 family of microprocessors.
2. Explain RISC and CISC processors distinguishing between the two, elaborating the advantages and disadvantages of each type. Name a few processors of each type.
3. Find out about the 'Apple Mac' – its history, its features and why and how it is considered a 'great' PC.
4. Write the decimal equivalent of the following numbers:
 - a) 31.3H
 - b) 1100.101B
 - c) A32.3H
 - d) 100101B
5. Convert the following numbers to binary form:
 - a) 34
 - b) 200
 - c) 90
6. Convert to hexadecimal format.
 - a) 3454
 - b) 4523
 - c) 789

7. Write the binary values for:
 - a) 34ADH
 - b) 78FH
 - c) 407BH
8. Write the hexadecimal values of:
 - a) 11000101010110001B
 - b) 10011111100001010B
9. Find the packed BCD representation of the following decimal numbers:
 - a) 45
 - b) 4678
 - c) 802345
10. Represent the packed BCD of the following numbers in hex:
 - a) 235
 - b) 9123
11. What is the ASCII of each of the following?
 - a) 7
 - b) 8
 - c) 0
 - d) A
 - e) Z
 - f) y
 - g) d
 - h) *
 - i) &
12. Find the two's complement representation of the following numbers in 8 bits:
 - i) -45
 - b) -90
 - c) -12
 - d) -34
13. Represent the following negative numbers using 16 bits:
 - a) -267
 - b) -4
 - c) -5676
 - d) -675
14. Perform binary addition for the following numbers:
 - a) 34 and 56
 - b) -52 and -70
 - c) -47 and +120
15. Convert to packed BCD and add,
 - a) 46 and 23
 - b) 55 and 67
 - c) 34 and 49
 - d) 99 and 44
16. Subtract after converting to binary form,
 - a) -20 from -75
 - b) +49 from +97
 - c) E5H from A4H

17. Add the following signed numbers:
 - a) F3H and 3245H
 - b) AH and F45H
 - c) B2H and 123EH
18. How many bytes constitute
 - a) 5 MB
 - b) 4 KB
 - c) 32 MB
 - d) 32 KB
 - e) 8 GB

1 THE ARCHITECTURE OF 8086



In this chapter, you will learn

- The internal architecture of the 8086.
- The registers of 8086 and their specific functions.
- The operation of the flags of 8086.
- The memory segmentation technique used by the x86 family.
- How to perform memory address calculations.
- The different addressing modes of 8086.

We will start our learning of the x86 family of microprocessors, by understanding the architecture of the 8086, on which is based the architecture of the whole family. The 8086 is a processor catering to 16-bit data (D0 to D15) and having a 20-bit address (A0 to A19).

1.1 | Internal Block Diagram of the 8086

The 8086 processor is a 16-bit processor internally as well as externally – which means that its data bus width as well as the bit size of its internal data registers is 16. Thus it can access an external source 16 bits at a time through its data bus. However it also has the capability to access and work on 8-bit (byte) data.

As the first step in understanding the architecture of the 8086, let us examine the internal block diagram of the 8086. Fig 1.1a shows the simplified block diagram and Fig 1.1b shows it in more detail. Let us examine the simplified diagram first. Though not specifically shown in the block diagram, it is to be noted that all actions in the processor are synchronized by a system clock, which provides the basic timing. There is a control unit which provides the control signals for the overall functioning of the processor.

We see that the internal block diagram has been partitioned into two logical units, namely, the Bus Interface Unit (BIU) and the Execution Unit. They interact directly with each other through the internal bus, but perform separately as two units with well-defined functions for each unit. Let us explore the internal configuration of these units, one by one.

1.2 | The Execution Unit

It contains the arithmetic and logic unit, the control unit, an internal bus, plus a few registers. Let us get a bird's eye view of the EU starting with the register set.




Figure 1.1a | Internal block diagram of the 8086 (simplified model)

Why Registers? All processors have **internal registers**. An internal register is actually an **on-chip RAM**. Some are **program visible**, while others are **not directly visible to programmers**. The more the number of program visible registers, the better it is for the programmer. This is because internal registers **are used as temporary storage for operands**, and if the operand is **already in a register**, it takes less time for execution of the associated instruction. Such registers are called **general-purpose** or **scratchpad** registers.

However, registers constitute on-chip RAM, which is costly. Also, more the number of registers, more are the number of bits required to identify and address a particular register. This can cause an increase in the length of **opcodes** in modes involving registers. Thus the number of registers a processor can have is a trade-off between a number of conflicting factors.

Register Set of 8086 Let us see how the register set of 8086 has been organized. The 8086 has **data registers**, **address registers**, **segment registers** and also a **flag register**. We will first examine the registers in the Execution Unit.

1.2.1 | The Scratchpad Registers

They are called so because they are **used for temporary storage**, just as we jot down a few things temporarily in a scratchpad. As shown in Fig 1.2, the 8086 has **four 16-bit general-purpose registers** labeled as **AX, BX, CX and DX**. Each of these registers can also be used as two separate




Figure 1.1b | Internal block diagram of the 8086 (detailed)




Figure 1.2 | Register model of the 8086

and exclusive 8-bit registers also i.e., AX has two parts AH and AL, where H and L stands for the high and low portions respectively. When used as 8-bit registers, AH and AL are separate, but they can be concatenated to form the 16-bit register AX. Similar is the case for the registers BX, CX and DX. Each register, while being **general-purpose**, has specialized uses as well. As such, let us understand these features clearly.

AX This is a 16-bit data register, but it can also be used as two unrelated 8-bit registers AH and AL. AL/AX is sometimes called the **accumulator**, but the relevance of the accumulator is less for 8086 compared to the earlier 8085 in which one operand is implied to be in the A register for many instructions. However, for the 8086, only for certain specialized instructions is the AL or AX register implied to contain one of the operands. In most other instructions, the source and destination must be specifically mentioned in the instruction. (The destination is the register or memory location in which the result of any operation is to be put.)

BX, CX and DX These are the other working registers of the 8086, which means that temporary data storage, arithmetic calculations and data manipulation can be done with these registers. They also can be used as 8-bit or 16-bit registers as mentioned earlier. Still, each of them has special functions.

Base register BX is frequently used as an **address register** in many based addressing modes. **Counting register CX** is used as a **counter** in many instructions. **Data register DX** is used in I/O instructions as a pointer to data by storing the address of the I/O port. Also, multiply and divide operations involving large numbers assume DX and AX as a pair.

Pointer and Index Registers SP, BP, SI and DI are **address registers**, and can be used only as **16-bit registers**. They are very important in programming, as they facilitate the use of various addressing modes.

BP and SP They are the **Base pointer** and **Stack pointer** respectively. **SP** always points to the top of the stack, while **BP** can point to any location in the stack. (We will discuss the concept of the stack shortly.) Thus normally, whenever we speak about data referred to by BP and SP, it is implied that we are talking about data in the stack.

SI and DI These are **Index registers**, labeled as **Source Index** and **Destination Index** respectively. They function as **address registers** in various addressing modes, but they have a **specialized function** as being the default registers in **string instructions**.

1.2.2 | Flag Register

It is a 16-bit register, of which 7 bits are unused. Out of the remaining bits, 6 bits are used as **conditional flags**. The others are **control flags**. A **conditional flag** is a single bit flip flop which is set or reset according to the result of an arithmetic or logic operation.

The conditional flags available are the **Carry (CF)**, **Zero (ZF)**, **Parity (PF)**, **Overflow (OF)** and the **Sign Flag (SF)**. Since conditional flags are an important ingredient in computation using assembly language, we will not defer our discussion on them.

Carry flag (CF) The carry flag (CF) gets set if there is a carry out from the most significant bit during a calculation. For example, when 8-bit addition causes the result to be greater than 8 bits, there is a carry out from the MSB (D7), which causes this flag to be set. For 16-bit operations, the carry will be from bit D15, and CF will be set.

Zero Flag (ZF) When the result of an arithmetic or logic operation is zero, the zero flag gets set. For example, if we keep on decrementing the contents of a register (say the CX register, which

D15	D14	D13	D12	D11	D10	D9	D8	D7	D6	D5	D4	D3	D2	D1	D0
U	U	U	U	OF	DF	IF	TF	SF	ZF	U	AF	U	PF	U	CF

U = unused

CONDITIONAL FLAGS

- OF = Overflow flag
- SF = Sign Flag
- ZF = Zero Flag
- AF = Auxiliary Carry Flag
- PF = Parity Flag
- CF = Carry Flag

CONTROL FLAGS

- DF = Direction Flag
- IF = Interrupt Flag
- TF = Trap Flag

Figure 1.3 | 8086 Flag register

is used as a count), it will finally become zero. At this instant, the zero flag gets set i.e., $ZF = 1$. Another case is when two numbers are compared. Comparison is achieved by subtraction. If the numbers compared are equal, the zero flag is set ($ZF = 1$) indicating equality of the operands.

Parity Flag (PF) The setting of this flag indicates the presence of an even number of bits in the least significant byte (lower 8 bits) of the destination. For example, after a particular arithmetic or logic operation, if the destination contains the number 11100111, the parity flag is set to indicate even parity. Even if the operation is on 16-bit data, only the parity of the lower 8 bits is verified. This flag does not find much applications in current programming practices.

Sign Flag (SF) After an arithmetic or logic operation, if the result contains a negative number, the sign flag is set. In essence, it contains the MSB of the result, which can be interpreted as the sign bit in signed arithmetic operations.

Auxiliary Carry Flag (AF) This flag functions similar to the Carry flag, except that the overflow is from bit D3 into D4. It thus indicates a carry out from the lower 4 bits. The need for this flag is for the **Decimal Adjust instruction**, which is important in BCD number calculations. There are no other instructions that directly test the state of this flag and no conditional branching is associated with this flag.

Overflow Flag (OF) This flag is set under one of the following conditions

- i) there is an overflow into the MSB (8^{th} or 16^{th} bit) from the bit of lower significance, but no carry out from the MSB,
- ii) there is a carry out from the MSB, but no carry into the MSB.

This flag indicates that the result of a signed number operation is too large, causing the higher order bit to overflow into the sign bit, thus changing the sign bit.

Note: To understand the use of flags we will use a few simple instructions, with adequate explanation. But the instruction set will be dealt in detail, only from Chapter 2 onwards.

Example 1.1

Find the status of the flags CF, SF, AF after the following instructions are executed.

```
MOV AL, 35H
ADD AL, 0CEH
```

Solution

The format of a MOV instruction is MOV destination, source.

In the above two instructions, a number 35H is first moved to AL. In the next line, 0CEH is added to AL. The sum will be in AL, the destination register.

$$\begin{array}{r}
 35\text{H} \qquad \qquad 0011 \ 0101 \ + \\
 + \ \underline{\text{CEH}} \qquad \underline{1100 \ 1110} \\
 \hline
 103\text{H} \qquad \underline{1 \ 0000 \ 0011}
 \end{array}$$

After computation, AL will contain 0000 0011 and

CF = 1 since there is a carry out from D7.

SF = 0 since the sign bit (MSB) of the 8-bit destination is 0.

AF = 1 since there is an overflow from D3 to D4.

Example 1.2

Show the effect of the following instructions on the CF, ZF and OF bits of the flag register.

```
MOV BX, 45ECH
ADD BX, 7723H
```

Solution

This is the case of 16-bit addition

	45ECH	0100 0101 1110 1100
+	7723H	0111 0111 0010 0011
	<u>BD0FH</u>	<u>1011 1101 0000 1111</u>

The sum will be in BX, which is the destination register.

CF = 0 since there is no carry out from D15.

ZF = 0 since the destination is not zero.

OF = 1 since there is an overflow into the MSB D15.

Example 1.3

Assuming we are adding **signed** 8-bit numbers, how is the result of the following addition to be interpreted?

```
MOV AL, 125
ADD AL, 75
```

Solution

Both operands here are decimal numbers.

This program causes AL to be loaded with +125 and then +75 to be added to AL. The result is to be in AL.

Signed numbers consider the MSB as the sign bit. With 8 bits, the maximum range of signed numbers is -128 to +127.

When adding the given signed numbers,

+	125	0111 1101
+	75	0100 1011
+	<u>200</u>	<u>1100 1000</u>

There has been an overflow from D6 to D7 (i.e., to the MSB). In this case we find OF = 1. The sum is interpreted as a negative number with value of the sign bit = 1.

The setting of the OF flag indicates that the magnitude of the result has caused an inversion of the sign bit, indicating a mistake. This interpretation is done, because the maximum range of **signed positive** numbers is only +127 for 8-bit representation.

Incidentally the Sign Flag is also set (SF = 1).

Control Flags Besides the conditional flags, there are three control flags, namely, the **Trap flag**, **Direction flag** and **Interrupt flag**. The control flags have to be **deliberately set or reset according to the requirements of the program**. The **Trap flag (TF)** is **set to perform step by step execution, during debugging**. The **Interrupt flag (IF)** is **set to enable interrupts**. The **Direction**

flag (DF) is used in string operations. Details of the functioning of these flags will be discussed in later chapters.

Arithmetic Logic Unit It is the part of a computer that performs all arithmetic and logic computations. The ALU is the most important unit of the processor. Instructions that are fetched and decoded, are executed in the ALU. Thus the ALU has direct access to the general-purpose registers and flags.

1.3 | Bus Interface Unit

This unit is responsible for address calculations, pre-fetching instructions for the queue and sequencing instructions one by one. Let us examine each function.

1.3.1 | The Instruction Queue

Instructions are found in memory, from where they are fetched and decoded as and when they need to be executed. However in 8086, there is a queue which fetches instructions ahead of the execution time and places them in a 6-byte first-in-first-out (FIFO) queue. This pre-fetching is done when the buses are free i.e., not being used for the execution of the current instruction.

The advantage of pre-fetching is that when a particular instruction is to be executed, there is a good chance of finding it in the queue (which is on-chip), rather than having to go to memory to fetch it. However the queue will have to be emptied in the event of a branch instruction, i.e., when the next instruction to be executed is not the next one in the sequence. After this, the queue needs to be re-loaded from the new 'branch' address. This pre-fetching belongs to a class of ideas called pipelining, which means that both execution and fetching take place at the same time i.e., while the execution of one instruction is going on, the fetching of another one can be done. Pipelining greatly speeds up processing.

1.3.2 | Memory Segmentation

The 8086 has a 20-bit address bus, but let us remember that we have not seen any 20-bit registers so far. It is understandable that the general-purpose registers are used as 16-bit or 8-bit registers, as this is the way data is organized. However we see that all address registers also are 16-bit. An ingenious idea called segmentation has been used to take care of this situation. A segment is just an area in memory. Thus the total memory size can be considered to be divided into segments of various sizes. The idea of dividing memory this way is called segmentation.

In memory, data is stored as bytes. Each byte has a specific address. With 20 address lines, the memory that can be addressed is 2^{20} bytes. This will amount to 1,048,576 bytes (1 Megabyte). Thus the 8086 can access a physical memory with addresses ranging from 00000 to FFFFFH. A 20-bit physical address is to be placed on the address bus to access any location in memory.

Memory as a whole is considered to have four different types of segments labeled as data, code, stack and extra. One idea of segmentation is to keep data and code separately. Thus the code segment contains code only. Data is stored in the data segment, and the extra segment is just another type of data segment used in a special way. The stack segment is an area of memory which functions and is accessed differently from the other three segments.

1.3.2.1 | Segment Registers

Each of these segments is addressed by an address stored in corresponding segment registers. These registers are all 16-bit in size. Each register stores the base address of the corresponding

IMPORTANTE!!! Es pregunta de parcial/final:
Qué almacena el registro de segmento?

segment. A base address is the starting address of a segment. However, since the segment registers cannot store 20 bits (of the base address), they store only the upper 16 bits. The least significant nibble is implied to be zero.

How is a 20-bit physical address obtained if there are only 16-bit registers? The 20-bit address of a byte is called its **physical address**. However it is specified as a **logical address** which is in the form of two 16-bit numbers in the format **base address: offset**. See Fig 1.4, for a typical example.

We consider that a data byte is stored in a data segment, whose base address is 22220H. Then the data segment register (DS) will contain the number 2222H. The data at any location within this segment is referenced by an offset (displacement) with respect to the base address. Thus, if a data at a location has a logical address specified as 2222H : 0016H, the number 0016H is the offset or displacement with respect to the base address. For physical address calculation,




Figure 1.4 | Calculation of a physical address from the logical address for a data segment




Figure 1.5 | Segment registers and corresponding segments

the BIU appends the segment register content with four binary zeros (or one hex zero) on the right. The physical address is thus calculated as

$$\begin{array}{r} 22220\text{H} \\ + \quad 0016\text{H} \\ \hline 22236\text{H} \end{array}$$

Similar is the case for the other segments. All offsets are limited to 16 bits which means that the maximum size possible for a segment is only 2^{16} which is 65,536 bytes or 64K bytes. We can conclude that every memory reference by the 8086 will use one of the segment registers (i.e., DS, ES, SS, or CS) that is combined with an offset (usually given in the instruction) to determine the physical address being accessed. A data byte has the logical address of the form DS: offset.

1.3.2.2 | The Code Segment and the Instruction Pointer

The code segment is the area of memory where code alone is stored. The offsets within the code segment are referenced using the Instruction Pointer (IP), which is a 16-bit register. The IP sequences the instructions, and always points to the next instruction to be executed. Whenever an instruction byte has to be fetched from memory, the bus interface unit (BIU) performs the address calculation using the contents of CS register and the IP. This 20-bit address is then placed on the address bus and the instruction byte is fetched. Thus the logical address for an instruction byte is of the form CS : IP.

1.3.2.3 | The Stack Segment and the Stack Pointer

The stack is an area of memory that is used in a special way. Data is generally pushed in and popped out of the stack. The stack of the 8086 is a last-in-first-out (LIFO) stack, which means that the last data that was pushed in, is the first one that can be popped out. This also means that data can only be taken out or put in only from the top of the stack. However the ‘stack top’ changes as data is put in or taken out.

There is a 16-bit register called Stack Pointer which points to the top of the stack. The stack segment is like any other segment, and the upper 16 bits of its base address is available in the




Figure 1.6 | Stack and physical address of the top of stack

SS register. A stack address of the form 4466H : 0122H means that the SS register contains 4466H, the stack pointer (SP) contains the number 0122H, and the physical address of the top of the stack is 44660H + 0122H = 44782H.

The stack is used mostly to keep aside address and data temporarily, when a subprogram is called. This will be taken back when the subprogram ends. The 8086 has a stack which grows downwards (i.e., to lower memory addresses). This means that each push instruction causes the SP to be decremented (by two for 8086), and data is pushed deeper downwards into the stack segment. We will discuss the operation of the stack in more detail in Chapter 4.

Even though a stack is usually accessed only from the top, the 8086 has another register named BP (base pointer), which is used to reference data anywhere within the stack.

Thus the format SS : BP is also a logical address. This has certain interesting applications.

1.3.2.4 | The Data Segment and Extra Segment

Both these segments store data, but in certain special cases (string instructions), it may be necessary to list them separately. There is an Extra Segment (ES) register to store the upper 16 bits of the base address of the extra segment. The offset within the data segment is also termed as an 'effective address'. The effective address calculation depends on the mode of addressing. We will discuss this very soon. Table 1.1 shows the segment registers and the corresponding registers that can be used for offsets.

Example 1.4

The content of DS is 345BH. The amount of data that is to be stored in the data segment is 12K bytes. Where in memory, will this segment be located?

Solution

DS contains the number 345BH.

This number corresponds to the upper 4 hex digits of the starting (base address) of the data segment.

The base address of the segment is 345B0H.

The last address will have an offset of 12K from the base address.

1K = 1024 bytes.

12K = 12 × 1024 = 12288 = 3000H.

The last address of the segment = 345B0H + 3000H = 375B0H.

Hence the data segment occupies the memory addresses from 345B0H to 375B0H.

Table 1.1 | Segments and Corresponding Registers Used for Specifying Offsets

Segment	Offset Registers	Function
CS	IP	Address of the next instruction
DS	BX, DI, SI	Address of data
SS	SP, BP	Addresses in the stack
ES	BX, DI, SI	Address of destination data (for string instructions)

Example 1.5

The contents of the following segment registers are as given.

CS = 1111H, DS = 3333H, SS = 2526H.

IP = 1232H, SP = 1100H, offset in data segment = 0020H.

Calculate the corresponding physical addresses for the addressed byte in a) CS b) SS and c) DS.

Solution

- The base address of the code segment is 11110H. The address of the next instruction to be executed is referenced by CS and IP which is given by $11110H + 1232H = 12342H$.
- The current top of the stack is referenced by SS and SP. The base address of the stack segment is 25260H. The corresponding physical address is $25260H + 1100H = 26350H$.
- The data that needs to be accessed is given by DS and the offset. The base address of the data segment is 33330H. The physical address of this data is calculated as $33330H + 0020H = 33350H$.

Segmentation: General Concerns

- Segments begin at paragraph boundaries, where a paragraph is 16 bytes. i.e., **Base addresses are divisible by 16 (because the lowest nibble of all base addresses are to be zero)**.
- Segments can overlap with a new segment starting every 16 bytes. This also means that there can be more than one logical address for the same memory location (physical address).** For example, 0000 : 0100, 0001 : 00F0, and 0010 : 0000 all refer to physical address 00100H.
- We can have **many segments with the same base address, but the offsets within the segment will have to be taken care of to prevent overwriting of data.** For example, in string operations, we might use the same segment as the data and extra segments. This is achieved by loading the same number in the ES and DS registers.
- A segment can have a maximum size of 64K.**
- A program need not have all four segments, but it will have at least the code segment.**
- The addresses used in instructions are called logical addresses and are translated by the BIU to physical addresses.**

Advantages of Segmentation

- It allows all address registers to have the same size as the data registers (16-bit), while allowing the use of 20-bit physical addresses.
- All addresses in memory are **re-locatable**. This means that any program or data can be loaded in any address in memory. A re-locatable program is one which can be placed in any area of memory and executed without change. Data is also re-locatable. Changing the base address of the corresponding segment is the only action we need to perform, in order to re-locate. All addresses within the program are relative to the base address, as they are of the form Base address: offset. All processors in the x86 family have this kind of segmentation. The higher order processors also have **protection** added to this structure.

Memory Organization

We have seen that **data is stored in memory as bytes. Each byte location has an address. If a word (two bytes) is to be stored, two locations of consecutive addresses are needed.** Consider that the

ADDRESS	DATA
60088H	E6H
60089H	34H

Figure 1.7 | Little endian format

number 34E6H is to be stored. Then 34H is stored in one location and E6H is in the next location. Intel's processors use what is termed the **little endian format**, which means **the lower byte is stored in the lower address and the upper byte in the higher address**. Thus, if the above word has an address of 60088H, the storage is as shown in Fig 1.7. This is in contrast to certain other processors (like Motorola) which uses the **big endian format**. (The origin of the word 'endian' is supposed to be from Gulliver's travels, referring to a controversy about whether an egg should be broken at the little end or the big end.)

Since the data in this case is a word, two memory locations with two addresses have to be used to access it. The hardware connection between the processor and memory is designed in such a way that **both these locations are accessed and a word (16 bits) access is obtained, even though only one logical address is mentioned in the instruction**.

1.4 | Addressing Modes

For computations in assembly language, we need an opcode and operands. **Opcode stands for operation code** – i.e., **the code which specifies the operation to be performed**. Opcodes **operate on data, which are then called the operands**. In the **8086**, **an instruction can have either one, or two operands (not more)**. For example, an addition operation needs two operands, while a shift operation operates on a single operand. **The way in which operands are specified in an assembly language instruction is called its addressing mode**. Let us attempt to get a clear understanding of the different addressing modes allowed for this processor. We will use the MOV instruction for understanding these modes. This has the format

MOV destination, source

which causes the source data to be copied into the destination.

The basic assumptions in this context are

- i) the operands can be in registers, in memory, or may be in the instruction itself. However the 8086 does not have an addressing mode in which both operands are in memory locations,
- ii) in the case of two operands, one of them can be in memory, but the other will have to be placed in a register,
- iii) **data types should match** – i.e., the source and destination should both be either bytes or words.

1) Register Addressing

Here both the **source and destination are registers**. **No memory access is involved**. See the following instructions.

MOV AL, AH	;copy the content of AH to AL
MOV CH, BL	;copy the content of BL to CH

MOV SI, BX	;copy the content of BX to SI
MOV ES, AX	;copy the content of AX to ES

Note that the first two are **byte** operations, while the other two are word operations.

MOV AX, BL	;gives an error as AX is 16 bit and BL is 8 bit
MOV BL, AX	;gives an error for the same reasons

2) Immediate Addressing

In this mode, the source will be a constant data

MOV AL, 45H	;copy 45H to AL
MOV BX, 34E3H	;move the hex number 34E3H to BX
MOV CL, 'Q'	;move the ASCII value of Q to CL
MOV PRICE, 40	;move the hex number 40 to the memory location with label PRICE
MOV NUMS, 0FC6H	;move the hex number 0FC6H to the memory location NUMS

Segment registers are not allowed to use this mode of addressing.

MOV DS, 2300H	;gives an error as DS is a segment register
---------------	---

3) Direct Addressing

Here either the source or the destination will be a memory address.

MOV AX, [2345H]	;move the word in location 2345H into AX
MOV [1089H], AL	;the byte in AL is moved to location 1086H

It is to be remembered that the addresses in the instructions are offsets within the data segment (i.e., the logical address). We need to know the content of DS to calculate the physical address. The size of the registers indicate the size of the operand. Hence, in the first instruction, a data word is referred to; while in the second, a data byte is moved. The square brackets are necessary to indicate that the number is an address and not data. However, we may use labels for addresses, and re-write the above two instructions.

MOV AX, PRICE
MOV COST, AL

We must then ensure that the referred addresses have been defined earlier with these labels. We will see this aspect later.

4) Register Indirect Addressing

In this mode, the address of the data is held in a register. The register acts as a pointer to the data. The registers must be enclosed in square brackets to indicate that they function as pointers. We also use the term **effective address** for the address of the operand. For this mode of addressing, the address registers allowed are BX, SI and DI.

EA = {[BX]/[DI]/[SI]}	
MOV AL, [BX]	;move into AL the content of the location whose address is in BX
MOV [SI], CL	;move the content of CL to the address pointed by SI

MOV [DI], AX ;move the content of AX to the address pointed by DI

In the third instruction, AX contains two bytes. Hence the content of AL will be moved to the address pointed by DI. The content of AH will be moved to the address [DI + 1].

Example 1.6

Show the location of data in memory, after the execution of each of these instructions, if the content of registers are as given

DS = 1112H, AX = EE78H and BX = 3400H

- i) MOV [0422H], AL
- ii) MOV [0424H], AX
- iii) MOV [BX], AX

Solution

i) This is a case of direct addressing. The destination is a memory location, which is specified directly in the instruction. AL contains 78H

DS = 1112H

Data segment base address is 11120H

The offset in the instruction is 0422H

Physical address is 11120H +
 0422H
 11542H

After the instruction execution, the content of the addressed location is as shown

PHYSICAL ADDRESS	CONTENT
11542H	78H

ii) This is also a case of direct addressing. The physical address is calculated as

11120H +
 04224H
 11544H

However, since AX contains two bytes of data, the lower byte (AL) is stored in the address 11544H and the higher byte (AH) in the next address (recollect the 'little endian' concept)

PHYSICAL ADDRESS	CONTENT
11544H	78H
11545H	EEH

iii) This is a case of register indirect addressing. BX is a pointer to the address 3400H, which is an offset. The corresponding physical address is

11120H +
 3400H
 14520H

The content of AX is moved to this address. Hence the corresponding memory will be as shown

PHYSICAL ADDRESS	CONTENT
14520H	78H
14521H	EEH

5) Register Relative Addressing

In **relative addressing mode**, a **number or displacement is part of the effective address**.

$EA = \{[BX]/[DI]/[SI]/[BP]\} + 8\text{-bit or } 16\text{-bit displacement}$

MOV CL, 10[BX] ;move the content of the address specified by adding the content of BX and 10.

Thus the effective address is $[BX + 10]$. **Once the effective address is computed, the physical address is calculated as the sum of the segment base address and the effective address.** The displacement **can be a 16-bit signed/unsigned number or an 8-bit sign extended number.** However the displacement should not be so large as to make the effective address go beyond the range of the maximum size of a segment. The above instruction can also be written as

MOV CL, [BX + 10]

MOV CL, [BX] + 10 or

MOV CL, [BX][10] or

MOV CL, PRICE [BX]

In the last case, PRICE has to be defined earlier as a displacement of 10.

6) Based Indexed Mode

In this mode, an **index register and a base register together carry the effective address**. The content of these two registers are added and called the effective address.

MOV AL, [BX][SI] ;move the content of the effective address pointed by [BX] and [SI] into AL

The **effective address is obtained by adding the content of BX and SI**. Since the destination register is an 8-bit register, **this is a byte operation**. The following is a **word** operation as the source register CX is 16-bit in size.

MOV [BX][DI], CX ;move the content of CX to the effective address pointed by [BX] and [SI]

7) Relative Based Indexed Mode

This is the case when the **'effective address' is specified with a base register, an index register as well as a displacement**. The **'effective address' is the sum of the two registers and the displacement**. For example, the following use the relative-based indexed mode of addressing

MOV DL, 5[BX][DI] ;EA = 5 + BX + DI

MOV 5[BP][SI], AX ;EA = 5 + BP + SI

MOV CL, COST[BX][SI] ;EA = COST + BX + SI

;COST has to be defined as a displacement earlier

Example 1.7

Find the address of physical memory for the following instructions if the content of the required registers are as given below

SS = 2344H, DS = 4022H, BX = 0200H, BP = 1402H, SI = 4442H

i) MOV CL, 1234H[SI]

ii) MOV AL, 5[SI][BP]

Solution

i) MOV CL, 1234H[SI]

This is a case of register relative addressing. The effective address is obtained from the instruction, to be the sum of the displacement and SI

$$\begin{array}{r} \text{Effective address} = 1234\text{H} \quad + \\ \quad \quad \quad \quad \quad 4442\text{H} \\ \hline \quad \quad \quad \quad \quad 5676\text{H} \end{array}$$

The segment base address is obtained from DS to be 40220H

The physical address is the sum of the segment base address and the effective address. i.e.,

$$\begin{array}{r} 40220\text{H} \quad + \\ \quad \quad \quad 5676\text{H} \\ \hline \quad \quad \quad 45896\text{H} \end{array}$$

ii) MOV AL, 5[SI][BP]

This is a case of relative based indexed mode. The effective address is calculated as the sum of the displacement and the contents of the registers SI and BP.

$$\begin{array}{r} \text{Effective address} = 0005\text{H} \quad + \\ \quad \quad \quad \quad \quad 4442\text{H} \\ \quad \quad \quad \quad \quad 1402\text{H} \\ \hline \quad \quad \quad \quad \quad 5849\text{H} \end{array}$$

The physical address is the sum of the segment base address and the effective address. In this case, as **BP** is one of the address registers, the segment referred is the **stack segment**.

Physical address is 23440H (base address of the stack segment)

$$\begin{array}{r} + \quad 5849\text{H} \\ \hline \quad 28\text{C}89\text{H} \end{array}$$

Segment Override

Table 1.2 shows the default segments to be used in the various memory-based addressing modes. However, **this can be overridden by a segment override prefix**. i.e., We can use segments other than those specified in the table. The **format** of this **override prefix** is shown below. MOV AL, ES:[BX] overrides the fact that an instruction of this sort implies that the data segment is to be used. Here, it says that ES is to be used instead. Hence, physical address calculations use the content of the CS register as the base address. Other examples are

MOV DS : [BP + 7], BL	;DS to be used instead of SS
MOV AX, CS : [BX]	;CS to be used instead of DS

Table 1.2 | Effective Address and Referred Segments for Various Memory Based Addressing Modes

SI No.	Addressing mode	Effective address	Segment
1	Direct	offset	DS
2	Register indirect	[BX] [SI] [DI]	DS DS DS
3	Register relative	Disp + [BX] Disp + [SI] Disp + [DI] Disp + [BP]	DS DS DS SS
6	Based Indexed	[BX] + [SI] [BX] + [DI] [BP] + [SI] [BP] + [DI]	DS DS SS SS
7	Relative Based Indexed	Disp + [BX][SI] Disp + [BX][DI] Disp + [BP][SI] Disp + [BP][DI]	DS DS SS SS

Note that BP cannot be used without a displacement. See it being used in 'register relative' mode but not in register indirect mode.

KEY POINTS OF THIS CHAPTER

- The architecture of 8086 is representative of the architecture of all members of the x86 family.
- The data bus width of 8086 is 16, and internal registers also have the same bit size.
- The internal block diagram shows two units named the 'execution unit' and the 'bus interface unit' with well-defined and separate functions.
- The BIU and EU communicate through the internal bus. Both have a common timing and control unit.
- The execution unit takes care of computations, and thus it contains the necessary registers and flags.
- The conditional flags are set or reset according to the result of certain arithmetic/logic operations.
- The 8086 has 20 address lines, which means it can access 1 MByte of memory.
- Memory is segmented into four types of segments.
- Code and data are kept in separate segments.
- The upper 16 bits of the base address of any segment are available in the corresponding segment registers.
- The stack segment is a type of segment that grows downwards as data is pushed into it.
- The BIU converts the logical addresses to physical addresses.
- Segmentation causes data and code to be re-locatable.

- The addressing mode dictates where operands are available.
- 8086 follows the 'little endian' scheme of data storage in memory.

QUESTIONS

1. Name the 8-bit registers of 8086.
2. What is meant by saying that the 8086 is a 16-bit processor?
3. Which are the first and last memory addresses that an 8086 can address?
4. Which unit (EU or BIU) is responsible for performing arithmetic calculations?
5. How will it be helpful for a processor if it has a large number of internal registers?
6. What are the minus points of the presence of too many scratchpad registers?
7. Which are the registers catering to string instructions?
8. What is meant by a LIFO stack?
9. What is the difference in operation between the carry flag and the overflow flag?
10. Which register is frequently used as a counter?
11. Where in memory is the program code located?
12. BP is a register which refers to the extra segment. True or False?
13. BP cannot be used without a displacement. True or False?
14. What is meant by the term 'program relocation'?
15. What happens to the pre-fetch queue when a branch instruction is encountered?
16. What is the format of a logical address?
17. If the physical address of a byte is 12345H, suggest three logical addresses that it can have.
18. Which is the default segment when IP is used?
19. Can we have segments of size less than 64K?
20. Write an example instruction that contains a segment override prefix.

EXERCISE

1. Find the content of the destination and the status of the carry flag (CF) and zero flag (ZF) after the execution of the following set of instructions.
 - a) `MOV AL, 09H`
`ADD AL, CFH`
 - b) `MOV BX, 0876H`
`MOV AX, 0034H`
`ADD BX, AX`
 - c) `MOV AL, 0FFH`
`ADD AL, 01`
2. Add the numbers 100 and 89 using the following instructions.
`MOV AL, 100`
`ADD AL, 89`
Which of the conditional flags are set and why?

3. Show the binary number that will be available in AL after the execution of each of the following instructions.
 - a) MOV AL, 'C'
 - b) MOV AL, 25
 - c) MOV AL, 34H
4. Find the physical address of the top of the stack if SS = 0777H, and SP = 1234H.
5. If IP contains the number 0034H, and CS = 5555H, where will the next instruction be fetched from?
6. If CS contains the number 0FC2H, and the volume of code is 2400 bytes, which is the last address in this code segment?
7. Write the addressing mode of the following instructions.
 - a) MOV [4560H], AX
 - b) ADD BL, 89H
 - c) ADD BX, [DI]
 - d) MOV 6[BP][DI], AL
 - e) MOV CX, [BP + 9]
 - f) ADD CL, BL
8. Find the physical address of the memory locations referred to in the following instructions if DS = 0223H, DI = 0CCCH, SI = 1234H.
 - a) MOV [DI], AL
 - b) MOV [SI][56H], BL
9. Which segments will be accessed in the following cases?
 - a) MOV [6][BX], CX
 - b) MOV [BP] + 4, CL
 - c) PUSH AX
 - d) MOV CS : [BX], DL
 - e) MOV AL, ES : [45H]
10. The memory addresses starting from the logical address 0124H : 0056H store the numbers 0978H, CE45H, 45H and 7809H in consecutive locations. Draw a diagram showing the physical address and the corresponding data stored.

2 PROGRAMMING CONCEPTS - I



In this chapter, you will learn

- The working principle of an assembler.
- The various assemblers available for x86.
- The features and use of a popular assembler.
- To write programs using simplified memory models.
- To write programs using the full segment definition.
- To understand the techniques of instruction design for 8086.

Assembly Language Programming

In Chapter 0, we have seen why assembly language is very important. To code efficiently in assembly language for a particular processor, the prerequisites are a good knowledge of the internal architecture of the processor and addressing modes, which means that understanding the contents of Chapter 1 is absolutely necessary.

When we want to write and execute an ALP, we first type it using an editor. Then we assemble, link and run it. There are many assembly language program development tools available for each of these actions. We shall examine them one by one.

2.1 | The Assembly Process

The way an assembler is designed depends heavily on the internal organization of the processor for which it is used. Architectural features such as memory word size, number formats, internal character codes, index registers, and general-purpose registers affect the way assemblers are written and the way an assembler handles instructions and directives.

In assembly language programming, coding is done in symbolic language (called mnemonics). The word *mnemonic* comes from a Greek word meaning *pertaining to memory*; it is a memory aid. The mnemonic specifies the operation to be done. We need an assembler which translates this symbolic language to machine code which the processor will understand. As such, we guess that this 'translation' is indeed the main task of the assembler.

Let us first get an idea of the assembly process. An interesting definition for an assembler is that it is a translator that translates source instructions (in symbolic language) into target instructions (in machine language), on a **one to one** basis. The point is that each source instruction is translated to exactly one target instruction. This is indeed the main task of the

translator. However, an assembler is more than just a translator. It is an important program development tool, the quality of which can affect the ease of programming. Since assembly language programming is more prone to errors (than HLLs); the assembler should give appropriate error messages to guide the programmer at every step during programming.

2.1.1 | Features of Assemblers

When code is written in assembly language, it is easy if we can use labels (names) for memory locations. All constants and variables are given names so that they can be referred by name. Managing and utilizing these labels (called symbols) efficiently reflects on the quality of an assembler.

Assemblers that support macros are called Macro Assemblers. A macro is a name given to a sequence of instruction lines. Once a macro is defined, its name may be used in place of this set of lines. In the assembly process, each macro is replaced by the text lines. Thus a macro seems to the programmer to be like a mnemonic, and he can use it as frequently as he wants, without incurring any of the overheads that a procedure or function may cause. The use of macros can help to make programs structured.

2.1.2 | Instructions and Directives

An assembly language program contains pseudo-instructions (directives) along with instructions. Instructions get translated to machine codes, but directives do not. i.e., instructions produce executable code, but directives do not. They direct the assembler to perform in certain ways that we would like it to. The directives perform as 'help' for the assembler for deciding certain other aspects of the assembly process. Thus the assembler takes our **source code** (which is our program containing instructions and directives) and converts it into **object code** which contains the machine code. The object file is the one that will be loaded into memory and executed after linking it with other necessary files.

2.1.3 | The Forward Reference Problem

A typical line in a program could be

```
BEGIN: MOV AX, COST           ;move the content of the location COST to AX
```

In this example, BEGIN is a label, which corresponds to the address at which this line of code is stored. MOV is a mnemonic. This instruction means that the content of COST is to be moved to register AX. COST is a label (for a memory location) that acts as one operand in this two operand instruction. A label when defined is called a symbol. A symbol can only be defined once but it can be used many times. There will be a memory address corresponding to all symbols. In this instruction, the other operand is AX. The text coming after ';' is a comment.

All current assemblers allow the use of labels, rather than numeric addresses of memory locations. However, finally, the assembler must translate these labels to memory addresses. As such, an important issue that an assembler will have to sort out is what is called the 'forward reference problem'.

See the code below

```
BEGIN: MOV AX, BX
        JMP AGAIN
        -----
        -----
        -
        -
        -
AGAIN: -----
```

As the assembler reads the second line of this program, an undefined label AGAIN is encountered. The assembler has no knowledge of it as yet, but the value of the symbol AGAIN is needed before it can be defined. This is also called the future symbol problem or the problem of unresolved references. This problem is solved in various ways depending on the type of the assembler.

2.1.4 | Two Pass Assemblers

There are many types of assemblers – one pass, two pass and multi pass assemblers. Each pass optimizes the assembly process. Two pass assemblers are the commonest, so let us discuss those. A pass implies ‘a reading’ of the code from beginning to end. A two pass assembler does two passes over the source file (the second pass will be over a file generated in the first pass). In the first pass it looks for label definitions and inserts them in the symbol table after assigning them addresses. Modern assemblers keep more information in their symbol table, which allows them to resolve addresses in a single pass. Known addresses (backward references) are immediately resolved. Unknown addresses (forward references) are “back-filled”, once they are resolved.

Memory is allocated sequentially, and a ‘location counter’ is incremented at each step. At the end of this pass, the symbol table should contain definitions of all the labels used in the program. This means that during the first pass, the only duty of the assembler is assigning memory addresses to labels.

In pass 2, the actual translation of assembly code to machine code is done. Errors are also reported after this. The assembly process produces a ‘re-locatable’ object file. This object file can be loaded anywhere in memory, when the program is to be executed. A ‘loader’ is a program which does this. Please note that loading can be done anywhere in memory and that it is not necessary to load it at the same location each time the program execution is done.

2.2 | Assemblers for x86

Assembly language is specific to a particular processor. As such, we should concern ourselves with assemblers for the x86 family of processors, which includes the whole set from 8086 to Pentium. A large set of assemblers is available, but the code running on one assembler will not run on another. Another matter to be taken care of is the OS under which it can run. A86 was the first assembler to appear on the scene. It could run only under DOS. Now there are assemblers which run under DOS, Windows and Linux. The names of a few popular assemblers are NASM, FASM, MASM, TASM and HLA. FASM and NASM can run under all the three operating systems mentioned.

Two popular assemblers are TASM and MASM. TASM is Borland's Turbo Assembler, while MASM is Microsoft's Macro Assembler. It is found that code written on TASM runs on MASM, though the reverse need not be true. Both of them run under DOS and Windows, which means they can use 16-bit and 32-bit coding. Both of them are classified as high level assemblers, which means that we can use high level control structures (such as IF, WHILE, REPEAT...UNTIL) that make programming easy. HLA is another assembler which uses high level control structures, high level data types and high level procedure calls.

2.2.1 | Why MASM?

For programs in this book, we will choose MASM because it is one of the best assemblers and also because of its popularity. A popular assembler gives us the chance to get easy access to references and material. Also MASM has many constructs and features which makes ALP easy and fun for beginners. Another reason is that the documentation available for it is very good. MASM was a commercial product for a very long time, and is a sort of industry standard. Microsoft has written considerable documentation for this assembler and many third parties have written assembly language reference manuals for MASM. The versions of MASM 6.0 and above have a lot more features (aimed at simplification in writing code) than previous versions.

Incidentally, MASM is a proprietary product of Microsoft. We will now use the 16-bit version which runs under DOS, because we can use DOS and BIOS interrupts in this version. There is a 32-bit version which runs under Windows. This is freely available on the net (the 32-bit version does not allow the use of DOS and BIOS interrupts). In later chapters, we will use MASM32 for the 32-bit x86 processors. This approach gives a very good understanding of assembly language in general. The DOS version used for programs in this book is MASM 6.14. These programs will run in all versions of MASM from 6.0 upwards.

2.2.2 | Assembly Language Programming

Now that we have understood the assembly process and have decided the assembler to use, let us get into the nitty-gritty of programming. The steps involved are:

- i) Write the code with the help of an editor. We can have an editor within MASM or can use notepad or wordpad. Save it as filename.asm in the BIN directory of MASM. Note that instructions, directives and labels are **not case sensitive**.
- ii) Open the DOS command window and go to the directory containing MASM. From this, go to the BIN directory.
- iii) Assemble the code. Syntax errors (if any) will be reported at this stage. This will generate the object file with the name filename.obj. Other files are also generated and we can opt to view them (list file, configuration file and so on). Details of these are given in Appendix B.
- iv) Link the object file. The linking involves combining separately assembled modules into one executable file (filename.exe).
- v) Run the exe file. This will cause the program to be loaded in memory and executed.
- vi) The results of execution will be in memory or in registers. If we want to view it on our display unit, we will have to use special codes for it (this will be discussed later). Otherwise, we

can examine the memory and register contents using the debugger. The debugger has all the commands needed to do this, and the full set of commands is detailed in Appendix B. DOS commands are also listed here.

2.3 | Memory Models

To write programs, we know we have to define segments, which means that the segment registers must be initialized. This will take us to the conventional **full segment model**. However, versions of MASM 6.0 and above have incorporated certain shortcuts which help to make programming simple. These are called the dot models. They have the format:

.MODEL MODEL NAME

The different models tell the assembler how to use segments and to provide sufficient space for the object code. The simplest of these are the tiny model and the small model. We will start with these models. The tiny model is used when code, data and stack will all fit into one segment with a maximum size of 64 K. The small model can have one data segment and one code segment each of which has a maximum size of 64 K.

2.3.1 | The Tiny Model

Let us now write a program using the tiny model.

Example 2.1

.MODEL TINY	;choose single segment model
.CODE	;start of code segment
.STARTUP	;start of program
MOV AL, 67H	;move 67H to AL
MOV BL, 45H	;move 45H to BL
ADD AL, BL	;add BL to AL
MOV DL, AL	;copy AL to DL
.EXIT	;exit to DOS
END	;program end

Let us enter this code using an editor and save it as `tinym.asm` in the BIN directory of MASM6.14. Open the DOS command window. Go to the directory MASM 6.14/BIN. In the BIN directory, assemble and link the file using the command `ml tinym.asm`. Example 2.2 shows the DOS command window corresponding to this. The `ml` command causes assembling and linking to be done at a go. Now to get the list file, type the command `ml/FI tinym.asm`. The list file will be saved in the BIN directory as `tinym.lst` (more commands and options are available in Appendix B). Earlier versions of MASM required two steps for the same actions. To assemble, the command was `masm filename.asm`. Linking required the command `link filename.obj`. The list file was obtained by using `masm/l filename.asm`. Examples 2.2 shows the DOS command window as we do assembling and linking. Example 2.3 shows the list file.

Note If you think this part is taxing, rush through these topics fast and come back to it after you have learned more aspects of programming and are ready to do programming using this assembler.

Example 2.2

```
C:\masm6.14\BIN>ml tinym.asm
Microsoft (R) Macro Assembler Version 6.14.8444
Copyright (C) Microsoft Corp 1981–1997. All rights reserved.

Assembling: tinym.asm

Microsoft (R) Segmented Executable Linker Version 5.60.339 Dec 5 1994
Copyright (C) Microsoft Corp 1984–1993. All rights reserved.

Object Modules [.obj]: tinym.obj/t
Run File [tinym.com]: "tinym.com"
List File [nul.map]: NUL
Libraries [.lib]:
Definitions File [nul.def]:
```

Example 2.3

```
Microsoft (R) Macro Assembler Version 6.14.8444
tinym.asm

                                .MODEL TINY          ;choose single segment model
0000                            .CODE                ;start of code segment
                                .STARTUP             ;start of program

0100 B0 67                      MOV AL, 67H          ;move 67H to AL
0102 B3 45                      MOV BL, 45H          ;move 45H to BL
0104 02 C3                      ADD AL, BL           ;add BL to AL
0106 8A D0                      MOV DL, AL           ;copy AL to DL

                                .EXIT                ;exit to DOS
                                END                  ;assembler to stop reading
```

Example 2.3 shows the listing corresponding to our tinym program. (Only a part of the list file is shown here. See Appendix B for the additional details.) A lot of information can be obtained just by looking at the list file. The list file is one of the outputs obtained in the assembly process. Examine Example 2.3.

We see first the details of the version of MASM that is used. The complete assembly language program along with the comments that we had typed in, is seen next. On the left hand side, we see the offsets within the code segment in which the code is saved. These offsets have been generated by the assembler. We also find the opcodes corresponding to each instruction. For example in the first line:

```
0100 B0 67                      MOV AL, 67H          ;move 67H to AL
```

0100H is the offset in the code segment where the first instruction is stored, and B0 67 is the opcode of MOV AL, 67H. If we add a directive .listall, the list file is changed as shown in Example 2.4.

Example 2.4

```

                                .MODEL TINY
0000                            .CODE
                                .STARTUP
0100        *@Startup:
                                .listall

0100  B0 67                    MOV AL, 67H           ;move 67H to AL
0102  B3 45                    MOV BL, 45H           ;move 45H to BL
0104  02 C3                    ADD AL, BL            ;add BL to AL
0106  8A D0                    MOV DL, AL            ;copy AL to DL

                                .EXIT                ;return to DOS
0108  B4 4C        *      mov    ah, 04Ch
010A  CD 21        *      int     021h
                                END

```

.Listall has listed the instructions corresponding to .EXIT

.EXIT is a shortcut which is actually transformed to two instructions:

```

MOV AH, 4CH
INT 21H

```

These two instructions terminate program execution and bring control back to DOS. END is a directive. It tells the assembler to stop reading, as there are no instructions beyond this. The .STARTUP command is available only in versions of MASM 6.0 and above. It is a shortcut, whereby segment registers get initialized automatically, which means that programmers need not write instructions for that.

Now let us use the debugger to get more information about these aspects. The debugger allows us to view memory and registers, along with various other functions (see Appendix B). To enter the debugger, type **debug tinym.com**. We get an underscore as the prompt. On typing 'r', we can see the contents of the registers, before execution of the program. Now type 'u', which is the command for unassembling or disassembling. Disassembling is the reverse of assembling. It gets the mnemonics back from the machine codes. Example 2.5 gives the result of these two commands.

Example 2.5

```

C:\masm6.14\BIN>debug tinym.com
-r
AX=0000 BX=0000 CX=010C DX=0000 SP=0000 BP=0000 SI=0000
DI=0000 DS=13AD ES=13AD SS=13BD CS=13BD IP=0100 NV
UP EI PL NZ NA PO NC
13BD:0100 B067          MOV     AL, 67

-u
13BD:0100 B067          MOV     AL, 67
13BD:0102 B345          MOV     BL, 45
13BD:0104 02C3          ADD     AL, BL
13BD:0106 8AD0          MOV     DL, AL

```

13BD:0108 B44C	MOV	AH, 4C
13BD:010A CD21	INT	21

The r command gives us the content of the registers, before the execution of the instruction MOV AL, 67H. We find that the content of the CS register is CS = 13BD. On typing the ‘u’ command, we find that the logical address of the first instruction is 13BD:0100. The value of CS need not be the same every time the program is loaded into memory. This is because the object file generated by the assembler is **re-locatable**.

2.3.2 | COM and EXE Files

For this program, there is only one segment, which is the code segment. The run file generated for the tiny model is a Command (.com) file, rather than an Executable (.exe) file. This can be seen in Example 2.2. in the line:

Run File [tinym.com]: “tinym.com”

Later, it will be seen that run files will be obtained as ‘executable’ with .exe as the extension. However, the tiny model always assembles to be a com file. This is a special case, when we need a compacted form of data and code, all fitting into just 64K size for memory efficiency. Other models which we will see later all assemble to .exe files. Quantifying the features of COM files:

- i) Size is limited to 64K.
- ii) Only one segment, which is the code segment.
- iii) Data is defined in this code segment.
- iv) Code starts at offset 0100H, just after the PSP (program prefix segment) of DOS.
- v) Smaller file compared to exe files, because it does not have the 512-byte header block.

From Examples 2.2 to 2.5, we have seen the tactics of running and analyzing a single segment assembly language program running on MASM. Now, see the listing corresponding to another program which uses the tiny model.

Example 2.6

	.MODEL TINY	;choose single segment model
0000	.CODE	;start of code segment
	.STARTUP	;start of program
0100 8B C8	MOV CX, AX	;copy contents of AX to CX
0102 8B C2	MOV AX, DX	;copy contents of DX to AX
0104 8A C3	MOV AL, BL	;copy contents of BL to AL
0106 8A E1	MOV AH, CL	;copy contents of CL to AH
0108 8A F2	MOV DH, DL	;copy contents of DL to DH
010A 8E DA	MOV DS, DX	;copy contents of DX to DS
010C 8E D3	MOV SS, BX	;copy contents of BX to SS
	.EXIT	;exit to DOS
	END	;end of program

Example 2.6 shows a sequence of instructions that copy various data between 16- and 8-bit registers. The act of moving data from one register to another changes only the destination register, never the source.

Example 2.7

```

0000          .MODEL TINY          ;choose single segment model
              .CODE                ;start of code segment
              .STARTUP             ;start of program
0100  B8 0000  MOV  AX, 0          ;place 0000 into AX
0103  B0 43    MOV  AL, 'C'        ;place 'C' into AL
0105  B9 5674  MOV  CX, 05674H     ;place 5674H into CX
0108  B2 2D    MOV  DL, 45         ;place 45 (decimal) into DL
010A  B6 45    MOV  DH, 45H        ;place 45H into DH
              .EXIT                ;exit to DOS
              END                  ;end of program

```

Example 2.7 is a listing and it shows various assembly language instructions that use immediate addressing. Only 8-bit data can be placed into an 8-bit register. AL, DL and DH are 8-bit registers. 16-bit data can be copied to 16-bit registers. Though in the editor, the second instruction was written as MOV AL, 'C', the assembler translates ASCII 'C' to its hex equivalent 43H. All data written in memory will be in the hex format (i.e., the compressed form of binary representation). Thus the decimal 45 is found to 2DH in the listing.

Note A data (byte, word or longer) starting with the hex characters A, B, C, D, E or F must be preceded by a 0. Otherwise the assembler will give an error i.e., MOV AL, EFH gives an assembly error. It must be re-written as MOV AL, 0EFH. Similarly MOV BX, C456H is wrong. Rewrite it as MOV BX, 0C456H.

2.3.3 | Definition of Data Types

Before we go on to using a two segment memory model, we need to understand a few directives of the assembler which define and describe different kinds of data. Data which is used in a program can be bytes or words or of greater length. We have to define data and assign labels to their corresponding addresses. Assigning labels eases the programmer's burden as he does not have to concern himself with numerical values. The location counter in the assembler keeps on incrementing as labels are encountered.

Defining data implies allocating space for data. Data is defined accordingly using directives. Thus DB defines data byte, while DW defines data word. This is the traditional way of defining data. The newer versions of MASM (6.0 and higher) allow the use of the directive BYTE for DB, and WORD for DW. Table 2.1 shows the definitions of data of different lengths.

Table 2.1 | Data Definitions Used by MASM

Definition	Traditional directive	New directive
Byte	DB	BYTE
Word	DW	WORD
Double word	DD	DWORD
Quad word	DQ	QWORD
Ten bytes	DT	TBYTE

Example 2.8

```

                                .MODEL TINY
0000                                .CODE
0000  50                        NUM1 DB    50H           ;place 50H into NUM1
0001  210A                    NUM2 DW    210AH          ;place 210AH into NUM2
0003  0789                    NUM3 DW    0789H          ;place 0789H into NUM3

                                .STARTUP

0100  A0 0000 R                MOV AL, NUM1            ;move NUM1 to AL
0103  8B 1E 0001 R            MOV BX, NUM2            ;move NUM2 into BX
0107  8B C8                    MOV CX, AX              ;copy AX to CX
0109  8B 1E 0003 R            MOV BX, NUM3            ;move NUM3 into BX
                                .EXIT
                                END

```

Example 2.8 is a listing, which shows data being placed in the code segment itself. In the tiny model, we can have data and code in the same segment, with the restriction that the size of the segment should not exceed 64Kbytes. NUM1, NUM2 and NUM3 are locations which store data. NUM1 is a byte location, while NUM2 and NUM3 are word locations.

2.3.4 | The Small Model

Now we are in a position to write a program using the small model, which has one data segment and one code segment. Example 2.9 shows the same program as in Example 2.8, except that a separate data segment is used to store the data.

Example 2.9a

```

                                .MODEL SMALL           ;choose small model
0000                                .DATA               ;start data segment

0000  50                        NUM1 DB    50H
0001  210A                    NUM2 DW    210AH
0003  0789                    NUM3 DW    0789H

0000                                .CODE               ;start code segment
                                .STARTUP                ;start of program

0017  A0 0000 R                MOV AL, NUM1
001A  8B 1E 0001 R            MOV BX, NUM2
001E  8B C8                    MOV CX, AX
0020  8B 1E 0003 R            MOV BX, NUM3
                                .EXIT
                                END

```

The salient points of this program listing are:

- i) The word R on code lines with addresses says that the addresses are **re-locatable**. This can be observed in the listing of Example 2.8 as well.
- ii) In the data segment, data locations have been given labels. Data will be stored at offsets as shown in Fig 2.1 (recollect the **little endian** concept).

Now let us see the same program listing with the **.LISTALL** directive inserted before the startup directive. This will show the conversion of **.STARTUP** and **.EXIT** to a number of instructions of 8086. What is to be noted is that the startup directive corresponds to instructions necessary to initialize the segment registers (except CS, which is done by the OS alone). Compare this with Example 2.12. However, if all this confuses you, leave it aside for the time being and come back to it when you become more proficient in assembly language programming.

Example 2.9b

```

                                .MODEL SMALL           ;choose small model
0000                            .DATA                 ;start data segment

0000 50                        NUM1 DB    50H
0001 210A                      NUM2 DW    210AH
0003 0789                      NUM3 DW    0789H

                                .LISTALL
0000                            .CODE                 ;start code segment
                                .STARTUP               ;start of program

0000                            *@Startup:
0000 BA ---- R                *   mov     dx,  DGROUP
0003 8E DA                    *   mov     ds,  dx
0005 8C D3                    *   mov     bx,  ss
0007 2B DA                    *   sub     bx,  dx
0009 D1 E3                    *   shl     bx,  001h
000B D1 E3                    *   shl     bx,  001h
000D D1 E3                    *   shl     bx,  001h
000F D1 E3                    *   shl     bx,  001h
0011 FA                      *   cli
0012 8E D2                    *   mov     ss,  dx
0014 03 E3                    *   add     sp,  bx
0016 FB                      *   sti

0017 A0 0000 R                MOV AL, NUM1
001A 8B 1E 0001 R            MOV BX, NUM2
001E 8B C8                  MOV CX, AX
0020 8B 1E 0003 R            MOV BX, NUM3

                                .EXIT
0024 B4 4C                  *   mov     ah,  04Ch
0026 CD 21                  *   int     021h
                                END

```

LABELS	OFFSET	DATA(H)
NUM3+1	0004	07
NUM3	0003	89
NUM2+1	0002	21
NUM2	0001	0A
NUM1	0000	50

Figure 2.1 | Data segment with labels and offsets corresponding to Example 2.9

2.3.5 | The DUP Directive

This directive is used to replicate a given number of characters. For example, we may want to fill up a number of locations in the data segment with the same word or byte.

NUMS DB 10 DUP(0) fills up the 10-byte locations starting with the label NUMS.

STARS DB 5 DUP('#') fills up 5-byte locations starting at location STARS, with the ASCII value of the character #.

BLANK DB 10 DUP(?) reserves 10-byte spaces starting from location with the label BLANK, but these are not initialized, which implies that whatever data is there will remain. We can overwrite these locations with new data when we choose to. WRDS DW 4 DUP(FF0FH) fills up 4 word locations with the word FF0FH.

2.3.6 | The EQU Directive

This directive allows us to equate names to constants. The assembler just replaces the names by the values mentioned. Examples are:

```
TEMP EQU 34
PRICE EQU 199
```

Example 2.10

```
                                .MODEL SMALL
0000                            .DATA
0000 45                        ONE DB 45H
0001 00                        TWO DB ?
0000                            .CODE
                                .STARTUP
= 0025                        FACTR EQU 25H
0017 A0 0000 R                MOV AL, ONE           ;copy ONE to AL
001A 04 25                    ADD AL, FACTR         ;add FACTR to AL
001C A2 0001 R                MOV TWO, AL          ;move the sum to TWO
                                .EXIT
                                END
```

Example 2.10 shows the use of the EQU directive. From this list file, we see that the label FACTR is replaced by 25H when the program is being executed.

2.3.7 | The ORG Directive

ORG is a directive which means ‘origin’. In the context of assembly language programming, it can change the location of storage of data or code in memory i.e., the programmer gets the freedom to decide the offset of data/code when it is stored. The format is ORG d16, where d16 means a 16-bit displacement. Let us try out the usage of this directive.

Example 2.11a

```
.MODEL SMALL
.DATA

ORG 0010H                ;origin to be at 0010H
NUMS DB 20H, 40H
ORG 0030H                ;origin to be at 0030H
NUMS1 DW 8907H, 0FDH
.CODE
.STARTUP
ORG 000AH                ;org is a directive
MOV AL, NUMS              ;this is an instruction
MOV BL, NUMS+1            ;this is an instruction
ORG 0020H                ;org is a directive
MOV DX, NUMS1             ;this is an instruction

.EXIT
END
```

Example 2.11b

	.MODEL SMALL
0000	.DATA
	ORG 0010H
0010 20 40	NUMS DB 20H,40H
	ORG 0030H
0030 8907 00FD	NUMS1 DW 8907H,0FDH
0000	.CODE
	.STARTUP
	ORG 000AH
000A A0 0010 R	MOV AL,NUMS
000D 8A 1E 0011 R	MOV BL,NUMS+1
	ORG 0020H
0020 8B 16 0030 R	MOV DX,NUMS1
	.EXIT
	END

Example 2.11b is the list file corresponding to the program in Example 2.11a. Compare it to any of the list files seen earlier. Now, neither the data nor the code segments start at the default

offset of 0000. The data and code coming after the ORG statements have offsets as specified in the ORG directive. It has thus left out a lot of space in both the data and code segments.

2.3.8 | Other Models

These are other models available to be used according to our requirements. Table 2.2 will give an idea of those models and their specifications. There is another model called the flat model, which uses 32-bit addressing and un-segmented memory and is used in the protected mode. From the above, it is clear that all the models are not for 8086 alone, but are the features of MASM which caters to the whole x86 family.

2.3.9 | Full Segment Definition

Now let us see the traditional model of MASM. In the simplified memory model, it is left to the loader software to initialize the segment registers. In the traditional model, we use directives to define segments, and instructions to initialize the segment registers. This is called the **Full Segment** definition. The CS and SS registers are automatically initialized by the loader, but the DS and ES (if the extra segment is being used) will have to be initialized by the programmer. Let us re-write Example 2.9 using the full segment definition.

Example 2.12

```

DAT                                SEGMENT                                ;start the data segment
NUM1          DB          50H
NUM2          DW          210AH
NUM3          DW          0789H
DAT                                ENDS                                ;end the data segment

COD                                SEGMENT                                ;start the code segment
                                ASSUME          CS:COD,DS:DAT
                                MOV  AX, DAT                                ;move DAT to AX
                                MOV  DS, AX                                ;copy AX to DS
                                MOV  AL, NUM1
                                MOV  BX, NUM2
                                MOV  CX, AX
                                MOV  BX, NUM3
                                MOV  AH, 4CH
                                INT  21H
COD                                ENDS                                ;end the code segment
                                END                                ;end of program

```

The salient features of this model are:

- i) The data segment has been given a name DAT. The data within the segment is enclosed between the SEGMENT and ENDS directives which are similar to parentheses for a segment.
- ii) Similarly, the code segment is named COD and the contents of this segment also have been enclosed between the same directives (segments can have any name).
- iii) The DS register has been initialized by the first two instructions. DS cannot use immediate

Table 2.2 | The Different Memory Models Used by MASM

SI No	Name	How many code segments?	How many data segments?
1	Tiny	1	Nil
2	Small	1, max size 64K	1, max size 64K
3	Medium	Any number, any size	1, max size 64K
4	Compact	1, max size 64K	Any number, any size
5	Large	Any number, any size	Any number, any size
6	Huge	Any number, any size	Any number, any size

addressing. Hence DAT, which corresponds to a number, has to be loaded to AX (or any other 16-bit register) and then transferred to DS.

- iv) The value of DAT is loaded into DS. This means that DAT corresponds to the base address of the data segment. However, only when the loader loads this program in memory, will the actual base address (i.e., the upper 16 bits of the base address) be loaded into the DS register. At this point, the assembler just recognizes the 'reference' to a segment named DAT.
- v) The CS register need not be initialized in the program. This is done by the loader by default.
- vi) The program is terminated by using the INT 21H instruction with function number 4CH (DOS function calls will be discussed in detail later).
- vii) ASSUME is a directive which relates each segment to its segment register, by specifying the name.

However, the simplified dot models will be used throughout this text book, except when special reasons call for the use of the full segment definition. In older versions of MASM, before the advent of the .startup directive, segment registers had to be initialized with the instructions as below:

```
MOV AX, @DATA
MOV DS, AX
```

where DATA is used at the start of the data segment.

2.3.10 | General Rules for Writing Assembly Language

Now, let us summarize the general format for an assembly language line, with an example.

Example 2.13

LABEL	OPCODE	OPERAND	COMMENT
NUM1	DB	56H	;define NUM1 as a byte of 56H
NUM2	DW	0557AH	;define NUM2 as a word of 0557H
START:	MOV	BL, NUM1	;copy NUM1 into BL
	MOV	AL, BL	;copy BL into AL

An assembly language line has four fields, namely, label, opcode, operand and comment. A label is positioned at the left of a line and is the symbol for the memory address which stores that line of information. There are certain rules regarding labels that are allowed under MASM.

- i) All labels should begin with a letter or the special characters @, \$, _ and ?
- ii) A label can have 1 to 31 characters which may be digits, letters in upper or lower case or the special characters at (@), dollar (\$), underscore (_), dot (.) or question mark (?).
- iii) No reserved words of MASM may be used.
- iv) Each label must be unique.

The second field is the opcode or instruction field. The third is the operand field, and the last is the comment field which starts with a semicolon. The use of comments is advised for making programs more readable.

2.4 | Instruction Design

Until now, we were discussing assemblers, the best assemblers available and how to run our programs using assemblers and other program development tools. Now, let us go a bit deeper into the core of the processor and investigate the process of how machine codes have been designed. In doing so, we will realize why the assembly process is a 'one to one' process.

2.4.1 | Manual Coding

Another idea to think of is the possibility of doing manual or hand coding, as it is called – taking an assembly instruction, looking up or finding out its machine code and feeding it directly to the processor. Why not? It looks quite feasible, doesn't it? Actually it is, but very soon it becomes obvious that it is too cumbersome to think of programming as fun.

During the time that the 8085 was a popular processor, hand coding was the rule rather than the exception. This was because 8085 had a limited instruction set and a look up table was sufficient for hand coding. The scenario has changed with the 8086, which has more registers and more instructions – but the most notable fact is that it has many, many more modes of addressing – which implies that each instruction can be expressed in too many modes to facilitate the use of a look up table. With the more advanced of the x86 processors coming to the scene, the difficulty has multiplied many fold. As such, we will conclude that using a standard and good assembler is the best solution.

However, it will be very instructive and informative to understand how and why a particular mnemonic has been converted to a particular code. Let us understand the underlying principles of the Instruction Set Architecture of 8086.

2.4.2 | Instruction Set Architecture (ISA)

The following statement gives an almost appropriate definition of what is meant by ISA. An instruction set, or instruction set architecture (ISA), is defined as the part of computer architecture related to programming, including the native data types, instructions, registers, addressing modes, memory architecture, interrupt and exception handling, and external I/O. An ISA includes a specification of the set of opcodes (machine language), i.e., the native commands implemented by a particular CPU design.

It must be noted that processors with the same ISAs need not have the same internal designs (micro architectures) – they only need to share a common instruction set. That is why we

find that AMD64 and P64 have the same ISA but they are not the same internally. The ISA is an interface between a higher level language and the machine language.

2.4.3 | Instruction Set Design of 8086

Now let us get to the 8086 processor, which is a CISC processor. One feature of CISC is that its instruction size varies – all instructions do not have the same size. 8086 has instruction size varying from one byte to five bytes. This makes the processes of assembly, disassembly and instruction decoding complicated because the instruction length needs to be calculated for each instruction. An instruction should have the complete information for fetching and executing an instruction. It should thus have the following information:

- i) Opcode corresponding to the operation to be carried out.
- ii) Size of the operands.
- iii) Addressing mode of each operand mentioning one or more of the following:
 - a) General-purpose register.
 - b) Value of an immediate operand.
 - c) Address of operand.
 - d) Base register.
 - e) Index register.
 - f) Combination of a base and index register.
 - g) Displacement in combination with base/index/[base + index] registers.

See the general format of the first three bytes of an instruction (see Fig 2.2).

Prefix

This is an optional byte and need to be used only to change the operation, e.g. segment override prefix.

First byte

This is considered to be the first byte of an instruction. In this byte, there is the operation code (opcode) which is 6 bytes long. This is the code which defines the operation to be carried out. The other two bits are the D and W bits.




Figure 2.2 | Format of three bytes of an instruction

W (1-bit) – operand size. W = 1, means word operand; W = 0, means byte operand.
D (1-bit) – Direction bit. D = 1, means register is destination; D = 0, means register is source.

Second byte

MOD (2-bit) – Register bits.
REG (3-bit) – the identifying code of the register used.
R/M (3 bits) – Specifying a register or memory operand.
The MOD and R/M bits together specify the **addressing mode** of the instruction.

Note All instructions need not have the W and D bits. In such cases, the size of the operand is implicit, and the direction is irrelevant.

2.4.4 | Designing a Code

The design of a code requires more information than has been presented so far. We need the identifying codes of the registers, obviously. Each 16-bit and 8-bit register has a unique code. These codes are given in Tables 2.3 and 2.4. In addition, a table showing the MOD and R/M bits corresponding to various combinations of addressing modes is also given. Using these, code design will be quite simple. We also need the Intel manual for the opcodes and formats of instructions. This is given in Appendix A.

Table 2.3 | Codes of General-Purpose Registers

REG	W = 0	W = 1
000	AL	AX
001	CL	CX
010	DL	DX
011	BL	BX
100	AH	SP
101	CH	BP
110	DH	SI
111	BH	DI

Table 2.4 | Codes of Segment Registers

Segment regs	Code
ES	00
CS	01
SS	10
DS	11

Table 2.5 | Codes Pertaining to Addressing Modes – MOD and R/M Bit Patterns

MOD R/M	00	01	10	11	
				W = 0	W = 1
000	(BX) + (SI)	(BX) + (SI) + d8	(BX) + (SI) + d16	AL	AX
001	(BX) + (DI)	(BX) + (DI) + d8	(BX) + (DI) + d16	CL	CX
010	(BP) + (SI)	(BP) + (SI) + d8	(BP) + (SI) + d16	DL	DX
011	(BP) + (DI)	(BP) + (DI) + d8	(BP) + (DI) + d16	BL	BX
100	(SI)	(SI) + d8	(SI) + d16	AH	SP
101	(DI)	(DI) + d8	(DI) + d16	CH	BP
110	d16 (direct address)	(BP) + d8	(BP) + d16	DH	SI
111	(BX)	(BX) + d8	(BX) + d16	BH	DI
MEMORY MODE			REGISTER MODE		

Note d8 and d16 are 8-bit and 16-bit displacement respectively.

With the help of the above tables, let us try designing a code for a few simple instructions.

i) MOV AX, BX

Coming to our task, first look in the Intel manual for the opcode of MOV, with the specification ‘Register/Memory to/from Register’. It is seen to be 100010. This will correspond to the first 6 bits of the first byte. W = 1, as the operand is a word. The direction bit D has to be defined. In this case, both the source and destination are registers. Let us view it as a word being moved to AX i.e., register AX is thus the destination. Hence D = 1. Then we must use the code of AX in the register field of the second byte. The first byte is

OPCODE						D	W
1	0	0	0	1	0	1	1

Byte 1

To design the second byte, first look in Table 2.5 for the MOD and R/M bits. We find that the MOD bits are 11 (last column of the table) for register addressing, the reg bits are 011 for AX (from Table 2.3) and the R/M bits are treated as the ‘second register’ field. Here, the second register is BX, the code of which is 011 (Table 2.5 4th row). Thus the second byte is

MOD		REG			R/M		
1	1	0	0	0	0	1	1

Byte 2

There is no other information needed in coding this instruction. This is a two-byte instruction with the code 8B C3 H.

ii) ADD BL, CL

OPCODE						D	W
0	0	0	0	0	0	1	0

Byte1

MOD		REG			R/M		
1	1	0	1	1	0	0	1

Byte 2

The opcode of ADD with the specification of ‘Register/Memory with Register to either’ is 00000. W = 0, since it is a byte transfer. D = 1 considering BL as the destination register. For the second byte, MOD = 11. Register code for BL is 011, R/M bits corresponding to CL is 001. The code thus is 02 D9 H.

iii) MOV CX, 0213H

This is an instruction in the immediate mode. Two bytes are necessary for the immediate data, which is 16-bit. Now let us look at Intel’s manuals on the format of MOV in the immediate mode ‘to register’. It shows as:

1011 W Reg	data	data if W=1
------------	------	-------------

Figure 2.3 | Format of a MOV instruction in the immediate mode

Here W = 1, Reg (CX) = 001. Hence, the coding of this instruction is 1011 1001 as the first byte followed by the word data in the little endian scheme as 13 02 H. The complete instruction is B9 13 02H.

iv) MOV DX, [23CDH]

OPCODE						D	W
1	0	0	0	1	0	1	1

Byte 1

MOD		REG			R/M		
0	0	0	1	0	1	1	0

Byte 2

The first byte is 8BH. The second byte has MOD = 00 (Column 2 in Table 2.5) and R/M bits are 110 (corresponding to direct addressing in Table 2.5). Thus, the second byte is 16H. This is followed by the direct address in little endian form. The 4-byte machine code is 8B 16 CD 23 H.

v) MOV [BX], DH

OPCODE						D	W
1	0	0	0	1	0	0	0

Byte 1

MOD		REG			R/M		
0	0	1	1	0	1	1	1

Byte 2

For the second byte, MOD = 00 corresponding to [BX]. The register bits are 110 for DH and R/M = 111 from Table 2.5. Thus the instruction code is 88 37H.

vi) MOV AL, [SI][BX]

OPCODE						D	W
1	0	0	0	1	0	1	0

Byte 1

MOD		REG			R/M		
0	0	0	0	0	0	0	0

Byte 2

By following the same procedure as discussed for the previous instructions, the code for this is 8A 00 H.

vii) ADD AX, 53H [SI][BX]

OPCODE						D	W
0	0	0	0	0	0	1	1

Byte 1

MOD		REG			R/M		
0	1	0	0	0	0	0	0

Byte 2

The third byte of the instruction is 53H, being the 8-bit displacement given in the instruction. Thus the code is 03 40 53 H.

Segment Override Prefix

viii) MOV CS: [BX], DH

The segment override prefix is a byte with the format 001XX110. In place of XX, the code of the segment register which overrides the default segment has to be inserted. For the above instruction, the prefix byte is 00101110, using the code of CS which is 01 (Table 2.4). We have already coded the instruction MOV [BX], DH. Appending the prefix to it, the three-byte code is now 2E 88 37 H.

Now, to confirm if the ‘hand coding’ we have done is correct, let us put it all in a program and examine the list file. Example 2.14a is the program and 2.14b is the unassembled file from the debugger, which contains the assembly code as well as the machine code. On checking it, we find that our hand coding indeed agrees with the assembler’s coding.

Example 2.14a

```

                .MODEL  SMALL
                .DATA
ORG 23CDH
DAT DW 0000
                .CODE
                .STARTUP
MOV AX, BX
ADD BL, CL
MOV CX, 0213H
MOV DX, DAT
MOV [BX], DH
MOV AL, [SI][BX]
ADD AX, 53H[SI][BX]
MOV CS: [BX], DH
MOV AX, CS: [BX]
                .EXIT
END

```

Example 2.14b

0B56:0017	8BC3	MOV	AX, BX
0B56:0019	02D9	ADD	BL, CL
0B56:001B	B91302	MOV	CX, 0213
0B56:001E	8B16CD23	MOV	DX, [23CD]
0B56:0022	8837	MOV	[BX], DH
0B56:0024	8A00	MOV	AL, [BX + SI]
0B56:0026	034053	ADD	AX, [BX + SI + 53]
0B56:0029	2E	CS:	
0B56:002A	8837	MOV	[BX], DH

In the debugger, the unassembled instructions are shown above, with the addresses (offsets within the code segments) of the instructions also shown. For example, in the line

0B56:0019 02D9 ADD BL, CL

0B56 is the CS value, 0019 is the offset of the instruction, and 02D9H is the machine code of the instruction ADD BL, CL.

Notice how the address of the next instruction depends on the size (number of bytes) of the current instruction.

Now that we have the procedure of coding for a few instructions, it can be concluded that manual coding is an alternative in the absence of an assembler, but it is not a good practice to rely on it. Using a good assembler is recommended for fast and efficient use of assembly code. However, knowing the methodology of coding does help us to get an idea of the way in which data, registers and instructions are handled by the processor.

Note All the instructions we have coded in this section are at least two bytes in length. However, there are some instructions that have only one byte – examples are STI (set interrupt flag) and CLD (clear direction flag). Try to find out why they are only one byte long.

KEY POINTS OF THIS CHAPTER

- An assembler is a software that translates assembly language into machine code on a one to one basis.
- Assemblers convert labels to memory addresses.
- An assembly language program contains instructions and directives – the former produces executable code but the latter only gives directions to the assembler.
- For an assembler, having to know the memory address corresponding to a label that has not yet been defined is called the ‘forward reference’ problem.
- Assemblers can be one pass, two pass or multi pass. Each pass refines the optimality of the assembly process.
- A number of assemblers are available for x86, of which MASM is currently one of the most popular ones.
- MASM 6.14 and MASM 32 are used in this book.
- The tiny model is a memory model which is compact, and allows only one segment. The run file generated herein is a com file.
- The small model needs a data segment as well. It generates an exe file as the run file.
- Data types have to be defined as bytes, words and so on, before labels are used for data.
- Programs can also be written using the conventional ‘full segment definition’.
- Hand coding of instructions is possible if all the information pertaining to instruction formats, mode, register codes and opcodes are available.

QUESTIONS


1. The translation from assembly code to machine code is a one-to-one process, but the translation from a high level language to machine code is not so. Explain why.
2. List the steps involved in converting source code to executable code, using an assembler and other program development tools.
3. What are the roles of the linker and loader in the generation of the run file?

4. How is a com file different from an exe file?
5. 'The assembled code is re-locatable'. What is the implication of this statement?
6. Why does a .com program have the origin at 0100H?
7. Is MASM a high level assembler (HLA)? Comment.
8. How does a debugger help in the programming process?
9. What exactly does the .EXIT statement do?
10. List the rules regarding labels that can be used in MASM.

EXERCISE

1. Using the tiny model, write a program that does the following.
 - a) Moves 1020H to AX, 056H to BL and C76H to CX.
 - b) Copies the content of AX to DS, the content of BH to AH, and the content of CX to DX.
2. Write a program to copy 25 to AL, 'S' to BL and '*' to CL. Then move these to AH, BH and CH respectively.
3. Using the small model, store the bytes 56H and FCH, and words 0978H and CFD4H in the data segment.
4. Store three bytes and four words in the data segment. Copy one byte and one word of these to appropriate registers.
5. Using the small model, write a program which incorporates the following addressing modes.
 - a) register mode
 - b) immediate mode
 - c) register indirect mode.
6. Write a program which adds a constant correction factor to a temperature that is stored in the data segment. The corrected temperature should be stored in a new location in memory.
7. Using the DUP directive allocate space for 10 data words, 5 double words and 2 quad words in the data segment. Also fill the next 12 spaces with 0, the next 10 spaces with 'S' and the next space with 'S'.
8. Using the full segment definition, write a program to add the content of two words which are in memory, and store the sum in a third address.
9. Design the machine codes for the following instructions.
 - i) XOR AL, AL
 - ii) AND AX, BX
 - iii) OR AL, BL
 - iv) MOV DS, AX
10. Find the machine codes for the instructions
 - i) SUB AL, [5676H]
 - ii) CMP AX, 0CDE2H
 - iii) CMP 34H [DI][BX], DL
 - iv) ADD CX, [SI]

3 PROGRAMMING CONCEPTS-II



In this chapter, you will learn

- Good programming practices.
- How to use a few important DOS functions in programs.
- A part of the instruction set of 8086.
- To use data transfer instructions in different modes of addressing.
- The concept and importance of branch instructions.
- To use unconditional and conditional branch instructions.
- To use the arithmetic instructions catering to unsigned numbers.
- To use the multiply and divide instructions for interesting applications.
- The use of the logical, shift and rotate instructions of the 8086.

3.1 | Approaches to Programming

We are now in a position to feel that we have understood to a small extent, the tools needed for programming (in this case assembly programming). However, before getting down to the actual coding process, it might be worthwhile to spend some time trying to understand the various approaches to be used in programming. These ideas apply to all levels of programming i.e., whether we use high level languages or assembly level languages.

One common argument is whether programming is an art or a skill. To say that it is an art implies that good programmers are born, not made. This may be partly true, because it is seen that some people have a natural flair for programming. However, this applies, to all abilities we have. Even if programming skill is an inborn ability, one has to polish it and practice this craft to become good at it. There is no other way but hard work to become an expert.

Coming to the other side of the argument that programming is a skill – then also, the same rules apply. The craft of programming has to be learned with sound basics – we should develop our thought processes to become good problem solvers. For every problem, there is more than one way of solving it. Finding the optimal solution depends to a great extent on how much we have learned to build castles in the air – where these castles are the solutions to our problems.

Let us conclude that just anyone can become a reasonably good programmer, by getting a good grasp of the basics and following it up with keen studying and continuous practice.

There are various approaches to solve a programming problem. One is called the **top down** approach. This approach tends to look at the problem as one whole functional block. The solution to the problem is planned, after the whole problem is thoroughly understood.

This functional block is then broken down to smaller blocks, whose functions are to be defined along the way.

The other approach is the **bottoms up** approach. Here the basic functional blocks which are available or can be clearly defined, are joined together to get the bigger module. Thus, small function blocks are integrated to form the top level entity. Both these approaches have their advantages and disadvantages. Depending on the type of problem at hand, the methodology has to be decided.

Now, a Tip and Advice Before Getting Started

Something that is to be avoided is to jump to 'coding' as soon as a problem is presented. We must first build the solution to our problem, optionally in our mind – but the best way is to write a pseudo code or draw a flow chart before writing the actual code. In this way, our thought processes lead us to the algorithm to use. As we become more familiar with programming, we may gradually avoid pseudo codes or flow charts for small programs, but use them for the more complex ones. In this book, only the coding part is done – it is left to the student to visualize or write the algorithm, as he chooses.

Assembly Programming Tips

From the previous chapter, you might have got an insight into assembly programming using MASM. Now, we are in a position to get a better idea of the instruction set of 8086, in a phased manner.

In Chapter 2, a few programs were run, but we did not have any mechanism to display the outputs. There was also no mechanism to take in data through the keyboard i.e., interactive programming could not be done. There we used the debugger to access memory and view registers in order to verify if our results are as they should be. For entering data, we were supposed to store it in the data part of the segment being used. For a programmer, this approach seems cumbersome. Interactive programming is definitely more fun and instructive. Let us see what it is.

What is DOS?

We are now in the computer age when our learning of computers starts with using the latest version of Windows. However, that was not so in the beginning of computer evolution. A number of operating systems were tried, tested and discarded before the graphical interface that Windows is, came to common use. Earlier, DOS was a very popular OS. It is still used and is very useful.

DOS stands for Disk Operating Systems which could be an acronym for any OS, but it is most often used as shorthand for MS-DOS (Microsoft disk operating system). Originally developed by Microsoft for IBM, MS-DOS was the standard operating system for IBM-compatible personal computers.

The initial versions of DOS were very simple and resembled another operating system called CP/M. Subsequent versions have become increasingly sophisticated as they incorporated features of minicomputer operating systems. However, DOS is still a 16-bit operating system and does not support multiple users or multitasking.

DOS is considered insufficient for many of the graphical operations that we use our computers for, but all Windows have a strong dependency on DOS. So, newer versions of DOS are inbuilt in all Windows OS. DOS is considered to control hardware directly. It is this aspect of DOS which will prompt us to use DOS and BIOS interrupts in this book. A few basic commands of DOS are given in Appendix B. They will help you to get started on MASM 6.14.

3.1.1 | BIOS and DOS Function Calls

Until now, we were viewing the 8086 processor as being connected only to memory. For example, the physical memory – RAM. However, in reality, the processor is also connected to various input and output devices. For example, the PC has the keyboard and mouse as input devices. The video monitor is an output device. As programmers, we would like to input data into the processor using the keyboard and get the results displayed on the video monitor. To do this, we need to have some way of accessing these I/O devices, which contain a lot of extra hardware. Luckily, there is software available to handle this problem, and they come in the form of functions written to access Input and Output devices. The word BIOS may seem familiar. BIOS (Basic Input Output System) consists of a set of functions written for this purpose. Besides this, there is another set of functions called DOS functions or DOS interrupts. They are part of the DOS operating system. They are functions written for accessing input and output devices, which are called in the form of software interrupts. We have seen one such DOS interrupt which is the one that the .EXIT statement gets translated to:

```
MOV AH, 4CH  
INT 21H
```

In this example, 4CH is the function number which has to be loaded into the AH register before calling the interrupt with type number 21H. This function caused exiting the program and returning to the DOS prompt.

In Chapter 8, we will see the logic of interrupts and their way of functioning. However, in this chapter, we will use a few DOS interrupts to aid us in the process of programming.

Using these simple DOS interrupts will help us get the feel of programming as we are able to use the available input and output devices interactively. We will see that assembly language programming can be fun. In this chapter, a few DOS function calls will be introduced. BIOS interrupts will be used in later chapters.

3.1.2 | Using DOS Function Calls

We will start with four important function calls, all of which are of interrupt type 21H.

- i) Read the keyboard with echo

```
MOV AH, 01  
INT 21H
```

This call exits with the ASCII value of the key pressed, being available in AL. The key pressed is also **echoed** on the screen.

- ii) Read keyboard without echo

```
MOV AH, 08  
INT 21H
```

This call exits with the ASCII value of the key pressed being available in AL. The key pressed is **not echoed** on the screen.

- iii) Write a character to the standard display unit.

For this, the ASCII value of the character to be displayed should be in DL.

```
MOV DL, 'S'  
MOV AH, 02  
INT 21H
```

With this, the character S is displayed on the video screen.

iv) Display a character string on the standard display unit

The logical address DS : DX should point to the beginning of the string. This is to be followed by the following instructions:

```
MOV AH, 09
INT 21H
```

3.1.3 | The Instruction Set of 8086

The instruction set of the 8086 is reasonably large. Besides, each instruction can be used in various addressing modes making the whole set larger. Knowing all the instructions, their format and their features is a key factor in becoming an expert programmer. Let us list out the general features of instructions which will become clearer as they are explicitly earned.

- i) Most instructions have two operands i.e., the destination and the source. These can be registers or memory operands. However, both operands cannot be memory operands. Some instructions can also use immediate data as the source operand.
- ii) Certain instructions have only one operand mentioned which is either the destination or the source. These can be register or memory operands, but not immediate data.
- iii) There are instructions that have no operands explicitly mentioned in the instructions. In that case, some register is implicit in the usage of these instructions.

To learn the instruction set, the approach will be to list instructions based on their functionality, learn their format and function, and then write a few programs using them.

3.2 | Data Transfer Instructions

Table 3.1 gives all the data transfer instructions of 8086, except the input and output instructions, which will be introduced in Chapter 5. A few of these instructions will be discussed in detail and used in programs here. The rest will be used in later chapters.

3.2.1 | MOV – Move

Usage: MOV destination, source

This instruction causes the source content to be copied to the destination (copying does not change the contents of the source). Remember that the data types of source and destination should match – both should be either bytes or words.

MOV AX, CX	;copy CX to AX – this is a word operation
MOV AL, AH	;copy AH to AL – this is a byte operation
MOV AX, [BX]	;copy into AX the data in the address pointed by BX
MOV COST, DX	;copy the word in DX to the location labeled COST

The use of the MOV instruction in the register and immediate addressing modes was covered in Chapter 2. Hence, we will use MOV in other modes of addressing now.

Example 3.1

	.MODEL SMALL	;select small model
	.DATA	;start data segment
COSTP DB	67H	;store cost price

```

SELLP      DB      ?                ;space for selling price
                .CODE                ;start code segment
                .STARTUP            ;start program

PROFIT EQU    25H                ;define the profit

                MOV AL,COSTP          ;move COSTP to AL
                ADD AL, PROFIT        ;add PROFIT to AL
                MOV SELLP, AL        ;store the sum in SELLP

                .EXIT                ;exit to DOS
                END                  ;end of program

```

Table 3.1 | List of 8086 Data Transfer Instructions with Format and Function

Sl No.	Instruction format	Function performed	Flags affected
1	MOV dest, src (Move)	Copy the contents of the source to the destination	None
2	LEA reg16, memory (Load Effective Address)	Load the offset of specified memory location in the (16-bit) register	None
3	XCHG dest, src (Exchange)	Exchange the contents of the destination and source	None
4	PUSH source (Push)	Transfer one word from source to the stack top (SS : SP).	None
5	POP destination (Pop)	Transfer word at the current stack top (SS : SP) to the destination	None
6	LDS reg16, memory (Load to DS and register)	Load far pointer to DS and register	None
7	LES reg16, memory (Loads to ES and register)	Load far pointer to ES and register	None
8	LAHF (Load AH with flags)	Copy bits 0–7 of the flag register into AH	None
9	SAHF (Store AH with flags)	Transfer bits 0–7 of AH into the flag register	None
10	PUSHF (Push flags)	Push the flag register onto the stack	None
11	POPF (Pop flags)	Pop word from stack into the flag register	None
12	XLAT (Translate)	Replaces the byte in AL with byte from a user table addressed by BX	None

Example 3.1 shows the use of the MOV instruction in the direct mode of addressing. This mode uses a memory location as one operand. Here two memory locations labeled COSTP and SELLP have been defined in the data segment. COSTP has a data byte stored in it, but SELLP is only allocated space for later storage. The first MOV instruction gets the data byte in COSTP into AL. After adding a pre-defined constant PROFIT to it, the sum in AL is stored in SELLP. Both MOV instructions use **direct addressing**. At this stage, you need to remember that the assembler will do the conversion of the labels to memory addresses.

Example 3.2

```

                                .MODEL SMALL          ;select small model
0000                            .DATA                  ;start data segment

0000 10 20 30 40 50  ARRAY DB      10H, 20H, 30H, 40H, 50H
                                ;setup array
0000                            .CODE                  ;start code segment
                                .STARTUP                ;start program

0017 BF 0000                    MOV DI,0               ;load DI with 0
001A 8A 85 0000 R              MOV AL,ARRAY[DI]        ;get first element in AL
001E 04 07                      ADD AL,07H             ;add 07 to it
0020 83 C705                    ADD DI,05H             ;add 5 to DI
0023 88 85 0000 R              MOV ARRAY[DI], AL       ;move sum to memory

                                .EXIT                  ;exit to DOS
                                END                     ;end of program

```

Example 3.2 shows the list file of a program which uses the MOV instruction in the **register relative addressing** mode. Let us analyze the program. The data segment is first initialized with 5 bytes. The address of the first location is labeled as ARRAY. DI is used to address the element numbers. With DI = 0, the first element is accessed, and with DI = 5, the sixth element is accessed. Then the effective address (EA) is [ARRAY + DI]. When DI = 0, the instruction MOV AL, ARRAY[DI], transfers the data in the first memory location to AL. To this, 07 is added. With DI = 5, ARRAY[DI] becomes the effective address of the sixth location in memory. Then MOV ARRAY[DI], AL stores the sum in memory. The displacement for the relative mode is the value of the label ARRAY. From the list file, it is seen the offset corresponding to ARRAY is 0000. Since ARRAY is the address of the first location in the data segment, the data segment will look like this, after the program is executed. The sum will be in the location ARRAY + 5.

E A	CONTENT
ARRAY + 5	17H
ARRAY + 4	50H
ARRAY + 3	40H
ARRAY + 2	30H
ARRAY + 1	20H
ARRAY + 0	10H

Example 3.3

```

                                .MODEL SMALL          ;select small
                                model
0000                                .DATA              ;start data
                                                segment
0000  34 87 56 05 07 ARRAY  DB  34H, 87H, 56H, 05H, 07H
                                                ;setup array
0000                                .CODE              ;start code
                                                segment
                                .STARTUP              ;start program

0017  BB 0000 R      MOV  BX, OFFSET ARRAY  ;address ARRAY
001A  BF 0000      MOV  DI, 0              ;load DI with 0
001D  8A 01      MOV  AL, [BX + DI]        ;get first data
001F  04 35      ADD  AL, 35H              ;add 35H to it
0021  BF 0005      MOV  DI, 05              ;address new
                                                location
0024  88 01      MOV  [BX + DI], AL        ;move sum here

                                .EXIT              ;exit to DOS
                                END              ;end program

```

Example 3.3 is the list file of a program that performs the same function as the previous example, but uses the **based indexed mode** of addressing i.e., the effective address is the sum of a base register and an index register. The instruction `MOV BX, OFFSET ARRAY` causes the offset of `ARRAY` (which is 0000 as seen from the list file) to be loaded into `BX`. Thus, `BX` now is a pointer to the location labeled as `ARRAY`. However, to point to each element one by one, another register `DI` is also used. With `DI = 0`, the instruction `MOV AL, [BX + DI]` causes the data in the first location to be copied to `AL`. With `DI = 5`, `[BX + DI]` accesses the sixth location in the data segment.

3.2.2 | LEA – Load Effective Address

Usage: `LEA reg16, memory`

This instruction causes the offset of the specified memory location to be loaded in the indicated register.

```

LEA SI, COSTP      ;load the offset of memory location COSTP in SI
LEA BX, ARRAY      ;load the offset of memory location ARRAY in BX

```

You may have noticed that the above is a replacement for the instruction `MOV BX, OFFSET ARRAY` used in Example 3.3.

Now let us use this instruction in an interesting application. We would like to display two character strings on two lines on the video monitor. It was discussed in Section 3.1.2 that DOS function calls are needed for this. The following is the list file for the program which does this.

Example 3.4a

```

                                .MODEL SMALL
                                .DATA
0000                                MSG1      DB      "HELLO$"
0000 48 45 4C 4C 4F 24          MSG2      DB      0AH, 0DH, "I AM SAM$"
0006 0A 0D 49 20 41 4D          20 53 41 4D 24
                                .CODE
0000                                .STARTUP

0017 8D 16 0000 R              LEA DX, MSG1      ;load offset of MSG1 in DX
001B B4 09                      MOV AH, 09H      ;DOS function number for
001D CD 21                      INT 21H          ;displaying a string

001F 8D 16 0006 R              LEA DX, MSG2      ;load offset of MSG2 in DX
0023 B4 09                      MOV AH, 09H
0025 CD 21                      INT 21H

                                .EXIT
                                END

```

Example 3.4b

```

HELLO
I AM SAM
C:\masm6.14\BIN>

```

Now let us analyze the program that does this, the listing of which is shown in Example 3.4(a). In the data segment, two messages are stored at offsets with labels MSG1 and MSG2. The character strings are enclosed in double quotes and terminated by the \$ sign. This sign is mandatory to end a character string which is to be printed on the display device. Notice that the listing in Example 3.4a shows the corresponding ASCII values for each character of the string. There is an ASCII value for 'space' which comes between each word of the second message i.e., I AM SAM. It is seen to be '20H'. Note that '24H' is the ASCII value of \$ and it is seen at the end of both strings.

On the second line, we also see two characters 0AH and 0DH. These are the ASCII values for line feed (which advances the cursor to the next line) and carriage return (sets cursor to the left position of screen). These are necessary for printing to occur on the second line, starting from the left.

To print the first message, DX points to the offset MSG1. This is done by using the LEA instruction, which loads the offset of MSG1 in DX. Then the DOS function with number 09 is called. Similarly for the second message, DX should point to the offset of MSG2. The two messages as they are printed on the command window are seen in Example 3.4(b). The .EXIT command causes control to return to DOS prompt

3.2.3 | XLAT – Translate a Byte in AL

Usage: XLAT

This instruction translates a byte in AL with a byte from a look-up table in memory. XLAT is an instruction which can handle look up tables and access data depending on the value of the pointing number. There should be a one-to-one relationship between the number and the data in the table. Before using XLAT, the look up table must be in memory and the offset of the table must be in BX.

Requirements Load into BX, the offset of the look-up table.
Then use XLAT
Then into AL, load the number which will extract the required data. Refer Table 3.2. For example, if you want to get the ASCII value of 5, load AL = 5. It will get the sixth entry from the table. See the look up table which relates a decimal number to its ASCII value. The ASCII values such as '0', '1', '2' have numerical values 30H, 31H, 32H. Refer to the ASCII table in Section 0.6.5.

Example 3.5

The following program gets into the memory location starting from VALUES, the ASCII values of 1, 3, 5, 4, 0 ... in that order. The look-up table is stored in the array named ASC.

```
.MODEL SMALL
.DATA
ASC DB  '0', '1', '2', '3', '4', '5', '6', '7'
NUM DB  1, 3, 5, 4, 0
VAL DB  ?
.CODE
.STARTUP
LEA BX, ASC           ;let BX point to offset of the table
LEA SI, NUM           ;let SI point to NUM
LEA DI, VAL           ;let DI point to VAL
MOV CX, 5             ;CX = 5
AGN: MOV AL, [SI]      ;move the content of NUM to AL
XLAT                  ;translate
MOV [DI], AL          ;move translated value from AL to VAL
INC SI                ;increment the pointer to NUM
INC DI                ;increment the pointer to VAL
LOOP AGN              ;repeat until CX = 0
.EXIT
END
```

Table 3.2 | Look-Up Table for Decimal to ASCII Conversion

Number	ASCII
0	'0'
1	'1'
2	'2'
3	'3'
4	'4'
5	'5'
6	'6'
7	'7'
8	'8'
9	'9'

XLAT is equivalent to the following instructions:

```
MOV AX, 0
MOV SI, 0
MOV AL, [BX + SI]
```

SI or DI will be okay for showing the equivalence.

Try using XLAT for various look up tables, for example:

- i) An integer and its square.
- ii) An integer and its cube.
- iii) A BCD number and its seven segment code.
- iv) A key and its ASCII or seven segment code.

3.2.4 | PUSH and POP

Usage: PUSH source

POP destination

For the 8086, only 16-bit data can be pushed and popped.

Note CS is not a valid destination for POP.

PUSH BX	;save the contents of BX to stack
PUSH [BX]	;save to stack the contents of the word pointed by BX
PUSH DS	;save to stack the contents of the segment register DS
PUSH COSTP	;save to stack the contents of the word location COSTP
POP CX	;load the contents of the stack top to CX
POP [SI]	;load the contents of the stack top to the memory location pointed by SI

The PUSH and POP instructions are relevant only for the stack segment. **So scroll down and read about the operation of the stack before attempting to understand the PUSH and POP instructions.**

The stack is a region of temporary storage. When a subprogram is called, the register contents of the main program at that point of time and the current value of IP and CS are pushed onto the stack, and retrieved on returning from the sub program. These things are done automatically as part of a CALL instruction. Besides this, a programmer can opt to store in stack temporarily, the contents of some registers, so that he can use those registers for other computations. He can get back the previous contents of these registers, after he completes his current computation. These operations are accomplished by the PUSH and POP instructions.

3.2.4.1 | Operation of the Stack

The 8086, as mentioned in Chapter 1, uses a Last-In First-Out (LIFO) stack. A stack is an area of memory, in which the lowest address is the base address of the segment. The upper 16 bits of this base address will be in the SS register. The highest address will be in the SP register – or rather, the SP will contain the offset of the highest address with respect to the base address. This is why it is called the TOP of STACK. For example, if the stack is defined from addresses 20000H to 203E8H, SS = 2000H and SP = 03E8H, this stack then has a size of 3E8H bytes (1000 bytes). In the 8086, only a word (16 bits) can be pushed on to stack. A push thus causes two memory locations to be accessed.




Figure 3.1 | PUSH operation in a stack




Figure 3.2 | POP operation in a stack

The operation of the stack is different from normal memory operations. It is a push-down memory. For example, the value of SP mentioned above, let us say, the PUSH BX instruction is to be executed. Let BX = CC99H. Recollect that SP = 03E8H. The operation of pushing is as follows. SP is first decremented by two. The new data will be put in addresses (offsets) 03E7H (SP-1) and 03E6H (SP-2), adhering to the ‘little endian’ concept i.e., the higher byte of BX will be the higher address (03E7H) and the lower byte in the lower address (03E6H). The new value of SP will be 03E6H. This shows that new data to be pushed onto the stack will be pushed in the lower two addresses with respect to the current stack top. Always, the stack top (SP) has the offset corresponding to the last byte that was pushed in.

The reverse will be the case after a POP operation. Assume a POP CX operation at a time when the SP value is 0312H. The content of the stack will first be loaded into the CX register, and SP is incremented by two. See Fig 3.2.

3.2.4.2 | Defining a Stack

Now, how will we use the stack in programming? Obviously we have to define elements including the stack size, base address of segment and the stack top. If we are using the tiny segment, this is done for by DOS – but in other memory models, we will have to do it ourselves. In case we do not, a warning ‘No stack segment’ will appear on assembling. This can be ignored if the

stack size is less than 128 bytes. Otherwise, the stack that gets defined by DOS may erase the PSP (program segment prefix) and cause the system to crash.

Example 3.6a

```
.MODEL SMALL
.STACK 400H                ;specify stack size
```

Example 3.6b

STACKSEG	SEGMENT STACK
	DW 200H (?)
STACKSEG	ENDS

Example 3.6 shows two ways of defining a stack. Example 3.6a uses the simplified memory model and it specifies the size of the stack as 400 bytes. Example 3.6b uses the full segment definition. 200 data words are to be set aside as stack, as per this definition. Here, both definitions create stacks of the same size.

Example 3.7

Explain what is done in this program. Assume SP = 0310H, when the stack was initialized.

```
.MODEL SMALL
.STACK 300H
.CODE
.STARTUP
MOV AX, 4567H
MOV BX, 0ACEH
PUSH AX
PUSH BX
POP AX
POP BX
.EXIT
END
```

Solution

This program defines a stack of 300 bytes by the use of the stack directive. The program first loads two numbers in AX and BX. These numbers are pushed into stack using the PUSH operations. Followed by, POP operations. What we need to remember is that popping occurs in the reverse direction. What is pushed in last will be popped out first. The last PUSH was PUSH BX. It is obvious then that by the POP AX operation, the content of AX is replaced by the content of BX. Similarly the second POP causes BX to get the content of AX. Thus, this program exchanges the content of these two registers, using the stack as a temporary data storage area. After execution, the contents of the registers are: BX = 4567H and AX = 0ACEH.

Note Exchanging register content is not really an application for the stack. We can exchange these register contents easily by using the instruction XCHG AX, BX. However, this program shows the relationship between POP and PUSH operations.

3.3 | Branch Instructions

Branch instructions are very important, because they carry on them the full power of a computer. Branching can be conditional or unconditional. Unconditional branching becomes necessary when looping back or forward is done infinitely, like when we use software to generate a square wave continuously. However, most of the time, conditional branching is what may be more important. Taking decisions based on the result of a computation is what gives computers their power and versatility. Most instructions that we will use subsequently, become meaningful only when conditional branching is used. In this chapter, we will discuss only the jump and loop instructions. Other branching instructions like call and return will be discussed in Chapter 4. We will start with the jump instruction.

3.3.1 | JMP – Jump

Usage: JMP destination

The destination is a memory location, which can be expressed in different ways as will be discussed presently. The jump destination is frequently referred to as ‘target’ as well. A jump instruction breaks the normal sequence of program execution, and takes control to a different location in the code segment. There are many ways and restrictions in specifying the target value, and based on this, we can have different types of jumps. Fig 3.3 could be a typical scenario. When the ADD instruction is being executed, IP (Instruction Pointer) points to the next instruction which is an **unconditional** JMP instruction. On decoding the JMP instruction, the EU (Execution Unit) realizes that the next instruction should be taken from the location with label AGAIN. It is obvious then, that the IP content must be changed in such a way that a new sequence is to be followed. The salient points regarding jump are:

- i) The jump location may be in the same code segment, in which case it is **near** jump. For a near jump, the value of CS remains the same, but the value of IP will change. This is also called an intra-segment (within segment) jump.
- ii) The jump location may be in a different code segment – then it is called a **far** jump. In this case, both CS and IP will have to take on new values. This is an inter-segment (between segments) jump.
- iii) The label of the destination address should be followed by a colon.
- iv) Jumping can be backward or forward in the program sequence.



Figure 3.3 | Control flow using the jump instruction

- v) Jumping can be unconditional. Then the format of the jump instruction is `JMP label`, where label is the address of the destination.
- vi) Besides unconditional jumps, there are a number of conditional jumps available. The conditions are based on the state of the flags which get affected by arithmetic or logic instructions.

3.3.1.1 | The Near Jump

In the **direct mode** of addressing, the near jump has the format `JMP label` where the label is the address of a memory location. The destination can be a 16-bit signed number, which means that the range of jumping is within $\pm 32\text{K}$ bytes. In direct jump, the jump location is a signed displacement from the current value of IP to the destination address i.e., the assembler calculates this displacement, and adds it to the current IP value. Thus, the jump address is relative or rather, re-locatable. For a forward jump, this displacement will be positive, while for a backward jump, this will be negative. It uses three bytes as instruction size – one byte for the opcode, and two bytes for the 16-bit destination offset.



$$\text{NEW IP} = \text{CURRENT IP} + \text{OFFSET (16-bit)}$$

Figure 3.4 | Format of the near jump instruction

3.3.1.2 | The Short Jump

The short jump is a special case of a near jump with the format `JMP destination`. For short jumps, the displacement must be an 8-bit signed number, which means that the jump destination can only be -128 or $+127$ bytes displaced from the current value of IP.



$$\text{NEW IP} = \text{CURRENT IP} + \text{OFFSET (8-bit)}$$

Figure 3.5 | Format of the short jump instruction

A directive `SHORT` can be used if we are sure that our jump destination is within the above specified range. This instruction is two bytes long – one-byte for the opcode and one-byte for the destination offset. Now let us see an address calculation for a short jump instruction.

Example 3.8

0000	B4 01	BEGIN:	MOV AH, 01
0002	CD 21		INT 21H
0004	8A D0		MOV DL, AL
0006	EB 07		JMP SHORT THERE
-----,skipped instructions>-----			
000F	80 C2 02	THERE:	ADD DL, 02
0012	B4 02		MOV AH, 02
0014	CD 21		INT 21H
0016	EB E8		JMP SHORT BEGIN

Example 3.8 shows how ‘short’ jump instructions pass control from one part of the program to another.

0006 EB 07

JMP SHORT THERE

See the above line which is taken from the list file of the corresponding program. It shows the address (0006) at which this instruction is stored. The opcode of this direct JUMP instruction is EB and the ‘displacement’ is 07.

During the time that this instruction is decoded, the IP value is 0008 (address of the next instruction). At the beginning of the first pass, the assembler reserves one-byte space in memory for the address of THERE. The assembler knows that only one-byte is required for THERE, because of the directive SHORT. Otherwise it would have reserved two bytes space, as for a general near jump. If the address of the next instruction (0008) is added to the sign extended displacement (0007) of the jump instruction (this displacement is calculated by the assembler), the address of THERE is at location $0008H + 0007H = 000FH$. Thus, control branches to location 000FH.

Now, see the other direct jump instruction in the program

0016 EB E8

JMP SHORT BEGIN

The opcode is EB and the displacement is E8. JMP BEGIN is seen to be a backward jump as BEGIN is a label appearing much before this instruction line. The assembler calculates this displacement (which is negative) and then adds this displacement to the value of IP at that point of the instruction sequence. Then it jumps to the destination BEGIN. In this case the negative displacement is E8H. When JMP SHORT BEGIN is executed, the IP content will be 0018H. For jumping to BEGIN, the address calculation is $0018 + FFE8 = 0000$, which is the address of the instruction with the BEGIN label. **Note that the negative number E8H, when used in 16-bit calculations, has to be sign extended. So E8H is FFE8H in 16-bit sign extended form.**

Now let us write and run a fun program using an unconditional jump instruction. Example 3.9a is a simple program named UPLOW which uses DOS interrupts to enter a character from a keyboard. If this character is an upper case alphabet, it converts it to lower case. The ASCII character table shows that a lower case character can be obtained by adding 20H to the upper case code i.e., the ASCII of ‘A’ is 41H and that of ‘a’ is 61H, and so on.

Example 3.9a

```

                                .MODEL TINY
                                .CODE
                                .STARTUP
START:    MOV AH, 01H           ;AH = 01
          INT 21H              ;enter a key
          ADD AL, 20H          ;add 20H to it
          MOV DL, AL           ;move the sum to DL
          MOV AH, 02           ;AH = 02
          INT 21H              ;use this function to display
          MOV DL, 0AH          ;0AH is the ASCII for newline
          MOV AH, 02H          ;display a new line
          INT 21H
          JMP START            ;go to start to repeat
          .EXIT
          END

```

Example 3.9b

```
C:\masm6.14\BIN>UPLow
Aa
Bb
Cc
Dd
Ee
Ff
Gg
Hh
^C
C:\masm6.14\BIN>
```

INT 21H with function number 01 is used to enter a key. Then 20H is added to the ASCII value of the key, and this is displayed in the same line. To get further displays on separate lines, the ASCII character 0AH is displayed. Remember that 0AH corresponds to 'newline'. The JMP START instruction is used to repeat this whole sequence indefinitely. Since there is no stopping condition for this loop, we have to use the 'ctrl c' combination of the keyboard to terminate the program. The displayed output is shown in Example 3.9b. Note that, for any key input other than an upper case character, the output is some random character.

3.3.2 | Other Forms of the Unconditional Jump Instruction

There are other modes of addressing for the near jump.

i) JMP reg16

This is called an indirect jump. The 16-bit register contains the destination address (offset). This jump is not a relative jump. IP gets replaced by the value in the register.

JMP BX	;jump to the destination address in BX. Make IP = BX
JMP SI	;jump to the destination address in SI. Make IP = SI

ii) JMP [reg 16]

This is like a double indirect jump. The register points to an address which contains the jump location.

JMP [SI]	;jump to the address which is stored in the memory location pointed by SI
----------	---

If SI = 0670H, the destination address is taken from the data segment at locations 0670H and 0671H. IP is replaced by the value at these addresses.

3.3.2.1 | Conditional Jumps

Conditional jumps are the best part of the idea of control transfer. They change the sequence of program execution based on the result of a computation which causes flag bits to be set or reset. However, there is one case (JCXZ) where the register content is checked.

Note All conditional jumps are short jumps.

Usage: J(condition) destination.

Table 3.3 contains the jumps used for unsigned data, and those directly testing flags and registers.

Note

- i) Certain mnemonics are equivalent – for example, JZ and JE mean the same as they test the same flags. So, do JC and JB, JP and JPE and so on. More can be identified from the list.
- ii) JCXZ tests the CX register but other Jump instructions test conditional flags only.
- iii) **There is another set of conditional jump instructions which are specially meant for signed arithmetic. The list of those instructions and their usage will be dealt with in Chapter 4.**
We will use a few of the conditional jumps in the forthcoming examples.

3.3.2.2 | Far Jump

A far jump is an intersegment jump, which means that the destination address is in a different code segment. This will be a 5-byte instruction, the first byte being the opcode, the second and

Table 3.3 | List of Conditional Jump Instructions which Cater to Unsigned Arithmetic and which Directly Address Flags or Registers

Sl No.	Mnemonic	What it means	Condition tested
1	JA	Jump if Above	CF = 0 and ZF = 0
2	JAE	Jump if Above or Equal	CF = 0
3	JB	Jump if Below	CF = 1
4	JBE	Jump if Below or Equal	CF = 1 or ZF = 1
5	JC	Jump if Carry	CF = 1
6	JCXZ	Jump if CX Zero	CX = 0
7	JE	Jump if Equal	ZF = 1
8	JNA	Jump if Not Above	CF = 1 or ZF = 1
9	JNAE	Jump if Not Above or Equal	CF = 1
10	JNB	Jump if Not Below	CF = 0
11	JNBE	Jump if Not Below nor Equal	CF = 0 and ZF = 0
12	JNC	Jump if No Carry	CF = 0
13	JNE	Jump if Not Equal	ZF = 0
14	JNP	Jump if No Parity	PF = 0
15	JNZ	Jump if Not Zero	ZF = 0
16	JP	Jump if Parity	PF = 1
17	JPE	Jump if Parity Even	PF = 1
18	JPO	Jump if Parity Odd	PF = 0
19	JZ	Jump if Zero	ZF = 1

OPCODE	IP LOW	IP HIGH	CS LOW	CS HIGH
--------	--------	---------	--------	---------

Figure 3.6 | Format of the far jump instruction

third, the new value of IP, and the fourth and fifth, the new values of CS. To specify that a jump is to a different code segment, a **far pointer** can be used in the format.

```
JMP FAR PTR AGAIN ;PTR stands for 'pointer'
```

Defining an address as external and then using a jump to this address is also possible. These ideas will be discussed in Chapter 5.

3.3.3 | The LOOP Instruction

Usage: LOOP label.

There is a very useful and frequently used instruction called LOOP. This combines jump with a counter. The register CX is assigned to decrement every time LOOP executes. When CX = 0, the looping is exited.

```
MOV CX, N
MORE: .....
      .....
      .....
      .....
      LOOP MORE
      .....
      .....
```

Initially CX is loaded with a number, which we call 'count'. See the above program sequence. What is to be understood here is that, the sequence of instructions up to the line containing the LOOP instruction, will execute once. Then, when the LOOP instruction is encountered, CX is decremented and then tested. If CX is found to be equal to zero, looping is exited and the next instruction in the sequence is taken up. If CX! = 0, control returns to the label MORE.

LOOP can be combined with other conditions, to make it a conditional instruction. LOOPNE/LOOPNZ and LOOPE/LOOPZ can be used. These instructions test the zero flag, as well as the value of CX. Examples 3.10a and b show the use of an unconditional LOOP instruction. In this program, three arrays have been declared in the data segment. Corresponding bytes from the first two arrays are to be added and the sum is to be saved in the third array. To point to the data items in the three arrays, three registers are used. Data is accessed using register indirect addressing. After one round of these actions, the pointer registers are incremented using the INC instruction. This instruction adds 1 to the registers. This sequence of operations is repeated six times by using the LOOP instruction. Once CX = 0, the loop is exited and the program ends.

Example 3.10a

```
.MODEL SMALL
.DATA
NUMS1 DB 45H, 67H, 89H, 65H, 34H, 23H
NUMS2 DB 09, 09, 12H, 13H, 08, 02
NUMS3 DB 6 DUP(0)
```

```

        .CODE
        .STARTUP
        LEA BX, NUMS1    ;use BX to point to NUMS1
        LEA SI, NUMS2    ;use SI to point to NUMS2
        LEA DI, NUMS3    ;use DI to point to NUMS3
        MOV CX, 06       ;CX = 6
REPEA:  MOV AL, [BX]      ;take into AL data from the array NUMS1
        ADD AL, [SI]      ;add to AL the data in array NUMS2
        MOV [DI], AL      ;save the sum in the third array
        INC BX            ;increment pointer
        INC SI            ;increment pointer
        INC DI            ;increment pointer
        LOOP REPEA        ;repeat the sequence until CX = 0
        .EXIT
        END

```

This problem can be done with less number of address registers if the ‘register relative mode’ of addressing is used. See Example 3.10b. Here only the register BX is used. Its content is added to different addresses so as to access the three different arrays.

Example 3.10b

```

        .MODEL SMALL
        .DATA
NUMS1 DB 45H, 67H, 89H, 65H, 34H, 23H
NUMS2 DB 09, 09, 12H, 13H, 08, 02
NUMS3 DB 6 DUP(0)
        .CODE
        .STARTUP
        MOV CX, 6
        MOV BX, 0
REPEA:  MOV AL, NUMS1[BX] ;point to the first array
        ADD AL, NUMS2[BX] ;add numbers of the second array
        MOV NUMS3[BX], AL ;save in the third array
        INC BX
        LOOP REPEA        ;repeat until CX = 0
        .EXIT
        END

```

Now, we will digress to the discussion of arithmetic instructions, but you can see that branch instructions in various forms will continue to be used.

3.4 | Arithmetic Instructions

The complete list of arithmetic instructions is given in Table 3.4. Instructions from 1 to 9 will be discussed in detail and used in this chapter. The rest will be dealt with, in Chapter 4 for signed number arithmetic.

Table 3.4 | Full Set of Arithmetic Instructions

SI No.	Instruction format	Function performed	Flags affected
1	ADD dest, src	Add dest and src	AF CF OF PF SF ZF
2	ADC dest, src	Add dest, src and CF	AF CF OF PF SF ZF
3	INC dest	Add 1 to the destination	AF --- OF PF SF ZF
4	SUB dest, src	Subtract src from dest	AF CF OF PF SF ZF
5	SBB dest, src	Subtract source and CF from dest	AF CF OF PF SF ZF
6	DEC dest	Subtract 1 from dest	AF --- OF PF SF ZF
7	CMP dest, src	Subtracts source from destination and updates the flags but does not save result. Only flags are affected	AF CF OF PF SF ZF
8	MUL src	Unsigned Multiply	-- CF OF --- -- --
9	DIV src	Unsigned binary divide	--- -- -- -- -- --
10	CBW	Convert byte to word	None
11	CWD	Convert word to double word	None
12	DAA	Decimal Adjust for Addition	AF CF --- PF SF ZF
13	DAS	Decimal Adjust for Subtraction	AF CF --- PF SF ZF
14	NEG dest	Two's Complement Negation	AF CF OF PF SF ZF
15	IMUL src	Signed Multiply	-- CF OF --- -- --
16	IDIV src	Signed Integer Division	--- -- -- -- -- --
17	AAA	ASCII Adjust for Addition	AF CF --- -- -- --
18	AAS	ASCII Adjust for Subtraction	AF CF --- -- -- --
19	AAM	ASCII Adjust for Multiplication	--- -- -- PF SF ZF
20	AAD	ASCII Adjust for Division	SF --- -- PF --- ZF

Note

- i) In the list, dest means destination and src means source. After an operation, the result is available in the destination.
- ii) As mentioned earlier, the destination and source may be both registers or a memory location and a register. The source can be immediate data, unless explicitly mentioned (for certain instructions) that immediate mode of addressing is not allowed.
- iii) Flags are affected by most of the arithmetic instructions. In certain cases, all flags are affected, but some flags are undefined. The --- notation for the corresponding flags is used in such cases meaning 'undefined' (in the table).
- iv) In certain cases, flags are not affected. This is specified by 'none'.

3.4.1 | Flag Control Instructions

There are certain instructions which can be used to set/reset flags. This is very relevant in the case of control flags like direction flag, interrupt flag and so on. However, there is one conditional

flag which can be set/reset using instructions, and that is the carry flag. The relevant instructions for it are:

- i) CLC – Clear carry.
- ii) STC – Set carry.
- iii) CMC – Complement carry.

Now, we will see the format and usage of some of the important and commonly used arithmetic instructions.

3.4.2 | Addition Instructions

ADD – Add.

Usage: ADD destination, source.

This instruction adds the destination and source and puts the sum in the destination. All conditional flags get affected.

ADD AH, AL	;add AL to AH, sum in AH
ADD AL, COSTP	;add the byte in COSTP to AL, sum in AL
ADD BX, 0987H	;add the number 987H to BX, sum in BX
ADD CX, [BX]	;add the word in the location pointed by BX to CX sum in CX

ADC – Add with carry.

Usage: ADC destination, source.

This instruction adds CF and source to the destination, and puts the sum in the destination. There are three operands, of which the third is the carry. All conditional flags get affected.

ADC AH, 0	;AH = AH + 0 + CF
ADC [BX], AL	;add the byte pointed by BX with AL and CF, put sum in the location pointed by BX
ADC AX, [BX][SI]	;add to AX, CF and the word with EA = BX + SI sum in AX

INC – Increment.

Usage: INC destination.

This instruction adds 1 to the destination. All conditional flags, **except the carry flag**, get affected.

INC BX	;add 1 to the content of BX
INC [BX]	;add 1 to the content of the memory location pointed by BX
INC AH	;add 1 to the content of AH

The PTR directive

This is a convenient point to introduce the PTR (pointer) directive. When the size of the operand is not implicit in the instruction, a pointer is used to indicate whether the operand is a byte, a word, or a double word. For example, see the instruction.

INC [BX]

Here, BX is a pointer to a location whose content is to be incremented. It is not clear whether the data being pointed to, is a byte or a word or a double word. To make it clear, it will be helpful to modify the instruction as:

```

                INC BYTE PTR [BX]           ;byte pointer
or             INC WORD PTR [BX]          ;word pointer
or             INC DWORD PTR [BX]         ;double word pointer

```

as the case may be. We will encounter various instances where this pointer is used.

Example 3.11

The following is a program which adds two bytes stored in memory. The sum of the bytes is likely to be greater than FFH. Hence, a word space needs to be allocated for the sum.

```

                .MODEL SMALL
                .DATA
NUMS DB 95H, 0FCH           ;store the two bytes
SUM  DW ?                  ;allocate a word space for the sum

                .CODE
                .STARTUP
MOV AX, 0           ;AX = 0
CLC                ;clear the carry flag
MOV AL, NUMS        ;move the first number to AL
ADD AL, NUMS + 1     ;add the second number to AL
ADC AH, 0           ;add AH and CF and 0
MOV SUM, AX         ;the sum in AX is moved to memory
                .EXIT
                END

```

First AX is loaded with 0. Observe how a byte addition causes a result to be a word. The sum of two bytes which are added to AL is too large to fit in AL. Hence, it generates a carry. The next instruction adds CF, 0 and AH which contains a 0. So we find that the carry bit is accommodated in AH. Thus, AH-AL i.e., AX contains the sum which is a word.

Example 3.12 adds 10 bytes stored in memory. The sum of these ten bytes will not fit into a byte location. So a word location is allocated for the sum. It must also be noted that the sum will definitely fit into a word location. [If there are ten bytes, the maximum number possible for the sum of ten bytes is only $255 \times 10 = 2550$, which is less than 65,536 the maximum decimal number that a word can hold.]

Example 3.12

```

                .MODEL SMALL
                .DATA
NUMS DB 245, 178, 190, 167, 56, 178, 250, 89, 150, 235
SUM  DW 0

                .CODE
                .STARTUP
CLC                ;clear the carry flag

```

```

        LEA BX, NUMS           ;BX to point to the numbers
        MOV CX, 10             ;move the count to CX
        MOV AX, 0              ;make AX = 0
REPEA:  ADD AL, [BX]           ;add the numbers, sum in AX
        ADC AH, 0              ;add the carries into AH
        INC BX                 ;increment the pointer
        LOOP REPEA            ;repeat if CX is not equal to 0
        MOV SUM, AX           ;store AX in SUM
        .EXIT
        END

```

The above is a simple program which gives a lot of information. For one thing, note that data is written in the decimal form, and not as hexadecimal bytes. However, if the list file is checked, it will show that the data is converted to hexadecimal format (in the machine code part of the listing). The program uses the LOOP instruction to carry out cumulative addition. CX is used as a counter, which decrements with each addition. The first ADD instruction adds two bytes. It is possible that a carry is generated out of this addition. This carry is added to AH, the other operand of which is 0. Thus, the sum of the carries of the ten addition operations will be in AH, which will form the upper byte of the sum. This sum is thus available as a word in AX which is saved in the word location SUM. It is important to clear the carry flag before using the ADC instruction, to avoid the possibility of any residual carry leading to a wrong result.

Example 3.13 is an interesting program which adds two double words. A doubleword is two words long i.e., 4 bytes long.

Example 3.13

```

        .MODEL SMALL
        .DATA
LONGNUM1 DD 0F8FC6768H        ;first double word data
LONGNUM2 DD 0C6EF2109H        ;second double word data
SUM       DD 0                 ;sum
SUMC      DB 0                 ;space for carry out
        .CODE
        .STARTUP
        LEA SI, LONGNUM1       ;SI as pointer to first operand
        LEA DI, LONGNUM2       ;DI as pointer to second operand
        LEA BX, SUM             ;DX as pointer to the sum
        CLC                    ;clear carry
        MOV CX, 2               ;CX to count the number of words
REPEA:  MOV AX, [SI]            ;move to AX the first word
        ADC AX, [DI]            ;add the second operand and carry
        MOV[BX], AX             ;the lower word of result stored
        INC SI                  ;increment SI twice to point to
        INC SI                  ;next word operand
        INC DI                  ;increment DI twice to point to
        INC DI                  ;second word operand
        INC BX                  ;increment BX twice to point to

```

```
INC BX                ;second word of the sum
LOOP REPEA           ;go back if CX is not 0
MOV DL, 0             ;make DL = 00
ADC DL, 0             ;add the carry to DL
MOV SUMC, DL          ;move the carry to SUMC
.EXIT
END
```

There are a number of notable features for this program. Let us list out these features.

- i) The double words are stored in two DD locations named LONGNUM1 and LONGNUM2. They will be stored in the little endian format.
- ii) The sum of two double words can be more than a double word. This can be accommodated as a carry, added to a register, and can be stored in memory as the uppermost byte of the sum. Three registers are used as pointers to each double word. However, addition can be done only as words, as the maximum size of an operand is only 16 bits (called a word), for the 8086. Hence, first the lower order words of the two double words are added. There is a possibility of a carry due to this addition, which may have to be added to the upper word. Since, addition is looped, the ADC instruction is used so as to accommodate the addition of the carry.
- iii) The address pointers have to be incremented by 2 to point to the upper word of the double word. **The INC instruction is preferred over the ADD instruction, as it does not affect the carry flag.** It is important to preserve the carry bit out of the second word operation. This final carry is added to DL (which is first ensured to be 0). The numbers used herein are two double words whose sum is a double word, plus a byte (01 BFEB 8871H).

Table 3.5 | Relevant Portion of the Data Segment After the Execution of Example 3.12

Offsets in memory	Data	
000C	01	} SUMC
000B	BF	
000A	EB	} SUM
0009	88	
0008	71	
0007	C6	} LONGNUM 2
0006	EF	
0005	21	
0004	09	
0003	F8	} LONGNUM1
0002	FC	
0001	67	
0000	68	

- iv) The memory locations and the corresponding data, after program execution is as shown (assuming the origin of the data segment is 0000).
- v) The three arrays in memory can be accessed by using ‘register relative addressing’ as in Example 3.10b. Try to see if it gives any advantage in terms of reducing the number of instructions/registers used.

3.4.3 | Subtraction

SUB – Subtract.

Usage: SUB destination, source.

This instruction subtracts the source from the destination. The result is in the destination. All conditional flags are affected.

SUB AX, BX	;subtract BX from AX
SUB AL, [BX]	;subtract the byte pointed by BX from AL
SUB COST[SI], CX	;subtract CX from the word with EA = COST + SI
SUB AX, 8956H	;subtract 8596H from AX
SUB CL, BYTE PTR [SI]	;subtract from CL the byte pointed by SI

SBB – Subtract with borrow.

Usage: SBB destination, source.

This instruction subtracts the source and the carry flag from the destination. All conditional flags are affected.

SBB CH, 7	;subtract from CH, 7 and CF – result in CH
SBB AX, [BP + 2]	;subtract from AX, the word pointed by [BP + 2]. Since BP is used, the data is taken from the stack segment. The result is put in AX

DEC – Decrement.

Usage: DEC destination.

This instruction subtracts 1 from the destination. All conditional flags, except the **carry flag**, are affected.

DEC CL	;subtract 1 from CL
DEC WORD PTR [SI]	;subtract 1 from the word pointed by SI
DEC BYTE PTR NUMB[BX]	;subtract 1 from the byte pointed by the effective address NUMB + BX

Example 13.14

Explain what occurs on execution of the following instructions:

MOV CL, 0B5H
SUB CL, 0FCH

Solution

B5H	181	1011 0101	–
– FCH	– 252	1111 1100	
<u>B9H</u>	<u>–71</u>	<u>1011 1001</u>	

This is a subtraction of a bigger number from a smaller number. Hence, the carry flag is set to indicate that borrowing was required in the subtraction. The difference is -71 (decimal), which is represented as $B9H$ which is the two's complement representation of decimal -71 . The carry flag is set to indicate that the subtraction has used a 'borrow'. The sign flag is set to indicate that the MSB of the result is 1, which in this case is interpreted to be a negative number.

Now let us use the SUB instruction in an interesting example. Division is an operation which can be achieved by repeated subtraction. Many early microprocessors (8085, for example) had no divide instruction. The 8086, has a divide instruction. However, here, let us use repeated subtraction to divide say, 100 by 9. We know that the quotient is 11 ($0BH$) and the remainder is 1.

Example 3.15

```

.MODEL SMALL
.DATA
QUOTIENT DB 0
REMAINDER DB 0

.CODE
.STARTUP
MOV AL, 100          ;copy dividend to AL
MOV CL, 0             ;CL, the quotient register = 0
REPEA: SUB AL, 9      ;subtract 9 from 100
JC OVER              ;stop when CF = 1
INC CL               ;increment CL if CF is not set
JMP REPEA            ;repeat subtraction if CF is not set
OVER: ADD AL, 9       ;add 9 to AL to get remainder
MOV REMAINDER, AL    ;store remainder
MOV QUOTIENT, CL     ;store quotient
.EXIT
END

```

The steps of the program are as follows:

- i) Subtract 9 from 100 repeatedly until a negative number is obtained. In the program, start with $AL = 100$. When AL becomes negative, the carry flag will be set (due to borrow). This is the stopping condition for subtraction.
- ii) In this problem, after subtracting 9 from 100 eleven times, AL will contain 01 . One more subtraction will cause the content of AL to become -8 ($F8H$ in 2's complement form). The carry flag is found to be set, and further subtraction is stopped.
- iii) To get the remainder, add 9 (the divisor) to AL . We get 1. This is the remainder.
- iv) Every time a subtraction is performed unhindered, increment a count in CL . This will give the quotient.
- v) The quotient and remainder are stored in their allocated space in the data segment.

The SBB instruction can be used in instances when multibyte subtraction is done. For example, to subtract two double words or quad words, SBB will have to be used, just as ADC is used (Example 3.13) for multi-byte addition.

Table 3.6 | Flag Settings After a Compare Instruction

If	CF	ZF
destination > source	0	0
destination < source	1	0
destination = source	0	1

We have seen so far, addition and subtraction operations using numbers in the hexadecimal (essentially binary) format. However, numbers are also represented in other formats such as BCD and ASCII. There are special instructions like DAA and DAS that cater to BCD numbers, and instructions like AAA, AAS which cater to ASCII operations. In Section 4.4.2, these instructions will be used in arithmetic calculations using such numbers.

3.4.4 | Compare Instruction

CMP – Compare.

Usage: CMP destination, source.

This instruction compares the two operands and causes the conditional flags to be affected, but neither the destination nor the source changes. Comparison is done by a subtraction operation, and the flags are set/reset according to the result of this. However, only two flags really matter – they are the Zero flag and the Carry flag.

Consider the instruction CMP destination, source. The condition of the flags will be as shown in Table 3.6. It is thus obvious that following up a compare instruction with a JNC/JC or JNZ/JZ will do the trick. If the instruction is, say CMP AL, BL the way to picture it is

if AL > BL, CF is reset
if AL < BL, CF is set
if AL = BL, ZF is set

If two unsigned numbers are to be compared, it will be clearer if we use JA/JB as the mnemonic following the compare instruction, instead of JC or JNC. This will make it easier to comprehend the meaning of the result of the comparison. If we want to compare AX and BX we could think of it as: Jump to target if AX is above BX. The instructions to use are:

CMP AX, BX
JA Target

Similarly, when testing for equality, we could use JE/JNE instead of JZ/JNZ.

Note These tips are only for enhancing the readability of the program. Now, let us write a few programs illustrating the use of the compare instruction.

Example 3.16

Two unsigned words are stored in the data segment in locations WORD1 and WORD2. The program compares the two numbers and displays a message stating which number is bigger.

```
.MODEL SMALL
.DATA

WORD1 DW  ----- ;store the first number
WORD2 DW  ----- ;store the second number
MES1  DB  "WORD1 IS BIGGER$" ;first message string
MES2  DB  "WORD2 IS BIGGER$" ;second message string
```

```

                .CODE
                .STARTUP
                MOV AX, WORD1           ;copy WORD1 to AX
                CMP AX, WORD2           ;compare AX with WORD2
                JB SECOND                ;If AX < WORD2, go to SECOND
                LEA DX, MES1            ;point DX to MES1
                JMP DISP                ;jump to DISP
SECOND:         LEA DX, MES2            ;point DX to MES2
DISP:          MOV AH, 09              ;AH = 09 for string display
                INT 21H                ;DOS function INT21H is called
                .EXIT
                END

```

This is a very simple program in which the WORD1 is brought into AX. It is then compared with WORD2. The jump (JB) instruction following the comparison, directs control to the appropriate message. Remember that DOS function call with AH = 09 will display a string on the console.

Note To run this program, actual data will have to be written in the data segment at locations WORD1 and WORD2.

Example 3.17 finds the biggest of 10 bytes stored in memory.

Example 3.17

```

                .MODEL SMALL
                .DATA
NUMS            DB 56H, 38H, 09H, 98H, 99H, 0C7H, 07H, 0BCH, 0CH, 0ECH
BIGGEST         DB ?                  ;store the biggest number
                .CODE
                .STARTUP
                LEA BX, NUMS           ;BX is the pointer to the array
                MOV CL, 0AH            ;CL acts as the counter
                MOV AL, 0              ;AL = 0
REPEA:          CMP AL, [BX]           ;compare each number with AL
                JAE AGAIN              ;if AL is above or equal, jump
                MOV AL, [BX]           ;if below, bigger no. to be in AL
AGAIN:          INC BX                ;increment pointer
                DEC CL                 ;decrement counter
                JNZ REPEA              ;repeat the sequence if count! = 0
                MOV BIGGEST, AL        ;the biggest number will be in AL
                                        ;store it in the location BIGGEST
                .EXIT
                END

```

The logic of this program is as follows:

- i) The numbers are stored as an array in memory named NUMS. BX is used as a pointer to this array.
- ii) CL contains the count of the numbers (here CL = 0AH).

- iii) Initially AL = 0.
- iv) AL is compared with each number. If AL is not above or equal to the number in the array, the bigger number is brought to AL.
- v) If AL is above the number in the array, the content of AL is retained as it is.
- vi) This is one round of comparison. Then the pointer (BX) is incremented and count (CL) is decremented.
- vii) When the count reaches zero, ZF is set and the looping is terminated because of using the JNZ mnemonic.
- viii) The biggest number is now in AL which is stored in memory in the space allocated for it.

Example 3.18 is another interesting program, which searches a character string for the presence of a particular character, and displays appropriate messages.

Example 3.18

```

.MODEL SMALL
.DATA
STRIN DB "HOWAREYOU MYBOY"      ;string of characters
LEN    DW 0EH
MSG1   DB 0AH, 0DH, "CHARACTER FOUND$"
MSG2   DB 0AH, 0DH, "CHARACTER NOT FOUND$"

.CODE
.STARTUP
    LEA BX, STRIN      ;BX to point to the string
    MOV CX, LEN        ;copy the string length to CX

    MOV AH, 01         ;DOS function call for
    INT 21H            ;character input with echo
                        ;the entered character is in AL
REPEA:  CMP [BX], AL    ;compare AL with string bytes
        JE FOUND       ;if equality is found, display
                        ;appropriate message
        INC BX         ;point to the next byte in array
        LOOP REPEA     ;if count is not 0, repeat
        LEA DX, MSG2   ;character not found, go to MSG2
        JMP DISP

FOUND:  LEA DX, MSG1    ;display MSG1
DISP:   MOV AH, 09      ;DOS function call for display
        INT 21H        ;of string
        .EXIT
    END

```

A character string is stored in location labeled STRIN. The length of this string is specified in a word location named LEN. The character whose presence (or absence) we want to check for, is entered through the keyboard using the DOS function call with AH = 01. This character is compared with each of the characters in the string. Once equality is found, the comparing loop

is exited and MES1 is displayed. Otherwise, MES2 is displayed, after all the characters in the array have been compared. Note that both the messages (MESG1 and MESG2) are preceded by the newline (0AH) and carriage return (0DH) characters. This is to ensure that the display is obtained in a different line, after the input character (the one that is entered through the keyboard) is echoed on the screen.

3.4.5 | Unsigned Multiplication

MUL – Multiply.

Usage: MUL source.

This instruction multiplies a number in AL or AX by the source (where the source can be a register or a memory location, **but not an immediate number**). All the conditional flags are affected, but only the CF and ZF are defined as meaningful for the result. The destination depends on the size of the operand. There are two ways of performing multiplication.

i) Byte by byte

In this, one of the operands must be in the AL register, and the source can be a byte in a register or memory location. The product (a word) will be in AX.

MUL BL	;multiply BL by AL – product in AX
MUL BYTE PTR[SI]	;multiply the byte pointed by SI, by AL – product in AX
MUL BIG	;multiply the content of BIG by AL – product in AX

ii) Word by word

In this, one of the operands must be in the AX register, and the source can be a word in a register or memory location. The product will be in DX and AX, with the upper word in DX.

MUL CX	;multiply CX by AX – product in DX and AX
MUL WORD PTR [DI]	;multiply the word pointed by DI with AX – product in DX and AX

iii) Word by byte

This is only a special case of the word $\times$ word multiplication. The byte must be extended to be a word by making the upper byte to be 0. If the byte is in AL, extend it to be a word by making the content of AH to be zero. Thus, AX now contains one of the operands.

Flags Affected

For the multiply instruction, all conditional flags are affected, but only the carry (CF) and overflow (OF) flags have any significance. They will be set or reset according to the size of the product. Recall that a byte $\times$ byte multiplication can produce a result of a word size. However, if the operands are small, the product itself is only a byte, and needs only the AL register for the product. In that case, both the carry and overflow flags are found cleared, but if the product occupies a word size, both these flags are set (CF = 1, OF = 1). This concept can be extended to the word $\times$ word multiplication as well. We can summarize that if the product has a size equal to the size of the operands, both these flags are reset. i.e., CF = 0, OF = 0. However, if the product is large enough to occupy the registers assigned for it, these flags are set i.e., CF = 1, ZF = 1. See this multiplication:

```
MOV AL, 78H
MOV CL, 0F9H
MUL CL
```

This program segment will cause the product to be 74B8H. Thus, the result of the multiplication is AX = 74B8H, CF = 1 and OF = 1.

However, if the program segment is

```
MOV AL, 13H
MOV CL, 09H
MUL CL
```

The result is AX = 00ABH, CF = 0 and OF = 0.

In example 3.19, two bytes stored in the data segment are multiplied. The result of multiplication is available in AX, which is then moved to the location PROD, a word location.

Example 3.19

```

.MODEL SMALL
.DATA
MULT DB 0AH           ;multiplier
MULP DB 0F6H          ;multiplicand
PROD DW ?             ;space allocated for product
.CODE
.STARTUP
MOV AL, MULP           ;move the multiplicand to AL
MUL MULT               ;multiply with the multiplicand
MOV PROD, AX           ;product moved from AX to memory
.EXIT
END

```

Now, let us use the multiply instruction in a more complex application. Example 3.20 is a program to find the factorial of a number N. For 8086, the maximum size of an operand for multiplication is only a word. This places a limitation on the value of N that can be used. It can be verified that N is to be less than 9, for the program to give the correct result. In this program, the value of N should be entered from the keyboard, remembering that N is to be less than 9. Since, any value that is entered from the keyboard is in ASCII form, 30H is subtracted from it, to get the binary value of N to be used in the computation.

Example 3.20

```

.MODEL SMALL
.DATA
FACT DW 0             ;space allocated for the factorial
.CODE
.STARTUP
MOV AH, 01
INT 21H               ;enter N from the keyboard
SUB AL, 30H           ;convert ASCII to binary
MOV AH, 0             ;convert N in AL to a word in AX
MOV BX, AX            ;move it to BX
MOV AX, 1             ;AX = 1, to start the iteration
CMP BX, 0             ;compare BX (= N) to 0

```

```

                JZ FINAL          ;if N = 0, jump to find 0!
REPEA:          MUL BX           ;for N not 0, multiply with AX
                DEC BX           ;decrement BX
                CMP BX, 0        ;compare with 0
                JNE REPEA        ;repeat if BX is not 0
FINAL:          MOV FACT, AX     ;AX = 1, hence 0! = 1
                .EXIT
                END

```

The computation is done iteratively. First, $AX = 1$. The value of N is copied to BX . For $N = 0$, the factorial is 1. For this case, the content of AX is transferred to the location $FACT$ and corresponds to 0! If N is not zero, then it is multiplied with AX . BX is decremented and cumulatively multiplied with the product in AX . As N becomes higher, $N!$ increases steeply, and only for numbers up to 8, with the factorial fit into a word location. Further multiplication becomes impossible because one of the operands has a size greater than a word. The value of $N!$ is finally stored in the memory location labeled $FACT$.

3.4.6 | Unsigned Division

DIV – Divide

Usage: DIV source

This instruction divides AX or $DX - AX$ by the source, where the source can be a register or a memory location, but not an immediate number. All the conditional flags are affected, but undefined – hence they do not give any interpretation or information about the result. The destination depends on the size of the operand. There are two ways of performing division:

i) Divide a word by a byte.

Here, the dividend must be a word placed in AX and the source must be a byte. The result of division causes AL to contain the quotient, and AH to contain the remainder.

```

                DIV BL           ;divide AX by the byte in BL
                DIV BYTE PTR [BX] ;divide the word in AX by the byte pointed by BX
                DIV DIG          ;divide the word in AX by the byte in DIG

```

See this division

```

                MOV AX, 0C678H
                MOV CL, 0F9H

```

This program segment will cause AH (remainder) = $0CH$ and AL (quotient) = CCH .

Dividing a byte by a byte.

This is just a special case of a division of a word by a byte. In this case, convert the dividend byte to a word by loading the dividend in AL and 0 in AH . Thus, AX will be the word that acts as the dividend.

ii) Divide a double word by a word.

In this case, the dividend has to be in AX and DX (the upper word in DX). The divisor should be a word. The result of this division causes the quotient to be in AX and the remainder to be in DX .

```

DIV BX           ;divide the double word in DX-AX by the word in BX
DIV WORD PTR [SI] ;divide the double word in DX-AX by the word pointed by SI
DIV ANGLE        ;divide the double word in DX-AX by the word in ANGLE

```

Dividing a Word by a Word.

Similar to the previous case, if we want a word by word division, extend the word in AX to be a double word by loading 0 in DX to get the dividend to be a double word in AX and DX.

Divide by Zero Error.

For division, if the divisor is zero, the quotient becomes undefined. In attempting such a division, the 8086 will exit from this program and generate an interrupt. This state is said to be a 'divide by zero error'. An interrupt generated by an error, is termed an exception. However, division by zero is not the only condition to cause such an error. If the quotient register is too small to accommodate the quotient, then also this happens. For example, if the dividend is a large number and the divisor is very small, such a condition is possible.

```

MOV AX, 09876H
MOV CL, 25H
DIV CL

```

The above program segment will give a quotient of 41EH, which obviously cannot be accommodated in AL. When such an error occurs, program execution is aborted and the assembler displays the message 'divide overflow error'. If this problem had been foreseen, the solution would be to convert the dividend to a double word, and the divisor to a word. Then the quotient will be in AX, which will be big enough for it.

Example 3.21

```

                                .MODEL SMALL
                                .DATA
FIRST      DB  89H
SECOND     DB  0CAH
AVG        DB  ?
                                .CODE
                                .STARTUP
MOV AL, FIRST    ;copy FIRST to AL
MOV AH, 0        ;zero-extend
MOV BL, SECOND   ;copy SECOND to BL
MOV BH, 0        ;zero-extend
ADD AX, BX       ;add the zero extended words
MOV CL, 02       ;load divisor to CL
DIV CL          ;divide AX by CL
MOV AVG, AL      ;save quotient to AVG
                                .EXIT
                                END

```

Example 3.21 gives a complete program for finding the average of two bytes stored in memory locations labeled FIRST and SECOND. These are copied to registers AL and BL and zero extended to convert them to words. Addition of the numbers is then done as words. This sum, which is the dividend, is in the AX register. The divisor 2 is copied to the CL register. After

division, the quotient is available in AL and the remainder in AH. For this problem, only the quotient is important, and this is saved to location AVG.

Now, we will see a very useful and interesting application of the division operation. This program converts a 16-bit hexadecimal number to decimal and displays the decimal number on the console after converting it to ASCII form. To see the logic of this program, let us take a number, say, 246.

- i) Divide this by 10 to get a quotient of 24 and a remainder of 6. Push the remainder on to the stack.
- ii) Next, divide the quotient again by 10 to get 2 as the quotient and 4 as remainder. Push this remainder also on to the stack.
- iii) One more division causes the quotient to be zero. This is the stopping condition in the division loop.
- iv) The number of division operations done is counted along with all these steps. Now, pop out the remainders from the stack. This will be in the reverse order of the push operations. Thus, 2 will be popped out first and 6 last. This can be displayed in this order.
- v) For displaying, each digit must be converted to its corresponding ASCII. This is done by adding 30H to each digit.

Now, examine Example 3.22. The program converts any 16-bit hexadecimal number to decimal and then displays it. Here, the hexadecimal number is placed in the data segment. Note that the division here is accomplished as a **double word by word** division.

Example 3.22

```

.MODEL SMALL
.DATA
NUM DW 0CEF6H      ;the hexadecimal number
COUNT DB 0        ;count of the number of divisions done
.CODE
.STARTUP
MOV AX, NUM        ;get the hex number into AX
MOV DX, 0          ;DX = 0, thus DX-AX is the dividend
MOV CX, 10         ;the divisor 10 is loaded into CX
REPEA: DIV CX       ;divide by CX
PUSH DX            ;the remainder in DX is pushed to stack
MOV DX, 0          ;DX = 0 again to get the dividend in DX-AX
INC COUNT          ;increment the count
CMP AX, 0          ;check if the quotient (in AL) is zero
JNE REPEA          ;if AL is not zero, repeat the division
DISP: POP DX        ;pop out the remainders for displaying
ADD DL, 30H        ;add 30H to convert to ASCII
MOV AH, 02         ;DOS function call for displaying
INT 21H            ;a character
DEC COUNT          ;decrement the count
JNZ DISP           ;if count = 0, it means all remainders
                  ;have been popped out

.EXIT
END

```

Table 3.7 | List of Logical Instructions and Functions Performed by them

Sl No.	Instruction format	Function performed	Flags affected
1	AND dest, src	Logical AND of the two operands returning the result in the destination	CF OF PF SF ZF (AF undefined)
2	OR dest, src	Logical inclusive OR of the two operands returning the result in the destination	CF OF PF SF ZF (AF undefined)
3	XOR dest, src	Performs a bitwise exclusive OR of the operands returning the result in the destination	CF OF PF SF ZF (AF undefined)
4	NOT dest	Inverts the bits of the “dest” operand forming the one’s complement	None
5	TEST dest, src	Performs a logical AND of the two operands updating the flags register without saving the result	CF OF PF SF ZF (AF undefined)

3.5 | Logical Instructions

Table 3.7 gives the complete list of logical instructions and the function performed

Example 3.23

Find the result and the state of the flags CF, ZF and OF due to the following instructions.
Given AX = 008CH, BX = 345EH, CX = 67EBH

- i) AND BL, CL
- ii) OR AH, BH
- iii) XOR AL, CH
- iv) TEST AH, BL

Solution

- i) AND BL, AL

BL =	5EH	0101 1110
CL =	<u>EBH</u>	<u>1110 1011</u>
	<u>4AH</u>	<u>0100 1010</u>

After this operation, BL contains 4AH. CF = 0, OF = 0, ZF = 0.

- ii) OR AH, BH

AH =	00	0000 0000
BH =	<u>34H</u>	<u>0011 0100</u>
	<u>34H</u>	<u>0011 0100</u>

After this operation, AH = 34H. CF = 0, OF = 0, ZF = 0.

iii) XOR AL, CH

AL =	8CH	1000 1100
CH =	<u>67H</u>	<u>0110 0111</u>
	<u>EBH</u>	<u>1110 1011</u>

After this operation, AL = EBH, CF = 0, OF = 0, ZF = 0.

iv) TEST AH, BL

AH =	00	0000 0000
BL =	5EH	0101 1110

After this operation, the ZF = 1, OF = 0, CF = 0.

Masking There is a word ‘masking’ associated with the AND operation. Masking is used to select the part of a word or a byte needed, while making the unwanted bits to be zero. For example, observe the value in AL after the following instructions are executed.

```
MOV AL, 78H
AND AL, 0FH
```

This gives AL = 08H in the AL register. The upper nibble of 78H has been masked. As another example, see below where 16-bit data is ANDed with 0FFFH, so as to mask the upper 4 bits. This is a particular case when only 12 bits of this data are needed.

```
MOV AX, 9876H
AND AX, 0FFFH
```

AX now has a content of 0876H.

Typical Applications of Logical Instructions Now, let us use these logical instructions in some simple applications.

AND To convert an ASCII number to a binary number, mask the number with 0FH.

```
MOV AH, 01           ;get in the number through the keyboard
INT 21H
AND AL, 0FH          ;mask the upper nibble
```

OR To convert an 8-bit binary number (from 0 to 9) to ASCII, OR with 30H.

```
MOV AL, 9             ;AL = 09
OR AL, 30H             ;AL = 39H
```

XOR To clear a register, use XOR.

```
XOR BL, BL            ;BL = 0
```

TEST To test whether a bit is set or not use the TEST instruction.

```
TEST BL, 01H
```

This tests whether bit D0 of BL is 1 or not. If the bit under test is reset, the AND operation corresponding to the TEST operations causes ZF to be set (ZF = 1).

```
TEST CX, 8000H
```

This tests whether D15 of CX is set or not. The result of this that ZF = 0 as D15 = 1.

3.6 | Shift and Rotate Instructions

Table 3.8 gives the complete list of the shift and rotate instructions.

3.6.1 | Shift

- i) Shift instructions do arithmetic or logical shifting.
- ii) They also shift right or left.
- iii) Arithmetic shift is used for signed number operations, while logical shift caters to unsigned numbers.
- iv) Shifting right causes a divide by 2, for each bit position shifted, while shifting left corresponds to a multiplication by 2.
- v) The count means the number of bit positions, by which shifting is to be done.
- vi) If the count >1, load it in CL, otherwise use '1' in the immediate mode.

SAL/SHL – Shift Left Arithmetic/Shift Left Logical.

Usage: SAL/SHL dest, count.

This instruction shifts the destination left by 'count' bits and zeroes are shifted in to fill the vacant positions on the right. The Carry flag contains the last bit shifted out. Because, '0' is shifted from the right side. **Left shift 'logical' and 'arithmetic' perform the same operation and hence the mnemonic SAL or SHL can be used interchangeably.**

Modifies Flags: CF OF PF SF ZF (AF undefined).




Table 3.8 | List of the Shift and Rotate Instructions .

Sl No.	Instruction format	Function performed	Flags affected
1	SHL dest, count	Shift logical left by 'count' bits	CF OF PF SF ZF (AF undefined)
2	SAL dest, count	Shift arithmetic left by 'count' bits	CF OF PF SF ZF (AF undefined)
3	SHR dest, count	Shift logical right by 'count' bits.	CF OF PF SF ZF (AF undefined)
4	SAR dest, count	Shift arithmetic right by 'count' bits	CF OF PF SF ZF (AF undefined)
5	RCL dest, count	Rotate through carry left by 'count' bits.	CF OF
6	RCR dest, count	Rotate through carry right by 'count' bits	CF OF
7	ROL dest, count	Rotate left by 'count' bits	CF OF
8	ROR dest, count	Rotate right by 'count' bits	CF OF

SHL BX, 1	;shift left logical by one position, the word in BX
SAL AL, CL	;shift left arithmetic, AL by the count specified in CL
SHL DATA2, CL	;shift left logical, the content of memory DATA2 by the count specified in CL
SHL BYTE PTR [BX][DI], 1	;shift left (logical) once, the byte with EA = BX + DI
SAL WORD PTR [DI], CL	;shift left (arithmetic) the word pointed by DI, by the count specified in CL

Example 3.24 shows a program segment for finding the weight of the modulo-2 sum of two numbers. Modulo-2 sum essentially is finding the positions in which the two numbers are different – which is an XOR operation. The weight of a number is the number of 1s in it.

In this problem, the modulo-2 sum is found by XORing the two numbers in AL and BL. The result of this is then shifted left. The bit shifted out will be in the carry flag. On checking the carry bit, if it is found to be 1, the register CH is incremented. Along with this, the register CL is decremented. CL contains the length of the operand. Here CL = 8. At the end of the program, we get the ‘weight’ of the modulo-2 sum of AL and BL in CH. For the numbers used in this program, CH = 6 is the result.

Example 3.24

```

.MODEL TINY
.CODE
.STARTUP
MOV AL, 89H      ;move first number to AL
MOV BL, 0F7H     ;move second number to BL
XOR AL, BL       ;EX-OR the two numbers
MOV CL, 8        ;CL = 8
BACK:  SHL AL, 1  ;shift left AL once
      JNC REPEA   ;repeat if there is no carry
      INC CH      ;increment CH if there is a carry
REPEA: DEC CL     ;decrement count
      JNZ BACK    ;repeat the shifting until CL = 0
      .EXIT
      END


```

SHR – Shift Right Logical.

Usage: SHR dest, count.

This instruction shifts the destination right by ‘count’ bits and fills with zeroes the vacant positions on the left. The Carry flag contains the last bit shifted out.

Modifies Flags: CF OF PF SF ZF (AF undefined)



SHR DX, 1	;shift right by one position, the word in DX
SHR WORD PTR [SI] THERE, CL	;shift right the word with EA = THERE + SI, by the count specified in CL
SHR CH, CL	;shift right the content of CH, by the count in CL
SHR BYTE PTR [BP][SI], 1	;shift right by one the byte pointed by EA = BP + SI

Example 3.25 shows an example of a shift operation. When a decimal number is represented as a 4-bit binary nibble, and these nibbles are packed in a byte, it is called packed BCD representation. For example, the packed BCD representation of 56 is 0101 0110. Unpacking it means separating it as two bytes 0000 0101 and 0000 0110.

Example 3.25

Convert a packed BCD byte to two unpacked bytes.

Solution

```

.MODEL TINY
.CODE
.STARTUP
MOV AL, 95H           ;copy the packed BCD byte to AL
MOV BL, AL            ;copy the same byte to BL
AND AL, 0FH           ;mask the upper nibble of AL
AND BL, 0F0H          ;mask the lower nibble of BL
MOV CL, 04            ;move a count of 4 into CL
SHR BL, CL            ;shift BL right 4 times
.EXIT
END

```

In this problem, a packed BCD byte is in AL. It is to be unpacked and placed in two registers. The operand is in AL. Keep a copy of it in BL too. For unpacking, first the upper nibble of AL is masked. This gives one unpacked BCD byte. This is now in AL. Next, the lower nibble of BL is masked, and the result is shifted right 4 times to bring the upper nibble data to the lower nibble position. Now, the two unpacked bytes are available in AL and BL.

Note The operation and practical use of the SAR instruction will be discussed in Section 4.7.4, where signed number arithmetic is dealt with.

3.6.2 | Rotate Instructions


Rotate instructions have the same format as the shift instructions. There are two types of rotate – rotate through carry or without taking the carry bit as a part of the rotate act.

ROL – Rotate Left.

Usage: ROL dest, count.

This instruction rotates the bits in the destination to the left ‘count’ times with all data pushed out at the left re-entering on the right. The Carry flag will contain the value of the last bit rotated out.

Modifies Flags: CF OF



ROL SI, CL


;rotate left the word in SI, by the count in CL

ROL BYTE PTR [DI][BX], 1

;rotate left the byte pointed by EA = [DI + BX] once

ROR – Rotate Right

Usage: ROR dest, count



This instruction rotates the bits in the destination to the right ‘count’ times with all data pushed out at the right side re-entering on the left. The Carry flag will contain the value of the last bit rotated out.

Modifies Flags: CF OF

ROR AL, 1

;rotate right once the byte in AL

ROR DX, CL

;rotate right the word in DX the count in CL

Let us examine what Example 3.26 will do.

Example 3.26

```
.MODEL TINY
.CODE
.STARTUP
MOV BX, 5634H
MOV CL, 8
ROL BX, CL
MOV AL, 0C6H
MOV CL, 4
ROR AL, CL
.EXIT
END
```

In the first instance, there is a 16-bit register BX which contains the value 5634H. By using ROL, and rotating 8 times, the value in BX changes to 3456H. Similarly with AL = C6H, using ROR with a count of 4, AL will have the new value of 6CH. Thus, we can switch exchange nibbles or bytes using the rotate instructions with the appropriate count in CL. Note that, for this application, it does not matter if ROL or ROR is used.


It is easy to check that Example 3.25 can be re-written replacing the SHR instruction with either the ROL or ROR instruction.

RCL – Rotate Through Carry Left.

Usage: RCL dest, count.

This instruction causes the left most bit (LSB) to enter the Carry flag, and the CF enters through the left end (MSB). Thus, due to one shift operation, the MSB enters the CF, and the CF gets into the LSB position.

Modifies Flags: CF OF



RCL BL, CL


;rotate left through carry, BL, by the count in CL

RCL CX, 1

;rotate left through carry the word in CX once

RCR – Rotate Through Carry Right.

Usage: RCR dest, count.



Modifies Flags: CF OF

RCR BYTE PTR [SI], 1

;rotate right through carry, the byte pointed by SI, once

RCR WORD PTR [DI][BX], CL

;rotate right through carry the word pointed by EA = SI + BX, by the count specified in CL

Example 3.27

Find the values in the destination for each line of this program segment.

```
STC
MOV AX, 5485H
RCR AL, 1
MOV CL, 03
RCL AX, CL
MOV CL, 05
ROR AX, CL
ROL AX, 1
```

Solution

Instructions

STC

MOV AX, 5485H

RCR AL, 1

Result after execution

CF = 1

AX = 0101 0100 1000 0101

AL = 1100 0010

MOV CL, 03	CL = 03
RCL AX, CL	AX = 1010 0110 0001 0101
MOV CL, 05	CL = 5
ROR AX, CL	AX = 1010 1101 0011 0000
ROL AX, 1	AX = 0101 1010 0110 0001

KEY POINTS OF THIS CHAPTER

- BIOS and DOS function calls facilitate the use of input and output devices in programming, making it interactive and more fun.
- The instruction set of 8086 can be divided into groups based on functionality.
- MOV is the most frequently used instruction and it is part of the data transfer group.
- LEA is an instruction that copies the offset of a memory address to a 16-bit register. It can be used to make the register be a pointer to an array of data items in memory.
- PUSH and POP can be used only with the stack segment
- The SS register contains the upper 16 bits of the lowest (base) address in the stack segment. SP contains the offset (with respect to the base address) of the highest address in the stack.
- Branch instructions cause the program sequence to change.
- JUMP is an important branch instruction which can be far or near, and unconditional or conditional.
- Jumps using the direct mode of addressing are re-locatable and 'relative'.
- LOOP is a jump instruction which causes branching as well as decrementing a count in CX. A LOOP instruction can be conditional, as well.
- There are 20 arithmetic instructions for 8086, of which a few specifically cater to signed numbers, a few to BCD numbers, and a few to ASCII numbers.
- The ADC (Add with carry) and SBB (Subtract with borrow) instructions find use in multi-byte additions and subtractions.
- The INC and DEC instructions do not affect the Carry flag.
- The CMP (Compare) instruction subtracts the source from the destination and sets/resets flags, but neither the source nor the destination is altered.
- MUL is the mnemonic for unsigned multiplication. For a byte by byte multiplication, one operand is implied to be in AL and the product in AX. For a word by word multiplication, one operand is to be in AX and the product will be in DX-AX.
- DIV is the mnemonic for unsigned division. For a word by byte division, the dividend is to be in AX. After division, the quotient will be in AL, and the remainder in AH. For a double word by word division, the dividend is to be in DX-AX. Then the quotient will be in AX and the remainder in DX.
- If the quotient is too large to fit in the register assigned for it, a 'divide by zero' error will be generated.
- The logical instructions of 8086 are AND, OR, XOR, NOT and TEST.
- The TEST instruction does logical ANDing, but only the flags are affected.
- Shifting of a data can be done left or right, and also by as many positions as required.
- Arithmetic shifting is meaningful for signed numbers, while logical shifting pertains to unsigned arithmetic.
- Rotation can be done left or right and can be direct or through the Carry flag.

QUESTIONS

1. Distinguish between the top-down and bottoms-up approach in programming.
2. What is the use of DOS function calls in assembly programming?
3. How many operands do each of the following instructions have?
 - a) ADD
 - b) ADC
 - c) INC
4. Which conditional flag is not affected by the DEC instruction?
5. If AX = 5600H and BX = 0C07H, find the contents of CX and DX after the following three sets of instructions.
 - a) PUSH AX
PUSH BX
POP CX
POP DX
 - b) PUSH AX
PUSH BX
POP DX
POP CX
 - c) PUSH BX
PUSH AX
POP CX
POP DX
6. Why is the stack pointer called the TOP OF STACK?
7. Distinguish between a near and a far jump.
8. Specify two ways of specifying an unconditional jump in the indirect mode of addressing.
9. Why does the far jump have five bytes of instruction length?
10. Name a conditional jump instruction that does not test a conditional flag.
11. What are the conditions under which a LOOPZ instruction exits looping?
12. What does the instruction STC do?
13. Can the ADD instruction use CS as a destination?
14. The instruction INC [SI] seems ambiguous. Why? How can the ambiguity be corrected?
15. What does the instruction CMP AX, BX do?
16. What is the role of OF and CF in multiplication?
17. What happens if the DIV instruction is used for a dividend of 1938H and divisor of 05?
18. What is meant by masking? Give an example of how it is used.
19. How can multiplication and division be achieved by shift operations?
20. What is the difference between the instructions RCL and ROL?

EXERCISE

1. Write a program to print on two separate lines, the messages 'Hi, folks' and 'I am Samantha'.
2. Enter a character through the keyboard without echoing it on the video monitor. Compare it with a number stored in memory and display the message 'EQUAL' or 'NOT EQUAL' as the case may be.

3. Display each character of the word PRADO on a separate line.
4. Enter 'N' characters through the keyboard without echo. Save it in memory as an array. Then display it as a character string.
5. Indicate what is wrong (if anything is wrong) with each of these instructions.
 - a) MOV CX,DL
 - b) ADD DATA1, 0978H
 - c) MOV BYTE PTR [SI][BX], DX
 - d) MOV 045FH, AX
 - e) MOV [DX], AL
 - f) MOV [SI][CX], AX
 - g) MOV DS, 3453H
6. Write a program that adds two quad words and stores the result in memory.
7. Write a program that subtracts two double words and stores the result in memory.
8. Find the status of the CF and ZF flags after the execution of the following sets of instructions.
 - a) MOV AX, 9078H
 - b) CMP AX, 0C089H
 - c) XOR AL, AL
 - d) MOV AL, 29H
 - e) CMP AL, 0
9. Find the biggest number in an array of
 - a) bytes
 - b) wordsDisplay the biggest number in each case, as a decimal number.
10. There are 10 unsigned bytes stored in memory. Arrange these bytes in
 - a) ascending and,
 - b) descending order.
11. Find the average of 20 bytes stored in memory. Display the average as a decimal number.
12. Add the sum of the first 20 natural numbers. Display the sum.
13. Write a program that counts the number of 1s in a binary number.
14. Write a program to add two $N \times N$ matrices.
15. Find the number of times a particular character is present in a character string stored in memory.

4 PROGRAMMING CONCEPTS-III



In this chapter, you will learn

- To use string instructions for various applications.
- The concept and use of procedures and call instructions.
- To distinguish macros from procedures and learn to write and use macros.
- To write programs to convert between the commonly used number formats.
- Signed number arithmetic using 8086 instructions.
- To write programs using the high level language constructs of MASM.

4.1 | String Instructions

The 8086 has a set of instructions for handling blocks of data in the form of bytes or words. They are called 'string' instructions. A string is an array of data of the same type – for example, a character string or a byte string. Table 4.1 gives the list of string instructions/prefixes which are used in string manipulations. The usefulness of string instructions can be seen when in the memory, data has to be moved, searched or compared in blocks.

Consider the case of 100 (say) words or bytes in a particular memory area that is to be moved to another memory area. This can very well be done using pointer registers and looping using a counter. However, this whole process can be automated with lesser number of instructions if we use string instructions. Similarly, we may need to compare two blocks of data for equality, or search a data block for a particular data. In all these cases, string instructions make our task easier and our code shorter. However, before using these instructions, we have to include a few initialization steps in our code.

Pre-requisites for Using String Instructions

- Two segments are to be defined i.e., the data segment and the extra segment. This means that the corresponding segment registers DS and ES have to be initialized and used. The data segment is the source segment and the extra segment is the destination segment.
- The DI registers and SI registers should act as pointers to the data segment and extra segment respectively. This means that, initially, SI should contain the address (offset) of the first location in the data segment. Similarly, DI should contain the address (offset) of the first location in the extra segment.

Table 4.1 | List of Instructions/Prefixes Used in String Operations

Sl No.	Instruction format	Function performed	Flags affected
1	MOVSb/MOVSW	Move byte or word string	None
2	CMPSb/CMPSW	Compare byte or word string	AF CF OF PF SF ZF
3	SCASb/SCASW	Scan byte or word string	AF CF OF PF SF ZF
4	LODSb/LODSW	Load byte or word string	None
5	STOSb/STOSW	Store byte or word string	None
6	CLD	Clear direction flag	DF
7	STD	Set direction flag	DF
8	REP (Prefix for a string instruction)	Repeat execution of string instructions while CX is not 0	None
9	REPE/REPZ (Prefix for a string instruction)	Repeat execution of string instructions while CX is not 0 and while Zero flag is set	None
10	REPNE/REPNZ (Prefix for a string instruction)	Repeat execution of string instructions while CX is not 0 and while Zero flag is not set	None

Note In string instructions, 'SB' stands for 'string byte' and 'SW' for 'string word'.

- iii) There is a control flag called the direction flag which is used exclusively for string operations. Its purpose is that in string operations, if the flag is set, the pointer registers get automatically decremented and if reset, the reverse occurs. So whenever string instructions are being used, the direction flag (DF) should be set or reset depending on the direction the addresses are to be modified after each operation.
- iv) The counter CX should be loaded with the count of the number of operations required.

4.1.1 | The MOVS Instruction

Example 4.1 shows the use of the string instruction MOVSb with the necessary pre-requisites incorporated. Let us examine how this has been done. The problem in hand is to transfer 10 bytes (the character **) from an area of memory designated as DATA1 to an area named DATA2. Thus, the source has to be the data segment, and the destination, the extra segment. How do we incorporate the extra segment?

If we use the full segment model, we can define two segments – the data segment (DAT) and the extra segment (EXTR). See how it is done. When both the segments registers are

defined, the segments get defined. Then the program instructions cause the data in the data segment to be copied to the extra segment. Example 4.1a uses the full segment model for this.

Example 4.1a

```

DAT                SEGMENT
DAT1 DB    10 DUP('')                ;store '' in 10 locations
DAT                ENDS

EXTR               SEGMENT
DAT2 DB    10 DUP(0)                ;allocate 10 locations
EXTR               ENDS

COD                SEGMENT
                ASSUME CS:COD, DS:DAT, ES:EXTR
                MOV AX, DAT
                MOV DS, AX                ;initialize the DS register
                MOV AX, EXTR
                MOV ES, AX                ;initialize the ES register
                LEA SI, DAT1                ;point SI to Source
                LEA DI, DAT2                ;point DI to destination
                MOV CX, 10                ;load count in CX
                CLD                        ;clear Direction flag
REP  MOVSB                ;move the byte string
                MOV AH, 4CH                ;AH = 4CH to return to DOS
                INT 21H
COD                ENDS
                END

```

In Example 4.1a, two segments DAT and EXTR are defined and two segment registers DS and ES are initialized. The addresses SI and DI point to the source data and destination data addresses and DF is cleared for auto incrementing the pointer registers.

Now let us do the same program using the simplified model. The small model can have only one data segment, so how can we have an 'extra segment'? The solution is to have both data areas in the data segment itself, but after the data segment is defined, copy the value of DS into ES – this makes the assembler believe that the extra segment is available – though it is the same as the data segment. In this example, the destination data area is chosen to be at org 0200H. This is done just to space the source and destination data areas and is not mandatory. Thus, the first pre-requisite of having both the data and extra segments is satisfied.

Next, the SI register is made to point to the address DATA1 and the DI register to DATA2. Then the counter register is loaded with 10, the count of data transfer operations required. The instruction CLD is used to clear the direction flag (DF). This causes the value of SI and DI to be incremented after each move operation. Here the data is in the form of bytes. Hence, SI and DI are incremented only by one each time. The string instruction used here is MOVSB (prefixed by REP). This causes the data in the location pointed by SI to be moved to the location pointed by DI. After each such move operation, two actions occur. One is that CX decrements by 1 and the other is that the pointer registers are incremented. This sequence continues until CX = 0. Example 4.1a and 4.1b illustrate these steps.

Example 4.1b

```

                .MODEL SMALL                ;select small model
                .DATA                      ;start data segment
DAT1 DB 10 DUP('*')                      ;store '*' in 10 locations
      ORG 0200H                          ;define destination area
DAT2 DB 10 DUP(?)                        ;set aside 10 locations
                .CODE
                .STARTUP
      MOV AX, DS                          ;move DS into AX
      MOV ES, AX                          ;move AX to ES
      LEA SI, DAT1                        ;point SI to Source
      LEA DI, DAT2                        ;point DI to destination
      MOV CX, 10                          ;load count in CX
      CLD                                ;clear Direction flag
      REP      MOVSB                       ;move the byte string
                .EXIT
                END

```

For the above program, if the data involved is in the form of **words**, the only change would be to replace the line `REP MOVSB` with `REP MOVSW`. In this case, `CLD` will cause `SI` and `DI` to be incremented by 2.

4.1.2 | The CMPS Instruction

Now, let us use the next string instruction – `CMPSB/CMPSW`. This instruction is for string comparison, byte by byte or word by word as the case may be. The conditional flags are modified according to the result of the comparison. String comparison has to be accompanied by the use of the conditional `REP` prefix. **Since string comparison checks only for equality, the Zero flag is made use of automatically.** Example 4.2 compares two character strings (of length six) which are saved in the data segment. One of two messages is to be displayed, depending on whether `WRD2` is the same as `WRD1`.

Example 4.2

```

                .MODEL SMALL
                .DATA
WRD1 DB "SAMSON"                          ;store the first word
WRD2 DB "SAMRON"                          ;store the second word
GUD  DB "SAME$"                           ;store the message for equality
BAD  DB "NOT SAME$"                       ;store the message for inequality
                .CODE
                .STARTUP
      MOV AX, DS
      MOV ES, AX                          ;copy DS to ES
      LEA SI, WRD1                        ;point SI to source
      LEA DI, WRD2                        ;point DI to destination
      MOV CX, 6

```

```

        CLD                      ;clear direction flag
REPE    CMPSB                    ;repeat comparison if equal (ZF = 1)
        JNZ MSG                  ;if ZF is reset, go to MSG
        LEA DX, GUD              ;point DX to the first message
        JMP DISP
MSG:    LEA DX, BAD              ;point DX to second message
DISP:   MOV AH, 09               ;use the function for string display
        INT 21H                 ;use DOS interrupt 21H
        .EXIT
        END

```

Here SI and DI point to the two character strings, CX stores the count and the direction flag is cleared for auto incrementing the address pointers. REPE is used as the prefix for comparing the string bytes. REPE CMPSB means that the corresponding string bytes of the source and destination are compared as long as they are equal. Equality also implies that the Zero flag is set. The comparison is stopped as soon as an **inequality** is encountered. One of the following two conditions cause the comparison loop to be exited.

- i) An inequality is seen i.e., the Zero flag gets reset.
- ii) Normal exiting when CX = 0.

In the example, the two character strings are shown to be different, which causes the message 'NOT SAME' to be printed. If WRD1 and WRD2 are the same, the message 'SAME' is printed.

4.1.3 | The SCAS Instruction

The SCAS instruction scans a byte or word string to ascertain the presence of a specific byte or word. This specific data is loaded into AL or AX. **The string which is to be scanned has to be in the extra segment, and is to be pointed by DI. This is mandatory and cannot be overridden.** Example 4.3 illustrates a typical scenario for the use of this instruction. Suppose that an ASCII string labeled LIST is stored in memory. The program searches the string for the presence of the ASCII character 'S'. DI addresses the string to be searched. The string to be searched is placed in the Extra segment, by making the ES have the same value as DS. Thus the extra segment is the same as the data segment. In this example, the SCASB instruction has an REPNE prefix (repeat while not equal) prefix. The REPNE prefix causes the SCASB instruction to repeat until either the CX register reaches 0, or until an equal condition is indicated (ZF = 1) as the outcome of the SCASB instruction's comparison operation. The scanning of the string continues as long as 'S' is not found in the string. If 'S' is found, the search is terminated immediately.

Example 4.3

```

        .MODEL SMALL
        .DATA
LIST    DB "ANBHTSHFGTRIUUJEHNPBGDVBTH"
PR      DB "CHARACTER PRESENT$"
NOTP    DB "CHARACTER NOT PRESENT$"

```

```

.CODE
.STARTUP

MOV AX, DS           ;copy DS to AX
MOV ES, AX           ;copy DS to ES
LEA DI, LIST +26     ;point SI to the last character
STD                  ;auto-decrement
MOV CX, 26           ;load count in CX
MOV AL, 'S'          ;load the character in AL
REPNE SCASB          ;repeat if not equal
JZ MSG
LEA DX, NOTP
JMP DISP
MSG:  LEA DX, PR
DISP: MOV AH, 09
      INT 21H
      .EXIT
      END

```

In the above, a variation (from the previous examples) has been used. DI has been made to point to the last character of the string, and the direction flag is set, such that the address auto decrements after each scanning operation. Refer to Example 3.18, where the same problem has been done without using string instructions.

4.1.4 | The STOS and LODS Instructions

i) STOS

The STOS instruction is the mnemonic for ‘storing’ a string in memory. As such, we need to define a memory area in which ‘storing’ is to be done. **This memory area is defined to be the extra segment, and it is addressed by DI as it is the destination segment.** The data to be stored is placed in the AL or AX register. An area in memory can be filled with the required data, with this instruction. Example 4.4 illustrates the use of the STOS instruction. Here AX is loaded with 0001, and this is moved to a location named AREA, which has been defined as an array of 50 words. This array is filled with the word 0001 by the STOSB instruction, using the REP prefix and the counter CX initialized with 50. Since STOS is specified for the extra segment, ES register has to be initialized. Notice that, as done earlier, a data segment is first defined (by the .DATA directive) and the DS register content is then moved to ES. Note also that the memory AREA is addressed using the DI register. **SI register cannot be used.**

Example 4.4

```

.MODEL SMALL           ;select small model
.DATA                 ;start data segment
AREA DW 50 DUP(?)     ;define AREA
.CODE                 ;start code segment
.STARTUP              ;start program
MOV AX, DS            ;copy DS to AX
MOV ES, AX            ;copy AX to ES

```

LEA DI, AREA	;address AREA with DI
MOV CX, 50	;load count in CX
CLD	;clear direction flag
MOV AX, 0001	;load AX with 0
REP STOSW	;repeat until CX = 0
.EXIT	;exit to DOS
END	;end of program

It is obvious that the STOS instruction is used to fill up an area in memory with the same data.

ii) LODS

The last of the string instructions is LODS. This is an instruction for ‘loading’. Loading always means the act of taking data from memory and putting it into a register. Here, the source memory is the data segment and the pointer to it is SI. The data segment is the source segment and the destination register is the AL or AX register. There is no sense in using the REP prefix for the LODS instruction as data can be loaded to AL/AX only once.

4.2 | Procedures

In high level languages (i.e., C, C++) you might have come across ‘functions’. A function is a program which does a specific task. When this task is to be done repeatedly, the function is used again and again. When a ‘main’ program considers this as a subsidiary task, the function (or sub routine) is ‘called’ whenever its service is required. This is applicable to assembly language programming also. Borrowing Intel’s terminology, we shall however, call it a ‘procedure’.

Figure 4.1a shows two cases. One is the case of a main program calling many different procedures. The second is a program calling the same procedure repeatedly. Thus there is a main program which can call a procedure anywhere in its body. The former is the ‘calling program’ and the latter is the ‘called program’. The main program is in a particular code segment and the




Figure 4.1 | a Main program calling different procedures b Main program calling the same procedure repeatedly

procedure may be written in the same code segment (in which case it is a ‘near’ procedure), or it may be in a different segment (‘far’ procedure).

Let us now see the sequence of actions taken by a processor when using a procedure. For now, let us assume that the procedure is a ‘near’ procedure. In the course of the action of execution of the main program, a CALL instruction is encountered. This signals that a procedure written elsewhere has to be executed. This is essentially a branching operation, as the normal program execution sequence is disturbed. At the time the CALL instruction is being executed, the IP (instruction pointer) will be pointing to the next instruction in the main program. The steps taken by the processor automatically are:

- i) It saves the current IP content on the stack (this is the ‘return’ address for coming back to the main program after executing the procedure).
- ii) The CALL destination (specified in the CALL instruction) will be the address of the procedure. The IP is now loaded with this address and execution proceeds from that location.
- iii) The procedure is executed until a RET (return) instruction in the procedure is encountered.
- iv) Then, the old value of the IP is retrieved from the stack and control returns to the main program at the return address.

Fig 4.2 shows a CALL encountered in a main program. Then the procedure MULT is taken up for execution. At the end of the procedure, there is the RET instruction, executing which causes control to return to the main program. The instruction after the CALL instruction will now be taken up.

4.2.1 | Writing a Procedure

Like any assembler, MASM also has specifications for writing a procedure. The procedure should begin with the procedure name followed by the directive PROC. We also use the directives NEAR or FAR which specify the ‘type’ of the procedure. The procedure should end with the procedure name and ENDP. Example 4.5 shows a program, which enters 10 single digit numbers through the keyboard, finds their squares and stores the squares in memory. Only byte locations have been allocated in memory for the squares as the square of the highest single digit number is 81, which will fit into a byte space. In the procedure, the register




Figure 4.2 | Call and return

assigned for the product is AX. However, the product will occupy only AL. Hence, only the value in AL is moved to memory.

Now, let us examine the salient features of this program.

- i) A procedure named SQUARE is used to calculate the squares.
- ii) Entering data from the keyboard and storing the squares in memory is done in the main program.
- iii) The procedure is in the same code segment. Hence it is designated as 'near'.
- iv) The procedure ends with a RET (return) instruction which takes control back to the main program.
- v) The .EXIT statement is the last statement in the main program. We know that it gets converted to the instructions for returning control to DOS. Thus, it is the last instruction which gets executed.
- vi) The last line in the code is the END directive which tells the assembler to stop reading.
- vii) We know that the stack segment is necessary for the working of a procedure. For now, we are not defining any stack segment but using the stack defined automatically by DOS. When we need more stack size, we will define our own stack.

Example 4.5

```

.MODEL SMALL
.DATA
NUMBERS DB 10 DUP(0)    ;allocate space for the squares
.CODE
.STARTUP
MOV CX, 10               ;load count in CX
LEA SI, NUMBERS          ;let SI point to NUMBERS
REPEA: MOV AH, 01         ;key board input
INT 21H
SUB AL, 30H              ;convert from ASCII to binary
CALL SQUARE              ;call the procedure for squaring
MOV [SI], AL             ;store the square in memory
INC SI                   ;increment pointer
LOOP REPEA               ;decrement CX and repeat if CX!=0
.EXIT                    ;return to DOS prompt

SQUARE PROC NEAR         ;define the procedure
MOV BL, AL               ;move multiplicand to BL
MUL BL                   ;multiply
RET                       ;return to main program
SQUARE ENDP              ;define the end of the procedure
END

```

In Example 4.5, the square of the numbers has been calculated and saved in memory. It may be interesting to display them on the monitor. This can be done by referring to Example 3.22. The program for converting any hexadecimal number to binary and displaying can be made another procedure and called here.




Figure 4.3 | Format of the direct near CALL instruction

4.2.2 | Call and Return Instructions

Now that we have seen the general format and structure of a program with procedures, let us examine the CALL instruction in greater detail. This instruction is obviously a ‘branch’ instruction because it causes control to move to a different address (called the target or destination), which may be in the same or a different code segment.

4.2.2.1 | Intrasegment or ‘Near’ Call

i) Direct CALL

Usage: CALL label (see Fig. 4.3).

The direct call is like a direct jump instruction, and is three bytes long. It is relative and the destination can be –32,768 bytes to +32,767 bytes from the address of the instruction following the call (this will be the current content of the IP). This means that the offset can be a 16-bit signed number. When this call is executed, the new value of IP = old IP + offset, where the second and third bytes of the instruction give the offset. The assembler must calculate the offset during its first pass. Example 4.5 uses this type of the call instruction – the direct near call.

```
CALL NEW_WAY      ;using the label, the assembler calculates the offset
CALL MULTI        ;calls a procedure named MULTI
```

ii) Indirect CALL

Usage: CALL reg16, CALL [reg16]

In this case, the destination is specified in a 16-bit register or in a memory location pointed by a register. This is not a ‘relative’ call. The content of the referred register or memory location is loaded into IP for using the procedure.

```
CALL BX            ;the procedure’s address is in BX
CALL WORD PTR[BX] ;the address of the procedure is in memory and
                  ;is pointed by BX
```

The following program illustrates the use of an indirect near call. Example 4.6 illustrates the use of the CALL register instruction to call procedures that begin at offset addresses ENT and DISP. (These calls could also call the procedures directly as CALL ENT and CALL DISP.) The offset address ENT is placed in BX and DISP in SI, and then the CALL BX and CALL SI instructions call these procedures. The ENT procedure allows keyboard entry with echo. If an upper case letter is entered by ENT, the calling program converts into lowercase and the DISP procedure displays the corresponding lowercase letter.

Example 4.6

```
.MODEL TINY      ;select tiny model
.CODE           ;start code segment
```

```

        .STARTUP                ;start program
        LEA BX, ENT              ;offset of ENT in BX
        LEA SI, DISP             ;offset of DISP in SI
        CALL BX                  ;call the ENT procedure
        ADD AL, 20H              ;upper case to lower case
        MOV DL, AL               ;transfer lower case to DL
        CALL SI                  ;call the DISP procedure
        .EXIT

ENT PROC NEAR
    MOV AH, 01                  ;procedure for entering key
    INT 21H
    RET
ENT ENDP

DISP PROC NEAR
    MOV AH, 02                  ;procedure for displaying
    INT 21H
    RET
DISP ENDP
END
```

4.2.2.2 | Intersegment or Far Call

Direct Far Call A far call is an intersegment call, which means that the destination address is in a different code segment. This will be a 5-byte instruction, the first byte being the opcode, the second and third bytes being the new value of IP, and the fourth and fifth, the new values of CS. This is not a relative call. When the procedure is called, the IP and CS values are replaced by the corresponding values in the call instruction as shown in the Fig. 4.4.

Indirect Far Call For an indirect call, the destination address is not in the instruction – rather, it is in a register or memory. For a far call, four bytes are needed to specify a destination. Obviously, a register cannot specify it. Hence the four bytes needed to specify a destination are stored in memory and pointed by a register. As an example CALL DWORD PTR [SI] can be a far call instruction. [SI] and [SI + 1] gives the new value of IP and [SI + 2] and [SI + 3] gives the new value of CS.

Far calls and procedures are specified using the ‘far’ directive when defining the procedure. In Chapter 5, we will see how the EXTRN directive is used in this context.

4.2.3 | The RET Instruction

i) Usage: RET
When a procedure is called, the current value of IP is pushed on to the stack. This naturally means that when the procedure has ended, control should go back to the main program at the point



Figure 4.4 | Format of the far jump instruction

where it had branched off. This means that the value of IP (and CS for a 'far' call) will have to be retrieved. This is done by the 'return' instruction which is the last instruction in the procedure. The execution of RET causes the return address to be popped from the stack to IP or IP and CS. (Whenever a procedure is defined, it is known whether it is a far or near procedure. From this, it is determined whether the stack has saved just the old IP value or both IP and CS.)

ii) Usage: RET n

This is another form of the RET instruction. This adds the number 'n' to the stack pointer (SP) after the return address has been popped off the stack, on return from a procedure. We will see the use of this in Example 4.10.

4.2.4 | The Use of the Stack in Procedure Calls

The stack is used in procedure calls for saving the return address. There is another context when a stack is necessary. The main program may be using a number of general-purpose registers. Since the number of registers is limited, what happens if the procedure too needs a few of those registers? If the procedure uses the same registers, obviously their content will be changed. To avoid this problem, the contents of these registers (including the flag register) may be pushed on to the stack before calling the procedure, and popped back from stack on returning from the procedure. This allows both the main program as well as the procedure, the use of the same registers without losing their content. It is up to the programmer to write these push and pop instructions (either in the main program or in the procedure). The only point to remember would be to push and pop in reverse order. However, now MASM 6.x has a construct 'USES' which automatically pushes the specified registers onto the stack. The popping will be automatic on returning from the procedure. Example 4.7a illustrates this.

Example 4.7a

```

        .MODEL SMALL
        .DATA
        -----
        .CODE
        .STARTUP
            ..
            ..
            ..
        CALL MULTI
            ..
            ..
            ..
MULTI    PROC NEAR USES DX CX
        MUL BL
        MOV DX, 4509H
        ADD AX, DX
        MOV CX, 0608H
        ADD AX, CX
        RET
MULTI    ENDP
        END

```

Now see Example 4.7b. The main program is not shown. What is illustrated is that it uses registers DX and CX. These registers are used in the procedure MULTI as well. The USES construct at the start of the procedure will be converted to PUSH DX and PUSH CX when the procedure is called, and to POP CX and POP DX when returning from the procedure. This can be verified by ‘debugging’ the program or by using the directive .list all. Then the list file will show the listing of the procedure as shown below. Note the push and pop DX and CX which we have not written specifically in the program, but have been caused by the USES construct.

Example 4.7b

```

002B          MULTI PROC NEAR USES DX CX
002B 52      *      push    dx
002C 51      *      push    cx
002D F6 E3          MUL BL
002F BA 4509      MOV DX, 4509H
0032 03 C2      ADD AX, DX
0034 B9 0608      MOV CX, 0608H
0037 03 C1      ADD AX, CX
                        RET
0039 59      *      pop     cx
003A 5A      *      pop     dx
003B C3      *      ret     00000h
003C          MULTI ENDP
                        END

```

4.2.5 | Passing Parameters To and From Procedures

When procedures are called, they have to be given data on which to work on, and then these procedures will have to return results to the calling program. There are various ways of doing this. One way would be to place it in registers which can be accessed by the main program as well as the procedure. Another way would be to place data in memory and access it by name, or by using pointers. Still another method is to use the stack for this. Data or addresses made available to both the calling and called programs are called ‘parameters’ and now we will discuss various ways of ‘passing parameters’ to and from procedures.

4.2.5.1 | Passing Parameters Through Registers

In this case, data is placed in registers by the main program, and these registers are used by the procedure. Example 4.8 is a program which calculates and places in memory the Nth term of an arithmetic progression. The formula for the Nth term of an A.P. is $A + (N - 1) D$, where A is the first term and D is the common difference. Here the values of A, N and D are placed in the data segment from where the main program takes them, loads them into registers and passes the values to the procedure. The procedure does the computation and passes the result to the main program through register BX.

Example 4.8

```

.MODEL SMALL
.DATA

```

```

A DB 4                ; A = 4
D DB 5                ; D = 5
N DB 20               ; N = 20
NTH DW?

        .CODE
        .STARTUP
MOV BL, A              ;copy A to BL
MOV BH, 0              ;BX = 0 to have A in BX as a word
MOV AL, D              ;copy D to AL
MOV CL, N              ;Copy N to CL
DEC CL                ;decrement CL to get N-1
CALL NTH_TERM_AP      ;call the procedure
MOV NTH, BX            ;save the nth term in memory
        .EXIT

NTH_TERM_AP PROC NEAR
    MUL CL              ;multiply (N-1)x D, product in AX
    ADD BX, AX          ;Nth term = A+(N-1)D will be in BX now
    RET
NTH_TERM_AP ENDP
END

```

4.2.5.2 | *Passing Parameters Through Memory*

In Example 4.9, the data which is used by the procedure is accessed from memory through the pointer register BX, and data is put back in memory in the same way. Example 4.9 is a program for arranging a set of numbers (stored in memory) in descending order. The method involves comparing numbers pair-wise and repeating this $N - 1$ times, if there are N numbers to be sorted. The steps are:

- i) Suppose the numbers are 6, 8, 34, 0 and 67. The first two numbers are compared. If the first number is greater than the second, no change is made to the ordering. Otherwise exchange them, such that the bigger number comes first in the pair. For our set, now the ordering becomes 8, 6, 34, 0 and 67.
- ii) Next, the second and third numbers (i.e., 6 and 34) are compared. The new ordering is 8, 34, 6, 0 and 67.
- iii) This is repeated $N - 1$ times, so that due to pair-wise ordering, the bigger number comes first in the new ordering of pairs. In Example 4.9, the procedure `one_set` does this.
- iv) This sequence of steps is repeated $N - 1$ times, such that we get all the numbers sorted in descending order, replacing the unsorted array.

Example 4.9

```

        .MODEL SMALL
        .DATA

N      WORD 0009H
NUMS  BYTE 10H, 08, 89H, 78H, 67H, 07, 2, 99H, 0FEH
        .CODE

```

```

        .STARTUP
        LEA BX, NUMS      ;BX to point to the number array
        MOV CX, N          ;move the array size to CX
        DEC CX             ;CX must contain N-1
        MOV SI, BX         ;SI to have copy of starting address
AG:      MOV DX, CX         ;DX = N-1
STRT:    CALL ONE_SET      ;the procedure for pair-wise ordering
        DEC DX             ;keep a count of the numbers of calls
        JNZ STRT           ;repeat until DX = 0
        MOV BX, SI         ;put the starting address back in BX
        LOOP AG            ;loop back until CX = 0
        .EXIT

ONE_SET PROC NEAR
        MOV AL, [BX]       ;get the first number in AL
        INC BX             ;increment the pointer
REPEA:   MOV AH, [BX]      ;get the second number in AH
        CMP AL, AH         ;compare the numbers
        JB EXCH            ;jump if first number is below
BACK:    RET              ;otherwise return
EXCH:    MOV [BX], AL       ;put AL in second address
        MOV [BX-1], AH     ;put AH in first address
        JMP BACK          ;then return
ONE_SET ENDP              ;end the procedure
        END

```

In the above program, nine numbers (bytes) are stored in the data segment and sorted. Any number of bytes can be sorted this way, limited only by the maximum value of N. For sorting in ascending order, only one instruction in the procedure need be changed (which one?). After sorting is done, if we write a procedure for converting hexadecimal numbers to ASCII, the sorted numbers can be displayed as well.

4.2.5.3 | Passing Parameters Through the Stack

There may be reasons where the stack, which is an area in the memory, can be used to store data. The procedure can take data from the stack for its computation. Example 4.10 is a simple example which also elaborates how the register BP is used to access data in the stack. Here, four words which are in the data segment are pushed on to the stack by the main program. The procedure accesses the data using the indexing features of BP. Remember that BP is a register associated (by default) with the stack segment, and that BP cannot be used without an offset. Example 4.10 calculates $(A + B) - (E + D)$ where A, B, E and D are words stored in data memory. The result is also to be saved in the data segment.

Note We do not use C as a label because 'C' is a reserved word in MASM.

Example 4.10

```

.MODEL SMALL
.STACK 100H

```

```

        .DATA
A DW 0987H           ;data values in the data segment
B DW 678H
E DW 0A5EH
D DW 0034H
RESULT DW?
        .CODE
        .STARTUP
        PUSH A           ;push to stack
        PUSH B
        PUSH E
        PUSH D
        CALL COMPUTE     ;call the procedure
        MOV RESULT, AX   ;save AX in memory
        .EXIT

COMPUTE PROC NEAR
        MOV BP, SP       ;copy SP to BP
        MOV AX, [BP + 10] ;move the top most stack value to AX
        MOV BX, [BP + 8]  ;move next word in stack to BX
        ADD AX, BX
        MOV CX, [BP + 6]  ;move the next word to CX
        MOV DX, [BP + 4]  ;move the next word to DX
        ADD CX, DX
        SUB AX, CX
        RET 8             ;return and add 8 to SP
COMPUTE ENDP
        END

```

Note The salient features of the program are:

- i) A stack of 100 bytes has been defined in the beginning.
- ii) The push instructions are in the main program.
- iii) In the procedure, the SP value is copied to BP.
- iv) The data in the stack is accessed by adding offsets to BP.
- v) The return instruction is RET 8. This causes SP to have the value it had before the four PUSH operations are executed. Otherwise, what would happen is that each time this program is run, the stack pointer decrements by 8 bytes (for the four PUSH operations). Thus the stack size gets reduced by that much for each procedure call which will finally cause the stack size to reduce to zero, thus causing system crash.
- vi) RET 8 adds 8 to SP after returning from the procedure. In effect, it erases the data that had been pushed on to the stack.

Refer to Figure 4.5. Let us say that before the first PUSH, the value of SP = 0120H. This figure shows the content of the stack after the four PUSH operations. The address of SP now is 0118H. When the CALL is executed, SP = 0116 H, because the content of IP is pushed on the stack. Before returning from the procedure, SP = 0116 H. RET 8 causes IP to be popped. So, SP = 0118 H. Adding 8 to this value of SP causes SP to become 0120 H again.

Stack address	Stack content	
011FH	09	A
011EH	87	
011DH	06	B
011CH	78	
011BH	0A	E
011AH	5E	
0119H	00	D
0118H	34	

Figure 4.5 | Stack operation for the PUSH and RET 8 instructions

It is important to be very clear about stack operations if you plan to use the stack. Any wrong alteration of the stack can cause a system crash.

Stack Overflow and Underflow

SP can have a maximum value of FFFFH. For each PUSH operation, the SP value decrements by 2, and in the limiting case, it can go to SP = 0000.

Any PUSH operation beyond this will cause a ‘stack overflow’. This creates a condition when there is no space in the stack for new data.

Stack underflow is the other case when POP operations cause SP to have values beyond the defined top of stack.

4.3 | Macros

The name of our assembler is ‘Macro assembler’. As such, we see that it is expected to be able to handle an item called ‘macros’. What is a macro? It is like an opcode – when used, it executes. A macro when called by name, executes the instructions that are listed under that name. Thus essentially, a macro is a short hand notation for a number of instruction lines. It essentially makes assembly language coding more readable and helps to avoid repetitive coding.

Usually, a macro, like a procedure, is defined for performing a specific function. However, the ‘overheads’ involved in invoking a procedure are not incurred here. A procedure call causes pushing and popping of addresses/data in stack. A macro when invoked just expands the code by putting in all the instructions corresponding to the called macro. It does not have to ‘call’ and ‘return’. Thus, when we write our code with macros, it may look small, but when assembling the code, each macro is replaced by the full set of instructions it consists of. Thus we can say that macros execute faster but the assembled code takes more memory. This is because a procedure is written only once in memory, but the macro statements are written as part of the code every time the macro is invoked.

4.3.1 | Writing a Macro

Whenever we have a set of code lines that we may use frequently, it could be convenient to make a macro out of it. A macro has the following format:

```
MACRO NAME MACRO [parameter list]
    Instructions (body of the macro)
ENDM
```


Let us think of a simple task that could be defined as a macro. If we frequently want to enter data through the keyboard, it could be written as a macro.

```

    ENTR MACRO      ;the macro for entering a character with echo
    MOV AH, 01H
    INT 21H
    ENDM

```

This macro does not have parameters to be passed to it. Another similar macro is for displaying a character as shown below.

```

DISP MACRO          ;the macro for displaying a character
    MOV AH, 02H
    INT 21H
ENDM

```

When writing programs using macros, make sure the macros are defined before they are used. Example 4.11 is a simple illustration of these principles.

Example 4.11

```

    .MODEL TINY
    .CODE
    .STARTUP

    ENTR MACRO      ;define the macro for entering a key
    MOV AH, 01H
    INT 21H
    ENDM

    DISP MACRO      ;define the macro for displaying a character
    MOV AH, 02H
    INT 21H
    ENDM

    ENTR            ;invoke macro ENTR
    MOV DL, AL      ;copy ASCII code of key from AL to DL
    DISP            ;invoke macro DISP

    EXIT
    END

```

When we run the above program, we get the pressed key displayed twice. Why? Now let us define and use a macro using parameters.

```

DISPLAY MACRO STRINGDAT
    MOV AH, 09
    LEA DX, STRINGDAT
    INT 21H

```

The above is a macro which displays the string whose address is loaded into DX. The parameter is 'STRINGDAT'. It is also called a dummy variable. When the macro is invoked, the actual name of the string to be displayed may be used. Example 4.12 illustrates these ideas.

Example 4.12

```

.MODEL SMALL
.DATA
MSG1 DB "MY NAME IS ANTONY GONSALVES$"
MSG2 DB "I AM SAMANTHA$"
.CODE
.STARTUP
DISPLAY MACRO STRING      ;define a macro for string display
MOV AH, 09
LEA DX, STRING
INT 21H
MOV DL, 0AH              ;DL = ASCII data for newline
MOV AH, 02
INT 21H
MOV DL, 0DH              ;DL = ASCII data for carriage return
MOV AH, 02
INT 21H
ENDM

DISPLAY MSG1              ;invoke macro for displaying MSG1
DISPLAY MSG2              ;invoke macro for displaying MSG1
.EXIT
END

```

In the above program, the `DISPLAY` macro is first defined. Then the macro is invoked twice with two different parameters `MSG1` and `MSG2`. Note that the macro not only displays a string, but also has instructions for displaying new line and carriage return. This ensures that after displaying one line, the cursor moves to the left side of the next line. Thus, in Example 4.12, the two messages are displayed on separate lines.

4.3.2 | Using the 'Local' Directive in Macros

Whenever macros use labels, they should be declared as local to the macro. Otherwise if the same macro is used many times, they will create assembly errors, as the label will correspond to different addresses each time. So, just after the macro name is declared, the labels must be written preceded by the word 'LOCAL'. Any number of labels can be declared in this way.

It would be a good idea for you to examine the list file of the following program to see how the label 'AGAIN' is being managed.

Example 4.13

```

.MODEL TINY
.CODE
.STARTUP
STAR MACRO N              ;define macro STAR with parameter N
LOCAL AGAIN               ;declare AGAIN to be a local label

```

```

        MOV CL, N           ;move count to CL
        MOV DL, '*'        ;move '*' to DL
AGAIN:  MOV AH, 02          ;AH = 2 for displaying a character
        INT 21H
        DEC CL             ;decrement the count
        JNZ AGAIN          ;repeat the displaying until CL = 0
        ENDM              ;end the macro

NEW     MACRO              ;define a macro for moving to next line
        MOV DL, 0AH
        MOV AH, 02
        INT 21H
        MOV DL, 0DH
        MOV AH, 02
        INT 21H
        ENDM

        STAR 5             ;invoke macro STAR with N = 5
        NEW               ;invoke macro NEW
        STAR 4             ;invoke macro STAR with N = 4
        NEW
        STAR 3
        NEW
        STAR 2
        NEW
        STAR 1
        NEW

        .EXIT
        END

```

Example 4.13 shows such a case. Two macros have been defined – STAR and NEW. The former contains a label AGAIN which has been declared as LOCAL. This macro causes the character '*' to be printed on a line N times. N is a parameter whose value is passed to this macro, when called. The macro NEW just moves the cursor to the left side of the next line. The output of Example 4.13 is as shown below.

```

*****
****
***
**
*
C:\masm6.14\BIN>

```

4.4 | Number Format Conversions

We know that computers do all their calculations using binary arithmetic, but we are accustomed to doing calculations in decimal form. We see numbers printed out in the decimal form, and enter data using the keyboard as decimal numbers. Our conclusion obviously is that through


```

        LEA SI, BCD           ;load the offset of BCD in SI
        LEA BX, ASC           ;load the offset of ASC in BX
        MOV CX, 4             ;count of BCD bytes
RPT:    CALL BCD_TO           ;call the conversion procedure
        INC SI                ;increment pointer
        INC BX                ;increment pointer
        LOOP RPT              ;continue until CX = 0
                                ;with the above loop, the
                                ;conversion is
                                ;done and only displaying is left
        MOV CX, 08           ;count of ASCII bytes
        DEC BX                ;to point to address of last
                                ;ASCII digit
STRT:   MOV DL, [BX]          ;move ASCII value to DL for displaying
        MOV AH, 02           ;function to display a character
        INT 21H
        DEC BX                ;point to next ASCII digit
        LOOP STRT             ;repeat until CX = 0
        .EXIT

BCD_TO  PROC NEAR USES CX AX
        MOV AL, [SI]          ;move first BCD byte to AL
        MOV AH, AL            ;copy it to AH as well
        AND AL, 0FH           ;mask the upper nibble of AL
        AND AH, 0F0H          ;mask the lower nibble of AH
        MOV CL, 04            ;count of shift operations
        SHR AH, CL            ;shift AH 4 times
        OR AX, 3030H          ;convert AX to two ASCII bytes
        MOV [BX], AL          ;move one ASCII value to memory
        INC BX                ;BX to point to next locations
        MOV [BX], AH          ;store the next digit
        RET
BCD_TO  ENDP
        END

```

Converting from ASCII to packed BCD is just a reverse of all these processes.

4.4.2 | BCD Calculations

i) Addition

DAA – Decimal adjust AL after addition.

Usage: DAA

For BCD, none of the digits should have a value greater than 9, but we know that a nibble can contain a value of 0FH. This causes problems when BCD numbers are added. To sort out this problem, there is a specific instruction DAA which stands for ‘Decimal Adjust Accumulator’. This assumes that the sum of the BCD addition is in AL.

The need for this instruction is because direct addition of BCD numbers causes errors, and a correction is usually needed. The details of the correction done are discussed in Section 0.7.2.

Now, see the following example which adds two multi byte BCD numbers, to get the result in BCD itself. The two BCD numbers are stored as double words (DD).

Note When they are stored thus, the lowest byte of the number will be in the lowest address. Now, BX and SI point to the BCD numbers to be added. For adding, only one byte each of the numbers is accessed and added. The sum is stored in the byte location SUM. Here too, the LSB of the sum will be in the lowest address.

Example 4.15

```

.MODEL SMALL
.DATA
BCD1 DD 45679834H ;store the first BCD byte as a double word
BCD2 DD 93870989H ;store the second BCD byte as double word
SUM DB 5 DUP(0) ;allocate 5 byte memory for the sum
.CODE
.STARTUP
MOV AH, 0 ;clear AH
CLC ;clear carry flag
LEA BX, BCD1 ;BX to point to first BCD byte
LEA SI, BCD2 ;SI to point to second BCD byte
LEA DI, SUM ;DI to point to SUM location
MOV CX, 04 ;count for the number of bytes
RPT: MOV AL, [BX] ;copy the bytes of the first number to AL
MOV DL, [SI] ;copy the bytes of the first number to AL
ADC AL, DL ;add with carry AL and DL
DAA ;do the BCD correction
MOV [DI], AL ;store the result of each byte addition
INC BX ;increment pointer BX
INC SI ;increment pointer SI
INC DI ;increment pointer DI
LOOP RPT ;loop until CX = 0
ADC AH, 0 ;add carry, AH and 0;sum in AH
MOV [DI], AH ;store AH in memory
.EXIT
END

```

On adding 4 BCD bytes, there is a possibility of a carry, which corresponds to the 5th BCD byte of the sum. To preserve this final carry, AH register is used as shown. The sum of the given two numbers is 0139559823.

ii) Subtraction

DAS – Decimal adjust AL after subtraction.

Usage: DAS

For BCD subtraction, there is an instruction DAS, which operates on the data in AL. Refer Section 0.7.5 for the intricacies of BCD subtraction. DAS is used (like DAA) to make the necessary corrections. The use of this instruction is left to you.

4.5 | ASCII Operations

The processor has a number of instructions relating to the processing of ASCII numbers.

4.5.1 | ASCII Addition

AAA – ASCII Adjust AL after addition

This instruction is used after addition of two ASCII numbers. It checks the lower nibble of AL for one of the two conditions:

- i) whether it is within A and F.
- ii) whether the auxiliary carry flag (AF) is set.

Depending on this, it performs the required corrections. See the three cases shown.

Add '5' and '4', get the sum in AL.

Case 1

```

35H
34H
69H

```

In this case the lower nibble is not within A and F, and the AC flag is not set i.e., AF = 0. Then the instruction just resorts to making the upper nibble of AL to 0. Thus if the AAA instruction is used after the addition, the sum becomes 09 i.e., 0000 1001. To convert it back to ASCII, OR it with 30H.

```

MOV AL, '5'           ;AL = 35H
ADD AL, '4'           ;AL = 35H + 34H = 69H
AAA                   ;AL = 09H
OR AL, 30H            ;AL = 39H, the ASCII value of 9

```

we get 39H in AL which is the ASCII code of 9. Also AH = 0.

Case 2

Add '7' and '6', get the sum in AL.

```

37H  +
36H
6DH

```

In this case, the right most nibble is greater than 9 – within A and F – it is no longer representing a decimal digit. Hence, 6 is added to the right most nibble (like in BCD correction, the number 6 is added because that is the difference between decimal and hexadecimal), it also clears the upper nibble of AL, sets CF and AF, and adds 1 to AH. Thus, if the AAA instruction is used after the addition, the sum in AL is now 03 and AH = 01 (if AH had been cleared earlier).

```

MOV AL, '7'           ;AL = 37H
ADD AL, '6'           ;AL = 37H + 36H = 6DH
AAA                   ;AL = 03, AH = 01, CF = 1
OR AX, 3030H          ;AX = 3133

```

The AX register will contain 3133H after this, which is the ASCII value of 13. These facts must be taken into consideration when adding multibyte ASCII numbers.

Case 3

Add '9' and '8', get the sum in AL.

$$\begin{array}{r} 39\text{H} \quad + \\ 38\text{H} \\ \hline 71\text{H} \end{array}$$

Here, the right-most nibble is not within A and F, but AF = 1. Hence, 6 is added to the right most nibble, and it also clears the upper nibble of AL, sets CF and AF, and adds 1 to AH. Thus, if the AAA instruction is used after the addition, AL = 7, AH = 1, CF = 1 and AF = 1. In all cases, AAA clears the upper nibble of AL. These facts must be taken into consideration when adding multibyte ASCII numbers.

4.5.2 | ASCII Subtraction

AAS – ASCII adjust AL after subtraction.

Usage: AAS

This function is similar to the AAA instruction. After the subtraction of two ASCII numbers, AAS checks the lower nibble of AL. If this number has a value between A and F, or if the AF flag is set, 6 is subtracted from AL. Also 1 is subtracted from AH and AF and CF are set. In all cases, the upper nibble of AL is cleared.

Case 1

Subtract '5' from '8'.

$$\begin{array}{r} 38\text{H} \quad - \\ 35\text{H} \\ \hline 03 \end{array}$$

If AAS is used after this subtraction, the content AL is 03.

Case 2

Subtract '8' from '5'.

$$\begin{array}{r} 35\text{H} \quad - \\ 38\text{H} \\ \hline \text{FDH} \end{array}$$

After AAS, AH = FFH, AL = 07.

Actually the answer should be -3 but we get FF07, which is the ten's complement $-10 + 3 = 7$. Let us subtract '4' from '15'. We get the ASCII value of '11' in the AX register with the following program.

MOV AX, '15'	;AX = 3135H
SUB AX, '4'	;AX = 3101H
AAS	;AX = 3101H
OR AX, 3030H	;AX = 3131H

Now, it is up to you to find out what happens when '9' is subtracted from '4'. Also, note that the instructions AAA and AAS can be used after adding or subtracting unpacked BCD numbers as well.

4.5.3 | Multiplication and Division

AAM – ASCII adjust AX after multiplication.

Usage: AAM

For multiplication and division of ASCII numbers, we need to convert them to unpacked BCD. Let us deal with ASCII multiplication first. In the following program, the numbers to be multiplied are 6 and 9. The product has to be 54. The ASCII numbers are first unpacked, and multiplied. After that, the AAM instruction is used. This causes the product to be in the unpacked BCD format. If the result of multiplication (MUL BL) is compared with the result of AAM, it is obvious that AAM accomplishes the conversion by dividing AX by 10. Following this, if ORing with 3030H is done, the ASCII value of the product is available in AX. In the program, this value is copied to CX, and displayed using the DISP macro.

Note The most significant digit is displayed first.

Example 4.16

```

.MODEL TINY
.CODE
.STARTUP
DISP MACRO                                ;define macro for display
MOV AH, 02
INT 21H
ENDM                                        ;end of macro
MOV AL, '6'                               ;AL = 36H
AND AL, 0FH                               ;AL = 06H
MOV BL, '9'                               ;BL = 39H
AND BL, 0FH                               ;BL = 09H
MUL BL                                    ;AX = 0036H;decimal value is 54
AAM                                       ;AX = 0504H, the unpacked BCD of 54
OR AX, 3030H                             ;AX = 3534H, the ASCII value of 54
MOV CX, AX                               ;copy AX to CX
MOV DL, CH                               ;DL = CH to display the MSD
DISP                                     ;call macro for display
MOV DL, CL                               ;DL = CL to display the LSD
DISP                                     ;call macro for display

.EXIT
END

```

AAD – ASCII adjust AX before division.

Usage: AAD

This instruction works by converting the unpacked BCD in AX to binary (hex) before division. It multiplies AH by 10, adds the product to AL and clears AH. This is the hex representation of the BCD number in AX. After division, we find the quotient in AL and the remainder in AH. See the following program segment.

```

MOV AX, '82'           ;AX = 3832H
AND AX, 0F0FH          ;AX = 0802H

```

AAD	;AX = 0052H, the hex value of the dividend
MOV BL, '9'	;BL = 39H, the ASCII value of 9
AND BL, 0FH	;BL = 09H
DIV BL	;AX = 0109H, the quotient in AL, remainder in AH
OR AX, 3030H	;AX = 3139H, the corresponding ASCII values

4.6 | Conversions for Computations and Display/Entry

We have seen various number formats so far, but by now it should be obvious that using BCD and ASCII numbers for computation is cumbersome, especially when large numbers are involved. A more convenient format for all arithmetic calculations is the binary format, which is compactly written in hexadecimal. Another point is that, once the computation is done, the natural format for display is the ASCII format. For displaying, we convert the binary number to unpacked BCD and then to ASCII. When entering data through the keyboard also, the ASCII number format is used. All this implies the necessity of converting binary numbers to ASCII and vice versa.

4.6.1 | Converting ASCII Numbers to Binary Form

Let us see how we convert a decimal number (say 6754) to binary form.

$$\begin{array}{r}
 6754 = 6 \times 1000 + 7 \times 100 + 5 \times 10 + 4 \times 1 \\
 4 \times 1 = 4 \quad + \\
 5 \times 10 = 50 \quad + \\
 7 \times 100 = 700 \quad + \\
 6 \times 1000 = \underline{6000} \\
 \underline{6754}
 \end{array}$$

This is the algorithm to be used for conversion. Let us write a program which does this. Here, the multiplicands 1, 10, 100 and 1000 are stored in the data segment. The binary value also is finally stored in data memory.

Example 4.17

```

.MODEL SMALL
.DATA
ASC DAT DB '6754'           ;ASCII value of 6754
MULTI  DW 01, 0AH, 64H, 3E8H
BINAR  DW 0

.CODE
.STARTUP
DIG EQU 4                   ;DIG specifies the number of digits
MOV CX, DIG                 ;move into CX the number of digits
MOV DI, DIG-1               ;DI to point to LSD of the ASCII no.
LEA SI, MULTI               ;use SI to point to the multiplicands

```

```

RPT:    MOV AL, ASCDAT[DI]           ;bring first ASCII data to AL
        AND AX, 000FH               ;mask AX
        MOV BX, WORD PTR[SI]       ;get the multiplicand to BX
        MUL BX                     ;multiply AX by BX
        ADD BINAR,AX               ;add to binary value
        DEC DI                     ;decrement byte pointer
        INC SI                     ;increment word pointer
        INC SI                     ;two increments required
        LOOP RPT                   ;repeat until CX = 0
        .EXIT
        END

```

The above program is a direct application of the algorithm mentioned. We get the converted binary value in the word location BINAR. We can extend this for more number of digits too.

4.6.2 | Converting Binary Numbers to ASCII Form

The method for this is repetitive division. This has been done in Example 3.22 for a 16-bit binary number.

4.7 | Signed Number Arithmetic

In all our discussions and problems, whenever we dealt with numbers, we assumed the numbers to be unsigned. Now, let us bring signed numbers also into our domain.

Negative numbers are represented in the two's complement format. When the data length is one byte, the maximum value of decimal numbers it can represent is -128 to $+127$. With 16 bits this is from $-32,768$ to $+32,767$. If the result of any operation exceeds this word size, the overflow flag is set. This indicates that our result has to be re-interpreted and corrected. The conditions under which the overflow flag is set can be stated as:

- i) When there is a carry from D6 to D7, or a carry from D7 outwards, but not both. This is for byte operations.
- ii) For word operations, when there is a carry from bit D14 to D15, or a carry from D15 outwards, but not both.

Let us see a few cases of signed number operations.

Example 4.18

```

        .MODEL TINY
        .CODE
        .STARTUP
MOV AL, +34
ADD AL, -23
        .EXIT
        END

```

For the above program in which +34 (22H) and –23 (E9H) are decimal numbers, the calculation in binary form:

$$\begin{array}{r} 0010\ 0010 \\ \underline{1110\ 1001} \\ 1\ 0000\ 1011 \end{array} \quad +$$

Thus the sum is 0BH (11 in decimal) which is the right sum. CF = 1, OF = 0, SF = 0.

Example 4.19

```

.MODEL TINY
.CODE
.STARTUP
MOV AL, -28      ;AL = E4H
ADD AL, -78      ;AL = B2H
.EXIT
END

```

Computation:

$$\begin{array}{r} 1110\ 0100 \\ \underline{1011\ 0010} \\ 1\ 1001\ 0110 \end{array} \quad + \quad \text{OF} = 0, \text{SF} = 1, \text{CF} = 1$$

The answer of this computation is 96H (–106) which is correct.

In both the above problems, the overflow flag is found cleared which indicates that the signed number computation has produced the right result. Now, see the next program.

Example 4.20

```

.MODEL TINY
.CODE
.STARTUP
MOV AL, +100
ADD AL, +75
.EXIT
END

```

Computation:

$$\begin{array}{r} 0110\ 0100\ (64\text{H}) \\ \underline{0100\ 1011\ (4\text{BH})} \\ 1010\ 1110\ (\text{AFH}) \end{array} \quad \text{OF} = 1, \text{SF} = 1, \text{CF} = 0$$

In the above case, the overflow flag will be set because there has been an overflow from bit D6 to D7, and the sign bit is set. Even though two positive numbers have been added, the sign bit indicates a negative sum. Thus AFH has to be thought of as a negative number, which is not correct. The problem is that the sum is +175 which is too large to fit in an 8-bit register. Since the overflow flag is set, it indicates that the register AL has **overflowed the size allotted for signed operations in 8-bit format.**

Understanding that the problem that has caused the error is one of inadequacy of the size of the destination, it can be solved by extending the bit size of the operands. This can be done by sign extending a byte to a word and a word to a double word. The 8086 has two specific instructions for it.

i) **CBW** – convert byte to word.

This instruction works on the data in **AL**. It copies the sign bit (D7) of **AL** to all the bits of **AH**. Thus the byte in **AL** is extended to a word in **AX**.

ii) **CWD** – convert word to double word.

This instruction works on the data in **AX**. It copies the sign bit (D15) of **AX** to all the bits of **DX**. Thus, **DX – AX** is the new double word. Let us now use these instructions to avoid overflow in the case of signed number operations.

Example 4.21

```

.MODEL TINY
.CODE
.STARTUP
MOV AL, +100
CBW
ADD AX, +75
.EXIT
END

```

Computation:

0000 0000 0110 0100	+	
0000 0000 0100 1010		
<u>0000 0000 1010 1110</u>		OF = 0, SF = 0, CF = 0

In this case, the overflow flag is not set and hence the sum in **AX** is considered to be correct. The sign bit also indicates a positive number.

Note In a general case, we are not sure if the result of adding/subtracting signed numbers will exceed the register size allowed for it. To accommodate for a correction as in Example 4.21, the instruction **JO** (Jump on overflow) can be used. The program can be written in a way that after the addition or subtraction (as the case may be) operation, the **JO** instruction tests the overflow flag and the **CBW** or **CWD** instruction is used only if the overflow flag is found to be set. Then re-computation using the extended word or double word can be done.

4.7.1 | Comparison of Signed Numbers

Comparison of two unsigned numbers is based on their numerical value only, and the sign does not come into the picture. For example, 7 is greater than 3 numerically. However, when these numbers are signed the picture changes. Think of the numbers -7 and -3 . We now say that -3 is greater than -7 . With this picture, we use the terms 'above' and 'below' for unsigned numbers, and 'greater than' and 'less than' for signed numbers. The mnemonics used and the flags used also are different.

In Chapter 3, Table 3.2 lists out a number of conditional jump instructions, catering to unsigned number operations. Now, let us have a look at the jump instructions for signed number operations. See Table 4.2.

Table 4.2 | Conditional Jump Instructions Catering to Signed Numbers

Sl No.	Mnemonic	What it means	Conditions tested
1	JE	Jump if Equal	ZF = 1
2	JG	Jump if Greater	ZF = 0 and SF = OF
3	JGE	Jump if Greater or Equal	SF = OF or ZF = 1
4	JL	Jump if Less	SF != OF
5	JLE	Jump if Less or Equal	ZF = 1 or SF != OF
6	JNG	Jump if Not Greater	ZF = 1 or SF != OF
7	JNGE	Jump if Not Greater Nor Equal	SF != OF
8	JNL	Jump if Not Less	SF = OF
9	JNLE	Jump if Not Less Nor Equal	ZF = 0 and SF = OF
10	JNO	Jump if No Overflow	OF = 0
11	JO	Jump if Overflow	OF = 1
12	JNS	Jump if Not Signed	SF = 0
13	JS	Jump if Signed	SF = 1

Note i) The mnemonics JE/JZ and JNE/JNZ are applicable for both signed and unsigned numbers and the condition of the Zero flag being set or reset accordingly applies here as well.

ii) != is the symbol for 'not equal to'.

The list in Table 4.2 shows the conditions of the sign, overflow and zero flags and undoubtedly causes a certain amount of confusion. Let us test a program for the various situations possible and check the condition of these flags.

Example 4.22

```

.MODEL TINY
.CODE
.STARTUP
MOV AL, +40
CMP AL, -60      ;after this ,OF = 0, SF = 0
CMP AL, +70      ;after this ,OF = 0, SF = 1
MOV BL, -20      ;BL = -20
CMP BL, -70      ;after this ,OF = 0, SF = 0
CMP BL, -10      ;after this ,OF = 0, SF = 1
CMP BL, +10      ;after this ,OF = 0, SF = 1
CMP BL, -20      ;after this ,OF = 0, SF = 0, ZF = 1
.EXIT
END

```

The result of each comparison of the above program will help to clear the confusion.

- i) When the destination is greater than the source, OF = SF.
- ii) When the destination is less than the source, OF != OF.
- iii) When the destination is equal to the source, ZF = 1.

Example 4.23 finds the smallest of three signed numbers and stores it in a memory location. It compares the first number with the second number. If the first number is less, it is compared with the third number. Otherwise, the second number is compared with the third number and finally the smallest number is identified and stored in memory. In this case, the negative number -90 is found to be the smallest number.

Example 4.23

```

.MODEL SMALL
.DATA
NUMS DB +45, -76, -90
SMALLEST DB 0
.CODE
.STARTUP
    LEA BX, NUMS           ;use BX as pointer to the numbers
    MOV AL, [BX]           ;get the first number in AL
    INC BX                 ;increment the pointer
    CMP AL, [BX]           ;compare first no. with the second
    JLE NEXT               ;if AL is less than or equal, jump
    MOV AL, [BX]           ;else, load second no. in AL
NEXT: INC BX               ;increment the pointer
    CMP AL, [BX]           ;compare second no. with third
    JLE FINAL              ;if the first no. is less, jump
    MOV AL, [BX]           ;move the smallest number to AL
FINAL: MOV SMALLEST, AL    ;store the smallest no. in memory
    .EXIT
    END

```

4.7.2 | Signed Multiplication and Division

There are two instructions to be used for signed multiplication and division.

i) IMUL

All that was discussed about the MUL instruction (Section 3.4.4) holds true in the case of IMUL too. The only difference is that since the operands are signed, the product also is signed. Another point is that when the product is not large enough to use all the registers allocated for it, the sign bit is copied to the upper register (AH or DX as the case may be) and CF = OF = 0.

ii) IDIV

The signed division operation is similar to unsigned division, as discussed in Section 3.4.5, except that the operands and result are signed. Because of this, if the quotient register is AL, the value of the quotient is to be between -128 and +127. If the quotient register is to be AX, this

value is to be between $-32,768$ and $+32,767$. If this is violated, a 'divide overflow' error message is displayed on the screen.

Now, let us use the signed multiply and divide instructions in a program. Let us compute $(XY - Z^2)/L$ and store the quotient alone. The values of X, Y, Z and L are put in the data segment as shown in the program. For multiplication, a procedure named MULT is defined and used.

Example 4.24

```

.MODEL SMALL
.DATA
X DB -6
Y DB 25
Z DB 19
L DB -9
QUO DB?

.CODE
.STARTUP

MOV AL, X           ;copy X to AL
MOV BL, Y           ;copy Y to BL
CALL MULT           ;call the procedure to multiply X and Y
MOV CX, AX          ;copy the product to CX
MOV AL, Z           ;copy Z to AL
MOV BL, AL          ;copy Z to BL as well
CALL MULT           ;call the procedure to get Z2
SUB CX, AX          ;compute XY-Z2
XCHG CX, AX         ;exchange CX and AX, to get dividend in AX
IDIV L              ;signed division by L
MOV QUO, AL         ;copy the quotient to memory
.EXIT

MULT PROC NEAR      ;define the MULT procedure
IMUL BL             ;multiply AL by BL, product in AX
RET                ;return to calling program
MULT ENDP           ;end the procedure
END

```

In this program, the operands are positive as well as negative numbers. Let us make a few observations.

- i) With the values of X, Y, Z and L as given in the program, the product is -511 (FE01H) and the quotient is $+56$, which is found in AL as 38H.
- ii) If the divisor is now changed to $+9$, the quotient is -56 , which is found in AL as C8H.
- iii) If the divisor is changed to -3 , we get a message **divide overflow error** and the program execution is terminated. This is because the quotient ($+170$) is too large to fit in AL. Remember that the largest positive number that AL can hold is $+127$ only.

4.7.3 | Arithmetic Shift

A shift operation that makes a difference in the result depending on the sign of the operand, is arithmetic shifting. There are two instructions pertaining to this.

i) SAR destination, count – Shift right arithmetic.

This is similar to the shift right logical (SHR) discussion in Section 3.5.2, but here when shifting right, the sign bit is filled into the MSB. This instruction can be used to divide a signed number by 2, for one shifting each.

```
MOV BL, -24          ;AL = 1110 1000
SAR BL,1             ; AL = 1111 0100, which is -12
```

ii) SAL destination, count – Shift left arithmetic.

This is the same as shift left logical (SHL). As shifting left does not involve any sign bit, logical and arithmetic shifting are the same. The two mnemonics (SAL and SHL) mean the same and can be used interchangeably.

4.8 | Programming Using High Level Language Constructs

MASM has become easier and very convenient to use ever since it was appended with a number of high level language constructs. Such constructs are available only for MASM 6.0 and higher versions. They are designated as ‘dot commands’ as they are preceded by a dot. These constructs make MASM a ‘high level assembler’. A few of these constructs are listed below:

- i) IF, ELSEIF and ELSE ... ENDIF
- ii) REPEAT ... UNTIL ... ENDW
- iii) REPEAT ... WHILE ... ENDW
- iv) BREAK
- v) CONTINUE
- vi) REPEAT ... UNTILCXZ

Except the last one, these constructs function as they do in a high level language (for example, C and C++). For the last construct, the loop repeats until register CX = 0.

Now, we will endeavor to use these constructs in a few interesting programs. Example 4.25 is a simple program which takes in data through the keyboard (without echo) and verifies three conditions.

- i) Is the value of the number between 0 and 5?
- ii) Is the value of the number between 6 and 9?
- iii) Is a key pressed ‘not a number’?

These three questions are resolved using IF, ELSEIF and ELSE statements and appropriate messages are displayed. The symbol ‘&&’ is used for ‘and’ just as in a high level language.

Example 4.25a

```
.MODEL SMALL
.DATA
FIRST DB "THE VALUE IS WITHIN 0 AND 5$"
```

SECOND DB "THE VALUE IS WITHIN 6 AND 9\$"

OTHER DB "NOT A NUMBER\$"

```

.CODE
.STARTUP

        MOV AH, 08                      ;read in key without echo
        INT 21H

    .IF AL> = '0' && AL< = '5'          ;first IF condition
        LEA DX, FIRST                  ;display the first message
        MOV AH, 09
        INT 21H

    .ELSEIF AL > = '6' && AL< = '9'      ;second IF condition
        LEA DX, SECOND                 ;display the second message
        MOV AH, 09
        INT 21H

    .ELSE                                ;ELSE
        LEA DX, OTHER                  ;display the third message
        MOV AH, 09
        INT 21H

    .ENDIF

        .EXIT
        END

```

Let us see what exactly these high level language constructs are. Obviously, they get converted to assembly language statements. We use the .LISTALL directive to find out. Example 4.25b shows the listing for the code segment alone up to the .ENDIF statement.

We find the equivalent assembly code for the high level language constructs, which in this case are compare and jump instructions.

Example 4.25b

```

0017 B4 08                      MOV AH, 08
0019 CD 21                      INT 21H

                                .IF AL> = '0' && AL< = '5'

001B 3C 30                      *   cmp     al, '0'
001D 72 0E                      *   jnb     @C0001
001F 3C 35                      *   cmp     al, '5'
0021 77 0A                      *   jnb     @C0001
0023 8D 16 0000 R              LEA DX, FIRST
0027 B4 09                      MOV AH, 09
0029 CD 21                      INT 21H

                                .ELSEIF AL > = '6' && AL< = '9'

002B EB 1A                      *   jmp     @C0004
002D                                *@C0001:
002D 3C 36                      *   cmp     al, '6'
002F 72 0E                      *   jnb     @C0005

```

```

0031 3C 39          *      cmp      al, '9'
0033 77 0A          *      ja       @C0005
0035 8D 16 001D R          LEA DX, SECOND
0039 B4 09          MOV AH, 09
003B CD 21          INT 21H

                        .ELSE
003D EB 08          *      jmp      @C0008
003F              *@C0005:
003F 8D 16 003A R          LEA DX, OTHER
0043 B4 09          MOV AH, 09
0045 CD 21          INT 21H

                        .ENDIF

```

Next, we will do a program using the REPEAT ... WHILE construct. Example 4.26 shows how the WHILE statement is used to repeat execution of a loop while a condition is satisfied, and to exit the loop when the condition no longer holds.

As long as CX is not equal to zero, the loop repeats. The end of the loop is specified with a .ENDW construct. In this program, a string consisting of all the letters of the alphabet, stored in location starting from DAT is to be reversed and displayed. After the reversed string is stored in location DAT1 onwards, it is appended with a '\$' so that it can be displayed with the INT 21H function number 9. The message REVERSED STRING: is displayed before the reversed string is displayed.

Example 4.26

```

.MODEL SMALL
.DATA

DAT DB 'ABCDEFGHIJKLMNOPQRSTUVWXYZ$'
MES1 DB 0DH,0AH,'REVERSED STRING: $'
DAT1 DB 27 DUP(?)

.CODE
.STARTUP
MOV BX,OFFSET DAT +25      ;address end of DAT
MOV SI,OFFSET DAT1        ;address DAT1
MOV CX,26                 ;load count in CX
.WHILE CX! = 0             ;loop while CX! = 0
    MOV AL,BYTE PTR[BX]    ;move byte to AL
    MOV BYTE PTR[SI],AL    ;save AL in DAT1
    INC SI
    DEC BX
    DEC CX
.ENDW                      ;end of while loop
MOV BYTE PTR[SI + 1], '$'  ;append string with $
MOV DX,OFFSET MES1
MOV AH,09
INT 21H
MOV DX,OFFSET DAT1        ;display the reversed string

```

```

MOV AH,09
INT 21H

.EXIT
END

```

The output of the above program is as shown below.

REVERSED STRING: ZYXWVUTSRQPONMLKJIHGFEDCBA
C:\masm6.14\BIN>

For both the above problems, the high level language constructs are written indented to the left. This is done only to improve the readability of the program.

The above problem can be worked out using the REPEAT ... UNTIL construct. The loop is repeated until CX = 0. CX becoming zero is the stopping condition for the loop.

There is also a UNTILCXZ construct available that can be used in this example. The UNTILCXZ instruction uses the CX register as a counter to repeat a loop a fixed number of times. This example can be rewritten using a REPEAT UNTILCXZ loop also.

Now, let us see a case of nested loops. Example 4.27 shows the case of an IF loop nested within a repeat-until loop. In this, a character is to be searched within a string and the number of occurrences of that character in the string is to be found out. To make the program interactive for the user, messages asking for the character and the string are displayed using DOS INT 21H function 9. Both the character and the string are entered through the keyboard. The string is terminated when the 'enter' key is pressed. This corresponds to the 'carriage return' character 0DH and this is the termination condition of the repeat until loop.

Example 4.27

```

.MODEL SMALL
.DATA
MES1 DB 0DH, 0AH, 'ENTER THE CHARACTER: $'
MES2 DB 0DH, 0AH, 'ENTER THE STRING: $'
MES3 DB 0DH, 0AH, 'NUMBER OF OCCURENCES: $'
STRING DB 50 DUP(' ')
.CODE
.STARTUP
MOV DX, OFFSET MES1      ;address MES1
MOV AH, 09                ;display message
INT 21H                   ;for the character
MOV AH, 01
INT 21H                   ;enter character
MOV BL, AL                ;copy character to BL
MOV DX, OFFSET MES2      ;address MES2
MOV AH, 09                ;display message
INT 21H                   ;for entering string
MOV CL, 0                 ;initialize counter
.REPEAT                   ;repeat loop
MOV AH, 01                ;enter string, one character
INT 21H                   ;at a time

```

```

. IF (AL == BL)                ;compare with reference
    INC CL                    ;if matching, increment CL
.ENDIF                        ;end IF loop
.UNTIL AL == 0DH               ;exit when 'enter' is pressed
    ADD CL, 30H               ;convert to ASCII
    MOV DX, OFFSET MES3       ;address MES3 location
    MOV AH, 09
    INT 21H                   ;display message
    MOV DL, CL                ;move count to DL
    MOV AH, 06                ;display the count
    INT 21H

    .EXIT
    END

```

The output of the Example 4.27 for a string and the character 't', entered through the keyboard is shown below.

```

ENTER THE CHARACTER: t
ENTER THE STRING: tyu8iolhthtttfeerttt
NUMBER OF OCCURENCES: 8
C:\masm6.14\BIN>

```

Now that we have seen the use of the high level language constructs, it is obvious that many of our previous programs can be re-written using these constructs.

KEY POINTS OF THIS CHAPTER

- String instructions are used to simplify coding when bulk data is moved, scanned, and compared.
- Then, SI and DI registers are used as pointers to the data segment and extra segment, and the direction flag is used for auto increment/decrement.
- The conditional/unconditional REP prefix is used for repetitive string operations.
- Procedures are subprograms which are 'called' by the main program for performing specialized operations.
- Procedures may be intra-segment or inter-segment and this is clarified when defining a procedure, by using the 'near' or 'far' prefix.
- Parameters to and from procedures may be passed via registers, memory or stack.
- Macros are functionally similar to procedures, but act in a different way.
- Number format conversions are necessary as BCD, ASCII and binary numbers are used frequently.
- 8086 has specialized instructions for handling ASCII and BCD number operations.
- Signed number arithmetic involves the use of special instructions.
- The overflow flag is very important in signed operations.
- The high level language constructs of MASM have made assembly language programming easier.

QUESTIONS

1. What is the role of the direction flag in string instructions?
2. What is the difference in the functioning of the instructions STD and CLD?
3. How does the instruction CLD function in the case of the following two instructions?
 - a) REP MOVSB
 - b) REP MOVSW
4. Which is the pointer register to be used when using the extra segment in string operations?
5. Why does the LODS instruction not use the REP prefix?
6. Do the following instructions function differently? In what way?
 - a) REPNE CMPSB
 - b) REPE CMPSB
7. For REPZ CMPSB, what are the conditions under which the loop is exited?
8. Distinguish between a near and a far call.
9. What is the difference in the functioning of the RET and RET n instructions?
10. What is meant by the term 'parameter passing' in the context of procedures?
11. What is the benefit of using macros?
12. Compare procedures and macros.
13. How are local variables in a macro taken care of?
14. What is a dummy argument?
15. Explain how the DAA instruction functions.
16. Explain the action of the following instructions:
 - a) AAM
 - b) AAS
 - c) AAD
 - d) AAM
17. Write down the steps required for the following conversions:
 - a) binary to ASCII,
 - b) ASCII to binary.
18. When is the CWD instruction used?
19. Is it justified to call MASM a high level assembler?

EXERCISE

1. Write a program to move 100 words from 'FROM' to 'TO' which are two areas in the data segment separated by a 200-byte space.
2. Write a program to scan the name SAHABUDDIN and replace S by B and B by H.
3. Store a password in memory. Enter another password through the keyboard and verify if it matches the stored password. The password entered should not be displayed as such, but each letter should be displayed as [^]%. '.
4. Write a string program for an application which requires the pointer registers to be auto decremented.
5. Write a program to search for a word in a block of N words.

6. Write a program which contains the following macros:
 - a) for calculating the Fibonacci series for N,
 - b) for entering the value of N,
 - c) for displaying the numbers.
7. Do the above problem using procedures.
8. Display the factorial of three numbers. Solve the problem using macros as well as procedures.
9. Obtain the list file of the above program in both cases and observe the differences.
10. Enter a string of characters through the keyboard. Save it in memory. Verify if the string is a palindrome.
11. Using macros/procedures, write a program to add 2 BCD numbers entered through the keyboard and display their sum.
12. Write a program where nested macros are used.
13. There are 16 ASCII numbers stored in consecutive locations in memory. Convert this to an eight-bit packed BCD number.
14. Write a program to add two ten-byte BCD numbers. The sum is to be displayed.
15. Two BCD numbers 4093H and 2986H are stored in memory. Subtract the second number from the first. Use the DAS instruction.
16. Write a program to add and display two ASCII numbers '5678' and '3498'.
17. Divide '90' by '9'. Use AAD and necessary programming steps.
18. Divide '3476' by '5'. Display the quotient alone.
19. Use AAM to help in converting a 16-bit binary number in AX to a four-digit ASCII character.
Hint Divide the content of AX-DX by 100, then use AAM.
20. Write a program which checks the number of ASCII digits present in the given number, and converts it to the corresponding binary equivalent.
21. Write a program to convert binary numbers to ASCII format for display. The program should have the following features:
 - a) Display the number.
 - b) Convert the ASCII number and store it in memory.
 - c) A procedure, DISPLAY should access the memory to display the numbers.
22. There are 10 signed bytes stored in memory, some of which are negative and some of which are positive. Find the largest number of the lot.
23. Re-write Example 4.26 using REPEAT ... UNTIL and REPEAT ... UNTILCXZ constructs.
24. Re-write Example 4.27 using .CONTINUE and .BREAK constructs along with .IF and .REPEAT constructs.

5 PROGRAMMING CONCEPTS - IV



In this chapter, you will learn

- The difference between peripheral I/O and memory mapped I/O.
- The use of the input/output instructions of 8086.
- The importance of modular programming.
- The use of directives which facilitate modular programming.
- To write programs in different modules and link them.
- To use assembly modules in a C programming environment.

Introduction

We know that the processor we use is connected to the I/O as well as the memory. So far we have only dealt with data being read from and written to memory. We have accessed I/O devices like the keyboard and the video monitor, but for that we used DOS interrupts which are functions already written and tested. Now, let us try to deal with I/O directly using the instructions of 8086. In practice, I/O devices are not connected directly to the processor, but are connected through interfacing chips which provide control signals for the processor to communicate with the peripheral (I/O device). A lot of hardware is involved in this, which will be discussed in detail in Chapters 9, 10 and 11.

5.1 | Input/Output Programming

There are many I/O devices connected to the processor, using which, data is taken in (input device) or data is given out (output device). Thus in I/O processing, there appears the concept of a port which means a register or storage area for data, from/to which, data is transferred to the processor. Thus an I/O device is frequently referred to as a 'port' with an associated address. Any I/O device has a unique address, known as 'port address' or simply mentioned as 'port'.

Now see Figure 5.1 which shows an input and output device connected to the processor through the address and data bus and a control pin. Thus, when its address is placed on the address bus and the $\overline{WR}$ signal is activated, data from the processor is written to the output port. Similarly, data from the input port is read into the processor when the input port is addressed and the $\overline{RD}$ signal is activated. There are two schemes for connecting an I/O device to the processor.




Figure 5.1 | I/O ports with address bus and data bus and the read and write control signals

5.1.1 | Memory Mapped I/O

In this scheme, the I/O device has the same address space and addressing scheme as memory. For example, given the instruction `MOV AL, [3456H]`, the address 3456H may be a memory location or the address of an I/O device. That means, that if there are 50 I/O devices in a system, the address space available to memory gets reduced by 50. However, that also means that I/O does not need any special instructions for access, and also that data from/to I/O can be transferred to any general-purpose register of the processor. This kind of approach is used for some RISC processors where the idea is to reduce the number of instructions. This also implies that there is no need of extra hardware control signals to differentiate between I/O and memory.

5.1.2 | Peripheral or Isolated I/O

This scheme is sometimes referred to as 'I/O mapped I/O' as well. Here, there are special instructions catering to input and output devices, and the address space is disjoint and separate from the memory address space. Thus for 8086, the full 2^{20} addresses are available as memory addresses. If there are 50 ports in a system, there are 50 I/O addresses. Since I/O does not use the `MOV` instruction, there will be no conflict between memory and I/O accesses. However, extra control signals are required in this case. The hardware aspects of I/O and its connection to the processor will be discussed in Chapter 7. We will deal with the Peripheral I/O scheme alone here, as that is the scheme most commonly used for 8086 based systems.

5.2 | I/O Instructions

The 8086 has two instructions to communicate with I/O.

i) `IN`

Usage: `IN` accumulator, port

This is the case when data is taken from the I/O device – this implies that the device is an input device. With this instruction, an **I/O Read** machine cycle is initiated and the data in the input device is transferred to the accumulator of the processor.

ii) OUT

Usage: OUT port, accumulator

In this case data is given to an output device from the accumulator with the initiation of an **I/O Write** machine cycle. In both cases, ‘port’ implies the port address i.e., the address of the I/O device. The source has always to be the ‘accumulator’ which means the AL register if an 8-bit data is being transferred, or the AX register if a 16-bit register is used. No other register may be used. I/O addresses are limited to a maximum bit size of 16 bits. Thus there can be a maximum of 2^{16} ports – there can be 2^{16} input ports and 2^{16} output ports, as the instruction used for input and output is different. There are two formats available for I/O instructions.

5.2.1 | Fixed Port Addressing

This is used only **when the address of an I/O device is 8 bits wide**. Here the address of the port is directly mentioned in the instruction. Note that data can be 8 bits or 16 bits, depending on the data bus width of the I/O device.

Examples

IN AL, 45H	;move 8-bit data into AL from an input port with address 45H
IN AX, 12H	;move 16-bit data into AX from a port with address 12H
OUT 34H, AL	;move 8-bit data from AL to an output port with address 34H
OUT 0FCH, AX	;move 16-bit data from AX to an output device with the address FCH

5.2.2 | Variable Port Addressing

This is used when the addressed port has a **16-bit address**. Here, the port address is to be loaded into DX and then only the I/O instruction can be used. Only the DX register can be used for this. For example, if the address of an input device is 9876H, the steps are:

MOV DX, 9876H	;load the address of the port in DX
IN AL, DX	;move 8-bit data into AL from an input port whose address is in DX

Similarly,

MOV DX, 0FC6H	;load the 16-bit address of the port in DX
OUT DX, AX	;move 16-bit data from AX to output port whose address is in DX

Example 5.1

An 8086 is connected to two ports with 8-bit addresses. The address of the input port is 67H and the address of the output port is FEH. The input port sends a 16-bit data to the processor. Only the lower 12 bits are to be transferred to the output port. Write the relevant program lines for this.

Solution

```
IN AX, 67H           ;get to AX, 16-bit data from input port 67H
AND AX, 0FFFH        ;mask the upper nibble
OUT 0FEH, AX         ;send content of AX to output port FEH
```

The complete program can also be written using the EQU directive.

```
                .MODEL TINY
                .CODE
PORTA EQU 67H           ;use labels for port numbers
PORTB EQU 0FEH
                .STARTUP
                IN AX, PORTA
                AND AX, 0FFFH
                OUT PORTB, AX
                .EXIT
                END
```

Example 5.2

An 8086 is used to monitor temperature at 20 points and send out alarms if the temperature at any point is greater than the maximum allowed. Hence, there are 20 input ports to get the temperature value (from the corresponding sensors), and 20 output ports to send alarms. The input ports have addresses ranging from EC00H to EC13H, and the output port addresses range 0200H to 0213H. The maximum allowed temperature is MAX and this is stored in memory (here we take MAX = 65). Sending an alarm is done by transferring the number 01 to the corresponding output port.

Solution

```
                .MODEL SMALL
                .DATA
TMPS DB 20 DUP(0)
MAX DB 65
                .CODE
                .STARTUP
                MOV CX, 20           ;count of the number of ports
                MOV BX, 0           ;BX = 0
                MOV DX, 0EC00H      ;address of first i/p port in DX
                IN AL, DX           ;get data from input port to AL
```

```

TAKE_IN:      MOV  TMP5[BX] , AL      ;transfer data to memory
              INC  DX                  ;increment port address
              INC  BX                  ;increment memory address pointer
              LOOP TAKE_IN            ;repeat until CX = 0
              MOV  AH, MAX             ;copy maximum temperature to AH.
              MOV  DX, 0200H          ;get address of first o/p port
              MOV  AL, 01             ;get data for alarm, in AL
              MOV  CX, 20             ;CX = count
              MOV  BX, 0              ;BX = 0
AGAIN:        CMP  TMP5[BX] , AH      ;compare data in memory with MAX
              JA   ALARM              ;if data>MAX, go to ALARM
OTHER:        INC  DX                  ;increment port address
              INC  BX                  ;increment address pointer
              LOOP AGAIN              ;repeat until CX = 0
              JMP  EXEET              ;jump to exit point
ALARM:        OUT  DX, AL             ;output alarm data to o/p port
              JMP  OTHER
EXEET:
              .EXIT
              END

```

In the above program, the data from the 20 input ports are read, one at a time using the IN instruction, each time changing the content of DX. These are stored in the data memory. Then each of these data values is compared with MAX and if the value stored is greater than MAX, a value (01) is transferred to the corresponding output port which causes an alarm to become ON. The access to input and output ports is through DX, the content of which is changed as required.

For now, we have discussed the input/output instructions of 8086. In Chapters 9, 10 and 11, we will discuss other aspects of input/output programming and associated hardware aspects.

5.3 | Modular Programming

When a large programming problem is to be handled, it is quite logical to think of breaking it up into modules, testing the individual modules separately, and then integrating the modules together to form the complete solution. It may also be that a simple problem at hand can be solved easily by using program modules that have already been written and tested. Another scenario would be when various teams work on different modules of the same problem and finally integrating it all for a final and complete solution. All these indicate the importance of what we would term 'modular programming'.

In this case, different assembly files are individually tested, but they have to be linked together. There are many issues to be resolved for successful linking and execution.

- i) The different modules that constitute the solution may be in different code segments.
- ii) The data which is to be used by one module may have to be accessed by many different modules and the 'permission' for this must be indicated.

- iii) Some labels used in a module may not be found therein. In that case, there must be some indication that these are defined in some other module to which linking is possible and will be done.

To resolve these issues, MASM has directives which are to be included in programs. Let us use some of these.

5.3.1 | PUBLIC

When a data item or a procedure is to be allowed to be accessed by other modules, it is declared as PUBLIC.

For example,

```
PUBLIC num1, num2, num3           ;declares three data items to be public
```

This may also be declared as

```
PUBLIC num1
PUBLIC num2
PUBLIC num3
```

A procedure named MULT can be allowed access from other modules if it is declared as

```
PUBLIC mult
```

5.3.2 | EXTRN

When a module needs to use data and code which have been defined elsewhere, it should use the directive EXTRN meaning that the labels being used in this module are external to the module. Otherwise, the assembler will give an error message of 'undefined symbol'. The general format of using this directive:

```
EXTRN name1:type                 ;one item name alone is declared
EXTRN name1:type, name2:type     ;many items names are declared together
```

For example,

```
EXTRN num1:byte
EXTRN num2:word,num3: word
EXTRN multi :far                ;a far procedure is declared
```

Note Remember that MASM is not case sensitive. The use of upper case and/or lower case in all the above cases is only to enhance readability.

Let us now run a simple program which consists of two modules. A file 'moda.asm' is a program which uses data and a procedure defined in another module. Hence the data items that it uses have been declared with EXTRN, as they are not defined in this module. After computation, the result is stored in the data segment of this module.

The second module is another file named 'modb.asm'. This module contains the data and the procedure that the first module uses. Hence they are declared as PUBLIC. What the two modules together do is add two bytes – one taken in through the keyboard, and one stored in the data segment of the second module. The sum is stored in the data segment of the first module. (These data segments may be found to be the same, finally.)

Now, one point to ponder is whether both these modules are in the same code segment. If the small memory model is used, they will be in the same code segment (on assembling, this can be verified – the CS value will be the same for the instructions of both modules). If we want

the second module to be in a different code segment, we have to use memory models other than tiny, small or compact (refer Table 2.2). Let us use the large model here. These files are assembled and linked together as shown. The filenames corresponding to the two modules are named MODA.ASM and MODB.ASM.

Example 5.3

MODA.ASM

```

.MODEL LARGE      ;define large model
.DATA
SUM DB ?          ;allocate space for sum
EXTRN DATA1: BYTE ;DATA1 defined in another module
EXTRN DATA2: BYTE ;DATA2 defined in another module
.CODE
EXTRN REAP: FAR    ;the procedure in another segment
.STARTUP
CALL REAP          ;call the external procedure
MOV AL, DATA1     ;get external data DATA1 to AL
ADD AL, DATA2     ;add AL with external data DATA2
MOV SUM, AL        ;sum stored in the data segment
.EXIT
END

```

MODB.ASM

```

.MODEL LARGE      ;define large model
.DATA
PUBLIC DATA1     ;declare that DATA1 is PUBLIC
PUBLIC DATA2     ;declare that DATA2 is PUBLIC
DATA1 DB ?        ;define DATA1
DATA2 DB 26H      ;define DATA2
.CODE
PUBLIC REAP
.STARTUP
REAP PROC FAR     ;define the procedure REAP
MOV AH, 01        ;for entering character
INT 21H           ;through keyboard
SUB AL, 30H       ;convert ASCII to binary
MOV DATA1, AL    ;store it in data segment
RET              ;return
REAP ENDP         ;end the procedure
.EXIT
END

```

The modules can be assembled and linked together using ml.exe. The command is

```
ml filename1.asm filename2.asm filename3.asm ...
```

See the output on the command window corresponding to the command

```
ml moda.asm modb.asm
```

(the first filename should correspond to the module that calls the other modules).

Example 5.4

```
C:\MASM6~1.14\BIN>ml moda.asm modb.asm
```

```
Microsoft (R) Macro Assembler Version 6.14.8444
```

```
Copyright (C) Microsoft Corp 1981–1997. All rights reserved.
```

```
Assembling: moda.asm
```

```
Assembling: modb.asm
```

```
Microsoft (R) Segmented Executable Linker Version 5.60.339 Dec 5 1994
```

```
Copyright (C) Microsoft Corp 1984–1993. All rights reserved.
```

```
Object Modules [.obj]: moda.obj+
```

```
Object Modules [.obj]: "modb.obj"
```

```
Run File [modul1a.exe]: "moda.exe"
```

```
List File [nul.map]: NUL
```

```
Libraries [.lib]:
```

```
Definitions File [nul.def]:
```

```
LINK : warning L4021: no stack segment
```

```
C:\MASM6~1.14\BIN>
```

The list files of both modules can be obtained by the command

```
ml/F1 moda.asm modb.asm.
```

Note In the list file, the address of lines of code which use external data/program have the letter 'E' in it. For debugging use the command **debug moda.exe**. Then the sequence of steps of going from one code segment to the other and returning can be clearly observed.

Now let us examine another program, which is a password checking program. In Example 5.5, there is a stored password (PW1). Another password is entered through the keyboard (PW2) and saved. The storing of the first password and entering of the second one using a procedure named READ, is in one module (5.5b). Thus the two passwords and the procedure READ are declared as PUBLIC in module (5.5b) meaning that they can be accessed from any module. The second module (5.5a) accesses these passwords, compares them and authenticates the entered password if it matches with the first one. Appropriate messages are displayed. Note that the password PW2 is entered without echo.

Example 5.5a

```

.MODEL HUGE ;use huge model
.DATA
MES1 DB 'WRONG PASSWORD$'
MES2 DB 'PROCEED$'
```

```

        EXTRN PW1:BYTE          ;PW1 is defined in another module
        EXTRN PW2:BYTE          ;PW2 is defined in another module
                                .CODE
        EXTRN READ:FAR           ;READ is defined in another module
                                .STARTUP
        MOV AX, DS                ;copy DS to AX
        MOV ES, AX                ;copy AX to ES
        LEA SI, PW1               ;let SI point to PW1
        LEA DI, PW2               ;let DI point to PW2
        MOV CX, 6                 ;CX=6, the number of characters
START:   CALL READ                ;call procedure for entering PW2
        CLD                      ;clear DF
        REPE CMPSB                ;repeat comparison until CX=0
        JCXZ CORRECT             ;if CX=0,go to displaying message
        MOV DX, OFFSET MES1       ;point DX to MES1
        MOV AH, 09                ;function number for display
        INT 21H
        JMP EXEET
CORRECT: MOV DX, OFFSET MES2       ;point DX to MES2
        MOV AH, 09
        INT 21H
EXEET:

                                .EXIT
                                END

```

Example 5.5b

```

                                .MODEL HUGE ;use huge model
                                .DATA
        PUBLIC PW1                ;declare PW1 as public
        PUBLIC PW2                ;declare PW2 as public
        PW1 DB 'SAMSON'           ;define PW1
        PW2 DB 6 DUP(0)           ;allocate space for PW2

                                .CODE
        PUBLIC READ                ;declare READ as public
                                .STARTUP
READ    PROC FAR                  ;define the READ procedure
        LEA BX, PW2               ;BX to point to PW2
        MOV DL, 6                 ;DL = 6, the number of characters
START:  MOV AH, 08                ;enter character without echo
        INT 21H
        MOV [BX], AL              ;save the character in memory
        INC BX                    ;increment memory pointer
        DEC DL                    ;decrement the count
        JNZ START                 ;jump if DL!=0
        RET                       ;return

```



```

                .CODE
                PUBLIC CONV                ;declare the procedure as public
                .STARTUP
CONV  PROC FAR                ;conversion from binary to ASCII
        MOV AX, NUM
        MOV DX, 0
        MOV CX, 10
REPEA:    DIV CX
        PUSH DX
        MOV DX, 0
        INC COUNT
        CMP AX, 0
        JNE REPEA
DISP:    POP DX
        ADD DL, 30H
        MOV AH, 02
        INT 21H
        DEC COUNT
        JNZ DISP
        RET
CONV  ENDP
        .EXIT
        END

```

Note The explanation of the steps in this program is given along with Example 3.22.

Thus, we see that one important application of modular programming is the use of frequently used procedures which have already been written, tested and stored in a different module.

5.4 | Programming in C with Assembly Modules

We have discussed in Chapter 0, the importance and advantages of assembly language programming compared to high level language programming. However, we find high level languages being used more commonly, because of being portable and platform independent (generally, but not always). However, because of the efficiency of assembly language, there are advantages when both these are used in a mixed fashion. C is a very popular language and in this section, we shall use assembly language modules within C programs. This is called ‘inline assembly’ and most features (but not all) of assembly language can be used here. Depending on the version of C used, the constructs within the assembly module can vary. We will use Borland’s Turbo C++ compiler here. To understand the programs well, you need to have some experience in C/C++ programming. It is best to use lower case to write the code, so that there is no conflict with some of the reserved words of C which are in capitals.

The following examples 5.7a and 5.7b uses the C++ shell with an assembly module inserted within parentheses. Observe the way the assembly module is written. It starts with `asm {`. The opening bracket ‘{’ should be in the same line as the word `asm`. This assembly module simply gets a character from the keyboard with `echo`. It then uses carriage return and newline characters to go to the next line. In the next line, the character entered is displayed. Thus, the same character is displayed on two consecutive lines.

Example 5.7a

```

#include<iostream.h>
#include<conio.h>
int main()
{
    clrscr();
    asm {
        mov ah,01
        int 21h                ;enter a character from the keyboard
        mov cl, al             ;move it to cl
        mov dl, 0ah            ;character for 'newline'
        mov ah, 02             ;display newline
        int 21h
        mov dl, 0dh            ;character for 'carriage return'
        mov ah, 02             ;display 'carriage return'
        int 21h
        mov dl, cl              ;display the character which was preserved in cl
        mov ah, 02
        int 21h
    }
    getch();
    return 0;
}

```

Another way of writing assembly modules is to use the keyword 'asm' before each line of the in-line assembly code. Example 5.7b is the re-written version of Example 5.7a, in this format.

Example 5.7b

```

#include<iostream.h>
#include<conio.h>
int main()
{
    clrscr();
    {
        asm mov ah, 01
        asm int 21h                ;enter a character from the keyboard
        asm mov cl, al             ;move it to cl
        asm mov dl, 0ah            ;character for 'newline'
        asm mov ah, 02             ;display newline
        asm int 21h
        asm mov dl, 0dh            ;character for 'carriage return'
        asm mov ah, 02             ;display 'carriage return'
        asm int 21h
    }
}

```

```

asm mov dl, cl           ;display the character which was preserved in cl
asm mov ah, 02
asm int 21h
}
getch();
return 0;
}

```

In all the programs here, the screen is cleared with the function `clrscr()`.

The scratchpad registers (A, B, C and D) usually are not directly used by the C compiler. However, the other registers may be in use, and hence, it would be safe to push them on stack (and pop it later) before embarking on running the assembly module.

Now, let us see another program that defines two characters 'a' and 'b'. A character (char) defines an 8-bit number, and hence, in the assembly module, it can be loaded into an 8-bit register. In this program when two single digit numbers are entered through the keyboard, the sum is displayed. For example, for numbers 4 and 5, the display is $4 + 5 = 09$. For numbers 7 and 8, it will be $7 + 8 = 15$.

Example 5.8

```

#include<iostream.h>
#include<conio.h>
int main()
{
clrscr();
char a;
asm {
mov ah,01
int 21h           ;enter the first number with echo
mov a,al          ;move it to a
mov dl, '+'       ;display '+'
mov ah, 02
int 21h
mov ah, 01        ;enter the second number with echo
int 21h
mov ah, 0         ;ah = 0
add al, a         ;add the two ASCII numbers
aaa              ;adjust ASCII after addition
add ax, 3030h     ;convert the sum back to ASCII
mov bx, ax        ;save the ASCII sum in bx
mov dl, '='       ;display '='
mov ah, 02
int 21h
mov dl, bh        ;display the upper ASCII character
mov ah, 02
int 21h
mov dl, bl        ;display the lower ASCII character
mov ah, 02

```

```

int 21h
}
getch();
return 0;
}

```

We see that the constructs of C such as ‘for’ and ‘while’ can be used in the C shell. The next program uses a loop. The assembly module performs only the input and output functions of entering and displaying. See Example 5.9.

Example 5.9

```

#include<iostream.h>
#include<conio.h>
int main()
{
clrscr();
char i,j = 2;
for (i = 0; i <= 5; i++)
asm {
    mov al,j                ;part for displaying
    add al,30h
    mov dl,al
    mov ah,02
    int 21h
}
    getch();
return 0;
}

```

The next program shows how a character string is displayed.

Example 5.10

```

#include<iostream.h>
#include<conio.h>
int main()
{
clrscr();
char const*MESS ="This is my house\n$";
asm{
mov ah,09
mov dx,MESS
int 21h
}
    getch();
return 0;
}

```

KEY POINTS OF THIS CHAPTER

- In the use of an input port, $\overline{RD}$ is the control signal of importance. The machine cycle initiated is the I/O Read machine cycle.
- In the use of an output port, $\overline{WR}$ is the control signal of importance. The machine cycle initiated is the I/O. Write machine cycle.
- An I/O port can have 8 or 16 as the width of its data bus.
- An I/O port can have 8 or 16 as the width of its address bus.
- There are two schemes for I/O interfacing i.e., memory mapped I/O and I/O mapped I/O.
- There are two types I/O addressing – fixed port and variable port.
- If the address is 8 bits, fixed port addressing is to be used; if the address is 16 bits, variable port addressing is to be used.
- An I/O port can have only a maximum address size of 16 bits.
- The 8086 can address 2^{16} input ports and 2^{16} output ports. There are directives which allow programs in different modules to be linked together.
- Modular programming is very useful when it is necessary to use the service of frequently used modules.
- In modular programming, the directive PUBLIC indicates that the referred data item or procedure is accessible to other modules.
- The directive EXTRN shows that the referred item is in another module.
- The inline assembler is used to insert assembly sequences within a CC++ shell.

QUESTIONS

1. Differentiate between memory mapped I/O and isolated I/O.
2. What is meant by fixed port addressing?
3. Why are interfacing chips used between the processor and peripherals?
4. Why is modular programming important?
5. How can a data item be made visible to other modules?
6. What is the relationship between the EXTRN and PUBLIC directives?
7. How do we link two modules using MASM?
8. Why do not we use the small model when we have intersegment calls?
9. If we do not use the directive EXTRN for a label defined elsewhere, what error message is expected and why?

EXERCISE

1. In an 8086 temperature monitoring system, four points are connected as an input port with address 56H. If the data in these lines goes above the number 1000 (binary), an alarm has to be sent. This is done by sending the character 'A' to an output port with address 9FC3H. Write a program for this.
2. Test one data pin of an input port continuously. If it is low, send a character 'N' to an output port with address 78H, and continue monitoring the pin. When the input pin goes high, send the

character 'Y' to the output port and stop monitoring the pin of the input port. The address of the input port is 90CDH.

3. Write a program which pushes two values onto the stack (define stack here). Call a procedure in a different module which accesses the values from the stack, adds them and also multiplies them (refer Example 4.10). Then use a procedure in another module to display the sum and product. Appropriate messages may be incorporated in the display.
4. Write modules to solve the following problem. The first module enters from the keyboard two 'two-digit' decimal numbers. These numbers are to be multiplied, and the product is to be displayed in decimal form (i.e., ASCII). Thus this module calls two external procedures stored in two different modules – the first for converting ASCII to binary form (refer Example 4.17), and the second for converting binary to ASCII and display (refer Example 3.22).
5. Solve this problem using different modules that contain appropriate procedures.
 - a) Enter 10 two-digit decimal numbers through the keyboard.
 - b) Convert them to binary and store them in memory.
 - c) Find the biggest number.
 - d) Display the biggest number.
6. Refer to problem no. 5 above. Sort the above numbers in descending and ascending order and display the sorted numbers.
7. Use assembly instructions within C shell for the following programs:
 - a) Enter two digits through the keyboard and add them. If the sum is greater than 10, output a message indicating that, otherwise display the sum. The display and entering programs should use assembly instructions while the comparison (if scheme) uses C programming.
 - b) Enter a number of digits through the keyboard. Find the sum and display the sum.
 - c) Enter a six-digit number through the keyboard. Display it. Then display it reversed.

6 THE HARDWARE STRUCTURE OF 8086



In this chapter, you will learn

- The differences between 8086 and 8088.
- The pin functions of 8086 in the minimum and maximum modes.
- The use of address latches and data buffers in an 8086 based system.
- The concept of machine cycles and the associated bus timings of 8086.
- The reason for using the bus controller IC in a maximum mode system.
- How an 8086 is used in the maximum mode.
- To calculate instruction timing and thus create delay loops.

We have so far dealt exclusively with the programming aspects of 8086. It is now time to take a look into the hardware aspects of the processor.

6.1 | Pin Configuration

We will try to understand the processor in terms of its hardware, which includes the pin configuration, functions of the pins, memory bus cycles, connections to other chips, modes of operations and so on. It was mentioned in Chapter 0 that the first PC was built with the 8088 processor. This processor is in most ways very similar to the 8086. Following are the minor differences:

- 8088 has an external data bus of 8 bits, while 8086 has a 16-bit data bus.
- The instruction queue of 8088 is 4 bytes, while 8086 has a 6-byte queue.
- For 8088, pin no. 28 (in minimum mode) is an $\overline{M}/IO$ pin, while for 8086, the polarity of the pin is $M/\overline{IO}$.
- Pin No. 34 is $\overline{SSO}$ for 8088, while it is $\overline{BHE}/S7$ for 8086.

The pin diagrams of both these processors are given in Figs 6.1a and 6.1b. We see that both these processors are packaged as 40-pin DIP (dual-in-line package). It is also seen that some pins have dual functions, and this corresponds to two modes of operation – minimum and maximum modes. The minimum mode is used when the 8086/8088 is used in single processor systems, and the maximum mode is used when the system is a multi processor system, in which the 8086/8088 is one of the processors. Now, neither 8086 nor 8088 is used in PCs, because of the advent of the




Figure 6.1a | 8088 CPU pin diagram **b** 8086 CPU pin diagram

advanced processors of the same family. However, we will discuss the hardware aspects of 8086, because it is the basic processor in the x86 family and the features of the advanced processors are upward extensions of the 8086 features. As such, understanding the features of 8086 is very fundamental to the understanding of the advanced processors of the x86 family.

6.1.1 | Minimum Mode Pins

Let us discuss the minimum mode operation of 8086. In Fig 6.1b, pins 24 to 31 are seen to have dual functions. The functions in brackets indicate the pin designation in the minimum mode, for those pins. The chip with minimum mode pin designations is re-drawn in Fig 6.2. For minimum mode operation, pin no. 33 $\overline{MN/\overline{MX}}$ should be at logic high. It was mentioned earlier that 8086 has a 16-bit data bus and a 20-bit address bus. Let us locate the address and data pins, in the pin diagram. We see a set of 16 pins designated as AD. i.e., AD0–AD15. These are the address/data pins. The notation ‘AD’ means that these pins are used for address as well as data. They are multiplexed for data and address, which means that at a particular time they carry address and at other times they carry data. The reason for having multiplexed pins was necessary to reduce the number of pins of the chip and this restriction was more stringent in the early days of microprocessor development. Remember, 8086 has only a total pin count of 40.

Since the data bus and address bus have to be separate and because the ‘address’ has to be available at all times, the address-data lines have to be ‘de-multiplexed’ i.e., separated. This is done by the use of latches, also availing the timing features of bus cycles.




Figure 6.2 | 8086 pins in the minimum mode

6.1.2 | De-multiplexing the Address /Data Bus

Let us assume that memory is to be read from or written to. Recollect the sequence of steps for reading and writing memory (Section 0.7). The first step in (say) reading memory is ‘placing the address on the address bus’. This is, in practice, the first part of the bus cycle. Next, the data is to appear on the data bus, which means that the bus must be made free to carry data. For de-multiplexing, during the time that the address is on the address/data bus, it is latched on to a ‘latch IC’ whose clock is the ALE (Address Latch Enable) signal supplied by 8086. The signal ALE goes high and functions as a clock for a latch that is used to save the address values, because the address has to be available throughout the bus cycle. Once the address is latched on to the latches, it can be removed from the AD lines, and then these lines are free to carry data.

However, this accounts for only 16 address lines. 8086 has a 20-bit address. In the pin diagram, pin numbers 35 to 38 also cater to address, and they have the designations A_{16} to A_{19} . However, they are also multiplexed with status signals, and they also have to be de-multiplexed. Besides these, there is one more pin which carries address information, and that is the active low line $\overline{BHE}$ (Bus High Enable), which is multiplexed by the status signal S_7 . The function of this signal is explained in Chapter 7. Now, with all this information included, the address/data de-multiplexing block diagram is as shown in Fig 6.3.

6.1.3 | Minimum Mode Pin Functions

Now, let us get to know the functions of each of the pins of 8086. Table 6.1 lists the pin numbers, pin designations and pin functions, in the minimum mode.




Figure 6.3 | Block diagram of the de-multiplexed address /data bus

Table 6.1 | Minimum Mode Pin Functions

Pin No.	Designation	Function	Type															
16 to 2, 39	AD0–AD14, AD15	16 multiplexed address / data lines which carry address when ALE is high, and later functions as data lines D0–D15, when ALE is low	Bidirectional															
35 to 38	A_{19}/S_6 – A_{16}/S_3	Address lines A_{19} to A_{16} , which are multiplexed with status signals S_6 to S_3 . The status bits function are as follows: S_6 – always 0 S_5 – condition of the interrupt flag (IF) S_4 and S_3 show the current segment in use as below: <table><tr><td>S_4</td><td>S_3</td><td></td></tr><tr><td>0</td><td>0</td><td>Extra Segment</td></tr><tr><td>0</td><td>1</td><td>Stack Segment</td></tr><tr><td>1</td><td>0</td><td>Code or no Segment</td></tr><tr><td>1</td><td>1</td><td>Data Segment</td></tr></table>	S_4	S_3		0	0	Extra Segment	0	1	Stack Segment	1	0	Code or no Segment	1	1	Data Segment	Output
S_4	S_3																	
0	0	Extra Segment																
0	1	Stack Segment																
1	0	Code or no Segment																
1	1	Data Segment																
32	$\overline{RD}$	When this signal is low, data can be received from memory or input devices	Output															
29	$\overline{WR}$	When this signal is low, it is an indication that the data on the data lines are available for writing into memory or outputting to output devices	Output															

(continued)

Table 6.1 | (continued)

Pin No.	Designation	Function	Type
19	CLK	This is the clock pin to which a clock with at least 33% duty cycle is to be supplied	Input
21	RESET	This is an active high signal which signals the microprocessor to reset itself provided the pin is held high for at least 4 clock periods	Input
22	READY	For the bus cycle to proceed normally, the READY pin should be found to be at logic high when it is sampled. If it is at logic low, WAIT states are inserted into the current bus cycle	Input
23	$\overline{TEST}$	This pin is used usually when an arithmetic co-processor is in the system. This pin is tested by the WAIT instruction. If the pin is at logic zero, the 'WAIT' instruction becomes a NOP instruction. Otherwise, the processor waits until this pin becomes logic zero	Input
25	ALE	'Address Latch Enable' – this signal goes high in the beginning of a bus cycle and indicates that the multiplexed address bus contains address information	Output
26	$\overline{DEN}$	'Data Enable' – This active low signal functions as an activation signal for the external data bus buffers	Output
27	$DT/\overline{R}$	'Data Transmit/Receive' – The logic value of this signal indicates whether the data is received (in a read cycle $DT/\overline{R} = 0$) or transmitted (in a write cycle $DT/\overline{R} = 1$). Thus, it is used as a direction pin for external data bus buffers	Output
28	$M/\overline{IO}$	For I/O access this pin is low, and for memory access, it is high	Output
31	HOLD	This is a signal from a peripheral requesting direct memory access (DMA). If the signal is high, the processor issues a Hold Acknowledge signal and tri-states its data, address and control bus	Input
30	HLDA	'Hold Acknowledge' indicates the acknowledgement of the HOLD request	Output

(continued)

Table 6.1 | (continued)

Pin No.	Designation	Function	Type
33	$MN/\overline{MX}$	This pin is used to select the mode of operation – minimum or maximum. For minimum mode, the pin is to be connected to the 5 V supply	Input
34	$\overline{BHE}/S7$	The ‘Bus High Enable’ is a logic low signal which is used to enable the ‘high memory bank’ of the data bus – the D_8-D_{15} lines become active then. The status of this pin is latched along with the address information. This is multiplexed with the status pin S_7 , which is always 1	Output
1, 20	GND	‘Ground’ – The common point is to be connected to two pins. Two Ground pins are used so as to prevent having to connect them together internally, due to possible noise in the internal routing of the pins	—
40	VCC	The power supply must be +5V +/–10%	—
18	INTR	Interrupt Request – this is used by an external device to interrupt the processor, which responds, only if the interrupt flag (IF) is set, by lowering the $\overline{INTA}$ line, and initiating an ‘interrupt acknowledge’ machine cycle input	Input
24	$\overline{INTA}$	‘Interrupt Acknowledge’ is an active low signal acknowledging the interrupt request placed on the INTR pin	Output
17	NMI	Non maskable interrupt request which is placed by an external device – similar to INTR, but the Interrupt Flag (IF) does not have to be set for it to be serviced. It is a high priority interrupt	Input

6.1.4 | The De-multiplexed Address/Data Bus

6.1.4.1 | The Address Bus

The idea of de-multiplexing has already been discussed in Section 6.1.2. Now, let us get into the finer details. To ensure that the address is available on the memory or I/O address lines throughout the complete read/write bus cycle, the 20 lines of address information and the line $\overline{BHE}$ have to be latched, and the output of the latch will constitute the address bus. Thus 21 lines are to be latched. Obviously, external latch ICs will have to be used. One very popular latch IC is the three state octal transparent latch 74LS373. Since one IC can cater only to 8 lines, 3 such latches are needed for de-multiplexing. Fig 6.4 shows the pin diagram of this IC.




Figure 6.4 | Pin diagram of the octal latch IC 74LS373

The data sheet of this IC specifies that ‘the eight latches of the 74LS373 are transparent D-type latches meaning that while the enable (G) is HIGH the Q outputs will follow the data (D) inputs. Thus the ALE signal is connected to the G pin of the IC. The $\overline{OC}$ (output control) is grounded for normal operation. If not, the output lines are tri-stated. The IC also provides a buffered output, meaning that it has the capability to drive high capacitance buses that are normally encountered and the current driving capability of the address lines is raised so that more number of TTL loads may be driven.

Note The LS (Low Power Schottky) series of ICs offer a significant reduction in power as well as increases in speed over TTL.

Fig 6.5 shows the complete diagram of the 8086, connected to three latch chips. The outputs of the latches constitute the address bus.

6.1.4.2 | The Data Bus

In Fig 6.5, the data lines D_0 – D_{15} are shown available directly from the pins AD_0 to AD_{15} . However, in practice, these lines do not have sufficient current driving capability for many (more than 8 to 10) TTL loads and the capacity to drive capacitive loads. Most memory devices are MOS chips, which have capacitive inputs, which have to be charged or discharged when logic level changes occur. If the pin fan out is high, the processor cannot supply the necessary current for this. Hence, buffers/line drivers which raise the current driving capability of these pins are used. Moreover, these lines also have to be isolated when not in use. Hence, the buffers have to be tri-state buffers.

The data lines are bidirectional – the direction depending on whether the data is being read or written. Fig 6.6 shows a popular octal bidirectional tri-state buffer, the 74LS245. For 16 data lines, two such ICs are needed. Two pins of the 8086 act as enable and direction pins to be connected to the buffer. The $DT/\overline{R}$ line is high for data write and low for data read. This pin of 8086 is connected to the DIR (Direction) pin of the buffer. This fixes up the direction of data transfer. Another pin $\overline{DEN}$ of the processor is connected to the Enable line ($\overline{G}$) of the buffer. Only if this pin is at low level, will the output lines of the buffer be enabled, otherwise the lines are tri-stated. The bus cycle timing will ensure the activation of these signals at the appropriate time for data transfer. We will soon get to the discussion of bus cycles.

6.1.5 | Control Signals for Read and Write

Before getting to the discussion on bus timing, it will be worthwhile to mention that the associated control signals $\overline{RD}$, $\overline{WR}$ and $M/\overline{IO}$ may also be buffered, but since these are output signals sent out by 8086, these lines are unidirectional, and hence the buffer IC74LS244 will be sufficient for use here. This is a




Figure 6.5 | The buffered and de-multiplexed address bus




Figure 6.6 | Pin diagram of the octal tri-state bi-directional buffer 74LS245

unidirectional tri-state octal buffer. Thus, the diagram of the processor with the buffered data bus and control signals is given in Fig 6.7.

6.1.6 | Memory and I/O

The signals $\overline{RD}$, $\overline{WR}$ and $M/\overline{IO}$ have to be used suitably to get the memory read/write signals $\overline{MEMRD}$ and $\overline{MEMWR}$ and the I/O read/write signals $\overline{IORD}$ and $\overline{IOWR}$. See the circuit diagram in Fig 6.8. With this, I/O and memory have separate lines for access. Generation of these control signals may be done with any logic gates.

6.1.7 | Clock Generation

It has been emphasized that all activities within a processor are synchronized with a clock. A clock is a square wave signal of fixed frequency. All timings of activities done by the processor are calculated based on this clock period. One cycle of the clock is called a T state, and all timings and delays are multiples of this T state duration.

Thus, it is obvious that the 8086 has to have a clock signal. There is no circuitry inside the processor for this, and so an external clock generator IC is used. Intel has provided the clock generator IC 8284A which is compatible with 8086/8088. This generates a clock of frequency depending on the crystal connected between two of its terminals. The crystal frequency must be three times the desired frequency for the microprocessor. 8086 operates at frequencies of 5 to 15 MHz; depending on the requirement the appropriate crystal may be used. The CLK output generates a 33% duty cycle designed to drive the 8086 processor directly. The pin diagram of the clock generator IC is given in Fig 6.9.

The clock generator IC performs a few more functions than just supplying the clock frequency to the processor. We will discuss only a few basic functions that are in use for a minimum mode system. See the connection diagram in Fig 6.10. Note that the crystal is connected between the X_1 and X_2 pins.

6.1.8 | Ready

The READY output pin of 8284A is connected to the READY input pin of 8086. This pin is used to overcome the timing inconsistencies that are possible when a slow peripheral/memory device is connected to the processor. The slow device may not be able to complete its read or write within the normal timing period of the processor.




Figure 6.7 | Buffered data bus and control pins




Figure 6.8 | Generation of memory and I/O read/write signals




Figure 6.9 | Pin configuration of the clock generator 8284A




Figure 6.10 | Connections of the clock generator IC to 8086

In that case, it asks for extra time to complete its activity i.e., extra ‘wait’ states are inserted into the bus cycle. These extra wait states are inserted when the READY signal is low. For slow peripherals/memory, a wait state generator can be added in the system. If wait states are not needed, READY is found to be high when it is sampled. However, why would READY be connected to the clock generator IC? This is to synchronize it with the clock signal. READY may appear at any time, but the 8284A sends it to the 8086 only at the trailing edge of the clock.

6.1.9 | Reset

An active low $\overline{RES}$ signal from the control bus is sent to the 8284 which synchronizes it with the trailing edge of the clock. Most systems include a line that goes to all system components and possibly controlled by an operator push button (or just after power on), which causes a low signal on the $\overline{RES}$ pin of the clock generator. The reset circuitry of 8284A makes sure that the

timing requirements of reset of 8086 are met i.e., the RESET line of 8086 must remain high for at least 4 clock periods. The 8086 will terminate operations on the high-going edge of RESET and will remain dormant as long as RESET is HIGH. A high on the RESET pin of the processor causes all system components to be reset, and inside it, the instruction queue, PSW, DS, SS, ES and IP are cleared. CS gets a value of FFFFH and with IP = 0000, the first instruction will be executed from the location FFFF0H. However, normally, this address will be found in a read-only section of memory, and the code here will be a jump instruction to a program for loading the operating system. Such a program is called a bootstrap loader.

6.1.10 | Power on Reset

When power is first switched on, the 8086 should be reset. This is called power on reset and Fig 6.11 shows the 'power on reset' circuit. Observe the circuitry at the pin $\overline{RES}$. There is an RC circuit and a push button switch too. Assume the push button switch to be in the open condition now. When the 8086 is first switched on, power is applied. Then the capacitor in the circuit starts charging as shown in Fig 6.12. Thus V_C increases and the voltage at the $\overline{RES}$ increases. Inside the 8284 chip, this $\overline{RES}$ is the input to an inverted Schmitt trigger as shown in Fig 6.13. Till the voltage V_C increases to V_H , the Schmitt trigger output is high. Once it crosses the value of V_H , it goes low. The high pulse from the Schmitt trigger output is the RESET pulse to be applied to the 8086. In Fig 6.13, it is applied to the D input of a flip flop and hence it appears at the flip flop output, in synchronism with the clock.

It is specified that the 8086 reset pulse should be high and should have a period of at least 4 clock cycles. How is that ensured? It is ensured simply by designing the time constant of the RC circuit. Let's try a simple calculation.

The Schmitt trigger switches to low, once its input crosses 1.7V (as per its specifications). The equation for the charging of the capacitor is:

$$V_C = V_{CC} [1 - \exp(-t/RC)].$$

If V_C is to be at least 1 V, $t = 50 \mu\text{secs}$ and $V_{CC} = 4.5 \text{ V}$, then $RC = 188 \mu\text{secs}$.

So, if $C = 10 \mu\text{F}$, then $R = 18 \text{ K}$.

These are close to the values we have used in Fig 6.11.

By this, we are designing a reset pulse of $50 \mu\text{secs}$, which is higher than the minimum requirements of the reset pulse for the clock frequency of 5 MHz shown in Fig 6.11. The reset




Figure 6.11 | 'Power on reset' circuit




Figure 6.12 | Charging of the capacitor and generation of reset pulse




Figure 6.13 | The generation of the reset pulse using a Schmitt trigger (inside the 8284 IC)

output of the clock generator can be used for resetting other devices in the system along with the 8086 reset, and its time period $t = 4T$ (at least) where T is the clock period of the processor.

Why the push button switch? To start a new reset cycle – if the switch is pressed, the capacitor discharges, and to re-start, the whole process should start again.

When using a crystal for frequency generation, the $F/\overline{C}$ pin of the 8284 IC must be grounded. If this pin is high, the EF1 pin should be connected to an external oscillator and a pin CSYNC is used to provide synchronization in this case. Since, we are using a crystal in our circuit, the pins $F/\overline{C}$ and CSYNC are grounded in Fig 6.11.

6.2 | Clock

It has been mentioned that all activities on the system bus is synchronized by the system clock which provides the basic timing. Now, what are the activities on the bus?

Let us list it out.

- i) Reading from memory/IO
- ii) Writing to memory/IO

These are obviously two fundamental bus activities. Depending on the type of activity, different control signals are activated. Any read or write cycle is called a bus cycle. For 8086, a bus cycle takes 4 T states, where one T state is defined as the 'period' of the clock. If the clock frequency is 10 MHz, one T state = $0.1 \mu\text{secs}$ or 100 nsecs. A bus cycle is also called a machine cycle. During a machine cycle, a specific operation – say, reading or writing is accomplished. So, we can have the following basic machine cycles:

- i) Memory Read
- ii) Memory Write
- iii) I/O Read
- iv) I/O Write

There is also the 'Interrupt acknowledge machine cycle', but that will be discussed in Chapter 8.

6.2.1 | Read Machine Cycle

In Section 0.2.5, the steps involved in a typical **read machine cycle** were listed. Let us recollect them.

- i) Place on the address bus, the address of the location whose content is to be read. This action is performed by the processor.
- ii) Assert the **read** control signal which is part of the control bus.
- iii) Wait until the content of the addressed location appears on the data bus.
- iv) Transfer the data on the data bus to the processor.
- v) De-activate the read control signal. The read operation is over and the address on the address bus is not relevant anymore.

This sequence is followed in the read machine cycle of 8086 as well. However, here, we have to examine the exact signals of the processor. Many control signals are necessary and are shown in the timing diagram. Refer to the read machine cycle timing diagram in Fig 6.14 when reading the steps given below. This diagram is relevant for both 'memory read' and 'I/O read'. In the diagram, time is marked along the horizontal direction, and the activities on the pins of 8086 at corresponding points in time are shown. Some of the lines are single signals like ALE, $\overline{RD}$, $\overline{WR}$ etc, while others are groups of signals like the address bus and data bus. $M/\overline{IO}$ is a single line, which has one of two possible states in a machine cycle. '**Crossing lines**' show that the data in those lines become valid at that point in time, and also that the logic levels in the lines change. The T states are labeled as T_1 , T_2 , T_3 and T_4 . Each T state is defined from the 50% level of the trailing edge of the clock to the next corresponding point in the clock train.

- i) In T_1 , the address of the location to be accessed is placed on the lines AD_0 to AD_{15} and A_{16}/S_3 to A_{19}/S_6 . BHE is high or low depending on the type of data (byte/word) to be accessed.




Figure 6.14 | Read machine cycle timing diagram

- ii) In T_1 , the ALE signal goes high for one clock cycle, and on its trailing edge, the address information is latched and available at the output of the latches (Refer Section 6.1.2).
- iii) If it is a 'memory read', the $M/\overline{IO}$ signal is high from T_1 to T_4 . If it is 'I/O read', this signal is low. Also the $DT/\overline{R}$ is low from T_1 to T_4 indicating that data is to be 'received' by the processor (since it is a read cycle).
- iv) In T_2 , the address information is removed from AD_0 to AD_{15} and the bus is tri-stated (floated). The address information is removed from A_{16}/S_3 to A_{19}/S_6 and $\overline{BHE}/S_7$ also, and these lines now carry status information.
- v) In T_2 , the $\overline{RD}$ signal (which is the READ control signal) is made low. The slanting line shown to the data bus is indicative of a 'cause-effect' relationship – here, when the READ control signal goes low, the output lines of memory come out of the 'float' state.
- vi) For a normal read cycle, valid data appears on the data lines after the period defined as the 'access time' of the memory/IO.
- vii) In T_2 , $\overline{DEN}$ goes low to enable the data bus buffer outputs.
- viii) At the end of T_2 , the READY signal is sampled and if it is high, the bus cycle proceeds normally. This is the case shown in this diagram.
- ix) The clock cycle T_3 is to allow 'access time' for bringing data from memory/IO and put it on the data bus. Thus, after the data appears on the data bus, it is transferred to the 8086 (to the relevant registers) at the end of T_3 .
- x) In T_4 , all the bus signals are de-activated in preparation for the next bus cycle. This machine cycle ends with T_4 , and the next machine cycle is scheduled to start.

6.2.2 | Wait Cycles

If the access time for a device is longer than that permitted by 8086 timing, extra clock cycles termed ‘wait states’ have to be inserted in the bus cycle. This is applicable for reading and writing. The arrangement for this is to sample the READY line at the end of T_2 . If the READY signal is found low, an extra T state is inserted into the bus cycle, between T_3 and T_4 , which is designated as T_w . All signals on the bus remain unchanged during this extra wait state (T_w). In the middle of T_w , once again the READY signal is checked. If it is at logic 1, the next T state will be T_4 – otherwise another wait state T_w will be inserted. Thus bus cycles can be lengthened to accommodate slower devices in the systems. In such cases, a ‘wait state generator’ is connected to the READY pin of the clock generator to control the READY pin. Refer to the Intel manual in Appendix A, for a more detailed timing diagram of READ with exact time values given.

If no wait states are needed, then it will be sufficient to keep the READY pin of 8086 high. Another way would be to connect the RDY1 and RDY2 pins to low and the $\overline{AEN1}$ and $\overline{AEN2}$ to high which will permanently disable this function. For more details on these pins, refer to the data sheet of the clock generator IC 8284A. To insert wait states, the READY can be made low at any time (Fig 6.15), but it should be found to be low at the end of T_2 .




Figure 6.15 | Sampling the READY signal and inserting wait states




Figure 6.16 | Insertion of a wait state generator to add wait states

Example 6.1

What is the duration of the bus-cycle in an 8086 based microcomputer, if the clock frequency is 12 MHz and three wait states are inserted?

Solution

The period of a 12 MHz microprocessor is $T = 1/f = 83\text{ ns}$.
Thus, the duration of the bus-cycle, without any wait-states is given by;
 $T_{\text{bus-cycle}} = 4 * T = 4 * 83\text{ ns} = 332\text{ ns}$.
Duration of the wait-states is, $T_w = 3 * T = 249\text{ ns}$.
So the extended bus-cycle $T_{\text{bus-cycle}} + T_w = 332 + 249 = 581\text{ ns}$.

6.2.3 | Write Machine Cycle

The steps in a write machine cycle are (refer Section 0.2.5):

- i) Place on the address bus, the address of the location to which data is to be written.
- ii) On the data bus, place the data to be written.
- iii) Assert the **write** control signal which is part of the control bus.
- iv) Wait until the data is stored in the addressed location.
- v) De-activate the memory write signal. This ends the memory write operation.

For the 8086, the control signals for a write machine cycle are the same as in a READ cycle except the $\overline{WR}$ is used in place of $\overline{RD}$. Also the $DT/\overline{R}$ signal will be high (for ‘data transmit’) for writing. This being the case, the complete write machine cycle is not shown here – only the timing of placing of the address and data on the buses, and the lowering of the write control signal is shown here. See Fig 6.17. The write bus cycle also uses 4 T states normally, but can be extended by adding wait states if required.

6.3 | Other Processor Activities

6.3.1 | Interrupt Lines

The word ‘interrupt’ implies that something that is being done is to be temporarily stalled, so as to take up another activity. This is the case for a processor too. When the processor is performing




Figure 6.17 | Simplified write machine cycle timing

some activity, an external device can ask for attention by placing a signal on an ‘interrupt request’ line. Then the processor stops whatever it is doing and attends to the request of the interrupting device. When the processor gets an interrupt request on the INTR line, it responds by lowering the *INTA* (Interrupt Acknowledge line) and enters an Interrupt acknowledge machine cycle.

NMI (Non Maskable Interrupt) is another line on which an external device can place an interrupt request. One difference between the two is that for an interrupt request on the INTR line to be acknowledged and taken up, the interrupt flag should be set (IF = 1), while this is not necessary for a request on the NMI line. The details of interrupt processing will be discussed in Chapter 8. Look at the pin diagram (Fig 6.2) for these pins. (We will also discuss why the RESET pin can also be called an interrupt pin.)

6.3.2 | DMA

DMA stands for ‘direct memory access’. So far, we have seen the processor communicating with memory or I/O i.e., the communication always involved the processor. Is it possible for data in memory to be sent directly to I/O or vice versa without involving the processor? For example, we might need to print a large chunk of data which is in memory. This data can be sent to the output device (printer) from memory. In this case, there is no necessity to involve the processor in the data transfer. This is called ‘direct memory access’. However, since the processor is in the system, it must be isolated from this process. This means that, when DMA is to take place, the connection of the processor to memory and I/O is blocked i.e., the buses of the processor are to be tri-stated (in the high impedance state), leaving the path open for data to be transferred directly between I/O and memory. Fig 6.18 shows the typical scenario. Data transfer occurs between the memory and I/O, in either direction. The processor cannot do any bus-related operation, as its buses are in the float state. This whole process is managed and controlled by a DMA controller, which we will discuss in detail in Chapter 11. At this point, what needs to be understood is that when DMA operation is to be done, the DMA controller places a DMA request on the HOLD pin of 8086. This is essentially a request for permission to take control of the system bus. Since DMA is a high priority service, the processor stops whatever it is doing and acknowledges the request by raising the HLDA (hold acknowledge) line and DMA is initiated. See Fig 6.2 for these pins. Fig 6.19 shows the timing associated with the HOLD and




Figure 6.18 | Concept of direct memory access




Figure 6.19 | Timing of the HOLD and HLDA signals

HLDA lines. The HOLD pin is sampled at every rising edge of the clock, and the acknowledge signal is activated at the end of the current machine cycle i.e., at T_4 . In this state, the address, data and control pins of the 8086 are in the high impedance state, and this continues until the HOLD request is taken away by the requesting device.

6.3.3 | $\overline{TEST}$

This pin can be used as an alternative to interrupts. However, usually, this pin is not used in the minimum mode of operation where 8086 operates in a single processor environment. In the maximum mode, when an arithmetic coprocessor is also present in the system, this pin is needed to synchronize the activities of both the processors. The two processors normally can execute separate instructions, but if the 8086 needs a result from the co-processor before it can proceed further, it executes a WAIT instruction, which samples the $\overline{TEST}$ pin. This is because this pin is connected to the BUSY pin of the coprocessor, which is high when the coprocessor is executing instructions. Then the 8086 goes into an idle state where it simply waits until the $\overline{TEST}$ becomes low (this happens when the required result is obtained from the co-processor). After that, the 8086 can continue its normal sequence of operations. If the $\overline{TEST}$ pin is low, the WAIT instruction is a NOP for 8086. The communication between the 8086 and the co-processor is discussed in detail in Chapter 13.

6.3.4 | Bus High Enable ($\overline{BHE}$)

This active low signal is used to activate the higher order memory bank, the details of which are discussed in Chapter 7. Now that all the signals in the minimum mode have been discussed, the minimum mode configuration can be observed and understood from Fig 6.20.

6.3.5 | Halt Machine Cycle

Another machine cycle is the HALT machine cycle. The processor enters this machine cycle in response to a HLT (Halt) instruction. In the minimum mode, the processor issues an ALE signal with no other control signals. In the maximum mode, the processor signal causes the bus controller (will be discussed in Section 6.4) to generate the ALE signal. Since no other signals are generated, the bus cycles are halted. To bring the processor out of this state, an interrupt or a RESET signal must be issued.




Figure 6.20 | 8086 in the minimum mode configuration

6.4 | Maximum Mode

Now let us discuss the maximum mode pins. It is necessary to use the maximum mode if the processor is to be used in multiprocessor configurations. The most important issues in a multiprocessor environment are inter-processor communication and bus contention. In this mode, 8086 has special pins for resolving these issues. To use the maximum mode, pin No. 33 $MN/\overline{MX}$ must be connected to ground. Then, pin nos. 24 to 31 have different designations from what we saw in the minimum mode. See Fig 6.21. In this mode, some of the most important control functions are not obtained from the processor, but will have to be generated from an external chip called bus controller. The functions of $\overline{INTA}$, ALE , $DT/\overline{R}$, $M/\overline{IO}$, and $\overline{WR}$ are to be generated externally.

This mode of operation was designed by Intel for allowing the 8086 to communicate with other processors like the arithmetic co-processor (8087) and the input/output processor 8089. It was also used to allow 8086 to be used in large loosely coupled multiprocessor systems. This mode was dropped from Intel's designs from 80286 onwards. Later processors (80486 onwards) had the arithmetic co-processor integrated on the processor chip itself.

Now, let us see how the lost signals of the processor are generated by the bus controller. The bus controller 8288 was designed by Intel to be compatible with the 8086.

6.4.1 | Bus Controller

In the maximum mode configuration, pin nos. 26, 27 and 28 are designated as $\overline{S}_0$, $\overline{S}_1$ and $\overline{S}_2$. On these pins, the 8086 outputs a code that specifies the type of bus cycle to be initiated.

The listing of the codes, corresponding machine cycles and the important control signals generated by the 8288 is shown in Table 6.2.




Figure 6.21 | Maximum mode pins of 8086

Table 6.2 | Control Signals Generated by the Bus Controller

Status signals			Machine cycle	Control signal generated by 8288
S2	S1	S0		
0	0	0	Interrupt acknowledge	INTA
0	0	1	I/O read	IORC
0	1	0	I/O write	IOWC
0	1	1	Halt	NONE
1	0	0	Instruction fetch	MRDC
1	0	1	Memory read	MRDC
1	1	0	Memory write	MWTC
1	1	1	Inactive	NONE




Figure 6.22 | Pin diagram of the bus controller IC 8288

The connection diagram between the 8086 and 8288 is shown in Fig 6.23.

Note that the clock generator applies the same clock to both the ICs. Thus, according to the code on $\overline{S}_0$, $\overline{S}_1$ and $\overline{S}_2$, the appropriate signals are generated and can be used to read/write memory or I/O, and provide the necessary signals for latches and transceivers.

Note that the bus controller generates the read and write control signals for memory and I/O, the ALE signal for the latches, the $DT/\overline{R}$ and DEN signals for the transceivers and the $\overline{INTA}$ signal for acknowledging interrupts.

There are a few more pins for the bus controller IC, but we shall not discuss those, because such maximum mode systems have become obsolete. However, to complete our understanding of the 8086, let us see the rest of the maximum mode pins.

6.4.2 | Request/Grant pins

Pin numbers 30 and 31 are designated as $\overline{RQ}/\overline{GT}$. These are special pins, in the sense that they are bidirectional. As the name implies, they cater to ‘request’ and ‘grant’. Each pin acts similar to the HOLD and HLDA pin. In the minimum mode, HOLD is for bus request, and HLDA is for bus grant. In the maximum mode, one pin performs both functions. Here the request/grant pins are more likely to be used for communicating with other processors, than for DMA operations. Remember that this mode is used in multiprocessor configurations i.e., systems in which there are other processors besides the 8086. In multiprocessing systems, resources like memory and I/O are likely to be shared and the system bus is common to all processors. When one processor is using the bus, another processor may request the use of the bus and the current bus master may relinquish the bus (subject to conditions). Thus, requests from another processor will be placed on these lines and the grant signal is also placed on these pins, but in the opposite direction, obviously. That is how these pins are ‘bidirectional’. There are timings associated with requests and grants, and are shown in Fig 6.24.

Three pulses are needed to complete a request/grant/release of the bus. The $\overline{RQ}/\overline{GT}$ pins are examined at the rising edge of each clock pulse. If a request (low logic for one T state) is detected, and if the conditions are conducive for granting the request, the grant pulse appears on the same pin immediately following the next T_4 or T_1 state. This allows the requesting processor to use the bus. When this processor is ready to return the bus after use, it sends the release pulse




Figure 6.23 | Connection between the 8086 and 8288




Figure 6.24 | Timing associated with the $\overline{RQ}/\overline{GT}$ pin

on the same line. Incidentally there are two such pins and in a maximum mode system, if bus requests are received on both the pins, the pin $\overline{RQ}_0/\overline{GT}_0$ will get the priority. This can happen if there are two other processors in the system besides the 8086.

6.4.3 | Queue Status pins QS₀, QS₁

These pins are inputs to the 8086. It becomes useful when an arithmetic coprocessor is the second processor in the system. Since the co-processor is expected to work in step with 8086, the co-processor can interrogate the 8086 about its queue status, on these lines, and decide its course of action accordingly. More details are given in Chapter 13.

6.4.4 | $\overline{LOCK}$

This is an output signal from the 8086. If this processor wants to prevent any other bus masters from accessing the bus, this signal will be asserted low. The 8086 can use an instruction with a LOCK prefix to cause the hardware $\overline{LOCK}$ signal to be asserted. For example, if the 8086 wants

to retain the bus until a string transfer is completed fully, it can use the instruction (say) LOCK REP MOVSB. So the processor does not have to relinquish the bus after one bus cycle, as may be the case if the LOCK prefix is not used. Instead, the bus is retained until the complete string operation is over.

We have discussed all the maximum mode pins of 8086 and the diagram, with all the control signals, is shown in Fig 6.25

6.5 | Instruction Cycle

Now, let us decide what we mean by the word ‘instruction cycle’. The time taken by the processor to execute an instruction is called an instruction cycle, and it is specified in terms of the number of clock cycles needed to do it. However, this is not as direct as it seems. Let us examine the important issues in this. The operation of the CPU is just ‘fetch, decode and execute’. Once an instruction is fetched and is ready for execution, it will be decoded immediately, and after that, execution can be set in motion. The fetch-execute cycle can be decomposed into 6 stages.




Figure 6.25 | Maximum mode configuration of the 8086

FI – Fetch instruction
DI – Decode instruction
CO – Calculate operand addresses
FO – Fetch operands
EI – Execute instruction
WO – Write or store result in memory

The time for all these activities should constitute the ‘instruction cycle’. If, for example, we fetch an ADD instruction which is a 2-byte instruction, the first 8 bits specify this operation. The other bits identify the source and destination operands of our instruction. The ADD instruction may have a memory operand or an immediate operand. If a memory address or immediate operand is specified, there will be optional bytes following the instruction which will have to be fetched. There may also be the requirement of having to calculate the operand address. The BIU will separate the operation type from the operands and fetch any operands from memory if required, and after that, execution is done. In the execution phase, any instruction involving loading from or storing into memory, will cause a read or write cycle to occur. An instruction which moves data between two registers only, or must execute a more complicated instruction with operands available within the processor, is done in the Arithmetic and Logic Unit (ALU) alone.

Now, let us look at the ‘instruction fetch cycle’ once again. The instruction is in memory and fetching it requires a ‘memory read’ cycle which takes 4 T states. However, in 8086, ‘fetch and execute’ are overlapped i.e., instructions are fetched during the execution phase of an earlier instruction, and thus the current instruction is normally available in the instruction queue when it is ready for execution. The result is that, in most cases, the time normally required to fetch instructions “disappears”, because the Execution unit (EU) executes instructions that have already been pre-fetched by the BIU. Thus, the instruction cycle time does not include the time required to fetch the instruction. However, for instructions which have memory operands not specified directly, there is a time involved in calculating the ‘effective address’. Take the case of an addressing mode, for example, based indexed mode or relative based indexed mode. The BIU takes time to do the address calculation, and this time is also included in the instruction cycle.

Another factor for an extra delay to occur is, if an operand which is a word whose address is aligned with an odd address, has to be accessed. In this case, one more memory read/write cycle may be incurred, in comparison to a word which is at an even address (the mechanics of this is discussed in Chapter 7). This is not taken into account in the ensuing calculations.

Now, look at some execution times in Table 6.3. In the first three instructions, all the operands are register operands, and the execution time is just the time required to perform the operation in the ALU. Note that MUL requires quite a lot of time compared to ADD, because it is a complex instruction involving many steps. The fourth instruction is in the immediate addressing mode. In the fifth instruction, there is a memory operand and the execution time is given as $9 + EA$, where EA is the time for calculating the effective address. The EA value for different addressing modes is given in Table 6.4. For instruction 6 and 7, the execution times are $16/4$ and $17/5$. These instructions are conditional branch instructions. The first figure gives the execution time if the branch is taken and the second number corresponds to the case when branching is exited. When a branch is taken, the instruction queue is to be cleared and reloaded. That is why more time is expended when a ‘branch’ is taken, in contrast to the case of execution in the normal sequence.

In short, if all such information is provided, we will be in a position to calculate the ‘instruction cycle’ of an instruction, and note that the time for ‘instruction fetch’ does not need to be

Table 6.3 | Execution Times of Some Sample Instructions

No.	Instruction	No. of clock cycles for execution
1	ADD AX, BX	2
2	MOV AX, BX	2
3	MUL BX	133
4	MOV BX, N	4
5	CMP AX, [BX][SI]	9 + EA
6	JNZ label	16/4
7	LOOP label	17/5

Table 6.4 | Number of Cycles Expended in Calculating the ‘Effective Address’

No.	Addressing mode	No. of clocks for calculation of EA
1	Direct	6
2	Register indirect	5
3	Register relative	9
4	Based indexed with BP as the base register	8
5	Based indexed with BX as the base register	7
6	Relative based indexed with BP as the base register	12
7	Relative based indexed with BX as the base register	11

included herein. The instruction cycle times for each instruction is included along with the instruction set, in Appendix C.

6.5.1 | Delay Loops

Now that we have the means to calculate the time required to execute any instruction of the processor, let us use the execution time for an interesting application. We see that a certain amount of time or rather a ‘delay’ is associated with the execution of an instruction. Thus, instruction execution gives us a means of generating a delay. See the instructions below.

```
MOV CX, 100      4 cycles
HERE:  LOOP HERE  17/5 cycles
```

How many clock cycles are necessary for executing the above program segment? The MOV instruction takes 4 clocks. The LOOP instruction takes 17 cycles for repeating, and 5 when exiting the loop. Thus, totally we have $4 + (17 \times 100) - 12 = 1692$ cycles (12 is subtracted, because the last execution of the LOOP instruction takes 5 cycles rather than 17). In a system with a clock of 12 MHz, one clock period is 0.083μ secs, which makes for a total delay of

140 μ secs delay. Also note that most of the delay is incurred in the repetition of the LOOP instruction. Thus, the delay caused by the peripheral instruction (MOV CX, 100) may be neglected.

This idea may be used to fix up the value of N so as to get a desired value of delay. To increase the delay, we can use the NOP (No operation) instruction which has no operands and no function except to execute within 3 clock cycles. The NOP instruction is usually used to reserve space in programs, for instructions which may needed to be added later. Here we use it to create a delay – any number of NOP instructions may be added.

Example 6.2

Write a program to create a delay of 1 msec.

Solution

	Instruction	No. of cycles
	MOV CX, N	4
HERE:	NOP	3
	LOOP HERE	17/5

As most of the delay occurs within the loop, the total cycles of delay is $= [(3 + 17) \times N] - 12$.

Total delay time = 1 msec = $20 N \times 0.083 \mu$ secs

For 1 msec delay, the value of $N = \frac{1\text{m sec}}{(20 \times 0.083 \mu \text{ secs})} = 602$ or 25AH

This value of N if inserted into the program will give us a delay of 1 msec. There is a small % of error in this, because of having to leave out peripheral instructions, but for most situations this can be neglected, as the delay that we get is anyway only approximate. The maximum value that can be loaded into a register is FFFFH, and if larger delays are needed, a nested delay loop can be used as shown below.

Example 6.3

Write a delay loop with appropriate values of the count to get a delay of 1 second.

Solution

The program is

	Instructions	Cycles
	MOV BX, N1	4
THERE:	MOV CX, N2	4
HERE:	LOOP HERE	17
	DEC BX	2
	JNZ THERE	16

The inner loop is that which corresponds to the LOOP instruction. It repeats N2 times, which is the count in the CX register. The LOOP instruction plus a few overheads (caused by the instructions MOV CX, N2 and DEC BX) repeat N1 times, which is the count of the outer loop.

Thus a large delay can be obtained. To calculate N1 and N2, the best way would be to make N2 = FFFFH and then get a corresponding value for N1.

For a total delay of 1 second,

Total cycles = $N1 \times [17 \times N2 + 4 + 2 + 16] = N1 [22 + 17 N2]$

If N2 = FFFFH i.e., 65,535 the calculation comes to

$1114117 \times N1 \times .083 = 1000\ 000\ \mu\text{ secs}$

$92471.711 \times N1 = 1000000$

$N1 = 10.81$

Thus, a value of 11 (0BH) can be used in the place of N1 in the program.

Note No values are subtracted to take into account the decreased number of cycles when a loop is exited. Thus, this calculation is only approximate, but the error is negligible.

6.5.2 | Why Delay Loops?

What application is there for the exercise that we have just done? Where are these delays to be used? Generating delays in this manner is called 'software delay'. One can generate a square wave using a software delay. Suppose we have an output port with an address of 78H. See Example 6.4.

Example 6.4

Generate a square wave of frequency 1 KHz at the output port with address 78H

Solution

```
AGAIN:  MOV AL, 0FFH
        OUT 78H, AL
        CALL DELAY_1MS
        MOV AL, 00
        OUT 78H, AL
        CALL DELAY_1MS
        JMP AGAIN
```

The above program assumes that a program named DELAY_1MS has been written and used as a procedure along with this. This procedure is like the programs in Examples 6.2 or 6.3, with a count which gives a delay of 1 msec.

If a CRO is connected to any pin of the data bus of the output port, a continuous square wave can be observed.

This is an application of software delay, but in practical cases, we are more likely to use dedicated timer chips for generating square waves of any frequency, which is more easy to program and gives more accurate delays.

In Chapter 5, Example 5.2 gets 20 temperature values from 20 input ports. Usually there should be a time delay between each measurement, and a software delay can be used there. It is only necessary to write a procedure for a specific delay and call this delay between each measurement. Finer details of how exactly this is done, will be discussed in Chapter 9.

KEY POINTS OF THIS CHAPTER

- 8088 and 8086 processors are similar in most ways, except for three differences.
- There are two modes of operation for the 8086, which are designated as minimum mode and maximum mode.
- When in the minimum mode, the 8086 functions in a single processor environment.
- The address bus of the processor is multiplexed with the data bus and hence, to use these buses separately, de-multiplexing must be done.
- The address bus, data bus and control bus must be buffered to increase the driving capability of the buses.
- An external clock generator is used to provide the clock and to synchronize the reset and ready signals with the clock.
- A machine cycle is a bus activity during which a specific operation is performed.
- Reading and writing of memory and I/O are important bus cycles.
- A typical read/write machine cycle of 8086 is 4 T states long.
- Wait states are inserted into machine cycles to accommodate slow memory or peripherals.
- In the minimum mode, all control signals are obtained from the processor.
- In the maximum mode, an external bus controller IC generates the important control signals.
- In the maximum mode, there will be other processors in the system.
- There are pins on the 8086 on which other processors can request the service of the system bus.
- It is important to know how many cycles are expended for the execution of each instruction.
- The delay caused by instruction execution can be used to create delay loops.

QUESTIONS

1. List out the differences between 8086 and 8088.
2. How is the mode of operation of the processor selected and what is the purpose of each mode?
3. Why are buffers/line drivers used in the address, data and control lines?
4. What is the specification of the clock to be applied to the CLK pin of the processor?
5. What is the specification of the RESET signal to be given to the processor?
6. At what address does the 8086 wake up?
7. Why are RESET and READY signals applied to the clock generator IC?
8. When $DT/\bar{R}$ is high, what is indicated?
9. What is the difference between the 74LS244 and 74LS245 ICs?
10. When is the READY pin sampled?
11. What is the difference in the interrupt signals INTR and NMI?
12. What are the functions of the HOLD and HLDA pins?
13. When is the HOLD pin sampled?
14. How many wait cycles can be inserted in a bus cycle?

15. How can the processor be taken out of the HALT state?
16. What is ALE expected to do? What is the duration of ALE?
17. What is the purpose of the $\overline{LOCK}$ pin?
18. Which are the control signals which are not available from the 8086 in the maximum mode?
19. Why is it that the time for fetching an instruction is not included in the instruction cycle time of an instruction?
20. What is the use of the NOP instruction?

EXERCISE

1. Draw the de-multiplexed and buffered data bus, address bus and the read/write control signals with latches and buffers connected.
2. Draw a logic diagram using NAND gates to get the signals $\overline{IORD}$, $\overline{IOWR}$, $\overline{MEMRD}$ and $\overline{MEMWR}$ from the three processor signals $M/\overline{IO}$, $\overline{RD}$ and $\overline{WR}$.
3. Draw a memory write machine cycle showing the state of all important signals.
4. Redraw the above bus cycle with two wait states included.
5. Explain with relevant timing, the operation of the $\overline{RQ}/\overline{GT}$ pins.
6. In a typical write machine cycle, list out the activities occurring during each of the four T states.
7. Write a program for getting a delay of 100 msecs.
8. Write a program to generate a square wave at the LSB of the data bus of an output port with address F767H.
9. In a pressure monitoring environment, 10 pressure sensors have been connected. The various pressures from these sensors have to be read at intervals of 5 msecs. Write a program to read the sensor values from input ports having address 0FF0H to 0FF9H.

7 MEMORY AND I/O DECODING



In this chapter, you will learn

- The pin functions of a typical memory chip.
- The technique of memory address decoding.
- The reason and logic of using partial address decoding.
- The reason for organizing memory in banks.
- The principles of I/O address decoding.

Introduction

In the previous chapter, we discussed the hardware structure of 8086 and a few ICs which have to be used along with it. Now we will delve a bit deeper – the focus of this chapter will be on the way memory and I/O are connected to the processor.

Let us start with memory. Whenever we talked about reading or writing from memory, an address was referred to. Assuming that memory is byte organized, it is thought that every location in memory which can store a byte of data, has a unique address. The total address space of 8086 is from 00000 to FFFFFH. How is it that each location gets a unique address?

Another issue is that the total address space of 1MB is not obtained as a single chip of RAM – also it is not entirely RAM. ROM is also part of memory and a number of RAM and ROM chips together constitute the total memory i.e., the total address space.

Also, when the processor is used, its address space is partitioned such that certain address ranges are standardized for certain specific applications. Thus the IBM PC has a standard memory map – if it is not standardized, incompatibility issues are likely to arise. This chapter attempts to find answers to all these doubts and queries. The assumption here is that memory space and I/O space are disjoint (refer Chapter 5) i.e., the scheme of I/O interfacing used is 'peripheral or Isolated I/O'. This allows memory to use the full 1MB of space for itself.

7.1 | Memory Device Pins

The memory that we are talking about can be RAM or ROM – the difference is that ROM can only be read from, so the $\overline{MEMWR}$ signal from the processor does not have any relevance for ROM. A typical RAM has the pin lines as shown in Fig 7.1. As shown, it has data lines D_0 to D_{M-1} . If it is a byte organized memory, the data lines are D_0 to D_7 . The number of address lines (A_0 to A_{N-1}) depends on the number of locations it contains. For example, if it




Figure 7.1 | Typical RAM with control pins

is specified to be a 256×8 RAM, it means that it has 8 data lines and a storage capacity of 256 locations – thus it needs 8 address lines as $256 = 2^8$. Think of the following cases then – a $1K \times 8$ RAM needs 10 address lines, as $1K = 1024 = 2^{10}$, a $2K \times 8$ RAM has 11 address lines and so on. Thus, it is the capacity of the chip which decides the number of address lines. Coming to the other pins of the chips, the active low signals $\overline{WE}$ is to be connected to the $\overline{MEMWR}$ signal from the processor side. Only if this pin is activated, can the write operation be done in the addressed location. $\overline{WE}$ enables the input tri-state buffers of the data lines of the RAM. For reading, the pin $\overline{MEMRD}$ (from the processor side) is to be connected to the $\overline{OE}$ (output enable) pin of the memory chip for reading. When this pin is low, the output lines are activated; otherwise they remain tri-stated. Thus $\overline{OE}$ is meant to enable the output tri state buffers of the RAM.

Some RAM devices have only a single pin for reading or writing, which is $R/\overline{W}$. This is fine, as it is obvious that only one of the activities (read or write) can occur at a time. The pin $\overline{CS}$ (chip select) or $\overline{CS}$ (chip enable or S (Select)) is the pin which enables the memory chip. No activity is possible if this pin is inactivated, as the chip remains turned off. How does this pin get activated? It will be activated only if the address placed on the address bus of the processor is one of the addresses in the address range of this chip. When this condition is satisfied, a select pulse is obtained from the ‘address decoder’ output and the memory chip is turned ON for reading or writing. It is during a memory read or write cycle that the select pulse is obtained i.e., when reading or writing to the particular chip is required and the address is placed on the address bus.

A ROM chip is similar except that it does not have the $\overline{WE}$ pin, as it can only be read from. The ROMs that are used are usually EPROM, which have a pin to enable the chip to be programmed (PGM).

Why Active Low Control Signals?

You might have noticed that most of the control signals discussed here and in the previous chapters are active low. This is a TTL concept. The stray capacitances of the control pins can get charged from noise voltages, and this may cause the signals to cross the threshold of ‘high’ level as defined for TTL – thus if the control signals are active high, it may cause wrong triggering. On the other hand, an **active low signal** trigger will happen only when the line is pulled low deliberately by the controller.

Example 7.1

How many address and data lines are needed for the memory chips with the following organization?

- i) 256×4
- ii) 512×8
- iii) $1\text{K} \times 16$
- iv) $32\text{K} \times 8$
- v) $128\text{K} \times 8$

Solution

- i) 256×4 means, it has 4 data lines, and 256 memory locations, each of width 4 bits. $256 = 2^8$. Thus, 8 address lines are needed.
 - ii) 512×8
This chip has 8 data lines.
 $512 = 2^9$. Thus, it has 9 address lines.
 - iii) $1\text{K} \times 16$.
This chip has 16 data lines.
 $1\text{K} = 1024 = 2^{10}$. It has 10 address lines to access each of the 16-bit data words.
 - iv) $32\text{K} \times 8$
It has 8 data lines.
 $32\text{K} = 2^5 \times 2^{10} = 2^{15}$.
It has 15 address lines
 - v) $128\text{K} \times 8$
It has 8 data lines.
 $128\text{K} = 2^7 \times 2^{10} = 2^{17}$
It has 17 address lines.
-

7.2 | Memory Address Decoding

See Fig 7.2 in which a $2\text{K} \times 8$ memory chip is shown. This memory chip has 11 address lines, which are directly connected to the 11 lower lines of the address bus of the processor. Now the remaining 9 lines of the address bus of the processor are connected to the inputs of a NAND gate whose output pin feeds the $\overline{CS}$ pin of the memory chip. Thus, it is obvious that the chip is selected (enabled) only if all the input lines of the NAND gate are high – which means that $A_{11} - A_{19}$ of the address has to be high for the memory chip to be selected and made active. The NAND gate thus functions as the address decoder for the memory. It fixes up the address range of the chip.

Example 7.2

- i) Find out the range of addresses that the memory chip in Fig 7.2 contains.
- ii) Repeat the same when the logic of the NAND gate is changed as in Fig 7.3.




Figure 7.2 | Memory with address decoding




Figure 7.3 | NAND gate decoder

Solution

The address lines from A_{10} to A_0 can vary from 000 0000 0000 to 111 1111 1111. The address lines A_{19} to A_{11} must always be 1111 1111 1 for this chip to be selected. As such, the lower and upper range of the addresses in the chip are seen as:

A_{19}	A_{18}	A_{17}	A_{16}	A_{15}	A_{14}	A_{13}	A_{12}	A_{11}	A_{10}	A_9	A_8	A_7	A_6	A_5	A_4	A_3	A_2	A_1	A_0
1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1

The first address on the chip:

1111	1111	1000	0000	0000
------	------	------	------	------

The last address on the chip:

1111	1111	1111	1111	1111
------	------	------	------	------

Thus, the address range of the chip is FF800H to FFFFFH i.e., 2048 bytes i.e., 2 K bytes.

If the logic on the NAND gate is modified as in Fig 7.3, the address range will be:

0000	0000	0000	0000	0000
------	------	------	------	------

to

0000	0000	0111	1111	1111
------	------	------	------	------

For instance, 00000 to 007 FFH (0 to 2047 bytes) which is 2 K bytes.

7.2.1 | Address Decoding Concepts

- Thus the basic idea of address decoding is to decode the extra unused address lines of the processor to specify the address range.
- When more chips are to be interfaced, decode the extra address lines to a different range for each group.
- Any logic /logic gate can be used to perform address decoding. From the address decoder output, the right logic value should be obtained to turn the chip on. Thus the 'select' pulse obtained from the address decoding logic can be high or low, depending on the requirement of the chip select pin of the memory chip.
- In general, address decoders can be built using:
 - Random logic (simple gates)
 - Block decoders (e.g. 2×4 , 3×8 ...)
 - Programmable logic (PLAs, CPLDs, FPGAs, ...)

The first two types of address decoders will be discussed in detail here. The last type follows the same principle – only, the devices used are 'programmable' or 're-programmable'. This brings in more flexibility to the design. The design of the address decoder is then done using advanced design techniques involving hardware description languages and CAD tools. This is beyond the scope of this book and hence, is not covered here.

Example 7.3

Design an address decoder using OR logic for a $32\text{ K} \times 8$ RAM. Find the address space of this memory chip.

Solution

Fig. 7.4 shows the address 'decoding logic'. Since it is a 32 K RAM, it has 15 address lines. Thus the lower 15 lines of the address bus of the processor are connected to the address lines of the RAM. The upper 5 lines A_{15} – A_{19} are used for address decoding. Here the lines are given as input to the OR gate. When the address on the address bus corresponds to A_{15} to A_{19} being all at logic 0, the decoder gives the *SELECT* pulse to the $\overline{CS}$ pin of the memory chip and thus it is selected.

The lowest address is:

0000	0000	0000	0000	0000
------	------	------	------	------

And the highest address:

0000	0111	1111	1111	1111
------	------	------	------	------

The address range of this memory chip is 00000 to 07FFFH.




Figure 7.4 | Address decoding using OR logic

Example 7.4

Find the address space of the following chips with the address decoding circuitry as shown in the Fig 7.5a and 7.5b.

- i) 8 K × 8 EPROM
- ii) 8 K × 8 RAM

Solution

Both the memory chips are of 8 K capacity. Hence, they have 13 address lines. The remaining 7 lines of the processor system bus are used for address decoding.

- i) In Figure 7.5a, for the EPROM chip to be selected, the values of A₁₉ to A₁₃ are to be 1111 101. Thus, the lowest and highest addresses in the chip:

1111	1010	0000	0000	0000
to				
1111	1011	1111	1111	1111

For example, FA000H to FBFFFH.

- ii) In Figure 7.5b, for the RAM chip to be selected, the values of A₁₉ to A₁₃ are 0001 111. Thus, the lowest and highest addresses in the chip:

0001	1110	0000	0000	0000
to				
0001	1111	1111	1111	1111

For example, 1E000 to 1FFFFH.




Figure 7.5a | EPROM with address decoder **b** RAM with address decoder

7.2.2 | Address Decoding Using Block Decoders

A very popular decoder is the 3 to 8 decoder (74LS138) whose block diagram is shown in Fig 7.6a. The output lines are active low and depending on the selection inputs, one output line alone will be active. To enable the decoder chip, it must be ensured that $\overline{G_2A}$ and $\overline{G_2B}$ are at logic level 0 and G1 is at logic 1. The decoding table of the chip is shown in Fig 7.6b. Consider that a $1\text{ K} \times 8\text{ RAM}$ chip uses a 3 to 8 decoder for memory decoding. The RAM has 10 address lines. Thus, the remaining 10 lines of the processor address bus can be used for decoding. Let us calculate its address space. Consider a hardware connection as shown in Fig 7.7. In this, the address lines A_0 to A_9 from the address bus of the processor are directly connected to the address lines of RAM. The logic on these lines can vary from 00 0000 0000 to 11 1111 1111. The remaining 10 address lines of the processor are used for address decoding. Thus, we see them connected to the pins of the decoder. To enable the decoder, $\overline{G_2A}$ and $\overline{G_2B}$ must be 0, and G1 must be 1. Also, to enable Y3, CBA must be 011. Thus, the logic on the lines from A_{19} to A_{10} should be 1111 1111 10. Thus, the upper and lower range of the address is as shown in the following table.

[illegible]




Figure 7.6a | Functional block diagram of the 3 to 8 decoder 74LS138

(b)

Inputs					Outputs							
Enable		Select										
G_1	G_2^*	C	B	A	Y_0	Y_1	Y_2	Y_3	Y_4	Y_5	Y_6	Y_7
X	H	X	X	X	H	H	H	H	H	H	H	H
L	X	X	X	X	H	H	H	H	H	H	H	H
H	L	L	L	L	L	H	H	H	H	H	H	H
H	L	L	L	H	H	L	H	H	H	H	H	H
H	L	L	H	L	H	H	L	H	H	H	H	H
H	L	L	H	H	H	H	H	L	H	H	H	H
H	L	H	L	L	H	H	H	H	L	H	H	H
H	L	H	L	H	H	H	H	H	H	L	H	H
H	L	H	H	L	H	H	H	H	H	H	L	H
H	L	H	H	H	H	H	H	H	H	H	H	L

Note $G_2^* = G_2A$ and G_2B

Figure 7.6b | Decoding table of 74LS138




Figure 7.7 | Address decoding of a RAM using a block decoder

The address range is

1111	1111	1000	0000	0000
to				
1111	1111	1011	1111	1111

i.e., FF800H to FFBFFH

Subtracting FF800H from FFBFFH, we get 3FFH which is 1023 which means that there are 1023 + 1 addresses on this chip. Thus, there are 1024 byte (1 K) locations in the chip.

Example 7.5

The figure shows a RAM and a ROM connected to an address decoder, which is a 3 to 8 decoder. Find the address space of each of these chips.

Solution

Only 4 lines of the address bus are to be used for address decoding, as the remaining 16 lines are to be connected to each of the 64K memory chips.

For the ROM, A_{19} to A_{16} are to be 1001 ($A_{19} = 1, CBA = 001$). Thus, the address range of this chip is 90000H to 9FFFFH. Similarly, for the RAM, A_{19} to A_{16} are to be 1100 ($A_{19} = 1, CBA = 100$). Thus, the address range of this chip is C0000H to CFFFFH.




Figure 7.8 | Address decoding of a RAM and ROM using a 3 to 8 decoder




Figure 7.9 | Partial address decoding




Figure 7.10 | Partial address decoding using two address lines

7.3 | Memory Banks

8086 has a 16-bit memory bus – which means that data transfer can occur at a maximum rate of 16 bits (one word) per bus cycle. However, sometimes only a byte needs to be accessed. This means that the processor must have both options – i.e., both byte and word transfer must be possible. We know that for a word transfer, two byte locations must be accessed i.e., two addresses are actually needed.

Keeping these concepts in mind, let us see how memory is organized for 8086. A 16-bit data can be obtained by accessing two memory chips in parallel, each having 8 bits each of the word (see Fig 7.11). This is the way memory is organized in 8086. A 16-bit word is obtained as the concatenation of two bytes in two ‘memory banks’ i.e., memory is organized as two banks – one, the upper bank or high bank with the data corresponding to the upper byte D_8-D_{15} , and the other, the lower bank or low bank which has the data lines D_0-D_7 . The upper bank is also called the odd memory bank, because it has the odd addresses mapped to it – in the same way, the lower bank is called the even bank (see Fig 7.12). Remember that a 16-bit word has one byte with an odd address and one byte with an even address. **Thus for reading / writing a byte, only one of the banks is to be accessed, but to access a word, both memory banks have to be accessed.** In the second case, the byte in the given logical address and the byte in the next address are accessed. For instance, if the instruction is MOV AX, [0002] the logical addresses involved are 0002 and 0003 – one byte in




Figure 7.11 | 16-bit memory realized using two 8-bit memory chips




Figure 7.12 | Memory banks of 8086

the even bank, and one byte in the odd bank. For an instruction `MOV AX, [0001]` the addresses are 0001 and 0002 – again, both banks have to be accessed to get the complete word. How are the two banks differentiated? An even address implies that the LSB of its address i.e., A_0 is low. Thus, A_0 can be used to enable the even bank. To enable the high (odd) bank, 8086 generates a signal $\overline{BHE}$ (**Bus High Enable**) which goes low whenever the addressed byte is in the high (odd) bank. In Fig 7.13, see how these two signals are used along with the address decoder output, to enable the two banks independently. $\overline{BHE}$ is enabled if and only if the address on the address bus is for accessing a byte in the high bank. That is why it is designated ‘BUS HIGH’ enable. A_0 is frequently called ($\overline{BLE}$) (Bus Low Enable) because it is used to enable the low bank.

Fig. 7.13 shows a case when two $32\text{K} \times 8$ RAMs form the 16-bit memory. Thus the total memory is $64\text{K} \times 8$ bytes. 32K RAM needs 15 address lines. Of the 20 address bits from the address bus of the processor, A_0 is used to enable the even bank. A_1 to A_{15} are connected to the address pins A_0 to A_{14} of each RAM chip. The rest of the address lines are used for address decoding. A_0 and $\overline{BHE}$ are used for selecting each of the banks separately.

Now see Table 7.1 and note how and when $\overline{BHE}$ goes low. What emerges from the table is that, for byte access, $\overline{BHE}$ goes high if the referred byte is in the high (odd) bank. **When a word whose logical address is even: both A_0 and $\overline{BHE}$ are low simultaneously.** Hence, for the reading or writing of such a word, only one machine cycle is needed.




Figure 7.13 | Memory banks with the respective bank enables signals, along with the decoder

Table 7.1 | Status of Control Signals for Memory Bank Access

Address	Typical instruction	Data type	$\overline{BHE}$	A_0	Bank accessed	No. of machine cycles
(even)	MOV AL, [0000]	byte	1	0	Low	1
(odd)	MOV AL, [0001]	byte	0	1	High	1
(even)	MOV AX, [0000]	word	0	0	Both	1
(odd)	MOV AX, [0001]	odd byte	0	1	High	I cycle
		even byte	1	0	Low	II cycle

7.3.1 | Odd Addresses Word

For accessing a word whose address is odd, two machine cycles are expended. See the example instruction MOV AX, [0001] in Table 7.1. This means that the effective addresses to be accessed are 0001 and 0002. When the address 0001 (physical address corresponding to this) is placed on the address bus, A_0 is not low. So, only the upper bank address is accessed in this machine cycle. Only one byte is read in this machine cycle. To get the next byte from the even bank, one more machine cycle must be used.

This is why it was mentioned (Section 6.6) that for accessing a word operand from an odd location, an extra machine cycle machine time has to be added to the instruction cycle time. Thus, it is advantageous (in terms of speed) to align all words at even addresses. Many optimized assemblers do this automatically. Otherwise, there is a directive named EVEN which does this.

Take the case when the data segment has 7 bytes and a word. The bytes are stored in location 0000 to 0006. Obviously, the word will be aligned at the next address, which is odd. If the EVEN directive is used, the location counter is incremented by 1, and the word is stored at location 0008. This is illustrated in the following data segment definition.

```

        .DATA
NUMS DB 7 DUP(05)
        EVEN                ;increments location counter to 0008
WDR  DW 2345H
```

However, this obviously wastes one byte space, and memory is fragmented. This situation cannot be avoided if speed is to be optimized.

7.3.2 | Why Memory Banks?

Thus, we see that arranging memory to be in two memory banks causes an additional burden. Then why is the memory system designed this way?

What is being attempted is that, when a byte alone is to be accessed, only the bank in which that byte resides, should be enabled – the other bank should remain disabled.

What is the problem if both the banks are enabled for any and every access?

The answer is that there will not be any problem in case of a read operation, but it may cause havoc for a write operation. See the following cases:

```
MOV AL, [0000]
```

Note the above instruction which initiates a read cycle. It is the case of reading a byte. Only the even (low) bank need be enabled to read. However, even if the high bank is also enabled, it does not matter. The data present on the lines D_8 to D_{15} will not be used by the processor. It takes into AL, only the data on lines D_0 to D_7 and ignores the data on the upper bus. However, the writing process creates a problem. Observe the case of the following write instruction:

```
MOV [0000], AL
```

This should cause only the even bank to be enabled. If the odd (upper) bank also gets enabled, the logic levels on the data lines D_8 – D_{15} will get written onto the enabled location in the upper bank, which is not admissible. **Thus, when a ‘byte’ alone is to be written into memory, it is important to ensure that only one of the memory banks is enabled.**

7.3.3 | Using Separate Write Strobes

Another approach is used to enable memory banks, and it is by generating separate write strobes. Now that the problem has been identified to be associated with ‘writing’, there is another way to solve it, and it is by generating separate write strobes.

We can use the logic shown in Fig 7.14. When there is the necessity to write a byte to the high order (upper) bank, only the $\overline{HWR}$ (High write) signal will be enabled – similarly for writing a byte to the lower bank, only that bank will be enabled, but for writing a word, both banks get enabled. Only the write control signal is taken into consideration – reading unwanted data is not a problem as the processor takes from the data bus, only the byte that it needs. In this scheme, the $\overline{LWR}$ and $\overline{HWR}$ signals are applied to the $\overline{WR}$ pins of the respective memory banks only. Separate bank select signals from the address decoding circuitry are not used.

7.3.4 | Memory Map of IBM-PC

Table 7.2 shows the memory map of **IBM PC**, the first version of the PC manufactured by IBM with the 8088 processor, which has a total memory address space of 1 MB. For a 4 GB (FFFFFFFFH) memory of the current x86 PC, the same memory mapping is followed, with RAM at the lower end and ROM area at the upper end.

7.4 | I/O Address Decoding

It was mentioned in Chapter 5 that there are two schemes for I/O port interfacing. The simpler scheme is ‘memory mapped I/O’, but this causes I/O to encroach on the memory address space. Hence, most systems prefer the other scheme which is designated as ‘Isolated I/O’, ‘Peripheral I/O’ or ‘I/O mapped I/O’. We will confine our discussion to this scheme alone. Even though our processor has 20 bits for address, I/O uses only 16 bits. Thus I/O port addresses can range from 0000 to FFFFH and the total I/O space is only 64 K i.e., we can use 65,536 different I/O addresses. However, since the instructions for input (IN) and output (OUT) are separate, we




Figure 7.14 | Using separate write strobes

Table 7.2 | IBM PC Memory Map

Address range	Size	Type	Allocation
00000 to 9FFFFH	640 K	RAM	User RAM – with some areas with special allocations like 1 K (00000 to 003FFH) for interrupt vector table, some portions for storing BIOS, some for DOS parameters and some for the operating system
A0000H to BFFFFH	128 K	VDR	Video display RAM with portions allocated for video buffers and other video applications
C0000H to EFFFFH	192 K	ROM	Memory expansion area for ROM
F0000H to FFFFFH	64 K	ROM	BIOS ROM, hard disk and other peripherals' ROM




Figure 7.15 | Input and output ports connected to the 8086

can have 64K input ports and 64K output ports. This is quite a lot, practically. Now, recollect (Section 5.2.1) that ports with 8-bit port addresses use ‘fixed port addressing’, while port addresses beyond that have to use the ‘variable port addressing scheme’. See Fig. 7.16, which shows that I/O addresses above 00FFH have to use ‘variable port addressing’. The lower range of I/O addresses from 00 to FFH can use ‘fixed port addressing’. Another point to remember is that any port can have an 8-bit or 16-bit data bus.

Fig 7.15 shows input and output ports connected to the data bus of the process. Note that data is sent ‘out’ to output ports, and data is taken ‘in’ through input ports.

7.4.1 | Output Ports

Output ports use the instruction ‘OUT’ for writing into it. Then the data in the accumulator (AL or AX) gets written into the output device. An output device may be as complex as a video display or as simple as LEDs, which have to light up. Data from the processor is sent to an output port during the I/O write cycle. This data will be available on the data bus only until the end of a machine cycle. For the output port to be able to use this data, it must be latched. Thus, a latch is mandatory for an output port set up. (For memory, there are latches inside the chip.) Fig 7.17 shows a basic output port.




Figure 7.16 | Schemes of port addressing




Figure 7.17 | Basic output port setup

Requirements of an Output Port

- i) To identify/select the specific device, an address decoder should give a select pulse.
- ii) To write into the output device, the $\overline{IOWR}$ signal must be active.
- iii) When both the above two conditions are true, the data on the data bus is latched.
- iv) All the above activities occur during an I/O write cycle.

Example 7.7

Design a system with 8 LEDs connected to the lower 8 bits of the data bus of the 8086 data bus. The LEDs must switch ON and OFF with a delay of 1 second between each switching action.

Solution

Fig 7.18 shows the hardware setup.

- i) The diagram with the address decoder, the select pulse and the $\overline{IOWR}$ signal is shown. The address decoder gives a high select pulse. This is ANDed with the inverted $\overline{IOWR}$ pulse, and applied to the G pin of the octal transparent latch 74LS373 (refer Section 6.1.4).




Figure 7.18 | Output port with 8 LEDs interfaced to it

The G pin needs a high level pulse to latch the data on its D inputs to its Q outputs. $\overline{OC}$ must be strapped to ground to enable the output pins of the latch.

- ii) Current limiting resistors are connected to the cathodes of the LEDs
- iii) The address decoder specifies an 8-bit address of F8H for the port.
- iv) The program for switching the LEDs ON and OFF:

```

STRT:    MOV AL, 0FFH           ;data to light up all the LEDs
          OUT 0F8H, AL           ;transfer it to the output port
          CALL DELAY_1SEC        ;call the delay procedure
          MOV AL, 00             ;data to switch off the LEDs
          OUT 0F8H, AL           ;transfer it to the output port
          CALL DELAY_1SEC        ;transfer it to the output port
          JMP START              ;repeat continuously
  
```

- v) For the delay procedure refer Example 6.3.

7.4.2 | Input Ports

Input ports use the IN instruction to receive data into the processor in the accumulator of the processor (AL or AX). An input port requires a tri-state buffer along with it to ensure that it




Figure 7.19 | Basic input port

is isolated from the bus when it is not selected. A simple switch can be an input port, and the alphanumeric keyboard used in PCs is also an input port. Fig 7.19 shows a basic input port.

Requirements of an Input Port

- i) To identify/select the specific device, an address decoder should give a select pulse.
- ii) To read from the input device, the $\overline{IORD}$ signal must be active
- iii) When the above two conditions are true, the corresponding tri-state buffer gets enabled and data is transferred to the processor.

Taking into consideration all these requirements, let us design a system with two sets of switches acting as two input ports. The tri-state buffer 74LS244 is used (refer Fig 7.20). This has two active low enable inputs $\overline{1OE}$ and $\overline{2OE}$. When both of them are low, data at the A inputs appear at the Y outputs. However, if the enable pins are high, the output is tri-stated.




Figure 7.20 | Functional pin diagram of the octal tri-state buffer 74LS244

Example 7.8

Find the addresses of the two input devices shown in Figure 7.21 and write a program to read data from each of these ports and move it to some other registers.

Solution

- i) Two sets of 8 switches S_0 to S_7 are connected to the data bus through two tri-state buffers and this setup functions as two input ports.
- ii) The address of the ports are 8-bit, as only address lines A_0 to A_7 have been considered for decoding. The address of the first port is 7EH and that of the second is 3EH.
- iii) The address on the address bus decides which of the input ports get selected. The tri-state buffer ensures that the port which is not selected, is isolated from the bus.
- iv) The program for inputting data from each of these ports is

```

STRT:      IN AL, 3EH           ;take in data from Port2
           MOV BL, AL          ;move data to BL
           IN AL, 7EH          ;take in data from Port1
           MOV DL, AL          ;move data to DL

```

Thus, after reading both ports, the switch settings are available in BL and CL. This data can be used for processing and/or display.

7.4.3 | Decoding 16-Bit I/O Addresses

We have used 8-bit addresses so far for I/O devices. However, personal computers use 16-bit addresses as well. In that case, only the decoding circuitry changes. Address lines A_0 to A_{15} should be used by the address decoder to generate the device select pulse. The other point to remember is to use 'variable port addressing' when writing I/O instructions for ports with 16-bit addresses.

7.4.4 | Ports with 16-Bit Data Bus

So far we have discussed only I/O ports with 8-bit data bus. However, 16-bit data bus is also possible for I/O ports. What becomes different then? The answer is that, the situation is similar to the case of memory banks. Recollect the concept of memory banks. If each of the 16-bit ports needed to be accessed as two 8-bit ports as well, we have to have 8-bit I/O banks. So we can have I/O banks too, with 8 bits in the upper bank and 8 bits in the lower bank. To read a word, both banks can be enabled, by using both A_0 and $\overline{BHE}$ bits. Fig 7.22 shows the upper and lower I/O banks.

Note Here the addresses can have only a maximum size of 16 bits, while for memory, the addresses are 20 bits.

Most of the I/O ports used for PCs have 8 bits wide data, but there are 16-bit wide ports such as video and analog to digital conversion ports. If a port is by definition, a 16-bit wide port, uses of separate bank select signals are not needed – the port can be selected directly by the address decoding and the I/O read/write signals. That is because data is always read or written as a 16-bit word, and byte access is not necessary.




Figure 7.21 | Two sets of switches acting as two 8-bit input ports




Figure 7.22 | I/O banks

KEY POINTS OF THIS CHAPTER

- Memory mapped I/O causes I/O space to encroach on the memory address space of the processor.
- Active low control signals prevent accidental triggering by noise voltages or open connections.
- 74LS138 is a very popular 3 to 8 decoder.
- The number of locations in a memory chip is 2^N where N is the number of address lines of the chip.
- Memory is always organized as 'banks'.
- The $\overline{BHE}$ signal goes low when the addressed byte is in the upper bank of memory.
- The LSB of the address bus i.e., A_0 is used to enable the low bank of memory.
- Memory chips have tri-state buffers and latches inside the chip.
- All input ports need a tri-state buffer to connect it to the system bus.
- All output ports need a latch to keep the outputted data, until it is changed by the next OUT instruction.
- An I/O address can be 16 bits in size, in which case 'variable port addressing' is necessary for writing I/O instructions.

QUESTIONS

1. How many address lines does a 256K memory chip need?
2. Name the important pins in a typical RAM and explain the function of each.
3. Find the numbers of two standard ROM chips and two EPROM chips, along with the way the chips are organized (how many address and data lines).
4. Why are active low signals preferred for use as control signals?
5. What is meant by the term 'address decoding' with respect to memory as well as I/O devices?
6. Why is memory organized in 'banks'?
7. Why is the address line A_0 also designated as $\overline{BLE}$?
8. Why are separate write strobes used for memory banks, but not read strobes?
9. Whenever LEDs are connected to the data output lines, it is preceded by latches. Why?
10. What is the range of address space for I/O in an 8086 based system?

EXERCISE

1. Design an address decoder for two RAM chips and two ROM chips each organized as $1K \times 8$ chips. Use gates for the decoding circuitry and specify the address range of each chip.
2. Repeat the above problem using the chip 74LS159 (2 to 4 decoder) as address decoders.
3. Using a 3×8 decoder, design the decoding setup for eight RAM chips each of size $8K \times 8$. Indicate the address space of each chip.
4. For the following (Fig 7.23a and 7.23b) decoding circuitry, find the range of the memory addresses.




Figure 7.23a




Figure 7.23b

5. Design a memory system using partial address decoding for 4 RAM chips of size 1K x 8 each.
6. Draw a diagram using memory banks for a 16K x 16 memory. Draw the decoding circuitry, and use write strobes for enabling the $\overline{WE}$ pins of the banks.
7. Design an I/O system in which there are two switches connected as an input port. The status of these switches is to be read and fed to corresponding LEDs connected at the output side. Draw the hardware diagram for this and write a program to do this. The address of the ports should be designed.
8. Design an I/O system in which an input counter counts from 0 to 9, then goes back to 0 and then repeats this sequence. This count has to be displayed on a seven segment LED. Design the setup with input and output ports.

8 THE INTERRUPT STRUCTURE OF 8086



In this chapter, you will learn

- The concept of an interrupt.
- The interrupt response of the 8086.
- The way the interrupt vector table is organized for 8086.
- The idea of software interrupts.
- The operation of hardware interrupts of 8086.
- The significance of DOS and BIOS interrupts.
- How to access the video display by the use of BIOS interrupts.
- How to access the video display by direct writing/reading to the video RAM.
- How the PC's keyboard and its controller are organized.
- How to use BIOS interrupts for keyboard access.
- The method of 'hooking' interrupts.
- How to write TSR programs.

Introduction

The word 'interrupt' brings to our mind many situations. It could be that we are at home doing something rather important and interesting and suddenly the doorbell rings. There is no option but to go to the door and attend to the person who has interrupted us – it may be that the visitor needs only a few minutes of our time – but it may also happen that the visitor needs us to go out and do some errand for him. In either case, the task we were at, is suspended temporarily. We spend some time attending to the task requested by the visitor, after which we get back to the suspended task. This is a real life situation that happens quite frequently, and being social beings, none of us ever get to do any of our tasks uninterruptedly

This is exactly the case with the processor too. Since the processor is connected to various peripherals and also because it has various programs stored in its memory, there is very little chance of it being able to perform a single task without interruption. Besides, performing a task only when 'requested' is the most efficient way of using the capabilities of the processor. Let us look at it this way. The keyboard is a peripheral connected to the processor – this means that the processor expects a key to be pressed at any time. One option is for the processor to do 'polling' i.e., the processor continually keeps checking if a key has been pressed. When a key is pressed, the processor goes on to the task of identifying the key. However, if the processor is 'polling' for a key press, it is obviously waiting and doing nothing else. This is just a waste of processor time. A better way of organizing the setup would be to let the processor do some task, and 'interrupt' it only when a key is pressed. When a 'key press' interrupt is

received, the processor performs the 'interrupt service routine' (ISR) which is the process of identifying the key that was pressed. This, in essence, is the philosophy of interrupts and it helps utilize the processor with maximum efficiency.

8.1 | Interrupts of 8086

Now that we have accepted that interrupts are a part of the normal working of a processor, let us see the interrupt processing mechanism of 8086. 8086 has hardware interrupts, software interrupts and error generated interrupts. For the three cases mentioned here, the interrupting mechanism is different, but the way the processor responds is similar. So let us first have a look at the way 8086 responds to any interrupt. The point to note is that, after an interrupt request is processed, the processor has to come back to its previous task which was left unfinished. Thus, an interrupt is not very different from a CALL instruction in its philosophy, but there are some differences in the way it is handled and processed.

8.1.1 | Interrupt Response of 8086

After every instruction cycle, the processor checks if any interrupt is awaiting service. If it finds an interrupt request and decides to acknowledge and service it, the response is the following sequence of steps:

- i) The flag register is pushed on to the stack.
- ii) The interrupt flag is disabled ($IF = 0$).
- iii) The trap flag is disabled ($TF = 0$).
- iv) The CS register is pushed on to the stack.
- v) The IP register is pushed on to the stack.
- vi) Control is transferred to the location in which the corresponding 'Interrupt Service Routine' (ISR) is stored. This in effect, would be a far jump.
- vii) The program corresponding to the ISR is executed. The last instruction in the ISR will be IRET.
- viii) Then IP is popped off the stack.
- ix) CS is popped off the stack.
- x) The Flag register is popped off the stack.
- xi) Control returns to the point at which it had left off.

Steps i to vi constitute the actions before getting to the **interrupt service routine**, and steps viii to xi occur on returning from the interrupt.

Let us go into the details of the steps in the above list. See Fig 8.1

- A. The first step is part of the action of saving the 'context' of the current program. First the flag status is saved in stack for later retrieval. This is because the ISR may change the status of flags.
- B. The second step is for ensuring that this interrupt routine does not get interrupted by other interrupts received on the INTR line. (We will see later that for interrupts on the INTR line to get acknowledged, it is mandatory for the interrupt flag to be set).
- C. The third step is for ensuring that the program taken up does not stop after the execution of each instruction (this happens if the Trap flag is set).




Figure 8.1 | Steps in processing an interrupt request

- D. The fourth and fifth steps are for saving on stack, the address of the next instruction in the current sequence, just as is done in the case of a far CALL. They will be taken back from the stack after the interrupt processing is over. Both IP as well as CS are saved, because the new program taken up will in a different code segment.
- E. The sixth step is that of control being transferred to a location in a different code segment. That is why it is designated as a 'far jump'.
- F. The interrupt service routine, which is the program at which, the jump has occurred, will now be taken up for execution. The last instruction of an ISR is an IRET instruction which will cause the reversal of the previous actions.
- G. In the eighth and ninth steps, the IP first and then CS are popped from stack. Thus the address of the next instruction in the normal sequence is back in the required registers.
- H. In the tenth step, the flag register is popped off, and thus the original status of all flags is restored.
- I. Finally, instruction execution resumes from where it had been diverted.

8.1.2 | Interrupt Service Routine and Interrupt Vector

Now, what is an interrupt service routine? When an interrupt occurs, the processor suspends the execution of its current task and takes on another task as required by the interrupting source. This program, or routine as it may be called, is designated as an 'interrupt service routine'. This routine corresponds to the request of a particular source of interrupt and it is also called its 'interrupt handler'. This means that for any interrupt that occurs, there is a particular interrupt service routine (ISR).

Now where is this ISR available? It is to be available in memory and must be accessed on the occurrence of the specific interrupt. For that, the address of the ISR must be obtained. The address of an ISR is called its '**interrupt vector**'. For an 8086 based system, any address of code is in the following form, CS:IP. Thus, the interrupt vector for any interrupt has 4 bytes – two for the CS value and two for the IP value. Thus, if the interrupt vector for a particular interrupt is obtained, control can be transferred to the new location by using the new values of CS and IP specified as the 'interrupt vector'.

8.1.3 | Interrupt Vector Table

It is obvious that the number of interrupt vectors in a system is the same as the number of interrupts that the system can process. The 8086 has 256 interrupt vectors and since each vector is specified by 4 bytes, it implies that $256 \times 4 = 1024$ bytes (1K) of memory are allocated to store the interrupt vectors. These 256 vectors are stored in a table called the ‘Interrupt Vector Table’ (IVT) in system RAM from locations 00000 to 003FFH i.e., upto 0000: 03FF. See Fig 8.2a. Each interrupt vector is numbered from 0 to 255, and as such, each interrupt is also numbered in the same way. These numbers thus turn out to be what can be called as the ‘type numbers’ of the interrupts. Thus, there are interrupts designated as Int 0, Int 1, Int 2.....Int 255. As is obvious from Fig 8.2a, these numbers show the position of the corresponding interrupt vector in the interrupt vector table. Thus, the vector of INT 0 is the 1st entry in the IVT that of INT 1 is the second entry ... that of INT 6 is the 7th entry and so on. Since each interrupt vector is 4 bytes long, the memory location corresponding to the vector of INT n is obtained as $n \times 4$ i.e., for INT 1, the address of the first byte of its vector is $1 \times 4 = 0004$, for INT 0 it is 0000 and so on, as seen in Fig 8.2a. In fact, this is how the 8086 reacts to an interrupt with a type number. **It multiplies the type number by 4, gets the address of its interrupt vector and loads IP and CS with these new values and starts executing the ISR it has located.**




Figure 8.2a | Interrupt vector table of 8086




Figure 8.2b | A typical interrupt vector (4 bytes)

Example 8.1

Find the address (in the IVT) of the interrupt vector of INT 61H. Find the physical address of the ISR corresponding to this interrupt if the vector is 0F00:9872.

Solution

The type number of the interrupt is 61H = 97 in decimal.
The address of the interrupt vector is $97 \times 4 = 388 = 184H$.
Thus, the interrupt vector is to be stored in the IVT in location 0000:0184 onwards.
For the ISR, the CS value is 0F00H and the IP value is 9872H. The part of the interrupt vector table which has these vectors stored is shown.

		0FH
0186H		00
		98H
0184H		72H
Address		

Now, How Does the Interrupt Vector Table Get These Entries?

It was mentioned in Chapter 6 that after ‘power on’, the processor wakes up at the address FFFF0H. This is the entry point to BIOS in ROM (Ref Table 7.2). BIOS contains a set of routines that provide device support. These routines check, verify and initialize various devices, and also establish two data areas, one of which is the interrupt vector table from locations 00000 to 003FFH.

How is an Interrupt Different from a ‘Call’?

An interrupt seems to perform actions similar to a far CALL instruction. It differs in a few points, however. An interrupt causes the flag register to be pushed on the stack in addition to CS and IP values. It clears the trap and interrupt flags. The action of clearing the flags is not associated with a CALL, but is part of the automatic response due to an interrupt. Also, note that the ISR ends with IRET rather than a RET instruction. The IRET instruction ensures that the flag register is also popped off the stack, rather than just the CS and IP.

8.2 | Dedicated Interrupt Types

Intel has dedicated certain interrupt types for specific applications directly related to CPU operations. These are listed below.

8.2.1 | INT 0 (Divide by Zero Error)

The interrupt with type number 0 is dedicated to the ‘divide by zero’ error. This interrupt is an ‘error generated’ interrupt (also called an ‘exception’). On division, if the quotient register is not large enough to contain the quotient, this interrupt is generated automatically. Dedicating Type 0 for this case means that the corresponding interrupt vector in the interrupt vector table is available at 0000:0000 (see Fig 8.2a). Thus, the ISR for this error generated interrupt is written

in the address specified in the IVT. What could be a possible ISR for this condition? One possibility is just to display a message indicating ‘divide overflow’ and expect the programmer to correct his data/program. Another possibility is to write as the ISR, a program which increases the size of the quotient register, so that the problem is corrected without user intervention. It is up to the system designer to decide how the error handler is to be written.

8.2.2 | INT 1 (Single Stepping)

This type number is dedicated for ‘single stepping’ or ‘trace’. Single stepping is an important idea in debugging. During logical debugging of our programs, we would like to stop after the execution of each instruction and check the contents of registers, memory and so on. We usually perform the action of ‘trace’ this way. Intel has provided the ‘Trap’ flag for this, and this flag has to be set to let this happen. The ISR for viewing the register and memory contents will be pointed by the vector of interrupt type 1. However, an important issue in this is how to set the trap flag. No such instruction has been encountered so far. Recollect the flag register configuration.

D ₁₅	D ₁₄	D ₁₃	D ₁₂	D ₁₁	D ₁₀	D ₉	D ₈	D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
U	U	U	U	OF	DF	IF	TF	SF	ZF	U	AF	U	PF	U	CF

The Trap flag is the D₈ bit of the flag register.

```
PUSHF                                ;push the flag register to the stack
POP AX                              ;pop it to AX
OR AX, 0000000100000000B             ;OR it to set bit D8 i.e., TF
PUSH AX                              ;push AX to stack
POPF                                ;pop it to the flag register
```

This program segment sets the trap flag which causes INT 1 to occur.

Note The Trap flag will be reset as part of the interrupt response, so the ISR can execute without stopping after each program line. However, on going back to the mainline program, when the flags are popped back, the set condition of the Trap flag will be retrieved, and therefore single stepping can be continued for the mainline program until the Trap flag is reset again.

8.2.3 | INT 2 (Non Maskable Interrupt)

This interrupt corresponds to the vector (pointer) of the hardware interrupt NMI. When an interrupt is received on the pin NMI (Non Maskable Interrupt) of the processor, a type 2 interrupt occurs – this means that the ISR for NMI must be written in the address pointed by the corresponding IVT content. We will deal with the NMI interrupt in detail in Section 8.4.1.

8.2.4 | INT 3 (Breakpoint Interrupt)

This is the breakpoint interrupt, which is useful for de-bugging. We will need to set breakpoints (stop after executing a group of instructions) and check the content of registers and memory after executing instructions up to the breakpoint. Since setting breakpoints is an important aspect of debugging, the breakpoint interrupt is special in that it is a single byte instruction with the code CCH. When we set breakpoints in debugging, what is really happening is that this code (CCH) is inserted and getting executed at the position in the program where the breakpoint is set. Just like in single stepping, the ISR will be written to dump the register and memory contents on the video display.

8.2.5 | INT 4 (Overflow Interrupt)

This interrupt corresponds to the overflow flag. If the overflow flag is set, this interrupt occurs, but not automatically. An instruction INTO (interrupt on overflow) must be written after the program segment which is likely to cause the overflow flag (OF) to be set.

```
MOV AL, NUM1
ADD AL, NUM2
INTO                ;interrupt on overflow
```

The last line of this program segment can pass control to the ISR written for INT 3, if the overflow flag is set by the result of the addition. Otherwise INTO acts as a NOP instruction.

8.2.6 | Allocation of Interrupt Type Numbers

Thus, we see that interrupts with type numbers 0 to 4 have been allotted pre-defined actions related to the CPU. See Table 8.1. Interrupts with type numbers 5 to 31 are reserved by Intel, for special uses and for future processors. As a matter of fact, a few of them are used for various hardware and I/O devices and some for error conditions of the higher processors in the x86 family. The rest of the interrupts can be used by the user. However, in the PC, many interrupt types have been used by BIOS and DOS. We will see the list of those interrupts and also use some of them, later.

8.3 | Software Interrupts

What is meant by the term ‘software interrupt’? When an interrupt is initiated by an instruction, it is called a software interrupt. The format of this instruction is:

```
INT type number.
```

The type numbers can vary from 0 to 255. This is a very important and interesting way of using procedures whose vectors are inserted into the IVT. Thus, the user can write any procedure as an ISR, store it in some address in memory and call it by using the instruction INT n. Before using it this way, the user must however insert the CS and IP values of his ISR in the IVT in the location corresponding to INT n. Before using it this way, the user must however insert the CS and IP values of his ISR in the IVT in the location corresponding to INT n.

Hardware interrupt service routines can be tested using software interrupts. Take the case of NMI which is a hardware interrupt. This is vectored to location 0000:0008 i.e., it corresponds

Table 8.1 | Interrupt Vector Table Allocation

INT type no.	Location in IVT (HEX)	Application
0	0000:0000	Divide by zero error
1	0000:0004	Single step interrupt
2	0000:0008	Non maskable interrupt
3	0000:000C	Breakpoint
4	0000:0010	Overflow
5 to 31		Reserved by Intel
32 to 255		Available to user

to type number 2. The routine corresponding to this hardware interrupt can be tested by using the instruction INT 2. In effect, a hardware interrupt on the NMI pin and the software instruction INT 2 takes us to the same absolute address. Similarly, the effect of a 'divide by zero' error can be simulated by writing the instruction INT 0 which takes control to the ISR which has been written for the 'divide by zero' error.

What Is the Size of an Interrupt Instruction?

The format of an interrupt instruction is INT n where n can only be a byte. The opcode for INT is 'CD', again a byte – for example INT 5 has the opcode CD 05. Thus, an interrupt instruction is 2 bytes long, except the breakpoint interrupt INT 3 which has the code of just CC. This feature is given to this interrupt because setting of breakpoints is very frequently done in debugging, and having one byte only, means a lot of memory space saving.

8.3.1 | DOS and BIOS Interrupt Routines

We have used software interrupts in the form of the DOS function INT 21H. DOS has a number of functions (procedures) to access input/output devices, and we can use these functions by knowing the interrupt type number. The absolute address of the function need not be known to us.

Besides DOS interrupts, there is another set of functions for I/O access and this is supplied by the system BIOS. BIOS is an acronym for 'Basic Input output System'. A collection of routines i.e., procedures are available here too, to access I/O. These functions are accessed by interrupts with type numbers specified for a particular peripheral (Refer Table 8.2). Thus, DOS and BIOS interrupts belong to the class of software interrupts and will be discussed in greater detail soon.

8.4 | Hardware Interrupts

Now we will examine the hardware interrupt pins of 8086. There are two pins on which interrupt requests can be received – they are the INTR pin and the NMI pin.

8.4.1 | NMI

This is a positive edge triggered interrupt, but it is also required to have a duration in the high state of more than two clock cycles. Any high going transition of NMI is latched on-chip and will be serviced at the end of the current instruction. This is a non maskable interrupt, in that it does not depend on the setting of the interrupt flag (IF). It cannot be masked or prevented from being activated. It is a Type 2 interrupt, meaning that its vector is obtained from the corresponding location in the interrupt vector table. NMI caters to applications of the highest priority, like power failure.

8.4.2 | INTR

This is the non-vectored interrupt pin of the 8086, which means that when an interrupt request is received on this pin, it does not get automatically directed towards any particular entry in the interrupt vector table. Another feature of this interrupt is that the interrupt flag (IF) is required to be set for an interrupt request on the INTR line to be honored. When the processor is reset, all flags are found to be cleared and so is the IF. It is to be ensured that the IF is set by the instruction STI if interrupts on the INTR lines are to be acknowledged.

INTR is a high level triggered interrupt. To be responded to, the line should be high during the clock period preceding the end of the current instruction. During any interrupt response,




Figure 8.3 | Interrupt acknowledge machine cycle of the 8086

further interrupts on the INTR line are blocked, because clearing the IF is part of the processor's interrupt response. Now let us see the sequence of actions following an interrupt request on this line. On being interrupted, the processor goes into an 'interrupt acknowledge' state which causes two 'interrupt acknowledge machine' cycles to occur. During the first machine cycle, the address/data buses are floated and the $\overline{INTA}$ (interrupt acknowledge) pin goes low in T₂ and remains low till T₄. This is an indication for the interrupting device that its interrupt request has been accepted by the processor. In the second machine cycle, once again the line $\overline{INTA}$ is low from T₂ to T₄, and then the interrupting device places the 'type number' of the interrupt on the lower data lines (D₀ to D₇), which comes out of the float condition. The processor accepts this number, multiplies it by 4, and accesses the interrupt vector table at that location.

In Fig 8.3, the ALE signal is also shown. Actually, this signal has no role in the INTA cycle, but is included just to show that in the first T state of any machine cycle, the ALE signal is activated.

8.4.2.1 | Generation of a Type Number

Fig 8.4 shows a simple circuit which is capable of sending a 'type number' to the processor. Consider an interrupting device interrupting the 8086 on its INTR pin. The processor responds by lowering the INTA line. This signal is connected to the enable pins of a tri-state buffer (74LS244). Eight switches are connected to the input pins of the TSB. When a switch is closed, the line is grounded corresponding to a '0' level on the pin. When a switch is open, it gets V_{cc} on it and hence is in the '1' state. In the figure, the switch configuration corresponds to 1000 0001 or 81H. This is because the first and last switches are open, and the rest are closed. When the enable pins (OE) of the TSB are low (i.e., when INTA becomes low), the data at the input of the TSB is transferred to the output of the TSB chip.

During the first Interrupt acknowledge machine cycle, the address/data lines of the processor are floated. This 'interrupt acknowledge' machine cycle signals the interrupting device that its interrupt request has been acknowledged, and that it is to get ready. In this cycle, even though




Figure 8.4 | Generation of a type number during the INTA cycle

the type number appears on the output of the TSB, the bus of the 8086 is in the high impedance state and so nothing happens. However, during the second machine cycle, when the INTA line is lowered, the lines D_0 to D_7 of the processor are enabled, and the switch settings of the TSB are transferred to the data lines D_0 to D_7 of the 8086, constituting the type number. Thus it is understood that this INTR interrupt is now vectored to type 81H (for Fig 8.4).

Thus, INTR is an interrupt pin using which, an interrupt can be funneled to any location in the IVT and thus, facilitate flexibility for use in any application. It is also possible that more than one device can place interrupt requests on this line, and each of them may have to be associated with different type numbers. All these necessitate the need of extra hardware to manage the possible interrupts arriving on this pin from various peripherals. This is done by a peripheral chip named Programmable Interrupt Controller (PIC) 8259, the details of which will be presented in Chapter 10. The INTR pin along with the PIC is used in PCs for managing various hardware devices.

Can We Consider the Reset Pin as an Interrupt Pin?

When the reset pin is activated, the processor resets and control branches to the absolute address FFFF0H . Thus reset has a particular pointer or vector – which is just the way an interrupt functions. Thus RESET is sometimes included in the list of hardware interrupts along with NMI and INTR .

For a PC, the process of resetting also occurs when the system is powered on. Let us note some points regarding ‘power on’. In fact, the address FFFF0H is designated as the ‘power on reset vector’. This address (FFFF0H) is in ROM and contains a far jump instruction to the beginning of the BIOS chip ‘power on reset code’. To find this vector, go to command prompt C drive, type ‘debug’ and get the contents of the reset location. This is shown below.

```
C:\>debug
-dfff:0000
FFFF:0000 EA 5B E0 00 F0 30 36 2F-32 30 2F 30 35 00 FC 00. [...06/20/05...
```

The first 5 bytes, constitute the assembly code for the jump instruction i.e., $\text{JMP (EA) F000:E05B}$. The next eight bytes were originally called the “RELEASE MARKER” by IBM,

as they always contain the release date of the code i.e., BIOS (which is 20th June 2005 for this PC). See the ASCII code of this marker:

30 36 2F 32 30 2F 30 35 i.e., 06/20/05 in the format mm/dd/yy.

8.5 | Priority of Interrupts

When many interrupts occur at the same time, which source gets its request honored first? The processor decides the priority. The order of priority is set in the following manner:

- i) Internal interrupts and Software interrupts – get the highest priority,
- ii) NMI,
- iii) INTR – gets the lowest priority.

Consider a case when the INTO, NMI and INTR occur simultaneously. Because of the inbuilt priority mechanism, the internal interrupt (INTO) will be taken up for servicing first. One step in the interrupt response would be to disable IF. Hence, INTR will not be taken up now. However, the NMI can interrupt the current ISR, as its functioning does not depend on the setting of the IF. Thus the NMI will interrupt the INTO ISR. The end result is that effectively NMI gets serviced before the internal interrupt. This is fine, because NMI caters to highest priority. After the NMI ISR is processed, the INTO ISR is taken up, and after it is done with, the request on the INTR is dealt with.

8.6 | Interrupt Type Allocation for Current PCs

The allocation of type numbers for the PC is shown in Table 8.2. This will be our reference when discussing video and keyboard interrupts.

8.6.1 | BIOS and DOS Interrupts

We have talked about these interrupts before and know that they correspond to routines already written by system designers. Both these classes of interrupts cater to dealing with input and output devices. DOS interrupts come along with the operating system (MS-DOS). BIOS routines are stored in ROM and thus are also called ROM BIOS. Since ROM access is rather slow, the current practice is of having the BIOS copied into RAM (which is faster) as well. This is called shadow memory, which is write protected to prevent tampering of the BIOS routines. In the beginning of the PC development era, there were problems due to non-standard BIOS but now standardization has been more or less agreed upon. Table 8.2 shows the type numbers of interrupts catering to specific devices. Even for a particular interrupt type, the routines that have been written have various functions and we can use them appropriately by passing parameters to designated registers. In comparing BIOS and DOS routines, it will be noticed that BIOS gives a greater degree of control than DOS routines, and are thus closer to the actual hardware. DOS routines are more 'high level' in comparison, and they use BIOS routines for peripheral access. Now we will use a few BIOS interrupts, the first of which is INT 10H which is for video applications, mostly pertaining to the video display.

8.6.2 | BIOS 10H Functions

To be able to use these functions effectively, we first have to make a study of the video display. We need to have an overall/idea (not all the intricate details) of its history such as components and standards. So let us now embark on the study of the video system of a typical PC.

Table 8.2 | Interrupt Allocation for Type Numbers 0 to 21 in the IBM PC

Interrupt type no. (HEX)	Function
0	CPU – Divide by zero (exception interrupt)
1	CPU – Single step
2	CPU – Non maskable interrupt (NMI)
3	CPU – Break point instruction
4	CPU – Overflow
5	BIOS – Print screen
6	CPU – Invalid op-code
7	CPU – Math coprocessor not present
8	System timer interrupt (IRQ 0)
9	Keyboard data ready (IRQ 1)
A	Reserved (IRQ 2) cascade from slave interrupt controller
B, C	Hardware interrupt for serial communication
D	Parallel port hardware interrupt (IRQ 5, LPT 2)
E	Diskette controller hardware interrupt (IRQ 6)
F	Printer hardware interrupt (IRQ 7, LPT 1)
10	BIOS – Software interrupt to use video
11	BIOS – Equipment check call
12	BIOS – Memory check call
13	BIOS – Software interrupt to use hard drive
14	BIOS – Software interrupt to use serial port
15	Not used
16	BIOS – Keyboard software interrupt
17	BIOS – Printer software interrupt
18	BIOS – ROM BASIC loader
19	BIOS – Bootstrap loader
1A	BIOS – System and real time clock software interrupt
1B	BIOS – Control-break handler
1C	BIOS – System timer tick handler
1D	BIOS – Video initialization parameter table pointer
1E	BIOS – Diskette initialization parameter table pointer
1F	BIOS – Graphics display character bitmap table
20	DOS – Program terminate
21	DOS – Function request services

8.6.3 | Video Adapter

The video display along with the associated hardware and software is one of the most complex parts of the PC. To be able to get a certain amount of accessibility to it, we need to develop a basic understanding of the video system of a PC. In any PC, there is the display unit and a 'video adapter'. This adapter is the hardware that provides the system its display capabilities. It is also designated by names such as video card and graphics card. It is the hardware that works between the processor and monitor. It relays the information received from programs and applications running on the system to the monitor and provides visual representation of the results.

Previously, the video adapter was plugged into a slot in the motherboard – now it is in the motherboard. Also, there is usually an additional video card with an inbuilt processor and RAM for improving the graphic capability of the PC. This is called a graphic accelerator.

Now, let us see the basic concepts of video adapters and displays. The video display usually has a 'controller' which takes care of the complexities of the display mechanism. Besides that, it needs RAM to store the data to be displayed. The IBM PC memory map (Table 7.3.) shows that locations from A0000H to BFFFFH (128K) is dedicated to video, and this area of memory is called video display RAM (VDR) or video memory or VRAM.

8.6.4 | History

In 1981, IBM released the MDA (monochrome display adapter) which could only give a black and white display. Subsequently, CGA (color graphics adapter) appeared which could process color too. Improvements over the CGA, for color display are the EGA, VGA and XGA. These higher level display systems can also 'emulate' the lower ones. Thus, if we have VGA in our PC, it can also be made to work in a mode corresponding to MDA or CGA. For color displays, CGA is the lowest common denominator and here, we will use the mode corresponding to this standard, even if the display we have on our PC is VGA or XGA. The mode setting can be done using BIOS interrupts with specified function numbers. See what each of these display designations denote.

CGA: Color graphics adapter

EGA: Extended graphics adapter

VGA: Video graphics adapter

SVGA: Super VGA

8.6.5 | Text and Graphics

The most basic classification of the display mode is 'text or graphics'. In text mode, a monitor can display only ASCII characters. In graphics mode, a monitor can display any bit-mapped image. The emphasis of our discussions here will be the text mode. In addition to the text and graphics modes, video adapters offer different modes of screen resolution and varying number of colors used.

CGA: This being the basic mode of color displays, we will concentrate on the CGA. Table 8.3 shows the text mode details for CGA. Only the text modes have been shown, and it corresponds to the CGA standard. VRAM shown in Table 7.2 specifies a memory space in RAM of 128 KB from location A0000H to BFFFFH. A part of this is used to store the data to be displayed. All the color modes in CGA have the video memory to start from B8000H. See the third row which corresponds to mode 2. It shows a screen to be composed of 25 rows and 80 columns, which means that 2000 elements can be displayed on the screen. Each element of text display is to have a data of 2 bytes – one byte for the ASCII character and one byte for

the color attribute. This will take 4 KB for one page of display. 16K memory is totally allotted for this, and thus 4 pages can be in use at any time. Each display begins on a 4K boundary, at address – B800:0000, B800:1000, B800:2000 and so on. Only one page can be displayed at a time which is called the active page, but it is possible to switch to other pages reasonably fast. Note that when the screen resolution is less, more pages are possible (25×40 of modes 0 and 1 have 8 pages). For monochrome display, the video memory allocation starts from B000:0000. In our discussions ahead, we will use the color screen with size 25×80 .

8.6.6 | Character Display

In video memory, for any character, the ASCII value is in the even address and the attribute byte is in the odd address. The attribute byte decides the color of the background and text. The order of storing the data is that in the first word location starting from B800:0000, all the 80 character words of the first row are stored, then the data of the second row, and so on; in other words, data is stored in rows. For 25 rows and 80 columns, the co-ordinates of the first and last character elements are (0, 0) and (24, 79). This is shown in Fig 8.5.

Table 8.3 | Text Mode Details of CGA

Mode	Type	Rows × Columns	Pages	Starting address in VDR	No. of colors
00	Color	25×40	8	B800:0000	16
01	Color	25×40	8	B800:0000	16
02	Color	25×80	4	B800:0000	16
03	Color	25×80	4	B800:0000	16
07	Monochrome	25×80	1	B000:0000	2




Figure 8.5 | Text mode video screen

8.6.7 | Color Configuration

The bit configurations for the various possible colors are given in Table 8.5 and 8.6. All this is for the attribute byte which has the format shown in Table 8.4. This byte defines the background and foreground colors. Foreground refers to the color of the character displayed. Background is the color that surrounds a character. It is well known that adding the primary colors red, green and blue (RGB) in suitable proportions can give a resultant color. By adding intensity (brightness) to this color, we get a lighter color.

Now refer to Table 8.5. This shows the 16 possible colors that can be obtained for the foreground. Adding R, G and B gives only 8 colors, but when intensity is another parameter, 16 colors are obtained. With intensity added (I = 1) and when intensity is not added (I = 0), two different shades of the same color are obtained. For example, in the fifth row of the table, the color specified is red with the I bit equal to 0. For the same value of RGB, when I = 1, the color becomes light red. For the background, only 8 colors are possible and this is given in Table 8.6.

Referring to Table 8.4, in the attribute byte, 'I' and 'B' apply to the foreground (text) only. 'I' means intensity and 'B' means blinking. What is blinking? Blinking happens when the background attribute is substituted for the foreground, every two seconds. Thus, in the attribute byte,

Table 8.4 | Format of the Attribute Byte

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
B	R	G	B	I	R	G	B
BLINK = 1 NORMAL = 0		BACKGROUND		FOREGROUND			

Table 8.5 | The 16 Colors of the Foreground–Bit Assignment

I	R	G	B	Color
0	0	0	0	Black
0	0	0	1	Blue
0	0	1	0	Green
0	0	1	1	Cyan
0	1	0	0	Red
0	1	0	1	Magenta
0	1	1	0	Brown
0	1	1	1	White
1	0	0	0	Gray
1	0	0	1	Light blue
1	0	1	0	Light green
1	0	1	1	Light cyan
1	1	0	0	Light red
1	1	0	1	Light magenta
1	1	1	0	Yellow
1	1	1	1	Bright white

Table 8.6 | The 8 Colors of the Background-Bit Assignment

R	G	B	Color
0	0	0	Black
0	0	1	Blue
0	1	0	Green
0	1	1	Cyan
1	0	0	Red
1	0	1	Magenta
1	1	0	Brown
1	1	1	White

if MSB = 1, it corresponds to ‘blinking’. Also, if the background and foreground colors are the same, the displayed character is invisible.

See the bit assignment of the attribute byte in Table 8.4.

Example 8.2

Find the attribute byte for the following cases of text mode display.

- i) Black background with bright white text (foreground)
- ii) White background, red text (foreground)
- iii) Blinking with cyan background with brown text (foreground)

Solution

Look at Tables 8.4, 8.5 and 8.6.

- i) For black background, the RGB values are 000, and since blinking is not mentioned, it is ‘normal’. Hence, the nibble corresponding to the background is 0000. This is to be the bits D_7 to D_4 of the attribute byte.
The lower four bits of the attribute is 0111
Thus, the attribute byte is 0000 0111 i.e., 07
- ii) No blinking. Hence, $D_7 = 0$.
White background needs 111
Red foreground (text) needs 0100
Thus, the attribute byte is 0111 0100 i.e., 74H
- iii) Blinking and cyan background needs 1011
Brown foreground (text) needs 0110
Thus, the attribute byte is 1011 0110 i.e., B6H

8.7 | BIOS 10H Functions

In Table 8.2, of interrupt vector allocation, it can be seen that the type number 10H has been allocated for video display applications. In this category, a number of actions are possible

using different function numbers to be loaded in AH. In addition, each function requires that various information bytes be loaded in specified registers. Appendix D gives the list of functions of INT 10H. In this section, we shall use some of them (but not all) for various applications.

i) BIOS 10H Function 0 – Set video mode.

This function is for setting the ‘video mode’ – the function number is placed in AH, and the video mode in AL. This function is for setting the video mode. It also clears the screen.

Example 8.3

```

.MODEL TINY
.CODE
.STARTUP

MOV AH, 0           ;set function number
MOV AL, 3           ;mode 3, standard text mode
INT 10H             ;the screen is found to be cleared
.EXIT
END

```

This program clears the screen and sets the display for mode 3 (CGA).

ii) BIOS 10H Function 02 – Set cursor position.

This is for positioning the cursor, by specifying the row and column co-ordinates. (The row co-ordinate varies from 0 to 24 and the column from 0 to 79). The following registers should contain the relevant information. BH = page number (default page is 0), DH = row, DL = column.

Example 8.4

```

.MODEL TINY
.CODE
.STARTUP

MOV AH, 0           ;set function number
MOV AL, 3           ;mode 3, standard text mode
INT 10H             ;the screen is cleared with this.

MOV AH, 02          ;function number for setting cursor position
MOV BH, 0           ;specify page (which is the default page)
MOV DH, 09          ;row co-ordinate
MOV DL, 35          ;column co-ordinate
INT 10H

.EXIT
END

```

iii) Function 06 – Scroll up screen.

This function scrolls up the screen by the specified number of lines and blank lines appear at the bottom. The registers involved and the values to be stored in them are:

AL = number of lines to be scrolled up (0 for a full blank screen)

BH = attribute value or pixel value, CH:CL = starting row: column,

DH:DL = ending row: column

See the following program segment.

```

MOV AH, 06          ;function number
MOV AL, 00          ;number of lines to be scrolled up
MOV BH, 30H         ;cyan background, black foreground
MOV CH, 0           ;starting row co-ordinate = 0
MOV CL, 0           ;starting column co-ordinate = 0
MOV DH, 24          ;starting row co-ordinate = 24
MOV DL, 79          ;starting column co-ordinate = 79
INT 10H

```

This scrolls up 0 lines, and a full blank screen of cyan color is seen. Any text written on the screen will be of black color. The starting point is 0:0 (left hand corner of the screen), and the end point is the right corner i.e., the point (24, 79).

This function number can also be used to create a window of a desired size given the co-ordinates of the window.

```

MOV AH, 06          ;function number
MOV AL, 9           ;number of lines to be scrolled up
MOV BH, 61H         ;brown background, blue foreground
MOV CH, 15          ;starting row co-ordinate = 15
MOV CL, 20          ;starting column co-ordinate = 20
MOV DH, 23          ;ending row co-ordinate = 23
MOV DL, 60          ;ending column co-ordinate = 60

```

The above program segment creates a window with the given co-ordinates.

Now, if two windows are to be created, repeat the above program segment with a different set of co-ordinates and a different color (for clarity of understanding). Because the attribute of a window remains unchanged until another operation changes it, we can set many different windows at the same time.

There is another function 'Scroll Down screen', with function number 07. This works similar to function 06, except that scrolling down the screen causes the bottom lines to scroll off and blank lines to appear at the top.

iv) INT 10 H Function 09 – Display character with attribute at cursor position.

This function displays a character at the cursor position.

AL = ASCII character, BH = page number, BL = attribute byte, CX = count

```

MOV AH, 09          ;function number
MOV AL, 'S'         ;display 'S' at the cursor position
MOV BH, 0           ;page 0
MOV BL, 16H         ;blue background, brown foreground
MOV CX, 9           ;display the character 9 times
INT 10H

```

The above program segment displays 'S', nine times starting at the position of the cursor.

Note This function does not advance the cursor. So if we try to display the word 'Hello' at the cursor position, the characters get displayed one on top of the other, unless we arrange to advance the cursor after the display of each character.

Example 8.5

The following program does the following:

- i) Sets the video mode and clears the screen.
- ii) Makes a window of a specified size and color.
- iii) Sets the cursor at a specified position within the window.
- iv) Displays 10 times the character '*' at the cursor position.

Solution

```

.MODEL TINY
.CODE
.STARTUP

MOV AH, 0           ;set video mode
MOV AL, 3           ;mode 3, normal text mode
INT 10H             ;clear screen

MOV AH, 06          ;function number for scrolling up
MOV AL, 10           ;no of lines = 10
MOV BH, 61H         ;brown background, blue foreground
MOV CH, 5           ;starting row co-ordinate = 5
MOV CL, 20           ;starting column co-ordinate = 20
MOV DH, 14          ;ending row co-ordinate = 14
MOV DL, 60          ;ending column co-ordinate = 60
INT 10H             ;create the scroll window

MOV AH, 02          ;function number to set the cursor
MOV BH, 0           ;page number = 0
MOV DH, 09          ;row co-ordinate of cursor
MOV DL, 39          ;column co-ordinate of cursor
INT 10H             ;cursor at (9, 39)

MOV AH, 09          ;function to display character
MOV AL, '*'         ;character to be displayed = '*'
MOV BH, 0           ;page number
MOV BL, 16H         ;blue with brown text
MOV CX, 10          ;count = 10
INT 10H             ;display the character 10 times

.EXIT
END

```

You may note that this program uses all the previously discussed functions of BIOS 10H for the purposes of clearing the screen (AH = 0), creating a scroll window (AH = 06), setting the cursor position (AH = 02) and then displaying the character '*' in the cursor position (AH = 09) 10 times.

Remember that the cursor position has not advanced by this. So how do we write a string of different characters? We must fix up an initial value for the cursor position and then advance the cursor for each new character to be displayed.

Example 8.6

Write a program to display 'HELLO' vertically downwards at the centre of the screen.

Solution

Here we are to display a string vertically downwards. For that we will do the following:

- i) Fix up the initial cursor position. Write the first character there.
- ii) Next, increment the row co-ordinate of the cursor, and display the next character there. Repeat this for the whole string.
- iii) To avoid writing a long program, let us use two macros – one for setting the cursor (SET_C) at different row co-ordinates (the column co-ordinate is kept constant at 38) and another for displaying the different characters at each cursor position (DISP).
- iv) The output of this program is a vertical display of HELLO, with blue text blinking on a brown background. The brown and blinking together will cause the background to appear yellowish. The cursor position is at (15, 38) after the display.

```

.MODEL SMALL
.DATA
COL DB 38                ;specify the column co-ordinate of
                           cursor
.CODE
.STARTUP

SET_C MACRO ROW          ;macro for positioning cursor
    MOV AH, 02           ;function for setting the cursor
    MOV BH, 0            ;page 0
    MOV DH, ROW          ;specify the row co-ordinate
    MOV DL, COL          ;specify the column co-ordinate
    INT 10H
ENDM

DISP MACRO CH            ;macro for displaying character
    MOV AH, 09           ;function for display
    MOV BH, 0            ;page 0
    MOV BL, 0E9H         ;blinking, brown background, blue text
    MOV AL, CH           ;the character to be displayed
    MOV CX, 1            ;display the character only once
    INT 10H              ;clear screen
ENDM

    MOV AH, 0            ;set video mode and clear screen
    MOV AL, 3            ;mode 3
    INT 10H

```

```

SET_C 11          ;set cursor to row 11
DISP 'H'          ;display 'H'
SET_C 12          ;set cursor to row 12
DISP 'E'          ;display 'E'
SET_C 13          ;set cursor to row 13
DISP 'L'          ;display 'L'
SET_C 4           ;set cursor to row 14
DISP 'L'          ;display 'L'
SET_C 15          ;set cursor to row 15
DISP 'O'          ;display 'O'

.EXIT
END

```

Example 8.7

Next, let us write a simple program to make the screen green in color on pressing the 'G' key.

Solution

```

.MODEL TINY
.CODE
.STARTUP
MOV AH, 01          ;DOS interrupt for keyboard entry
INT 21H             ;the code of the key pressed is in AL
CMP AL, 'G'         ;compare AL with the code of 'G' key
JNE LAST           ;if not equal, jump to LAST

MOV AH, 0           ;set video mode
MOV AL, 02          ;mode = 2
INT 10H             ;clear screen after setting the mode
MOV AH, 09          ;display function
MOV BH, 0           ;page 0
MOV BL, 22H         ;green background with green foreground
MOV CX, 2000        ;number of characters = 2000
INT 10H

LAST:
.EXIT
END

```

In Example 8.7, the DOS interrupt 21H with AH = 01, is used for keyboard entry. The ASCII value of the character of the key pressed will be obtained in AL, which is compared with 'G'. In the display function, we have not loaded any character in AL to be displayed. This is because the attribute byte in BH is defined with the same foreground and background. Thus, whatever the character in AL, it will not be visible. We give a count of CX = 2000, which is the number of characters on the screen i.e., 25×80 (2000) characters can be displayed.

8.8 | Addressing Video Memory Directly

Now, let us try to address video RAM which is just like any other RAM. We know that the starting address of video memory for CGA display is B800:0000. We will attempt to display our data without using BIOS, rather, we write data directly into video RAM which causes displaying of that information. Let us try to make the whole of page 0 display appear green. For this a 16-bit number with the attribute byte and ASCII character is loaded in AX. This word is written into the 2000 locations (25×80 characters) starting from the address B800H. Using string instructions, this display can be easily accomplished.

Example 8.8

```

.MODEL SMALL
.CODE
.STARTUP
MOV AH, 0
MOV AL, 3
INT 10H           ;this is used to clear the screen

CLD               ;clear direction flag
MOV AX, 0B800H    ;starting address of video memory
MOV ES, AX        ;make video memory the extra segment
MOV DI, 0         ;to point to first location in ES
MOV CX, 25*80     ;characters on the screen (2000)
MOV AH, 29H       ;attribute for green background
MOV AL, 20H       ;ASCII code of space bar key
REP STOSW         ;store this data in the 2000 locations

.EXIT
END

```

Let us discuss the scheme of Example 8.8.

- i) First, the whole screen is cleared using the INT 10H, function 00.
- ii) Next, we use string instructions to address video memory. The ES register is loaded with the starting address of video memory.
- iii) DI is loaded with the offset 0.
- iv) CLD is for auto-incrementing this pointer.
- v) The STOSW (store string word) is repeated 2000 times i.e., the content of AX is stored in all the locations of page 0 of the video memory. For using the STOSW instruction, the destination memory should be defined to be the 'extra segment' and should be pointed by DI. This has been done here.
- vi) The content of AX is the attribute byte (AH) and the space bar character (AL). The space bar character display is a 'blank' display. Hence, the screen has only the background color i.e., green.

In Example 8.9, the message 'HELLO WORLD' is displayed at the centre of the screen. The background color is set to green with blinking (signified by the number A9H). Because of 'blinking', the green of this screen can be observed to be different from the green of Example 8.10. 20H corresponds to the ASCII code of the space bar key (no character).

First, the whole screen is made green by setting the background color to be green and the foreground to be blank by filling it with the space bar code. Next, the message to be displayed is displayed starting in the middle of the screen, by locating the 2000th point (centre of the 25×80 screen). Remember that each character takes two bytes of storage.

Example 8.9

```

.MODEL TINY
.CODE
.STARTUP
MSG DB 'HELLO WORLD'

MOV AH, 0           ;set video mode
MOV AL, 03          ;mode = 2
INT 10H             ;clear screen after setting the above

    CLD              ;clear direction flag
    MOV AX, 0B800H    ;starting address of video memory
    MOV ES, AX        ;make video memory the extra segment
    MOV DI, 0         ;point to first location in ES
    MOV CX, 25*80     ;characters on the screen (2000)

    MOV AH, 0A9H      ;green background, blinking blue text
    MOV AL, 20H        ;ASCII code of space bar key
REP  STOSW            ;store this in the 2000 locations
                     ;the screen is now green in color

    MOV AX, 0B800H    ;starting address of video memory
    MOV ES, AX        ;make video memory the ES
    MOV DI, 2000      ;point to the centre of the screen
    MOV SI, OFFSET MSG ;SI to point to message
    MOV AH, 0A9H      ;green background, blinking blue text
    MOV CX, 11        ;characters in the message = 11
START:
    LODSB             ;load each character in AL
    STOSW              ;store each word in video memory
    LOOP START        ;repeat until CX = 11
    .EXIT
    END

```

The salient features of this program:

- i) The foreground and background of the screen are made green and blank just as was done in Example 8.10. This is the first part of the program.
- ii) Then, each character is brought to AL using the instruction LODSB. For using this instruction, the data segment should contain the source data (the message) and SI should point to it. This has been done here.
- iii) The attribute byte is already in AH. The character to be displayed, along with the attribute (in AH) is now in AX.

- iv) The content in AX is now stored in the video memory, using STOSW. DI should point to the 2000th location of the video memory, since that is where the storing of the content of AX should start.
- v) The LODSB and STOSW instructions are executed until CX = 0. Recollect that for using the LODSB instruction, the message should be in the data segment and should be pointed by SI. Here, the data segment is the same as the code segment as we are using the tiny model.
- vi) Thus, we find the string HELLO WORLD displayed on a green screen.

8.9 | Keyboard Interfacing

Next, let us use BIOS interrupts for the computer keyboard. For that, we first have to understand the computer keyboard, as well as the hardware associated with the keyboard.

8.9.1 | Computer Keyboard

All of us are familiar with the PC keyboard, which is more or less standard. However, the keyboard which is one of the most important, (but taken for granted) part of the PC did not always look like this. In the first IBM PC and PC-XT, there were 83 keys. Later the PC-AT keyboard had the same number of keys, but the arrangement of the keys was different. After taking into consideration, various design ideas, the current PC keyboard has been designed to have 101 keys, including the function keys F1 to F12. This is called the 'enhanced keyboard'. We are all very much used to the keyboard of the PC, but not everyone really knows the functions of each and every key. So, let us start with a discussion on the keys of the keyboard.

Character keys We have first, the character keys a to z, and above it there are the number keys. All these keys have dual function, as they are pressed along with, or without the shift key. These are the keys most commonly used, and most familiar to us.

Function Keys These keys marked F0 to F12 and located on the topmost row of the keyboard, were widely used with many older DOS programs. They are still very popular as keyboard shortcuts. Sometimes, they are used in combination with other keys like Ctrl, Alt or Shift keys. The help menu of any program will be seen to contain shortcuts with such combinations.

Enter (Return) Key This key is used to enter commands or to move the cursor to the beginning of the next line. If we make a choice, out of a set, this key specifies the choice.

Escape Key It is used to escape from a program, or exit to a previous screen. It is helpful also, to get out of difficult situations as many of us has experienced.

ALT (Alternate) Key This key is used in combination with other keys like character keys or function keys and is as defined, used in various programs. For example, in the TC++ compiler, Alt +F9 is the command for compile.

Ctrl (control) Key This key is also used in combination with other keys, and it is defined for certain programs. For example, in MS Word Ctrl+c copies data and Ctrl+v pastes the copied data. This is not necessary valid for all programs – for example, not for the TC++ editor, anyway.

Caps Lock Key The Caps Lock key is a toggle key. Pressing it once turns it on. Pressing it again turns it off. Some computer keyboards have a light or indicator that shows when the Caps Lock is on. When Caps Lock is on, every letter that is typed will be a capital letter. The Caps Lock key on a computer keyboard affects only letters. It has no effect on the number or symbol keys.

Num Lock and Numeric Keypad Most desktop computers have a numeric keypad. It is on this numeric keypad that the page Down, page Up and other arrow keys are seen. The Num Lock is a toggle key, which when turned on allows numbers to be typed in. When this key is toggled off, the other options of each key will be activated, for example Pg Dn and Pg Up. Note that there are other keys also on the numeric keypad.

Windows Key Most windows based computers have a 'Windows' key. It is marked with a Microsoft windows symbol and is in the bottom row of the keypad, and there may be two such keys on either side of the space bar key. Pressing this key will open the start menu. It can be used in combination with other keys for some short cuts. For example, Windows+d will minimize all open windows and one more such key press will maximize them.

Application Key Another key, which most of us might not have noticed at all is a key which has a symbol of a page with a number of lines. This is called the application key and it works like a right mouse click. Try using it in the middle of typing, using this key is much easier than having to get to the mouse and click it.

8.9.2 | Keyboard Hardware

Now, let us discuss the hardware associated with the keyboard of a modern PC. The keyboard along with the hardware involved is quite a complicated part, but here the attempt will be to present it in a simple form with important details covered and more details left to be referred from other sources.

The earlier PCs had a keyboard of 83 keys only, but now the standard keyboard has 101 keys including the alphanumeric characters. The keyboard we are assuming is the PS/2 keyboard and not the ones which use the USB connector (there are some differences in this case). Let us summarize some of the points regarding the keyboard.

8.9.2.1 | Keyboard Controller

The keyboard which is connected to the motherboard through a cable has a microcontroller in it, which belongs to the 8042 family. In the motherboard, there is another similar microcontroller (or its emulation, as in present times) and they communicate through a bidirectional serial communication protocol.

8.9.2.2 | Scan Code

The microcontroller inside the keyboard is programmed to detect whether a key has been pressed and identify the key, in case of a key being pressed. When a key is pressed, the microcontroller therein sends to the motherboard a unique 'scan code' for each key. A key press is generally followed by a key release and this 'release' causes another scan code to be sent, which is **different** from that of the key press of the same key. Key press and release are also called 'make' and 'break'. The make and break scan codes are different for a key. They differ by 80H. For example, if the make scan code for a character is 9, its break scan code will be 89H. **Incidentally, the scan code has nothing to do with the ASCII or other values of any key, but is a unique code defined by IBM.** (Tables 8.7 to 8.9, at the end of this chapter give the scan codes for the 101 keys of a PC).

Keyboard Status Bytes On checking this list of scan codes, it will be seen that the scan code of a key is the same whether the key is a lowercase or upper case alphabet. Also, it is the same whether the shift key has been pressed or not, or if the Caps lock is on or off. See the list of scan codes in Tables 8.7 to 8.10. The scan code of * and 8 which are on the same key is '09'. The scan code of g and G is 22H. How is that accounted for?

The answer is that there are two keyboard status bytes which are held in the BIOS data areas 40:17 and 40:18. The first status byte maintains the current status of the modifier keys on the keyboard. The bits have the following meanings as shown in Fig 8.6a. If a toggle key is on, the corresponding bit is 1. For example, for a caps lock key being toggled on, the corresponding bit is 1, otherwise it is zero. For the shift, alt and ctrl keys, '1' indicates that the corresponding keys have been pressed. The second keyboard status byte specifies if the corresponding keys are 'currently' down.

8.9.3 | Keyboard Interrupt Type 09


In Table 8.2, we see that the interrupt with type number 09 has been allocated to the keyboard on IRQ1 line. Thus, this is a hardware interrupt. The IRQ line stands for 'interrupt request line 1' of the 8259 PIC/(Refer Chapter 10) which is vectored to type 9 of the interrupt vector table. Thus, when a key is pressed, its scan code is sent to the motherboard in serial form, and therein converted to parallel form and this 8-bit scan code is presented to port A of the 8255 (refer Chapter 9) with I/O address 60H, and along with this, the IRQ1 line is activated, which has been vectored to INT 9.

8.9.3.1 | BIOS INT 09 Routine

What is the INT 09 routine?


- i) The ISR reads the scan code from Port 60H and identifies the key pressed. There are three possibilities for this.
 - a) The key may correspond to special keys like shift, alt, control etc. In that case, the corresponding bit in the **status byte** (Fig 8.6a) will be set.
 - b) The second case is that the key may correspond to an ASCII character and then the 8-bit scan code and the ASCII value are stored in the keyboard buffer (which is a 32-byte area in the BIOS data area starting from 0040:001E).
 - c) Another case is when there is no ASCII code for the key – this is for function keys F1 to F12. In this case, the ISR saves the scan code in the buffer and also puts 00 in place of the ASCII code.
- ii) When the key is released, it sends a different scan code, and this is verified by the ISR and confirmed to be a key release.
- iii) There is a possibility that the key continues to be kept pressed. If it is found so for more than half a second, the ISR identifies it as a new key press and sends its scan code accordingly. Repeating the same key is referred to as 'typematic' in IBM literature.

Next, let us use some software interrupts related to the keyboard. We have already used a few DOS interrupts for keyboard access. Now we will get acquainted with a few BIOS interrupts for keyboard functions. The software interrupt type 16H has been allocated for the keyboard.



Note The corresponding bit is '1' for the keys pressed or toggled.

Figure 8.6a | First keyboard status byte



Note The corresponding bit is '1' for the keys kept pressed.

Figure 8.6b | Second keyboard status byte

8.9.4 | BIOS 16H Functions

A few BIOS interrupts for keyboard control are presented here. Some of the early BIOS functions catered to the old 83-key keyboard only. Later, some more BIOS functions were added for the enhanced keyboard.

- i) AH = 0 or AH = 10H. The first function number is for the older keyboard. The second one is the equivalent one for the enhanced keyboard. This function checks for a character in the keyboard buffer. If available, the scan code is returned in AH, and ASCII value in AL. For

function keys (F0 to F12) which have no ASCII values, AL = 0. If no character is available in the keyboard buffer, the function waits for a key press.

- ii) AH = 01 or AH = 11H. The first function number is for the older keyboard. The second one is the equivalent one for the enhanced keyboard. This function is similar to the previous one, except that if no character is available in the keyboard buffer, it does not wait for a key press. It simply sets ZF (ZF = 1) and returns.
- iii) AH = 02 or AH = 12H. The first function number is for the older keyboard. The second one is the equivalent one for the enhanced keyboard. This function returns the first keyboard status byte in the AL register. This status byte is also available in the BIOS data area 0040:0017.

Example: 8.10

Explain what the following program does.

```

.MODEL TINY
.CODE
.STARTUP

MOV AH, 10H
INT 16H
MOV AH, 12H
INT 16H

.EXIT
END

```

Solution

There are two interrupt functions in the above program.

- i) The function with AH = 10H, waits for a key press. If the key pressed is 8, then AL = 38H and AH = 09. The value 09 is the scan code of 8 and 38H is its ASCII value of 8.
The key 8 also represents the “*” character which is displayed when the shift key is also pressed. Then AH = 09 and AL = 2AH which is the ASCII value of “*”. Thus, the scan code is the same for both the characters, but their ASCII values are different.
If the key pressed is F5, the content of AL = 0, because the function keys F0 to F12 have no ASCII values. The content of AH = 3FH which is the scan code of the F5 key.
- ii) The function with AH = 12H, provides the keyboard status byte in AL.
With no shift, alt or control keys pressed, AL = 0.
With CAPSLOCK on, NUMSLOCK on, and the left shift key kept pressed, the value in AL = 62H i.e., 0110 0010B.
Keep the left shift key pressed, and check the content of the BIOS data area 0040:0017.

```

-d0040:0017
0040:0010      62-00

```

We find the first keyboard status flag saved there.

Example 8.11

The following program fills up the screen with the key pressed. If no key is pressed, it keeps waiting for a key press.

```

.MODEL TINY
.CODE
.STARTUP
NO_KEY: MOV AH, 11H      ;function for waiting for key
        INT 16H
        JZ NO_KEY        ;if ZF = 1, no key press, keep waiting
KEY:    MOV AH, 10H      ;function for taking into AL the key
        INT 16H
        MOV CX, 2000     ;number of characters to fill the screen
DISP:   MOV DL, AL       ;move the ASCII character to DL
        MOV AH, 02       ;function number to display DL content
        INT 21H          ;DOS interrupt 21H
        LOOP DISP        ;display until CX = 0
        .EXIT
        END

```

The program of Example 8.11 uses INT 16H function 11H. This function checks whether there is a character in the keyboard buffer. If there is no character and no key press either, it sets the Zero flag. The state of the Zero flag is used to wait for a key press. When a key is pressed, this loop is exited and it takes into AL, the ASCII value of the key pressed (using BIOS INT 16H function 10H). Next, the content of AL is moved to DL and DOS INT 21H function is used to display it 2000 times, so as to fill the whole screen.

Now that we have had a reasonably good exposure to the intricacies associated with the video display, keyboard and the BIOS interrupts associated with them, we can use them to manipulate the interrupt vectors of the PC.

8.10 | Hooking an Interrupt

A term associated with interrupts is ‘hooking an interrupt’. This just means the mechanism of installing an interrupt vector in the IVT. The tricky part of this is that the type number we plan to use may already be in use for another interrupt routine. This ISR is temporarily kept aside when we install a new vector corresponding to the same type number. To install an interrupt vector, the assembler needs to address absolute memory and the easiest way to do it is using certain DOS interrupt functions. Now, we will go through the whole process of hooking interrupts.

8.10.1 | Terminate and Stay Resident

DOS uses a part of the 640K of RAM in the lower area of memory, and the exact amount of this varies from version to version. Whenever an application program runs, a portion of RAM is allocated to the application. After this, the program is abandoned and the allocated memory is freed and is marked as available for other programs. If this is not done, the amount

of free memory keeps on getting reduced. However, there are programs which are kept 'resident' in memory, even after it is run and abandoned. So, when it is to be run again, it need not be brought from the disk again. This is quite comfortable for the specific program, but the disadvantage is that free memory is eaten up by such programs. Such programs are titled 'Terminate and stay resident (TSR) programs'. There is a specific interrupt 21H function which makes a program resident. Most of these TSR programs are invoked by interrupts, and let us see how it is done.

8.10.2 | DOS Functions for Interrupt Hooks

There are three functions which we will use to hook an interrupt vector.

i) INT 21H, Function 35H – Get Interrupt Vector.

Get the segment-offset value of an interrupt vector.

Input: AL = interrupt number. Output: ES:BX = address of the interrupt handler.

When we attempt to get our new ISR pointed by a particular interrupt vector, there is the possibility that there was already a previous ISR using that type number. So we have to save the value of CS:IP of that interrupt handler. This will be necessary for restoring the previous state of the interrupt vector table. This is done by Function 35H.

ii) INT 21H, Function 25H – Set Interrupt Vector.

Set an entry in the Interrupt Vector Table to a new address. Input: DS:DX points to the interrupt-handling routine that will be inserted in the table; AL = the interrupt number

We then have to get our ISR to be pointed by the interrupt vector we have decided upon. This is done by function 25H.

iii) INT 21H, Function 31H – Terminate and stay resident.

Terminate the current program or process, but make it resident in memory and attempt to set the current memory allocation to the number of paragraphs specified in DX. Input: AL = return code, and DX = requested number of paragraphs.

The function 31H makes our ISR resident in memory. The number of paragraphs (16-byte chunks) that our program occupies is to be specified in DX.

The program for installation of an interrupt hook is expected to be a com program and hence we proceed as follows.

Example 8.12

Here, the new ISR named 'NEW60' is to be installed as a type 60H interrupt. The ISR is a program to make the full screen cyan in color with black foreground.

```

.MODEL TINY
.CODE
.STARTUP

JMP START
OLD DD?

NEW60 PROC FAR
    MOV AH, 06                ;function number
    MOV AL, 00                ;no character
    MOV BH, 30H               ;cyan background, black foreground
```

```

        MOV CH, 0                ;row co-ordinate = 0
        MOV CL, 0                ;column co-ordinate = 0
        MOV DH, 24               ;row co-ordinate = 24
        MOV DL, 79               ;column co-ordinate = 79
        INT 10H
        IRET
NEW60   ENDP

START:  MOV AX, 3560H             ;get interrupt vector of type 60H
        INT 21H
        MOV WORD PTR OLD, BX     ;save old IP in OLD
        MOV WORD PTR OLD + 2, ES ;save old CS in OLD + 2

        MOV DX, OFFSET NEW60     ;ISR pointed by 60H vector
        MOV AX, 2560H            ;AL = 60H, AH = 25H
        INT 21H

        MOV DX, OFFSET START     ;find number of paragraphs
        MOV CL, 4                ;divide by 16, by shifting 4 times
        SHR DX, CL               ;to the right
        INC DX                   ;add 1 to it
        MOV AX, 3100H            ;exit with AH = 31H rather than 4CH
        INT 21H                  ;this makes NEW60 resident
        END

```

Let us examine the salient features of Example 8.12.

- i) This program installs a new ISR corresponding to type 60H vector. It is at offset NEW60. This means that the CS and IP of this program (ISR) are to be installed in the corresponding position in the interrupt vector table.
- ii) The new ISR is a program named NEW60. It does the job of making the screen cyan in color. It ends with the IRET instruction and is defined as a far procedure.
- iii) The program for saving the old CS and IP and installing the new CS and IP begins at the label START.
- iv) The function AH = 35H gets the old IP and CS in BX and ES. These are saved in 'OLD' which is a double word location.
- v) The function AH = 25H installs the CS and IP of NEW60 into the IVT.
- vi) The function AH = 31H makes the program NEW60 resident.
- vii) Note that the program does not end with AH = 4CH but has to end with AH = 31H.

The next issue is how to invoke this program which we have now installed. Since we have installed it at the type 60H location of the IVT, we can invoke it by the instruction INT 60H as in Example 8.15.

Example 8.13

```

.MODEL TINY
.CODE
.STARTUP
MOV AH, 0      ;function number for setting video mode
MOV AL, 3      ;mode 3
INT 10H        ;this is for clearing the screen
INT 60H        ;run the ISR now installed as Type 60H

.EXIT
END

```

Example 8.13 will make the screen cyan in color, because INT 60H gets invoked.

Note

- i) One thing to note is that the installation of the new ISR will be lost if the system is closed or re-booted. If you exit the command window, you will have to re-install it again i.e., run the program of Example 8.14 again, if you want to use the program in Example 8.15.
- ii) Never use DOS INT 21H functions in a TSR program, because DOS is **not reentrant**, and its usage in TSR programs will cause havoc.

What is a Reentrant Program?

Re-entrancy is a useful, memory-saving technique for multi-programmed timesharing systems. A Reentrant Procedure is one in which multiple users can share a single copy of a program during the same period. A programmer writes a reentrant program by making sure that no instructions modify the contents of variable values in other instructions within the program. Each time the program is entered for a user, a data area is obtained in which to keep all the variable values for that user. The data area is in another part of memory from the program itself. When the program is interrupted to give another user a turn to use the program, information about the data area associated with that user is saved. When the interrupted user of the program is once again given control of the program, information in the saved data area is recovered and the program can be reentered without concern that the previous user has changed some instruction within the program.

8.10.3 | Hooking Into Hardware Interrupts

Now, let us do the same with hardware interrupts. Let us choose the hardware interrupt with type number 09 which has been allocated to the keyboard on IRQ1 line. We will replace the normal keyboard sequence associated with this hardware interrupt, with another program and that program will be activated with a hotkey.

What is a hot key? It is defined as a user-defined key sequence that executes a command or causes the operating system to switch to another program. We are going to use a hot key in Example 8.14, most of which is similar to Example 8.14. The important points are:

- i) It replaces the normal routine associated with INT 9.
- ii) The hot key used is Alt+F10. To verify this combination, first the INT 16H routine with AH = 12H is used, to check for the ALT key. This tests for the D₃ bit of the data in AL

(first keyboard status byte) which is to be '1' if the Alt key has been pressed. Then port A of the 8255 (with address 60H) is tested for the scan code of F10 (which is 44H). If both these conditions are true, the new ISR is taken up for execution.

- iii) If it is found that the hot key pressed is not the right one, control goes to OVER, which retrieves the old value of CS and IP for the type 9 interrupt. This is very important, because the keyboard is locked (until a reboot) if the normal keyboard routine is not restored.
- iv) After our TSR is executed, then also control goes back to OVER for the same reason, as stated above.
- v) The hot key activated ISR is similar to Example 8.5 which brings up a video window.

Note The scan codes of keys are listed in Tables 8.7 to 8.9.

Example 8.14

[illegible]

```

MOV WORD PTR OLD, BX      ;save old IP
MOV WORD PTR OLD+2, ES    ;save old CS
MOV AH, 25H               ;install new vector
MOV AL, 09                ;at type number 9
MOV DX, OFFSET NEWINT9    ;point DX to the new ISR
INT 21H

MOV DX, OFFSET START      ;count the number of paragraphs
MOV CL, 4
SHR DX, CL
INC DX
MOV AX, 3100H             ;make the program resident
INT 21H
END

```

Example 8.15 is a similar program, except that the hot keys to be invoked are Ctrl+F3. The ISR for this hot key is copied from Example 8.8. It accesses the video memory directly and causes the full screen to be green in color.

Example 8.15

```

.MODEL TINY
.CODE
.STARTUP

JMP START
OLD DD?

NEWINT9 PROC FAR
    PUSH AX                      ;hot key is Ctrl+F3
    MOV AH, 12H
    INT 16H
    TEST AL, 04                  ;test if Ctrl key is pressed
    JZ OVER
    IN AL, 60H
    CMP AL, 3DH                  ;test for the scan code of F3
    JNE OVER

    MOV AH, 0                    ;the new ISR for Int 9
    MOV AL, 3                    ;this makes the screen green
    INT 10H
    CLD
    MOV AX, 0B800H
    MOV ES, AX
    MOV DI, 0
    MOV CX, 25*80
    MOV AX, 2920H
    REP STOSW

OVER: POP AX
      JMP CS:OLD
NEWINT9 ENDP

```

```

START: MOV AH, 35H
      MOV AL, 09
      INT 21H
      MOV WORD PTR OLD, BX
      MOV WORD PTR OLD+2, ES

      MOV AH, 25H
      MOV AL, 09
      MOV DX, OFFSET NEWINT9
      INT 21H

      MOV DX, OFFSET START
      MOV CL, 4
      SHR DX, CL
      INC DX
      MOV AX, 3100H
      INT 21H
      END

```

The next is a TSR program activated by the hot key combination ctrl+alt+i. It prints the word 'ANURAG' in green color in a red window with size as specified in the program. Since DOS interrupts cannot be used, the string is not displayed using the DOS interrupt with function number 09. Instead a macro is used with BIOS functions to print single characters. See Example 8.16.

Example 8.16

```

.MODEL TINY
.CODE
.STARTUP

JMP START
OLD DD?
CSET MACRO                ;macro for setting character position.
MOV AH, 02                ;function no to set the cursor
MOV BH, 0                 ;page no = 0
INC DL
INT 10H
ENDM

DISP MACRO SYMB           ;macro for displaying a character
CSET
MOV AH, 09                ;to display character
                        ;at cursor position
MOV AL, SYMB              ;display SYMB at cursor position
MOV BH, 0                 ;page = 0
MOV BL, 04AH              ;attribute byte
MOV CX, 1                 ;display the character once
INT 10H
ENDM

```



```

NEWINT9 PROC FAR                                ;the ISR to be installed.
PUSH AX
MOV AH, 12H                                     ;AH = 12H, kb s/w interrupt
INT 16H                                         ;gets into AL, kb status byte
TEST AL, 08h                                   ;check whether D3 is high - alt key
JZ OVER                                         ;if not, go to OVER
TEST AL, 04h                                   ;check if D2 is high - ctrl key
JZ OVER                                         ;if not, go to OVER
IN AL, 60H                                     ;read parallel port
CMP AL, 17H                                    ;compare AL with the scan code of i
JNE OVER                                       ;if not equal, go to OVER
                                              ;the following is the pop up program

MOV AH, 0
MOV AL, 3
INT 10H                                         ;set video mode and clear screen
MOV AH, 06                                     ;function for a scroll window
MOV AL, 07                                     ;no of lines = 7
MOV BH, 4ah                                    ;attribute byte
MOV CH, 0
MOV CL, 18
MOV DH, 6
MOV DL, 61
INT 10H                                         ;video window from (0, 18) to (6, 61)
                                              ;display values with cursor position

MOV DL, 20
MOV DH, 2
DISP 'A'
DISP 'N'
DISP 'U'
DISP 'R'
DISP 'A'
DISP 'G'
CSET

OVER: POP AX                                   ;popping AX
      JMP CS:OLD
NEWINT9 ENDP                                   ;end of ISR
                                              ;main program
START: MOV AH, 35H                             ;save old interrupt vector
      MOV AL, 09                               ;the type number = 9
      INT 21H                                 ;DOS interrupt for getting vector
      MOV WORD PTR OLD, BX                    ;save old IP
      MOV WORD PTR OLD+2, ES                  ;save old CS
      MOV AH, 25H                             ;install new vector
      MOV AL, 09                               ;at type number 9
      MOV DX, OFFSET NEWINT9                  ;point DX to the new ISR
      INT 21H
      MOV DX, OFFSET START                     ;count number of paragraphs

```

```
MOV CL, 4
SHR DX, CL
INC DX
MOV AX, 3100H           ;make the program resident
INT 21H
END
```

The advantage of TSR programs that work with a hot key is that, as the program is resident in memory, it can be invoked while we are doing something else. The disadvantage of TSR programs in DOS is that it gives users direct access to hardware and peripherals which may cause havoc.

8.10.3.1 | PC Scan Codes[†]

Table 8.7 | PC Scan Codes for 83 Key PC/XT Keyboard

Hex	Key	Hex	Key	Hex	Key	Hex	Key
01	Esc	15	Y and y	29	~ and '	3D	F3
02	! and 1	16	U and u	2A	LeftShift	3E	F4
03	@ and 2	17	I and i	2B	and \	3F	F5
04	# and 3	18	O and o	2C	Z and z	40	F6
05	\$ and 4	19	P and p	2D	X and x	41	F7
06	% and 5	1A	{ and [	2E	C and c	42	F8
07	^ and 6	1B	} and]	2F	V and v	43	F9
08	& and 7	1C	Enter	30	B and b	44	F10
09	* and 8	1D	Ctrl	31	N and n	45	NumLock
0A	(and 9	1E	A and a	32	M and m	46	ScrollLock
0B	) and 0	1F	S and s	33	< and ,	47	7 and Home
0C	_ and -	20	D and d	34	> and .	48	8 and UpArrow
0D	+ and =	21	F and f	35	? and /	49	9 and PgUp
0E	back space	22	G and g	36	RightShift	4A	-(gray)
0F	tab	23	H and h	37	PrtSc and *	4B	4 and LeftArrow
10	Q and q	24	J and j	38	Alt	4C	5 (keypad)
11	W and w	25	K and k	39	Spacebar	4D	6 and RightArrow
12	E and e	26	L and l	3A	CapsLock	4E	+ (gray)
13	R and r	27	: and ;	3B	F1	4F	1 and End
14	T and t	28	" and '	3C	F2	50	2 and DownArrow
						51	3 and PgDn
						52	0 and Ind
						53	. and Del

[†]Table 8.7–8.9: Reprint courtesy of International Business Machines Corporation, copyright © International Business Machines Corporation.

Table 8.8 | Combination Key Scan Codes

Hex	Keys	Hex	Keys	Hex	Keys	Hex	Keys
54	Shift F1	60	Ctrl F3	6C	Alt F5	78	Alt 1
55	Shift F2	61	Ctrl F4	6D	Alt F6	79	Alt 2
56	Shift F3	62	Ctrl F5	6E	Alt F7	7A	Alt 3
57	Shift F4	63	Ctrl F6	6F	Alt F8	7B	Alt 4
58	Shift F5	64	Ctrl F7	70	Alt F9	7C	Alt 5
59	Shift F6	65	Ctrl F8	71	Alt F10	7D	Alt 6
5A	Shift F7	66	Ctrl F9	72	Ctrl PrtSc	7E	Alt 7
5B	Shift F8	67	Ctrl F10	73	Ctrl LeftArrow	7F	Alt 8
5C	Shift F9	68	Alt F1	74	Ctrl RightArrow	80	Alt 9
5D	Shift F10	69	Alt F2	75	Ctrl End	81	Alt 10
5E	Ctrl F1	6A	Alt F3	76	Ctrl PgDn		
5F	Ctrl F2	6B	Alt F4	77	Ctrl Home		

Table 8.9 | Extended Keyboard Scan Codes

Hex	Keys	Hex	Keys	Hex	Keys	Hex	Keys
85	F11	8E	Ctrl -	97	Alt Home	A0	Alt DownArrow
86	F12	8F	Ctrl 5	98	Alt UpArrow	A1	Alt PgDn
87	Shift F11	90	Ctrl +	99	Alt PgUp	A2	Alt Insert
88	Shift F12	91	Ctrl DownArrow	9A		A3	Alt Delete
89	Ctrl F11	92	Ctrl Insert	9B	Alt LeftArrow	A4	Alt /
8A	Ctrl F12	93	Ctrl Delete	9C		A5	Alt Tab
8B	Alt F11	94	Ctrl Tab	9D	Alt RightArrow	A6	Alt Enter
8C	Alt F12	95	Ctrl /	9E			
8D	Ctrl UpArrow	96	Ctrl *	9F	Alt End		

KEY POINTS OF THIS CHAPTER

- Whenever the processor is performing a task, it can be interrupted by other programs or peripherals requesting service.
- The 8086 responds to an interrupt in a standard way, irrespective of the source of the interrupt.
- Associated with any interrupt is an interrupt service routine and an interrupt vector.
- The 8086 has a table with size 1K for storing all its 256 interrupt vectors.
- Certain interrupt vectors are reserved by Intel for specific purposes.
- Many interrupt types are used by BIOS and DOS.
- BIOS 10H functions are commonly used for display functions.
- It is also possible to address video memory directly.

- For text mode display, each character needs two bytes, one for the ASCII code and one for the attribute.
- BIOS 16H functions are used for accessing the keyboard.
- The hardware interrupt 09 is used by the PC keyboard.
- It is possible to hook interrupts using some standard DOS functions.
- TSR programs are pop up programs that are resident in memory.
- TSR programs are useful, but since hardware and interrupt vector tables are accessed directly, they can create havoc if not used with caution.

QUESTIONS

1. Why are the interrupt and trap flags cleared as part of the interrupt response?
2. How many hardware and software interrupts can the 8086 support?
3. Explain the terms 'interrupt service routine' and 'interrupt vector'.
4. The hardware interrupt INTR is called a non-vectorized interrupt. Why?
5. List the interrupts reserved by Intel.
6. Explain the sequence of actions that occur in an interrupt acknowledge cycle.
7. What happens if a number of interrupts occur at the same time?
8. What is meant by the 'attribute' byte with reference to video display of text?
9. List a few BIOS functions used for display activation.
10. List a few BIOS functions used for keyboard testing.
11. What does the word TSR mean?
12. Why cannot DOS interrupts be used in TSR programs?

EXERCISE

1. Find the address in the IVT for the interrupts of types 102H and 200H.
2. Find the attribute byte and character byte for the following cases.
 - a) Displaying '&' in a blue background and red foreground.
 - b) Displaying 'O' with blinking red background and yellow text.
3. Write a program to create display windows of three sizes and colors, at the same time.
4. Write a program to write a string 'WELCOME HOME' horizontally on a blue screen. Write the program using
 - a) BIOS functions,
 - b) accessing video memory directly.
5. Modify the above program such that it is activated by pressing a specific key only.
6. Use BIOS 16H function for the above program for checking for a specific key.
7. Write a program to print the names of three students on three rows with their roll numbers, on the left of the names. This should be printed on a small red window on the top of the video screen.
8. Write TSR programs for getting the above display with different hot key combinations.

9 PERIPHERAL INTERFACING - I



In this chapter, you will learn

- The necessity of using dedicated hardware for peripheral interfacing.
- The architecture and programming features of the chip 8255.
- The method of interfacing an ADC to the 8086.
- The methods of generating different waveforms using a DAC.
- The way in which an LCD can be used to display ASCII characters.
- The principle of the stepper motor and how it is controlled by the CPU.
- Hex keyboard interfacing to the 8086.
- The interfacing of dynamic multiplexed LEDs with the 8086.
- That in all the above cases, the 8255 is used as the parallel interface controller.

Introduction

In Chapter 6, we made a study of the hardware features of 8086. In Chapter 7, the methods by which memory and ports are connected to the 8086 have been clarified. One point that should be clear by now is that the processor is connected to memory and I/O devices through the address/data and control lines. Remember that the processor is just a computing device. It computes what is expected of it and transfers it to the outside world – in essence, the outside world is external to the processor – memory and I/O devices are external to the processor. Memory management for 8086 is done by using the read and write control signals, and I/O is also accessed using these signals. However, I/O devices are of different varieties, functions, features, specifications and so on. For example, a keyboard is very different from a video monitor. The point is that the 8086 does not have the requisite hardware within it to manage each of them according to their special requirements and specifications. Thus, it is clear that I/O devices need some extra hardware to let them be interfaced to the processor. It is in this context that various interfacing chips have been designed. All these chips are programmable in the sense that they can be made to work in different ways as required, by ‘programming’ them, using 8086 instructions. Thus, there are separate chips for managing parallel data transfer, serial data transfer, keyboard, interrupts and timers for example. We will now go into the details of interfacing various I/O devices to the processor using these special purpose interfacing chips.

9.1 | Trainer Kit

Remember that current PCs do not use the 8086 processor and therefore we cannot use the PC for a first level understanding of interfacing. However, the PC uses many interfacing chips in its 'chipset'. Now, we will try to understand the idea of 'interfacing' by the use of a trainer kit which has been designed for educational purposes. As the name indicates – it will have all the hardware and software components for 'training'. The trainer kit used here has the following specifications (see Fig 9.1).

Features

1. Intel 8086 CPU at 4.77MHz clock speed.
2. 16 KB for monitor EPROM upgradable to 64KB.
3. 16 KB RAM expandable to 64KB.
4. Standard RS232C compatible serial port brought out to a 9-pin D-type male connector.
5. 3 channel 16-bit counter/timer using 8253 terminating in an I/O pin connector (one channel used for baud rate generation).
6. Fully buffered address, data and control signals terminated at a 50-pin header for interfacing.
7. 8 nos. of interrupt lines brought out to a 10-pin connector using 8259 programmable interrupt controller.
8. Serial communication package to link PC and the kit

What is to be understood from these specifications is that the clock frequency of the processor used is 4.77MHz and that its data, address and control lines are buffered and available in a 50-pin connector. RAM and ROM are available on the board. The trainer also has a programmable interrupt controller (8259) and a timer chip 8253. It has an AT keyboard. The kit can be connected to the serial port of a PC. Thus, the most convenient way would be to write our program in the PC (using an editor), assemble and link it using MASM, and then download the resultant hex file to the RAM on the trainer board. Then, the processor on the trainer board will execute the program and give the results appropriately. The trainer also has what is called a 'monitor program' to manage the whole setup.

With this brief introduction, we will start our discussion on the various peripherals and the corresponding interfacing chips. The sequence followed will be to present a chip, its features, its programming and then use it for different applications.



Figure 9.1 | The trainer kit

9.2 | Programmable Peripheral Interface (PPI)-8255A

This chip is also called a parallel port chip and it eases out the problems and issues related to parallel data transfer. Parallel data transfer is what we have done all this time. We have transferred 8/16 bits at one go either to/from memory or I/O devices. The I/O devices we have discussed (Chapter 7) were quite simple and needed only a read/write control signal and a select pulse generated during the read/write cycle. However, when I/O devices are more versatile and have more features, a PPI will be very useful, especially when more than one I/O device is to be interfaced to the processor. When this chip is used, its functions are sufficient to ensure that normally, no other extra hardware is required to interface peripheral devices that perform parallel data transfer. The chip 8255A is a more advanced version of the original 8255, but since the original one is obsolete, we will simply refer to this chip as the parallel port chip 8255.

Using the 8255 Let us be clear about why this chip is necessary in the first place and how it can be used (see Fig 9.2). The data bus of the 8086 is connected to the 8255. That means data transfer between the two chips is possible. Consider first, the case of data from the 8086 being sent to the 8255. Where does the 8255 keep the data? The answer is that it has registers called ‘ports’ – there are three ports here – A, B and C, and these ports have pins connecting it to external devices. Thus, port A has 8 pins PA0 to PA7 – so also ports B and C, as well. To these port pins, external devices like keyboards, displays, printers can be connected. Figure 9.2 shows a set of LEDs connected to Port C and a printer to Port B. This means that the data from the 8086 can be transferred to the output devices, routed through these ports of 8255. To do that, we must be able to ‘program’ Ports B and C as ‘output ports’. Port A has a keyboard connected to it. This means that Port A is to be programmed as an input port – then the keyboard can send data to it, which can be received by the 8086 when it chooses to. Thus, with this simple setup, we see that the 8086 has been ‘interfaced’ to three I/O devices using the 8255 as a temporary storage space. Thus, the 8255 has three sets of ‘8-bit parallel ports’ which acts as intermediary between a processor and a number of I/O devices. Note that in all cases, the I/O devices used have only 8-bit data capability – so only the lower data lines D0-D7 are connected to the 8255.

However, it is a ‘programmable chip’ as are all the chips we will learn in the next two chapters as well. Keep in mind that these ‘interfacing chips’ do not have any ‘processing’ capability. They can only be made to act in the way we want to, by ‘programming’ it using the instructions of




Figure 9.2 | The connections between an 8086, 8255 and three peripherals

8086 (i.e., the processor to which it is connected). However, they have the hardware to do the special functions that are needed for each specific purpose. Programming interfacing chips entails writing of ‘control words’ in their control/command registers. We will learn this idea, starting with the 8255.

9.2.1 | Pin Configuration and Internal Block Diagram

Figure 9.3 shows the pin configuration of the DIP (dual-in-line) version of the chip. We see that it has 40 pins, which consist of three 8-bit ports named Port A (PA), Port B (PB) and Port C (PC), each of which can be programmed as input or output ports. Figure 9.4 shows the internal block diagram of the chip. It is seen to consist of various functional blocks, and let us take a quick look at each of the blocks.

Data Bus Buffer There is a three state bi-directional 8-bit buffer which is used to interface the chip to the data bus of the system. Upon execution of the processor’s input or output instructions, data and control/status words are received or transmitted by the buffer.

Read/Write Control Logic It manages all data transfer between the chip and the processor, on accepting control signals from the control and address buses of the system.

Group A and B Controls Functionally this chip has been divided as Group A and B and they have their corresponding controls. We will soon see what this grouping is intended for. The way the chip is to act is decided by a register called the control/status register which can be written




Figure 9.3 | Pin diagram of 8255




Figure 9.4 | Internal block diagram of the 8255

to and read from. Now see a typical connection between the 8086 and an 8255 in Fig 9.5. The lower data lines of the chip can be connected to 8 bits of the data bus of the 8086. The $\overline{RD}$ and $\overline{WR}$ are connected to the $\overline{IORD}$ and $\overline{IOWR}$ generated from the 8086. If this chip (which is viewed as an I/O port by the processor) is to have only an eight-bit address, only the lower 8 lines of the address lines of the processor need to be involved in the address decoding process, which causes the $\overline{CS}$ line to be activated. Otherwise, 16 address lines may be used. (Recollect that no I/O port can have an address size greater than 16 bits.) Any two address pins of the processor are to be connected to the pins A_0 and A_1 of the 8255. This is because there are four separate entities associated with this chip, and each one needs a unique address. This is achieved by the four possible combinations of these two lines, as shown in Table 9.1.




Figure 9.5 | Address decoding and connections between the 8086 and the 8255

Table 9.1 | Port Selection for the 8255 Using Pins A₀ and A₁

$\overline{CS}$	A ₁	A ₀	Selected entity
0	0	0	Port A
0	0	1	Port B
0	1	0	Port C
0	1	1	Control register

Example 9.1

Find the addresses of Port A, Port B, Port C and the control/status register of the 8255 whose address decoding circuitry is as shown in Fig 9.5.

Solution

The NAND gate output is 0, when the address lines A₇ to A₃ and A₀ have the following bit status, and that is when the chip gets selected/enabled.

A ₇	A ₆	A ₅	A ₄	A ₃	A ₀
1	1	0	0	0	0

When, A₂ and A₁ are 00, then A₇ ... A₀ is 1100 0000 i.e., C0H. This is the address of Port A. Similarly (refer Table 9.1) the other ports have the address as shown in Table 9.2.

Table 9.2 | Address of the Ports of 8255 for the Connection in Fig 9.5

Entity	Address (Hex)
Control register	C6
Port A	C0
Port B	C2
Port C	C4

9.2.2 | Programming the PPI

How is the chip programmed and what are the options available?

Programming the chip involves only the writing of a particular word to the control register. The control register is an 8-bit register which can be written into. The bits of this word (called the control word) will decide the way the ports of the chip are to be configured. To understand this, let us have a look at the control word format (refer Fig 9.6).

As seen in this and also in the block diagram, the ports A, B and C are grouped into two – Groups A and B. Group A consists of Port A and upper 4 bits (PC4–PC7) of Port C. Group B then obviously includes Port B and Port C lower (PC0–PC3). The division into groups give us only one piece of information – it is that Group A ports can have three operational modes (0, 1 and 2), but Group B ports have only two modes of operation.




Figure 9.6 | Control word format of 8255

Bit Assignment of the Control Word First look at D7 – it suggests the options of I/O mode or BSR mode. To begin with, let us use the chip in the I/O mode. This is the normal mode of operation (BSR mode is a special mode catering to Port C alone which we will deal with, later). The next thing to do in this, is to decide the mode of operation of the two groups. It is acceptable for the two groups to be in different or same modes. After this is done, select the option of a particular port being either an input or output port. For i/p, the corresponding bit is to be 1, and for output it is to be 0.

9.3 | Modes of Operation

The 8255 can be made to work in three modes as decided by the system designer. They are:

Mode 0: Basic input/output

Mode 1: Strobed input/output

Mode 2: Bi-directional bus

Mode 0 This is the simplest and most widely used mode. In this mode, the two 8-bit ports A and B, and the 4-bit ports Port C upper and Port C lower, may be used independently. Here data is simply taken in from an input port or given to an output port.

Mode 1 This is the ‘handshaking’ mode. Handshake implies data transfer in which the communicating devices exchange request and acknowledge control signals with each other. For an 8255 operating in this mode, Ports A and B pins are used for data transfer, while 4 bits of each of Port C are used for generating the handshaking signals for each of the 8-bit ports.

Mode 2 This is the bidirectional mode. Only Group A can use this mode. Here Port A is used for transmitting as well as receiving data. Handshaking signals generated by the upper 4 bits of Port C maintain bus discipline for proper flow of data in the required direction. Now, let us use the different modes.

9.4 | Mode 0

The specifications of this mode are:

- i) Two 8-bit ports and two 4-bit ports.
- ii) Any port can be input or output.
- iii) Outputs are latched (ref Section 7.4.1 for the reason for this).
- iv) Inputs are not latched.
- v) 16 different input/output combinations possible.

Example 9.2

Design the control word to configure the ports of an 8255 chip in mode 0, with Port B and Port C upper (PC_U) as inputs and Port A and Port C lower (PC_L) as outputs.

Solution

Refer Fig 9.6.

Since we are using the I/O mode (rather than the BSR mode), $D_7 = 1$

Both groups are to be in mode 0. Hence, $D_6 D_5 = 00$ and $D_2 = 0$

For i/p, the corresponding bit is to be 1, and for output, it is to be 0

Since Port A is to be an output port, $D_4 = 0$

Since Port C_U is to be an input port, $D_3 = 1$

Since Port B is to be an input port, $D_1 = 1$

Since Port C_L is to be an output port, $D_0 = 0$

D_7	D_6	D_5	D_4	D_3	D_2	D_1	D_0
1	0	0	0	1	0	1	0

The control word is thus 8AH.

Example 9.3

It is necessary to read the setting of 12 switches which are connected to the 8255 and display it on LEDs. Draw the setup and write the program for the same.

Solution

We can use the 8255 with the addresses as in Table 9.2 and the control word as in Example 9.2.

Figure 9.7 shows how eight switches are connected to Port B and Port C_U . When a switch is closed, a ‘0’ is read in, and when it is open, a ‘1’ is read into the port lines. To the pins of Port A and Port C_L , connect 12 LEDs along with their current limiting resistors. The LEDs glow when the corresponding pins have a high logic on them.




Figure 9.7 | Connection of switches and LEDs to an 8255

Now, let us discuss the steps in the program that will be run:

- i) The ports and the control register addresses have been given labels and will be specified using the EQU directive. Since the port addresses are only 8 bits long, fixed port addressing can be used.
- ii) The switch settings at the input Port B can be read in directly and transferred to the output Port A.
- iii) The data at Port C_U can be read in by inputting Port C fully, but the switch settings will be available only at the upper 4 bits position and have to be shifted to the position of the lower 4 bits to output it to the LEDs.

The program is as follows:

```

CR EQU 0C6H           ;address of the control register
PA EQU 0C0H           ;address of Port A
PB EQU 0C2H           ;address of Port B
PC EQU 0C4H           ;address of Port C

MOV AL, 8AH           ;move control word to AL
OUT CR, AL            ;send it to the control register
IN AL, PB             ;get switch status from i/p port B
OUT PA, AL            ;output it to the LEDs at o/p port A
IN AL, PC             ;take in switch settings of Port C
MOV CL, 04            ;CL = 4
ROR AL, CL            ;rotate AL 4 times to the right
OUT PC, AL            ;the data is in the lower part of PC
                     ;output it to the LEDs there
END                   ;end of the assembly file

```

Note In the programs on I/O interfacing, we will not use BIOS or DOS interrupts, because they are specific to PCs only. We are using trainer kit to run these programs. We will not specify any memory models also – that can be decided by the type of software supported by the trainer. Only the program lines will be written here.

9.4.1 | Interfacing 16 Bits I/O Ports to the 8255

All the above is for a port to which only 8-bit data is to be transferred from D₀ to D₇ of the data bus of the processor. Since our processor (8086) has a data bus width of 16 bits, it is quite possible that it may sometimes need to send/receive 16 bits through a parallel port when the port has a data bus width of 16 bits. This will necessitate the need for two chips of the PPI (see Fig 9.8).

Recollect the discussion on memory banks and I/O banks. 16-bit data is to be arranged into two 8-bit banks, an upper bank and a lower bank. When data is to be accessed through the upper data lines, the processor makes arrangements to lower the $\overline{BHE}$ line. This is used to decode the address of the upper bank, as shown in the figure. Thus, to transfer a 16-bit data, both the PPIs will get enabled. With the same decoding logic as in Example 9.2, the addresses of the ports of each of the PPIs may be enumerated and is left for the interested reader. All the I/O devices we will discuss further are ones with only 8-bit data bus, and we will use Fig 9.5 as our reference.




Figure 9.8 | 16-bit port interfacing using two PPIs

Example 9.4

Generate a square wave of 33% duty cycle at the lowest bit of Port A. The clock frequency of the processor is 4.77 MHz

Solution

A 33% duty cycle means that the square wave should be high for one third of the cycle. Let us consider the ON time to be 1 msec. Then the OFF period will be 2 msecs.

We can solve the problem this way – create a delay of 1 msecs and call the delay procedure once during the ON time and twice for the OFF time.

For delay calculations, refer to Example 6.2. The delay loop is:

```

MOV CX, N
HERE:  NOP
      LOOP HERE

```

Since the clock frequency is 4.77 MHz, one clock period = $0.21 \mu\text{secs}$

The delay is to be 1msec

Total delay time = 1 msec = $20 N \times 0.21 \mu\text{secs}$

For 1 msec delay, the value of $N = [1 \text{ msec} / (20 \times 0.21 \mu\text{secs})] = 238$ or EEH

The complete program is as shown below. Note that we need the square wave only on the lowest pin of Port A, but it is available at all pins. This is because the pins of Port A are not 'bit addressable' i.e., we cannot manipulate Port A bits, one at a time. For this case, it is okay to have the square wave generated from all the pins – but we need to use only one pin to observe the square wave. Connect a CRO probe to one pin and observe the waveform on the CRO. Also, for this application, Port A should be made an output port. All other ports are 'don't care'. The control word is written in mode 0, and bit $D_4 = 0$ specifically i.e., 1000 0000 is the control word.

```

                PA EQU 0C0H           ;address of Port A
                CR EQU 0C6H           ;address of control register

                MOV AL, 80H            ;control word for Port A to be output
                OUT CR, AL             ;send it to control register
BACK:           MOV AL, 0              ;AL = 0
                OUT PA, AL             ;send it to Port A
                CALL DELAY             ;call a delay for one msec
                CALL DELAY             ;call a delay for one msec

                MOV AL, 0FFH           ;AL = FFH
                OUT PA, AL             ;send it to Port A
                CALL DELAY             ;call a delay of 1 msec
                JMP BACK               ;repeat

DELAY          PROC NEAR              ;delay program
                MOV CX, 0EEH           ;N = EEH
HERE:           NOP
                LOOP HERE
                RET
                ENDP
                END

```

9.4.2 | Bit Set Reset Mode

It will be more useful to discuss the BSR mode before we go on to the discussion of modes 1 and 2. This is a special mode and is applicable only for the bits of Port C. In the control word format, if the MSB is made 0 ($D_7 = 0$) the bit set/reset (BSR) mode takes effect. In this mode, any bit of Port C can be set or reset by specifying the bit which has to be set or reset. However, at a time, only one bit can be addressed – and that bit is to be either set or reset. The corresponding control word has to be decided, **and moved to the control register**. This is because, even though we are manipulating the bits of Port C, it is a 'control word' that is being written. See the format of the BSR control word in Fig 9.9. D_3 to D_1 specifies the bit to be set/reset.

Example 9.5

Write the BSR control words for the following cases:

- i) PC_0 to be set
- ii) PC_7 to be reset
- iii) PC_1 to be set




Figure 9.9 | Format of the BSR control word

Solution

	D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀	
i) PC ₀ to be set	0	0	0	0	0	0	0	1	= 01
ii) PC ₇ to be reset	0	0	0	0	1	1	1	0	= 0EH
iii) PC ₁ to be set	0	0	0	0	0	0	1	1	= 03

Example 9.6

Write a program to do the following. Get data from input Port A. If the data is above 7FH, send a high value on PC₇. Otherwise, make PC₇ to be zero. Use the port addresses in Example 9.1.

Solution

The first step in solving this problem is to configure Port A to be an input port and Port C_{upper} to be an output port (since a signal has to generated on it), in mode 0.

The point is that, though we might be using a BSR mode, we must first ‘configure’ the ports as input/output.

The control word format for this is 1 0 0 1 0 0 0 0 = 90H

The BSR control word for resetting PC₇ is 0 0 0 0 1 1 1 0 = 0EH

The BSR control word for setting PC₇ is 0 0 0 0 1 1 1 1 = 0FH

The program is

```
CR EQU 0C6H      ;address of control register
PA EQU 0C0H      ;address of Port A
PC EQU 0C4H      ;address of Port C
```

```

        MOV AL, 90H           ;control word
        OUT CR, AL           ;send it to the control register
        IN AL, PA            ;take in data through Port A
        CMP AL, 7FH          ;compare it with 7FH
        JA LOW               ;if AL > 7FH, go to LOW
        MOV AL, 0EH          ;AL < 7FH, write a word for PC7 = 0
        OUT CR, AL           ;send it to the control register
        JMP XIT              ;go to exit
LOW:    MOV AL, 0FH          ;AL > 7F, BSR word for PC7 = 1 in AL
        OUT CR, AL           ;send it to the control register
XIT:
        END                  ;end the program

```

Thus, we see that with the BSR mode, we are able to do 'bit addressing' for port pins, but only for Port C, in the output mode alone.

9.5 | Mode 1

This is also called the handshaking mode. The features of this mode are that Ports A and B can be used for data input or output, and the bits of Port C are used as handshaking signals to control these data transfers. This mode is frequently referred to as strobed input/output mode also. The word 'strobe' is indicative of signaling – classically, by the use of light.

Recollect the grouping into groups A and B.

- i) Each group contains one 8-bit port and one 4-bit port. The 4-bit port is used for 'handshaking'.
- ii) The 8-bit port can be either an input or an output.
- iii) Input and output are latched.
- iv) The 4-bit port is used for control and status of the 8-bit data port.
- v) Interrupt logic is supported.

To understand it completely, we need to separate the cases of input and output.

9.5.1 | Strobed Input Mode

Figure 9.10 shows the signal definitions of Port C when Ports A and/or B are used as inputs. The definitions of the handshaking signals are as given:

- i) **$\overline{STB}$ (Strobe):** The peripheral sends a data byte to the 8255 and indicates this by lowering the strobe line. This active low signal is given by an input device indicating that it has sent a byte of data to the 8255, which will be latched therein.
- ii) **$\overline{IBF}$ (Input Buffer Full):** This is the signal the 8255 sends as acknowledgement to indicate that the input latch has received the byte sent by the peripheral. This is reset when the 8086 reads the data.
- iii) **$\overline{INTR}$ (Interrupt Request):** A high on this is used to interrupt the CPU. This is set to high if $\overline{STB}$ is high, $\overline{IBF}$ is high and the INTE (interrupt enable) FF is set. It is reset by the falling edge of $\overline{RD}$ as shown in Fig 9.11. For enabling the 'INTE' flip flop, use PC₄ (for INTE_A) and PC₂ (INTE_B). The setting/resetting of these FFs is done by the BSR mode.

Note PC_4 and PC_2 are used as strobe signals also, but remember that they are used in the BSR mode (for $INTE$). There is no effect of the BSR control word on the activities in Mode 1. Once the interrupt enable flip flops are set/reset, these pins are used as strobe signal pins.

iv) In Mode 1 input mode, two pins PC_6 and PC_7 are not used for handshaking. They may be used as input or output lines as desired.




Figure 9.10 | Mode 1: Strobed input – Pin configuration and control word format




Figure 9.11 | Timing waveforms for strobed input mode

9.5.2 | Mode 1: Strobed Output

In a similar way, Mode 1 output uses the pins of Port C as shown in Fig 9.12. The timing diagram for data transfer is as in Fig 9.13. Let us discuss the definitions of the handshaking signals here.

- i) **$\overline{OBF}$ (Output Buffer Full):** When the processor writes a data byte to the output latch of the 8255, this signal goes low. It indicates that this data can be read into the peripheral. It goes high when the peripheral takes the data and acknowledges it.
- ii) **$\overline{ACK}$:** This is the signal given by a peripheral when it accepts the data from 8255A. A 'low' on this input informs the 8255 that the data from the port has been accepted. In essence, a response from the peripheral device indicating that it has received the data output by the CPU.
- iii) **INTR (Interrupt Request):** This signal goes high after the acknowledge signal is removed. This can be used to prompt the processor to deliver the next data byte into the output latch of the 8255. The setting of this signal requires that $\overline{OBF}$, $\overline{ACK}$ and INTE (interrupt enable Flip flop) are all high.
- iv) **INTE (Interrupt Enable):** This is an internal flip which can be set/reset by the BSR mode of port C on PC_6 and PC_2 respectively. This is an internal flip flop used to enable or disable the generation of the INTR signal. $INTE_A$ is controlled by the set/reset of PC_6 . $INTE_B$ is controlled by the set/reset of PC_2 .




Figure 9.12 | Mode 1: Strobed output – Pin configuration and control word format

- v) In Mode 1 output mode, two pins PC₄ and PC₅ are not used as handshaking signals. They may be used as input or output lines as desired.

Status word Sometimes, it may be necessary to know the current status of the handshaking signals, and take decisions within the program accordingly. Since the handshaking signals are on Port C lines, it will be sufficient to read in the status of Port C, which will have the same signal definitions as the port pins, and the ‘status word’ is shown in Fig 9.14. This status word is read by using the instruction IN AL, PORT C. Thus, we get the status of Port C pins in AL, which can be used to take decisions. It is also possible to use the PPI in such a way in Mode 1, that one port is in strobed output and the other is in strobed input, mode. The two possibilities are shown here in Figs 9.15 and 9.16. Now we will examine a case in which Port A is in strobed input mode connected to a keyboard and Port B is connected to a printer, in the strobed output mode (as in Fig 9.15).




Figure 9.13 | Timing waveform for strobed output mode




Figure 9.14 | Status of Port C for input and output configurations of Mode 1




Figure 9.15 | Port A in strobed input and Port B in strobed output mode

Example 9.7

- i) Write the control word for setting Port A as an input port and Port B as an output port.
- ii) Port C upper and Port C lower are used as control signals.
- iii) Write instructions to find the status of the signal IBF_A and $\overline{OBF}_B$.

Solution

Refer to Fig 9.4 or Fig 9.15 for designing the control word.
The control word is thus

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
1	0	1	1	0	1	0	0

i.e., B4H

- iv) In this application, Group A pins corresponds to an input configuration, and Group B pins corresponds to an output. Hence the pins of Port C will have the status as designated (refer Fig 9.14).

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
I/O	I/O	IBF _A	INTE _A	INTR _A	INTE _B	$\overline{OBF}_B$	INTR _B

Thus the status of Port C will be as shown.
What is needed, is to read in Port C and test the condition of D₅ for knowing the status of IBF_A.
Test the condition of D₁ for the status of $\overline{OBF}_B$

```
IN AL, PORTC      ;read in the contents of Port C
TEST AL, 5        ;for testing the condition of D5, IBFA
TEST AL, 1        ;for testing the condition of D1,  $\overline{OBF}_B$ 
```




Figure 9.16 | Port A in strobed output and Port B in strobed input mode

Example 9.8

Consider a case in which the 8255 is used with Port A in the strobed input and Port B in the strobed output mode. The chip is connected to a keyboard as well as a printer as shown in Fig 9.17. Write instructions to get data from the keyboard and send it to the printer to be printed.

Solution

The scenario is that when a key is pressed, the corresponding ASCII data is transferred to the 8086 using Mode 1 handshaking, and it generates an interrupt. This interrupt calls the print ISR which is responsible for sending this data to the printer. The problem has two parts.

- i) Port A is used as the input port to get data from the keyboard. The keyboard sends a strobe $\overline{STB}_A$ signal to the 8255 on its PC_4 pin, along with sending the ASCII code corresponding to the key pressed. This is read into the microprocessor only after it is confirmed that the signal IBF_A (Input buffer full) has been sent by the 8255 on Pin PC_5 . The status of all the pins of Port C can be read by 'reading in' Port C. Once PC_5 is confirmed to be high, the data that has been latched into Port A of the 8255 can be read into the processor. Now, see the timing diagram of Fig 9.11. When $\overline{STB}_A$ signal goes high and IBF_A is high, $INTR_A(PC_3)$ is raised. To ensure this, the corresponding $INTE_A$ flip flop must be set. This is done by setting bit PC_4 , using the BSR facility. In Fig 9.17, the $INTR_A$ signal from pin PC_3 of 8255 is connected to the $INTR$ pin of the 8086. When the 8086 is interrupted, it runs the ISR 'PRINT' which sends the character that was read in from the keyboard.
- ii) In the second part, the PRINT ISR (Interrupt Service Routine) is executed. It does the following:

Port C status is read to confirm that OBF_B is high. The processor considers the printer to be ready to accept data only when this signal is high. Only then will data be outputted to Port B. The status of this signal is obtained in D_1 bit of Port C. If it is high, the data which was brought from the keyboard and placed in memory is transferred to Port B using an OUT instruction. When this is done, the $\overline{OBF}_B$ signal will go low and the printer should respond to it with a $\overline{ACK}_B$ signal which will cause $\overline{OBF}_B$ to go high again. (This will allow the system to send the next byte to be printed, when interrupted again.) The generation of the signals $\overline{OBF}_B$ will be




Figure 9.17 | Connecting a keyboard and a printer to the 8086 in Mode 1

done automatically by the 8255 in Mode 1 output, once the $\overline{WR}$ goes low corresponding to a write operation by the processor.

Steps in the Program

- i) Write the control word in Mode 1, with Port A as input and Port B as output.
- ii) Set the interrupt enable FF ($INTE_A$) by setting bit PC_4 .
- iii) Set the interrupt flag of the 8086 processor.
- iv) Read in the status of Port C.
- v) Test the status of $\overline{IBF}_A$.
- vi) When it is high, transfer data byte from Port A to AL of the processor.
- vii) Transfer the byte to a memory location.
- viii) Since the interrupt enable flag of the 8255 has been set, the $INTR_A$ will go high, and since the Interrupt flag of the 8086 is set, the 8086 gets interrupted and gets directed to the PRINT ISR.

Steps in the Print ISR

- i) Read in the contents of Port C.
- ii) Test the $\overline{OBF}_B$ pin.
- iii) If it is high, send the data to the printer.

```

CR EQU 0C6H           ;address of the control register
PA EQU 0C0H           ;address of Port A
PB EQU 0C2H           ;address of Port B
PC EQU 0C4H           ;address of Port C

LOC DB ?              ;space to store key data

        MOV AL, 0B4H   ;mode 1, Port A input, Port B output
        OUT CR, AL     ;send it to control register
        MOV AL, 09     ;BSR word for setting  $PC_4$ 
        OUT CR, AL     ;set bit  $PC_4$  for setting  $INTE_A$ 
        STI            ;set the interrupt flag of 8086
        IN AL, PC       ;read in Port C
REPEA:  TEST AL, 20H    ;test bit  $D_{20H}$  ( $IBF_A$ )
        JZ REPEA       ;if it is 0, keep testing
        IN AL, PA       ;if  $IBF_A$  is high, send Port A data to AL
        MOV LOC, AL    ;move it to memory

        PRINT ISR

        IN AL, PC       ;read in Port C
REAP:   TEST AL, 2       ;test bit  $D_1$  ( $OBF_B$ )
        JZ REAP        ;if low, test bit till high
        OUT PB, AL
        IRET

```

9.6 | Mode 2 (Strobed Bidirectional Bus I/O)

This functional configuration provides a means for communicating with a peripheral device or structure on a single 8-bit bus for both transmitting and receiving data (bidirectional bus I/O). 'Handshaking' signals are provided to maintain proper bus flow discipline in a similar manner to Mode 1. Interrupt generation and enable/disable functions are also available.

Mode 2 Basic Functional Definitions:

- i) Used in Group A only.
- ii) One 8-bit, bi-directional bus port (Port A) and a 5-bit control port (Port C).
- iii) Both inputs and outputs are latched.
- iv) The 5-bit control port (Port C) is used for control and status for the 8-bit, bi-directional bus port (Port A).

Since this is used only for very special applications, detailed discussion of this mode will not be attempted here. For more details of this mode, the data sheet of the chip may be referred.

9.7 | Centronics Printer Interface

Now, let us discuss a very popular application of 8255 in Mode 1. The Centronics printer interface is one such application. The printer parallel port has been in use for many years, but is getting to be more or less obsolete these days – this is because the USB port is now used in most applications replacing the parallel port as well as the conventional serial port. However, the parallel port is still maintained as a legacy port in many computers. Let us have a look at this application which has been in use for many, many years.

In the 1970s, a printer company named Centronics developed an electrical interface to its printer which soon became very popular. By 1981, when IBM wanted to interface its PC to a printer, they decided to modify the connector on the PC side. The end result was that the printer interface consisted of a 25-pin connector on the PC side and a 36-pin Centronics connector on the printer side. This is also called the Epson FX-100 standard. It might seem quite weird to find 36 pins devoted to sending data at just one byte at a time, though with a few handshaking signals as well. However, in actuality, most of the control pins and all the data pins have individual ground return pins, rather than a common ground. This has been done to reduce the effect of interfering electrical noise.

Table 9.3 shows the pin designations of the 36 pins of the Centronics interface as well as the corresponding 25 pins (DB-25) to which these 36 pins are connected to. Obviously, many of the 36 pins will be designated as 'no connection' (NC). Now, let us discuss how the PC and the printer communicate such that data bytes get printed one at a time. See, the Timing diagram Fig 9.18, of the data transfer between the PC and the printer. The steps involved in sending one byte of data to the printer are:

- i) The PC checks if the BUSY signal of the printer is asserted. If it is found low, it means the printer is not busy i.e., it is willing to accept a data byte for printing.
- ii) The PC then sends out one-byte of data on the data bus.
- iii) The PC also asserts the $\overline{STROBE}$ signal (makes it low) for at least 0.5 μ secs.
- iv) The timing diagram also shows that the data should remain at the data pins at least 0.5 μ secs after the $\overline{STROBE}$ is made high.
- v) Once the $\overline{STROBE}$ goes low, the BUSY signal goes high to indicate that it is busy in the process of printing this byte.

- vi) When the printing of this byte is done with, the printer sends an acknowledge signal ($\overline{ACK}$ to the PC, indicating that it is ready to accept a new byte for printing). The PC can use either the $\overline{ACK}$ or BUSY signal to verify if the next byte can be sent for printing.
- vii) The ‘minimum’ timing requirements are to be met.

Various other status and control signals are seen for communications between a PC and a printer. Their functions are listed in Table 9.3. The printer interface will not be discussed in more

Table 9.3 | List of Centronics and DB-25 Pins $\overline{SLCTIN}$

Centronics Pins	DB-25 Pin	Name	Direction	Description
1	1	$\overline{STROBE}$	PC-to-Printer	Strobe
2	2	D0	PC-to-Printer	Data Bit 0
3	3	D1	PC-to-Printer	Data Bit 1
4	4	D2	PC-to-Printer	Data Bit 2
5	5	D3	PC-to-Printer	Data Bit 3
6	6	D4	PC-to-Printer	Data Bit 4
7	7	D5	PC-to-Printer	Data Bit 5
8	8	D6	PC-to-Printer	Data Bit 6
9	9	D7	PC-to-Printer	Data Bit 7
10	10	$\overline{ACK}$	Printer-to-PC	Acknowledge
11	11	BUSY	Printer-to-PC	Busy
12	12	POUT	Printer-to-PC	Paper out
13	13	SEL	Printer-to-PC	Select
14	14	$\overline{AUTOFEED}$	PC-to-Printer	Autofeed
15	N/A	n/c	N/A	Not used
16	N/C	0 V	N/A	Logic ground
17	N/C	CHASSIS GND	N/A	Shield ground
18	N/C	+5 V PULLUP	Printer-to-PC	+5 V DC (50 mA max)
19–30	18–25	GND	N/A	Ground reference for signal pins 1–12, in most cables as twisted pairs
31	16	$\overline{RESET}$	PC-to-Printer	Reset
32	15	$\overline{FAULT}$	Printer-to-PC	Fault (Low when offline)
33	N/C	0 V	N/A	Signal ground
34	N/C	N/C	N/A	Not used
35	N/C	+5 V	Printer-to-PC	+5 V DC
36	17	$\overline{SLCTIN}$	PC-to-Printer	Select in (Taking low or high sets printer on line or off line respectively)




Figure 9.18 | Timing diagram of the Centronics interface

detail because, as mentioned earlier, it is more or less obsolete. However, before we conclude, a few tit-bits from the history of parallel ports are presented here.

Note N/A stands for not available and N/C for no connection.

9.7.1 | History

The word LPT is sometimes used to designate the parallel port of the PC. The word stands for 'Line Print Terminal', which was designed for line printers meant for printing text data. Because there is no written standard, the timing relationships between the handshaking signals vary widely among printers from different manufacturers. As we have just seen, the parallel port is unidirectional, because data was 'sent' for printing. However, printing is not the only function that a parallel port can do. A parallel port with unidirectional capabilities is called the 'standard parallel port'. In this standard, there are a few registers involved in data transfer, and specific timing too. The SPP has only a few differences from the original Centronics port.

The bidirectional port The IBM PS/2 (Personal System/2) computer enhanced the standard PC parallel interface by adding bidirectional drivers to the eight data lines. The I/O connector and signal assignments remained the same. A parallel port with bidirectional drivers is often referred to as an Extended Parallel Port (EPP). EPP was developed to provide for high-speed, bidirectional data transfers that are compatible with the register map of the existing standard parallel port. The EPP specification assigns traditional microprocessor bus signaling functions to the standard parallel port lines (address strobe, data strobe) to access adapter hardware directly.

Extended Capabilities Port (ECP) ECP is an extension to the bidirectional standard parallel port and was developed by Microsoft and Hewlett Packard in 1992. The specification defines automatic hardware handshaking, command and data cycles, and DMA transfers. The handshaking signals for data transfers have the same timing relationships as defined for standard parallel ports. In 1994, the IEEE 1284 standard was released. It included the two specifications for parallel port devices, EPP and ECP. In order for them to work, both the operating system and the device must support the required specification. This is seldom a problem today since most computers support SPP, ECP and EPP and will detect which mode needs to be used, depending on the attached device. If you need to manually select a mode, you can do so through the BIOS on most computers.

9.8 | Interfacing an Analog to Digital Converter to the 8086

Real world applications use sensors to get a signal voltage. For example, for measuring a temperature, a thermocouple may be the sensor used. This gives an analog voltage corresponding to the temperature value. Various other sensors measure quantities like humidity and speed, and give the measured values as analog voltages. To use these values in a computer, these voltages have to be converted to digital numbers. This is the function of an analog to digital converter. There are various methods employed in this conversion, which you might have learned in your basic digital circuitry course. Each method has its merits and demerits. There are various standard ICs available as ADCs. They vary in the number of digital output bits, the number of analog inputs (multiplexed) and speed of conversion. These parameters will have to be verified from the data sheet of the IC before choosing a particular ADC for an application.

Here our interest will be to interface an ADC to the 8086 using the 8255 as an interface. Also, ADCs may be ‘parallel’ or ‘serial’. This is an important application for the parallel port IC 8255. When an analog voltage is given as an input to an ADC, it gets converted to a digital number which is transferred to the 8086. The digital value can be stored in the RAM of the system and may be displayed or used in further computations. See the block diagram of such a setup in Fig 9.19 with the 8255 acting as an intermediary in this. What is the role of the 8255 chip in the conversion and data transfer? The answer is that the converted digital data from the ADC is transferred to the 8255 with one of its ports acting as an input. This data is sent to the processor through the data bus. Moreover, the ADC needs a few control signals for its operation. This is provided by selectively using the other port pins of the 8255 PPI.

ADC 0808/0809 There are many different methods of analog to digital conversion, and many different standard ADC chips. We will choose the ADC0808/ADC0809 which is 8-bit parallel ADC and microprocessor compatible. Its functional pin diagram is shown in Fig 9.20. It is designated as an ‘8-Bit μ P Compatible A/D Converter with 8-Channel Multiplexer’. It uses the successive approximation technique for analog to digital conversion.

Its key specifications are given as:

- | | |
|---------------------------|---|
| 1. Resolution | 8 Bits |
| 2. Total Unadjusted Error | + / ₋ ½ LSB and + / ₋ 1 LSB |
| 3. Single Supply | 5 VDC |
| 4. Low Power | 15 mW |
| 5. Conversion Time | 100 μ secs |




Figure 9.19 | General block diagram of the connection between an ADC, PPI and 8086




Figure 9.20 | Functional pin diagram of the ADC 0808/0809

Table 9.4 | Channel Selection Logic

Selected analog channel	C	B	A
IN0	0	0	0
IN1	0	0	1
IN2	0	1	0
IN3	0	1	1
IN4	1	0	0
IN5	1	0	1
IN6	1	1	0
IN7	1	1	1

It is an 8-input multiplexed ADC – which means that it has 8 input analog signal lines, though only one of them can be operational at a time. This is selected by three address inputs A, B, C. Table 9.4 shows the address bit configuration for selecting specific input channels. For example, if IN0 is to act as the input, the address lines C, B and A all have to be low; for IN1, the values of CBA is to be 001 and so on. The first requirement in using the ADC is to select an input channel by giving the appropriate logic on the address pins. To latch this on to the chip, a signal called ALE (Address Latch Enable) is to be supplied on the ALE pin. ALE is to be a low to high transition (ref the timing diagram in Fig 9.21). After the address is latched and the analog input is available at the selected input line, the ADC must be signaled to start conversion. This is a low to high pulse of minimum specified duration (as mentioned in the data sheet). The ADC requires a clock and the speed of conversion depends on the clock rate. The maximum clock frequency will be specified in the data sheet. The clock of the processor can be divided to get the right frequency for the ADC.




Figure 9.21 | Timing diagram for the ADC 0808/0809

The ADC takes a finite time to complete the conversion and then it notifies this fact by lowering the pin called EOC (End of Conversion). This should be brought to the notice of the processor. The EOC signal can be used to interrupt the microprocessor, and allow the converted data to be transferred to the 8086 or this signal can be polled continuously. To receive the digital data, the output lines of the ADC are to be activated. This is done by making high the line OE (Output Enable). Once, the output lines are activated, the converted digital data, can be transferred to the 8086. The above is the sequence of actions necessary to use the ADC chip 0809 to perform analog to digital conversion and then to transfer the digital data to the micro-processor.

Let us now use the pins of 8255 for the purposes specified above. Make the connections between the 8086, 8255 and the ADC as shown in Fig 9.22. The salient points regarding the connection are

- i) Port A is used in the input mode to get the converted digital data from the ADC to 8255.
- ii) The port pins PB_0 to PB_2 are used in the output mode as the address selection pins A, B, C of the ADC
- iii) The port pin PC_0 is used as ALE it is to be an output pin.
- iv) The port pin PC_1 is used to give the start conversion (SC) pulse to the ADC. Hence it is to be an output pin.
- v) The port pin PC_7 is used in the input mode, to receive the End of Conversion (EOC) signal from the ADC.
- vi) Port pin PC_2 is used as OE for the ADC. It is defined as an output pin.




Figure 9.22 | ADC with its control and data pins connected to an 8086

Example 9.9

Write a program to interface the ADC 0809/0808 to the 8086 using 8255 as an interfacing chip.

Solution

Refer to Fig 9.22 for the connection diagram. Also refer to the timing diagram (Fig 9.21) to understand the following steps:

- First the control word in Mode 0 is to be finalized. Port A is to be input, Port B output, Port C_L output and Port C_U input. Hence, the control word is 1001 1000 i.e., 98H.
- The ALE signal has to be pulsed i.e., it should have a low to high transition and for that pin PC₀ can be used in the BSR mode. The BSR control word for clearing and setting PC₀ are 0000 0000 i.e., 00 and 0000 0001 i.e., 01.
- To start conversion, the SC signal, which uses pin PC₁ should be made high and then low. The corresponding BSR control words are 02 and 03.
- After a delay, both ALE and SC are to be made low. The minimum delay necessary can be checked from the data sheet.
- Then the EOC signal is tested. This is done by rotating Port C left and getting PC₇ into the carry bit. When the carry bit is found to be low, the OE signal on pin PC₂ is raised. This enables the output pins of the ADC. The BSR control words for raising and lowering the OE pin are 07 and 06 respectively.

```

DAT DB ?                ;memory location for converted data
CR EQU 0C6H             ;address of the control register
PA EQU 0C0H             ;address of Port A
PB EQU 0C2H             ;address of Port B

```

```

PC EQU 0C4H      ;address of Port C
MOV AL, 98H      ;control word in AL
OUT CR, AL       ;send it to control register
MOV AL, 00       ;address corresponding to IN0
OUT PB, AL       ;send it to Port B
MOV AL, 00       ;BSR control word to clear ALE (PC0)
OUT CR, AL       ;send it to the control register
MOV AL, 01       ;BSR control word for setting ALE (PC0)
OUT CR, AL
MOV AL, 03       ;BSR control word to set SC(PC1)
OUT CR, AL
CALL DELAY       ;call a delay
MOV AL, 00       ;BSR control word for clearing ALE(PC0)
OUT CR, AL
MOV AL, 02       ;BSR control word for clearing SC(PC1)
OUT CR, AL

AGN:  IN AL, PC    ;read in the contents of Port C
      RLC         ;rotate left to check if PC7 is high
      JC AGN      ;if carry high, continue the checking
      MOV AL, 07   ;if low, EOC got. Load word for OE
      OUT CR, AL
      IN AL, PB    ;read in converted data from PB
      MOV DAT, AL  ;save it in memory
      MOV AL, 06   ;BSR control word to lower OE(PC2)
      OUT CR, AL   ;send it to control register

      DELAY PROC NEAR ;delay procedure
          MOV CX, N   ;the value of N can be decided by

      NXT:  LOOP NXT ;referring the data sheet for
          RET        ;required delay
      DELAY ENDP
          END

```

9.9 | Interfacing to a Digital to Analog Converter

Converting a digital number is an important application in electronics – you must have learned the basic methods of this conversion. Here, we will use a DAC chip, and connect it to the 8086 using an 8255 as the interface. A digital number sent from the processor is converted to an analog current/voltage by this arrangement. DACs of different resolutions are available – those with 8, 10, 12 or 16-bit inputs – the more the number of bits, the better the resolution. This is understandable because an 8-bit DAC gives 2^8 or 256 levels of voltage at the output – similarly the others give 2^{10} , 2^{12} and 2^{16} discrete current levels respectively. Adding these current values give the resultant current which is converted to a voltage by an I to V converter (using

an operational amplifier). For an 8-bit DAC, the output current I_o is a function of the binary numbers at the data input pins D_0 to D_7 .

$$I_o = I_{ref} \times \left(\frac{D_7}{2} + \frac{D_6}{4} + \frac{D_5}{8} + \frac{D_4}{16} + \frac{D_3}{32} + \frac{D_2}{64} + \frac{D_1}{128} + \frac{D_0}{256} \right) \quad (9.1)$$

where I_{ref} is the reference input current corresponding to the V_{ref} (pin no:14) applied to the chip. Figure 9.23 shows the pin diagram of the chip DAC0800/DAC0802 which is an 8-bit high speed digital to analog converter with a settling time of 100ns. We will now use the 8-bit data from the 8255 as the input to the DAC. The minimum connections are as shown in the Fig 9.24. Additional components like Zener diodes may be used for protection purposes. The




Figure 9.23 | Pin diagram of the DAC 0800




Figure 9.24 | Connections between the 8255 and DAC

output of the DAC is a current which is converted to a voltage by the opamp at the output. Since the opamp is connected in a simple inverting configuration, only positive output voltages can be obtained here. However, if the opamp is used in a difference configuration, both positive and negative values may be obtained.

Example 9.10

For the following binary numbers applied to a DAC, calculate the output analog voltage obtained given that $I_{\text{ref}} = 2 \text{ mA}$ and $R = 5.6 \text{ K}$

- i) 11000011
- ii) 00010111

Solution

Equation 9.1 can be used to find the value of the analog voltage for each case.

$$I_{\text{ref}} = 2 \text{ mA}$$

$$\begin{aligned} \text{i) For } 11000011, I_o &= I_{\text{ref}} (1/2 + 1/4 + 0 + 0 + 0 + 0 + 1/128 + 1/256) \\ &= 2 \times (195/256) \\ &= 1.52 \text{ mA} \end{aligned}$$

$$V_o = 5.6 \times 1.52 = 8.52 \text{ V}$$

$$\begin{aligned} \text{ii) For } 00010111, I_o &= I_{\text{ref}} (0 + 0 + 0 + 1/16 + 0 + 1/64 + 1/128 + 1/256) \\ &= 2 (23/256) \\ &= 0.1796875 \text{ mA} \\ V_o &= 5.6 \times 0.1796 \\ &= 1.006 \text{ V} \end{aligned}$$

Next, we will generate waveforms of various shapes using a DAC.

Example 9.11

Generate a triangular waveform using the DAC 0800.

Solution

The method is to increment the data outputted to the DAC, from 0 to FFH. On conversion to analog form, it will give a gradually increasing output from 0 to the maximum output voltage. Then decrement the number from FFH to 0. This will give the high to low triangular transition.

Note As the clock frequency of the 8086 is reasonably high, the time for incrementing the number 0 to FFH and decrementing from FFH to 0 will be quite small. Thus, the waveform generated will have a very short duration. To obtain a slower increase and decrease in analog voltage, a good idea will be to have a small delay after each increment/decrement of the digital number. The delay must be small, such that steps are not seen in the increasing/decreasing portions of the analog voltage. The amount of this small delay can also be used to control the periods of the generated waveforms. This has been done in this program.

```
CR EQU 0C6H          ;address of the control register
PA EQU 0C0H          ;address of Port A

MOV AL, 80H           ;control word for Port A as output
OUT CR, AL            ;send word to the control register
```

```

                MOV AL, 0                ;AL = 0
REPEA:          OUT PA, AL              ;send AL to port A
                CALL DELAY              ;call a delay
                INC AL                  ;increment AL
                CMP AL, 0FFH            ;compare AL with FFH
                JNZ REPEA               ;if not equal, increment
AGN:            OUT PA, AL              ;otherwise decrement
                CALL DELAY              ;call a delay
                DEC AL                  ;decrement AL
                CMP AL, 0               ;compare AL with 0
                JNZ AGN                 ;if not 0, continue decrementing
                JMP REPEA               ;if AL = 0, increment

                DELAY PROC NEAR          ;delay procedure
                MOV CX, 03FH            ;value of N for a small delay
HOW:            LOOP HOW
                RET
                DELAY ENDP

```

Example 9.12

Generate a staircase waveform with 5 steps.

Solution

A staircase waveform has steps in the increasing portion of the analog voltage. Once the voltage reaches the maximum value, it falls to zero in zero time. Here, we are to generate a staircase waveform with 5 steps. After each step, a delay is called. The method is as follows.

Divide the total range of 0 to 255 into 5 parts. Each increment can correspond to 51. The digital number range is 0, 51, 102, 153, 204 and 255. The program is as follows, and the output waveform is in Fig 9.25.

```

                CR EQU 0C6H             ;address of the control register
                PA EQU 0C0H             ;address of Port A

                MOV AL, 80H             ;control word for Port A as output
                OUT CR, AL              ;send word to control register
STRT:           MOV AL, 0               ;AL = 0
                OUT PA, AL              ;send AL to Port A
                CALL DELAY              ;call a delay
REPEA:          ADD AL, 51              ;add 51 to AL
                OUT PA, AL              ;send AL to Port A
                CALL DELAY              ;call a delay
                CMP AL, 255             ;check whether AL = 255
                JNZ REPEA               ;if not, repeat the addition of 51
                JMP STRT               ;if AL = 255, go to START

```




Figure 9.25 | Staircase waveform generated

```

    DELAY PROC NEAR                ;delay procedure
        MOV CX, 200                ;value of N for 'Delay' in Fig 9.25
AGN:    LOOP AGN
        RET
DELAY    ENDP                    ;end of procedure
        END

```

Example 9.13

Generate a saw tooth waveform using a DAC.

Solution

For a saw-tooth waveform, the content of AL is incremented from 0 to FFH. Thus, the analog voltage varies slowly from 0 to the maximum value.

To enable a 'slow' variation, a delay is used every time an increment in AL is done. Once it reaches the maximum, it is made to go to 0, and then the whole procedure is repeated.

```

        CR EQU 0C6H                ;address of the control register
        PA EQU 0C0H                ;address of Port A

        MOV AL, 80H                ;control word for Port A as output
        OUT CR, AL                 ;send word to control register
AGN:    MOV AL, 0                   ;AL = 0
START:  OUT PA, AL                 ;send AL to Port A
        CALL DELAY                 ;call a delay
        INC AL                     ;increment AL
        CMP AL, 0FFH               ;compare AL with FFH
        JNZ START                  ;if AL ≠ FFH, go to start
        JMP AGN                    ;if AL = 0, go to AGN to start

        DELAY PROC NEAR            ;delay procedure

```

```
MOV CX, 03FH
HOW: LOOP HOW
RET
DELAY ENDP
```

9.10 | Interfacing Liquid Crystal Displays to the 8086

Liquid crystal displays called LCDs are very popular with their qualities of low power dissipation and ease of use. The only problem normally encountered is the problem of the viewing angle. The display is not equally clear at all viewing angles. Now, LCD modules of many different specifications (mostly differing in the number of lines, number of characters per line and so on) are available. An LCD module has registers writing into which, the display can be easily programmed and controlled. Here, we will use a 16 × 2 character LCD which looks as shown in Fig 9.26.

9.10.1 | Pins of the LCD

The LCD that we have selected has 16 pins as shown in Table 9.5. We see that DB0 to DB7 correspond to the data pins. The others are the pins for control signals and the power supply. VEE is a pin used to adjust the contrast of the display. It is usually connected to a potentiometer, so that contrast can be adjusted. Besides that, there are the VCC and ground pins. There are two pins for backlight adjustment, if necessary. Backlighting means extra lighting behind the LCD panel (usually LEDs) so that the display is visible in the dark also.

RS Register Select This pin selects between a command register and a data register. RS = 0 corresponds to the command register and RS = 1 to the data register. The data to be displayed is to be sent to the data register.

R/W: Read/Write This pin allows the user to write to or read from the display. When there is no necessity for reading the display, this pin is grounded. If both reading and writing is required, this pin is made programmable.




Figure 9.26 | A 16 × 2 LCD module

Table 9.5 | Pins of the 16 × 2 LCD Module

No	Symbol	Function
1	V _{ss}	Power supply ground (0 V)
2	V _{cc}	Power supply (5V)
3	VEE	Power supply for adjusting contrast
4	RS	Register select signal
5	R/W	Read write select signal
6	E	Enable signal
7	DB0	Data bus line
8	DB1	Data bus line
9	DB2	Data bus line
10	DB3	Data bus line
11	DB4	Data bus line
12	DB5	Data bus line
13	DB6	Data bus line
14	DB7	Data bus line
15	AV _{EE}	Positive voltage for back light
16	K	0 V for back light

E: Enable The Enable pin has to be given a high to low pulse, which is maintained high for at least 450 ns (may be different for other LCD modules).

DB0 to DB7 These are the data pins of the LCD. Data to be written is to be sent through these pins, and data to be read will be received from the LCD through these pins. The data to be written for display are sent as ASCII characters. For writing into the command registers, there are predefined codes for the LCD. The codes for the LCD we are using, are given in Table 9.6.

Busy Flag It is seen that there is a minimum time required to latch one data on the LCD and get it displayed. Suppose we want to give a new data for display, the simplest way would be to introduce a small delay between sending the two display data (which can be given only one character at a time). However, another method for sending consecutive characters is to check what is called the 'Busy' flag of the LCD. For testing the busy flag, make RS = 0 first. The 'Busy' flag is DB7 and can be read when R/W = 1 and RS = 0. If DB7 is found high, it means that the LCD is busy doing its operations, and will not accept any new information. Keep checking this flag until it is low. Then, the next data can be written to it.

Note If the busy flag is to be read, the R/W pin has to be made programmable.

Now, let us do some display activities using a 16 × 2 LCD. Data and commands are sent from the 8086, using the 8255 as the interface.

Backlight: There is a lamp here instead of reflected light. If backlighting is provided by LEDs as in the case of many 16 × 2 LCDs, connect pin 16 to ground, and pin 15 to V_{cc} through a 100Ω resistor.

The connections are to be done as follows:

- V_{SS} and R/W is connected to ground.
- V_{CC} is connected to 5 V supply.
- V_{EE} is connected through a 10 K pot to the supply for contrast adjustment.
- RS is connected to PC0 and E is connected to PC1.
- Pin 7-14 (DB0 to DB7) of the LCD module are connected to port A.
- Pins 15 and 16 of the LCD are used for backlight adjustment (not shown in Fig 9.27).

Table 9.6 | Command Codes of the LCD

Code (Hex)	Command
01	Clear display screen
02	Return home
04	Shift cursor left (decrement cursor)
05	Shift display right
06	Shift cursor right (increment cursor)
07	Shift display left
08	Display off, cursor off
0A	Display on, cursor on
0C	Display on, cursor off
0E	Display on, cursor blinking
0F	Display off, cursor blinking
10	Shift cursor position to left
14	Shift cursor position to right
18	Shift the entire display to the left
1C	Shift the entire display to the right
80*	Force cursor to the beginning of the first line
C0*	Force cursor to the beginning of the second line
38	2 lines and 5 x 7 matrix

* For a 16 x 2 line display, the addresses of the cursor positions are 80 to 8F for the first line, and C0 to CF for the second line




Figure 9.27 | The LCD connected to the 8255

Example 9.14

This program displays 'YA' on the LCD. The generation of control signals necessary for sending the commands to the command register are put in a procedure named 'COMMAND', and similarly there is a procedure named 'DAT' which takes care of the control signals necessary for sending to the data register, the characters which are to be displayed.

Note that first the LCD is initialized, then cleared, and then the cursor is positioned. This is done by sending command words to the LCD command register. For latching these words to the command register, RS has to be made low, and E has to be given a high to low pulse. These are done in the COMMAND procedure.

Similarly, for displaying a character, the ASCII value of the data is to be sent to the data register and RS should be made '1'. Also a high to low transition should occur at E. These issues are taken care of in the DAT procedure.

```

PA EQU 0C0H           ;address of Port A
CR EQU 0C6H           ;address of the control register

MOV AL, 80H           ;control Word for port A as output
MOV CR, AL            ;send it to the control register
MOV AL, 38H           ;initialize LCD, 2 lines, 5x7 matrix
CALL COMMAND          ;call the COMMAMD procedure
CALL DELAY            ;call a delay
MOV AL, 0EH           ;command for display on, cursor on
CALL COMMAND          ;call the COMMAMD procedure
CALL DELAY            ;call a delay
MOV AL, 01            ;command to clear LCD
CALL COMMAND          ;call the COMMAMD procedure
CALL DELAY            ;call a delay
MOV AL, 80H           ;cursor at line 1, positon 1

CALL COMMAND          ;call the COMMAMD procedure
CALL DELAY            ;call a delay
STRT: MOV AL, 'Y'      ;move to AL, the ASCII of 'Y'
CALL DAT              ;call the DAT procedure
CALL DELAY            ;call a delay
MOV AL, 'A'           ;move to AL, the ASCII of 'A'
CALL DAT              ;call the DAT procedure
CALL DELAY            ;call a delay

COMMAND PROC NEAR     ;procedure named COMMAND
OUT PA, AL            ;send data in AL to port A
MOV AL, 0             ;BSR word to make PC0 (RS) low
OUT CR, AL            ;RS = 0 for command register
MOV AL, 03            ;BSR word to make PC1 (E) high
OUT CR, AL            ;this is to make E = 1
CALL DELAY            ;call a delay
MOV AL, 02            ;BSR word to make E (PC1) low
OUT CR, AL            ;this is to make E = 0
RET

```

```

COMMAND ENDP                                ;end the procedure

DAT     PROC NEAR                            ;procedure named DAT
OUT PA, AL                                  ;send data in AL to port A
MOV AL, 01                                  ;BSR word to make PC0(RS) high
OUT CR, AL                                  ;RS = 1 is to select data register
MOV AL, 03                                  ;BSR word to make PC1(E) high
OUT CR, AL                                  ;this is to make E = 1
CALL DELAY                                  ;call a delay
MOV AL, 02                                  ;BSR word to make E(PC1) low
OUT CR, AL                                  ;this is to make E = 0
RET
DAT     ENDP

DELAY PROC NEAR                              ;delay procedure
MOV CX, 0FFFH
AGN: LOOP AGN
RET
DELAY ENDP

```

Example 9.15

Write a program which rotates the word HELLO from right to left, displayed on the LCD.

Solution

Refer Table 9.5. For rotating the display to the right, the command '18H' is used.

For rotating the display in the opposite direction, use the code 1CH.

The program is the same as Example 9.14, except that the characters to be displayed are written as an array in memory and accessed one by one. Also, one more command code is added, that is for rotation of the display. Since everything else is the same, only the additional lines of the program, are shown here.

```

MSG DB 'H', 'E', 'L', 'L', 'O'
-----
-----
      ----
      ---

REEP:  MOV AL, 18H                            ;command to shift display left
        CALL COMMAND
        CALL DELAY

        LEA BX, MSG                          ;point to the characters
        MOV CX, 5                            ;CX = number of characters
REPEA: MOV AL, [BX]                          ;move to AL the character
        CALL DAT
        CALL DELAY_ROLL
        INC BX

```

```

      LOOP REPEA      ;repeat this until CX = 0
      JMP REEP        ;infinite loop for continuous rotation
-----
-----
      ----
      ----

```

Note The DELAY_ROLL procedure is used for deciding the rate at which rotation occurs. For slow rotation, this delay can be made high.

9.11 | Interfacing a Stepper Motor to the 8086

Introduction to Stepper Motors A stepper motor is an electromechanical device which converts electrical pulses into discrete mechanical movements. When electrical pulses are applied to it, the shaft of the motor rotates in steps and this type of movement gives the motor its name.

Principle of Operation Stepper motors operate differently from normal DC motors. A DC motor rotates continuously when voltage is applied to its terminals. Stepper motors, have multiple ‘toothed’ electromagnets arranged around a central gear-shaped piece of iron (see Fig.9.28). The electromagnets are energized by an external control circuit which sends pulses to the motor.




Figure 9.28 | Cross-section of a variable reluctance (VR) motor

To turn the motor shaft, one of the electromagnets is given power first, which makes the gear's teeth magnetically attracted to the electromagnet's teeth. When one tooth of the gear is thus aligned to the energized electromagnet (Electromagnet 'B' and tooth '6' are aligned in Fig 9.28), others are slightly offset from the corresponding electromagnets. When the next electromagnet is turned on and the first is turned off, the gear rotates slightly to align with the next one, and from there the process is repeated. Each of those slight rotations is called a **step**, with an integral number of steps making a full rotation. In this way, the motor can be turned by a precise angle.

The rotation of the motor is related to the sequence of the input pulses:

- i) The order in which a particular sequence is applied, decides the direction of rotation (clock-wise or anti-clockwise).
- ii) The speed of stepping depends on the frequency of the pulses applied i.e., higher the frequency, faster the stepping motion.

We can use stepper motors for movement which needs to be finely controlled. The fine control is obtained because these motors move in steps, and the steps can be quite small in size. For example, one step can be 2 degrees and, for one complete (360 degree) rotation, 180 steps are obviously needed. In order to obtain 90 degree rotation for such a motor, we must write a program to supply only 45 pulses to it. This can be used to advantage when it is needed to control aspects such as rotation angle, speed, position and synchronism. As such, they are used in applications such as printers, plotters, high end office equipment, hard disk drives, medical equipment, fax machines and automotive and industrial applications where precise and controlled rotation is required.

9.11.1 | Stepper Motor Types

There are three basic stepper motor types. They are

Variable reluctance

Permanent magnet

Hybrid

Variable Reluctance This type of motor consists of a soft iron multi-toothed rotor and a wound stator. When the stator windings are energized with DC current, the poles become magnetized. Rotation occurs when the rotor teeth are attracted to the energized stator poles (Fig 9.28).

Permanent Magnet Such motors have permanent magnets added to the motor structure. The rotor has no teeth, but is made with alternating north and south poles of permanent magnets, situated in a straight line parallel to the rotor shaft (see Fig 9.29). Applying current to each phase in sequence will cause the rotor to rotate by adjusting to the changing magnetic fields. Although it operates at fairly low speed, the PM motor has a relatively high torque characteristic, as the magnetic flux intensity due to the permanent magnet is much higher.

Hybrid The hybrid stepper motor combines the best features of both the PM and VR type stepper motors. The rotor is multi-toothed like the VR motor and contains an axially magnetized concentric magnet around its shaft (see Fig 9.30). The teeth on the rotor provide an even better path which helps guide the magnetic flux to preferred locations in the air gap.

The permanent magnet and the hybrid types are the most popular types with the latter being more expensive, but capable of being designed for greater resolutions i.e., smaller step sizes.




Figure 9.29 | A permanent magnet stepper motor




Figure 9.30 | Cross-section of a hybrid stepper motor

9.11.2 | Two-Phase Stepper Motors

There are two basic winding arrangements for the electromagnetic coils in a two-phase stepper motor: bipolar and unipolar.

Unipolar Motors A unipolar stepper motor has logically two windings per phase, one for each direction of magnetic field. In this arrangement, since a magnetic pole can be reversed without

switching the direction of current, the commutation circuit can be made very simple for each winding. Typically, given a phase, one end of each winding is made common: giving three leads per phase and six leads for a typical two phase motor. Often, these two phase commons are internally joined, so the motor has only five leads.

Bipolar Motors Bipolar motors have logically a single winding per phase, reversed in order to reverse a magnetic pole, so the driving circuit is more complicated. There are two leads per phase, none are common. A unipolar motor can be converted to a bipolar one (by discarding the common) but not vice versa.

Universal Motors A universal stepper motor has 8 leads, while the unipolar has six and the bipolar has four. Usually we use a universal motor, using only its four leads to which we feed in our electronic pulses. In our discussions, we will consider the application of a sequence of pulses to the four leads of such an arrangement. We will use a simple, 90 degree PM motor with two phases. Applying current to each phase in sequence will cause the rotor to rotate by adjusting to the changing magnetic fields, and thus rotor alignment is done by pulsing.

9.11.3 | Driving a Stepper Motor

9.11.3.1 | Full Step Drive (two phases on)

This is the usual method for full step driving of the motor. Both phases are always on. The motor will have full rated torque. This is achieved by the sequence of ones and zeros as shown in Table 9.7 which is to be repeatedly applied. Reversing the order in which the sequence is applied gives anti-clockwise rotation. In short, for clock wise rotation, the sequence to be applied repeatedly is 09, 0CH, 06, 03 ... For anti-clock wise rotation, it is 03, 06, 0CH, 09 ...




Figure 9.31 | a Unipolar motor

b Bipolar motor

9.11.3.2 | Wave drive

In this drive method only a single phase is activated at a time. It has the same number of steps as the full step drive, but the motor will have significantly less than rated torque. This sequence is 8, 4, 2, 1 for clockwise and 1, 2, 4, 8 for anti-clockwise rotation.

9.11.3.3 | Half stepping

When half stepping, the drive alternates between two phases on and a single phase on. This increases the angular resolution, but the motor also has less torque at the half step position (where only a single phase is on). The advantage of half stepping is that the drive electronics need not change to support it. The step-angle is half that of the previous two cases – thus the stepping resolution is increased. For anticlockwise rotation, the order of the above sequence should be reversed.

Table 9.7 | Driving Sequence for a Stepper Motor-Full Step Drive

Step No.	Clockwise			
	A	B	$\bar{A}$	$\bar{B}$
1	1	0	0	1
2	1	1	0	0
3	0	1	1	0
4	0	0	1	1

Table 9.8 | Driving Sequence for a Stepper Motor-Wave Drive

Step No.	Clockwise			
	A	B	$\bar{A}$	$\bar{B}$
1	1	0	0	0
2	0	1	0	0
3	0	0	1	0
4	0	0	0	1

Table 9.9 | Driving Sequence for a Stepper Motor-Half Stepping

Step No.	Clockwise			
	A	B	$\bar{A}$	$\bar{B}$
1	1	0	0	1
2	1	0	0	0
3	1	1	0	0
4	0	1	0	0
5	0	1	1	0
6	0	0	1	0
7	0	0	1	1
8	0	0	0	1

9.11.4 | Using the 8255 to Interface a Stepper Motor

Since we are trying to run the stepper using a sequence generated by the 8086, we also use the 8255 to send the sequence through one of the ports. The motor, however, cannot be driven directly by the parallel port, because the motor requires a current much more than can be supplied by the 8255. (The exact current requirement depends on the specifications of the particular motor being used). As such, current drivers are needed between the 8255 port lines and the leads of the motor. Transistors with high current capability (e.g. Darlington pair or power transistors) can be used. Besides this, there are special motor driving ICs available. One such IC is the ULN 2003 driving IC whose pin diagram is shown in Fig 9.32. This IC contains an array of seven Darlington pair transistors.

Example 9.16

Write a program to rotate a stepper motor in the clockwise direction using the full step drive scheme.

Solution

Let us use the sequence shown in Table 9.7, for rotation in the clockwise direction. The connection is as shown in Fig 9.33. Only the lower four pins of port B need to be used. The sequence is obtained by rotating right, the data 0110 0110. On rotating this, the lower four bits are 06, 03, 09, and 0CH this repeats to get a clockwise rotation of the motor.

```

                                CR EQU 0C6H           ;address of the control register
                                PB EQU 0C2H           ;address of Port B

                                MOV AL, 80H           ;control word for Port B as output
                                OUT CR, AL            ;send it to the control register
                                MOV AL, 66H           ;load sequence in AL
BACK:                          OUT PB, AL            ;send it to Port B
                                ROR AL, 1            ;rotate right once
                                CALL DELAY            ;call a delay
                                JMP BACK              ;repeat the sequence

                                DELAY PROC NEAR        ;delay procedure
                                MOV CX, 0FFFFH
AGN:                          LOOP AGN
                                RET
                                DELAY ENDP

```

A delay is given between applying the sequence to the coils of the motor. The delay is needed for the coils to be energized to their new excitation sequence – also, **less the delay; more is the speed of the motor. Thus, the speed of rotation is varied according to the frequency of the pulse stream.**

9.11.5 | Step Angle

Now, think of a motor with a step-angle of 1.8° . Suppose a rotation of just 90° is required. How many steps are required? Obviously,

No. of steps = $90 / 1.8 = 50$. This motor takes 200 steps to cover 360° .




Figure 9.32 | Functional pin diagram of the driver IC ULN2003




Figure 9.33 | Connection diagram of a stepper motor using the ULN 2003 IC

Example 9.17

Write a program to rotate a motor by 90° in the clockwise direction, and then rotate it 180° in the anticlockwise direction. The step-angle is 1.8° .

Solution

For rotating 90° in the clockwise direction, 50 steps are needed.

For rotating 180° in the anticlockwise direction, 100 steps are needed, and the sequence has to be given in the reverse order. So the data 66H is shifted left in this case.

```
CR EQU 0C6H      ;address of control register
PB EQU 0C2H      ;address of Port B
```

```

        MOV AL, 80H           ;control word for Port B as output
        OUT CR, AL           ;send it to control register
        MOV CX, 50           ;CX = 50
        MOV AL, 66H          ;load sequence in AL
BACK:   OUT PB, AL           ;send it to Port B
        ROR AL, 1            ;rotate sequence once, to right
        CALL DELAY          ;call a delay
        LOOP BACK           ;repeat until CX = 0

        MOV CX, 100          ;CX = 100
        MOV AL, 66H          ;load sequence in AL
BACK1:  OUT PB, AL           ;send it to AL
        ROL AL, 1            ;rotate it once, to the left
        CALL DELAY          ;call a delay
        LOOP BACK1          ;repeat until CX = 0

        DELAY PROC NEAR      ;delay procedure
        MOV BX, 0FFFH
        AGN: DEC BX
        JNZ AGN
        RET
        DELAY ENDP

```

Example 9.18

Write a program to rotate a stepper motor in the wave drive 4-step sequence in the clockwise direction.

Solution

```

        CR EQU 0C6H          ;address of control register
        PB EQU 0C2H          ;address of Port B

        SEQ DB 8, 4, 2, 1    ;store the sequence
        MOV AL, 80H          ;port B to be an output port
        OUT CR, AL           ;send it to control register
RAP:    MOV CX, 4             ;CX = 4
        LEA BX, SEQ          ;point BX to memory location SEQ
START:  MOV AL, [BX]          ;load the sequence values to AL
        OUT PB, AL           ;send it to Port B
        CALL DELAY          ;call a delay
        INC BX               ;increment the pointer value
        LOOP START           ;repeat until CX = 0
        JMP RAP              ;start all over again

        DELAY PROC NEAR      ;delay procedure
        MOV DX, 0FFFH
        AGN: DEC DX

```

```
JNZ AGN
RET
DELAY ENDP
```

9.11.6 | Other Issues Regarding Stepper Motors

An important issue to take care of when using stepper motors is that there is a chance of back emf being produced during the energization of the coils. This can damage the circuits producing the sequence and hence diodes are connected which block these spikes. Figure 9.34 shows the diodes connected in a circuit which uses Darlington pair transistors for producing the high current required to drive the motors. Note that the diodes are not shown in Fig 9.33 where the motor driving IC ULN 2003 is used. This is because such a diode is in-built within the IC. See Fig 9.34 where the drive for the stepper motor is given by high current Darlington pair transistors TIP 120. Diodes IN4001 are used to suppress the back emf spikes.




Figure 9.34 | Current driver Darlington transistors and back EMF suppressing diodes in a stepper circuit

9.11.7 | Optocouplers

An optocoupler is another device helpful in damping the back EMF produced in a motor circuit. An optocoupler or optical isolator is a device that uses a short optical transmission path to transfer a signal between elements of a circuit, while keeping them electrically isolated. Typically, they come in a small 6-pin or 8-pin IC package, but are essentially a combination of two distinct devices: an optical transmitter, typically a gallium arsenide LED (light-emitting diode) and an optical receiver such as a phototransistor or light-triggered diac. The two are separated by a transparent barrier which blocks any electrical current flow between the two, but does allow the passage of light. The basic idea is shown in Fig. 9.35. Usually the electrical connections to the LED section are brought out to the pins on one side of the package and those for the phototransistor or diac to the other side, to physically separate them as much as possible. Figure 9.36 shows the pin diagram of a quad optocoupler ILQ74. Figure 9.37 shows this optocoupler used in a stepper motor interfacing circuit, along with a driving IC ULN 2003.




Figure 9.35 | Principle of operation of an optocoupler




Figure 9.36 | An optocoupler IC




Figure 9.37 | Optocoupler used in a stepper motor driving circuit

9.12 | Hex Keyboard Interfacing

In Chapter 8, we discussed the ideas of interfacing a keyboard to the PC. It was seen that specialized microcontrollers are used to handle all the issues associated with keyboard interfacing. As such, the main processor of the PC is free of the task of identifying the key pressed. However, still it is important to understand the mechanism of how a keypress is detected and how the pressed key is identified. For this purpose we use a hexadecimal keyboard, which has 16 keys with the characters 0 to F. This keyboard is interfaced to the 8086 through the port lines of the 8255. The program identifies the key that has been pressed and displays it.

9.12.1 | Hex Keyboard

See the diagram of the hex keyboard. This is a 4×4 matrix keyboard connected to the 8255. For our purpose, we will use two ports of the 8255, Port A as an output port connected to the row lines, and Port B as an input port connected to the column lines. Only the lower 4 bits of the two ports are needed. See Fig 9.38 for the connection, with the key positioned at the interconnection of a row and a column. Only when a key at a junction is pressed, a path is established between the corresponding row and column.

9.12.2 | Detecting a Keypress

Remember that the rows are connected to an output port and the columns to an input port. Now, send zeros from the four row outputs, and read in the contents of the columns. If no key has been pressed, there is no connection between a row and column and hence only '1's are received from all the columns. We get 1111 from PB3 to PB0. This should be clear from Fig 9.38. If a key is pressed, the '0' from the row will pass through the closed switch to the column and will be read in from Port B. Say the key '9' is pressed. Then a '0' is received from that column i.e., the PB3 to PB0 will read as 1101. An information we get from this input port reading is that a key has been pressed. So we say that a 'keypress has been detected'.




Figure 9.38 | Hex keyboard connected to an 8255

9.12.3 | Identifying the Key

The next task is that of identifying the key that has been pressed. For the case of key ‘9’ being pressed, looking at the number received from Port B makes it clear that a key in column 1 is the pressed key. However, we have no clue to the identity of this key, since we do not know the row in which the key resides. Thus, identifying the row of this key is the solution to this problem. How do we go about this?

The technique is to ground the rows one by one and read in the columns, starting with Row 0. If the data read from the columns is 1111, it is obvious that this particular row does not contain the key that was pressed – if a number less than 1111 is received, the row gets identified, since we know which row is grounded now. Using this information along with the column information identifies the key. If the key ‘9’ was pressed, we get an input reading from Port B of 1101 only when the row grounded is Row 2. Thus, we identify that the pressed key is the one at the junction of the second row and first column, which is ‘9’. If a display mechanism is available, we can arrange it to display ‘9’.

Note The rows and columns are numbered from 0 to 3

9.12.4 | Key Debounce

The key that we are using has mechanical contacts and is susceptible to what is called ‘key jitter or bounce’. That is, when a key is pressed, the key moves back and forth between the contacts for a certain time before it settles. To ensure that the key settles before we read the key, we allow a certain time for it – and this is called the key debounce time. It may range from 10 ms


```

ROW0:    MOV AL, 0FEH          ;send '0' to Row0 alone
          OUT PA, AL
          IN AL, PB            ;read in the columns
          AND AL, 0FH          ;mask the upper nibble
          CMP AL, 0FH          ;compare with 0FH
          JZ ROW1              ;no keypress in Row0, check Row1
          LEA SI, ROW_0        ;not equal, SI to point to Row0 data
          JMP COL_ID           ;go to finding the column

ROW1:    MOV AL, 0FDH          ;send '0' to Row1 alone
          OUT PA, AL          ;repeat steps
          IN AL, PB
          AND AL, 0FH
          CMP AL, 0FH
          JZ ROW2
          LEA SI, ROW_1
          JMP COL_ID

ROW2:    MOV AL, 0FBH          ;send '0' to Row2 alone and
          OUT PA, AL          ;repeat steps
          IN AL, PB
          AND AL, 0FH
          CMP AL, 0FH
          JZ ROW3
          LEA SI, ROW_2
          JMP COL_ID

ROW3:    MOV AL, 0F7H          ;send '0' to Row3 alone and
          OUT PA, AL          ;repeat steps
          IN AL, PB
          AND AL, 0FH
          CMP AL, 0FH
          JZ KEYP
          LEA SI, ROW_3

COL_ID:   SHR AL, 1            ;to identify column, shift right
          JNC FOUND            ;if no carry, column is found
          INC SI                ;otherwise point SI to next row
          JMP COL_ID           ;go back to finding the column

FOUND:    MOV AL, [SI]         ;get data pointed by SI into AL
          MOV KEY, AL          ;this is the pressed key. Store in KEY

          DELAY PROC NEAR
          MOV CX, 6000         ;delay for key debounce

AGN:      LOOP AGN
          RET

DELAY    ENDP
          END

```

9.13 | Interfacing LED Displays

We know that LEDs are very popular display devices. They are easy to use and give very bright and pleasing display which can be viewed equally well from any viewing angle, unlike LCDs. The only drawback with LED displays is the high amount of current they need, unlike LCDs which need very low power. Figure 9.39 shows that a single LED can be connected to a positive power supply as shown. The value of the current limiting resistor depends on the current rating of the LED. Suppose we need a number of LEDs for displays, we usually use only one power source for all of them. In that case, they are connected together either as ‘common anode’ or ‘common cathode’ LED displays. In Fig 9.40b the cathodes of the three LEDs are connected together. If we want to light up only the first and third LED i.e., apply a ‘1’ only at A1 and A3. In Fig 9.40a, which is ‘common anode’, K1 and K3 alone should be grounded, for the same result.

9.13.1 | Seven Segment LED

However, more important applications of LEDs are as alphanumeric displays. In that case, seven segment LEDs are used, in which seven LEDs are arranged as the segments of a display arranged in a particular shape which when selectively lighted up, give the display of alphanumeric characters (see Fig 9.41). By lighting up all the segments, we get ‘8’ displayed. We can have one more segment in this display and it is for the decimal point. In Fig 9.41, there are eight segments, including the segment for the decimal point. In spite of this, we still designate such displays as ‘seven’ segment displays. Such LED modules also can be used in the common anode or common cathode, configuration. To light up a seven segment display LED of the common cathode type, ensure that the cathode is grounded, and give a ‘1’ to the segments which are to



Figure 9.39 | A single LED




Figure 9.40 | a Common anode LEDs

b Common cathode LEDs




Figure 9.41 | A seven segment LED display unit

be lighted. The first information (grounding the cathode) corresponds to the 'digit' data and the second is the segment data. By selectively lighting up segments, we can have all alphanumeric characters displayed.

9.13.2 | Static Seven Segment Displays

Now, suppose we want to use such a module to display a single character. Let us think of a scenario wherein, the number to be displayed is sent from an 8086 through an 8255, to the display module. We just send the code (called seven segment code) corresponding to the segments of the LED. This code gives the information as to which of the segments are to be lighted up for the display of a specific character.

Suppose we use only a one digit display module. If it is a common cathode type, we ground the common cathode and then we send the seven segment code directly to activate the required segments. This causes some of the segments to be ON and some to be OFF. Either way, as long as the display is ON, the module draws its required current from the power source. This may be in the range of 5 to 30 mA for a single segment to be lighted up. The display is ON all the time, and hence it is called a 'static' display. Suppose we need an eight digit display. If we use the same kind of static display, the current drawn is multiplied by 8, and this becomes quite a large amount. Multiply $7 \times 25 \times 8$ mA. This gives a value of 1.4 A, which is much too large for an electronic circuit. For this reason, static displays are not preferred for multiple digit displays.

9.13.3 | Dynamic Display

When there is an array of digit display units, say 8 seven segment LEDs arranged in digit form in a row, a continuous display can be obtained by lighting up just one digit at a time. The next instant this digit is switched off and the next one is lighted up. This is done continuously and cyclically from digits 1 to 8 and repeated at a rapid rate. Because of the property of persistence of vision of the eyes, an illusion of continuous display is obtained. This is also called a multiplexed display.

The important points to note here are:

- i) The common anode/cathode of a digit is to be activated for a digit to be active.
- ii) At a time, only the segments of one digit are 'ON'.
- iii) After a specified delay, this digit is switched off and the segments of the next digit are ON. The information displayed here is different from that of the previous case.
- iv) Thus, for display multiplexing, consecutive digits should be switched on in a cyclic fashion, and for each digit, the segment information should be supplied.




Figure 9.42 | Eight digit dynamic display connected to two ports of an 8255




Figure 9.43 | Segments of the display and the data byte

Now, let us use this concept in a system in which an 8086 handles a dynamic display using the 8255 as an interfacing chip (see Fig 9.42). This is an 8-digit display, of the **common cathode** type. The ports of 8255 are used in such a way that Port A supplies the digit information and Port B supplies the segment information. Digit information is to select which digit is being activated at a particular time. For segment information, the seven segment code of each digit should be sent as a byte. To identify which segment corresponds to which bit in a byte, see Fig 9.43. This is just the way that the port pins are connected to the leads of the seven segment module in this example, meaning that it is not a standard connection.

Example 9.20

Assuming a common cathode type of display, find the seven segment codes to be sent to Port B for displaying
i) 8, ii) A, iii) b.

Solution

Since it is a common cathode type, the common cathode of the LEDs of a digit are to be grounded. Then, supply the segment information to the anodes.

For displaying 8, the segments to be lighted up are a, b, c, d, e, f, g.
 Since, we are using a common cathode type of display, the data to light up a segment is '1'.
 Hence, the bits from D_0 to D_6 is '1'.

Thus, the seven segment code for the display of 8 is 0111 1111, i.e., 7FH.

i) Similarly for 'A' it is 0111 0111 i.e., 77H.

ii) For 'b' it is 0111 1100 i.e., 7CH.

9.13.4 | An Eight-Digit Dynamic, Multiplexed Display

In Fig 9.44, the connection between the LED, the port lines and other extra components are shown. Port A is the digit driving port. It gives the information about which digit is to be ON at a particular time. This is achieved by switching on one of the digit driving transistors Q_1 to Q_8 . These are PNP transistors and are turned on if a '0' is applied at their bases. This '0' goes to the emitters of the transistors, which are connected to the cathodes of all the segment LEDs. At a time, Port A gives a '0' only on one of its port lines. To understand this clearly, observe Fig 9.44. The most significant digit (or the left most digit of the display) is activated by a '0' on D_7 of Port A. Thus, when Port A outputs the binary sequence 0111 1111, the left most digit is scheduled to be on, but the segment information also is to be obtained from Port B. If Port B gives the data 77 H, the first digit displays 'A'. This is the method for displaying one digit.




Figure 9.44 | Eight digit dynamic display with segment and digit drivers

This is to be repeated for all digits as long as we want this data to be displayed. The steps are listed as follows.

1. Select the first digit to be displayed and send a suitable word to Port A to activate this digit.
2. Send the segment code through Port B.
3. Call a delay of say, 1.5 msecs.
4. Switch off this digit and repeat steps 1 to 3.
5. Repeat this sequence for all eight digits.
6. Then start again from the first step.

With 8 digits and 1.5 msecs delay, we can get back to the first digit every 12 msecs. This corresponds to a refresh rate of around 80 ms which is sufficient to fool the eye into believing that all the digits are ON at the same time.

Now, let us write a program to use the above circuit for displaying data.

Example 9.21

Display the word 'attitude' on the 8 LEDs.

Solution

Here, for activating the leftmost digit, the byte 0111 1111 is sent. For the second digit, this sequence is rotated right once. The segment code for each character is stored in the location named TABLE.

```

CR EQU 0C6H           ;address of the control register
PA EQU 0C0H           ;address of Port A
PB EQU 0C2H           ;address of Port B

TABLE DB    5FH,78H,78H,04,78H,3EH,5EH,7BH

MOV AL, 80H           ;CW with ports A and B as outputs
OUT CR, AL           ;send it to the control register

STRT:  LEA SI, TABLE  ;let SI point to the display data
        MOV CX, 8      ;CX to contain the number of digits
        MOV BL, 01111111 ;bit sequence for leftmost digit
AGN:   MOV AL, BL       ;copy it to AL
        OUT PA, AL      ;send it to Port A

        MOV AL, [SI]    ;copy display data to AL
        OUT PB, AL      ;send it to Port B
        CALL DELAY      ;call a delay
        INC SI          ;increment pointer
        ROR BL, 1        ;activate next digit
        LOOP AGN        ;repeat until CX = 0
        JMP START       ;repeat the sequence of display

DELAY PROC NEAR        ;delay loop
        MOV DX, 300     ;calculate delay value

```

```

BACK:  DEC DX
        JNZ BACK
        RET
DELAY ENDP

END

```

Example 9.22

Write a program for moving the character 'U' from right to left continuously.

Solution

The data corresponding to the character to be displayed is stored in the location DAT. When the first digit is switched on, it displays 'U'. Next, after a delay, this same character is displayed only on the second position. The important thing here is that the delay between successive activation of the digits should be high enough, so that the persistence of vision does not apply. Thus, the letter 'U' is seen to move from the left to right continuously on the display. The delay should be very high (compared to the case in Example 9.21) to create such an effect.

```

CR EQU 0C6H           ;address of the control register
PA EQU 0C0H           ;address of Port A
PB EQU 0C2H           ;address of Port B

DAT DB 3EH            ;data to be displayed - 'U'

MOV AL, 80H           ;CW for ports A and B as outputs
OUT CR, AL            ;send it to the control register
STRT:  MOV CX, 8        ;CX = 8, the number of digits
        MOV BL, 01111111 ;sequence for one digit to be 'on'
AGN:   MOV AL, BL        ;copy it to AL
        OUT PA, AL       ;send it to Port A
        MOV AL, DAT       ;data of 'U' is loaded in AL
        OUT PB, AL        ;send it to Port B as segment data
        CALL DELAY        ;call delay
        ROR BL, 1         ;rotate BL to turn on next digit
        LOOP AGN          ;repeat for 8 digits, until CX = 0
OK:    JMP STRT           ;repeat the display sequence

DELAY PROC NEAR        ;procedure for a long delay
        MOV DX, 30000
BACK:  DEC DX
        JNZ BACK
        RET
DELAY ENDP

```

This program rolls one character along the display. A very interesting program to try will be to have a message roll from left to right. Suppose we want to rotate the sequence ABCDEFGH,

Table 9.10 | Display Output for Rolling Display of ABCDEFG

	D1	D2	D3	D4	D5	D6	D7	D8
T1	A	B	C	D	E	F	G	H
T2	H	A	B	C	D	E	F	G
T3	G	H	A	B	C	D	E	F
T4	F	G	H	A	B	C	D	E
T5	E	F	G	H	A	B	C	D
T6	D	E	F	G	H	A	B	C
T7	C	D	E	F	G	H	A	B
T8	B	C	D	E	F	G	H	A

the display will be as shown in the Table 9.10. D1, D2 ... shows the digit positions and T1, T2 ... corresponds to the time instants at which the corresponding characters are seen on the display. The delay must be large enough to create a feeling of movement. The program for this problem is left as a challenge to interested readers.

KEY POINTS OF THIS CHAPTER

- A 'trainer kit' with an 8086 and associated circuitry helps us to understand the ideas of hardware interfacing to the 8086.
- For interfacing the processor to various peripherals, specialized peripheral controller chips are used.
- One chip which acts as an intermediary in parallel data transfer is the programmable peripheral interface (PPI) chip 8255.
- The PPI has three 8-bit ports A, B and C which are programmable as input/output ports.
- The PPI has three modes of operation.
- Mode 0 is the simple I/O mode while the modes 1 and 2 are handshaking modes.
- The parallel port of the PC uses mode 1 of the 8255 for all parallel data transfer, but the parallel port is getting to be more or less obsolete.
- An analog to digital converter can be interfaced to the 8086 using an 8255 as the intermediary.
- A DAC can be used to get analog waveforms generated by the 8086.
- LCDs are very popular display devices. Data and command words can be sent to its internal registers to get displays in different formats.
- Stepper motors are used to get precise and controlled rotation. They can be fed pulse sequences from the 8086 channeled through the PPI.
- The stepper motor needs more current than can be supplied by the 8255. Hence, current driver transistors /ICs are used in such circuits.
- For understanding how keyboards operate, a hex keyboard can be used.
- LEDs are popular displays, but they draw a lot of current. Hence, seven segment digit LEDs are used in a multiplexed fashion.

QUESTIONS

1. What is the special feature of Port C when used in
 - a) Mode 0
 - b) Mode 1
2. What does the term handshaking mean to you?
3. What is the importance of the BSR mode of the 8255?
4. Distinguish between the SPP and EPP protocols.
5. How can we operate an ADC in the interrupt mode?
6. How is the current output of a DAC converted to a voltage?
7. How can a sine wave be generated by a DAC?
8. What is meant by a backlit LCD?
9. What is the function of the BUSY pin of an LCD module?
10. How do we make the LCD display move from left to right?
11. Name a few applications of the stepper motor. Why cannot dc motors be used here?
12. Distinguish between the full step drive and half step drive sequences of the stepper motor driving sequences.
13. Why cannot we directly connect the stepper motor from the port pins of the 8255?
14. Why are optocouplers popular in stepper motor circuits?
15. What is meant by the step angle of a stepper motor?
16. Why are dynamic displays preferred in multi-digit displays?
17. Distinguish between common anode and common cathode displays.

EXERCISE

1. Design the control word of the 8255 with:
 - a) all ports as input ports,
 - b) Port A and B as input and port C as output.
2. Write the BSR control word to set PC3 and PC2 and clear PC5 and PC4.
3. With a 5 MHz clock, write a program to get a symmetric square wave from port B of the 8255.
4. Write a program to get an asymmetric square wave of 20% duty cycle from port pin PC2.
5. Write the control word for mode 1 operation with Port A in strobed output and Port B in strobed input.
6. Write instructions to test $IO/\overline{B}_B$ and IBF_B for the previous problem.
7. Find the name of any standard 12-bit ADC and list its control signals. Draw the figure showing how it can be connected to an 8086 through a parallel port.
8. Interface a DAC to the 8086 with a clock frequency of 5 MHz and write programs to generate
 - a) a triangular waveform of $T = 1$ msecs,
 - b) a ramp of $T = 3$ msecs,
 - c) a staircase waveform of 6 steps each of time period 0.2 msecs.
9. Use a 16×2 LCD module and write programs to do the following:
 - a) Display HELLO and WORLD on different lines and move the entire display back and forth.
 - b) Write four strings on different lines of the LCD which move.

10 PERIPHERAL INTERFACING - II



In this chapter, you will learn

- The architecture and programming features of the timer chip 8253/8254.
- The method of generating waveforms of different frequencies using this chip.
- The use of the different modes of operation of the timer chip.
- The use of a dedicated chip 8279 for keyboard and display interfacing.
- The connections between the 8279 and the display circuits.
- How to use this chip in a circuit with a keyboard and a display.
- How interrupts are managed by the programmable interrupt controller chip 8259.
- The use of initialization and operational control words of 8259.
- The method of changing the priority settings using this chip.
- The use of 8259 in PCs.

Introduction

In Chapter 9, we used the parallel port interfacing IC 8255 for many applications. This was done to illustrate the fact that data is available for transfer in the parallel form – so it can be sent to various external peripherals each of which has specific functions. In all these applications, the 8255 PPI mostly was used as a data transfer intermediary, along with the capability of generating signals for controlling the individual peripherals.

In this chapter and the next, we will familiarize ourselves with a few more peripheral chips with specialized applications. The first is the 8253/8254 timer chip.

10.1 | The Programmable Interval Timer 8253/8254

So far, whenever we wanted to create delays, we used the idea of software delay wherein the processor was kept in a loop repeatedly executing a set of instructions, thus creating a delay loop. This idea was extended to create square waves of required frequencies.

However, this is not a very good solution – first, because the processor is held up in this loop unable to do anything else and second, delays so created are not very accurate.

The use of dedicated hardware for generating delays and generating waveforms of required frequencies and pulse widths has solved the above problems.




Figure 10.1 | Pin diagram of 8253/8254

Table 10.1 | Address of the Control Register and Counters

$\overline{CS}$	A ₁	A ₀	Selected entity
0	0	0	Counter 0
0	0	1	Counter 1
0	1	0	Counter 2
0	1	1	Control register

The chip 8253 is a programmable timer chip used for timing applications. Another chip 8254 is also available, which is the same as 8253, except that it is capable of being operated at higher frequencies. In fact, 8254 is a superset of 8253 and all programs written for 8253 can be used for 8254. Both these timer chips provide three independent 16-bit timers, but 8253 can handle clock inputs of only up to 2.6 MHz, while 8254 can be used up to 10 MHz. Both chips are available as 24-pin DIP, and the pin diagram is as shown of 8253/8254 in Fig 10.1. In the IBM PC, the 8253 was used initially, but now in the PC-AT, the timer chip used is the 8254 which is pin compatible with 8253.

Now let us have a good look at the internal details of the chip. As the internal block diagram shows, there are three independent counters and the operation of these is managed by a control register. For selecting these four entities, two address pins A₁ and A₀ and the chip select line $\overline{CS}$ is used as shown in Table 10.1.

10.1.1 | Counters

The block diagram (Fig 10.2) shows that each counter has three pins associated with its operations – CLK, GATE and OUT.




Figure 10.2 | Functional block diagram of 8253/8254

CLK This pin corresponds to the pin on which an input frequency is received and all delays are calculated based on this frequency f . $T = 1/f$ is the time period of this input frequency and we will use T in all our calculations. This frequency should not be higher than 2.6 MHz for 8253, but can be as high as 10 MHz for the 8254.

GATE This functions as a control pin for each of the timers. This has to be high (Gate = 1) in most of the modes of operation, but is used as a trigger pin in some modes and then it should be a 0 to 1 pulse.

OUT Each counter has an output pin and the waveforms generated by programming the specific counter can be observed at this pin.

Data and Control Pins The other pins of the chip are the data pins D_0 to D_7 and the $\overline{RD}$ and $\overline{WR}$ pins. The data pins are connected to the lower order data pins of the processor. The read and write control pins are activated by the I/O read and I/O write cycles respectively, of the processor.

10.1.2 | Programming the Chip

All operations are decided by the control word loaded into the control register. For each of the counters, there is a count-register which is 16 bits in size. A number is written into this register and stored in the counter block. The maximum size of this number is FFFFH i.e., 16-bit, but since the counter's internal register is only 8 bits long, this number must be written as two 8-bit chunks. The operation of the counter occurs by creating a delay by decrementing this number down to 0 – the rate at which this occurs depends on the input clock frequency. Hence, this number decides the factor by which the input frequency gets divided to give an output frequency.

Suppose the number loaded is 3. Then during each input time period ' T ', the number decrements to 2, 1, and then 0. Thus, it creates a delay to divide the input frequency by 3. This is the way with any number ' N ' loaded in the counter register of a particular counter. The maximum delay is obtained by loading '0' as the number – it will decrement to 65,535 and then back to 0. Figure 10.3 gives an idea of what constitutes the hardware associated with a **single**




Figure 10.3 | Internal architecture of a counter

counter. The control register is common to all the counters, but the other blocks are replicated for each counter. CR stands for count-register, the subscripts H and L are for high byte (MSB) and low byte (LSB) respectively. CE is the counting element and OL stands for output latch. This block diagram will come in handy for understanding the procedure for reading the content of a count-register. This is discussed in Section 10.1.13.

10.1.3 | Control Word

Figure 10.4 shows the control word and the definition of each bit of the word. D₇ and D₆ (SC1 and SC0) are used to select the counter intended to be used. D₅ and D₄ (RL0 and RL1) point to 4 options. One option is ‘latch’. This will be needed only when we want to ‘read’ the contents of the count-register. The other three options are for writing a number ‘N’ into the count-register. This number can be 8 or 16 bits long. There might be confusion about why there are three ways of loading the count in the counter. One possibility is that the count will be two bytes in size. In that case use RL1, RL0 = 11 in the control word and then load the count. So load the LSB first, followed by the MSB. Later, if the LSB byte alone needs to be changed then the control word should be changed and reloaded with RL1 and RL0 as 01, or as 10 if the MSB alone need to be changed. After configuring the control word in this way, the count can be loaded in the count-register. D₃ to D₁ (M2 to M1), select the mode in which the counter is to be used. The chip has 6 modes, which will be explained in detail soon. D₀ is the bit which informs the counter of the format of the ‘count’. If it is BCD in form, this bit is set to be 1, if binary,




Figure 10.4 | Control word format of the 8253/8254

it is '0'. In binary, the count can be from 0 to FFFFH – in BCD, the number can vary only from 0 to 9999.

Recollect that the maximum BCD count can only be 9999, which is to be written as 9999H (refer Section 0.6.4), while the maximum binary count can be FFFFH.

Example 10.1

Configure the control word for the following cases.

- a) Counter 0 in mode 3, count in BCD, specify that a 16-bit count is written with LSB first.
- b) Counter 2 in mode 0, count in binary format, count loaded as in the previous case.

Solution

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
0	0	1	1	0	1	1	1

The control word is written referring to Fig 10.4 and is 37H

Solution

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
1	0	1	1	0	0	0	0

The control word in this case is B0H

Example 10.2

Find the addresses of the control registers and the three counters for the decoding logic shown (Fig 10.5).

Solution

In the circuit, we see that the address lines A₂ and A₁ of the 8086 are connected to A₁ and A₀ of the 8253. The other address lines (A₇ to A₃ and A₁) are used for address decoding as shown. The address of various entities, within the chip are as follows:

Entities	A ₇	A ₆	A ₅	A ₄	A ₃	A ₂	A ₁	A ₀	Address (Hex)
Control Register	1	1	0	0	1	1	1	0	CE
Counter 0	1	1	0	0	1	0	0	0	C8
Counter 1	1	1	0	0	1	0	1	0	CA
Counter 2	1	1	0	0	1	1	0	0	CC




Figure 10.5 | Decoding logic for the chip

Now, let us use the timer chip in its various modes. The different modes are listed as:

Mode No.	Mode bits			Operation
	M2	M1	M0	
0	0	0	0	Interrupt on terminal count
1	0	0	1	Programmable one shot
2	X	1	0	Rate generator
3	X	1	1	Square wave generator
4	1	0	0	Software triggered strobe
5	1	0	1	Hardware triggered strobe

To get a feel about the use of the timer chip, we will first use it in Mode 3 which is the simplest and direct of the modes. Mode 3 is used to generate square waves of a desired frequency.

Example 10.3

The input frequency to the timer chip is 1.5 MHz. Generate a square wave of 1.5 KHz using Timer 0.

Solution

For an input frequency of 1.5 MHz, $T = 1 / (1.5 \times 10^{-6}) = 0.66 \mu\text{secs}$
To get an output frequency of 1.5 KHz, the time period = 0.66 msec
N is the number of clock cycles of the input frequency to get 0.66 msec time period for the output signal.

$$N = 0.66 \text{ msec} / 0.66 \mu\text{secs} = 1000 \text{ cycles.}$$

Thus, the number to be loaded in the count-register is 1000 i.e., 1000 is the number that divides 1.5 MHz to get 1.5 KHz. What this means is that, when the counter starts running, this number decrements down to zero, and 1000 cycles of the input signal occur during this period. Before writing the program, the control word is to be designed (refer Fig 10.4). The control word is

0 0 1 1 0 1 1 0 i.e., 36H

This word selects counter 0 and uses mode 3. It also loads the count-register with LSB first and then the MSB. The count is loaded as a binary number. Hence D₀ is '0'.

Program

```
CR      EQU 0CEH
CNT0    EQU 0C8H

MOV     AL, 36H           ;control word
OUT     CR, AL            ;send control word to control register
MOV     AX, 3E8H          ;load count = 1000 in AX
OUT     CNT0, AL          ;send lower byte of count to counter 0
MOV     AL, AH            ;move higher byte of count to AL
OUT     CNT0, AL          ;send higher byte of count to counter 0
END
```

The output signal from the terminal OUT 0 divides the input frequency by 1000. It is seen that the input frequency is divided by N in this mode. For this mode, if the number N is even, the high and low portions are of the same period. If N is odd, the high pulse is $(N + 1) / 2 \times T$ and the low pulse is $(N - 1) / 2 \times T$. Note that T is the period of the input frequency. It is mandatory to keep GATE = 1 for the counter to run. We load a number into the count-register which decrements down to 0. After that, this number gets re-loaded into the count-register and that is how a continuous square wave is obtained.




Figure 10.6 | Output waveform corresponding to Example 10.3

10.1.4 | Using all the Counters

Note There are three channels in this chip i.e., three different counters with individual clocks, gates and output pins. Hence three different signals can be obtained from the chip simultaneously. If, say, we need three square waves of different frequencies, we can program all the three counters in mode 3. Try the following program.

Example 10.4

Generate three different frequencies from the 8253 timer. Use a clock frequency of 1.5MHz for all the counters.

Solution

```

CR      EQU      0CEH      ;address of control register
CNT0    EQU      0C8H      ;address of counter 0
CNT1    EQU      0CAH      ;address of counter 1
CNT2    EQU      0CCH      ;address of counter 2

MOV     AL, 37H            ;control word to use counter 0
OUT     CR, AL             ;send control word to control register
MOV     AX, 10H            ;load count = 10 in AX
OUT     CNT0, AL           ;send low byte of count to counter 0
MOV     AL, AH             ;move high byte of count to AL
OUT     CNT0, AL           ;send high byte of count to counter 0

MOV     AL, 77H            ;control word to use counter 1
OUT     CR, AL             ;send control word to control register
MOV     AX, 100H           ;load count = 100 in AX
OUT     CNT1, AL           ;send low byte of count to counter 1
MOV     AL, AH             ;move high byte of count to AL
OUT     CNT1, AL           ;send high byte of count to counter 1

MOV     AL, 0B7H           ;control word to use counter 2
OUT     CR, AL             ;send control word to control register
MOV     AX, 1000H          ;load count = 1000 in AX
OUT     CNT2, AL           ;send low byte of count to counter 2
MOV     AL, AH             ;move high byte of count to AL
OUT     CNT2, AL           ;send high byte of count to counter 2
END

```

With this program, counter 0 divides the input frequency by 10, counter 1 by 100 and counter 2 by 1000. We have used the BCD of bit $D_0 (= 1)$ of the control word, and so used the numbers 10H, 100H and 1000H as the count for the three cases. It should be clear that to write BCD in hex, 1000 should be written as 1000H, 100 as 100H and so on (refer Section 0.6.4). To use the binary option, we should write the counts as 0AH, 64H and 3E8H after ensuring that the D_0 bit in the control word is made '0', for each of the cases. The control words of the three counters would then be 36H, 76H and B6H.

Now, let us examine all the modes of the timer chip. Keep in mind that, associated with each of the three counters, there are three pins – the clock, the gate and the out. The gate

is a control input. In some modes, it is mandatory to keep $GATE = 1$, but in certain other modes, the $GATE$ should have a low to high transition for the operation in that mode to take effect.

10.1.5 | Mode 0: Interrupt on Terminal Count

This can be used to generate an event after a programmed time period. $GATE = 1$ should be maintained for this mode to function. What we need to do to use in this mode, is to load a number N in the count-register. The output is normally low and will go high after the terminal count is reached i.e., when N rolls down to 0. The width of the low pulse is $N \times T$ where T is the duration of the input clock pulse. This mode can be used to interrupt the microprocessor after a delay period.

Note If, in between, the $GATE$ pin goes low, counting will be stopped and the low duration will be extended by the amount of time that the gate is kept low.

Example 10.5

Use counter 0 in mode 1 to generate a delay of 10 msec. After this delay, the OUT 0 pin must go high. The clock frequency used for the counter chip is 1.5 MHz.

Solution

Calculation to generate a delay of 10 msec

$$T = 10 \text{ msec} / (0.66 \mu\text{s}) = 15151 = 382 \text{ FH}$$

Referring to Fig 10.4, write the control word for mode 0.

CW is 00110000

```
CR      EQU    0CEH      ;address of control register
CNT0    EQU    0C8H      ;address of counter 0

MOV     AL, 30H           ;control word for mode 0, counter 0
OUT     CR, AL            ;send it to control register
MOV     AX, 382FH         ;load count in AX
OUT     CNT0, AL          ;send LSB of count to counter 0
MOV     AL, AH            ;move AH to AL
OUT     CNT0, AL          ;send MSB of count s to counter 0
END
```

The output will become high after 10 msecs.

This mode can be used as an event counter, and signal a processor when the number of events has reached a predefined number. Suppose it is needed to have a signal when 100 bottles have passed through a conveyor belt. An optical sensor can sense this, generate a pulse for each bottle, and the output of the sensor is fed to the clock input of counter 0 in mode 0. A count of 64H is loaded in the count-register. When the counter counts down to 0, the OUT pin of the counter goes high, and this signals the event of 100 bottles passing the conveyor belt. If the OUT pin is connected to the INTR pin of the 8086, appropriate action can be programmed in the interrupt service routine of the processor.




Figure 10.7 | An example of a practical application of mode 0

10.1.6 | Mode 1: Programmable One Shot

This mode is also called the **hardware re-triggerable one-shot**. After the count-registers are loaded, the GATE must be given an L to H transition, which constitutes the trigger for the one-shot. Once the trigger is obtained, the OUT pin goes low and remains low until $N \times T$ where N is the count loaded in the count-register and T is the period of the input clock. Meanwhile, if the GATE is given another L to H transition, the count is re-loaded and counting starts again. This **re-triggering capability** is what gives it the name of a ‘re-triggerable one shot’. Refer Fig 10.8a and 10.8b.




Figure 10.8a | Output waveform corresponding to mode 1




Figure 10.8b | Output waveform when a new L to H transition is given to the GATE before $N \times T$ is reached

Example 10.6

Consider a gate input waveform of 1 KHz, clock input of 1.5 MHz ($T = 0.66 \mu\text{secs}$).

The following program gives a low output pulse of 0.2 msec, from the output of counter 0 only when the gate is given a low to high transition. To calculate the count value, $N \times 0.66 \mu\text{secs} = 0.2 \text{ msec}$. $N = 303 = 12 \text{ FH}$.

Apply a 1 KHz signal at the gate input and draw the output waveform at OUT 0. The low pulse is not produced when the count is loaded into the count-register – it starts only when a low to high transition occurs on the gate.

Solution

Control word is 0011 0010 = 32H

```

CR      EQU    0CEH      ;address of control register
CNT0    EQU    0C8H      ;address of counter 0

MOV AL, 32H              ;control word for mode 1, counter 0
OUT CR, AL               ;send it to control register
MOV AX, 12FH             ;move count to AX
OUT CNT0, AL             ;send LSB to counter 0
MOV AL, AH               ;move MSB to AL
OUT CNT0, AL             ;send it to counter 0
END

```




Figure 10.9 | Output waveform of Example 10.6

10.1.7 | Mode 2: Rate Generator

Here if $\text{GATE} = 1$ and a count is loaded into the count-register, the output will be high for the duration $N \times T$, and will then go low for a period of one clock cycle. Thus the period of the output waveform can be considered to be $(N + 1) T$ where it is high for the duration $N \times T$ and low for one cycle. The count is reloaded automatically and the counter continues to produce the output waveform. This mode is also called a divide by N counter. In Fig 10.10, WR corresponds to the time in which the count is written into the counter.




Figure 10.10 | Output waveform corresponding to mode 2 operation

Example 10.7

```

CR      EQU    0CEH      ;address of control register
CNT0    EQU    0C8H      ;address of counter 0

MOV AL, 35H              ;control word for counter 0, mode 2
OUT CR, AL                ;send it to control register
MOV AX, 0005              ;load count N = 5 in AX
OUT CNT0, AL              ;send LSB to counter 0
MOV AL, AH                ;move MSB to AL
OUT CNT0, AL              ;send MSB to counter 0
END

```

This program generates a low pulse of one T at OUT 0, after a period of 5 T cycles.

Example 10.8

Generate a square wave of 10 KHz from counter 0, and use counter 2 to get an asymmetric output waveform.

Solution

Apply a clock of 1.5 MHz to CLK0. Connect OUT 0 to CLK2 pin, and observe the output at OUT2 pin.

```

CR      EQU    0CEH      ;address of control register
CNT0    EQU    0C8H      ;address of counter 0
CNT2    EQU    0CCH      ;address of counter 2

MOV AL, 37H              ;control word to use counter 0, mode 3
OUT CR, AL                ;send control word to control register
MOV AX, 0150H            ;load count in AX
OUT CNT0, AL              ;send low byte of count to counter 0
MOV AL, AH                ;move high byte of count to AL
OUT CNT0, AL              ;send high byte of count to Counter 0
MOV AL, 0B4H             ;control word to use counter 2 in mode 2
OUT CR, AL                ;send control word to control register
MOV AX, 0003              ;load count in AX

```

```

OUT CNT2, AL      ;send low byte of count to counter 2
MOV AL, AH        ;move high byte of count to AL
OUT CNT2, AL      ;send high byte of count to counter 2
END

```




Figure 10.11 | OUT 0 and OUT 2 outputs

The signal from OUT 0 is a symmetric square wave obtained by dividing CLK0 (1.5 MHz) by 150. Its frequency is 10 KHz and its T is 0.1 msec. Counter 2 is programmed to run in mode 2 with a count of $N = 3$. Its input frequency is 10 KHz. It will generate an output signal which is high for $3T$ (0.3 msec) and low for one period of its input clock i.e., 0.1 msec. Hence, an un-symmetric output waveform is obtained, as shown in Fig 10.11.

10.1.8 | Mode 3: Square Wave Generator

This mode has been discussed earlier in Example 10.3 and Section 10.1.4.

10.1.9 | Mode 4: Software Triggered Mode

In this mode, if GATE = 1, the output will be high for the duration $N \times T$ and will go low for one clock period. It then goes high again. Repetition of the sequence is possible only after reloading the count-register. This mode is similar to Mode 2, except that here, the counter is not re-loadable.

10.1.10 | Mode 5: Hardware Triggered Mode

In the previous mode, counting started when the count-register was loaded. However, here the counting starts only when an L to H transition is given to the GATE. In Fig 10.12, WR is the time at which the count is written into the count-register.




Figure 10.12 | Output waveform for mode 5, starting from the L to H transition of the GATE

10.1.11 | Reading a Count

In many applications, it may be important to read the contents of the count-register. This can be done by stopping the counting. To stop the counting, the method is to make the GATE pin equal to 0, or to inhibit the clock. However, the better method is to do the reading ‘on the fly’ which does not cause any interference on the circuit operation. Referring to Fig 10.3 shows that there is a latch associated with the counting element. To read a stable value from a counter, the method is to latch the count and then read the content of the latch. Refer to Fig 10.4 which shows the control register format – RL1 and RL0 being 0 gets the count to be latched. After that, the contents can be read. The following are the steps for reading the count of counter 0. This program segment can be used anywhere in the program once counting starts, which means that it performs reading ‘on the fly’.

```

MOV AL, 00000000    ;control word for latching counter 0
OUT CR, AL          ;send it to control register
IN AL, CNT0          ;input the LSB of counter 0
MOV AH, AL           ;move it to AH
IN AL, CNT0          ;input the MSB of counter 0
XCHG AH, AL          ;exchange the contents of AH and AL
MOV COUNT, AX        ;save the count in memory

```

Now the current count is available in AX, which is to be moved into COUNT, a memory location.

10.2 | The Programmable Keyboard Display Interface – 8279

In Chapter 9, you had an exposure to the techniques of interfacing keyboards as well as displays to the 8086. What should have been noticed is that when a key is expected to be pressed, the processor is fully involved in verifying this and then identifying the key that has been pressed. Similarly, for displaying data in multiplexed LED displays, the processor is fully involved in this work of refreshing individual digits, sending segment information and so on.

We would like to have a peripheral chip which will take care of these activities, such that the processor has only to give commands for the actions. One chip which takes up the responsibility of managing these functions, is the 8279 keyboard display interface. This chip has two portions – one for providing scanned display interface for LED, incandescent and other popular display technologies in alphanumeric form, and the other for interfacing with keyboards or an array of sensors.

We will discuss the interfacing of keyboards and displays separately, and then use them jointly. Our approach will be to discuss the features of the chip, with the help of an example in which the 8279 is interfaced to a keyboard as well as a display. Figure 10.13 shows the functional block diagram of the chip. $\overline{BD}$ is used to blank the display during digit switching or by a display blanking command. We will discuss the important pins when we use them in practical circuits, as we go ahead.

10.2.1 | Display Interface

Let us consider using the chip for interfacing to a multiplexed seven segment LED of 8 digits. This will help to understand the use and capabilities of the chip. Refer to Fig 10.14. However, do not be intimidated by the seemingly complicated diagram. It is actually very simple and easy to understand.




Figure 10.13 | Functional pin diagram of the 8279

The circuit is for using the 8279 for interfacing an 8-digit, multiplexed dynamic display. As discussed in Section 9.13.2, for a multiplexed display, two signal sequences are required.

- i) Only one digit is to be activated at a time, and so a sequence which switches 'ON', only one of the digit transistors is needed. The scan lines SL0 to SL3 of the 8279 generate a continuous sequence of numbers from 0000 to 1111. In the case that we are discussing, since only eight digits are to be displayed, SL0 to SL2 are connected to a three to eight decoder 74LS138. This gives a low signal only at one of its eight outputs, and thus each digit transistor is switched on in sequence. Only one of the eight transistors Q9 to Q16 will be ON at a time. The LEDs are connected as common cathode, and when, say, transistor Q9 is ON, it grounds the cathode of the LEDs of the corresponding digit. So when segment information is given at the anodes, this digit can display the character corresponding to the segment information.
- ii) When a particular digit is activated, the segment information corresponding to that digit will get its output from the chip. This is done at the output lines A0 to A3 and B0 to B3 of the chip. For increasing the current driving capability (to give a bright display) transistors Q1 to Q7 are used between the output lines and the segments of the LED. Thus a segment gets lighted up only when a 'high' appears on the collectors of the driving transistors (note that the transistors are PNP type). **For that to happen, the output lines A0 to A3 and B0 to B3 must be low – which means, that for a segment to be lighted up, the bit from the 8279 must be a '0'.** However, note that this is not a feature of the 8279; rather, it is due to the connections used between the chip and the display.




Figure 10.14 | Display interface part of the 8279

- iii) The chip has a 16-byte display RAM into which the seven segment codes of the display characters are to be written. The chip automatically cycles through a process of sending out the code for one digit, turning on the digit for a short time and then repeating the procedure for the next. In between the displaying of one digit and the next, the chip also sends a ‘blanking code’ which turns off all the segments, so as to prevent ‘ghosting’ of information from one digit to the next.

10.2.2 | Segment Definition

The segment definitions are as shown in Fig 10.15. The table below shows the correspondence between the data bus and the output port bits of the 8279. This definition is only because of the way the lines A0 to A3, and B0 to B3 have been connected in this case.

Data bus	D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
8279 outputs	A3	A2	A1	A0	B3	B2	B1	B0
Segments	d	c	b	a	dp	g	f	e

Example 10.9

What data should be outputted from the 8279, to display the following characters, with reference to the connection in Fig 10.14?

- i) A
- ii) 5
- iii) blank

Solution

For the connections in Fig 10.14, for any segment to be lighted so, the corresponding bit is to be ‘0’.

- i) For A, the segments to be activated by a ‘0’ are a, b, c, e, f, g. From Fig 10.15, the corresponding byte is 1000 1000 i.e., 88H.
- ii) For 5, the segments to be lighted up are a, c, d, f, g. The byte is 0010 1001 i.e., 29H.
- iii) For no segment to be lighted up, the byte is 1111 1111 i.e., FFH.




Figure 10.15 | Correspondance between the segments and the data bus lines

10.2.3 | Keyboard Section

The chip has the capability to be used with keyboards as well as with sensor matrices. However, here we will discuss only the former case. Let's consider a matrix keyboard and recollect how a key press is identified (Section 9.12). '0's are sent on one dimension of the matrix and the other dimension is read in – the lines that are read in are called 'return lines (RL0 to RL7 here). However, here we have to send out a binary sequence. That is done by the scan lines SL0 to SL3. There are two options for using these scan lines. There are two kinds of keyboard scans – encoded scan and decoded scan.

Encoded Scan For the keyboard also, the scan lines SL0 to SL3 of the chip are used. In the encoded scan mode, these lines generate a binary sequence from 0000 to 1111. A decoder is then used to generate a low on only one of its output lines. This is shown as is used in our example.

Decoded Scan In this case, a decoder is not needed. The pins SL0 to SL3 generate a sequence of 'low' pulses directly on the lines. This can be used if there are only 4 digits to refresh.

To detect a keypress and identify the key, the lines RL0 to RL7 are the 'return' lines. There are return buffers inside the chip which buffer and latch the return lines. In the keyboard mode, these return lines are scanned looking for key closures in that row. If a key press is sensed, a debounce circuit (inside the chip) waits for 20 msecs to confirm the keypress. If key closure is confirmed, the address of the key (in terms of its row and column position and status of SHIFT and CONTROL pins) is stored in a FIFO RAM.

Let us understand this by using the keyboard shown in Fig 10.16. Here we use a hex-keyboard with 16 keys, besides two keys for CNTL (control) and SHFT (shift). The 16 keys are arranged in two rows and eight columns. The 3 to 8 decoder (used for the display interface as well) gives two signals S0 and S1 (in response to the sequence produced in SL0 to SL2 – Ref Fig 10.14 also) and these two lines are connected to the rows of the keyboard, while the columns are connected to the pins RL0 to RL7 of the chip. The RLs are 'return' input lines




Figure 10.16 | Keyboard section




Figure 10.17 | Return key code

from the columns to the chip. The idea is that the scan lines cause ‘lows’ to be outputted on the rows and the status of the columns are returned to the chip. Remember that this is the same as the principle of keyboard interfacing discussed in Section 9.12, but here, the 8279 does this automatically in response to the control words we use.

The return key code which is produced has a format as shown in Fig 10.17. This code is stored in an eight-byte FIFO RAM where FIFO stands for ‘First-In First-Out’. This RAM thus stores the key codes in the order in which the keys have been pressed.

A valid keypress also causes an interrupt to be generated on the *IRQ* line of the chip. This can be connected to the interrupt line of the processor, so as to operate the keyboard in an interrupt driven mode. Another thing that a valid keypress causes, is the incrementing of a FIFO count in a status register. This information can be used to verify if the FIFO contains a valid key code.

10.2.4 | System with Keyboard and Display Interfaced

Figure 10.18 shows the block diagram of the system that we are using i.e., 8-digit multiplexed LED display and a 2 × 8 hex keyboard with SHFT and CNTRL keys. With this minimum information, let us get to the task of programming the chip for our application.

CLK This is connected from the system clock to generate timing.

A0 The pin A0 is used to select between data and command. If A0 is low, it indicates that the data bus contains data; otherwise it is to be considered as command.

10.2.5 | Control Words

It was mentioned earlier that when A0 is high, the byte sent from the processor is a command. Thus, it is obvious that the control register has one address only (and the data register has another). However, because of the many functions that this chip performs, a number of control words are to be written. **The upper three bits of the control register specify the type of the control word.** We will discuss the control words in the order in which they are to be used for a particular application. Only the control words necessary for the application will be discussed. For more information, the datasheet of the chip may be referred.

i) Keyboard / Display Mode Setup

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
0	0	0	D	D	K	K	K




Figure 10.18 | Connections between the 8279, an 8 digit display and a 2 × 8 keyboard

In this word, DD stands for display and KKK for keyboard.

DD has the following options:

- 0 0 Eight 8-bit character display – Left entry
- 0 1 Sixteen 8-bit character display – Left entry
- 1 0 Eight 8-bit character display – Right entry
- 1 1 Sixteen 8-bit character display – Right entry

The display can be initialized to 8 or 16 character display. In our example, since we have only an 8-digit display, DD can be set for 8 character display. This will cause the refresh lines SL0 to SL2 to change from 000 to 111.

Left Entry This is the simplest display format and is like a typewriter. The first character is displayed in the left most position and the second one will be displayed right of the first one and so on – just like the typewriter types on paper.

Right Entry This is the way most electronic calculators display data. The first entry is placed on the right most position. For the second entry, the earlier one is shifted left and the new character is placed again at the right most position. This is repeated for all the characters.

Since the characters come from the display RAM, see Fig 10.19 for the way the contents of the display RAM are displayed for each case. T_1 to T_4 are the time instants at which the characters stored in the display RAM, appear on the display. The seven segment code of the characters A, B, C, D is in the RAM at address 0, 1, 2, 3. KKK has the following variations

- 0 0 0 Encoded Scan Keyboard – 2 Key Lockout
- 0 0 1 Decoded Scan Keyboard – 2 Key Lockout
- 0 1 0 Encoded Scan keyboard – N Key Rollover
- 0 1 1 Decoded Scan Keyboard – N Key Rollover
- 1 0 0 Encoded Scan sensor matrix
- 1 0 1 Decoded Scan sensor matrix
- 1 1 0 Strobed input, Encoded Display Scan
- 1 1 1 Strobed input, Decoded Display Scan

Two key lockout happens when the keys are such that only if the first key is released, will the second key be considered at all. In N key rollover, if, say, two keys are pressed almost simultaneously, both of them are debounced and their codes will be stored in FIFO RAM in the order in which they were pressed.

A sensor matrix is a matrix of switch type sensors, the condition of which can be stored in the FIFO RAM ($8 \times 8 = 64$ switch states).




Figure 10.19 | Contents of the display RAM displayed for right entry and left entry

Example 10.10

Write the keyboard/display control word for the case of an 8-character display with right entry and encoded scan keyboard – 2 key lockout.

Solution

Refer the keyboard/display mode setup control word. The control word is

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
0	0	0	1	0	0	0	0

i.e., 10H

ii) Program Clock

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
0	0	1	P	P	P	P	P

All timing and multiplexing signals for the 8279 are generated by an internal pre-scaler. This pre-scaler divides the external clock by a programmable integer. Bits PPPPP determine the value of this integer which ranges from 2 to 31. Choosing a divisor that yields 100 KHz will give the specified scan and debounce times. For instance if the 8279 is clocked by a 2MHz signal, PPPPP should be set to 10100 to divide the clock by 20 to yield the proper 100 KHz operating frequency. For dividing by 30, the control word is 0011 1110.


iii) Clear Display

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
1	1	0	C _D	C _D	C _D	C _F	C _A

C_D C_D C_D– The lower two C_D bits specify the blanking code to be sent to the segments to turn them off while the display is switching from one digit to the next.

C_F – If C_F = 1, FIFO status is cleared, interrupt and output lines are reset, and sensor RAM pointer is set to Row 0.

C_A – Clear all bit has the combined effect of C_D and C_F. It uses the C_D clearing code on the display and clears the FIFO status. It also synchronizes the internal timing chain.



Example 10.11

Find the clear display control word for the circuit of Fig 10.14.

Solution

We have discussed that for this circuit, the bit required to blank a segment is '1'. Hence, D_3 and D_2 are 11. To enable clear display, $D_4 = 1$. C_F is not relevant and may be 0. We can have $C_A = 1$.

Hence, the required control word is 1101 1101 i.e., CCH

iv) Write Display RAM

D_7	D_6	D_5	D_4	D_3	D_2	D_1	D_0
1	0	0	AI	A	A	A	A

To display data, it must be entered into the display RAM. After this word is written, all subsequent writes will be to the display RAM. Remember that commands are written using an address which ensures that pin $A_0 = 1$. Data is written with $A_0 = 0$.

AI – Auto Increment. If AI = 1, the row address selected in the display RAM will be incremented after each write.

AAAA – Selects one of the 16 rows in the display RAM – thus it varies from 0000 to 1111.

Example 10.12

Write the display RAM control word

- i) with auto-increment
- ii) without auto-increment

Solution

Here, let us select the first location in the display RAM.

- i) With AI = 1, the control word is 1001 0000 i.e., 90H.
- ii) With AI = 0 (no auto-increment), the control word is 1000 0000 i.e., 80H

v) Read FIFO/Sensor RAM

D_7	D_6	D_5	D_4	D_3	D_2	D_1	D_0
0	1	0	AI	X	A	A	A

This control word is for setting up the FIFO/SENSOR RAM

AI – Auto Increment. Irrelevant for scanned keyboard mode, but used in sensor RAM mode.

AAA – Irrelevant for scanned keyboard mode, but used in sensor RAM mode.

Example 10.13

Write the control word for setting up the FIFO so as to read in keys that are pressed.

Solution

The control word is

0100 0000 i.e., 40H

Example 10.14

Write a program to display '5' in the first LED position of the LED display used, as per the block diagram of Fig 10.14. The display consists of 8 LEDs to be used in the left entry mode. The keyboard is of type 'encoded scan keyboard – 2 key lockout'. The address of the data register is 0C0H and of the control register is 0C2H.

Solution

The control words for setting up the display, clearing the display, and for writing into the display RAM have been designed in Examples 10.10 to 10.12 and may be used here.

```

DATR EQU 0C0H
CNTR EQU 0C2H

MOV AL, 00          ;control word for 8 digit, left entry display
OUT CNTR, AL        ;send it to the control register
MOV AL, 3EH         ;word to divide clock frequency by 30
OUT CNTR, AL        ;send it to control register
MOV AL, 0CCH        ;control word for clearing display
OUT CNTR, AL        ;Send it to control register
MOV AL, 90H         ;word to write display RAM with auto-increment
OUT CNTR, AL        ;Send it to control register
MOV AL, 29H         ;seven segment code to display '5'
OUT DATR, AL        ;Send it to data register
MOV AL, 0FFH        ;seven segment code for no display
MOV CX, 07          ;7 display LEDs to be blanked
NEXT: OUT DATR, AL   ;send the data to the data register
      LOOP NEXT      ;repeat for seven digits

```

This program displays '5' in the first digit location. The rest of the digits are not lighted up because the data given to these digits is FFH, which corresponds to no segment being lighted.

Example 10.15

Write a program to display a rolling message 'H' continuously moving left.

Solution

Here the seven segment code for 'H' is in the location named TABLE. Then the codes of 'no display' i.e., FFH is written 3 times and this is repeated once more. A counter of 8 causes this

sequence to be displayed once. The display seen is 'H' moving left followed by three blanks. In effect, it shows 'H' moving from right to left continuously.

```

        DATR EQU 0C0H
        CNTR EQU 0C2H
TABLE DB 98H, 0FFH, 0FFH, 0FFH, 98H, 0FFH, 0FFH, 0FFH

STRT:   LEA SI, TABLE           ;let SI point to the address of table
        MOV CX, 07              ;CX to act as a counter
        MOV AL, 10H             ;word for 8 digit, right entry display
        OUT CNTR, AL            ;send it to control register
        MOV AL, 3EH             ;word to divide clock frequency by 30
        OUT CNTR, AL            ;send it to control register
        MOV AL, 0CCH            ;word to clear display
        OUT CNTR, AL            ;send it to control register
        MOV AL, 90H             ;write to display RAM with auto-increment
        OUT CNTR, AL            ;send it to control register
NEXT:   MOV AL, [SI]             ;get data in table to AL
        OUT DATR, AL            ;send it to the data register
        CALL DELAY              ;call a delay
        INC SI                  ;increment pointer
        LOOP NEXT               ;repeat until CX = 0
        JMP STRT                ;start all over again

        DELAY PROC NEAR         ;delay procedure
        MOV DX, 0AFFFH
AGN:    DEC DX
        JNZ AGN
        RET
        DELAY ENDP

```

10.2.6 | Detecting a Keypress

Whenever a key is pressed, the code of the key enters the FIFO RAM. This action also increments a count in a **status** register which gives the information regarding the number of characters in the FIFO and gives, besides other things, the indication as to whether the FIFO is full. This status register may be read to get the required information. Figure 10.20 shows the bit configuration of the status register, which has the same address as the control register. The difference is that one 'writes' into a control register, but 'reads' from a status register. Observe that the lowest three bits indicate the number of characters in the FIFO RAM. If these bits are zero, it is understood that no valid key presses have occurred. There is bit D_3 to indicate that the FIFO is full. The two error conditions it indicates are the overrun and underrun errors. The former will be set when writing into a full FIFO is attempted. What it means is that more than 8 characters have been written in the FIFO RAM, and hence the oldest data has been lost. Under run error implies an attempt to read an empty FIFO.




Figure 10.20 | Format of the status word

Example 10.16

Write a program to read a key that is pressed and store the key code in memory.

Solution

Refer to Fig 10.17 which shows the way a key code is specified. This code will be entered into the FIFO RAM when the key is pressed. To read the key, the steps are (after pressing the key):

- Read the status word, and check if the lowest three bits are zero. If they are, it means that the FIFO RAM is empty meaning that no valid key press has been detected. The status register has the same address as the control register, the difference being that a status word is 'read' while a control word is 'written'.
- If a valid key press is identified, read the code from the FIFO RAM, after writing the 'read FIFO RAM' control word into the control register.

```

DATR EQU 0C0H
CNTR EQU 0C2H
DAT DB ?           ;the code of the key should be here

NEXT:  LEA SI, DAT      ;let SI point to a location in memory
        IN CNTR, AL     ;read in the status register
        TEST AL, 01     ;check whether D0 = 1
        JZ NEXT        ;if not, continue checking
        MOV AL, 40H     ;otherwise, write the word to read FIFO
        OUT CNTR, AL    ;send it to the control register
        IN AL, DATR     ;read the FIFO RAM
        MOV [SI], AL    ;send the content to the location DAT

```

In the above program, while reading the status register only the D_0 bit is tested. This is because it is assumed that the FIFO was empty earlier and so, a key press makes the count equal to 1.

In the location DAT, we get a key code C6 if '6' was pressed, CE if 'E' is pressed and so on. The key code is obtained as shown in Fig 10.17 and it depends on the positioning of the keys in the key matrix. This code corresponds to the keyboard used here.

Example 10.17

Write a program which displays the key that has been pressed.

Solution

Here, the seven segment code corresponding to the keys is stored in memory and the program accesses these codes and displays it, depending on the key pressed. The seven segment codes are stored in the order 0, 1, 2 ... F for the 16 keys. No provision has been made to account for pressing the shift and control keys.

When a key is pressed, and a key code is obtained, the upper nibble of this byte is masked off. Thus, for the key 6, the key code is C6, which is 06 after the masking operation. The seven segment code of this key is available in the location TABLE + 6.

```

        DATR EQU 0C0H
        CNTR EQU 0C2H
TABLE DB 0CH, 9FH, 4AH, 0BH, 99H, 29H, 28H, 8FH,
        08, 09, 88H, 38H, 6CH, 1AH, 68H, 0EH

        LEA SI, TABLE      ;pointer to the seven segment codes
        MOV CX, 08          ;counter
        MOV AL, 0           ;word for display RAM, right entry
        OUT CNTR, AL        ;send it to control register
        MOV AL, 3EH         ;word to divide clock frequency by 30
        OUT CNTR, AL        ;send it to control register

        MOV AL, 3EH         ;word to divide clock frequency by 30
        OUT CNTR, AL        ;send it to control register
        MOV AL, 0CCH        ;word to clear display
        OUT CNTR, AL        ;send it to control register
        MOV AL, 90H         ;word to write display RAM
        OUT CNTR, AL        ;send it to control register
        MOV AL, 0FFH        ;data to blank out the display
AGN:    OUT DATR, AL         ;send it to the data register
        LOOP AGN            ;repeat for 8 digits
NEXT:   IN AL, CNTR         ;read in the status register
        TEST AL, 07         ;test if  $D_0 = 1$ 
        JZ NEXT            ;if not, continue testing
        MOV AL, 40H         ;otherwise write the word to read FIFO
        OUT CNTR, AL        ;send it to control register
        IN AL, DATR         ;read in the FIFO
        AND AL, 0FH         ;mask upper nibble of key code

```

```

MOV BL, AL           ;copy it to BL
MOV BH, 0            ;BH = 0
ADD SI, BX           ;add BX to SI
MOV AL, [SI]         ;take corresponding display code
OUT DATR, AL         ;send it to data register
JMP NEXT            ;continue for more keys

```

This program displays every key that is pressed. After 8 entries, the FIFO drops the earliest entries and continues the displaying of the keys pressed.

10.3 | The Programmable Interrupt Controller (PIC) 8259

10.3.1 | Introduction

Before trying to understand the description and working of the programmable interrupt controller (PIC), it will be worthwhile to read Section 8.1 to get a precise and clear understanding of the mechanism of interrupts in general and also the specific details of the interrupt structure of the 8086. This will go a long way in the understanding of the use and functionality of this chip which is dedicated to the management of the hardware interrupts of the system.

Recollect that there are only two hardware interrupt lines for the 8086. An external device can place its interrupt request on either of these pins. One pin called NMI is a 'vectored' interrupt, in the sense that its interrupt handler is pointed to, by a particular type number and that is '2'. Thus, NMI is a Type 2 interrupt, and since it is non-maskable it is used for important and high priority services. As such, this interrupt pin is not available for general interrupt requests from various peripherals. That leaves just one pin for accepting hardware interrupt requests, and if many peripherals want to use this pin, obviously a number of issues have to be resolved. It is in this context that a dedicated interrupt controller is necessary.

10.3.2 | 8259 – Features

The first interrupt controller 8259 was designed for the 8080 microprocessor only. Later the 8259A was made available that gives support to the x86 processor as well. However, we will use the word 8259 itself for 8259A, because the original 8259 is obsolete now. This is the standard chip which acts as an interrupt manager for a system – depending on the size and requirements of the system, there can be one or more interrupt controllers. See the basic connection between the 8086 and the 8259 in Fig 10.21.

The PIC has eight interrupt request lines IR0 to IR7, on which peripherals can place their interrupt requests. If one interrupt request alone is received on the chip, that request is channeled to the INTR line of the processor. As is the usual sequence, the processor sends back an acknowledge request in the form of the $\overline{INTA}$ signal. Then the PIC (on behalf of the peripheral which has made the request) sends to the 8086 the interrupt type number corresponding to that specific interrupt. This implies that the PIC is to have been programmed to send a specific type number for a particular interrupt request.

Another activity of the chip is to resolve priorities. Suppose interrupt requests arrive at more than one interrupt input of the chip. Which one of these requests should be channeled to the processor? This decision is made by the PIC and it has to be programmed for this. There should also be the provision to change the priorities.




Figure 10.21 | Basic connections between a PIC 8259 and the 8086

Suppose there are more than 8 possible interrupt sources, more PICs can be used, which means that the chip has pins for cascading more of such chips, and a particular chip can be a master or a slave. With these ideas in mind, let us see the finer details of the chip.

10.3.3 | Block Diagram

The Intel 8259A Programmable Interrupt Controller handles up to eight vectored priority interrupts for the CPU. It can cascade up to 64 vectored priority interrupts without additional circuitry. It is packaged in a 28-pin DIP, uses NMOS technology and requires a single 5 V supply. Circuitry is static, requiring no clock input. The 8259A is designed to minimize the software and real time overhead in handling multi-level priority interrupts. It has several modes, permitting optimization for a variety of system requirements.

The block diagram in Fig 10.23 shows the different functional parts of the chip. It has 8 interrupt request lines IR0 to IR7, which are asynchronous inputs lines. An interrupt request is executed by raising an IR input (low to high), and holding it high until it is acknowledged (Edge Triggered Mode), or just by a high level on an IR input (Level Triggered Mode).

Int (Interrupt) This output goes directly to processor interrupt input and causes the processor to send back the $\overline{INTA}$ as an acknowledgement, to the 8259.

Interrupt Request Register (IRR) and In-Service Register (ISR) The interrupts at the IR input lines are handled by two registers in cascade, the Interrupt Request Register (IRR) and the In-Service Register (ISR). The IRR is used to store all the interrupt levels which are requesting service; and the ISR is used to store the interrupt levels which are being serviced.

Priority Resolver This logic block determines the priorities of the bits set in the IRR. The highest priority is selected and strobed into the corresponding bit of the ISR (in service register) during the $\overline{INTA}$ pulse.

Interrupt Mask Register (IMR) The IMR stores the bits which mask the interrupt lines to be masked. The IMR operates on the IRR. Masking of a higher priority input will not affect the interrupt request lines of lower priority.




Figure 10.22 | Functional block diagram of the PIC




Figure 10.23 | Internal block diagram of a programmable interrupt controller (PIC)

CAS0–CAS2 These three bi-directional lines can be used to cascade several such chips to expand the number of interrupts up to 64, in a master or slave mode. The associated three I/O pins (CAS0–2) are outputs when the 8259A is used as a master and are inputs when the 8259A is used as a slave.

$\overline{SP}/\overline{EN}$ This stands for Slave Program/Enable Buffer. This is also a bidirectional pin. In the buffered mode, it can be used to control buffer transceivers ($\overline{EN}$). In other modes, it indicates whether the 8259 is a master ($SP = 1$), or whether it is a slave ($SP = 0$).

R/W LOGIC Just as in the case of any other peripheral chip, the decoding logic on the $\overline{CS}$ and A0 determine the addresses of the registers within the chip, and in conjunction with the $\overline{RD}$ and $\overline{WR}$ signals, the registers can be written into and read from (to get status information).

10.3.4 | Interrupt Sequence for an 8086 Based System

The 8259 has been designed to interface the processors 8080, 8085 as well as the x86 family of microprocessors. The interrupt processing sequence of the 8080 and 8085 is different from that of the x86 processors. Here, we will discuss only the sequence for the x86 processors.

Let us list out the sequence of operations that occur when a peripheral places an interrupt request in an x86–8259 system. (Refer Fig 10.24.)

- 1. One or more of the interrupt request lines (IR0 to IR7) are raised high. The 8 individual interrupt request lines are first passed through the Interrupt Mask Register (IMR) to see if they have been masked or not. If they are masked, then the request is not processed any further. However, if they are not masked, they will register their request with the Interrupt Request Register (IRR) setting the corresponding IRR bit(s).
- 2. The Interrupt Request Register will hold all the requested IRs until they have been dealt with appropriately. The Priority Resolver simply selects the IR of highest priority.
- 3. The 8259 then sends an INTR signal to the CPU.
- 4. The CPU acknowledges the INTR signal and responds with an $\overline{INTA}$ pulse.




Figure 10.24 | Interrupt Sequence and the registers involved

- 5. Upon receiving the processor’s interrupt acknowledge, the IR (Interrupt Request) which the PIC is processing at the time, is stored in the In Service Register (ISR) which, as the name suggests, shows which IR is currently in service. This IR bit is also reset in the Interrupt Request Register (IRR) as it is no longer requesting service but is actually getting service. The 8259 does not drive the data bus during this cycle.
- 6. The 8086 will initiate a second $\overline{INTA}$ pulse. During this pulse, the 8259 releases an 8-bit pointer (type number) onto the data bus where it is read by the CPU and the interrupt service routine of the selected device thus gets executed.
- 7. This completes the interrupt cycle.
- 8. Once the ISR has done everything it needs, it sends an End of Interrupt (EOI) to the PIC, which resets the In-Service Register.

10.3.5 | Using the 8259

The chip must be programmed for our use, and the programming sequence starts with writing the necessary control words into the control register. There are two types of control words:

- i) Initialization Command Words (ICWs) which are needed for starting the operation. These are a sequence of 4 bytes, some of which are optional.
- ii) Operational Control Words (OCWs) are the words which command the 8259 to operate in various interrupt modes.

10.3.6 | Initialization Command Words (ICWs)

Note that there is only one address line i.e., A_0 to communicate with the chip (See Table 10.2). Thus the initialization control words are selected using this pin, and it is obvious that ICW1 uses a port address different from the other three initialization control words. See Fig 10.25 for the initialization flow chart.

10.3.7 | Initialization Sequence

10.3.7.1 | ICW1 (Initialization Command Word1)

For initialization, always ICW1 should be sent first. Whenever a command is issued with $A0 = 0$ and $D4 = 1$, this is interpreted as Initialization Command Word 1 (ICW1). ICW1 starts the initialization sequence. This control word carries the information required to distinguish between ICW2, ICW3 and ICW4. The important thing is that all these words may not be necessary.

Table 10.2 | Value of A_0 for the Different ICWs

$\overline{CS}$	A_0	ICW
0	0	ICW1
0	1	ICW2
		ICW3
		ICW4




Figure 10.25 | Flow chart showing the initialization sequence for the PIC

- D₀ indicates whether ICW4 is necessary.
- D₁ indicates whether the system is operating in a slave or master mode. If D = 1, it is a slave mode and hence ICW3 is not needed.
- D₂ = 0 always for the x86 processor.
- D₃ selects the options of edge or level triggering for the interrupt lines IR0 to IR7.
- D₄ is high always in the ICW1 and D₅ to D₇ is low for x86 processors.

10.3.7.2 | ICW2 (Initialization Command Word2)

This word allows us to decide the type of numbers to be associated with the interrupt input lines of IR0 to IR7 (See Fig 10.25). This word is sent with A0 = 1. The bits D₂ to D₀ vary from 000 to 111, and together with the upper bits D₇ to D₃, decide the type numbers of the interrupts (IR0 to IR7). Thus we have to write the bits D₇ to D₀. Figure 10.27 shows how this is done.




Figure 10.26 | Format of the initialization command words 1 and 2

Content of Interrupt Vector Byte
for 8086 System Mode

	D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
IR7	T7	T6	T5	T4	T3	1	1	1
IR6	T7	T6	T5	T4	T3	1	1	0
IR5	T7	T6	T5	T4	T3	1	0	1
IR4	T7	T6	T5	T4	T3	1	0	0
IR3	T7	T6	T5	T4	T3	0	1	1
IR2	T7	T6	T5	T4	T3	0	1	0
IR1	T7	T6	T5	T4	T3	0	0	1
IR0	T7	T6	T5	T4	T3	0	0	0

Figure 10.27 | The method of deciding interrupt type numbers for each IR

Example 10.18

Design the initialization control word 2 (ICW2) for a system to associate type numbers 78H onwards for the interrupt inputs IR0 to IR7.

Solution

The method is that ICW2 should be assigned the value corresponding to the type number of IR0. Here, for this to be done, ICW2 should be designed with D₇ to D₃ to be 01111. Thus ICW2 can be

0	1	1	1	1	0	0	0
---	---	---	---	---	---	---	---

i.e., 78H

This will assign the type number 78H to IR0, 79H to IR1 and so on. The idea is that the type number assigned to IR0 should have D₂ to D₀ to be 000. It cannot be assigned a type, say 76H because D₂ to D₀ are 110.

10.3.7.3 | ICW3 (Initialization Control Word3)

This control word is necessary only when the two or more 8259s are cascaded. A single 8259 can be connected to 8 slave chips thus giving the possibility of upto 64 hardware interrupts. In cascade modes, there are separate ICW3 words for the master and the slave. For the master, this word specifies which IR input has a slave connected and another ICW3 informs the slave which IR input of the master is connected to it. See Fig. 10.27.

10.3.7.4 | Initialization Command Word 4 (ICW4)

D₀ indicates whether the processor is 8080/85 (D₀ = 0) or x86 (D₀ = 1).
D₁ specifies whether an AEOI (automatic end of interrupt) will be used (D₁ = 1) or whether EOI commands should be deliberately inserted at the end of an interrupt service routine (D₁ = 0).
D₂ and D₃ specify the options possible when the data bus has bi-directional transceivers (Section 6.1.4.2) and when it does not.
D₄ indicates the use of the special fully nested mode (D₄ = 1). This mode is to be used when the chip acts as a master.
D₅ to D₇ are always '0' in this control word.

Example 10.19

Assume that address decoding is done such that when ICW1 is to be used, the port address is C0H and for the other ICWs, the address is C2H.

Write the initialization instruction for setting up the system with the following specifications:

- i) The system has a single 8259, with edge triggered interrupt inputs and interrupt type numbers from 90H, 91H and 92H for IR0 to IR2.
- ii) It is also to operate in slave buffered mode with normal EOI.

Solution

First the ICWs are to be designed. It is designed with D₀ = 1 i.e., ICW4 is needed, so as to be able to use the second set of specifications (i.e., slave buffered mode with normal EOI) ICW1 is as follows: (Ref Fig 10.26)

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
0	0	0	1	0	0	1	1

i.e., 13H

ICW2 is designed with D₇ to D₃ to be 10010 for specifying the interrupt type numbers.

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
1	0	0	1	0	0	0	0

i.e., 90H

Since the system has only a single 8259, the control word ICW3 is not needed. ICW4

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
0	0	0	0	1	0	0	1

i.e., 09

The port address to send ICW1 is C0H and for the other ICWs is C2H.

Program

```
MOV AL, 13H           ;value of ICW1
OUT 0C0H, AL          ;send it to port C0H
MOV AL, 90H           ;value of ICW2
OUT 0C2H, AL          ;send it to port C2H
MOV AL, 09            ;value of ICW4
OUT 0C2H, AL          ;send it to port C2H
```

10.3.8 | Operational Command Words

Now that we know how to initialize the chip, the next thing will be to understand its different operating modes. For this, it is necessary to use the control words called ‘Operational Command Words’. We will discuss the modes and features of the chip step by step, along with the operational command words that are used for these modes. The OCWs are sent with A₀ = 1 for OCW1 and A₀ = 0 for OCW1 and OCW2. See Table 10.3.

10.3.8.1 | OCW 1 (Operational Command Word 1)

This word is used to mask any of the interrupt requests IR0–IR1. For masking, the specific bit is ‘1’. See Fig 10.28. Note that this word has the interrupt mask structure. To know the current

Table 10.3 | Value of A₀ for the Different OCWs

$\overline{CS}$	A ₀	OCW
0	0	OCW2
0	1	OCW3
		OCW1




Figure 10.28 | ICW4 and ICW3

mask status, OCW1 can be read. For example, to mask the interrupt requests on IR0, IR3 and IR7, OCW1 has to have the value.

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
1	0	0	0	1	0	0	1

i.e., 89H

Example 10.20

Write instructions to read the current status of the IMR. After that, change the masking structure such that only IR0 is masked.

Solution

The status of the IMR is to be read, read the status of OCW1. The instruction is,

```
IN AL, 0C2H      ;C2H is the address we have used in
                  Example 10.19 and it is the same as that of
                  ICW2 (with A0 = 1)
MOV AL, 01       ;mask bit for IR0 is set
MOV BL, AL
OUT 0C2H, AL     ;send to OCW1
```

10.3.8.2 | OCW2 (Operational Control Word2)

This word is used to assign priorities to the interrupt requests. There is a default priority assigned on initialization of the chip. This can be changed dynamically by the bit patterns of OCW2. Refer Fig 10.29 for the bit patterns for the various priority assignment schemes. Also, let us have a look at the various priority options available.


Fully Nested Mode This mode is entered after initialization until another mode is programmed. The interrupt requests are ordered in priority from 0 through 7 (0 highest). When an interrupt is acknowledged, the highest priority request is determined and its vector placed on the bus. Additionally, a bit of the In-Service register is set. This bit remains set until the microprocessor issues an End of Interrupt (EOI). When an interrupt is being serviced, all lower priority interrupts are inhibited, but higher priority interrupts can be acknowledged, provided the IF is ensured to be re-enabled using software.

After the initialization sequence, IR0 has the highest priority and IR7 the lowest. Thus, if IR3 and IR4 requests occur simultaneously, IR3 will be serviced first. However, IR3 can be interrupted by a request on IR2, provided the interrupt flag has been enabled by software (IF = 1).


Automatic Rotation mode This mode is useful when a just and equitable distribution of service is meant to be allotted for devices which may be considered to be of equal importance, and thus of equal priority. The idea is that, once an interrupt for a device is serviced, it goes to the status of lowest priority, until all other requests are serviced.

Specific Rotation Mode This mode has similarities to the previous mode, in that the one that has been serviced will be made to go to the lowest priority status. The difference is that the priorities can be

OCW 1



OCW 2



OCW 3




Figure 10.29 | Operational Control Words 1 to 3

made to follow a specific sequence. The programmer can change priorities by programming the bottom priority and thus fixing all other priorities i.e., if IR3 is programmed as the bottom priority device, then IR4 will have the highest priority device connected to it. Priorities are now IR4, 5, 6, 7, 0, 1, 2, 3 ... in that order of decreasing levels.

In Fig 10.29, D₂ to D₀ are used to assign new priority schemes – if the highest priority is to be IR4, D₂–D₀ should be 100. D₄ to D₃ must always be 0 for OCW2. D₇ to D₅ can have values according to the options needed.

Example 10.21

Find the value of OCW2 for the following options:

- i) Bottom priority level IR5.
- ii) Specific EOI command with rotating priority.

Solution

The word is

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
1	1	1	0	0	1	0	1

i.e., E4H

D₇ to D₅ = 111, because of using the mode ‘specific rotation’

D₂ to D₀ = 101, because IR5 is the bottom priority.

EOI Command This is the ‘end of interrupt’ command. At the end of any ISR, this command should be present, so that the bit corresponding to that interrupt request level in the In-Service Register (ISR) is reset. To clarify this point, remember that when an interrupt on line IR2, say, it taken up for servicing, a bit in the in-service register is set to indicate the ‘interrupt servicing’ status of that request. Once that is over and done with, the bit is to be reset and this is done only on receiving an EOI command. If this is not done, new interrupt requests cannot be placed on this line.

10.3.8.3 | OCW3 (Operational Control Word 3)

This word can be read to get the status of the In-Service register and the Interrupt Request Register. Also it is used to enable some high level features of the chip.

Example 10.22

Write a program to read the contents of the following registers of the 8259.

- i) IMR
- ii) IRR
- iii) ISR.

Assume that initialization has been done earlier.

Solution

```

IN AL, 0C2H          ;to read IMR, READ OCW1
MOV AL, 6AH          ;write OCW3 for reading IRR
IN AL, 0C0H          ;read IRR
MOV BL, AL           ;save IRR in BL
MOV AL, 69H          ;write OCW3 for reading ISR
IN AL, 0C0H          ;read ISR
MOV CL, AL           ;save it in CL

```

Note The 8259 is a complex chip and for finer details, the data sheet of the chip is to be referred to.

10.4 | Cascade Mode

The 8259 can be easily interconnected in a system of one master with up to eight slaves to handle up to 64 priority levels. The master controls the slaves through the 3 line cascade bus. The cascade bus acts like chip selects to the slaves during the $\overline{INTA}$ sequence.

In a cascade configuration, the slave interrupt outputs are connected to the master interrupt request inputs. When a slave request line is activated and afterwards acknowledged, the master will enable the corresponding slave to release the type number corresponding to the device interrupt handler.

The cascade bus lines are normally low and will contain the slave address code from the trailing edge of the first $\overline{INTA}$ pulse to the trailing edge of the third pulse. Each 8259 in the system must follow a separate initialization sequence and can be programmed to work in a different mode. An EOI command must be issued twice: once for the master and once for the corresponding slave. An address decoder is required to activate the Chip Select $\overline{CS}$ input of each 8259. The cascade lines of the Master 8259 are activated only for slave inputs, non-slave inputs leave the cascade line inactive (low). Figure 10.30 shows three 8259s cascaded together.

10.4.1 | Programmable Interrupt Controllers in the PC

In Table 8.2, a list of some of the standard PC interrupts was seen. Some of them are software interrupts and some are hardware interrupts. The hardware interrupts are routed through an 8259 and each one is denoted as an IRQ (Interrupt Request). In the IBM PC and the original PC-XT, there was only one 8259 chip, and thus only eight IRQ lines. However, users soon needed more devices to be interrupt driven, and thus an additional 7 IRQ's were added to the PC. This involved attaching another 8259 to the existing one. Compatibility always causes problems as the new configuration still had to be compatible with old hardware and software. The new configuration is shown in Fig 10.31.

The CPU only has one interrupt line, thus the second controller had to be connected to the first controller, in a master/slave configuration. IRQ2 was selected for this. By using IRQ2 for the second controller, no other devices could use IRQ2. What happened to all these devices using IRQ2? Nothing, the interrupt request line found on the bus, was simply diverted into the IRQ9 input. As no devices yet used the second PIC or IRQ9, this could be done.

As shown in Fig 10.31, IRQ0 through IRQ7 are the master 8259's interrupt lines, while IRQ8 through IRQ15 are the slave 8259's interrupt lines. The actual names on the pins on an 8259 are IR0 through IR7. IRQ0 through IRQ15 are the names of the ISA bus lines to




Figure 10.30 | Cascaded connection of three PICs

which the 8259's are attached. The devices connected to the IRQ lines are shown below. This is standard for the PC-AT.

Master 8259

- IRQ0 – Intel 8253 or Intel 8254 Programmable Interval Timer i.e., the system timer
- IRQ1 – Intel 8042 keyboard controller
- IRQ2 – not assigned in PC/XT; cascaded to slave 8259 INT line in PC/AT
- IRQ3 – 8250 UART serial port COM2 and COM4
- IRQ4 – 8250 UART serial port COM1 and COM3
- IRQ5 – hard disk controller in PC/XT; Intel 8255 parallel port LPT2 in PC/AT
- IRQ6 – Intel 82072A floppy disk controller
- IRQ7 – Intel 8255 parallel port LPT1/spurious interrupt




Figure 10.31 | PICs used in the PC-AT

Slave 8259 (PC/AT and later only)

- IRQ8 – real-time clock (RTC)
- IRQ9 – no common assignment
- IRQ10 – no common assignment
- IRQ11 – no common assignment
- IRQ12 – Intel 8042 PS/2 mouse controller
- IRQ13 – math coprocessor
- IRQ14 – hard disk controller 1
- IRQ15 – hard disk controller 2

Initially IRQ7 was a common choice for the use of a **sound card**, but later IRQ5 was used when it was found that IRQ7 would interfere with the printer port (LPT1). The **serial ports** are frequently disabled to free an IRQ line for another device.

KEY POINTS OF THIS CHAPTER

- The programmable interval timer chips 8253 and 8254 are similar in most ways except that the latter can be used at higher operating frequencies.
- The timer chip has three internal counters each of which has a clock input, a gate and an output pin.
- The count in any count-register can be read 'on the fly' by latching it into a latch.
- There are six operating modes for the timer, and writing suitable words into the control registers will enable any of these modes.
- The 8279 is a keyboard display interface chip which does automatic keyboard scanning as well as refreshing digits in a dynamic display.
- The 8279 has a display RAM into which a seven segment code can be written for displaying a character.
- The 8279 also has a FIFO RAM which stores the key codes corresponding to the keys pressed.
- A number of control words are needed to operate this chip.
- The third interfacing chip discussed in this chapter is the programmable interrupt controller 8259.
- This chip is designed to handle all the issues related to hardware interrupts arriving at the INTR pin of the 8086.
- It has a number of internal registers to handle interrupt requests, interrupt masks and so on.
- A number of initialization command words and a number of operational command words are needed to use this chip.
- In the current PC, two PICs are used to handle standard hardware interrupts.

QUESTIONS

1. In what way is the 8253 different from the 8254?
2. In which mode can symmetric square waves be obtained from the 8253?
3. What is meant by reading a count 'on the fly' and how is it done?
4. Explain the use of mode 1 of the timer chip.
5. What is the maximum BCD number that can be used as a count?
6. What are the actions performed by the 8279 to interface it to a dynamic display?
7. What are the actions performed by 8279 to let it be interfaced to a keyboard?
8. Differentiate between left and right entry display.
9. What information does a 'key code' hold?
10. What is the role of the decoder 74LS138 in the figure 10–17?
11. Briefly explain how the 8259 channels an interrupt request from a peripheral to the 8086.


12. Discuss the roles of the following registers in handling multiple interrupt requests:
 - a) IRR
 - b) IMR
 - c) ISR
13. How is a type vector installed for a particular interrupt request?
14. How can interrupt priorities be changed?
15. Name two features which can be set using operational command words.

EXERCISE

1. Using the 8253 chip, generate square waves of frequencies 1 KHz, 12 KHz and 500 Hz assuming that the input frequency is 2 MHz.
2. Using any two modes of the chip, generate asymmetric square waves.
3. Write a program with a count of 8888H loaded into count-register 2, and draw the output waveform for the waveform, in Mode 1. What additional condition is required to facilitate the functioning of this mode?
4. Use an 8254 timer chip with a clock of 10 MHz, and program it to operate in Mode 2, with three different values of N. Draw corresponding output waveforms.
5. Referring to Fig 10.14, write a program to display ABCD on the display.
6. Referring to Fig 10.14, write a program to get a rolling display.
7. Interface a keyboard to the 8279 as in Fig 10.17, write a program that allows a key press and displays the corresponding key. Now modify it such that when 1 is pressed, 6 is displayed, when 2 is pressed, 7 is displayed and similarly for all the keys.
8. For an 8259 PIC, design an ICW to fix up interrupt type numbers of your choice for the interrupt requests.
9. Draw a decoding circuit for an 8259, and write the corresponding ICW to make all the interrupt requests to be level triggered.
10. Write instructions to read the contents of the ISR and IRR.

11

PERIPHERAL INTERFACING - III



In this chapter, you will learn

- The principles involved in serial communications.
- The terminology and standards of serial communications.
- RS-232 pins and connections.
- The use of the USART 8251 as an interface to computers.
- The modes and programming of 8251.
- The concept of direct memory access.
- The necessity and functions of a DMA controller in a computer system.
- Different types and modes of DMA.
- How to use the DMA controller IC 8237.

11.1 | Serial Communication Principles

Transferring data is what we mean when we use the word ‘communication’. Data is to be sent from the source to the destination, and it is necessary for the source and destination formats to be similar, for compatibility between them to be ensured.

We have seen such data transfer in many contexts and in all these cases, a number of bits are sent together i.e., 8 bits, 16 bits or 32 bits. This is ‘parallel communication’. All the bits are sent and received together. Data transfer between registers in a processor is done this way. This is fine as long as the source and the destination are in close proximity, but when they are placed far apart physically, like, say two computers in two different buildings, a lot of problems crop up. If we are sending 8 bits, eight long wires would be required, and the problem aggravates as the bit size of the data increases. To circumvent this and other related problems, ‘serial communication’ is chosen in preference to the parallel communication just discussed. Serial communication is when we send digital data one bit at a time, one after the other. Thus 8 bits need 8 times the time required, compared to the previous case (assuming the same sending rate), but the direct advantage is that only one physical wire is required for transmission. This is a very great advantage when dealing with larger data word sizes, but along with this, a new set of issues arrive, which need to be dealt with effectively.

First, let us discuss the various terms likely to be encountered while learning serial communication.

11.2 | Simplex, Half Duplex and Full Duplex Communication

A Simplex Connection is a connection in which the data flows in only one direction, from the transmitter to the receiver. This type of connection is useful if data does not need to flow

in both directions (for example, from the computer to the printer or from the mouse to the computer.)

A Half-duplex Connection (sometimes called an alternating or semi-duplex connection) is a connection in which data flows in one direction or the other, but not both at the same time. At a time, transmission is done only in one direction and thus the full bandwidth of the line can be used for the transmission (like when talking on a radio).

A Full-duplex Connection is a connection in which data flows in both directions simultaneously. Each end of the line can thus transmit and receive at the same time, which means that the bandwidth is divided by two for each direction of data transmission if the same transmission medium is used for both directions. Talking over a telephone line is a typical case of full duplex transmission. The serial port on the PC is a full-duplex device meaning that it can send and receive data at the same time. In order to be able to do this, it uses separate lines for transmitting and receiving data.

11.2.1 | Synchronous vs Asynchronous Communications

Now, let us assume we have a block of bytes to be transmitted and that the transmitter and receiver are connected by a serial line. Serial data transfer depends on accurate timing in order to differentiate the bits in the data stream. This timing can be handled in one of two ways: asynchronously or synchronously. **In asynchronous communication, the scope of the timing is a single byte at the maximum. In synchronous communications, the timing scope comprises one or more blocks of bytes.** In both cases, the transmitter and the receiver should know the state of each other. In synchronous transmission, synchronizing information is sent between the transmission of each block of data. The size of this 'block' depends on the system. Even when real data is not being sent, a constant flow of bits allows each device (transmitter or receiver) to know the state of the other at any given time. That is, each character that is sent is either actual data or an idle character. Synchronous communications allows faster data transfer rates than asynchronous methods, because additional bits to mark the beginning and end of each data byte are not required, but only between the large blocks.

11.2.2 | Asynchronous Communications

In asynchronous transmission, the beginning and end of each byte of data must be identified by start and stop bits. The start bits indicate when the data byte is about to begin and the stop bits signal when it ends. The requirement to send these additional bits causes asynchronous communication to be slightly slower than the synchronous type.

An asynchronous line that is idle is identified with a value of 1 (also called a mark state). By using this value to indicate that no data is currently being sent, the devices are able to distinguish between an idle state and a disconnected line. When a character is about to be transmitted, a start bit is sent. A start bit has a value of 0 (also called a space state). Thus, when the line switches from a value of 1 to a value of 0, the receiver is alerted that a data character is about to be sent. Asynchronous communication is the prevailing standard in the personal computer industry, both because it is easier to implement and because it has the unique advantage that bytes can be sent whenever they are ready, as opposed to waiting for blocks of data to accumulate.

We will discuss this mode of communication in more detail. When the transmitter intends to send a byte, it sends a start bit to the idle line which is at the 'mark' or '1' state. The start bit is a 'space' character i.e., a '0' bit. This indicates that what follows is the actual data byte. **Note that the LSB is sent first.** After all the eight bits are sent, one or two stop bits are sent. The type and




Figure 11.1 | A typical asynchronous character frame

and number of stop bits must be defined for the system. The total character frame includes the data and the start and stop bits and parity bits for error checking and correction. A bit time can be defined as the time given for the transmission of one-bit, and the receiver samples the received voltage level at periodic intervals known as the bit time, to determine whether a 0-bit or a 1-bit is present on the line. Observe Fig 11.1. Note that one character can be 8 bits or less. A parity bit may be used for error control. The parity bit in each character can be set to none (N), odd (O) or even (E). However, the practice now is to have no parity bit, but instead to have some other communication error control protocol to be used. The data sent using this method, is said to be framed. That is, the data is framed between the Start and Stop bits. **In transmission, the start bit is sent first, then the data bits with LSB first, followed by the parity and stop bits.**

11.2.3 | Transmission Rate

There are two terms in common use, relating to transmission. One is bits per second. This is what we have seen in our discussion above. If the bps is specified to be 1200, it means 1200 bits are sent per second.

There is another term that is commonly used in serial communications, and that is baud rate. Many people use these words interchangeably, which is not correct, but to understand the difference, we need to discuss the equipment called 'modem'. For the case of direct serial data transmission from one computer to another in simple binary form, the baud rate and bps are the same.

11.2.4 | Modems

The need to communicate between distant computers led to the use of the existing phone network for data transmission. Most phone lines were designed to transmit analogue information – voice, while computers and their devices work in digital form – using pulses. So, in order to use an analog medium, a converter between the two systems is needed. This converter is the **modem**, which performs Modulation and Demodulation of data. It accepts serial binary pulses from a device and modulates some property (amplitude, frequency, or phase) of an analog signal, in order to send the signal in an analog medium. It also performs the opposite process, enabling the analog information to arrive as digital pulses at the computer or device on the other side of connection.



Figure 11.2 | A computer to computer communication system using a modem

Wireless modems convert digital data into radio signals and back. Early modems used frequency shift keying wherein one frequency was used for digital ‘1’ and another frequency for ‘0’. In order to create faster modems, modem designers had to use techniques far more sophisticated than frequency-shift keying. First they moved to phase-shift keying (PSK), and then to quadrature amplitude modulation (QAM).

11.2.5 | Baud Rate vs Bps

The difference between the two is complicated and intertwining. They are dependent and inter-related. However, the simplest explanation is that a bit rate is the number of data bits transmitted per second. Baud rate is the measurement of the number of times per second a signal in a communications channel changes per second.

Bit rates measure the number of data bits (i.e., 0’s and 1’s) transmitted in one second in a communication channel. A figure of 2400 bits per second means 2400 zeros or ones can be transmitted in one second, hence the abbreviation **bps**.

A baud rate, by definition, means the number of times a signal in a communications channel changes state or varies. For example, a baud rate of 2400 means that the channel can change states up to 2400 times per second. The term ‘change state’, means that it can change from 0 to 1 or from 1 to 0 up to X (in this case, 2400) times per second. It also refers to the actual state of the connection, such as voltage, frequency or phase level.

The main difference between the two is that one change of state of the communication signal can transmit one bit – or more or less than one bit and that depends on the modulation technique used. So the bit rate (bps) and baud rate (baud per second) have this connection:

$$\text{bps} = \text{bauds per second} \times \text{the number of bits per baud}$$

The number of bits per baud is determined by the modulation technique. See these examples. When FSK (‘Frequency Shift Keying’, a transmission technique) is used, each baud transmits one bit; only one change in state is required to send a bit. Thus, the modem’s bps rate is equal to the baud rate.

When we use a baud rate of 2400, and use a modulation technique called quadrature phase shift keying (QPSK) that transmits four bits per baud, we get the result that $2400 \text{ baud/second} \times 4 \text{ bits per baud} = 9600 \text{ bps}$. Such modems are capable of 9600 bps operation.

11.2.6 | RS-232 Standards

We have been talking about sending data as bits as either ‘1’ or ‘0’. As per TTL standards, these levels correspond to 5 V and 0 volts. However, during transmission over short distances (without a modem), say from one room to another, the TTL level signals will get easily corrupted by noise.

However, this problem does not occur normally, because we don’t send or receive at TTL levels, instead we use a standard called RS-232C, which defines a different set of voltages/ currents. RS-232 stands for Recommend Standard number 232 and C is the latest revision of

the standard. By this standard, before being transmitted, the TTL voltage levels are changed to a level between -3 to -25 V for a '1', and $+3$ to $+25$ V for a '0' (the voltages between -3 V and $+3$ V is undefined). This means that before sending, the bits should be changed to this level and reconverted to TTL levels on receiving. There are some standard chips available for doing this.

11.2.7 | RS-232 Level Converters

Almost all digital devices that we use require either TTL or CMOS logic levels. Two common RS-232 Level Converters are the 1488 RS-232 Driver and the 1489 RS-232 Receiver. Each package contains 4 inverters of one type, either Drivers or Receivers. See Fig 11.3 which shows how and where these converters are used, to connect two computers in close proximity, using an RS-232 cable. A better IC for the same purpose is the MAX232 which can be used for conversion in both directions (see Fig 11.4).

11.2.8 | RS-232 Connectors

The serial ports on most computers use a subset of the RS-232 standard. The full RS-232 standard specifies a 25-pin D connector of which 22 pins are used. Most of these pins are not needed for normal PC communications, and indeed, most new PCs are equipped with male D type connectors having only 9 pins. Figure 11.5 shows the DB-25 and DB-9 connectors and Table 11.1 and 11.2 describe the pin functions of these connectors.

11.2.9 | DCE and DTE Devices

The two terms that are commonly seen in serial communication literature are DTE and DCE. DTE stands for Data Terminal Equipment, and DCE stands for Data Communications Equipment. These terms are used to indicate the pin-out for the connectors on a device and the direction of the signals on the pins. The computer and terminals that send and receive data are DTE devices, while other devices like the modem that transfer data are usually DCE devices. To avoid confusion, remember that a typical DTE is a PC and a typical DCE is a modem.




Figure 11.3 | Serial communication using RS 232 cable and line drivers/receivers




Figure 11.4 | MAX 232 chip pin diagram and a typical connection




Figure 11.5 | RS-232 connectors **a** DB-25 male connector **b** DB-9 male connector

The RS-232 standard states that DTE devices use a 25-pin male connector, and DCE devices use a 25-pin female connector. It is possible therefore to connect a DTE device to a DCE using a straight pin-for-pin connection as in Fig 11.6.

However, to connect two like devices, say a DTE with another DTE device, like two PCs, a null modem cable is to be used. What is a null modem? See Fig 11.7. Null modem cables

Table 11.1 | Pin Functions of the DB-25 Connector

Pin No.	Description	Pin No.	
1	Protective ground	13	Secondary clear to send
2	Transmitted data (T × D)	14	Secondary transmitted data
3	Received data (R × D)	15	Transmit signal element timing
4	Request to send ($\overline{RTS}$)	16	Secondary received data
5	Clear to send ($\overline{CTS}$)	17	Receive signal element timing
6	Data set ready ($\overline{DSR}$)	18	Unassigned
7	Signal ground	19	Secondary request to send
8	Data carrier detect ($\overline{DCD}$)	20	Data terminal ready ($\overline{DTR}$)
9	Data test pin	21	Signal quality detector
10	Data test pin	22	Ring indicator
11	Unassigned	23	Data signal rate select
12	Secondary data carrier detect	24	Transmit signal element timing
		25	Unassigned

Table 11.2 | Pin Functions of the DB-9 Connector

Pin No.	Description
1	Data carrier detect ($\overline{DCD}$)
2	Received data (R × D)
3	Transmitted data (T × D)
4	Data terminal ready ($\overline{DTR}$)
5	Signal ground (GND)
6	Data set ready ($\overline{DSR}$)
7	Request to send ($\overline{RTS}$)
8	Clear to send ($\overline{CTS}$)
9	Ring indicator (RI)




Figure 11.6 | Typical DTE–DCE connection




Figure 11.7 | DTE–DTE connection using a null modem connection




Figure 11.8a | DTE–DCE connection using 3 pins **b** DTE–DTE connection with a null modem

cross the transmit and receive lines in the cable, so that the transmit pin of one is connected to the receive pin of the other and so on. In addition to transmit and receive, the signals $\overline{DTR}$ and $\overline{DSR}$, as well as $\overline{RTS}$ and $\overline{CTS}$ are also crossed in a null modem connection.

We can say that serial communication using RS-232 standards is possible using 25 pins, 9 pins and also just three pins (See Fig 11.8). In the last case, transmission and reception is done without additional handshaking. $T \times D$ is the line for transmitted data and $R \times D$ for received data.

11.2.10 | Handshaking in the RS232C Protocol

A number of signals seem to be involved in the RS232C protocol, some of which are control signals. They are used for handshaking and let us see the sequence in which these signals occur. Assume that a DTE equipment (for example, a COM port of a PC) and a DCE equipment (serial port of a modem) are to perform data transfer. The handshaking signals involved in this data transfer are as follows. Refer to Fig 11.6 to understand the direction of the signals.

$\overline{DTR}$ (Data Terminal Ready) When the DTE terminal is turned on, after its self test, it sends out the $\overline{DTR}$ indicating the robustness of the COM port and its readiness for communications.

$\overline{DSR}$ (Data Set Ready) When the DCE terminal is turned on, it sends out the $\overline{DSR}$ signal after its self-test. In this case, the modem sends this signal to the PC's Com port, which also indicates that it has made a connection to the network (telephone or wireless network). Together $\overline{DTR}$ and $\overline{DSR}$ are used to confirm that both devices are connected and turned on.

RTS (Request to Send) Now that both the equipments are ready, data transmission can occur. If the DTE has a byte to send, it asserts the RTS signal, which is received by the DCE equipment (modem in this case).

CTS (Clear to Send) The receiving equipment (the modem in this case) responds to the RTS signal by this signal. CTS makes it known to the transmitting equipment that buffer space is available for the incoming data. Together RTS and CTS are used for hardware flow control, it regulates the flow of data. If no buffer space is available in the receiver side, the transmission must wait until the buffer is cleared.

DCD (Data Carrier Detect) This is sent by the modem to indicate that a valid carrier has been detected.

RI (Ring Indicator) This is an output from the modem that the telephone is ringing. The modem toggles (goes on and off) the state of this line when an incoming call rings the phone. The Data Carrier Detect (DCD) and the Ring Indicator (RI) lines are only available in connections to a modem. Because most modems transmit status information to a PC when either a carrier signal is detected (i.e., when a connection is made to another modem) or when the line is ringing, these two lines are rarely used. The RI signal is not shown in Fig 11.6.

11.3 | The Programmable Serial Communication Interface

Since we have had an overview of serial communications, let us dig a bit deeper. We know that within the processor and between the processor and memory of most I/O devices, data is transferred in parallel. So, if we need to take it out from such a system in serial form, it is obvious that a parallel to serial conversion is necessary. Similarly, if serial data is received from a serial line, it needs to be converted to parallel form before the processor can process it. Another point is that all the serial communication rules and protocols must be ensured before we try to send and receive serial data. All these matters will have to be taken care of by a dedicated hardware and there are several specialized ICs available for this purpose. Listing out a few, we have:

8250 – Universal Asynchronous Receiver Transmitter (UART)

8251 – Universal Synchronous Asynchronous Receiver Transmitter (USART)

16550 – Universal Asynchronous Receiver Transmitter (UART)

Now, let us try to understand the features of the USART 8251.

11.3.1 | Universal Asynchronous/Synchronous Receiver/Transmitter – USART 8251

The 8251 USART is designed to be directly compatible with the x86 processors as well with the other Intel family processors. It caters to synchronous as well as asynchronous transmission. It can accept data from the processor in parallel form and convert it into a continuous serial stream for transmission. Simultaneously, it can also receive serial data streams and convert them to parallel format for use by the processor. It signals the processor whenever it is ready to accept a new character for transmission and also whenever it has received a character for the CPU. It inserts and deletes the framing bits as necessary in the communication mode it is set. In effect, the intricacies of the communication protocols being used are not felt by the programmer beyond the difficulty of writing into command registers. Additionally, it allows the CPU to read its complete status at any time.




Figure 11.9 | Pin diagram of the USART

11.3.2 | Functional Block Diagram

Let us examine the functional block diagram of the chip.

11.3.2.1 | Data Bus Buffer

It is the block to which the data bus (three state bi-directional and 8-bit) of the processor is connected. It is through the data bus that control words and status words and data are sent from the processor to the chip and vice versa.

11.3.2.2 | Transmitter Buffer

This block accepts parallel data from the data bus buffer, converts it to a serial bit stream, inserts the appropriate extra bits (based on the communication technique) and outputs a composite serial stream of data on the T × D (Transmitter Data) output pin on the falling edge of T × C. On reset, the T × D line will be held in the mark state (which corresponds to an idle state). It will begin transmission upon being enabled if CTS = 0. When the transmitter is empty, or if CTS or T × Enable is OFF, the line will once again go to the idle state of 'mark'.

11.3.2.3 | Transmitter Control

This block manages all the activities associated with transmission. It accepts and issues signals required for this. The signals are:

T × RDY (Transmitter Ready) This is an output signal saying that the transmitter is ready to accept a data character to be sent. This can be used to interrupt the processor, or the processor can check its state by reading the status register, in the polled mode of operation.




Figure 11.10 | Internal block diagram of the 8251

$T \times \text{Empty}$ (Transmitter Empty) When the 8251 has no characters to send, this output pin will go high. It is cleared on receiving a character from the processor. It can be used to indicate the end of a transmission, so that next, the CPU can turn around for reception (in the case of half duplex mode, say).

$\overline{T \times C}$ (Transmitter Clock) This is the transmitter clock input pin, which controls the rate of transmission. In the synchronous transmission mode, the baud rate is equal to the $\overline{T \times C}$ frequency. In asynchronous transmission mode, the baud rate is a fraction of this frequency. This is selected by the mode control word. The falling edge of the $\overline{T \times C}$ clock shifts the serial data out of the 8251.

11.3.2.4 | Receive Buffer

The receiver accepts serial data, converts this serial input to parallel format, checks for data that are unique to the communication technique and sends an 'assembled' character to the processor. Serial data is taken as input on the $R \times D$ pin, and is clocked on the rising edge of the $\overline{R \times C}$.

11.3.2.5 | Receiver Control

This functional block manages all receiver related activities. Before starting to receive serial characters on the $R \times D$ line, a valid ‘1’ must first be detected after a chip master reset. Once this is obtained, a search for a valid low (Start bit) is enabled. This is so only in the asynchronous mode, and this is only done once for each master reset. The signals associated with receiver control are:

$R \times RDY$ (Receiver Ready) This output signal indicates that the 8251 contains a character that is ready to be input to the CPU. $R \times RDY$ can be connected to the interrupt structure of the CPU. If polling is being used, the CPU can check the status of this pin using a status read operation.

$\overline{R \times C}$ (Receiver Clock) This input pin controls the rate at which receiving is done. In the synchronous mode, the baud rate is equal to the actual frequency of $\overline{R \times C}$. In the asynchronous mode, the baud rate is a fraction of this frequency. Data is sampled in at the rising edge of the $\overline{R \times C}$ signal.

Note In most cases, the 8251 will be handling the transmission and reception of the same link. Hence the baud rates will be the same and so the $\overline{R \times C}$ and $\overline{T \times C}$ can be tied together to a single frequency source.

$Syndet/Brkdet$ (Sync detect or Break Detect) This is a bidirectional pin, and is used in the synchronous mode as SYNCH DETECT. It is used either as input or output, programmable using the control word. In the asynchronous mode, it is used as an output pin as ‘BREAK DETECT’. It will go high to indicate a ‘break’ or ‘disconnection’. It goes high whenever the receiver remains low through two consecutive stop bit sequences (including start bits, data bits and parity bits). Break detect may also be read as a Status bit.

11.3.2.6 | R/W Control Block

In this block, there are the read $\overline{RD}$, write $\overline{WR}$, chip select $\overline{CS}$, reset and $C/\overline{D}$. Only the last pin needs an explanation.

$C/\overline{D}$ (Control/Data) This is an input signal which, in conjunction with the $\overline{CS}$ can be used for selection of the data register or control/status register.

$\overline{CS}$	$C/\overline{D}$	Register selected
0	0	Data Register
0	1	Mode, command and status register

11.3.2.7 | Modem Control Block

The pins in this block perform the function of interfacing to a modem and they are the control/status signals used therein. The 8251 has a set of control input and outputs that can be used to simplify the interface to any modem. However, these signals are general-purpose in nature and can be used for functions other than modem control. Refer to Section 11.2.10 to know how these signals are important in serial communication, in general.

$\overline{DSR}$ (Data Set Ready) The $\overline{DSR}$ input signal is a general-purpose, one-bit input port. Its condition can be tested by the CPU using a status read operation. The $\overline{DSR}$ input is usually used to test modem conditions such as Data Set Ready.

$\overline{DTR}$ (Data Terminal Ready) The $\overline{DTR}$ output signal is a general-purpose, one-bit output port. It can be set low by programming the appropriate bit in the command instruction word. The $\overline{DTR}$ output signal is normally used for modem control.

$\overline{RTS}$ (Request to Send) This output signal is a general-purpose, one-bit output port. It can be set low by programming the appropriate bit in the Command instruction. The $\overline{RTS}$ output signal is normally used for modem control.

$\overline{CTS}$ (Clear to Send) A low on this input enables the 8251 to transmit serial data if the $T \times \text{Enable}$ bit in the command byte is set to a 'one'. If either a $T \times \text{Enable OFF}$ or $CTS \text{ OFF}$ occurs while the transmitter is in operation, the transmitter will transmit all the data in the USART written prior to the Transmitter Disable command, before shutting down.

11.3.3 | How This Chip Operates

Like any other interfacing chip, the chip functions according to the way we program it. A set of control words must be sent to it to support the communication mode in which it is to operate. Once the chip is programmed for use, it is ready for transmission or reception. For transmission, the $T \times RDY$ output is raised high to tell the CPU that it is ready to transmit the character the CPU sends. On receiving this affirmation, the CPU writes a character into the 8251 and then the $T \times RDY$ pin is reset.

In the case of reception, the 8251 receives a character from a modem or serial I/O device and then it raises the output pin $R \times RDY$ to communicate this to the CPU. The CPU then reads into it the received data and the $R \times RDY$ pin is reset. $\overline{T \times C}$ and $\overline{R \times C}$ are connected to the same clock source. $T \times RDY$ and $R \times RDY$ pins can be used in the interrupt mode or status checking mode (polling) for transmission and reception to occur as outlined above. $T \times D$ is the pin on which serial data is transmitted and $R \times D$ is the pin on which serial data is received.

11.3.4 | Programming the 8251

Prior to starting data transmission or reception, the 8251 must be sent a set of control words. This must be done after an external or internal reset. The control words are split into two formats.

- i) Mode control word
- ii) Command word

Just after reset, the mode word must be sent to the chip. Any word that comes after the mode word is considered to be the command word, and command words can be written any time during the operation of the 8251. To change the mode, we can write a new mode word but this can be done only after an internal or external reset. An internal reset can be initiated by setting the 'reset' bit in the command word (bit D_6).

The USART, as the name implies, can be used for synchronous as well as asynchronous communications. The communication format is different for the two cases and they should be considered separately. We are interested only in the asynchronous mode, and only that will be discussed here. The data sheet of the chip may be referred for details of the synchronous mode.

11.3.5 | Asynchronous Mode

Let us start by writing the control words. Figure 11.11 shows the format of the mode word.




Figure 11.11 | Mode control word

11.3.5.1 | Mode Instruction Format

The mode word, as shown, fixes up the baud rate, number of characters, and stop bits for transmission. Remember that when the line is idle, it is in the mark ('1') state. A start bit is always low. This should be followed by the data bits and then the stop bits. The insertion of start and stop bits along with a character, is called 'framing'. The baud rate is a fraction of the clock frequency ($T \times C$ and $R \times C$ are tied together for a link) and the factor can be 1, 16 or 64. If $T \times C = 9600\text{ Hz}$, the baud rate is 9600 bps for a factor of 1 and $9600/16 = 600$ for a factor of 16. If $D_1 D_0 = 00$, it means 'synchronous operation'. For using the chip in asynchronous mode, these bits should have values other than 00, which fix up the baud rate factor as well.

The character length can be 5, 6, 7 or 8 bits and it does not include the framing bits or parity bits. If the number of characters is less than 8, the least significant data bus bits will hold the data; unused bits are don't cares when writing data into the 8251, and will be zeros when reading. When parity is enabled, it is not considered as one of the data bits, and the parity bit cannot be read on the data bus. The upper most two bits specify the number of stop bits. Figures 11.12 and 11.13 show clearly the transmission and reception formats. It should be clear that a character is framed and sent. On reception, the framing bits are removed and only the 'character' is considered as the received byte.




Figure 11.12 | The framed character at T × D and R × D pins




Figure 11.13 | Transmission and reception formats

Example 11.1

Write the mode control word for a case of asynchronous transmission, with an 8-bit data format, one stop bit, odd parity and 10000 baud rate approximately. The $T \times C$ clock is 150 KHz.

Solution

The mode control word is defined referring to Fig. 11.11.

D_7	D_6	D_5	D_4	D_3	D_2	D_1	D_0
0	1	0	1	1	1	1	0

i.e., 5EH

$D_7 D_6 = 01$ for 1 stop bit

$D_5 D_4 = 01$, odd parity enabled

$D_3 D_2 = 11$, 8-bit character

$D_1 D_0 = 10$ baud rate factor of 16x.

The clock source gives a clock of 150 KHz, the baud rate factor is 16x i.e., 1/16.

Hence the baud rate is $(150 \times 10^3) / 16 = 10,000 = 10 \text{ K}$ approximately.

11.3.5.2 | Command Instruction Format

Figure 11.14 shows the command instruction format.

Example 11.2

For a simple operation, we have to enable transmit and receive, make DTR to be 0, and RTS to be 0 and enable error reset. Design the command word.

Solution

The command word is

D_7	D_6	D_5	D_4	D_3	D_2	D_1	D_0
0	0	1	1	0	1	1	1

i.e., 37H

11.3.5.3 | Status Word

The 8251 has a status word which enables us to read the status of the device during its operation. This can be read as and when necessary and used to perform transmission and reception in the polled mode.

D_7	D_6	D_5	D_4	D_3	D_2	D_1	D_0
DSR	SYNDET	FE	OE	PE	$T \times \text{EMPTY}$	$R \times \text{RDY}$	$T \times \text{RDY}$

Let us examine the important status bits:

$T \times \text{RDY}$ (Transmitter Ready) This indicates that the USART is ready to accept data or command from the processor to which it is connected.




Figure 11.14 | Command instruction format

$R \times RDY$ (Receiver Ready) This indicates that the chip has received a character and is ready to transfer it.

$T \times Empty$ (Transmitter Empty) This means that the transmitter buffer is empty. There are three error flags in the status word but their setting does not inhibit the operation of the 8251. They are reset by the command word.

PE (Parity Error) This bit is set when there is a parity error.

OE (Overrun Error) This flag is set when the processor does not read a character before the next one becomes available. In the case that this bit gets set, the previous overrun character is lost.

FE (Framing Error) This is applicable for asynchronous communication only. This flag is set when a valid stop bit is not detected at the end of every character.

Syndet–Synch Detect This pin is used in synchronous mode for sync detect and is used in asynchronous mode for break detect.

DSR (Data Set Ready) This bit indicates that DSR is at zero level.

11.3.6 | Programming the 8251

Now that we have got an idea of the important aspects of the chip, we can program the chip for transmission and reception.

Loop Back Mode We will first do transmission and reception in the same chip – send a character from the transmitter section of the USART and receive it in the receiver section therein. For that, the following connections are to be done (See Fig. 11.15). The data register has an address of C0H and the control/status an address C2H. The transmitter and receiver clocks are from the same clock source. T × D and R × D are to be interconnected, so also the $\overline{CTS}$ and $\overline{RTS}$ pins. Then, initialize the 8251 with the mode and command words, transmit a character, receive it, and save it in memory.

Example 11.3

Write the program for transmission and reception in the loopback mode.

Solution

The mode control word used is shown here. Parity is disabled.

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
0	1	0	0	1	1	1	0

i.e., 4EH

D₇ D₆ = 01 for 1 stop bit
D₅ D₄ = 00 parity disabled
D₃ D₂ = 11, 8-bit character
D₁ D₀ = 10 baud rate factor of 1/16

```
TRANS DB 'S'           ;the character to be transmitted
RECD DB?                ;space for received character

DATR EQU 0C0H
CSR EQU 0C2H

MOV AL, 4EH             ;mode word
OUT CSR, AL             ;send it to control/status register
MOV AL, 37H             ;command word
OUT CSR, AL             ;send it to control/status register
MOV AL, TRANS           ;move character to AL
OUT DATR, AL            ;send it to data register

IN AL, DATR             ;move received data to AL
MOV RECD, AL            ;save it in memory
END
```




Figure 11.15 | Loop back mode

Example 11.4

Next, we will connect two microprocessor kits serially and perform transmission and reception. The transmission through one kit will be through its RS-232 cable and reception by the other cable will also be similar. Both the kits are clocked from the same source.

Solution

The sequence should be such that the receiver program is run first. The receiver then is in a loop waiting for a start bit to be received. Then run the transmitter program. When the receiver senses the start bit, it starts reception and saves the received character string in memory.

TRANSMITTER PROGRAM

```

DATR    EQU    0C0H
CSR     EQU    0C2H

TRANS DB 'INDIAN'    ;location for character string
LEA BX, TRANS        ;point BX to character string
MOV CX, 06           ;CX= the number of characters
MOV AL, 4EH          ;mode word
OUT CSR, AL          ;send it to control register
MOV AL, 33H          ;command word to enable Tx
OUT CSR, AL          ;send it to control register
REPEA:  IN AL, CSR    ;read in status register
        AND AL, 01    ;test if D0 is 1. i.e., is Tx ready?
        JE REPEA      ;if not '1', keep testing
NXT:    MOV AL, [BX]   ;otherwise get character in AL
        OUT DATR, AL  ;send it to the data register
        INC BX        ;increment pointer
        LOOP NXT      ;continue until CX=0
END

```

RECEIVER PROGRAM

```

DATR    EQU    0C0H
CSR     EQU    0C2H

RECD DB 6 DUP(0)           ;location for received data

LEA SI, RECD                ;point SI to character string
MOV CX, 06                  ;CX=the number of characters
MOV AL, 4EH                 ;mode word
OUT CSR, AL                 ;send it to control register
MOV AL, 36H                 ;command word to enable receiver
OUT CSR, AL                 ;send it to control register
REPEA:  IN AL, CSR          ;read in status register
        AND AL, 02          ;test if D1 is 1.
                                i.e., is Rx ready?
        JE REPEA            ;if not '1' keep testing
NXT:    IN AL, DATR          ;get received character into Al
        MOV [SI], AL        ;move it to memory
        INC SI              ;increment pointer
        LOOP NXT            ;continue until CX = 0
END

```

The salient points of the programs are:

- i) The mode control word for both the transmitter and receiver are the same. CSR is the control /status register and DATR is the data register.
- ii) The command words enable transmission or reception and causes RTS pin to be low. It is sufficient to enable only the transmitter in the transmitter program and the receiver in the receiver program. Thus, the command word for the transmitter is 33H and for the receiver, it is 36H.
- iii) Each of the transmitted bytes are framed and sent. In the receiver, the framed character is received only if the start bit is sensed. This is done automatically by the 8251.
- iv) After the mode word and command word are sent to the control/status register, this register is 'read' to confirm if $T \times RDY$ is high (for transmission) and $R \times RDY$ is high for reception. Only after this is ensured, the data for transmission is sent to the data register (for transmission), and the received byte is read in (in the receiver program).
- v) The characters to be sent are stored in a memory location named TRANS at the transmitter. On reception, they are saved in a location named RECD at the receiver side.

11.4 | Internal Reset on Power Up

When power is first turned on, the 8251 is expected to come to the reset state but sometimes it does not do so. Instead it may come up in the mode, sync character or command format. For normal operation, we need it to be 'reset', then we send the mode word and then the command word. To ensure this to be done, we can get it into the 'command instruction format and from there cause an 'internal reset'. This is done by executing the worst-case initialization sequence (sync mode with two sync characters). Loading three 00s into the control register configures

sync operation and writes two dummy 00 sync characters. An internal reset command may then be issued to bring the chip to the idle state. The following sequence of instructions must precede the program lines of Examples 11.3 and 11.4

```

MOV AL, 00           ;AL = 00
OUT CSR, AL          ;send AL to CSR
OUT CSR, AL          ;send AL to CSR
OUT CSR, AL          ;send AL to CSR
MOV AL, 01000000B    ;command word with D6 = 1
OUT CSR, AL          ;send to CSR for internal reset

```

The Serial Port of the PC

As already mentioned, the PC supports only asynchronous serial communication. It uses the UART chip 8250 for interfacing. This chip is, incidentally, a more powerful chip than the 8251. The 8250 was improved along the years that the processors used in PCs became more and more powerful – now 16450, which is a faster version of 8250 is being used. Recently a better version NS16550 has appeared in the market and is starting to be used widely.

Is the Serial Port on the PC Obsolete?

The serial port used to be found on almost all PC's but not anymore. Thus, as of 2009, it is becoming obsolete, if not already obsolete. It is often called a 'legacy' port, but it is still used for hardware designed to connect to the serial port, especially for computers used as servers by companies. Laptops and Macs stopped being sold with serial ports several years before desktops did. However, if one needs a serial port, it is possible to buy one and install it, such as on the USB bus, and it is still found on older PC's. However, as is seen now, the serial port as well as the parallel port are becoming obsolete, because of being superseded by the USB port.

11.5 | Direct Memory Access

Data transfer between peripherals and memory is a frequent activity in any computer system. There are three methods by which this can be done – we have seen two such methods – polling and interrupt based data transfer. Polling is unacceptable, except for very small systems, because it keeps the processor in a waiting loop. Interrupt based data transfer is fine because only when data transfer is needed, is the processor called into action for this.

Now, there is a third method of data transfer and that is 'direct memory access'. A brief introduction to this was given in Section 6.3.2. The word 'direct' means that data can be transferred between a peripheral and memory without the intervention of the processor – a direct connection is established between these two (one a source and the other, a destination). Refer Fig 6.19. Since the processor is effectively isolated from this data transfer function, it is left free to do other processing activities, but only such activities that do not require the use of the system bus (which is now being used for DMA). Also, since the processor is left without access to the system bus, this sort of data transfer is justified only for transferring large chunks of data. Thus, DMA is used to transfer big blocks of data from/to memory and peripherals. Figure 6.19 is redrawn here to explain the operation of DMA.

HOLD and HLDA are two pins of the processor. Whenever a device wants a DMA operation to be initiated, it places a request on the HOLD pin of the processor. As the timing diagram indicates, HOLD is sampled at the middle of a clock edge. Then the processor completes the current bus cycle (not the current instruction cycle) and enters the hold state after sending a




Figure 11.16 | DMA cycle timings

HLDA signal (Hold Acknowledge) to the requesting device. HOLD has a higher priority than either the INTR or NMI interrupts. Also note that, **interrupts are acknowledged only after the end of the current instruction cycle, but DMA operations start after completing the current bus cycle.**

11.5.1 | DMA Controller

Let us review the concept of DMA. The idea is one of transferring data in one of the following cases.

- i) from memory to a peripheral
- ii) from a peripheral to memory
- iii) from memory to memory (a special case)

It is obvious that there is the need for a controller which gives information to the CPU about the source and destination addresses and also the number of bytes to be transferred. Thus, DMA operations require what is called a DMA controller.

DMA controllers vary as to the type of DMA transfers and the number of DMA channels they support. There are two types of DMA transfers, namely, flyby DMA transfers and fetch-and-deposit DMA transfers. There are also three common transfer modes – single, block, and demand transfer modes. These DMA transfer types and modes are described in the following paragraphs.

11.5.2 | Flyby DMA

The fastest DMA transfer type is referred to as a single-cycle, single-address, or flyby transfer. In a flyby DMA transfer, a single bus operation is used to accomplish the transfer, with data being read from the source and written to the destination simultaneously.

In flyby operation, the device requesting service asserts a DMA request on the appropriate channel request line of the DMA controller. The DMA controller responds by gaining control of the system bus from the CPU and then issuing the pre-programmed memory address. Simultaneously, the DMA controller sends a DMA acknowledge signal to the requesting device. This signal alerts the requesting device to drive the data onto the system data bus or to latch the data from the system bus, depending on the direction of the transfer. In other words, a flyby




Figure 11.17 | Timing associated with flyby DMA operation (DMA write)

DMA transfer looks like a memory read or write cycle with the DMA controller supplying the memory address and the I/O device reading or writing the data.

DMA terminology dictates that a DMA write occurs when data is moved from an I/O device to memory by activating $\overline{MEMWR}$ and $\overline{IORD}$ simultaneously. A DMA read is just the inverse operation when data is read from memory and written to I/O. Note that only the memory address can be placed on the address bus, while the I/O address should be implicit. Because flyby DMA transfers involve a single memory cycle per data transfer, these transfers are very efficient; however, memory-to-memory transfers are not possible in this mode. Figure 11.17 shows the flyby DMA transfer signal protocol.

11.5.3 | Fetch and Deposit DMA

The second type of DMA transfer is referred to as a dual-cycle, dual-address, flow-through, or fetch-and-deposit DMA transfer. As these names imply, this type of transfer involves two memory or I/O cycles. The data being transferred is first read from the I/O device or memory into a temporary data register internal to the DMA controller. The data is then written to the memory or I/O device in the next cycle. Although inefficient because the DMA controller performs two cycles and thus retains the system bus longer, this type of transfer is useful for interfacing devices with different data bus sizes. For example, a DMA controller can perform two 16-bit read operations from one location followed by a 32-bit write operation to another location. A DMA controller supporting this type of transfer has two address registers per channel (source address and destination address) and bus-size registers, in addition to the usual transfer count and control registers. Unlike the flyby operation, this type of DMA transfer is suitable for both memory-to-memory and I/O transfers. Fig 11.18 shows the timing associated with this.

11.5.4 | DMA Transfer Modes

In addition to DMA transfer types, DMA controllers have one or more DMA transfer modes. Single, block, and demand are the most common transfer modes.

Single Transfer Mode In this mode, one data value is transferred for each DMA request assertion. This mode is the slowest method of transfer because it requires the DMA controller to arbitrate for the system bus with each transfer.




Figure 11.18 | Timing for fetch and deposit DMA




Figure 11.19 | DMA concept

Block and Demand Transfer Modes These modes increase system throughput by allowing the DMA controller to perform multiple DMA transfers, once the DMA controller has gained control of the bus. For block mode transfers, the DMA controller performs the entire DMA sequence as specified by the transfer count register at the fastest possible rate, in response to a single DMA request from the I/O device. For demand mode transfers, the DMA controller performs DMA transfers at the fastest possible rate as long as the I/O device asserts its DMA request. When the I/O device de-asserts the DMA request, transfers are signaled off.

DMA Operation There should be a DMA controller to co-ordinate all these activities. The controller is a programmable peripheral chip which can be programmed to act as directed, as we will see in the following discussion. It is not a 'processor' which has execution capability.

However, it can be made to act to control and co-ordinate the DMA requirements of the system. See Fig 11.19. Initially the system works as a normal computer with all the switches in the upper position. In this state, the processor controls the address, data and control buses. Now, say, a peripheral device wants DMA service. It sends a request to the DMA controller on the pin DREQ. The DMA chip consults its various internal registers and issues a HRQ (Hold Request) to the CPU. It gets a HLDA (hold acknowledge) signal from the processor. This signals the start of DMA operations, with the processor being disconnected from the system. Now the switches in Fig 11.19 are toggled to the lower position, giving the DMA chip, the control of the buses. The DMA chip relays this information as a DACK (acknowledge) signal to the peripheral.

After this initial signaling schedule, actual data transfer occurs. After the complete transfer is over, the DMA controller deactivates its hold request (HRQ) and control returns to the processor which means that the switches are back in position 1.

11.5.5 | DMA Controller Features

In a system, there can be a number of I/O devices needing DMA action. Each of them will have to be supported by the DMA controller which should have a number of 'channels' each one catering to one DMA requesting device. For each channel in the DMA controller, certain components are replicated in the controller, while certain others are common to all the channels. Keep in mind that only one DMA operation is possible in the system at a particular time.

Now see Fig 11.20. There may be n channels in the DMA controller of which let the currently active channel be designated as channel 'N'. For DMA transfers, the starting address in memory, and the number of bytes to be transferred have to be specified. For each channel, the DMA controller saves the programmed address and count in the base registers and maintains copies of the information in the current address and current count registers, as shown in the figure.

Each DMA channel is enabled and disabled via a DMA mask register. When DMA is started by writing to the base registers and enabling the DMA channel, the current registers are loaded from the base registers. With each DMA transfer, the value in the current address register is driven onto the address bus, and the current address register is automatically incremented. The current count register determines the number of transfers remaining and is automatically decremented after each transfer. When the value in the current count register goes to 0, a terminal count (TC) signal is generated, which signifies the completion of the DMA transfer sequence.

11.6 | The DMA Controller – 8237

Now that we have discussed DMA operations in general, let us see the features of Intel's 8237A multimode DMA controller chip which has all the features mentioned above.

This is a 40-pin chip with four channels for transferring data, each of which caters to one device. For each channel, there are two signals, DREQ (DMA Request) and DACK (DMA Acknowledge). The former is a request from a peripheral device asking for DMA facility, and the latter is the reply from the controller informing the device that its request has been accepted and acknowledged. From the DMA controller to the processor, there is only one HOLD and HLDA signals, meaning that at least four devices can request DMA services, but only one of them will get service and that is decided by the DMA controller, according to the way priorities are programmed in the chip. However, each channel must be initialized separately for the base address and the count of bytes to be transferred. After this is done, a control word will enable and control the channels. See the functional block diagram of the 8237 chip in Fig 11.21.




Figure 11.20 | A general block diagram of a DMA controller

11.6.1 | Programming the 8237 DMA Controller

Just like any other interfacing chip, this chip also is programmed by writing suitable words into its internal registers. Let us see how this is done. Keep in mind that there are four channels, which need certain information to be supplied to them. There are two registers associated with each channel:

- i) Base address register.
- ii) Base word count register.




Figure 11.21 | Functional block diagram of the 8237 chip

Before DMA operation starts, these 16-bit registers should contain the starting address of memory, and the count of the number of bytes to be transferred. Thus, these registers are to be written into, and the content of these registers do not change during the DMA cycle. However, the content of these registers are also written into (automatically) into two other registers called ‘current address register’ and ‘current count register’. The content of the former is incremented after each transfer, while the content of the latter is decremented. Thus, the current status of the DMA operation can be known by reading these registers. The DMA chip has four pins A3 to A0 to select these eight registers.

Example 11.5

Find the addresses of the base address and base count registers for the four channels, when the decoding logic of Fig 11.22 is used.

Solution

The decoding logic needs A7 to A4 to be 1100 i.e., 0CH for selecting the chip. The other bits for selecting each register is given in column 4 of Table 11.3. Hence the address of the base address register and base count registers for each of the channels is given in column 5 of the table.

Note We ‘write’ into the base address register and base count register. Later, at any time, the current values of address and count are ‘read’ from the current address register and current count register. That should explain Column 3 of Table 11.3.

Table 11.3 | Addresses of the Registers of the Four Channels

Channel	Register name	Operation	A3	A2	A1	A0	Address (Hex)
0	Base Address	Write	0	0	0	0	C0
	Current Address	Read	0	0	0	0	C0
	Base Count	Write	0	0	0	1	C1
	Current Count	Read	0	0	0	1	C1
1	Base Address	Write	0	0	1	0	C2
	Current Address	Read	0	0	1	0	C2
	Base Count	Write	0	0	1	1	C3
	Current Count	Read	0	0	1	1	C3
2	Base Address	Write	0	1	0	0	C4
	Current Address	Read	0	1	0	0	C4
	Base Count	Write	0	1	0	1	C5
	Current Count	Read	0	1	0	1	C5
3	Base Address	Write	0	1	1	0	C6
	Current Address	Read	0	1	1	0	C6
	Base Count	Write	0	1	1	1	C7
	Current Count	Read	0	1	1	1	C7




Figure 11.22 | Address decoding

11.6.2 | Functional Block Diagram

Now observe the functional block diagram in Fig 11.21. Most of the pins seem familiar, but the functions of a few (AEN, ADSTB, $\overline{EOP}$) need more explanation. We will come to that soon. However, some features and some anomalies need to be noted.

- i) The base count register is only 16 bits long implying that only 64KB of data can be transferred by one DMA operation.
- ii) The base address register is only 16 bits in size – but for a processor like 8086, the address is 20 bits long, so how can the upper 4 bits be supplied?

- iii) The DMA chip has only an 8-bit address bus, so how does it place a 20-bit address on the address bus to address memory?

The first point is a feature of the chip, but the second and third are obviously inadequacies of the chip which have to be addressed. Let us get to that.

Example 11.6

Write into the base register and count register of channel 1, given that 4KB of data are to be transferred from a memory location 13230H to a peripheral.

Solution

The address of the base address register and the base count register for channel 1, are C2H and C3H respectively. The number of bytes to be transferred is 4K i.e., 4096 bytes.

MOV AL, 30H	;move the LSB of the base address to AL
OUT 0C2H, AL	;send it to the base address register
MOV AL, 32H	;move the MSB of the base address to AL
OUT 0C2H, AL	;send it to the base address register
MOV AX, 4096	;move the byte count to AX
OUT 0C3H, AL	;send LSB of count to the count register
MOV AL, AH	;move MSB of count to AL
OUT 0C3H, AL	;send MSB of count to the count register

Note that there is no provision to send the upper hex digit of the 20-bit address 13230H to the base address register, since it can store only 16 bits. Let us see how this problem is taken care of.

11.6.3 | Memory Addresses above 64 KB

In the program of Example 11.6, only the lower 16 bits of the 20-bit address can be sent to the base address register, which means that during a DMA cycle, when the DMA controller controls the system bus, only the 16-bit part of the starting address of the memory block can be placed on the address bus. So, who sends the upper 4 bits to the address bus? The answer is that an extra four bit latch should be included in the system hardware. The latch's input should be hardwired to carry the upper 4 bits of address, or it should be sent to the latch through program instructions. Thus, during a DMA cycle, the upper 4 bits of the address are obtained at the output of this latch.

Now look at Fig 11.21 once again. Note that the DMA chip has only 8 address lines and they are numbered as A_7 to A_0 . That means the chip has no lines to put the upper eight bits (A_{15} to A_8) of the address, during a DMA cycle. This necessitates the use of another latch which is 8-bit in size.

Now, let us review the address bus operation totally.

- i) Before DMA operation starts, the processor controls the address bus. It uses the address bus to write words into the DMA controller's control and mode registers for initialization and mode setting of the chip. Then the AEN (Address Enable) signal of the DMA chip goes high and after that, the processor is isolated from the system bus and the DMA controller is in charge of the bus.
- ii) When the DMA operation starts, the 20-bit address of the memory block starting address should be put on the system address bus. Remember that 16 bits of this address are available in the DMA chip's internal register (base address register). The chip puts lower 8 bits of this address on its pins A_7 to A_0 . It also puts on its data lines (DB_7 to DB_0), the upper




Figure 11.23 | Address generation for DMA operation

8 bits of the address available with it. These data lines are fed to the input of an 8-bit latch (See Fig 11.23) which is clocked by the ADSTB (Address Strobe) signal. Thus at the output of this 8-bit latch, the upper 8 bits of the 16-bit address is available and it can then be removed from the data lines of the 8237 chip.

- iii) During DMA operation, the address on the address bus goes on incrementing. The lowest 8 bits increments from 00 to FFH and then has to overflow to the upper 8 bits which have to be incremented by one, now. For example, for a 16-bit starting address of 1200H, the lower 8 bits roll from 00 to FFH, and then the address must change to 1300H. At this point, the number 13H must be sent from the DMA chip through its data pins (DB7 to DB) and this must be latched on to the latch IC. At this instant ADSTB goes high again acting as a clock to this latch. This is an additional operation during a DMA cycle, but since it occurs only once in 255 DMA cycles, it is not a big overhead.
- iv) What about the uppermost 4 bits of the address A_{19} to A_{16} ? See the 4-bit latch in Fig 11.23. When AEN goes high, these addresses are latched to the latch, and the output of this latch constitutes the uppermost 4 bits of the address bus. How does the 4-bit address reach the input of the 4-bit latch? For a particular application, if this part of the address is fixed, it can be hardwired. Or it can be sent from the processor from an output port, and latched by the use of AEN during DMA operation.

The EOP signal When the programmed number of bytes have been transferred i.e., when the terminal count has been reached, the DMA controller lowers the $\overline{\text{EOP}}$ (End of Process)

signal, de-asserts its HRQ pin, lowers the AEN signal and releases the system bus, so that the process can resume its normal operation.

11.6.4 | Initialization and Programming the 8237

Continuing with the idea of programming the 8237 chip, let us see the registers it has for command, status and modes. Since we have designed the logic of $\overline{CS}$ as in Fig 11.22, the addresses of these registers are as in Table 11.4.

Command/Status Register The command register is the register used for controlling the operation of the DMA chip. For that, the register must be written to. The same register is used as a status register when its contents are read from. First, let us see the format of the command register.

Example 11.7

Write the command word for using the 8237 DMA chip with the following requirements:

- i) memory to peripheral transfer,
- ii) compressed timing,
- iii) rotating priority,
- iv) late write,
- v) Both DREQ and DACK active high,

Solution

Refer to the command word format in Fig 11.24.

D ₇	D ₆	D ₅	D ₄	D ₃	D ₂	D ₁	D ₀
1	0	0	1	1	0	0	0

The command word is 98H.

The Status Register This register is available to be read out and it contains information about the status of the channels of the chip. See Fig 11.25. This information includes which channels have reached a terminal count and which channels have pending DMA requests. Bits 0 to 3 are

Table 11.4 | Addresses of Internal Registers

Register name	A3	A2	A1	A0	Address (Hex)
Status/Command Register	1	0	0	0	C8
Request Register	1	0	0	1	C9
Single Mask Bit Register	1	0	1	0	CA
Mode Register	1	0	1	1	CB
Clear Byte Pointer	1	1	0	0	CC
Master Clear/Temperory Register	1	1	0	1	CD
Clear Mask Register	1	1	1	0	CE
All Mask Register	1	1	1	1	CF




Figure 11.24 | Status of the command register

set every time a terminal count is reached by that channel or an external $\overline{\text{EOP}}$ is applied. These bits are cleared upon reset and on each status read. Bits 4–7 are set whenever their corresponding channel is requesting service.

Mode Register Each channel has an 8-bit mode register associated with it the format of which is shown in Fig 11.26. When the register is being written to, in the program condition, bits 0 and 1 select which channel's mode register is to be written to. The rest of the bits select various mode options. There are a few more registers for the DMA chip. For more details of programming the chip, the data sheets of the chip may be referred to.

11.7 | DMA and IBM-PC

In the original IBM-PC, only one 8237 chip was used, and its four channels were allocated to the following devices:

- i) Channel 0 for refreshing DRAM:
- ii) Channel 1 unused, but allowed for any I/O application.




Figure 11.25 | Status register format




Figure 11.26 | Mode register format

- iii) Channel 2 for floppy disk controller.
- iv) Channel 3 for hard disk controller.

Remember that the 8237 has only a 16-bit address bus capability; hence for the addresses beyond 64KB, the extra bits are provided by a page register, which is a latch for holding the extra address bits.

11.7.1 | DMA and PC-AT

With the advent of the ISA bus for PC-AT type computers, two DMA controllers came to be used. Two Intel 8237-DMA controllers were cascaded together, and each chip had four channels. Channel allocation was as follows:

- 0 DRAM Refresh (obsolete),
- 1 User hardware,
- 2 Floppy disk controller,
- 3 Hard disk (obsolete)
- 4 Cascade from XT DMA controller,
- 5 Hard Disk (PS/2 only), user hardware for all others,
- 6 User hardware,
- 7 User hardware.

11.8 | PCI Based Computers

The latest development in PCs is the PCI bus. This bus has a different philosophy as compared to the ISA bus. As far as DMA is concerned, the scenario is as follows.

A PCI architecture has no central DMA controller, unlike ISA. Instead, any PCI component can request control of the bus (become the bus master) and request to read and write from the system memory. More precisely, a PCI component requests bus ownership from the PCI bus controller (usually the *southbridge* in a modern PC design), which will arbitrate if several devices request bus ownership simultaneously, since there can only be one bus master at one time. When the component is granted ownership, it will issue normal read and write commands on the PCI bus, which will be claimed by the bus controller and forwarded to the memory controller using a scheme which is specific to each chipset.

KEY POINTS OF THIS CHAPTER

- Serial communication is the preferred method of data transmission over long distances.
- In synchronous transmission format, data can be 8 bits or less, with/without parity bits, one start bit and one or more stop bits.
- Transmission rate in serial communications can be expressed in bauds or bits per second, but both are not the same in all cases.
- Modems are used to connect computers to an analog transmission medium.
- RS-232 is a serial communication standard in which voltage levels are higher than TTL levels.
- Converters are needed for TTL to RS232 level conversion and vice versa.
- RS-232 connectors have 25 pins, but 9 pins are sufficient and so most computers are equipped only with DB-9 connectors.
- Two terms commonly used in serial communications are DTE and DCE.
- There are a number of handshaking signals in the RS-232 protocol.
- There are many chips used as serial communication interfaces, but the chip discussed here is the USART 8251.
- This chip caters to both synchronous and asynchronous communications.

- The method of programming this chip is to write into mode and command words.
- Direct memory access is a high priority activity in any computer system.
- Flyby DMA is much more efficient than fetch and deposit DMA.
- A DMA controller with a number of channels is necessary to manage DMA requests from various peripherals.
- The DMA controller chip 8237 has four channels and a number of registers to handle DMA operations.
- The chip has only the capability to handle 16-bit addresses, so extra hardware is required for processors with greater address widths.
- DMA chips are used in PCs, but the new PCI bus uses a different concept in bus mastering.

QUESTIONS

1. Distinguish between simplex and duplex communications.
2. Compare the merits and demerits of asynchronous communications with synchronous serial communications.
3. What is RS-232C ?
4. Explain how baud rate and bps may be different.
5. Where is the chip MAX 232 used, and for what purpose is it used?
6. What is the use of a null modem connection?
7. How many lines are needed to have serial communications between two computers?
8. Explain the functions of the pins $T \times RDY$ and $R \times RDY$ of the 8251.
9. What happens to the start and stop bits at the receiving point?
10. How is baud rate decided when using the 8251?
11. Why is DMA said to be a high priority activity?
12. What are the capabilities expected of a DMA controller?
13. What is meant by the word 'Terminal Count'?
14. When is memory-to-memory DMA required? Mention a specific case.
15. Why is it said that flyby DMA is very efficient?

EXERCISE

1. Write a program to send a string of characters between a transmission point using the 8251 is
 - a) loop back mode,
 - b) two different kits.
2. Write mode words for transmission with different baud rates, different parity settings and different character lengths.
3. Write instructions to read the following conditions:
 - a) $T \times RDY$
 - b) $T \times EMPTY$

- c) DSR
 - d) PE
4. For the DMA controller, design a decoding scheme and get a different set of addresses for the count and address registers.
 5. For the 8237, design a command word for the following conditions:
 - a) peripheral to memory transfer,
 - b) normal timing,
 - c) DREQ active low,
 - d) DACK active high,
 - e) fixed priorities.

12 SEMICONDUCTOR MEMORY DEVICES



In this chapter, you will learn

- The concepts associated with semiconductor memory.
- The internal architecture of SRAM.
- The principle of operation of DRAM memory.
- The control signals associated with DRAM.
- The reason why SDRAM is considered to be superior to DRAM.
- The different types of programmable memory.
- The necessity for having cache in a system.
- The different cache mapping techniques used.
- The cache structure of x86 processors.

Introduction

We have talked about ‘memory’ all along the chapters of this book. We talked about memory reading and writing, memory speed, main memory, secondary memory and virtual memory. However, we are yet to examine the technical constituents of memory. This chapter discusses the technology behind memory. Here, we will be talking only about main or primary memory – the memory that the processor accesses directly, the speed of which plays an important role in any computer system. This is semiconductor memory and a first classification of it is as RAM and ROM. To understand the complete hierarchy of memory as a memory system, the concept of ‘virtual memory’ is also to be made clear and that is done in Chapter 15.

RAM stands for ‘Random Access Memory’ which implies that all addressed locations are accessible and the access time is the same for all of them, but the quality that differentiates it from ROM is that it is writable as well as readable, and also it is volatile – the data that it contains disappears when power is removed. ROM, on the other hand is non volatile, and it can only be read from. Data is permanently stored in it and so in PCs, BIOS and certain startup information is put in ROM. One important aspect of ROM is that its access time is much higher than that of RAM, and so in PCs, the BIOS stored in ROM is copied to RAM for quick access and that is what is called ‘shadow RAM’. Having said that much about ROM, we will now go into the main component of our discussion – RAM.

When a user talks about ‘memory’, he usually means RAM. The amount of RAM available in a system has an impact on the speed of the system – the more the better is what is believed. This is because all programs currently running on a system must fit on the RAM, so if there is not sufficient amount of RAM, the processor will have to make trips to the secondary memory and that is painful, in terms of speed. However, the cost factor prevents a large amount of RAM

to be included in any system. (Of course, there is the ‘cache’ also which contributes to speed, but that is another story which we will soon come to.)

12.1 | Semiconductor Memory

However, first, let us discuss some general aspects of semiconductor memories. Data is stored in memory, and it is usually defined to be ‘byte oriented’ (in most cases). This means that one address corresponds to one byte of memory. Thus, when one address is accessed, one byte is either read or written into it.

Reading and writing takes a certain amount of time and that is termed as ‘memory access time’. For reading, this is the time interval from the instant the address is placed at the address pins to the time the data is available at the data pins. A similar definition applies to write access time as well. The access time of any memory device depends on the technology involved. Another term frequently encountered could be ‘memory cycle time’. This is the time interval between two consecutive memory accesses. These terms will be used frequently throughout this chapter. Now, let us discuss different types of RAM.

12.1.1 | Static RAM (SRAM)

This is the type of RAM in which data is held until power is removed from it. Each cell holds either a ‘0’ or a ‘1’. The circuit for an individual SRAM memory cell comprises typically of four transistors configured as two cross coupled inverters. In this format the circuit acts as a flip flop with two stable states, a logical 0 or a 1. In addition to the four transistors in the basic memory cell, an additional two transistors are required to control the access to the memory cell during read and write operations. This makes a total of six transistors, making what is termed a 6T memory cell. Sometimes more transistors are used to give either 8T or 10T memory cells. These additional transistors are used for functions such as implementing additional ports in a register.

Figure 12.1 shows a typical SRAM chip which has N address lines, M data lines and control signals for reading and writing. Only when the $\overline{CS}$ (Chip select) line is activated will the chip become usable (selected or enabled). For reading and writing, the processor generated




Figure 12.1 | Pins of a typical SRAM chip

signals $\overline{MEMRD}$ and $\overline{MEMWR}$ should be connected to the $\overline{OE}$ (actually a read control signal) and $\overline{WR}$, (which is the write control signal) respectively. Now, observe the read and write timings of a typical SRAM chip.

12.1.2 | Memory Read Cycle

The steps in a read cycle of SRAM are as listed:

- i) Place the address of the byte to be read, on the address bus.
- ii) Ensure that the chip is activated by making $\overline{CS}$ low.
- iii) Activate the $\overline{OE}$ pin. This ensures that data is read.
- iv) The required data then appears on the data bus.

In the timing diagram two timing figures are shown – one is t_{AA} , the **read access time**. This is the time measured from the instant the address is placed on the address bus to the point in time when the required data is available on the data pins. The other timing figure is t_{RC} , the **read cycle time**, which is the minimum time between two read cycles. These two timing figures can be equal for SRAM because, as soon as data is available at the data pins, it can be read by the processor, and a new read cycle may be initiated. (This is pointed out here to contrast it with DRAM timing, where we will see another element of delay).

12.1.3 | Memory Write Cycle

A write cycle has also a similar timing. The steps in writing are:

- i) Place the address of the byte to be read, on the address bus.
- ii) Ensure that the chip is activated by making $\overline{CS}$ low.
- iii) Place the data to be written on the data bus.
- iv) Activate the $\overline{WR}$ line. Only then the data is considered to be valid.

The data then gets written into the addressed location.




Figure 12.2 | Read cycle of a typical SRAM




Figure 12.3 | Write cycle of a typical SRAM




Figure 12.4 | Internal architecture of a typical SRAM chip

12.1.4 | Internal Architecture of an SRAM CHIP

How are memory cells organized inside a memory chip? One memory can store only one-bit of information. For memory arrays, these cells are arranged in a matrix fashion, and the address is internally split into rows and columns – there are row and column decoders inside the chip, which activate the right rows and columns for a particular address placed on the address bus. Fig 12.4 shows the internal organization of a typical memory chip, which should be quite clear

since we know all the signals associated with memory functioning. Data lines are bi-directional as data can be read or written – the drivers shown indicate the direction of data travel.

What are the merits and de-merits of SRAM?

To achieve low levels of power consumptions, CMOS is typically used for SRAM technology. It uses less power than DRAM (which we will talk about, soon), but at high frequencies, it can also consume significant amounts of power. Because each cell needs at least six transistors, SRAMs are more expensive and less dense compared to DRAMs. Hence, they are not used as the primary (main) memory in personal computers. They are as fast as typical CPUs because of using the same technology as the CPU and so find more important use as ‘cache memory’ where speed is the deciding factor. Because of being expensive, caches are naturally much less in size (storage capacity) than the main memory in PCs. We will talk about ‘cache’ in Section 12.5.

12.2 | Dynamic RAM

Another very popular RAM is ‘dynamic RAM’. It is designated as dynamic because its content does not remain unchanged or static as in SRAM, and hence frequent ‘refreshing’ is necessary. To understand this point, let us see what is contained in a typical DRAM cell (Fig 12.5). A DRAM memory cell consists as shown of a single field effect transistor (FET) and a capacitor. It is the amount of charge stored in the capacitor that decides whether the cell stores a ‘1’ or a ‘0’. One of the problems with this arrangement is that the capacitors do not hold their charge indefinitely as the charge in a capacitor ‘leaks off’ and needs to be replenished. This action of replenishing the charge that gets lost is done by ‘refreshing’ the cell at regular intervals. The data is sensed and written and this then ensures that any leakage is overcome, and the data is re-instated. The two lines, the Word Line and the Bit Line are connected as shown, so that the required bit within the memory can be selected to be read or written to.

In a DRAM chip, there are multitudes of such cells which form words consisting of bits. Refreshing for the cells is done at one go, in a particular sequence. Memory addresses are decoded and converted as rows and columns of the matrix arrangement that memory elements are arranged as.

12.2.1 | Read Cycle of DRAM

Let us see the steps involved in a typical memory read of a DRAM chip. Recollect that a processor when addressing memory sends the complete address on its address pins. Between the




Figure 12.5 | Typical DRAM cell

processor and a DRAM chip, there is a memory controller whose function is to split the address into two, as columns and rows. A DRAM has only half the number of address pins as the address supplied by the processor, because the address lines of the DRAM chip is multiplexed in time for the row and column addresses. The memory controller should also generate the signals necessary for reading and writing the DRAM (Fig 12.6).

See the diagram of a memory controller of a DRAM. Because the row and address information is placed on the same address lines (multiplexed in time) the pin count of the DRAM chip is reduced. DRAM chips are large, rectangular arrays of memory cells with support logic that is used for reading and writing data in the arrays, and refresh circuitry to maintain the integrity of stored data. Memory arrays are arranged in rows and columns of memory cells called wordlines and bitlines, respectively. Each memory cell has a unique location or address defined by the intersection of a row and a column.

Let us see the steps in a typical read cycle of DRAM. Refer to the internal diagram of a DRAM chip (Fig 12.7) and the timing diagram in Fig 12.8.

- i) The row address is placed on the rows and given sufficient time to stabilize and be latched.
- ii) The Row Address Strobe $\overline{RAS}$ signal is then activated.
- iii) The Row Address Decoder selects the proper row.
- iv) Next, the column address is placed on the same address lines and allowed to stabilize and be latched.
- v) The Column Address Strobe $\overline{CAS}$ signal is then activated.
- vi) The $\overline{CAS}$ pin also serves as the Output Enable, so once the CAS signal has stabilized, the sense amps place the data from the selected row and column, on the data bus.
- vii) With this, the data in the selected address is available at the output buffers of the chip, and it is transferred to the data bus.
- viii) Before the read cycle can be considered complete, CAS and RAS must return to their previous state.




Figure 12.6 | Memory controller for a DRAM




Figure 12.7 | Internal architecture of a DRAM chip




Figure 12.8 | Typical read cycle of DRAM

12.2.2 | Refreshing

What about the refreshing rate? It varies, but typically manufacturers specify that each row should be refreshed every 64 ms. This time interval falls in line with the JEDEC* standards for dynamic RAM refresh periods. How is refreshing done? There are many methods for refresh, and one commonly used method is called ROR (RAS Only Refresh). Practically, it is done by activating each row using *RAS*. The DRAM controller takes care of scheduling the refreshes and making sure that they do not interfere with regular reads and writes. So to keep the data in DRAM chip from leaking away, the DRAM controller periodically sweeps through all of the rows by cycling repeatedly and placing a series of row addresses on the address bus. This method is designated as ROR or RAS Only Refresh.

To reduce the number of refresh cycles, one method of design is to split the address such that there are fewer rows and more columns. So, the DRAM array is then a rectangular array, rather than a square one. Other methods of refreshing are:

- i) Using CAS before RAS Refresh (CBR)
- ii) Self refresh

12.2.3 | Fast Page Mode DRAMS

The DRAMs we discussed in Section 12.2.1 are called standard mode DRAMs and their characteristic is that they require two access components – row access and column access. Placing addresses one after the other i.e., having to multiplex the row and column addresses on the same pins, obviously doubles the access times.

In fast page mode DRAMs, the idea used is that since data access is sequential in nature, the row and column addresses need be supplied only once. After that, the row address is maintained the same, because it is the address of the same page, (a page is defined as all of the memory cells that have a common row address) and the column address alone is changed, till the end of the page is reached. Thus, the time for placing row address for each read cycle is done away with, and speed of access is increased.

12.2.4 | EDORAM (Extended Data out RAM)

In 1995, a newer type of memory called extended data out (EDO) RAM became available for Pentium systems. EDO, a modified form of FPM, is sometimes also referred to as Hyper Page mode. EDO memory consists of specially manufactured chips that allow a timing overlap between successive accesses. EDO DRAM is similar to Fast Page Mode DRAM with the additional feature that a new access cycle can be started while keeping the data output of the previous cycle active. **The data output drivers are not disabled when CAS goes high on the EDO DRAM, allowing the data from the current read cycle to be present at the outputs while the next read cycle begins, and resulting in a faster cycle time.**

The previous sentence needs a little more explanation to be understood. Check Fig 12.8 that shows the read cycle timing. Also read again the sentence (in Section 12.2.1) that says

* JEDEC (Joint Electron Device Engineering Council)

JEDEC Solid State Technology Association, formerly known as Joint Electron Device Engineering Council(s) (JEDEC), is the semiconductor engineering standardization body of the Electronic Industries Alliance (EIA), a trade association that represents all areas of the electronics industry in the United States. JEDEC has over 300 members, including some of the world's largest computer companies.

that $\overline{CAS}$ functions as the output enable as well. Some time after the $\overline{CAS}$ signal is activated, the data becomes available at the data bus, and the processor needs some time to read it. So the $\overline{CAS}$ should remain active until the processor reads it. If $\overline{CAS}$ is inactivated, the output buffers are disabled, and the data is no longer available at the output. However, only after the CAS signal goes high, can a new cycle start. Thus, delaying the inactivation of the $\overline{CAS}$ signal creates a small delay. If this delay can be avoided, the speed of the memory can be increased.

In EDO DRAMs, a new cycle can start by inactivating CAS as soon as the data in the previous cycle is available at the data bus. The data remains at the output for a time sufficient for the CPU to take it. This means that the $\overline{CAS}$ signal is obviously not the 'Output Enable' signal for the data buffers, for EDO DRAMs (ref Fig 12.9). This allows a certain amount of overlap in operation (pipelining), allowing somewhat improved performance. To be precise, EDO DRAM begins data output on the falling edge of $\overline{CAS}$, but does not stop the output when $\overline{CAS}$ rises again. It holds the output valid (thus extending the data output time) until either $\overline{RAS}$ is unasserted, or a new $\overline{CAS}$ falling edge selects a different column address.

12.3 | Synchronous DRAM (SDRAM)

You might have noted the word 'asynchronous' when we talked about the read timing cycle of DRAM. This meant that the access timing is not related to the system clock at all. In around 1996, a new type of DRAM started making headway in the memory arena and that technological innovation in DRAMs is called Synchronous DRAM. For this type of DRAM, accesses are synchronized with the system clock and SDRAM is currently 'the' RAM that is used as the primary (main) memory in computer systems. Now, let us see what SDRAM has to offer in terms of improvements over asynchronous DRAM. Technologically, they are similar, but SDRAM has incorporated certain new features and modes of operations. Let us try to understand each of these features.

Synchronous Operation All operations are synchronized to the leading edge of the system clock and thus controls are made easier.

Mode Register There is a command register for this RAM into which command words are written to specify various operating modes and also generate various control signals. This is an entirely new concept for memory chips and thus allows a level of control based on the level of performance needed from the RAM.

Interleaved Memory Architecture In previous chapters, we have talked about memory banks. Now let us check out on that idea once again. Refer to the memory bank architecture of 80386 (Section 15.5). We see that the total memory is organized as banks and each bank is enabled by its Bank Enable ($\overline{BE}$) signal. Now consider that we need to access a 32-bit data from location 000000 to 000003. All four banks of Fig 12.10 have to be simultaneously enabled and read. In two-way interleaving, the method is to divide the four banks into two sets. Getting the complete double word is done by first enabling the banks in Set A, and then in Set B. While the data in Set A is being read, the DRAM cells in set B are being pre-charged. Thus, the pre-charge time of Set B is hidden behind the access time of Set A, and for data being read from consecutive addresses, this works out to be an advantage. Interleaving can be done for any RAM system, but for asynchronous DRAMs, if interleaving is to be done, it must be done on the circuit board. In the case of SDRAMs, it is done within the chip itself, and there is also the provision for 'not to




Figure 12.10 | Interleaved memory banks

use interleaving’, which means that the command register can be programmed for ‘interleaved or not-interleaved’ access.

Burst Mode with Selectable Burst Length Data that we need usually reside in consecutive memory locations, so when data in a page is to be accessed, the row and column addresses need to be provided only once. After that, only the column addresses have to be incremented/decremented and data is continuously obtained. This is the burst mode of data transfer (also ref Section. 16.1 and Fig. 16.3), and it is the way that read/write occurs in FPM and EDO DRAMs. In SDRAMs, this is done but the command register in the chip gets the necessary control signals enabled. The burst length is the number of words that can be continuously input/output for a read or write operation, and this is selectable, using the command register). Thus, burst lengths of 2, 4, 256 or full page can be ‘set’ by suitably configuring the command word of the SDRAM.

Let us analyze a typical ‘burst read cycle’ taken from the data sheet of Micron Technologies. Fig 12.11 gives a lot of information about the points we have discussed earlier. A burst length of 4 has been specified – hence, with one read command, four sets of data outputs have been obtained. For reading again, one more read command has to be initiated. The commands are issued by loading suitable control words in the command register and it causes the necessary control signals to be activated. The NO OPERATION (NOP) command is used to perform a NOP to an SDRAM that is selected. This prevents unwanted commands from being registered. The DQ lines are the data input/output lines. The burst length is the number of words that can be continuously input/output for a read or write operation. Fig 12.11 shows a burst length of 4.

Selectable CAS Latency CAS Latency is the number of clock cycles that occur from the input of a command to the output of data. The number of clocks can be set in the Mode Register. CL is the delay, in clock cycles, between the registration of a READ command and the availability of the first piece of output data. The latency can be set to two or three clocks. If a READ command is registered at clock edge n and the latency is m clocks, the data will be available by clock edge $n + m$. The DQs will start driving as a result of the clock edge one cycle earlier ($n + m - 1$), and provided that the relevant access times are met, the data will be valid by clock edge $n + m$. For example, assuming that the clock cycle time is such that all relevant access times




Figure 12.11 | DRAM ‘read’ cycle with a burst length of 4

are met, if a READ command is registered at T0 and the latency is programmed to two clocks, the DQs will start driving after T1 and the data will be valid by T2, as shown in Figure 12.11 which shows a CAS latency of 2 ($n = 0, m = 2$).

Why is CAS Latency Important?

Think of it this way. One factor that we have ignored during all our discussions on memory and higher processors of the x86 family is the speed difference between the CPU and memory. Generally, CPU works at high speeds, while memory is slower. Typical SDRAMs have access times of 50 ns and less, but CPU clock speeds tend to be higher.

So, when a CPU issues a read command, it might have to wait to get the data. In the normal case, the READY pin of the CPU co-ordinates this, by making the CPU to issue wait cycles within its read (or write) cycles.

When using SDRAM, there is a way of using it such that the CPU does not have to wait – instead the SDRAM issues NOP commands in its read/write cycle. This means that the ‘latency’ of the SDRAM can be known and set in the mode register (as a number of clock cycles) of the SDRAM. Since the processor knows how much time will elapse before the data appears on the data bus, it can turn to other jobs instead of just waiting idly. For more details on this and other features of SDRAMs, please refer to the data sheets of SDRAMs.

12.3.1 | Synchronous vs Asynchronous DRAMS

Let us conclude by saying that in terms of the basic technology and principle of operation of a basic DRAM memory cell, both types are the same, but the SDRAM scores higher only because of the way it is used. Since asynchronous DRAM does not share any sort of common

clock signal with the CPU and chipset, the chipset has to manipulate the DRAM's control pins based on all sorts of timing considerations. SDRAM, however, shares the bus clock with the CPU. Commands can be placed (or, certain predefined combinations of signals) on its control pins on the rising clock edge.

A significant difference between conventional DRAM and SDRAM is the way in which memory access is executed. In a standard DRAM, the toggling of the external control inputs has a direct effect on the internal memory array. In an SDRAM, the input signals are latched into a control logic block which functions as the input to a state machine. Therefore, the state machine actually controls the memory access. Basic operations of the SDRAM, such as Read, Write and Refresh, are initiated by loading control commands into the device. The most common control commands of the SDRAM are:

- Row Address Strobe (RAS)
- Column Address Strobe (CAS)
- Pre Charge
- CAS-before-RAS (CBR) Refresh
- Self Refresh

Now that the concept of DRAMs is clear, let us find out what the words DDR, DDR-2 etc. mean. These are all SDRAMs, but because we keep hearing these terms frequently, knowing their special features will definitely contribute to our knowledge.

DDR This stands for 'Double Data Rate' SDRAM and the difference it has from regular SDRAM is that it can be made to transfer data at both the rising and falling edges of the clock, instead of just at the rising edge. This should double the data rate and hence the designation DDR. It achieves higher throughput by using differential pairs of signal wires to allow faster signaling without noise and interference problems. (Differential signaling means that the signal is taken to be the difference of a positive and a negative signal, like a differential amplifier. This reduces common mode signals, which are noise signals.) DDR SDRAM first came to market in the year 2000, but it did not really catch on until 2001 with the advent of mainstream motherboards and chipsets supporting it. DDR found initial support in the graphics card market and since then has become the mainstream PC memory standard. As such, DDR SDRAM is supported by all the major processors, chipset, and memory manufacturers.

DDR – 2 and DDR – 3 These are just faster versions of DDR SDRAMs and use special techniques for speed up, considering that the basic latencies of a DRAM cell can never be done away with completely. DDR2 is still double data rate just as with DDR, but the modified signaling method enables higher speeds to be achieved with more immunity to noise and cross-talk between the signals. The additional signals required for differential pairs add to the pin count of DDR2 and DDR3.

12.4 | ROM (Read Only Memory)

This is a very commonly used term and we know it stands for 'Read Only Memory'. Since it is 'firmware' a ROM does not lose its contents when power is switched off. ROM is a type of 'programmable' memory. It has internal fuses which when blown create a bit pattern which is permanent and hence can be read whenever needed. However, if it is an OTP (one time programmable) ROM, its contents can never be changed again. This statement implies that there are other kinds of ROM into which data can be re-written. This is EPROM.

12.4.1 | EPROM

EPROMs are ‘Erasable and Programmable’ – their contents can be erased by exposing them to ultraviolet radiation. Such ROMs have a window through which UV light penetrates the chip.

12.4.2 | EEPROM

This is ‘Electrically Erasable’ PROM, and erasure can be done while the chip is on the circuit board. The predominant feature of EEPROM is that the programmer can change the data embedded on the memory one byte at a time, giving him more control on how he enters the data. However, this method takes a very long time especially when erasing all the data in it. EEPROM was used largely in electronics that are programmed only a few times before shipping but can then be updated via patches. An example of this would be the chip that holds the BIOS (Basic Input Output System) of our computer. It can be reprogrammed with updates from the manufacturer in order to add further functionalities or to fix a bug that was not discovered at the time of shipping.

12.4.3 | Flash ROM

This is a special type of EEPROM that can be erased and reprogrammed in blocks instead of one byte at a time. It partitions memory into blocks. Although writing data into the flash memory is still done at the byte level, erasing the content would mean erasing the block as a whole. The block can be the whole memory itself. This feature gave flash memory the advantage of speed over EEPROM. Flash memory became very popular because it takes so much less power compared to a hard disk and is much more durable, capable of surviving excessive heat, pressure, and even can be submerged in water. Flash memory became the instant successor of the aging Floppy disks not only for its durability but also for its great capacity and relatively small size. The only drawback of flash memory is that like EEPROM, it can only last a certain amount of erase and write cycles before failing, but the number is quite high, in the range of 500,000 and more for flash and EEPROM.

Many modern PCs have their BIOS stored on a flash memory chip so that it can easily be updated if necessary. Such a BIOS is sometimes called flash BIOS. Flash memory is also popular in modems because it enables the modem manufacturer to support new protocols as they become standardized. Many microcontrollers have flash ROM on chip. This allows the programmer to change the burned program on line.

12.4.4 | NVRAM

It is an abbreviation of Non-Volatile Random Access Memory, a type of memory that retains its contents when power is turned off. One type of NVRAM is SRAM that is made non-volatile by connecting it to a constant power source such as a battery (usually a lithium battery). Another type of NVRAM uses EEPROM chips to save its contents when power is turned off. NVRAM is composed of a combination of SRAM and EEPROM chips. The CMOS memory in PCs which stores important settings and the system time is NVRAM.

12.5 | Cache Memory

Cache is a word you might have heard in various contexts, not necessarily in this book alone. Now that we have discussed various semiconductor memories and their relative merits and demerits, this is the right time to introduce the concept of ‘cache memory’ or ‘cache’ as it is

usually called. The literary meaning of ‘cache’ (as a noun), is ‘a hiding place’, and as a verb ‘to conceal’. We will find that in the context of computer architecture, this word would be more meaningful as ‘a place of temporary storage’.

In a computer system, there are various hierarchies in the storage of information. In the processor, registers store data. Then there is main memory, outside the processor but very close to it and also secondary memory, way down in the memory hierarchy. Where in this picture, does ‘cache’ fit in? See Fig 12.12. The block ‘cache’ is shown close to the processor, directly connected to it through the system bus, just as the main memory is. This means that data from the processor can be transported to/from the cache just as in the case of the main memory. However, which of them is to be considered to be higher in the hierarchy? The answer is ‘cache’. The cache is where information (data and programs) required for the execution of any program is expected to be found in. When the processor is to execute a task, it looks in the cache for the required program. If the item is found there, a cache ‘hit’ is said to have occurred. Otherwise, it is a ‘cache miss’. If a miss occurs, the processor goes to the main memory and brings the item to the cache and then the program execution is done. The point is that, when the required information is in the cache, speed of execution is very high.

Now let us discuss the ‘how and why’ of this matter. In a computer system, the fastest component is the processor. It has very little storage because only general-purpose registers are available therein. To have fast execution, there should be a storage space as fast as the processor. This is the desired ideal, and the aim of any computer design is to get as close as possible to the ideal. For that, we need very fast memories. The fastest memory available is SRAM whose technology is similar to the technology of the processor. However, SRAM is very expensive and having sufficiently large quantities of it is prohibitively expensive. It also dissipates a lot of power. So using SRAM as main memory is ruled out. The next best alternative is DRAM which is not as good, but sufficient (see Table 12.1 for a comparison between the two RAM technologies). DRAM has high packing density, low power and low cost. Thus, DRAM has been chosen as the technology of choice for the main memory.

Now, for the speed factor – to achieve higher speeds than is achievable by DRAM, a small size SRAM memory is inserted between the processor and the main memory. This is called the ‘cache’ and it is the choice of storage for the ‘current process’ – for the ‘program under execution’ at any time. This means that the memory system design should ensure that the data and code of a program to be executed, should be found in the cache at ‘all times’ – or that a ‘cache hit’ should be the rule rather than the exception. The design of most computer systems achieves a ‘hit rate’ above 90%, which is considered to be reasonably good. (The ‘hit rate’ is the percentage of hits, out of the total memory accesses.)




Figure 12.12 | Cache in the memory hierarchy

Table 12.1 | Comparison Between SRAM and DRAM

Feature	Dynamic RAM (DRAM)	Static RAM (SRAM)
Storage circuit	Capacitor	Flip-flop
Transfer speed	Slower than CPU	Same as CPU
Latency	High	Low
Density	High	Low
Power consumption	Low	High
Cost	Cheap	Expensive

Table 12.2 | Speed Comparison Between Different Memory Components

Memory component	Speed (approximately)*
Registers on CPU	5 ns
Cache	10 ns
Main memory	100 ns
Secondary memory	5 ms

*All these figures are reducing as technology is progressing. These figures are only for the purpose of comparison.

12.5.1 | Locality of Reference

Now that it is clear what we have set out to achieve in the design of a typical memory system, the next step is to quantify the goals and find ways of reaching it. The goal of an effective memory system is that the effective access time that the processor sees is very close to that of the fastest memory component in the system i.e., the access time of the cache. All accesses that the processor makes to the cache satisfy this goal. However, at times, the data required by the processor is not available in cache and then, the speed factor is affected. However, this should occur only for a very small percentage of accesses. That is how a figure close to the ideal can be aimed for. The achievement of this goal depends on many factors: the architecture of the processor, the behavioral properties of the programs being executed, and the size and organization of the cache.

Caches work on the basis of the principle of ‘locality of reference’ of program behavior. This means that when the CPU accesses a memory location, all the information it needs will be in sequential locations in and around this location. This is called ‘spatial locality’. There is also something called ‘temporal locality’ – this implies that it is highly likely that the CPU will access the same locations again and again. What does all this mean? This simply means that there are some programs which are very frequently used, and obviously the information for such programs will be in contiguous memory locations. The whole idea of ‘cache’ is based on the above mentioned principles. So, it is the duty of the ‘memory management system’ to use this principle to ensure a high percentage of cache hits.

12.5.2 | Cache Organization

It is already agreed upon that cache memory is made of SRAM. Based on affordability and performance requirements, the size of the cache for a system also has to be decided. More the

size of the cache, better is the performance – but only up to a limit – beyond that, the additional complexity that comes along with it, tends to defeat the advantage obtained.

However, more important than all these is to decide and finalize the ‘organization’ of the cache. An important point to note and remember at all times is that a cache stores only a ‘copy’ of a portion of the content of main memory. Compared to the size of the main memory, the cache is many times smaller. So, at any time, the ‘large main memory’ maps into a ‘small cache’. There are many mapping techniques for this and this constitutes the organization of the cache.

Let us start in the following way. The main memory is divided into blocks and the cache is also divided into blocks of the same size. (A block is a number of bytes, which is decided upon, at some time. It may be 16 bytes or 64 bytes or larger or smaller). The number of blocks in main memory is n times the number of blocks in the cache. For example, if a cache has 64 blocks, there should be ‘ $64n$ ’ blocks in main memory. Which 64 of the $64n$ blocks in main memory is also available in the cache? How are these blocks arranged? The answers to all these constitute the ‘organization’ of the cache i.e., the mapping.

12.6 | Mapping Techniques

There are three important mapping techniques in common usage and we will see the salient features of each of them.

12.6.1 | Direct Mapping

To understand this, we will start with small figures (numbers) and a simple explanation. Consider a cache that has space for 64 blocks to be stored. Main memory has $64n$ such blocks in it. The cache blocks are numbered as CB- 0, CB-1, CB- 2 up to CB- 63. In these blocks of cache, which 64 blocks of the main memory can be loaded? The main memory blocks are named MMB-0, MMB-1 MMB- $(64n-1)$. See Fig 12.13 for $n = 8$.

The ‘direct mapping’ technique fixes that, in cache block location CB-0, only MMB-0, MMB-64, MMB-128, MMB-192 etc. can be copied. Similarly in CB-1, only MMB-1, MMB-65, MMB-129 and so on can be copied. Fig 12.13 shows that in the cache blocks CB-0 to 63, main memory blocks MMB-0 to MMB-63 can be correspondingly mapped. The next 64 sets of main memory blocks MMB-64 to 127 can also be mapped to CB-0 to 63. This pattern repeats.

So to which cache block can main memory block MMB-125 be mapped to? The answer is $125 \bmod 64 = 61$ i.e., CB-61. This is fixed, and MMB-125 can only be copied to CB-61. The statement defining direct mapping is that:

If each main memory block has **only one place** it can appear in the cache, the cache is said to be direct mapped. The mapping is to the cache block number given by the calculation:

(Main memory Block address) MOD (Number of blocks in cache)

So MMB-128 is mapped to CB-0 as $128 \bmod (64) = 0$

MMB-221 is mapped to CB-29 as $221 \bmod (64) = 29$

12.6.2 | Fully Associative Mapping

This is the case when a main memory block can be copied to just any cache block. Figure 12.14 shows a diagram of a fully associative cache organization. This organizational scheme allows any block in main memory to be stored in any block in the cache. For example, the main




Figure 12.13 | Direct mapping




Figure 12.14 | Fully associative mapping

memory block MMB-0 of main memory is stored in CB-3 or CB-2 of the cache. However, this is not the only possibility; MMB- 0 could have been stored in any block within the cache. Since any cache block may store any memory block, the name ‘fully associative’ is used for such a mapping.

12.6.3 | Set Associative Mapping

If a block in main memory can be placed in a restricted **set** of positions in the cache, the cache is said to be set associative. A set is a group of blocks in the cache. A main memory block is first mapped onto a set, and then the block can be placed anywhere within that set.

Figure 12.15 shows the **2-way set associative cache mapping technique**. In this, the cache is divided into sets. For **2-way associativity**, each set has **two cache blocks** in it. Consider the case of 64 cache blocks, and thus 32 sets in the cache. The sets are numbered as Set 0 to Set 31.

Main memory blocks MMB 0 can be mapped to the cache blocks of Set 0. Likewise MMB 1 has the option of being mapped to any of the two cache blocks of Set 1. What about MMB-2? It can likewise be copied to Set 2, and so on. Thus, MMB-31 is copied to any one of the blocks of Set 31. What about MMB-32? It goes to Set 0. The set is usually chosen by the formula.

(Block address) MOD (Number of sets in cache)

MMB 32 is $25 \bmod (32) = 0$. It can be mapped to any cache block of Set 0.
MMB 125 is $125 \bmod (32) = 29$. It can be mapped to any cache block of Set 29.

The range of caches from direct mapped to fully associative is really a continuum of levels of set associativity. Direct mapping is simply one-way set associative and a fully associative cache with *m* blocks could be called *m*-way set associative. The vast majority of processor caches today are direct mapped, two-way set associative or four-way set associative.

Keep in mind that for 4-way associative mapping, there are four options for any main memory block. However, within each set, the mapping is fully associative. i.e., the MMB can map to any cache block within the set.

For 2-way associativity, there are two options for any main memory block and for direct mapping, there is only one ‘position’ for any main memory block – there is no other ‘option’




Figure 12.15 | Two-way set associative mapping

in the case of direct mapping, We have used the word 'cache blocks', but a more common terminology is 'cache line'.

Example 12.1

Consider the following case of a system with 2048 cache blocks and 8192 main memory blocks. Find, where in the cache, will the main memory blocks MMB-15 and MMB-2029 be placed for the mapping policies of:

- a) direct mapping
- b) fully associative mapping
- c) 4-way set associative.

Solution

- i) MMB – 15

- a) Direct mapping.

This block will be copied to CB-15.

- b) Fully associative.

This MMB block can be copied to any cache block.

- c) Four-way set associative.

In four-way set associative, the cache will be divided into sets of four blocks. Thus, there will be $2048/4 = 512$ sets in the cache.

$15 \bmod (512) = 15$ i.e., MMB-15 can be copied to any of the four cache blocks of Set 15.

- ii) MMB – 5029

- a) Direct mapping.

$5029 \bmod (2048) = 933$. This block will be copied to CB-933.

- b) Fully associative.

This block can be copied to any cache block.

- iii) Four-way set associative.

$5029 \bmod (512) = 421$.

This block can be copied to any of the four cache blocks of Set 421.

We have discussed the mapping technique using small number of cache blocks. In reality, the number of cache blocks will be quite high. Also, a lot of circuitry is necessary for checking for a hit or miss by comparing block addresses and so on. A detailed discussion on how these mapping techniques are practically implemented is beyond the scope of this book, but should be obtained from any book on computer architecture.

Comparison

In terms of complexity of implementation, the fully associative mapping has the highest amount of comparison to be done. However, it is the most flexible mapping technique, since any main memory block can be placed in any cache block position.

As the mapping becomes set associative, the complexity reduces, but so does the flexibility. For any system, the chosen mapping technique will be a compromise between these two conflicting factors.

12.6.4 | Cache Write Policies

Remember that a cache contains a **copy** of some part of the content of main memory. When a block of data is brought into the cache from main memory and processed, that content could get changed. However, the original content in main memory does not get changed. This will create an inconsistency, and to avoid it, the main memory content should be correspondingly changed. For this, caches have what is called a 'write policy'.

For any cache, policies regarding reading and writing are pre-defined. Here, we will touch upon the write policy alone. There is the write-back and write-through policy.

In write-back policy, the cache acts like a buffer. That is, when the processor starts a write cycle, the cache receives the data and terminates the cycle. The cache then writes the data back to main memory when the system bus is available. This can cause a situation that the content of the cache has changed, but not the content of the main memory. To make the system aware of this, there is a bit in the cache called the 'dirty' bit which is set to show this condition. At a later time, the system can copy the changed contents of cache to main memory and clear the dirty bit. This method provides the greatest performance by allowing the processor to continue its tasks while main memory is updated at a later time. The second method is the write-through policy. In this, when a cache content changes, the change will be updated in the cache as well as in the main memory concurrently.

12.6.5 | Cache Replacement Policies

We know that we cannot let any main memory block reside in memory permanently. Only the ones that are constantly being used are kept there. If a new block which is important needs to be transported to the cache and if the cache is found to be full, what is to be done? The least important of the blocks currently residing in the cache can be overwritten by the new data. However, the 'catch' here is how to determine the less important ones. There are various algorithms for this, and the most popular algorithm is the 'Least recently used' (LRU) algorithm. More algorithms for specific applications are also available – First-In First-Out (FIFO), Most-Recently Used (MRU), Least-Frequently Used (LFU), Most-Frequently Used (MFU) and so on. These matters are managed by the operating system.

12.7 | Cache and the x86 Family

These days, the word 'cache' is in common usage in the context of PCs. However, it was not always so. The early PCs had no cache. The first cache was used for '386 PCs. This cache was placed on the motherboard, but the processor ('386) offered support for its operation, and there was a cache controller on the board. The amount of available cache varied depending on the motherboard model and typical values at that time were 64 KB and 128 KB.

With the '486 processor, Intel added a small amount (8 KB) of cache inside the CPU. This internal cache was called L1 (level 1) or internal, while the external memory cache was called L2 (level 2) or 'external'. The amount and existence of the external memory

cache depended on the motherboard model. Typical amounts for that time were 128 KB and 256 KB.

Then there was the first Pentium processor with two internal caches – one for data and the other for instructions. Each was 8 KB in size. Later, the size of this cache was increased. However, the L2 cache was still external, and its size depended on the motherboard manufacturer. Gradually, sizes like 256 KB for L1 cache and 512 KB for L2 became standard. However, having an external cache amounted to less speed, as the processor clock is high speed (>200 MHz) while motherboard speeds were less (around 66 MHz).

However, even at that time, there was the Celeron processor which had no cache, and hence slower but cheaper – so there were buyers for this processor based PCs. However, now, Celeron processors also have internal caches.

With P6, things have changed. Both L1 and L2 caches are internal – on chip. This same architecture is used until today – both L1 and L2 caches are located inside the CPU running at the CPU internal clock rate. So the amount of memory cache you can have on your system will depend on the CPU model you have; there is no way to increase the amount of cache without replacing the CPU.

Figure 12.16 shows different caches for a typical CPU like x86. L1 D and L1 I caches are L1 data cache and L1 instruction caches respectively. More details of this are given in Chapter 16.

As more and more processors began to include L2 cache into their architectures, Level 3 cache became the name for the extra cache built into motherboards between the microprocessor and the main memory. Quite simply, what was once L2 cache on motherboards now became L3 cache when used with microprocessors containing built-in L2 caches. However, the latest is that with more CPUs becoming ‘multicore’ (refer Section 16.7.2) a tertiary level cache was added on to the CPU die, called the L3. It also became common to have the three levels be larger in size than the next lower level, so that it became not uncommon to find Level 3 cache sizes of eight megabytes. This trend appears to continue for the foreseeable future.

Figure 12.16 shows three levels of cache for the latest x86 processors. This is only schematic, and the arrows do not show the path of ‘data flow’.




Figure 12.16 | Caches for a current x86 processor

KEY POINTS OF THIS CHAPTER

- Semiconductor memories are used as main memory and cache, in a computer system.
- An SRAM memory for one-bit requires at least 4 transistors. So packing density of SRAMs is rather low.
- Reading and writing in an SRAM is done asynchronously.
- A typical DRAM cell for one-bit requires only one transistor and a capacitor.
- Since the charge in a capacitor tends to leak, DRAMs require refreshing.
- To reduce the pin count of a DRAM chip, the address from the processor are split into row and column address (by a memory controller), and then given to the pins of the DRAM in a multiplexed fashion.
- Besides the standard mode DRAMs, there are fast page and EDO DRAMs which are faster.
- Synchronous DRAM is similar in technology to DRAM, but is managed by writing to a command register, and this type of RAM is currently the most popular type in semiconductor memories.
- DDRs and DDR-2, DDR-3 etc. are just faster versions of SDRAM.
- There are different types of ROM available, some of which are erasable and re-programmable.
- NVRAM is used to store configuration settings and system time, in a PC.
- A 'cache' is a type of SRAM placed between the main memory and the processor.
- The cache is where working data and code are expected to be found in.
- The size of cache is small, and mapping techniques dictate which blocks of the large main memory are to be copied to the small cache.
- Direct mapping is the simplest and most economical type of mapping.
- Fully associative mapping is very efficient, but complex.
- Most systems use either direct mapping or 2 or 4-way associative mapping.
- Current x86 processors have L1, L2 and even an L3 cache.

QUESTIONS

1. List the pins of a typical SRAM chip with their functions.
2. How many active devices does a one-bit SRAM need?
3. Define 'read cycle time' for an SRAM.
4. Why are addresses split into row and column parts within an SRAM chip?
5. Why do DRAMs require refreshing?
6. Why is the read cycle time not the same as the read access time for a DRAM?
7. How do you account for the use of reduced address pins for a DRAM?
8. Compare FPM DRAM with EDO DRAM.
9. Which is the current memory device used as main memory in PCs?
10. What is meant by CAS latency?
11. What is meant by 'burst mode' data transfer?
12. What is the technology used for cache?

13. What is meant by 'mapping techniques' with reference to caches?
14. How much of cache should a system have?
15. Why is fully associative mapping not used commonly?
16. What is the penalty for a cache miss?

EXERCISE

1. List the numbers of 5 standard SRAM chips and find out their capacity and the number of output pins of each of them.
2. How many address pins are needed for each of these following SRAMs?
 - a) 8 KB
 - b) 128 KB
 - c) 1 GB
 - d) 4 MB
 - e) 128 MB
3. Find the names of three manufacturers of DRAMs and list the numbers of 5 chips and their capacity. State if each of them are standard DRAM, FPO or EDO DRAM.
4. List 5 SDRAM chips and from their specifications, find out their programmable features like number of burst cycles, CAS latency and so on.
5. What do the following ratings of SDRAM mean?
 - a) PC66,
 - b) PC100,
 - c) PC133
6. There is an SDRAM marked as 3-2-2. What does this x-y-z marking mean?
7. There is a type of RAM called RAMBUS RAM. What is it and how and where is it used now?
8. Find the amount of cache available in the 80486 and P-4 and P-6.

13 MULTIPROCESSOR CONFIGURATIONS



In this chapter, you will learn

- The principles and problems of multiprocessing.
- The different types of multiprocessing systems.
- The way 8086 handles multiprocessing.
- How the 8086 operates in a coprocessor configuration.
- How the 8086 works in conjunction with an I/O processor.
- The bus arbitration schemes used in loosely coupled systems.
- The use of the bus arbiter 8289 for bus arbitration.
- The relevance of IEEE 754 format for real numbers.
- Conversion from IEEE 754 number format to decimal numbers and vice versa.
- The architecture of the 8087 arithmetic coprocessor.
- To write programs using 8087 instructions.

Introduction

What is multiprocessing?

Multiprocessing is a term used to indicate that the system contains multiple processors, instead of a single processor (as is the case, usually). A 'processor' is defined as an active entity having an instruction set of its own and the capability to execute instructions by itself. Let us list out in clear terms the significant advantages of multiprocessing systems:

- i) The system can be designed such that tasks may be allocated to special purpose processors whose designs are optimized to perform certain types of tasks easily and efficiently.
- ii) It is seen that very high levels of performance are achievable when multiple symmetric processors execute simultaneously. This is termed 'parallel processing'.
- iii) Higher reliability can be obtained by isolating system functions such that a failure or error in one part of the system has only a limited effect on the rest of the system.
- iv) When an application can be broken up into smaller and more manageable tasks, parallel development and modifications of subsystems can be promoted.

Such systems have obviously a lot of advantages and extra features which a single processor system cannot possess, but along with it comes a host of new problems as well. The two basic issues to be sorted out can be classified as **inter processor communication and bus contention**. The first issue is concerned with passing information from one processor to the other, while the second is of sharing the same system bus. How these issues are effectively handled is a figure of merit for

a system. For that, the processors in the system must have special pins and functions and 8086 is a processor which has been designed with extra features usable for multiprocessing.

13.1 | Multiprocessor Systems

Before going into the details of 8086 based multiprocessor systems, let us review various types of multiprocessor systems in general. Broadly, they can be classified as tightly (closely) coupled and loosely coupled multiprocessing configurations.

13.1.1 | Tightly Coupled Systems

Tightly coupled multiprocessing systems are a group of processors that are dependent on one another (to a certain extent). The address, data and control lines of the processors are connected in parallel and called the local bus. They share the same resources like memory and I/O through a common shared system bus. For such systems, inter processor communication is provided through shared resources – one processor sends a message into the shared memory which the other processors can use, and also all processors can use the services of the common I/O and so on. In some cases where the processors are more dependent on each other, there is also the possibility of sending messages from one processor to the other through special processor pins.




Figure 13.1 | Tightly coupled multiprocessor system




Figure 13.2 | Loosely coupled multiprocessor system

13.1.2 | Loosely Coupled Systems

Loosely coupled configurations are relatively large. They may have a number of independent modular systems, each with local sources and independent existence, while being connected to similar or different modules through a global bus, through which global memory and I/O can be accessed. Each of the modules may be tightly coupled multiprocessor systems themselves or just single processor systems, or even just a processor having access to only global resources. These modules may be installed in a single cabinet or may be quite some distance apart. Figure 13.2 shows a loosely coupled system with three processing modules. The details of one processing module has been elaborated.

13.2 | Multiprocessing Using 8086

In Section 6.4, we discussed the maximum mode of the 8086 processor that provides multiprocessing capability for the processor. In this mode, the functions of certain pins get changed so as to provide the processor the ability to handle the basic problems of multiprocessing, one of which deals with bus issues. The issue may be elaborated and specified as 'bus arbitration' and 'mutual exclusion'.

Bus arbitration deals with finding a solution to the priority in which the processor or processing module is to be granted access to the bus, in the case of a conflict. This is solved by the request/grant ($\overline{RQ}/\overline{GT}$) logic pins of the 8086. This pin may be sufficient in tightly coupled configurations, but in larger systems, a bus arbiter IC is usually needed. For mutual exclusion, the processor has a $\overline{LOCK}$ pin, as well as a LOCK instruction which prevents other bus masters from grabbing the bus from the current master.

Now we will discuss, in detail, multiprocessing based on the 8086 processor. The Intel 8086 family architecture allows the inclusion of processors in either the tightly coupled or loosely coupled configurations. However, the architecture directly supports two types of processors – independent processors and coprocessors. The former class has 8086/8088 and 8089 (designated as an IO processor) as its members and these processors can work quite independent of each other. The latter i.e., the co-processor is different in the sense that its instructions are written in the stream of instructions written for the ‘host’ or master processor, usually an 8086. The coprocessor monitors this instruction stream, recognizes certain instructions as its own, executes them and returns the result to the host processor or saves it in their common memory. The coprocessor, in effect, extends the instruction set of the host processor. The arithmetic coprocessor 8087 is the one which has been specially designed to interface and interact directly with the 8086. Let us first discuss the 8086 and 8087 in a ‘coprocessor configuration’.

Please read carefully the topic ‘maximum mode’ (Section 6.4) before attempting to understand the following.

13.2.1 | The 8086 in a Coprocessor Configuration

The 8087 chip also designated as ‘arithmetic coprocessor’, ‘floating point processor’ and ‘numeric data processor’ is specially designed for complicated arithmetic operations.

Does that mean that arithmetic computations cannot be done by the 8086? We know it can, but it has a limited data size and instructions for only the basic arithmetic operations. However, some applications require extremely fast and complex math functions which are not provided by a general-purpose processor. Functions such as square root, exponentiation, sine, cosine, and logarithms are not directly available in a general-purpose processor. Software routines required to implement these functions tend to be slow and not very accurate. Integer data types and their arithmetic operations (add, subtract, multiply and divide) which are directly available on general-purpose processors, still may not meet the needs for accuracy, speed and ease of use. Providing fast and accurate, complex math can be quite complicated, requiring large areas of silicon on integrated circuits. Trying to incorporate all this into a general-purpose processor is a burden, because most general-purpose applications do not need such complex calculations. For such features, having a special numeric data processor is the best solution – one which is easy to use and has a high level of support in hardware and software.

The 8087 is a numeric data coprocessor capable of performing complex mathematical functions while the host processor (the main CPU) performs more general tasks. The fact that the 8087 is a coprocessor, means that it is capable of operating in parallel with the host CPU, which greatly improves the processing power of the system.

Now that it is understood that this coprocessor operates in parallel with the main or host processor (as the 8086 is called here), let us see how the two processors interact and communicate with each other.

As mentioned some time earlier, both processors have their own individual and distinct instruction sets. However, there is only one common instruction stream stored in memory and only the host processor can fetch instructions. The coprocessor monitors the instructions that are being fetched. As this goes along, at some point, an ESCAPE instruction is seen by the coprocessor.

All ESCAPE instructions start with the high-order five bits of the instruction opcode being 11011. They have two basic forms, the memory reference form and the non-memory reference form. Memory reference forms of the ESCAPE instruction allow the host to point out a memory operand to the coprocessor using any memory addressing mode. The non-memory form, initiates some activity in the coprocessor using the remaining bits of the ESCAPE instruction to indicate which function is to be performed.

Now that the coprocessor has been alerted by the ESC instructions, it goes on to the task that has been initiated by the ESC instruction. Depending on the rest of the bits in the ESC instruction, it performs either the non-memory based operation, or gets down to the task of performing the memory based operation. If the memory operand is very large, the 8087 grabs the bus from the host and accesses the whole memory operand (which may be as long as ten bytes), completes the operation, and stores the result in memory or its internal registers. When it is performing these activities, it keeps its BUSY pin high. This pin is tied to the $\overline{TEST}$ pin of the 8086 and thus the host knows that the coprocessor is executing some instruction. Once the coprocessor instruction execution is over and done with, the BUSY pin is pulled low.

Meanwhile, what is the host processor doing? After the ESC instruction, it continues with its sequence of fetching the next instructions in sequence. If they are 8086 instructions, the host executes those instructions. At this point, both the host and the slave processor (the coprocessor) are executing instructions in parallel, and one problem that may arise is that both of them may want the system bus to fetch some memory or I/O operand. This is the classic problem that arises in multiprocessing – ‘bus contention’. When both processors contend for the bus, only one of them gets it. Normally, the host is in control of the bus, but when the coprocessor asks for the bus (on one of the $\overline{RQ}/\overline{GT}$ pins), the host will relinquish it. The coprocessor uses it and returns it after use. This seems reasonable, but conflicts do arise when both of them want it at the same time. One of them has to wait, obviously.

Again, after a few 8086 instructions are executed, say, there appears in the instruction sequence, another instruction of the coprocessor – or probably, the host needs the result of the previous coprocessor instruction. In this case, the host has to take the result from the shared memory. To ensure that the result is taken only after the 8087 has stored it after the completion of its current computation, there should be a WAIT instruction in the instruction stream before this.

This instruction repeatedly checks the status of the $\overline{TEST}$ pin of the 8086 until it is found to be low. This ensures that the coprocessor has completed executing the instruction, whose result is needed by the host processor. Then only the next instruction in the sequence is executed by the host. The BUSY- $\overline{TEST}$ pin connection and the WAIT instruction effectively synchronize the activities of the two processors. See the state diagram in Fig 13.3 to get a feel of the activities of both the processors.




Figure 13.3 | Synchronized activities of the 8086 and the 8087

Table 13.1 | Queue Status Decoding

QS1	QS0	Queue operation
0	0	No instruction from queue
0	1	First byte from queue
1	0	Flush the queue, because of a branch instruction
1	1	Subsequent byte from queue

13.2.2 | Connections Between the Two Processors

You have been told that the coprocessor monitors the instructions being fetched by the host processor. How does it do that? The fetching of instructions is monitored via the data bus and the 8086 bus cycle status lines $\overline{S0}$ – $\overline{S2}$ while the execution of instructions is monitored via the queue status lines QS0 and QS1. The 8087 tracks the instruction execution of the 8086 by keeping an internal instruction queue which is identical to the processor’s instruction queue. Each time the host processor performs an instruction fetch, the 8087 latches the instruction into its own queue in parallel with the host. Each time the processor removes the first byte of an instruction from the queue, the 8087 removes the byte at the top of the 8087 queue and checks to see if the byte has an ESCAPE prefix. If it has, the 8087 decodes the following bytes in parallel with the processor, to determine which numeric instruction the bytes represent. If the first byte of the instruction is not an ESCAPE prefix, the 8087 discards it along with the subsequent bytes. Table 13.1 tells us what the queue status lines indicate.

Now, let us understand how the 8086 and the coprocessor 8087 are connected in a multi-processor system. Refer Fig 13.4. Both processors share the same clock, reset and ready signals (from the 8284 clock generator). The $\overline{RQ}/\overline{GT}$ lines of the two processors are interconnected. The 8086 has two such pins. They are bidirectional pins and the request will be sent to the current bus master (the processor which is currently using the bus) from the requesting processor. The queue status of the 8086 is monitored by the 8087 through the QS1 and QS0 pins. The BUSY pin of 8087 is tied to the $\overline{TEST}$ pin of the 8087.

In the system, an interrupt controller is also shown, as is available in the case of all computer systems. The INT line of the 8087 is connected to one of the interrupt request lines of the 8259(PIC) Interrupt controller chip. The reason for this is that when certain errors occur during numeric computations, the coprocessor generates an exception (error generated interrupt). Then, the computation will be aborted and an interrupt service routine should sort out the problem. The corresponding interrupt is routed through the interrupt controller to the host processor.

The ‘local bus’ in the figure pertains to the multiplexed address/data and control pins of the processor. The bus interface components are the latches, transceivers, and bus controller used in the maximum mode system. Refer Figs 6.6 to 6.8 and Fig 6.25 for full details. Since they have been discussed in detail earlier (Section 6.1), here they are bundled into one block, to make the figure simple and easy to understand. The system bus consists of the de-multiplexed and buffered address bus A_0 – A_{19} , the buffered data bus D_0 – D_{15} and the control bus consisting of the signals from the output of the bus controller IC, interrupt

lines, hold lines and so on. We will discuss the important features of the 8087 coprocessor in Section 13.5.1.

13.2.3 | Arithmetic Co-Processing Over the Years

What we have been discussing so far has been the case of the 8086 and its compatible arithmetic coprocessor. Over the years, many changes have occurred. The 8086 is no longer a high profile processor and many generations of more complex and computationally intensive processors have been designed by Intel. In spite of all this, arithmetic processing or complex floating point operations are still considered a specialized job. Just as 8087 was made to be directly compatible with 8086, the higher processors of Intel have compatible arithmetic coprocessors as shown in Table 13.2. However, the current trend is to have the coprocessor inbuilt in the general-purpose processor chip.




Figure 13.4 | 8086 and 8087 in a coprocessor configuration

Table 13.2 | x86 Processors and Compatible Co-processors

Host processor	Compatible Co-processor
8086/8088	8087
80186/80188	80187
80286	80287
80386	80387
80486	Inbuilt
Pentium	Inbuilt

13.3 | The 8086 and 8089 in a Tightly Coupled Configuration

The 8086 directly supports multiprocessing also with processors called ‘independent processors’. One such processor is the 8089 I/O processor. It is a processor specially designed by Intel to handle input/output processing including DMA (Ref. Section. 11.5). This processor is independent in the sense that it is not directed or dependant on the host processor. In a ‘local’ configuration, both the processors’ address, data and control signal lines are tied together. The only way of communication between the two is by messages stored in their shared memory space. The host processor sets up a message in memory and asks the IOP (I/O processor) to handle the I/O task as indicated in the message. The message is actually the program the IOP is to execute. How does the host get the attention of the IOP? It sends an OUT instruction – the IOP is treated like an I/O device with a specific address (as designed by an address decoding circuitry). The select pulse from the address decoder is sent to the CA (Channel Attention) pin of the IOP and thus it is awakened from slumber. Then the IOP takes up the assigned task, completes it and notifies this fact to the host by setting a status bit in the shared memory space or by generating an interrupt.

Thus, we can see that in this case, the two processors are more or less independent of each other. However, because most of the activities of an IO processor are bus-intensive, it needs the bus frequently. Then it asks for the bus on one of the $\overline{RQ}/\overline{GT}$ pins of the 8086, gets it, uses it and returns it. Figure 13.5 shows the simplified connections between the two processors. There is a select pin for the 8089. This is because it has two channels, and when the CA signal is activated, one of the two channels must be selected (SEL can be 0 or 1). Each channel corresponds to one message or task that has been set up in memory for the IOP to execute.




Figure 13.5 | The 8086 in a tightly coupled configuration with the 8089 IOP

It is also possible to have a multiprocessor system with three processors – an 8086, 8087 and 8089. Combining the connections in Fig 13.4 and Fig 13.5 will give the connections for such a configuration.

13.3.1 | Intel I/O Processor 8089

The Intel 8089 I/O processor is contained in a 40-pin integrated circuit package. Within the 8089 are two independent units called channels. Each channel combines the general characteristics of a processor unit with those of a direct memory access (DMA) controller. The 8089 is designed to function as an IOP in a microcomputer system where the Intel 8086 micro-processor is used as the CPU. The 8086 CPU initiates an I/O operation by building a message in memory that describes the function to be performed. The 8089 IOP reads the message from memory, carries out the operation, and notifies the CPU when it has finished.

It has 50 basic instructions that can operate on individual bits, on bytes, or 16-bit words. The IOP can execute programs in a manner similar to a CPU except that the instruction set is specifically chosen to provide efficient input-output processing. The instruction set includes general data transfer instructions, basic arithmetic and logic operations, conditional and unconditional branch operations, and subroutine call and return capabilities. The set also includes special instructions to initiate DMA transfers and issue an interrupt request to the CPU. It provides efficient data transfer between any two components attached to the system bus, such as I/O to memory, memory to memory, or I/O to I/O. The previous statement makes it seem to have DMA capability. It does have DMA capability, but because it has instruction execution capability as well, it is sometimes called an ‘intelligent DMA controller’.

The communication scheme between the two processors consists of program sections called ‘blocks’, which are stored in memory as shown in Fig 13.6. Each block contains control and parameter information as well as an address pointer to its successor block. The address of the control block is passed to each IOP channel during initialization. The busy flag indicates whether the IOP is busy or ready to perform a new I/O operation. The CCW (channel command word) is specified by the CPU to indicate the type of operation required from the IOP. The CCW in the 8089 does not have the same meaning as the command word in an interfacing chip. The CCW here is more like an I/O instruction that specifies an operation for the IOP, such as start operation, suspend operation, resume operation, and halt I/O program. The parameter block contains variable data that the IOP program must use in carrying out its task. The task block contains the actual program to be executed in the IOP.

The CPU and IOP work together through the control and parameter blocks. The CPU obtains use of the shared memory after checking the busy flag to ensure that the IOP is available. The CPU then fills in the information in the parameter block and writes a ‘start operation’




Figure 13.6 | Location of information in memory for I/O operations in an 8089 system

command in the CCW. After the communication blocks have been set up, the CPU enables the channel attention signal to inform the IOP to start its I/O operation. The CPU then continues with its own program. The IOP responds to the channel attention signal by placing the address of the control block into its program counter. The IOP refers to the control block and sets the busy flag. It then checks the operation in the CCW. The PB (parameter block) address and TB (task block) address are then transferred into internal IOP registers. The IOP starts executing the program in the task block using the information in the parameter block. The entries in the parameter block depend on the I/O device. This, in short, is how an IOP carries out its designated function.

The 8089 was used in many applications in the early years of microprocessor development. However, IBM did not use it for its PC, and the IOP is no longer manufactured by Intel. Because of obsolescence, we will not discuss it in greater detail.

13.4 | Loosely Coupled Configurations and Bus Arbitration

In the case of the two kinds of tightly coupled systems just discussed, it is clear that the system bus is a scarce resource which must be shared and multiplexed in time, as only one master can use it at any time. In the 8086, 8087 and the 8089, special pins have been made available for requesting and granting the bus. Now, let us examine the case of loosely coupled multiprocessor systems (Fig 13.2). We see that there can be a number of processing modules in such a system. Each of these modules may have their own private resources, but that is not mandatory. There can also be modules whose resources are just the global resources (memory and I/O) only.

At any rate, one or more of the processing modules will need the system bus at some time or the other, and the chances of conflict are very high. To solve this problem, there should be some sort of a bus arbitration scheme for any system. Let us find out what schemes are available and used – the merits and demerits of each, how they are implemented, current trends and so on.

13.4.1 | Bus Arbitration Schemes

For any system, the arbitration scheme designed to be used must balance two factors.

- i) Priority: The highest priority module must be serviced first.
- ii) Fairness: Even the lowest priority module should be able to get service.

In all schemes, there is a central arbiter which will act to control and restrict bus access. We will discuss three simple schemes for bus arbitration.

13.4.2 | Daisy Chaining

This scheme uses three lines – bus request, bus grant and bus busy as shown in Fig 13.7a. Each of these lines is shared by all the potential bus masters which are ‘daisy chained’ i.e., connected in a cascaded fashion. The priority of the modules is fixed by their physical connection – left to right, meaning that the modules closer to the central bus arbiter have the higher priorities. One or more modules may place a bus request on the common line which is received by the central arbiter. It sends out a bus grant signal, provided the ‘bus busy’ line is inactive. This bus grant signal propagates from left to right and is accepted by the first module which had asked for the bus. It activates the bus busy signal and uses the bus. Thus the bus grant signal does not propagate beyond this.




Figure 13.7a | Daisy chaining scheme of bus arbitration




Figure 13.7b | Parallel arbitration scheme of bus arbitration

This scheme is very simple, but is obviously not a fair scheme. The priority of the modules is fixed up by the physical connection which cannot be changed once wired up, and a low priority module may be locked up permanently. The delay in propagating the bus grant signal from the left to right also limits the number of modules that can be accommodated in the system.

13.4.3 | Parallel Arbitration Scheme

This is also called the independent requests scheme. Here, a master module may output the request (Req) only if the system bus is not busy. The centralized arbiter issues a Bus Grant




Figure 13.7c | Distributed arbitration by self selection

signal to the module originating the highest priority request. Thus, this module can acquire the control of the system bus and it then activates the common 'bus busy' line. This is depicted in Fig. 13.7b. The priority may be preset (in this case it can happen that low priority master modules cannot access the system bus at all, or after a long delay only), or it may use a round-robin scheme (the highest priority at a particular time instant is assigned to the right side neighbor of the master module currently using the system bus). In the latter case all master modules have equal chance in accessing the system bus in a longer time period. In the case of very large systems a distributed version of the arbiter is used usually (the centralized version shown is failure critical). The demerit of this scheme is each module needs one signal line for bus request and one for bus grant – two dedicated lines just for bus access.

13.4.4 | Distributed Arbitration by Self Selection

This is a type of polling scheme. See Fig. 13.7c. Each module has an arbitration number with an associated priority. When multiple modules ask for the bus, the one with the highest arbitration number gets the bus. When bus conflicts arise, it will be resolved in favor of the module with the highest arbitration number. Along with a request, a module will also make known its arbitration number to the others. Each requesting device will compare this number to its own number and the one with the highest number wins. This scheme should allow dynamic reallocation of arbitration numbers to make it a truly fair scheme.

13.5 | Bus Arbitration Using the 8289 Bus Arbiter IC

Now that we have seen the issues and possible solutions to bus conflicts in a loosely coupled multiprocessing scheme, it becomes obvious to us that the 8086 does not have the pins or the extra logic to implement the bus arbitration schemes discussed. It has a $\overline{LOCK}$ pin to prevent any other processor module from taking over the bus from the current bus master, but that is all. This means that specialized hardware is necessary to handle bus arbitration, and that comes in




Figure 13.8 | Internal block diagram of the 8289 bus arbiter

the form of the chip 8289, Intel's bus arbiter IC. This chip has the pins for sending bus requests, receiving bus grants, activating the busy signal and so on. However, for setting up priorities, changing priorities and so on, extra logic is still needed. Now, let us see the functional block diagram of the bus arbiter IC 8289. See Fig 13.8 which shows the internal block diagram of the 8289 bus arbiter.

13.5.1 | Using the Bus Arbiter in a Multiprocessing System

For each module in a system, there will be a bus arbiter chip through which this module can access the 'system bus' (recall that there can be a private bus and private resources for each module, but we are not talking about that). When the particular module is using the system bus, its address lines, data lines are active. Otherwise, they should be tri-stated. This is to be ensured by the bus arbiter.

The 8289 bus arbiter operates in conjunction with the 8288 bus controller to interface 8086/8088 processors to a multi-master system bus. The processor is unaware of the arbiter's existence and issues commands as though it has exclusive use of the system bus. If the processor does not have the use of the multi-master system bus, the arbiter prevents the bus controller (8288), the data transceivers and the address latches from accessing the system bus (all bus driver outputs are forced into the high impedance state). Since the command sequence is not issued by the 8288, the system bus will appear as 'Not Ready' and the processor in the module will enter wait states. The processor will remain in 'Wait' until the bus arbiter acquires the use of the multi-master system bus, whereupon the arbiter will allow the bus controller, data transceivers, and address latches to access the system. Typically, once the command has been issued and a data transfer has taken place, a transfer acknowledgement (XACK) is returned to the processor to indicate 'READY' from the accessed device. The processor then completes its transfer cycle. Thus the arbiter serves to multiplex a processor (or bus master) onto a multi-master system bus and avoid contention problems between bus masters, meaning that, only one bus master can use the bus at any given time.

Figure 13.9 shows one module of a loosely coupled system (assuming only global resources which this module tries to access using the system bus). In a loosely coupled system, there will be many such modules tied to the common system bus.

13.5.2 | Bus Arbitration Between Different Processing Modules

In general, higher priority masters obtain the bus when a lower priority master completes its present transfer cycle. Lower priority bus masters obtain the bus when a higher priority master




Figure 13.9 | One processing module in a loosely coupled system, shown with the bus arbiter

is not accessing the system bus. A strapping option (ANYRQST) is provided to allow the arbiter to surrender the bus to a lower priority master as though it were a master of higher priority. If there are no other bus masters requesting the bus, the arbiter maintains the bus so long as its processor has not entered the HALT State. The arbiter will not voluntarily surrender the system bus and has to be forced off by another master's bus request, the HALT state being the only exception. Additional strapping options permit other modes of operation wherein the multi-master system bus is surrendered or requested under different sets of conditions.

13.5.3 | Using the Bus Arbiter in Various Arbitration Schemes

Since there can be many bus masters on a multi-master system bus, some means of resolving priority between bus masters simultaneously requesting the bus must be provided. We have discussed these methods, in general. Where does the bus arbiter come into this picture?

The 8289 Bus Arbiter supports these resolving techniques. All the techniques are based on a priority concept that at a given time one bus master will have priority above all the rest. There are provisions for using parallel priority resolving techniques, serial priority resolving techniques, and rotating priority techniques.

13.5.4 | Implementation of the Parallel Arbitration Scheme

The parallel priority resolving technique uses a separate bus request line $\overline{BREQ}$ for each arbiter on the multi-master system bus (see Fig. 13.10). Each $\overline{BREQ}$ line enters into a priority encoder which generates the binary address of the highest priority $\overline{BREQ}$ line which is active. The binary address is decoded by a decoder to select the corresponding $\overline{BPRN}$ (Bus Priority In) line to be returned to the highest priority requesting arbiter. The arbiter receiving priority ($\overline{BPRN}$ true) then allows its associated bus master onto the multi-master system bus as soon as it becomes available (when the bus is no longer busy). When one bus arbiter gains priority over another arbiter, it cannot immediately seize the bus, but must wait until the present bus transaction is complete. Upon completing its transaction, the present bus master recognizes that it no longer has priority and surrenders the bus by releasing $\overline{BUSY}$. $\overline{BUSY}$ is an active low, 'wired OR' signal line which goes to every bus arbiter on the system bus. When $\overline{BUSY}$ goes inactive (high), the arbiter which presently has bus priority ($\overline{BPRN}$ true) then seizes the bus and pulls $\overline{BUSY}$ low to keep other arbiters off of the bus. See waveform timing diagram, in Fig 13.11.

Note All multimaster system bus transactions are synchronized to the bus clock $\overline{BCLK}$.

13.5.5 | Daisy Chaining Scheme Using 8289

The serial priority resolving technique eliminates the need for the priority encoder-decoder arrangement by daisy-chaining the bus arbiters together and connecting the higher priority bus arbiter's $\overline{BPRO}$ (Bus Priority Out) output to the $\overline{BPRN}$ of the next lower priority (see Fig 13.12). The number of arbiters that may be daisy-chained together in the serial priority resolving scheme is a function of BCLK and the propagation delay from arbiter to arbiter. Normally, at 10MHz, only 3 arbiters may be daisy chained.

Note We have discussed the use of the bus arbiter IC in a loosely coupled multiprocessing system where multiple modules may try to access the system. The bus arbiter IC provides signal lines to manage this situation very efficiently. There are various other pins of the IC and various modes of operation also. Interested readers are requested to refer to the datasheet for more details.




Figure 13.10 | Implementing the parallel arbitration scheme using an 8289




Figure 13.11 | Timing of a higher priority arbiter obtaining the bus from a lower priority arbiter

Note

1. Higher priority bus arbiter requests the Multi-Master system bus.
2. Attains priority.
3. Lower priority bus arbiter releases $\overline{BUSY}$.
4. Higher priority bus arbiter then acquires the bus and pulls $\overline{BUSY}$ down.

13.5.6 | Multibus

You must have noted the word ‘multibus’ in some of the diagrams related to multiprocessing. This is a standard bus defined by Intel and is a multimaster bus, which means that it caters to multiple processors or has multiprocessing capability. The multibus system architecture was developed in 1975 by Intel Corporation to make microprocessors easier to use. The original multibus contained three qualities that made it very popular in the older days of computer development. The first of these qualities is its standardization. The multibus specification (IEEE 796) is precise enough such that multibus boards from different vendors are fully compatible. The second quality is multiprocessing — the ability to have multiple CPU boards in one system. This ability is the basis for distributed processing and allows a complex design to be built from easily managed

Table 13.3 | Comparison of speed of execution between 8086 emulation and an 8086 – 8087 co-processor system

Execution Times for Selected 8086/8087 Numeric Instructions and Corresponding 8086 Emulation		
Floating point instruction	Approximate execution time (μs)	
	8086/8087 (8 MHZ Clock)	8086 Emulation
Add/Subtract	10.6	1000
Multiply (single precision)	11.9	1000
Multiply (extended precision)	16.9	1312
Divide	24.4	2000
Compare	5.6	812
Load (double precision)	6.3	1062
Store (double precision)	13.1	750
Square Root	22.5	12250
Tangent	56.3	8125
Exponentiation	62.5	10687

emulation i.e., the same arithmetic operation done by a program written using 8086 instructions. The speed factor is formidable, and this justifies the use of a dedicated processor for such complex calculations.

In section 13.2.1, we have done a detailed discussion of how the host processor (8086) and the coprocessor communicate and execute instructions concurrently. In the following sections, the attempt will be to understand the features of the 8087. To do that effectively, the pre-requisite is an understanding of the philosophy and implementation of floating point arithmetic. Intel designed the 8087 in the year 1980. At that time, there were no standards defined for floating point arithmetic, but in 1985, many computer design companies including Intel set out to finalize a standard and it is the IEEE 754 standard. Intel’s 8087 was developed previous to the standardization, but we find that 8087 was designed using many of the concepts which later became the standard. The coprocessor 80287 was the first one to implement the standard fully. So, to understand the functioning of arithmetic processors, it is imperative to know something about this standard.

13.6.1 | Computer Arithmetic

There are several ways to represent real numbers on computers. In computing, a **fixed-point number** representation is a real data type for a number that has a fixed number of digits after the decimal (or binary or hexadecimal) point. For example, a fixed-point with 4 digits after the decimal point could be used to store numbers such as 1.3467, 281243.4356 and 0.1000, but would round 1.0301678 to 1.0302 and 0.0000678 to 0.0001. However, very few computer languages include built-in support for fixed point values, because for most applications, floating-point representations are better, faster and more accurate. Floating-point representations are more flexible than fixed-point representations, because they can handle a wider dynamic range. Floating-point representations are also slightly easier to use, because they do not require programmers to specify the number of digits after the decimal point.

13.6.2 | Floating Point Number Representation

Floating point numbers are represented with three components – the sign, biased exponent and mantissa. To elaborate, let us consider a few examples.

The numbers 6789 and 0.006789 can also be represented as:

$6789 = 6.789 \times 10^3 = 67.89 \times 10^2$ and so on.
 $0.0006789 = 6.789 \times 10^{-4} = 67.89 \times 10^{-5}$ and so on.

Thus a fractional number can be represented in various forms, with different exponent values where the exponent value is the power of 10 in this notation. The power of 10 is an indication of the position of the decimal point. In the above examples, there are four significant digits in the representation. However, by convention, when the decimal point is placed to the right of the first non-zero significant digit, the number is said to be normalized. In the normalized form, these numbers are 6.789×10^3 and 6.789×10^{-4}

In computers, binary representation of numbers is used. So, let us consider a few binary numbers represented in base 2 notation, with a sign included.

$+111.00111 = +1.1100111 \times 2^2$ in normalized form
 $+0.001110111 = +1.111011 \times 2^{-3}$ in normalized form
Both these numbers are positive. A negative number could be -1.11×2^{-3} .

In such normalized representations, two things can be made ‘implied’: The base can be implied to be ‘2’ and the number of significant digits before the decimal point is also implied (to be one). The conclusion is that ‘a floating point representation has a sign, a number of significant digits called the mantissa, and an exponent which is the power of the implied base (2 in binary)’. In the IEEE standard, there are two defined formats for floating point representation of binary numbers. See Table 13.4 which shows the bits assigned for each component of the floating point representation.

13.6.3 | Single Precision

32 bits are used for this representation, one-bit for the sign (‘0’ for plus and ‘1’ for minus), 8 bits for the ‘biased’ exponent and 23 bits for the fractional part of the mantissa. Only the bits after the decimal point of the mantissa are specified, because due to normalization, it is known that there is a ‘1’ to the left of the decimal point. See Fig 13.13.




Figure 13.13 | Single precision representation

Table 13.4

Formats	Sign	Exponent	Fraction	Bias
Single precision	1 [31]	8 [30–23]	23 [22–00]	127
Double precision	1 [63]	11 [62–52]	52 [51–00]	1023

Note The word ‘biased’ exponent. To the exponent, a value of 127(7FH) is added. This is to ensure that the representation of negative exponents do not require a sign bit. So when the exponent value is -67 , the biased exponent is $127 - 67 = 60$ i.e., 111100. For an exponent of 120, the biased exponent would be $120 + 127 = 247$, i.e., 11110111. The biased exponent is allowed only a range of 1 to 254. The end values of 0 and 255 are used to indicate special values.

0	00111100	111100-----0
---	----------	--------------

The value of the number represented as above = $+1.1111 \times 2^{-67}$. The biased exponent is 60. So the actual exponent should be $60 - 127 = -67$.

1	11011100	100110010-----0
---	----------	-----------------

The value of the number is = $-1.10011001 \times 2^{93}$. The biased exponent is 220. So the actual exponent should be $220 - 127 = 93$.

13.6.4 | Double Precision

As shown in Fig 13.14, in this, 64 bits are used, one for the sign, 11 bits for the biased exponent, and 52 bits for the fractional part of the mantissa. The bias for the exponent is 1023(3FFH).

Intel uses the terms ‘short real’ and ‘long real’ for the above two types of representations. Now let us list the steps used to convert a decimal number to the binary formats to conform to the IEEE standard.

- i) Convert the decimal number to binary form
- ii) Perform normalization on the number.
- iii) Calculate the exponent.
- iv) Add the bias to the exponent.
- v) Place the sign, biased exponent and the fractional part of the mantissa in the appropriate positions. The mantissa part is also referred to as the ‘significant’.




Figure 13.14 | Double precision representation

Example 13.1

Perform the conversion of the following decimal numbers, as specified.

- i) 230.25 into short real and
- ii) -6765.1875 into long real
- iii) 85.27 into short real

Solution

- i) Short real corresponds to the single precision format. We will do the conversion using the steps outlined earlier in Section 13.6.4.

Number: 230.25 Sign bit: 0
Binary representation: 11100110.01
Normalized form: 1.110011001×2^7
Biased exponent: $7 + 127 = 134 = 10000110$
Single precision format is:

0	10000110	11001100100-----0
---	----------	-------------------

This must be put into 32 bits, which is

0100 0011 0110 0110 0100 0000 0000 0000, which is represented in hex form as
4 3 6 6 4 0 0 0.
i.e., 43664000H is the single precision (short real) representation.

- ii) Long real corresponds to the double precision format. We will do the conversion using the steps outlined in Section 13.6.4.

Number: -6765.1875
Sign bit: 1
Binary representation: 1101001101101.0011
Normalized form: $1.1010011011010011 \times 2^{12}$
Biased exponent: $12 + 1023 = 1035 = 10000001011$

63	62.....52	51.....0
1	10000001011	10100110110100110-----0

This must be put into 64 bits, which is 1100 0000 1011 1010 0110 1101 0011 0000 0000 0000 0000 0000 0000 0000 0000 0000 In hex, this comes to be.
C0BA 6D30 0000 0000 H.

- iii) Short real corresponds to the single precision format.

Number: 85.27
Sign bit: 0
Binary representation: 1010101.01000101000111101
The binary representation of 0.27 will not terminate, and can be continued. However, because of limited precision, we have to limit it to fit into the number format given.
Normalized form: $1.01010101000101000111101 \times 2^6$
Biased exponent: $6 + 127 = 133 = 10000101$

Single precision format is:

0	1000 0101	0101 0101 0001 0100 0111 101
---	-----------	------------------------------

This must be put into 32 bits, which is:

0100 0010 1010 1010 1000 1010 0011 1101
 4 2 A A 8 A 3 D

i.e., 42AA 8A3D

Now, we will consider the act of converting from the IEEE format to a decimal number. The steps are:

- i) Separate the number into the sign, biased exponent and mantissa.
- ii) Remove the bias of the exponent.
- iii) Write the number in the normalized format
- iv) Convert to a de-normalized binary format.
- v) Convert the above to decimal format.

Example 13.2

Convert the short real number 42E64000H to a decimal format.

Solution

Number in Hex : 43664000H

- i) Number in short real format: 0100 0011 0110 0110 0100 0000 0000 0000
- ii) Separating the sign, exponent and mantissa
 Sign: 0, Biased exponent: 1000 0110, Mantissa part: 110 0110 0100 0000 0000 0000
- iii) Remove the bias of the exponent:
 1000 0110 is 134. Subtracting 127, it is 7,
 The exponent is 7.
- iv) Number in normalized format: 1.110011001 (excluding all the extra zeros and adding a '1' before the decimal point).
 The number can be written as 1.110011001×2^7
- v) In de-normalized form, it is 11100110.01
- vi) Convert to decimal form i.e., 230.25

Thus the conversion done in this example has verified the conversion done in Example 13.1.

13.6.5 | Special Values

The IEEE standard has specified some special values, which may appear in the course of calculations and need special attention. Let us try to decipher some of the related terms.

Zero Zero is not directly represented in the straight format, due to the assumption of a leading 1 (would need to specify a true zero mantissa to yield a value of zero). Zero is a special value denoted with an exponent field of zero and a fraction field of zero. Note that -0 and $+0$ are distinct values, though they both compare as equal.

Table 13.5 | Special Operations Defined by the IEEE Format

Operation	Result
$n \div \pm \text{Infinity}$	0
$\pm \text{Infinity} \times \pm \text{Infinity}$	$\pm \text{Infinity}$
$\pm \text{nonzero} \div 0$	$\pm \text{Infinity}$
$\text{Infinity} + \text{Infinity}$	Infinity
$\pm 0 \div \pm 0$	NaN
$\text{Infinity} - \text{Infinity}$	NaN
$\pm \text{Infinity} \div \pm \text{Infinity}$	NaN
$\pm \text{Infinity} \times 0$	NaN

Denormalized If the exponent is all 0s, but the fraction is non-zero (else it would be interpreted as zero), then the value is a denormalized number, which does not have an assumed leading 1 before the binary point. Thus, this represents a number $(-1)^s \times 0.f \times 2^{-126}$, where s is the sign bit and f is the fraction. For double precision, denormalized numbers are of the form $(-1)^s \times 0.f \times 2^{-1022}$. From this you can interpret zero as a special type of denormalized number.

Infinity The values +infinity and -infinity are denoted with an exponent of all 1s and a fraction of all 0s. The sign bit distinguishes between negative infinity and positive infinity. Being able to denote infinity as a specific value is useful because it allows operations to continue past overflow situations. Operations with infinite values are well defined in the IEEE floating point format.

Not a Number The value NaN (Not a Number) is used to represent a value that does not represent a real number. NaN's are represented by a bit pattern with an exponent of all 1s and a non-zero fraction. There are two categories of NaN: QNaN (Quiet NaN) and SNaN (Signalling NaN).

A QNaN is a NaN with the most significant fraction bit set. QNaNs propagate freely through most arithmetic operations. These values pop out of an operation when the result is not mathematically defined. An SNaN is a NaN with the most significant fraction bit clear. It is used to signal an exception when used in operations. SNaNs can be handy to assign to uninitialized variables to trap premature usage. Semantically, QNaNs denote indeterminate operations, while SNaNs denote invalid operations

Special Operations Operations on special numbers are well-defined by IEEE. In the simplest case, any operation with a NaN yields a NaN result. Other operations are as mentioned in Table 13.5.

13.6.6 | Features of 8087

Now that we have briefly discussed the IEEE format, let us go back to our study of the 8087 features. The 8087 math coprocessor has 68 instructions, for complex math operations including logarithmic, arithmetic, exponential and trigonometric functions. Recall that 8086 has only 8-bit and 16-bit data processing capability. 8087, on the other hand has a register size of 80 bytes and allows the usage of the following number formats as shown in Table 13.6 and confirms to the IEEE floating point standard. It can use integers, BCD numbers, single and double precision

Table 13.6 | Number Representations Used by the 8087

Data formats	Range precision	Precision	Directive	Total number of bits
Word integer	10^4	16 bits	DW	16
Short integer	10^9	32 bits	DD	32
Long integer	10^{18}	64-bit	DQ	64
Packed BCD	10^{18}	18 digits	DT	80
Short real	$10^{+/-38}$	24 bits	DD/REAL4	32
Long real	$10^{+/-308}$	53 bits	DQ/REAL8	64
Temporary real	$10^{+/-4932}$	64 bits	DT/REAL10	80

real numbers, as well as a ‘temporary real’ format with a word length of 80 bits. It uses the directives DW, DD or REAL4, DQ or REAL8 and DT or REAL10 as shown in the table.

Recall from the previous discussion (Section.13.4.2), that for real numbers, an additional number of bits are used for representing the biased exponent. Thus ‘short real’ totally needs 32 bits, ‘long real’ needs 64 bits and ‘temporary real’ needs 80 bits. Also note that DD stands for ‘define double word’, DQ for ‘define quad word’ and DT for ‘define ten bytes’.

Before we go into the full architectural details, it will be worthwhile understanding the tricks of programming using the coprocessor. We will use the instructions of 8087, which is the basic coprocessor, keeping in mind that the higher order coprocessors have the same set of instructions plus a few additional ones – but the newer coprocessors are faster, because of the advancement in technology. We can run the coprocessor programs on a PC, because all PCs generally have a coprocessor, unless specifically disabled.

However, before we start programming, we need to know the register structure and addressing modes. We will not attempt to go through the complete instruction set; instead the technique will be explained using a few instructions. The complete set of instructions is available in Appendix E.

13.6.7 | Registers of 8087

We will start with the register structure of the processor. Figure 13.15 shows the control and status registers, as well as the instruction pointer and data pointer. The 8087 has also a set of eight, 80-bit registers, which function as the general-purpose registers. Any data which is to be used for computation is converted into the temporary real format and then these registers are used in the computation process. The reason why this is called the ‘temporary format’ is that intermediate results are represented in this 80-bit format and only the final result will be truncated or rounded, when it has to be stored in the memory.

So $(A \times B)/C$ will be computed by multiplying A and B – this temporary product will be in the 80-bit format, and after division by C, if it has to be stored in memory in a 16-bit format, it will then be truncated or rounded if the word length allowed for it is less. However, the intermediate product $A \times B$ has a large word length, which will improve the accuracy of the final result.

The registers are labeled as ST meaning stack top. This is because they operate like a stack of 8 locations. At any time, one of the registers is considered as the stack top and this one is then designated as ST(0). This is not the name of any specific register. It just indicates the stack top.




Figure 13.15 | Register structure of the 8087

On initialization, one register is the stack top (ST(0)). Any data loaded thereafter, will go to this register if no other destination is specified. The next data will be pushed to the next register, which will then be designated as ST(0). The earlier ST(0) now becomes ST(1). Thus, data when loaded into the registers gets pushed into the stack. So what happens when all the registers (8 of them) get used up? The stack operates in a circular fashion, and the earliest data gets overwritten.

See Fig 13.16 for the way these registers are used when the data NUM1, NUM2 and NUM3 are loaded one after the other (NUM1 being loaded first and NUM3 last). See a set of tag registers corresponding to each of the 8 registers in the set. The tag word marks the content of each register as shown in Fig 13.17. The tag words can be used to interpret the contents of the registers. The tag word (2 bits) indicates whether the register content is valid, zero, special or empty.

13.6.8 | Programming the 8087

13.6.8.1 | Addressing Modes

Data that is used by the co-processor may be in memory or in any register of the register set. The addressing modes dictate that either the source and destination may be both registers, or one of





Figure 13.16 | Using the general-purpose registers for loading NUM1, NUM2 and NUM3



TAG VALUES

- 00 = VALID
- 01 = ZERO
- 10 = SPECIAL
- 11 = EMPTY

Figure 13.17 | Tag word interpretations

them is a register while the other is a memory location. When registers are used, there are ways of addressing in which ST(0) is implicit and need not be mentioned in the instruction.

13.6.8.2 | Instruction Format

All the instruction mnemonics start with the letter F. So FADD indicates addition, and FMUL represents multiplication and so on.

Now let us see a few of them. We will start with the data transfer instructions which consist of ‘load and store’ instructions only. Loading corresponds to taking data from memory to a

register. Since data can be in the form of integer, floating point or BCD, the load operation has three variations:

- i) FLD source
- ii) FILD source
- iii) FBLD source

The first instruction takes a real data from the source (memory or register), converts it to temporary real format and loads it into ST(0), the current stack top. Thus, the destination is implicit. Now the register below it is the stack top. The next two instructions in the above list do the same things – the only difference is that the source data (if in memory) is an integer or BCD. Keep in mind that data in registers can be only in the ‘temporary real format’ (in the instruction, ‘T’ means integer and ‘B’ implies BCD data).

FLD NUM1	;load into ST (0) the real data in memory location NUM1
FLD ST(2)	;load into ST(0) the data in ST(2)
FILD NUM2	;load into ST(0) the integer data in memory location NUM2

For storing, there are similar instructions:

FST destination	;store into destination
FIST destination	;store into destination after converting to integer
FBST destination	;store into destination after converting to BCD

In this case, the processor converts from the temporary real format to short or long real, integer or BCD (as the case may be) before storing into the destination. The source is always implied to be the stack top.

Example 13.3a

The following program loads two numbers in memory locations, to the registers. They are added and the sum is placed in a third memory location. Only integer data are used. Hence, the instruction format contains the letter ‘T’ in it.

```

      .8087
      .MODEL SMALL
      .DATA
NUM1 DD 09
NUM2 DD 4
NUM3 DD ?
      .CODE
      .STARTUP
      FINIT                ;initialize the 8087
      FILD NUM1            ;load the integer in NUM1 to ST(0)
      FILD NUM2            ;load NUM2 to ST(0). Now ST(1)
                           contains NUM1
      FADD ST(0), ST(1)    ;add the two registers-sum in ST(0)
      FIST NUM3            ;save the sum in NUM3
      .EXIT
      END

```

Example 13.3b

```
13BD:0017 9B          WAIT
13BD:0018 DBE3          FINIT
13BD:001A 9B          WAIT
13BD:001B DB060000      FILD      DWORD PTR [0000]
13BD:001F 9B          WAIT
13BD:0020 DB060400      FILD      DWORD PTR [0004]
13BD:0024 9B          WAIT
13BD:0025 D8C1          FADD      ST,ST(1)
13BD:0027 9B          WAIT
13BD:0028 DB160800      FIST      DWORD PTR [0008]
13BD:002C B44C          MOV      AH, 4C
-
```

Example 13.3b shows the program in 13.3a, as observed in the debugger in the ‘trace’ mode.

Note The WAIT instruction gets automatically added. This is for synchronizing the operation of the coprocessor to the host processor (refer Section.13.2.1).

Example 13.3c

```
-d0000
13C0:0000 09 00 00 00 04 00 00 00-0D 00 00 00 00 00
-----
-----
```

Example 13.3c shows the content of memory after the addition.

Note The stored numbers are in bold. All the numbers are in the integer format.

Example 13.4a

Add two real numbers, which are in memory and store the result in memory.

Solution

```
.8087
.MODEL SMALL
.DATA
NUM1 DD 230.25
NUM2 DD 85.27
NUM3 DD ?
.CODE
.STARTUP
FINIT          ;initialize the 8087
FLD NUM1       ;load NUM1 to ST(0)
```

```

    FLD NUM2           ;load NUM2 to ST(0). Now ST(1) has NUM1
    FADD ST(0), ST(1)  ;add contents of two registers
    FST NUM3
    .EXIT
    END

```

Example 13.4b

Microsoft (R) Macro Assembler Version 6.14.8444

```

                                .8087
                                .MODEL SMALL
                                .DATA
0000                            NUM1 DD 230.25
0000 43664000                 NUM2 DD 85.27
0004 42AA8A3D                 NUM3 DD ?
0008 00000000
0000                            .CODE
                                .STARTUP
0017 9B DB E3                  FINIT
001A 9B D9 06 0000 R           FLD NUM1
001F 9B D9 06 0004 R           FLD NUM2
0024 9B D8 C1                  FADD ST(0), ST(1)
0027 9B D9 16 0008 R           FST NUM3
                                .EXIT
                                END

```

In Example 13.4a, we have used two short real numbers, which are added and the sum found. Example 13.4b shows the list file in which the hexadecimal forms of the two numbers are seen. We have used the same decimal numbers of Example 13.1. The list file shows that the number 230.25 has been converted to **43664000H** and the number 85.27 to **42AA8A3DH**.

The two programming examples highlight the use of the coprocessor instructions. To learn more about the control, status and other registers of the 8087, the data sheet of the chip may be referred to. To become proficient in programming, refer to Appendix E and run a number of programs.

KEY POINTS OF THIS CHAPTER

- When a system contains more than one processor, it is called a multiprocessing system.
- Such systems are broadly classified as tightly coupled and loosely coupled systems.
- The 8086 has a number of inbuilt features to support multiprocessing.
- The processor should be in the maximum mode, when multiprocessing features are to be enabled.

- When the 8086 is used along with the 8087 arithmetic processor, it is called a coprocessor configuration.
- The 8086 and 8087 work in close co-operation in the coprocessor configuration.
- The 8086 IOP is another processor which can be used in a system along with the 8086.
- For large systems, multiprocessing exists in the form of a closely coupled configuration.
- In such cases, many modules will contend for the bus and that will need some specific protocols to resolve the conflict.
- For large systems, there is a bus arbiter IC 8289 which is expected to handle such problems.
- The 8087 is a specially designed IC to handle arithmetic computations along with the 8086, in a system.
- For each processor of the Intel x86 family, there is a specific arithmetic coprocessor.
- For floating point operations, IEEE has defined a format and standards.
- Floating point numbers can be represented in the single precision or double precision formats.
- MASM can be used to run programs using the instruction set of 8087.

QUESTIONS

1. Distinguish between loosely coupled and tightly coupled multiprocessor systems.
2. How is communication done between the processing modules of a tightly coupled system?
3. Name the features that 8086 has for multiprocessing.
4. In a coprocessor configuration, the coprocessor is dependent on the host processor. Elaborate.
5. How do the BUSY-TEST pin connections between an 8086 and 8087 synchronize the operations of the 8086 and 8087?
6. How has arithmetic processing developed and changed over the years, with respect to Intel's processors?
7. Why are general-purpose processors not equipped with complex arithmetic instructions?
8. Why is the 8089 designated as 'an intelligent DMA controller'?
9. Why is it necessary to have a dedicated bus arbiter IC in a loosely coupled multiprocessing system?
10. What is meant by the term 'global resources'?
11. Which is the most efficient of the bus arbitration schemes that we have discussed?
12. What is 'Multibus' and what are its features?
13. Explain the necessity of special floating point representation of numbers and differentiate between single and double precision numbers.
14. Why is a bias added to the exponent in floating point number representation?
15. What is meant by the temporary real format and why does this format have a large word size?

EXERCISE

1. Draw a multiprocessing system with an 8086, 8087 and 8089.
2. Express the following real numbers in single precision and double precision format.
 - a) 3463.75
 - b) 89.125
 - c) 8945683.875

3. Do the reverse process of converting from floating point format to decimal format for the numbers of Question 2.
4. Read the data sheet of 8087 and list the exceptions produced by 8087.
5. Write programs using the instruction set of 8087.
 - a) To find the circumference of a circle of radius 45.26.
 - b) Find the area of this circle.
 - c) Use the above two programs as procedures to find the area of circles of any given radius.
6. Write a program to find the sine, cosine and tan of angle 45° .
7. Find the instruction of 8087 for random number generation and use it in a program.

14 80186 – THE EMBEDDED MICROPROCESSOR



In this chapter, you will learn

- The distinguishing features of the 80186 processor.
- The enhancements in the instruction set of this processor compared to 8086.
- The details of the peripheral blocks inside the chip.
- How each of these blocks operates.
- The complete details of the operation of the timer block inside the chip.

Introduction

Now that we have made a thorough and detailed study of the 8086 processor, it is time to learn about the higher order versions of the x86 microprocessor family. In this and the next two chapters, we will review the differences and enhancements of each of them.

Our first case for study is the 80186/80188 processor. In terms of bus widths, they are similar to the 8086/8088 processors, in the sense that both have an internal data bus width of 16 bits – but the 80186 has an external data bus width of 16 bits, while for the 80188, it is only 8 bits. Both have an address bus width of 20 bits. Having made clear this distinction, we will henceforth talk only about the 80186.

With the release of 8086, many system designers used it in embedded system design by adding timers, interrupt controllers and DMA controllers on its board. It was this fact that caused Intel to think in terms of releasing an embedded controller and the 80186 turned out to be the resultant chip. An embedded processor is to act like a ‘microcontroller’ which is a microprocessor with a lot of peripherals integrated on to the chip. For a microcontroller, the capability to manage peripheral devices is more important than its computing capability. Thus, the 80186 does not offer much more processing capability than the 8086 (except for 7 more new instructions and new operand types to three existing instructions.), but it does have a number of peripherals integrated in the chip. With the exception of integrated components, the Intel 80186 microprocessor is not very different from the 8086, and so, it may be considered as an embedded version of 8086.

This processor did not get used in PCs but became quite popular in the embedded market and many different versions of the chip are still available and in use. The 80186 series was generally intended for embedded systems, or as microcontrollers with external memory. It includes features such as a clock generator, interrupt controller, timers, wait state generator, DMA channels, and chip select logic – all these are available within the chip.

Chapter-opening image: A microcontroller chip.

History and Development

It was in 1982 that Intel introduced the 80186 family of embedded microprocessors. In 1987, the second generation i.e., 80C186 which was a CMOS version with smaller geometry and an enhanced feature set, was released, which was pin compatible with 80186. Later, the 80186 EB was introduced with a static, stand-alone module, known as the 80C186 Modular core, in which peripherals were redesigned with standard interfaces. The new 80C186 EB was the first to use this modular capability. The design trend was towards reducing power dissipation, which is a key parameter for an embedded product. Other models of the processor are 80C186XL, 80C186 EA and 80C186 EC. The last model has 14 on-chip peripherals. See Table 14.1.

14.1 | Additions in the Instruction Set

We will first discuss the new instructions of the 80186 processor, in comparison with its predecessor, the 8086.

- Data transfer instructions
PUSHA
POPA
- String instructions
INS
OUTS
- High-level instructions
ENTER
LEAVE
BOUND

Table 14.1 | List of 80186 Versions and Corresponding Features

Features	80C186 XL	80C186 EA	80C186 EB	80C186 EC
80286 like instruction set	Yes	Yes	Yes	Yes
Power-save	Yes	Yes	–	Yes
Power down mode	–	Yes	Yes	Yes
80187 interface	Yes	Yes	Yes	Yes
ONCE mode	Yes	Yes	Yes	Yes
Interrupt controller	Yes	Yes	Yes	Yes
Timer unit	Yes	Yes	Yes	Yes
Chip selection unit	Yes	Yes	Yes	Yes
DMA controller	Yes	Yes	–	Yes
Serial communications unit	–	–	Yes	Yes
Refresh controller	Yes	Yes	Yes	Yes
Watch dog timer	–	–	–	Yes
I/O Ports	–	–	Yes	Yes

14.1.1 | PUSH/POPA

PUSHA (Push All) and POPA (Pop All) allow all general-purpose registers to be pushed and popped respectively. The order of pushing is AX, CX, DX, BX, SP, BP, SI, DI. The Stack Pointer (SP) value pushed is the Stack Pointer value before the AX register was pushed. When POPA is executed, the Stack Pointer value is popped, but ignored. Note that this instruction does not save segment registers (CS, DS, SS, ES), the Instruction Pointer (IP), the Processor Status Word or any integrated peripheral registers.

14.1.2 | String Instructions

14.1.2.1 | *INS* source_string, port

INS (In string) performs block input from an I/O port to memory. The port address is placed in the DX register. The memory address is placed in the DI register. This instruction uses the ES segment register, which cannot be overridden. After the data transfer takes place, the pointer register (DI) increments or decrements, depending on the value of the Direction Flag (DF). The pointer registers change by one, for byte transfers and by two, for word transfers.

14.1.2.2 | *OUTS* port, destination_string

OUTS (Out string) performs block output from memory to an I/O port. The port address is placed in the DX register. The memory address is placed in the SI register. This instruction uses the DS segment register, but this may be changed with a segment override instruction. After the data transfer takes place, the pointer register (SI) increments or decrements, depending on the value of the Direction Flag (DF). The pointer registers changes by one for byte transfers and by two for word transfers.

14.1.3 | ENTER and LEAVE

These instructions are used with stack frames in the case where the stack is used to pass parameters to and from procedures. They are useful in high level language programming. It creates the stack frame required by most block-structured high-level languages.

The ENTER instruction is equivalent to
PUSH BP
MOV BP, SP

Thus, stack variables can be accessed through the BP register. The ENTER instruction has two operands – the number of bytes to reserve for variables on the stack, and the level of nesting of the procedure. The lexical nesting level determines the number of pointers to higher level stack frames copied into the current stack frame. This list of pointers is called the display. The first word of the display points to the previous stack frame. The display allows access to variables of higher level (lower lexical nesting level) procedures. This instruction is useful for block structured high level languages which use nesting levels to control access to variables of previously nested procedures.

For example, ENTER 6,0 reserves 6 bytes of storage in stack and specifies the nesting level as 0. It leaves BP at the top of the stack and subtracts 6 from SP to allow 6 bytes of storage. The LEAVE instruction reverses all the above. It is equivalent to the instructions:

MOV SP, BP
POP BP

14.1.4 | BOUND Register, Address

BOUND verifies that the signed value in the specified register lies within specified limits. If the value does not lie within the bounds, an array bounds exception (INT 5) occurs. BOUND is useful for checking array bounds before attempting to access an array element. This prevents the program from overwriting information outside the limits of the array.

BOUND has two operands. The first, **register**, specifies the register being tested. The second, **address**, contains the effective relative address of the two signed boundary values. Array index in source register is checked against upper and lower bounds in memory source. The first word located at 'address' is the lower boundary and the word at 'address +2' is the upper array bound. Interrupt 5 occurs if the source value is less than or higher than the register content.

Consider the instruction BOUND BX, COST where COST is an address. The content of BX should not be less than COST or greater than COST+2. If this condition is violated, a Type 5 interrupt occurs.

14.2 | Instruction Set Enhancements

As mentioned earlier, the 80186 adds new operand types to three existing 8086 instructions, and this can be considered as 'enhancements'.

i) PUSH immediate

PUSH (Push immediate) allows an immediate argument, data, to be pushed onto the stack. The value can be either a byte or a word. Byte values are sign extended to word size before being pushed.

ii) For the following shift and rotate instructions, it allows count to be an immediate operand

- a) SHL
- b) SAR
- c) SHR
- d) ROL
- e) ROR
- f) RCL
- g) RCR

Recollect that for 8086, we need to load the count of the shift or rotate operations in the CL register if it is not = 1. Here we can write instructions such as:

```
SHL BX, 3
RCL AL, 5 and so on
```

iii) IMUL *destination, source, data*

IMUL (integer immediate multiply, signed) allows a value to be multiplied by an immediate operand. IMUL requires three operands. The first, *destination*, is the register where the result will be placed. The second, *source*, is the effective address of the multiplier. The source may be the same register as the destination, another register or a memory location. The third, *data*, is an immediate value used as the multiplicand. The *data* operand may be a byte or word. If *data* is a byte, it is sign extended to 16 bits. Only the lower 16 bits of the result are saved. The result must be placed in a general-purpose register.

Now, let us try using some of these instructions in programs. MASM 6.x allows the use of all instructions of a specified processor by the directives i.e., .186/.286/.386, depending on the instructions being used.

Example 14.1

The processor is connected to an input port and an output port, each of which has only 12 data lines. The input port (with address 0C24H) is connected to the data bus lines D₃ to D₁₅, while the output port (with address 0C25H) is connected to the data lines D₀ to D₁₅. However, since data can be read in and written only as bytes or words, some processing must be done to remove the data read from the unconnected pins.

There are three parts to the program. The first part inputs 10 data words, one after the other from the input port using the INS instruction and this assumes that the data continually changes and is available for reading in. This data is stored in the data segment in location TAKE. The second part of the program takes each word, masks the lowest nibble and shifts it right by 4 positions. Thus, the relevant data is in the lower three nibble positions, and is stored in the location GIVE.

The third part of the program uses the OUTS instruction to output this data string, one at a time to the output device.

Solution

```

.MODEL SMALL
.186
.DATA
TAKE DW 10 DUP(0)
GIVE DW 10 DUP(0)
.CODE
.STARTUP
    LEA DI, TAKE      ;point DI to the destination
    MOV DX, 0C24H     ;load the input port address in DX
    CLD              ;clear direction flag
    MOV CX, 10        ;count of the data
    REP INSW          ;input the string words until CX = 0
                     ;the 10 words are available TAKE
    MOV CX, 10        ;count of the data
    LEA BX, TAKE      ;point BX to the address of the data
TRANS:  MOV AX, [BX]   ;move the first word to AX
        AND AX, 0FFF0H ;AND it, to mask the lower nibble
        SHR AX, 4      ;rotate it right
        MOV 20[BX], AX ;save in locations 20 bytes away
        INC BX         ;increment the pointer
        INC BX         ;increment again as data is a word
        LOOP TRANS     ;repeat until CX=0
                     ;the valid 10 words are at GIVE
    MOV SI, GIVE      ;point SI to GIVE
    MOV DX, 0C25H     ;load the output port address in DX
    CLD              ;clear the direction flag
    MOV CX, 10        ;count of the data
    REP OUTSW         ;output the string words until CX = 0
    .EXIT
    END

```


The program uses the new instructions `INS`, `OUTS` and `SHR, n`. If the `.186` directive is not used, the message of ‘invalid instruction’ will be obtained.

14.3 | Block Diagram of the 80186

Figure 14.1 shows the internal block diagram of the 80186. The external pins are also seen in the diagram. Since we have already discussed the features of 8086, our focus here will be to have a quick at the components within the 80186, which are not present in 8086. We will discuss the peripherals of this processor, which we have not encountered in the 8086.




Figure 14.1 | Internal block diagram of the 80186

14.3.1 | Clock Generator

Recollect that for the 8086, an external clock generator IC (8284) was needed. For the 80186, the clock generator is inside, and it provides the 50% duty cycle clock needed by the processor. The CLKOUT pin provides the processor clock signal for use by other chips, if needed. The chip can operate at frequencies of up to 25 MHz, which is good enough for many embedded applications.

14.3.2 | Reset Logic

There is a RES input pin to reset the processor and a synchronized RESET output pin for use with other system chips.

14.3.3 | Peripheral Control Block

In general embedded processor architectures; there are registers for controlling each of the peripherals. Thus, each peripheral is associated with a set of registers and writing specific bit patterns into these registers amounts to making the peripherals to act as we choose (from a list of choices). The same features are used in this processor as well. There is a set of 256 sixteen-bit registers and this set is called the peripheral control block (PCB). Many registers of this block are unused and termed as ‘reserved’. The base address of this control block is decided by a 16-bit re-location register contained within the control block at offset FEH from the base address. On reset, the re-location register is set to 20FFH. This locates the PCB at I/O addresses FF00H to FFFFH onwards. The PCB can be re-located by using the re-location register. In Fig 14.2, the offset of each peripheral register from the base address of the PCB is shown.

Relocation Register	OFFSET FEH
DMA Descriptors Channel 1	DAH DOH
DMA Descriptors Channel 0	CAH COH
Chip – Select Control Registers	A8H A0H
Timer 2 Control Registers	66H 60H
Timer 1 Control Registers	5EH 58H
Timer 0 Control Registers	50H
Interrupt Controller Registers	3EH 20H

Figure 14.2 | Peripheral control block of 80186




Figure 14.3 | Peripheral control register

14.3.4 | Chip Select Logic

The processor provides programmable chip select logic (refer chapter 7) for memory and peripherals, and can also allow wait state generation. The chip select signals become active in the respective memory or I/O read/write cycles. The number and partitioning of chip select lines depends on the model number of the 80186 being used, but in general, the traits are as follows. It provides 6 different select outputs for different memory areas – upper memory, lower memory and mid-range memory. In the block diagram, you can see the pins $\overline{UCS}$ and $\overline{LCS}$. These are the select pins for the upper and lower memory respectively. The range of the memory area defined as ‘upper’ and lower is programmable by the control registers in the chip select block. You can also see three pins titled $\overline{MCS}_{0-3}$. These are the 4 pins which select different parts of the mid range memory-again, the range and base address of the portions of memory we use for different applications is programmable. It also provides peripheral (I/O) chip selects for up to 7 peripherals. The pins which generate the select output are the lines $\overline{PCS}_{0-6}$.

Next, let us have a look at the internal peripherals of the processor. As can be see in the block diagram, they are the timer unit, DMA unit and the interrupt controller unit.

14.3.5 | Timer Unit

There are three internal 16-bit programmable timers. Two of them have two pins – one is the timer out pin, on which signals generated can be outputted, the other is the timer in, which is the input for counting actions. The third timer is not connected to any external pins and is useful for real-time coding and time delay applications. In addition, the timer can be used as a pre-scaler for the other two timers.

The EC model of the processor has an additional watch dog timer. This is an important timer for embedded applications where user intervention is impossible. In such cases, the watch dog timer is set to run for a particular period, and at the end of this period i.e., when the counter counts down to zero, if the system is found to be in a hang state or the program goes into an infinite loop, the watch dog timer causes a system reset or activates an interrupt. Under normal and correct operation of the system, the watch dog timer never counts down to zero.

14.3.6 | Interrupt Controller

In the 8086, there are two hardware interrupt lines – one is maskable, the other is not. If more interrupt sources are expected, an external chip i.e., the Priority Interrupt Controller 8259A is used to handle them – this chip has the capability to prioritize and sequence the interrupts received, return acknowledge signals to the interrupting devices and also send interrupt vectors to the processors.

In the 80186, there is an interrupt controller on the chip (internal to the chip) which expands its interrupt capabilities. The block diagram (Fig 14.1) shows it having one NMI input and four INTR inputs. Though not seen in the block diagram, the interrupt controller has the duty of managing the interrupts from the on-chip peripherals too.




Figure 14.4 | Interrupt control unit in master mode

The interrupt controller operates in two modes – the master and the slave. In the master mode, the interrupt control unit controls the interrupts from the on-chip peripherals (timers and DMA units) and the four external interrupt pins. The unit synchronizes and prioritizes all interrupt sources and presents the correct interrupt vector to the CPU. Simple systems need to use only this master mode of operation.

See Fig 14.4 of the interrupt control unit in the master mode of operation. It accepts interrupts from the internal peripherals i.e., the timers and DMA unit and channels then to the CPU.

In the slave mode of operation, an external 8259A controls the maskable interrupt to the CPU and acts as the master interrupt controller. The interrupt control unit then processes only the interrupts from the on-chip peripherals. Such a design is used in large systems.

This section was meant to see the important aspects of the 80186 chip – without going into a great amount of detail. Each of the peripherals work just as some of the interfacing chips we have seen earlier. The DMA unit has all the features that a dedicated DMA controller (8237) has, the interrupt control unit operates like the 8259 we have seen in Chapter 10, and the timer unit can be used just like any such unit present in a typical microcontroller like 8051, and like the dedicated timer chip 8254/8253 (Chapter 10). Thus, we can acknowledge that the different versions of 80186 have an application range far beyond that of the 8086.

Now, just to get a feel of the usefulness of the chip, let us go into the details of one of the internal peripherals i.e., the Timer.

14.4 | Programming the Timer Unit

Let us try to understand the timer unit in very simple terms – what it is, what it can do and how it can perform the functions attributed to it.

The timer unit consists of three independent 16-bit timers 0, 1 and 2 and they operate independent of the CPU. Functionally timers 0 and 1 are identical. Timer 2 is used as a DMA request source, as a prescaler for other timers or as a watch dog timer.

Timing Sequence Any timer has a 16-bit **count register**. Whenever a timing event occurs, the number (count) in the register increments by 1. Thus, on a continuous sequence of timing events, the count keeps on incrementing until it is equal to the number stored in a

‘maximum count register’. At this point, the count register rolls over to zero. What has been achieved is a ‘delay period’ which depends on the numbers stored in the count register and the maximum count register. The normal thing is to start with a content of zero in the count register.

Timing Events A timing event is one that causes the count register to increment its count. A timer can count internal or external events. When a low to high transition occurs on the ‘Timer in’ pins of timers 0 and 1, its count registers are incremented. This implies that the timing source is external to the processor – we say that the timers ‘count external events’. Such a feature is available only for timers 0 and 1. Note, that in the block diagram (Fig 14.1), only timers 0 and 1 have ‘TIMER in’ pins. The maximum frequency allowed for the external counting signal is limited to one-fourth the CLKOUT of the 80186 processor. If no external events need to be counted, the counting source can be the processor clock and its timing event is then $\frac{1}{4}$ CLKOUT frequency or can be prescaled by Timer 2. What all this implies can be stated this way: **When the timing event is from an external source, the timer acts as a counter, otherwise it is a timer.**

Timers 0 and 1, each have a single output pin. The Timer output can be either a single pulse, indicating the end of a timing cycle, or a variable duty cycle wave. These two output options correspond to single maximum count mode and dual maximum count mode, respectively. Interrupts can be generated at the end of every timing cycle, and can cause the taking up of an ‘interrupt service routine’.

Timers 0 and 1 are functionally identical. Figure 14.5 illustrates their operation. Each has a latched, synchronized input pin and a single output pin. Each timer can be clocked internally or externally. Internally, the timer can either increment at $\frac{1}{4}$ CLKOUT frequency or be prescaled by Timer 2. A timer that is prescaled by Timer 2 increments when Timer 2 reaches its maximum count value. Timer 2 has no input or output pins and can be operated only in single maximum count mode.




Figure 14.5 | Timer/counter – internal block diagram




Figure 14.6 | Dual maximum count mode

Compare Register Modes The way the count register increments, and compares its value to the number in the Maximum count register has already been touched upon. This is when using the timer in the ‘single maximum count mode’ which is applicable for all the three timers. The timer counting from its initial count (usually zero) to its maximum count (Maxcount Compare A or B) and resetting to zero, defines one timing cycle. A Maxcount Compare value of 0 (special case) implies a maximum count of 65,536. For Maxcount Compare value of 4, the timer counts from 0 to 3 and then resets to 0 i.e., it thus, has four states.

However, for Timers 0 and 1 alone, there is another mode called the ‘dual maximum count mode’. Note in Fig 14.6, the maximum count registers A and B for timers 0 and 1. This allows the generation of asymmetric square waves very easily. The *modus operandi* is as follows. In this mode, Maxcount Compare A and Maxcount Compare B are both used. The timer counts to the value contained in Maxcount Compare A, resets to zero, counts to the value contained in Maxcount Compare B, and resets to zero again. This process can be repeated. Using both the registers allows the timers to count up to 131,072. The Register In Use (RIU) bit in the Timer Control register indicates which Maxcount Compare register is currently in use.

Timer Output pins In single maximum count mode, the timer output pin goes low for one CPU clock period. This occurs when the count value equals the Maxcount Compare A value. If programmed to run continuously, the timer generates periodic pulses. This needs ALT = 1 to be selected in the control register.

In dual maximum count mode, the timer output pin indicates which Maxcount Compare register is currently in use. A low output indicates Maxcount Compare B, and a high output indicates Maxcount Compare A. If programmed to run continuously, a repetitive waveform can be generated if alternate mode (ALT = 1) is selected. Thus, it only requires the right numbers to be loaded in these two compare registers to generate square waves of any duty cycle. Refer Example 14.2.

14.5 | Programming

Associated with each timer are a number of registers which are in the peripheral control block, with addresses (offset in the PCB) as shown in Fig 14.7.

Timers 0 and 1 have the following registers:

- i) Timer Control register – TCON
- ii) Compare register A




Figure 14.7 | Offsets and contents of the timer registers

EN	INH	INT	RIU	R	R	R	R	R	R	MC	RTG	P	EXT	ALT	CONT
----	-----	-----	-----	---	---	---	---	---	---	----	-----	---	-----	-----	------

Figure 14.8 | Register bits of TCON register 1 and 2

EN	INH	INT	R	R	R	R	R	R	R	MC	R	R	R	R	CONT
----	-----	-----	---	---	---	---	---	---	---	----	---	---	---	---	------

Figure 14.9 | Bit configurations of TCON2

iii) Compare register B

iv) Count register

Timer 2 has only three registers. It does not have the compare register B.

Control Registers – TCON1 and TCON2 They are 16-bit registers, the bit settings of which can choose the various options available for the operation of the corresponding timer. Fig 14.8 shows the control register bits for Timers 0 and 1. Table 14.2 explains the relevance of each bit of the registers.

For Timer 2, the Timer Control Register TCON2 has the format as shown in Fig 14.9. This has only 5 bits with the meanings exactly the same as in Table 14.2, the rest of the bits are marked as reserved.

Table 14.2 | Bit Designations of the TCON Registers

Bit	Bit Name	Function
EN	Enable	Set to enable the timer. This bit can be written only when the INH bit is set.
INH	Inhibit	Set to enable writes to the EN bit. Clear to ignore writes to the EN bit. The INH bit is not stored; it always reads as zero.
INT	Interrupt	Set to generate an interrupt request when the Count Register equals a Maximum Count Register. Clear to disable interrupt requests.
RIU	Register in use	Indicates which compare register is in use. When set, the current compare register is Maxcount Compare B; when clear, it is Maxcount Compare A.
MC	Maximum count	This bit is set when the counter reaches a maximum count. The MC bit must be cleared by writing to the Timer Control Register. This is not done automatically. If MC is clear, the counter has not reached a maximum count.
R	Reserved	For use in future processors-must be cleared now.
RTG	Re-trigger	This bit specifies the action caused by a low-to-high transition on the TMR INx input. Set RTG to reset the count; clear RTG to enable counting. This bit is ignored with external clocking (EXT = 1).
P	Pre-scaler	Set to increment the timer when Timer 2 reaches its maximum count. Clear to increment the timer at ¼ CLKOUT. This bit is ignored with external clocking (EXT = 1).
EXT	External clock	Set to use external clock; clear to use internal clock. The RTG and P bits are ignored with external clocking (EXT set).
ALT	Alternate compare register	This bit controls whether the timer runs in single or dual maximum count mode. Set to specify dual maximum count mode; clear to specify single maximum count mode.
CONT	Continuous mode	Set to cause the timer to run continuously. Clear to disable the counter (clear the EN bit) after each counting sequence.

14.5.1 | Asymmetric Square Wave Generation

Square waves of various duty cycles are frequently required. Example 14.2 shows the use of Timer 1 in generating a square wave of 80% duty cycle. The method used is as follows. Compare register A stores the count for the high time, and Compare register B stores the count for the low time. The Timer1 count register is initialized to 0. The control register is configured so that a continuous square wave is generated (ALT and CONT pins are 11). Assuming that the clock frequency is 8 MHz, the HIGH_TIME is calculated as (required pulse_width*f)/4. We need a high pulse width of 0.8 msecs. So HIGH_TIME = 1600. For a low pulse width of 0.2 msecs, LOW_TIME = 400. These are loaded in the Maximum count compare registers A and B.

Then the count register T1COUNT is reset to zero, and the control word is loaded into TCON1 to start the timer. The control word is 1100 0000 0000 0011 i.e., EN and INH bits are 11 and ALT and CONT bits are 11. Thus, it is C003H and has been labeled as CONWRD. This is to be loaded into TCON1 labeled here as T1CONTRL. Because the ALT and CONT bits are 11, this will generate a continuous signal with the specified duration and duty cycle.

Observe that the address of each of the registers is calculated by adding its corresponding offset (Fig 14.7) to the base address of the PCB. Here the default base address, on reset, is FF00H.

Example 14.2

;this program segment generates a square wave of given frequency and duty cycle on Timer 1 output pin.

```

TICMPA      EQU    0FF5AH
T1CMPB      EQU    0FF5CH
T1COUNT    EQU    0FF58H
T1CONTRL    EQU    0FF5EH
CONWRD      EQU    0C003H      ;ALT and CONT bits are 11
HIGH_TIME   EQU    1600
LOW_TIME    EQU    400

MOV DX, TICMPA
MOV AX, HIGH_TIME      ;set high time
OUT DX, AX

MOV DX, T1CMPB
MOV DX, LOW_TIME       ;set low time
OUT DX, AX

MOV DX, T1COUNT
MOV AX, 0               ;clear count register
OUT DX, AX

MOV DX, T1CONTRL
MOV AX, CONWRD          ;start Timer 1
OUT DX, AX

```

Example 14.3 shows the generation of a low pulse of specified width, after a specified delay. Thus, it functions as a one-shot. The delay is specified in the Compare A register and the width of the mono pulse is specified by the Compare B register. The calculation for the values in Compare A and B registers are exactly as in Example 14.3.

Example 14.3

;this program generates a low pulse at Timer 0 output, after a delay.

```

T0COUNT    EQU    0FF50H      ;substitute register offsets
T0CMPA      EQU    0FF52H

```

```

T0CMPB      EQU    0FF54H
T0CONTRL    EQU    0FF56H
MCBIT       EQU    0020H
CONWRD      EQU    0C002H          ;ALT and CONT bits are 10

        MOV DX, T0CMPA
        MOV AX, 20,000             ;specify a delay of 10 msecs
        OUT DX, AX

        MOV DX, T0CMPB
        MOV AX, 200                ;width of pulse = 100 microsecs
        OUT DX, AX

        MOV DX, T0COUNT          ;clear counter register
        MOV AX, 0
        OUT DX, AX

        MOV DX, T0CONTRL
        MOV AX, CONWRD            ;control word for one pulse
        OUT DX, AX

RESETMC:   IN AX, DX              ;read in TOCONTRL
          TEST AX, MCBIT          ;test if MCBIT is set
          JZ RESETMC             ;if not, wait
          AND AX, NOT MCBIT       ;clear MCBIT
          OUT DX, AX             ;update TOCONTRL

```

In this case, only a single pulse is required. Hence, in the control word loaded in the TOCON register, the CONT bit (D_0) is 0. Hence, the control word is 1100 0000 0000 0010 i.e., C002H. When the maximum count is reached, the MC bit in the control word is set. This can be used as an indicator of the end of the timing. To be able to start a new timing cycle, MC bit must be cleared and this is done here. Note that MC bit is being polled to verify the end of a timing cycle. The timer can also be operated in the interrupt mode whereby the setting of the MC bit can cause a Timer interrupt.

Timer Interrupts All timers can generate internal interrupt requests. Although all three timers share a single interrupt request to the CPU, each has its own vector location and internal priority. Timer 0 has the highest interrupt priority and Timer 2 has the lowest. Timer interrupts are enabled or disabled by the Interrupt (INT) bit in the Timer Control register. If enabled, an interrupt is generated every time a maximum count value is reached. In dual maximum count mode, an interrupt is generated each time the value in Maxcount Compare A or Maxcount Compare B is reached.

Conclusion

With this, we conclude our discussion on the 80186. The discussion is neither comprehensive nor complete. The aim of the chapter is to ensure that this processor is understood to be different in many ways from the other members of the x86 family. Its application domain is thus very

different—it will not be found in any PC, new or old, but can be seen to be used in many embedded designs even now. To gain a comprehensive knowledge about all the peripherals within the chip, you may refer to the Intel user's manual of 80186. To test the peripherals, an 80186 trainer will also be needed.

KEY POINTS OF THIS CHAPTER

- The 80186/80188 processors were developed in 1982 as a solution to the need for a processor with embedded peripherals.
- This processor and its many variants contain a number of peripherals on chip – like timers, DMA controller, Interrupt controller and so on.
- The 80186 has a few new instructions compared to the 8086. It also has new modes for some old instructions.
- The block diagram shows all the peripherals inside the processor.
- There are three timers, of which only two can be used for counting external events.
- Using the compare registers and control registers, counters can be used to generate pulses, continuous square waves and the like.


QUESTIONS

1. For what applications is the 80186 processor used?
2. What motivated Intel to develop a processor with embedded peripherals?
3. How has the PUSH and POP operations been enhanced compared to those of 8086?
4. What do the instructions INS and OUTS do?
5. How can the peripheral control block be relocated?
6. What are the features provided by the chip select logic?
7. Which are the interrupts supported by the internal interrupt control logic?
8. How can a delay be created using the internal timers of the processor?
9. In what way is Timer 2 different from the other two timers?
10. What is the significance of the MC bit of the TCON register?

EXERCISE

1. Find out the application domain of the 80186.
2. With respect to the chip select logic, find out what is meant by upper memory and lower memory.
3. Write programs to generate symmetric waveforms of two different frequencies.
4. Write programs to get a pulse of duration 500 μ s and 120 μ s.
5. Write a program to generate an asymmetric waveform of
 - a) 20% duty cycle from Timer 0.
 - b) 40% duty cycle from Timer 1.

15 THE 80286 AND 80386 PROCESSORS



In this chapter, you will learn

- The impact that 80286 made to the realm of PCs.
- The major changes that the 80386 architecture offered.
- The programming enhancements of 80386.
- The features of PVAM and the concept of virtual memory.
- The mechanism of address translation in the protected mode.
- The importance and implementation of protection mechanisms in the 80386.
- The issues and methods of task switching.
- The use of the interrupt descriptor table and exceptions of 80386.

15.1 | The 80286 Processor

You know that the first Intel processor used in a personal computer by IBM was the 8088. The PC based on 8088 and manufactured by IBM came to be known as IBM PC. It had only a floppy disk drive internally, but could have an external hard disk. All PCs with similar hardware and standards came to be called IBM PC clones. Intel later released the PC-XT (extended technology) which included a standard hard disk internally.

The next important processor of Intel was 80186. However, we have seen in the previous section that the 80186 has never been used as a processor in a PC, but instead, found applications in the field of embedded systems.

Soon after that, in 1982, Intel released the 80286, which came to be used in PCs that came to be called PC-AT where AT stands for 'advanced technology'. This turned out to be a big leap for Intel and for PCs as well. The PC-AT used 80286 as the processor and its technology could run all the software written for its predecessor. It could be operated at frequencies of 6MHz, 8MHz, 10MHz and 12.5MHz. The 80286 turned out to be a big improvement over the 8086/8088. It could boast of a performance, double that of the 8086. It had a number of features that were entirely new to the realm of processors.

The major enhancements can be listed as:

- i) The data bus remained the same at 16 bits (internally and externally), but there was no need for the multiplexed address/data bus, as separate pins were made available for them. The pin count increased so that the packaging could no longer be a Dual in line package like the 8086. Instead the packaging was a pin grade array (PGA).

- ii) The address bus was increased to 24 bits, which could address 2^{24} bytes or 16 Mbytes of physical memory.
- iii) The memory cycle time was reduced from 4 cycles to 2 cycles which made a direct improvement to the speed factor. Calculation of the more complex addressing modes (such as base + index) had less clock penalty because it was performed by a special circuit. The 8086, its predecessor, had to perform effective address calculation in the general ALU, taking many cycles. Also, complex mathematical operations (such as MUL/DIV) took fewer clock cycles compared to the 8086.
- iv) The biggest change and enhancement, however, was in the form of the introduction of what was called the 'Protected Virtual Addressing Mode (PVAM)'. Let us try to understand what exactly this is. For the 8086, there is no 'protection' provided, between users or between system software and application software. MS-DOS allowed access to anyone to any programs or any hardware through the mechanism of interrupts. Also, there was no facility to use the processor in a multi-tasking environment. 80286 reversed all this – it provided new mechanisms and instructions which provided protection and multi-tasking. Also, it could use the concept of 'virtual memory' whereby the memory as seen by the user becomes much larger than the actual physical memory present. This is the essence of PVAM.

When the processor is powered on, it is in the 'real mode' where it behaves as an accelerated 8086 – it then has only 20 active address lines and 1 MB of physical memory. From this mode, a special instruction switches it into the protected mode where the active address lines become 24 in number, the physical memory becomes 16 MB and all the protection and virtual memory mechanisms become activated. With all these enhancements and changes, this processor was popular for a short while, but one complaint against it was that it could not be switched back to real mode, without resetting it.

But it was this processor that made the real change in the way a computer could be designed and used. The concepts in the protected mode allowed the first GUI (graphic user interface) based operating system of Microsoft i.e., windows 3.1 to be designed comfortably for a PC. The design of the 80286 was followed up in Intel leading to a better and bigger (in terms of data and address size) processor named as the 80386 (1985). All the features of the 80286 were used here, and more features were added as well – we can say that the features of 80286 are a subset of the features of that of the 80386. Thus, it will be to our advantage to go into the details of the architecture of the 80386, which will give us all the wisdom necessary to understand the '486 and Pentium which followed very soon.

15.2 | The 80386

In 1985, Intel released the 80386 processor which carried in it all the new features it had introduced in its previous processor, and also made one great leap – that was the doubling of the data bus width to 32 bits and the increase in the address bus width to 32 bits. Thus, the data band width doubled, and the addressable physical memory became 2^{32} i.e., 4 GB. Thus, 80386 was Intel's first 32-bit processor. Intel also manufactured a model called 80386 SX which was a 32-bit processor internally, but had an external data bus width of 16 bits. However, our discussion will be on the 80386 DX which is a 32-bit processor internally and externally.

When powered on, the '386 is in the real mode with 20 address lines and capability to access only 1 MB of physical memory. From this state, it can be taken to the protected mode using a special instruction, the details of which we will soon come to.

15.2.1 | 80386 Enhancements

The 80386 has all the features of 80286 – additional instructions, protected virtual addressing mode – plus a lot more. It has more instructions, higher memory addressing capability, double the data bus width and address bus, and also includes the use of memory as ‘pages’ in addition to segmentation. Let us list out the new features.

- The external data bus is now 32 bits in width.
- The registers and hence the internal data bus width is 32 bits in width.
- The address bus is 32 bits in size and hence the memory addressing capability is 2^{32} or 4 GB.
- All addresses are 32 bits long and hence the address registers are also 32 bits long.
- A new addressing mode called the scaled addressing mode is available.
- A set of bit manipulation instructions is available.
- The '386 design overcame the biggest complaint against its predecessor – this processor can switch from protected mode to real mode without a hardware reset.
- The processor has all the features of PVAM that its predecessor had – as well as an additional ‘paging’ mode.

Let us go into the details of this processor – first the internal architecture, and then the hardware features.

15.3 | Internal Architecture

15.3.1 | Registers

Let us first view the working registers of the processor – the ones used for computation, address calculation and the like (see Fig 15.1). The processor has 16 registers, which consist of scratchpad (general-purpose) registers, segment registers, a flag register and the instruction pointer. Comparing it with 8086, shows that the 16-bit registers AX to DX have been re-designated as EAX to EDX, making them 32-bit registers. These registers can be used as 8 bits or 16 bits as well, but only the lower portions of the registers can be used i.e., the register EAX can be used as EAX (32-bit), AX (16-bit), AH (8-bit) and AL (8-bit). Similar is the case for the registers EBX, ECX and EDX.

The registers EBP, ESI, EDI and ESP are 32-bit address registers, as all addresses are 32 bits in length. The instruction pointer EIP is 32 bits long as each memory address is 32 bits in size. The same is the case of the stack pointer which is now designated as ESP.

15.3.1.1 | Segment Registers

The segment registers are 16-bit, but we will see later that they have what is called a ‘hidden part’ which will then make them totally 48 bits long (this becomes effective only in the protected mode). The visible part of the segment registers are still 16 bits and you can see that there are two new segment registers FS and GS. They are used to store the base addresses of the GS and FS segments, which are data segments itself.

15.3.1.2 | Control and Status Registers

The flag register i.e., now EFLAGS, contains the values of the status flags (i.e., C and Z), the control flags (i.e., DF, IF and TF) plus new flag bits catering to the additional features of the processor. The flag register has all the flag bits that the 8086 has – the additional bits are for use in the protected mode. Fig 15.2 shows the flag register bits.




Figure 15.1 | General registers and address registers used in the real mode

15.4 | Programming Enhancements

The processor has a number of new instructions and some enhancements with respect to the instructions available for its predecessor. It has all the instructions that the 8086, '186 and '286 has, plus a few more. Here we will see the new instructions which can be used in the **real mode**.

- i) The '386 has an interesting enhancement with regard to what we have until now defined as 'pointer registers' which carry addresses. **Here all the general-purpose registers too can hold addresses that can be used in the register indirect, register relative and based indexed mode of addressing.** For example, the following instructions are allowed.

```
MOV EAX, [EBX],  
MOV CX, WORD PTR [EAX]  
MOV AL, BYTE PTR [ECX]
```




Figure 15.2 | Flag register of the 80386

ii) A new and powerful mode of addressing called scaled index addressing mode has been introduced for the 80386. In this mode, each of the 32-bit registers (except the ESP) can be used as a pointer that is multiplied by a scale factor.

Now, let us do a few programming exercises using the new instructions of 80386.

Note This processor has 32-bit registers, and hence all the programming for this will be done in MASM32 which is the 32-bit Windows based assembler. The installation and use of this assembler is described in Appendix F available online at www.pearsoned.co.in/lylabdas.

Example 15.1

Find the effective address of the operand for the following instructions, given ECX = 233411E2H, EDI = 10008967H and EAX = 78960900H

- i) MOV [ECX*2], AX
- ii) MOV EBX, [2341H + EDI*4]
- iii) MOV BL, [EAX*8]

Solution

Remember that the 'effective address' calculation does not need the base address of the segment. It gives the offset of the operand from the base address.

- i) The effective address is obtained by multiplying the value in ECX by 2, which gives 466823C4H. The word from AX is moved to this address (offset).

- ii) The effective address is obtained by multiplying the value in EDI by 4 and adding 2341H i.e., $4002259CH + 2341H = 400248DDH$. The double word in this address is moved to EBX.
- iii) The effective address is obtained by multiplying the value in EAX by 8, which is $3C4B04800H$. The byte in this address is moved to BL.

Example 15.2

Add a set of 32-bit words stored in memory assuming that the result will not exceed 64 bits.

Note This program has been run in MASM 32 and results have been observed using the Olly debugger, the details of which are available in Appendix F.

```
.386
.MODEL FLAT, STDCALL

.CODE
DATS DD 00123456H, 00EF3456H, 07860EACFH,
      0A45674FH, 08976ABCH

STRT: MOV EAX, 0           ;EAX = 0
      MOV EBX, 0          ;EBX = 0
      MOV ECX, 00000005   ;ECX = 5, the count
AGN:  ADD EAX, [DATS + EBX*4] ;add to EAX the contents of memory
      INC EBX             ;increment the pointer
      LOOP AGN            ;repeat until ECX = 0
      END STRT
```

Here, after adding the five 32-bit data stored in locations starting from DATS, the sum is found to be 8C3F 2586H which is in EAX. Note that for the instruction `ADD EAX, [DATS + EBX*4]`, the scale factor used is 4. This is because the data is 4 bytes long, and therefore the address of each data is displaced from the next data by 4.

If the problem was one of adding words (2 bytes), we would have used a scale factor of 2. Thus the scaled addressing mode is useful in adding tables stored in memory, when the content of the tables are either 2 or 4 bytes long.

15.4.1 | New Instructions of 80386

Bit Test and Modify Instructions This group of instructions operates on a single bit which can be in memory or in a general register. The location of the bit is specified as an offset from the low-order end of the operand. The value of the offset either may be given by an immediate byte in the instruction or may be contained in a general register. These instructions first assign the value of the selected bit to CF, the carry flag. Then a new value is assigned to the selected bit, as determined by the operation. OF, SF, ZF, AF, PF are left in an undefined state. Table 15.1 defines these instructions.

Table 15.1 | Bit Test and Modify Instructions

Mnemonic	Instruction	Carry flag after execution	Effect on the selected bit after execution
BIT	Bit test	CF = bit	None
BTS	Bit test and set	CF = bit	BIT = 1
BTR	Bit test and reset	CF = bit	BIT = 0
BTC	Bit test and complement	CF = bit	BIT = not BIT

Example 15.3

```
.386
.MODEL FLAT, STDCALL
.CODE
DATS    DW    0560H

STRT:   LEA ESI, DATS           ;let ESI point to the data
        MOV AX, [ESI]          ;AX = 0560H
        BTC AX, 1              ;AX = 0562H, CF = 0
        BTR AX, 1              ;AX = 0560H, CF = 1
        BTS AX, 0EH            ;AX = 4560H, CF = 0
        END STRT
```

The effect of each of the operations is shown in the comment field.

15.4.2 | Bit Scan Instructions

These instructions scan a word or doubleword for a ‘1’-bit and store the index of the first set bit into a register. The bit string being scanned may be either in a register or in memory. The ZF flag is set if the entire word is zero (no set bits are found); ZF is cleared if a ‘1’-bit is found. If no set bit is found, the value of the destination register is undefined.

BSF (Bit Scan Forward) This instruction scans from low-order to high-order (starting from bit index zero)

BSR (Bit Scan Reverse) This instruction scans from high-order to low-order (starting from bit index 15 of a word or index 31 of a doubleword).

Example 15.4

```
.386
.MODEL FLAT, STDCALL
.CODE

DATS    DW    0560H

STRT:   LEA ESI, DATS           ;let ESI point to the data
```

```

MOV AX, [ESI]           ;AX = 0560H
BSR BX, AX              ;BX = 000AH meaning that D10 = 1
BSF CX, AX              ;CX = 0005 meaning that D5 = 1
END STRT

```

15.4.3 | MOVZX and MOVZX Instructions

Two other new and useful instructions are the MOVZX and MOVZX instructions. In MOVZX, the sign bit of a register (or memory) can be extended to any register. MOVZX zero extends the contents of a register or memory location. The MOVZX instruction is used for signed arithmetic and MOVZX for unsigned arithmetic.

Example 15.5

```

.386
.MODEL FLAT, STDCALL
.CODE
DAT1 DW -269
DAT2 DB 56
DAT3 DB - 123

STRT:  MOV AX, DAT1           ;AX = FEF3H
        MOVZX ECX, AX         ;ECX = FFFFFFFEF3H
        MOVZX EBX, DAT2       ;EBX = 00000038H
        MOVZX ECX, DAT1       ;ECX = 0000FEF3H
        MOVZX AX, DAT3        ;AX = FF85H
END STRT

```

In Example 15.5, the numbers are expressed in the decimal format. In the hex format, they are
 DAT1 = FEF3H
 DAT2 = 38H
 DAT3 = 85H

Note The first number is stored as a word, and the next two are stored as bytes. The result of each of the operations is shown in the comment field of the program.

15.5 | Hardware Features of 80386

The 80386 has 132 pins and is packaged as a pin grade array (PGA). Fig 15.3 shows the processor with the functional grouping of pins. The pin numbers are not mentioned – only the pin functionally is discussed. Only new pins which are not present in 8086 will be discussed. The pins that are not discussed may be assumed to have the same functionality at they have in the 8086.

Pin Grade Array A pin grid array (PGA) refers to the arrangement of pins on the integrated circuit packaging. In a PGA, the pins are arranged in a square array that may or may not cover the bottom of the package. PGAs are primarily used in applications that require more pins than what older packages such as the dual-in-line package (DIP) provide. Fig 15.4 shows a PGA.




Figure 15.3 | Functional pin diagram of the 80386




Figure 15.4 | A pin grade array (PGA)

Data Bus The data bus is 32-bit wide and is designated as D_0 to D_{31} . It is not multiplexed, as in the 8086, and carries only data. Since data can be used as 8, 16 and 32 bits, the data bus is grouped into four 8-bit combinations, and access of each of these combinations is selected by a separate enable signals for the memory/IO banks of the system.

Address Bus The address of the processor is 32 bits wide, but the dedicated address lines are only A_2 to A_{31} i.e., 30 lines. The address lines A_0 to A_1 are absent but they are encoded as four bus enable signals $\overline{BE0}$ to $\overline{BE3}$ (active low) which separately enable four memory banks. Each memory bank has a data bus width of 8 bits and data in the banks can be accessed by selective activation of the banks. See Fig 15.5.

There are four bank enable signals $\overline{BE0}$ to $\overline{BE3}$. They are selectively enabled depending on the banks which contain the data byte, word or double word. If a double word has to be read, all four bank enable signals are to be activated.

Example 15.6

For accessing a word, given the physical address as FFFF FFFCH, which banks have to be enabled and what should be the status of the bus enable signals?

Solution

To read a word, two bytes have to be obtained i.e., two banks have to be enabled. The specified address and the next address will be activated.

FFFF FFFCH and FFFF FFFDH. These addresses are in Bank 0 and Bank1. Hence, the corresponding bank enable signals have to be made low. Thus, $\overline{BE0}$ and $\overline{BE1}$ have to be low and $\overline{BE2}$ and $\overline{BE3}$ are to be high (addresses with the lowest two bits 00 will be of bank 0, with 01 in bank 1, 10 in bank 2 and 11 in bank 3).




Figure 15.5 | Memory banks for the 80386

Bus Cycle Control In this set, only the $D/\overline{C}$ signal is new.

$D/\overline{C}$ (Data/Control) This signal being high indicates that the data bus contains data for memory or I/O. The low status of this pin indicates that the processor is halted or is executing an interrupt acknowledge cycle.

Bus Control These signals are used for efficient bus control circuitry.

$\overline{ADS}$: Address Data Strobe: This signal becomes active whenever the processor issues a valid memory or I/O address.

$\overline{NA}$: Next address: This signal causes the processor to output the address of the next instruction or data in the current bus cycle itself. This is used for address pipelining.

$\overline{B16}$: Bus size 16 selects either a 32-bit data bus ($\overline{B16} = 1$) or a 16-bit data bus ($\overline{B16} = 0$).

Coprocessor Pins These are pins receiving signals from the coprocessor, which is external to the 80386. The 80387 is a compatible coprocessor operating similar to the 8087 (arithmetic coprocessor) discussed in Chapter 13.

$\overline{PEREQ}$: Co-processor request is a request from a coprocessor to the 80386 to relinquish control of the bus so that the 80387 coprocessor gets a direct connection.

$\overline{BUSY}$: This is a direct connection between the 80386 and the coprocessor, which indicates that the coprocessor is busy. This signal is used by the WAIT instruction to wait until the coprocessor finishes its task.

$\overline{ERROR}$: This is a signal from the coprocessor that an error has been detected.

Clock This provides the timing for the processor. The system frequency is half the frequency of the clock connected to this pin. For the system frequency to be 33 MHz, the frequency to be given to the pin should be 66 MHz. There is no internal clock generator – so an external clock chip is to be used.

Reset This is a high level sensitive input signal. When reset, control goes to the location FFFF FFF0H. The processor wakes up in the real mode and the upper 12 bits of the address remain high until a far jump is executed. This pin is included in the list of interrupt pins, because resetting has a pre-defined vector to which control is branched.

The other pins that are shown in Fig 15.4 have the same functions as they have in the 8086, and have been discussed in Chapter 6.

15.5.1 | Real Mode Operation of the 80386

The 80386 can operate in the real mode as well as in the protected mode. (There is also a third mode called the virtual 8086 mode.) On start up, the processor is in the real mode. In the real mode, it acts as a faster 8086 processor. The speed increase is due to the increased clock speed and the hardware enhancements. It takes the '386 only two cycles to perform a read/write (with zero wait states), while the 8086 requires four cycles. Besides being faster, the '386 in the real mode can also use the enhanced instructions the processor provides. It can use the 32-bit registers and data bus, and so the bandwidth is double that of the 8086.

The main feature of the real mode is that it has access only to 1 MB of physical memory, and address calculation is done as in the 8086, with the base address being directly added to the offset. The register set available in real-address mode includes all the registers defined for the 8086 plus the new registers introduced by the 80386 i.e., FS, GS, debug registers, control registers, and test registers.

When the processor is reset or after power on, it is in the real mode, with the reset vector being the address FFFF FFF0H. The upper 12 address bits remain high until a FAR call or jump is executed. The first far JMP or CALL instruction causes A31-A20 to drop low and thus the 80386 can execute instructions in the lower one megabyte of physical memory (00000 to FFFFFFH). This automatic lowering of address lines A31-A20 allows systems designers to use a ROM at the high end of the address space (of 1 MB) to initialize the system. The first instruction has an address in the upper part of the 1 MB physical memory and that is usually ROM – thus the processor continues to be in the real mode, working within the physical memory of 1 MB.

We will now discuss the protected mode features of the processor, starting with memory management. Before we do that, it is necessary to understand a very important concept regarding memory.

15.6 | Virtual Memory

Virtual memory is a concept by which the processor (in effect, the user) is made to believe that it has much more memory than is physically available. In effect, there is this idea of an ‘enlarged address space’ which is called a ‘virtual address space’. This space maps into a physical address space, ultimately. When a user writes a program, he uses ‘virtual addresses’ or ‘logical addresses’. The idea is that an application does not get to see where the data is physically located. It is up to the address translation mechanism to map these virtual addresses to a physical address. This is an operating system concept which keeps the user free of the worry of exhausting the available memory, when he writes an application.

Recollect that whenever we executed DOS based programs, we used physical memory, which is RAM/ROM. We noted that for the 8086, with 20 address lines, the memory space available is 2^{20} bytes i.e., 1 MB. Similarly, for the 80286, with 24 address lines, the physical memory it can address is $2^{24} = 16$ MB. For the 80386, naturally, the physical address space is 2^{32} i.e., 4 GB.

Consider using the 80386 for an application – does it mean that we might have to limit our data and code to fit into 4 GB? We know that in the computers that we use now, we never have experienced a ‘memory crunch’ or the feeling that our application cannot be run because the memory capacity is insufficient, though there might have been times when we have felt that our application runs too slowly. All this leads us to the concept of what is called ‘virtual memory’.

In a computer system there is a memory hierarchy, and a two level hierarchy is shown in Figure. 15.6. Let us discuss the concept of virtual memory using just a two level hierarchy,

Semiconductor memory, which is reasonably fast (RAM/ROM), is what we usually designate as ‘primary memory’, ‘main memory’ or ‘physical memory’. Keep in mind that we will frequently use these terms interchangeably and all the three refer to the memory that is addressable by the address lines of the processor. We then designate the hard disk (and other memory types) as ‘secondary memory’. The processor directly communicates with the main memory only, but data movement between the secondary and main memory is possible. Whenever an application is to be run, the processor expects to find the associated data and code in the main memory.




Figure 15.6 | Memory hierarchy in a computer system

In case it is not found there, it should be brought thereto from the hard disk. This takes a small amount of time. This is the reason why we find our computer ‘slow’ if the RAM size is low. However, even if the RAM is of a large capacity, it is not necessary that our application files are currently present there. Managing the system such that important files are found in the main memory most of the time, is part of the ‘cleverness’ of the operating system.

Does this mean that the ’386 processor does all this? Definitely not, but the processor provides the extra hardware, extra instructions, and an intelligent ‘memory management unit’ that allows the OS to handle many of such tasks efficiently.

15.6.1 | Protected Virtual Addressing Mode (PVAM)

Now, we will go through the features of the 80386 processor which has made all this possible. Please be warned that understanding these features is not very easy and cannot be done in one reading. Repetitive reading and contemplating on the intricacies of the mechanisms involved, is the only way to get a clear understanding of this processor. Since the same features are also present in the 80486 and Pentium as well, it is important to get a good insight into all this, at this point itself. It is also important to keep in mind that the ’386 is an upgraded version of the basic 8086 – as such, it will have all the features of 8086 plus additional features. The ‘additional’ features in the protected mode are what we will discuss now.

The following are the topics we will cover:

- i) Memory Management
- ii) Protection
- iii) Multitasking
- iv) Exceptions and Interrupts

Even though we discuss each of them separately, they are not really disjoint – for example, protection is a part of memory management; memory management and protection have to be taken into consideration when multitasking is done, and so on.

15.7 | Memory Management Unit

The 80386 transforms logical addresses (i.e., addresses as viewed by programmers) into physical addresses (i.e., actual addresses in physical memory) in two steps:

- i) Segment translation, in which a logical address (consisting of a segment selector and segment offset) is converted to a linear address. Here, memory is considered to be divided into segments which are large units of storage.
- ii) Page translation, in which a linear address is converted to a physical address. This step is optional, and usable on the discretion of systems-software designers. In this step, memory is divided into smaller storage areas, each of which is designated as a page.



Figure 15.7 | Address translation using paging and segmentation

These translations are performed in a way that is not visible to applications programmers. The 80386 has both segmentation and paging. If both these schemes are active, the logical address given by the program is converted to a linear address by the address translation hardware of the segmentation unit. This linear address is interpreted as a page address and given to the paging unit which converts it to a physical address. **But paging can be switched off by certain register settings. In that case, only segmentation is effective and then the logical address is directly converted to a physical address.** We will deal with ‘segmentation’ and ‘paging’ separately.

15.7.1 | Segmentation

A typical instruction could be

```
MOV EAX, COST
```

where EAX is a 32-bit register and COST is a 32-bit offset in a segment.

This implies that the segment has a 32-bit base address, because all addresses are 32-bit long. However, to access the location called COST, we have to get the base address of the segment it is contained in. COST is the offset of the data from the base address. The physical address calculation involves just adding the offset to the base address. In the real mode (as in 8086), this address calculation is very direct (Fig 15.8), but it not so in the protected mode.

The reason for the additional levels of calculation in the protected mode is the necessity for incorporating virtual memory concepts into memory management and ensuring protection features. Our problem is now of getting a physical address when the program supplies a virtual address, which is the logical address in a virtual memory system. In the instruction cited as an example, the label COST is the logical address of the operand to be accessed. COST will be a 32-bit address – a number (say, 0987 7654H).

Any location that is to be accessed will be in a particular segment. A segment can be of different types and it also has many attributes. All this is put together in what is called the ‘descriptor’ of the segment. Thus, any segment has a descriptor, which gives all information about the segment i.e., it describes the segment. See in Fig 15.9, the structure of a typical descriptor of a data/code segment used in an application program. (For an operating system segment, some bits of the descriptor will be different from this.)

A descriptor is an 8-byte (numbered as 0 to 7) entity which describes the segment pointed by it. Let us see the contents of the descriptor in Fig 15.9

- i) It contains the 32-bit base address of the particular segment. This address is obtained as the concatenation of bytes 2, 3, 4, and 7 of the descriptor.




Figure 15.8 | Address calculation in the real address mode




Figure 15.9 | Descriptor of a data/code segment of an application program

- ii) Bytes 0, 1 and 4 bits of byte 6 constitute the limit of the segment. This is a 20-bit number and specifies the maximum extent of the segment. There are two cases for the limit and this is decided by the G (granularity) bit of byte 6. If $G = 0$, it imposes a maximum limit of 64 K for the segment size. Then the limit bits L_{16} to L_{19} are to be taken as zeroes, and the bits L_0 to L_{15} give the size of this particular segment. Remember that with 16 bits, the maximum segment size is 64 K.

But if the G bit = 1, the number L_0 to L_{19} is to be a multiple of 4 K (2^{12}). Thus, the maximum possible size of a segment is $2^{20} \times 2^{12} = 2^{32} = 4$ GB. This is how the maximum size of a segment in the protected mode becomes 4 GB. In the descriptor, the number corresponding to the limit, gives the size of this particular segment.

- iii) Byte 5 of the descriptor is called the access byte. It contains the bits which specify the type and access rights to the segment pointed by this descriptor.

The various bits in the access byte mean the following. Refer Fig 15.10.

Bit 0A: Accessed This bit is set when the corresponding segment is accessed. The operating system tests this bit periodically to find out the frequency of usage of this segment – this is necessary for algorithms that need to have a measure of the ‘importance’ of this segment, meaning that, if this is a segment which is accessed very frequently, it is best to retain it in the main memory.

Bit 4 S: Segment indicates that this is a data or code segment descriptor for an application program ($S = 1$), or is a system segment descriptor ($S = 0$).

Bit 3: Executable E This bit indicates if the segment is executable ($E = 1$) i.e., code segment, or non-executable ($E = 0$) i.e., data segment.

Bits 1, 2: Type The type value depends on the E bit (bit 3 of the access byte) and the corresponding definitions are given in Table 15.2.

Bits 5, 6 DPL This is a two-bit field which stands for ‘descriptor privilege level’, and is part of the ‘protection mechanism’ which we will discuss in greater detail in Section 15.10.

Bit 7 P: Present Bit The setting of this bit shows that the segment corresponding to this descriptor is now present in the main memory. If $P = 0$, it means that the corresponding



Figure 15.10 | Contents of the access byte in a descriptor

Table 15.2 | The Designation of the Bits of 'Type' Field of the Access Byte

Bit no.	If bit 3, E = 1 meaning executable (code segment)	If bit 3, E = 0 meaning non-executable (data segment)
2	C, Conforming, C = 1 Non conforming, C=0	ED Expand down. ED = 1, stack segment Expand up ED = 0, data segment
1	R, Readable R = 1, may be read R = 0, may not be read	W, writable W = 1, writable W = 0, may not be written to

segment is not in main memory. Hence, a type 11 interrupt is generated which initiates an ISR to bring the segment into the main memory.

Byte 6 Four bits are used as part of the limit field. Byte 6 has four bits other than the limit bits

G: Granularity This bit (discussed earlier) decides if the limit field has to be multiplied by 1 or 4 K. If G = 0, the multiplier is 1. If G = 1, the multiplier is 4 K.

D Selects the default instruction mode for code segment descriptors. If D = 0, the registers and memory are 16 bits wide as in the 80286; if D = 1, they are 32 bits wide as in the 80386.

AVL This bit can be used by the OS in any way it needs. It is often used simply to specify that the described segment is available. Bit 5 of byte 6 is to be 0.

15.7.2 | System Descriptors

Segments that are used by the operating system are separate and somewhat different from application program segments. They have a descriptor which is slightly different. Fig 15.11 shows the general format of a descriptor used for operating system uses. If Fig 15.11 is compared with Fig 15.9, what will be noted is that only the bits in the access byte and the four bits of byte 6 are different. The difference in the access byte is that, 4 bytes are used to specify the type of the segment being described. There are 16 possible types, of which a few have reference to the 80286, some refer to future uses, and only a few indicate the types of system segments of the 80386. It is not worth going to the details of these types, since it will only lead to cluttering of this discussion with unnecessary detail.

15.7.3 | Descriptor Tables

So, we see that each and every segment has a descriptor. However, who creates descriptors? Definitely, not the application programmer. When the programmer creates a segment, the system software creates the corresponding descriptor with the help of compilers and loaders and linkers – in effect; it is an operating system job.




Figure 15.11 | Descriptor for a system segment

In any system, there will be a number of segments of various types created for various applications – thus, there should be as many descriptors too. These descriptors are stored in tables called ‘descriptor tables’. There are two types of descriptor tables – the local descriptor table and the global descriptor table.

15.7.3.1 | Local Descriptor Table (LDT)

Suppose the system has 10 application tasks whose segments and corresponding descriptors have been created. For each of these tasks, there will be a descriptor table designated as the LDT of the particular task. Thus, we may name these tables as the LDT of task 1, task 2 and so on. There will be as many LDTs as there are application tasks.

15.7.3.2 | Global Descriptor Table (GDT)

There are segments and corresponding descriptors, which do not belong to any specific application but cater to the system as a whole – they are global in nature. These descriptors are stored in the GDT. **There will be only one GDT in a system.**

For executing a particular application task, say task A, segments from the LDT of task A are required and some operating system services are also required. Figure 15.12 shows how an application task needs the GDT as well as its own LDT.

The next question is – where are these tables stored? The answer is – in the main memory, naturally. Remember also that a descriptor table is a memory array of 8-byte entries and that the size of a table is variable. Thus these tables are also ‘segments’ and each of them also should have a descriptor. This may sound a bit confusing at first, but on going ahead, this confusion will soon be overcome.

This takes us to the discussion of the processor registers which are used in the protected mode. At a particular time, the descriptor of the GDT will be available in a processor register called the ‘Global Descriptor Table Register (GDTR)’. This is to be ensured when the system is switched to the protected mode. Since descriptors are 8-byte entities, these registers have to be 8 bytes in length, obviously. In fact, the loading of the GDTR with the descriptor of the GDT occurs when the system is put into the protected mode.




Figure 15.12 | LDT of a task and the GDT of the system

If an application task is running, the descriptor of the LDT of that task is available in the Local Descriptor Table Register (LDTR). We will soon get to the details of all the processor registers involved in the address translation mechanism.

15.7.4 | Selector

Remember that any logical address for the 8086 could be specified in the form
Segment base address: Offset

For the 80386 protected mode, the format of any logical address is modified to be
Segment Selector: Offset

The offset is 32 bits long as mentioned earlier, but the selector is only 16 bits long and has the structure as shown in Fig 15.13. It contains a 13-bit number called the index. This number is an indicator of where in the descriptor table, the particular segment’s descriptor is located. Thus, for each descriptor, the selector is a pointer to its position in the descriptor table. Since the index is 13 bits long, the number of descriptors it can point to is 2^{13} i.e., 8192, this means that any descriptor table can have 8192 descriptors stored in it.

In the selector, there is a bit called TI which stands for ‘Table Indicator’. This indicates whether the selector corresponds to a descriptor in the LDT (TI = 1) or the GDT (TI = 0).

RPL This is a two-bit field ‘Requested Privilege Level’ which is part of the protection mechanism, which will be discussed in detail in Section 15.10.4. The selector value is fixed up by linker or loader software.

Registers for Address Translation To understand the concept of address translation in the protected mode, we should have a clear idea of the registers which are involved in this. Corresponding to the type of segments in use, there are the segment registers as shown in Fig. 15.14. Each segment register has a two-byte (16-bit) visible portion and an 8-byte hidden part in which the corresponding descriptor is loaded, when the corresponding segment is being accessed.

Now, let us consider a simple case of address translation. Note the following points. The GDT and LDT are, as mentioned earlier, memory arrays or segments. So they too have descriptors.




Figure 15.13 | Format of a selector




Figure 15.14 | Address translation registers

The GDT contains the descriptors of the segments used in system programs – there is only one GDT for the whole system. As such, there is no selector for the GDT. On startup, the descriptor of the GDT will be loaded into the GDTR. Once this is set up, the system can start running application tasks. Fig 15.15 shows the descriptor of a GDT loaded into the GDTR. Note that the GDT can contain a maximum of 8192 descriptors and the limit field indicates the size of the GDT (here, a maximum sized GDT is shown).

There are as many LDTs as there are tasks in the system. The descriptors of all the LDTs are in the GDT. When an application task is taken up, the selector of its LDT will be loaded into




Figure 15.15 | GDT and GDTR

the visible part of the LDTR and then the corresponding descriptor is taken from the GDT and loaded in the hidden part of the LDTR. It is then that this task becomes active. (When this task is abandoned and a new one is taken up, the content of the LDTR will be changed to favor the new task).

15.8 | Converting a Logical Address to a Physical Address

Now consider an instruction for an application program, which is

```
MOV EAX, COST
```

This instruction needs an operand from a data segment to be moved to the 32-bit register EAX. The segment has a descriptor as well as a selector created by the system software. The needed data has an offset labeled COST which corresponds to, say, 0012 5674H. Looking at Figure 15.16, try to follow the steps outlined below. The steps in accessing data are:

- i) Since it is a data segment which is to be accessed, the selector of the data segment will be loaded into the visible part of the DS register i.e., the segment register part. The selector contains a 13-bit number – the index. This is multiplied by 8 and added to the base address of the descriptor table (GDT or LDT, depending on the value of the TI bit in the selector).
- ii) Note that the base address of the descriptor table is available in the GDTR register (GDTR), or LDT register (LDTR) of the processor. Thus the descriptor of the data segment is obtained from the descriptor table, and this is loaded into the hidden part of the DS register. Since an application program is considered here, the descriptor of the data segment will be taken from the LDT of the application task.
- iii) The descriptor of the data segment has the base address of the data segment. This is used to locate the data segment. This base address plus the offset mentioned in the instruction i.e., 0012 5674H (with label COST) are added to get the physical address of the operand, which is then moved to the register EAX.

The above explanation should be sufficient to understand the translation of a logical address specified in the instruction to a physical address in memory, in the case of segmentation.




Figure 15.16 | Address translation of a logical address to a physical address

15.9 | Calculating the Size of the Logical Address Space

In a segmented model of memory, the address space as viewed by an applications program (called the logical address space) is a much larger space of up to 2^{46} (64 terabytes). How is this figure obtained? Let us make a calculation.

A selector has 13 bits as index for selecting one of the descriptors from a descriptor table. Thus, it can access up to 2^{13} descriptors with just one descriptor table. Since two descriptor tables (the GDT and one LDT) are active at a time, 2×2^{13} segments can be accessed. Each descriptor corresponds to a segment with a maximum size of 4 GB or 2^{32} bytes. The total available memory space is now calculated to be $2 \times 2^{13} \times 2^{32}$ bytes = 2^{46} or 64 terabytes

($2^6 \times 2^{40} = 64 \text{ TB}$). Thus the logical or virtual address space is 64 TB, which is much larger than the actual physical memory of 4 GB.

This calculation makes clear the idea of how the virtual memory concept creates an enlarged memory size. The address translation mechanism discussed above has the duty of mapping this large virtual memory space to a much smaller physical memory. Applications programmers do not need to know the details of this mapping.

As a matter of fact, the user usually thinks that the memory available to him is infinite. This illusion is created by having a secondary memory and shuttling data between the main and secondary memories according to the needs of the user. This underlying mechanism is not the concern of the user – the OS has the duty of performing this task efficiently.

15.9.1 | Paging

The mechanism of paging is in many ways, similar to segmentation – however, a smaller chunk of memory is handled in this case. Remember that a segment can have a variable size with a maximum physical size of 4 GB, but the size of a page is fixed to be 4 KB. This small size is what makes paging so attractive.

Just as in segmentation, whenever we want to run an application, we expect the required page to be in the main memory – if not, it has to be brought to the main memory. The charm of paging is that it allows just one page (i.e., just 4 K) to be brought in (or swapped out, if necessary). In contrast, segmentation requires that the whole segment be handled.

Paging is not a new concept, but 80286 did not allow paging – the '386 allows both paging and segmentation – which implies that any segment can be considered to be composed of a number of pages, and handled as one segment or as a number of separate pages. This is quite logical because we know that sometimes we might need only a part of the available data/code segment for an application.

15.9.2 | Page Translation

The 80386 has segmentation as a basic concept in its memory management unit. The **logical address** from a program is converted into a '**linear address**' by the segmentation unit, and given to the paging unit. This **linear address** is to be translated to a **physical address**, and this entails two levels of translation. A linear address can refer directly to a physical address so the segment and offset can be located directly, as is done in Section 15.8. This is the case when only segmentation is used. If paging is also activated, a linear address indirectly specifies a physical address (refer Fig 15.7) by specifying a page table, a page within that table, and an offset within that page. A 32-bit linear address, as shown in Fig 15 17, is composed of the following fields:

- i) A 10-bit field called DIR which points to a page descriptor in a page directory.
- ii) A 10-bit PAGE field which is used as an index into the page table determined by the page directory.
- iii) A 12-bit offset which points to the required byte within the located page.



Figure 15.17 | Format of the 32-bit linear address




Figure 15. 18 | Two level address translation for paging

A page or page frame, as it may be referred to, is a 4K-byte unit of contiguous addresses of physical memory. See Fig 15.18 for understanding the two levels of address translation in paging. Note that the base address of the current page directory is to be in the CR3 register. This register is in the processor and is also called the Page Directory Base Register (PDBR). The page directory contains page descriptors, which are similar to segment descriptors, except that these are 4-byte entries i.e., 32 bits long. Thus the page directory contains the descriptors for page tables. A page table is also a page (4K in size) – thus it has a 4-byte descriptor which is in the page directory.

In the first level of address translation, the bits of the DIR field of the linear address is used as an index to locate a page descriptor in the current page directory. This will contain, among other things, the base address of a page table. In the next level of translation, the PAGE field of the linear address is used as an index to locate an entry in the page table. This entry contains the address of the required page in the physical memory. Adding the OFFSET field (in the linear address) to this address, gives the physical address of the required data. Note that this is a two level address translation mechanism, and the two levels are similar. In the first level, a page directory is accessed, and in the second, a page table is accessed. Both the page directory as well as the page table are 4K-byte long pages.

Now to get a feel about the amount of memory addressable with paging, let us do a bit of calculation. Go back to the format of the linear address. Since the DIR field is 10 bits long, 2^{10} or 1K page tables are possible. Each entry in the page table is indexed by the 10-bit PAGE field of the virtual address – hence it has 2^{10} entries of page descriptors – which means 2^{10} pages can be addressed. Each page can contain 2^{12} (4K) bytes. Thus the total memory is $2^{10} \times 2^{10} \times 2^{12} = 2^{32}$ bytes or 4 GB, which is the complete physical memory of the 80386 processor (and also the maximum size of a segment in the protected mode).

15.9.3 | Format of a Page Table Entry

Fig 15.19 shows the format of a page table entry. Let us discuss its contents.

Page Frame Address The page frame address specifies the physical starting address of a page. Because pages are located on 4K boundaries, the low-order 12 bits are always zero. In a page




Figure 15.19 | Page table entry

directory, the page frame address is the base address of a page table. In a second-level page table, the page frame address is the address of the page that contains the desired memory operand.

P bit This bit has the same function as the P bit in segmentation i.e., P = 1, indicates a valid page access. P = 0 indicates that this page is not in physical memory, and hence it generates an exception which initiates an ISR to bring the page to physical memory.

R/W Read/Write-used for protection.

U/S User/Supervisor-used in the protection scheme.

D Dirty is undefined for page directory, but used in the page table. It means the page has been modified.

A Accessed is set to 1, whenever the page directory is accessed. Certain bits are left to be 0 which indicates Intel reserved.

15.9.4 | Translation Lookaside Buffer

It is cumbersome and time consuming to calculate the physical address from the linear address for every memory access. A translation look-aside buffer (TLB) simplifies the process. A TLB is a table in the processor which stores the physical addresses of the 32 recently accessed page table entries. The TLB is a ‘hardware cache’ for the physical addresses of the most recently used virtual address. It is a four way set associative cache inside the ’386 chip.

The paging unit receives a 32-bit linear address from the segmentation unit. The upper 20 bits of the linear address is compared with all 32 entries in the TLB to check if it matches with any of the entries. If it matches, the 32-bit physical address is calculated from the matching TLB entry and placed on the address bus. See Fig. 15.20. If it does not match, the processor must fetch into the TLB, the page table entry from memory.

15.9.5 | Combining Paging and Segmentation

Since 80386 has both segmentation and paging, it is easy to visualize a system in which a segment can be composed of a number of pages. To access a data in a page, what is needed is to have a technique by which the logical address corresponding to segmentation locates a segment, and then this logical address is converted to a linear address. This linear address is converted to a physical address as we have just seen. This is shown in Fig. 15.21.




Figure 15.20 | Translation lookaside buffer inside the 80386 chip




Figure 15.21 | Segmentation and paging combined

15.9.6 | Flat Model

In a '386 based system, it is not mandatory to use paging and segmentation together. There is the option to turn off the paging mechanism that leaves a segmented architecture. However, a mechanism to turn off segmentation is not inherently present. However, this can be achieved by programming techniques.

When the 80386 is used to execute software designed for architectures that do not have segments, it may be expedient to effectively 'turn off' the segmentation features of the 80386. The 80386 does not have a mode that disables segmentation, but the same effect can be achieved by loading the segment registers with selectors for descriptors that have a base addresses of 0, and by setting the limits to allow access the entire 32-bit address space. Once this is done, there is no longer a need to change the segment registers. The 32-bit offsets used by the 80386 instructions, in protected mode, are sufficient to access the entire linear address space. This is called the 'flat' model. In a 'flat' model of memory organization, the applications programmer sees a single array of up to 2^{32} (4 GB) bytes.

15.10 | Protection

In the previous section, you have seen the word 'protection' being used without being told what exactly it is and how it is achieved. So the first thing to be done in this section, is to get a clear idea of the word 'protection', in the way it is used in the protected virtual addressing mode of 80386. There are many aspects to protection in this context.

Protection Between Tasks A system can have a number of tasks, some of which are active, and some which are not. One requirement is to protect user tasks from each other – this means that the code and data of one user is to be isolated from all other users – since code and data are in segments, this implies that segments of one task should not have access to the segments of other user tasks. This protection is implemented by having separate LDTs for each task. The descriptors of segments of each task are kept in separate local descriptor tables. At a time, only one user task is active and only one LDT can be used.

Protection Between User and System Tasks The system programs i.e., tasks associated with the operating system, manage the computer system as a whole. Thus these tasks are very important and should not be allowed to be tromped on by users i.e., application tasks. This is achieved by the concept of privilege levels. System tasks are allotted a high privilege level, while user tasks are at a lower privilege level – and direct access of segments of a higher privilege level task is not permitted. Thus, user tasks are prevented from accessing system level segments.

Capturing Address Violation Errors This aspect ensures that addresses beyond the limit of a segment are not accessed. Any such attempt is immediately aborted. It should be remembered that all these features are provided in the hardware of the processor, and when an operating system is designed using these hardware features, the design of the operating system becomes relatively easy. Now, let us get into the details.

15.10.1 | Protection Checks

The important aspects of protection in the 80386 are as follows:

- Type checking.
- Limit checking.
- Restriction in the accessible domain.

- Restriction in the entry points of procedures.
- Restriction in the usable instructions.

Address translation and memory access have already been discussed. Protection hardware is an integral part of that unit, which applies to both segmentation and paging. It is thus necessary to keep in mind that protection is discussed separately, only for ease of understanding. Each reference to memory is checked by the hardware to verify that it satisfies the protection criteria. All these checks are made before a memory cycle is started. Any violation prevents that cycle from starting, and results in an exception. Since the checks are performed concurrently with address formation, there is no performance penalty. Invalid attempts to access memory, results in an exception (an exception is an interrupt caused by an error).

The role of descriptors

The protection parameters are placed in the descriptor by system software at the time a descriptor is created. When a segment selector is loaded into a segment register, the descriptor of the segment is automatically loaded into the hidden portion of the segment register – thus all the protection parameters are available in the segment register inside the processor and so subsequent protection checks on the same segment do not consume additional clock cycles. The first check that needs to be done is on the P bit in the descriptor. If $P = 0$, it implies that the referred segment is not present in main memory, and an exception is generated, which initiates an ISR to bring the corresponding segment to the main memory.

15.10.2 | Type Checking

There is a type field for a descriptor within the access byte (bits 1, 2 and 3). Refer Fig 15.10. This field has two functions.

- i) It distinguishes between different descriptor formats.
- ii) It specifies the intended use of a segment.

Programming errors can occur if a wrong segment selector is loaded into a segment register. Keep in mind that a code segment alone is an executable segment. It can be allowed to be read, but cannot be written into. Data segments must be readable, and stack segments are to be writable. Thus, the following conditions are to be verified.

- i) The CS register can be loaded only with a selector of an executable segment.
- ii) Selectors of executable segments that are not readable cannot be loaded into data-segment registers.
- iii) Only selectors of writable data segments can be loaded into SS.

When an instruction with memory reference is to be executed, type checking is performed to confirm that

- i) no instruction may write into an executable segment,
- ii) no instruction may write into a data segment if the writable bit is not set,
- iii) no instruction may read an executable segment unless the readable bit is set.

15.10.3 | Limit Checking

There is a limit field in the descriptor and the limit of a segment depends on the value of the G-bit. Before allowing a memory access, this field is used to confirm that memory access

outside the segment does not occur. The value of the limit field has to be interpreted differently for expand-up segments (data), and expand-down segments (stack). Catching and preventing address violation errors goes a long way in avoiding system crash.

15.10.4 | Privilege Levels

The concept of privilege levels is all about trust and protection.

Trust When a task is to be performed, it is code that gets executed – to execute code, data from various segments are required – also, code from other segments may be required when a branching to a different segment is to be done. The task that does this has to be a ‘trusted’ task – the code segment performing the task has to be a trusted segment. Then only it can be allowed to access other segments. The privilege level accorded to this segment depends on how much of trust is placed on it. If it is accorded the highest privilege level, it means that is allowed access to any other segment that it needs. Such a high privilege level is accorded only to segments of operating system tasks i.e., tasks with supervisor status, which are the most trusted tasks. Application level tasks usually are at low privilege levels – they are not trusted enough to allow them to use all available segments. All this means that all segments are given a particular privilege level (usually a number), and a task at a low privilege level can only access segments at the same level or at a lower privilege level.

Protection Now think of a data segment or even a code segment which contains a procedure to be accessed from other code segments. Let us call it a target segment. What is the extent of protection it has? Can it be used by just any task? In that case, it is a segment at the lowest privilege level and hence the ‘least protected’. If it is highly protected, it will be accorded the highest privilege level and then, only tasks with the supervisor status can use it.

15.10.4.1 | Privilege Levels of 80386

The ‘86 processor defines 4 levels of privilege numbered as 0 to 3 – the smaller the number, the higher the privilege level. Fig 15.22 shows the 4 level privilege hierarchy. Note that the kernel of the OS has been accorded the highest level of 0, while application tasks are at level 3. OS level services are at privilege level 1, and utilities and custom services of lesser importance are at level 2. In the figure, three application tasks A, B and C have been shown. It is not mandatory that an 80386 based system use all the 4 privilege levels. The system software can reduce it to 3 or 2 as needed.

Three different types of privilege levels enter into the privilege level checks.

CPL (Current Privilege Level) CPL is stored in the selector of the currently executing code segment register (bits 0 and 1); **CPL represents the privilege level of the currently executing task** i.e., the program or procedure. Normally, the CPL is equal to the privilege level of the code segment from which instructions are being fetched. Thus, this is also the privilege level in the access byte of the descriptor of this code segment (except for what is called a ‘conforming’ code segment). The CPL is also designated as TPL or the Task Privilege Level.

DPL (Descriptor Privilege Level) DPL is the privilege level of an object, which is being attempted to be accessed by the current task. It is the privilege level of the target segment, and is contained in the access byte of the descriptor of the corresponding segment. The above discussion might seem a bit complex, but what is actually done is only a simple comparison of numbers. In simple terms, what is done is that the CPL of the task in execution is compared




Figure 15.22 | Privilege level hierarchy of 80386

with the DPL of the target segment. After comparison is done, access is allowed only if the CPL is higher or the same as the DPL of the target segment. Keep in mind that a higher privilege level corresponds to a lower number.

RPL (Requestor Privilege Level) It is seen that the lowest two bits of any selector are reserved for the RPL field – these bits have an effect on privilege validation. This RPL field of a target segment’s selector can be used to ‘weaken’ the CPL, if desired. If RPL of the selector is 2 while the CPL of the task is 1, the effective privilege level is the numerical higher of them i.e., 2. Thus the task becomes less privileged. The task is effectively excluded from accessing a data segment that the task might otherwise have been privileged to use. (This may seem unnecessary, but it is generally used to prevent inappropriate use of pointers that could corrupt the operation of more privileged code or data from a less privileged level.)

But the opposite effect of making RPL = 0 does not have any effect. If the selector’s RPL is numerically less than the CPL, it has no effect.

15.10.5 | Accessing Data Segments

Assume that a task needs to get data from a data segment. The privilege level checks are performed at the time a selector for the descriptor of the target segment is loaded into the data-segment register. (Keep in mind that the higher privileges are represented by smaller numbers and vice versa.) As Figure 15.23 shows, three different privilege levels enter into this type of privilege checking mechanism.

1. The CPL.
2. The RPL of the selector used to specify the target segment.
3. The DPL of the descriptor of the target segment.




Figure 15.23 | Privilege level checks for a task when accessing a data segment

Instructions may load a data-segment register (and subsequently use the target segment) only if the DPL of the target segment is **numerically greater** than or equal to the maximum of the CPL and the selector's RPL. In **other words**, a code segment can only access data that is at the same or less privileged level.

The data segment has a DPL in the access byte of its descriptor, and this number designates the privilege level of the data segment. Then the data segment must be at the same or lower privilege level in comparison with the CPL of the task. The accessible domain of a task varies according to its CPL. If the CPL of a task is zero, data segments at all privilege levels are accessible; when CPL is one, only data segments at privilege levels one through three are accessible; when CPL is three, only data segments at privilege level three are accessible. This property of the protection mechanism is used, for example, to prevent applications procedures from reading or changing tables of the operating system.

In short, the following numeric equation should be satisfied: $CPL \leq DPL$. However, CPL can be weakened by the RPL. So we define the Effective Privilege Level (EPL) as the numeric maximum of RPL and CPL. Thus the final condition for access is that $EPL \leq DPL$.

Figure 15.23 shows the concept of privilege level checking when a task attempts to access a data segment. If the required condition is satisfied, access is allowed, otherwise it is not.

15.10.6 | Privilege Checks for Control Transfers

We have been talking of code segments having to access segments that contain data. However, many times code segments have to access other code segments. This happens when a far call or jump or ret instruction is encountered. (It also occurs in the case of interrupts, but that will be discussed separately). There are a few issues in this type of control transfer.

- We can go to a code segment of lower privilege level, but on returning, it causes privilege level violations, so in effect this type of access is not possible.
- Going to code segments of higher privilege level is not permitted, but there is an indirect way of doing this and that is using 'gates'.

Before proceeding further, a clarification on 'conforming' and 'non-conforming' code segments needs to be given.

Conforming Code Segment A conforming code segment is one that does not have a privilege level of its own. It can be called by programs at any privilege level and then, it acquires the CPL of the calling program. Bit 2 of the access byte of a code segment descriptor specifies if it is a conforming (C = 1) or non-conforming (C = 0) segment. This mechanism is well suited to handle programs that share code but run at different privilege levels e.g., shared libraries.

It permits sharing of procedures, so they can be called from various privilege levels. They are usually used for math library and exception handlers.

For a conforming code segment, the selector's lowest two bits need not be the same as the CPL bits in the access byte of the code segment's descriptor.

Non-Conforming Code Segment But most code segments are not conforming. The basic rules of privilege mean that, for nonconforming segments, control can be transferred only to executable segments at the same level of privilege. There is a need, however, to transfer control to higher privilege levels – for example, when a low privilege application program wants to use a system service which is at a high privilege level. An example will be when an application program wants to write data to disk, which involves the use of a device driver. In this case, an indirect access is allowed through a 'gate' which is stored in the LDT or GDT.

There are three types of gates for this:

- Task gate
- Interrupt gate
- Call gate

In this section, only the call gate will be described.

To allow a program to jump to a more privileged level, it must go through Call gates, which basically defines the entry points for the privileged code. Corresponding security checks will be conducted at those entry points to decide whether the invoking code has sufficient rights. These security checks are enforced by operating systems.

15.10.7 | Call Gates

A call gate is an 8-byte entity which contains the following information:

- Segment selector of the code segment to be accessed.
- Entry point for a procedure in the specified code segment (offset in segment).
- Privilege level required for a caller trying to access the procedure (DPL).
- Parameter count – if a stack switch occurs, it specifies the number of optional parameters to be copied between stacks.

Access Control Policy for Call Gates

- In general, $CPL \leq DPL$ of the call gate (numerically).
- For a CALL instruction: DPL of the target code segment $\leq CPL$ (only calls to the more privileged code segment are allowed through a gate) – in numerical terms.
- For a JMP instruction: DPL of the code segment = CPL . Gates cannot be used for jumping to a different privilege level.

This indirect access mechanism permits access only if a certain check is done. The privilege level of the calling program (i.e., CPL) is compared with that of the call gate. If the former has a privilege level higher than or equal to the DPL of the gate, access is allowed – otherwise




Figure 15.24 | Structure of a call gate

it is denied. This can be elaborated this way – a call gate with a DPL of 2 can be used by a program with CPL 2 to access a system service at level 0 – but the same gate cannot be used by a program with a CPL of 3. Figure 15.24 shows a call gate.

How to Use Call Gates?

The instruction Call xxxxx may be used where xxxxx may be specified as a gate name, but usually generates a selector leading us to an entry in the GDT or LDT, where the gate is stored. In the call gate, the detail of the target code segment is available and using this, the calling program enters the code segment at a higher privilege, but only at the defined (in the gate) entry point. Remember that a gate cannot be used to enter code segments at lower privilege levels.

15.11 | Multi Tasking

What is meant by the word multitasking in the context of computers and processors? If a computer can handle multiple tasks at the same time, it has the capability of multitasking. However, the phrase ‘at the same time’ is a bit tricky. Can a single processor perform the task of executing multiple programs at the same time i.e., in parallel?

Unlikely. Most of us have observed Windows OS allowing multiple tasks to be attended to, simultaneously. Then, there are systems with just one processor, but multiple terminals and multiple users. It seems like all the users’ programs are being done concurrently. However, actually, that is not the case. A single processor can perform the execution of only one code sequence. When multiple tasks are to be handled, task switching is being done.

A program in execution is called a process, or task. Suppose there are 10 tasks that the system has to handle – the OS has to arrange a sequence of how to organize the execution of these tasks. This multitasking OS can use one of the many concepts available. There is the ‘time slice’ method in which one task is given a certain time to execute – after this, the task has to give way to the execution of the next task and so on. In this way, all the ten tasks are executed in one round, and the user/users feel that all tasks are running concurrently. In the next round, once again all the tasks get executed and this continues.

In priority based task scheduling, tasks which are more important are allowed to execute, while low priority tasks are made to wait. There is also the co-operative scheduling system in

which a task is allowed to execute until its completion. It is the OS which decides the task scheduling mechanism. Windows is a multiuser OS while DOS is a single user OS.

Task Switching When one task is temporarily abandoned and another is taken up, context switching or task switching is said to have occurred.

15.11.1 | Issues in Multi-Tasking Systems

One main issue in task switching is that, before abandoning a task temporarily, its 'context' is to be saved so as to resume execution at the point at which it had left off. The context consists of the content of the registers being used, the state of the stack and all the information necessary to resume the old task. This is similar to returning from an interrupt, but more variables and more parameters have to be saved before switching to a new task.

If a processor is said to have 'multitasking' capability, it means that it has hardware features to handle all the issues associated with task switching. It should have the necessary mechanisms, instructions, registers and necessary additional hardware. These features are available in the 80386 and are useful in designing a multitasking OS for this processor.

TSS and LDTs It was mentioned earlier that if there are 'n' tasks in a system, there will be n LDTs – i.e., for each task there is an LDT and each LDT holds the descriptors of the segments required to run the particular task. When a particular task is taken up for execution, the selector for its LDT is loaded into the LDTR in the processor. Thus, we need to remember that an LDT is associated with a particular task and that there are as many LDTs in the system, as there are tasks. The descriptors of all the LDTs of a system are stored in the GDT of the system.

Similar to this, there is another structure associated with a particular task and that is its 'Task State Segment' or TSS. Each task has a TSS. All the information pertaining to a particular task is stored in its TSS, which is also a memory array or segment stored in main memory.

The TSS has the following fields.

- i) A dynamic set which changes every time a task switch occurs. They are:
 - The general registers (EAX, ECX, EDX, EBX, ESP, EBP, ESI, and EDI).
 - The segment registers (ES, CS, SS, DS, FS, and GS).
 - The flags register (EFLAGS).
 - The instruction pointer (EIP).
 - The selector of the TSS of the previously executing task (updated only when a return is expected).
- ii) A static set that the processor reads but does not change. This set includes:
 - The selector of the task's LDT.
 - The register (PDBR) that contains the base address of the task's page directory (read only when paging is enabled).
 - Pointers to the stacks for privilege levels 0-2.
 - The T-bit (debug trap bit) which causes the processor to raise a debug exception when a task switch occurs.
 - The I/O map base.

Figure 15.25 shows the structure of a typical TSS.




Figure 15.25 | Structure of a typical TSS

The TSS is a segment and hence it has a descriptor that is stored in the GDT of the system (just as the LDT's descriptor is). The descriptor of a TSS is similar to a segment descriptor except that it has a 'B' or busy bit. The B-bit in the type field indicates whether the task is busy. The B-bit allows the processor to detect an attempt to switch to a task that is already busy.

15.11.2 | Task Register

The processor has a register called a task register (TR) which points to the current task. When a particular task is taken up, the selector of the TSS is loaded into the visible portion of the task register. The selector contains the index that locates the descriptor of the corresponding TSS in the GDT. Then the descriptor of the TSS of this task gets loaded into the hidden part of the register.

Two important points regarding the execution of the current task is that it has the TSS selector in the Task register (TR) and the LDT selector in the LDT register (LDTR). This initiates the actions necessary to access the contents of the LDT and TSS that belong to that task. When task switching occurs, the contents of these registers are changed.

15.11.3 | Task Switching

How can task switching be done?

- The current task executes a JMP or CALL that refers to a TSS descriptor.
- The current task executes a JMP or CALL that refers to a task gate.
- An interrupt or exception vectors to a task gate in the Interrupt Descriptor Table.
- The current task executes an IRET when the NT flag is set.

Let us discuss the above four cases in more detail.

- i) In the first case, using instructions like JMP taskname or CALL taskname can load the selector of the corresponding TSS into the task register (TR) which gets the TSS descriptor from the GDT.
- ii) In the second case, the task switching is indirect. There is a 'task gate' involved just like the call gate we had discussed in the previous section. In both cases, a new task is embarked upon.
- iii) Task switching can also occur by the occurrence of an interrupt. This is also logical, but it involves an 'interrupt descriptor table' which we are yet to discuss. In any case, we know that an interrupt or exception causes the current task to be abandoned and a new one to be taken up. Thus, it can lead to a task switching when there is the mechanism to access the TSS of a new task.
- iv) Whether invoked as a task or as a procedure of the interrupted task, an interrupt handler always returns control to the interrupted procedure in the interrupted task. If the Nested Task (NT) flag is set, however, the handler is an interrupt task, and the IRET switches back to the interrupted task.

What are the Actions that Follow a Task Switch?

- i) If the switching occurs by a JMP or CALL instruction, the usual rules of privilege levels apply for the TSS descriptor or task gate, and task switching occurs only if the conditions are satisfied. Checking of the P bit and limit bit of the TSS descriptor is also done.
- ii) The context of the current task is saved i.e., contents of the registers are copied to the TSS. The EIP register which is now stored in the TSS will be pointing to the next instruction after the one that caused the task switch. These actions correspond to 'saving the context' of the abandoned task.
- iii) Next, the actions for taking up the new task are begun. The Task Register (TR) is loaded with the new task's selector. The selector is either the operand of a control transfer instruction or is taken from a task gate which gets the TSS descriptor and then the TSS from memory.
- iv) Now the new task can be proceeded with. For that, its state is loaded from the TSS into the corresponding registers. The registers loaded are the LDT register, the flag register; the

general registers EIP, EAX, ECX, EDX, EBX, ESP, EBP, ESI, EDI, the segment registers ES, CS, SS, DS, FS, and GS, and PDBR.

- v) Every task switch also sets the TS (task switched) bit in the MSW (machine status word) i.e., CR0.

15.11.4 | Task Linking

The back-link field of the TSS and the NT (Nested Task) bit of the flag word together allow the 80386 to automatically return to a task that called another task or was interrupted by another task. When a CALL instruction, an interrupt instruction, an external interrupt, or an exception causes a switch to a new task, the 80386 automatically fills the back-link of the new TSS with the selector of the outgoing task's TSS and, at the same time, sets the NT bit in the new task's flag register. The NT flag indicates whether the back-link field is valid. The new task releases control by executing an IRET instruction. When interpreting an IRET, the 80386 examines the NT flag. If NT is set, the 80386 switches back to the task selected by the back-link field.

15.12 | Interrupts of 80386

The interrupts of 80386 function exactly in the same way as the 8086 interrupts and the categorization is also similar i.e., the hardware interrupts are NMI and INTR and the software interrupts can be used in the format INT n . There are interrupts generated due to errors and they are called exceptions. The interrupt response is also similar i.e., the interrupt vector is in a table. On reset, this table is at location 0 and has a maximum size of 1 K.

Now, for the differences – instead of the 'interrupt vector table', here it is the 'Interrupt Descriptor Table (IDT)' that has the interrupt vectors. There is one and only one IDT for a system. It can be located anywhere in the physical memory, and that is specified by the descriptor of this table which will be in the GDT. When the system is in protected mode, the descriptor of the IDT is loaded into the IDT register (IDTR).

IDT Descriptors There are three types of descriptors that may be placed in the IDT. They are:

- Task gates
- Interrupt gates
- Trap gates

Task gates have been discussed already. An interrupt can cause task switching if the target segment corresponds to a different task. That is why we can have task gates also in the IDT.

We know that a gate contains an entry point into a code segment. It contains a selector of the destination code segment and the offset of the allowed entry point. Call gates have already been discussed. An interrupt is similar to a call – the response is similar, but the initiation mechanism is different. Also an interrupt causes the flags to be saved on the stack, and the return instruction is IRET rather than RET. So, we do not have call gates in the IDT, but instead have trap and interrupt gates.

The difference between an interrupt gate and a trap gate is in the effect on IF (the interrupt-enable flag). An interrupt that vectors through an interrupt gate resets IF, thereby preventing other interrupts from interfering with the current interrupt handler. A subsequent IRET




Figure 15.26 | Interrupt vectoring

instruction restores IF to the value in the EFLAGS image on the stack. An interrupt through a trap gate does not change IF. The figure shows the channeling of an interrupt from the IDT to the interrupt vector. Because descriptors are involved in this sequence, privilege level checking is done for interrupts as well. The privilege rule that governs interrupt procedures is similar to that for procedure calls.

15.12.1 | Exceptions

Exceptions are caused by error conditions. For the case of the 8086, we know that a divide by zero error causes a type 0 interrupt. We can call it an exception. The 80386 has more error conditions under which interrupts can occur. Table 15.2 shows the list of the 80386 interrupts. Compare it to Table 8.2, which gives the interrupt listing for a PC.

15.13 | Privileged Instructions

The instructions that affect system data structures can only be executed in the protected mode and it is also to be ensured that they are used only at the highest privilege level i.e., when CPL is zero. These instructions include:

- CLTS — Clear Task Switched Flag
- HLT — Halt Processor
- LGDT — Load GDT Register
- LIDT — Load IDT Register
- LLDT — Load LDT Register

Table 15.3 | List of 80386 interrupts

Identifier	Description
0	Divide error
1	Debug exceptions
2	Nonmaskable interrupts
3	Breakpoint (one-byte INT 3 instruction)
4	Overflow (INTO instruction)
5	Bounds check (BOUND instruction)
6	Invalid opcode
7	Coprocessor not available
8	Double fault
9	(Reserved)
10	Invalid TSS
11	Segment not present
12	Stack exception
13	General protection
14	Page fault
15	(Reserved)
16	Coprocessor error
17-31	(Reserved)
32-255	Available for external interrupts via INTR pin

- LMSW — Load Machine Status Word
- LTR — Load Task Register
- MOV to/from CRn – Move to Control Register n
- MOV to/from DRn – Move to Debug Register n
- MOV to/from TRn – Move to Test Register n

In fact, all the instructions involving load/store of protected mode registers are privileged instructions.

15.13.1 | Switching to Protected Mode

It has been mentioned that when the processor is powered on, it is in the real mode. How is it switched to the protected mode?

To switch it to protected mode, a number of initialization operations need to be done. These can be done in the real mode itself. (These things can be done in the protected mode as well, but then it must be ensured that the initialization procedures must not use protected-mode features that are not yet initialized.) For example, the GDTR must point to a valid GDT, the IDTR must point to a valid IDT and so on – after all the initializations are done, the switching may be done.

Next, to switch it to protected mode, the PE (Protection Enable) bit of control register CR0 must be set. Those who are familiar with the 80286 will remember that there is a LMSW (Load Machine Status Word) instruction which was executed with PE = 1. It is this same register MSW that has been re-named as CR0 for the '386.

The difference between the '286 and '386 is that in the case of the former, it is not possible to return to the real mode without doing a hardware reset. However, in the case of the '386, returning to the real mode is possible by clearing the PE bit and loading it to CR0 again. Of course, various conditions will have to be arranged before switching back to the real mode. All these conditions can be put inside a procedure, with the final instruction being the clearing of the PE bit of CR0. Thus, the processor returns to real mode.

15.13.2 | Virtual 8086 Mode

We know that even though the protected mode is important, running programs in MS-DOS has its charms and to facilitate the use of both these modes, the '386 has a virtual mode which allows switching between the real and protected modes. Thus, a multitasking system can have users running programs in the protected mode as well as users running real mode 8086 programs. Time slices are allotted for all these programs and thus there is continuous switching between real and protected mode. Virtual 8086 mode divides the computer into multiple address spaces and maintains virtual registers for each virtual machine. This mode is entered by setting the VM (Virtual Machine) bit in the EFlags register.

The desire to allow execution of MS-DOS applications under the control of a protected-mode environment, (such as Windows) has led to the inclusion of virtual-mode to all of Intel's 32-bit processors. When the processor is running in the virtual machine mode, it behaves as if it were an 8086 equipped with protection, multitasking and paging support.

Note that the virtual 8086 mode has nothing to do with 'virtual memory', but is associated with the term 'virtual machine' which can be explained as follows. 'A virtual machine is a type of computer application used to create a virtual environment, which is referred to as virtualization. The main advantage of system virtual machines is that multiple OS environments can co-exist on the same computer, in strong isolation from each other'.

15.14 | Conclusion

With this, we come to the end of our discussion on 80386. Most of the important features have been explained, but to be able to write instructions to use all these features, more study and more information (from the programming manual) is required. The attempt here has to give an idea of the capabilities of one of the more sophisticated processors of the Intel family. The 80486 and Pentium have all these features, so this chapter has to be understood for appreciating the characteristics of Pentium.

KEY POINTS OF THIS CHAPTER

- The 80286 had many improved features compared to the 8086 and fuelled the development of the PC-AT and the first Windows based operating system – Windows 3.1.
- The 80386 was the next processor of Intel in which the major change is the increase in the number of address and data pins to 32.
- This processor can address 4 GB of physical memory.
- It has a number of new instructions, and, also enhancements to some existing instructions.
- The general-purpose registers of 80386 are 32 bits long.

- It has a number of additional pins compared to its predecessors.
- The idea of virtual memory is of having an enlarged memory space which is mapped to a smaller physical memory.
- The memory management unit of the 80386 has the capability of paging and segmentation.
- Descriptors are 8-byte entities in which the information corresponding to a segment are stored.
- Selectors are 16 bits long and point to the descriptor it is associated with.
- The paging unit of the MMU can be switched off, but not the segmentation unit.
- Protection has many aspects, including protecting tasks from each other, and also the necessity of protecting system level tasks from user tasks.
- The concept of privilege levels is used for prescribing the level of trust, and the level of protection for each segment.
- Call gates are used for accessing code segments at a higher privilege level.
- The 80386 is a multitasking CPU meaning that it can take care of multiple tasks.
- The tasks are not executed concurrently; the CPU time is allotted to different tasks.
- The interrupt vector table is replaced by an Interrupt descriptor table while in the protected mode.


QUESTIONS

1. What were the major changes incorporated in the 80286 compared to the 8086?
2. What was the major complaint against the 80286?
3. Name the major change that made the 80386 a big leap in processor technology.
4. What is the 'scaled indexed' mode of addressing?
5. Where can the MOVSB and MOVZX instructions be used?
6. Why is it that the '386 has address lines only from A2 to A31?
7. Does the 80386 have an inbuilt floating point processor?
8. What is the importance of the virtual memory concept for a user?
9. Why is address translation in the PVAM so complex?
10. What does a descriptor mean to a segment?
11. What is the role of a selector in address translation?
12. How many LDTs and GDTs does a system have?
13. Why do we say that a segment with a CPL = 0 is a highly trusted segment?
14. What is the role of RPL in allowing access to a segment?
15. Distinguish between a conforming and a non-conforming segment.
16. Under what condition is a call gate used?
17. What is the main issue to be taken care of in task switching?
18. What does the task register contain?
19. What is meant by the term 'back link to the previous TSS'?
20. What is an IDT? How many IDTs can a system have?

EXERCISE

1. Write a program that uses MOVZX and MOVSB instructions to
 - a) add an 8-bit negative number to a sixteen bit negative number,
 - b) add a 32-bit negative number to a 16-bit positive number.
2. Write a program to test bits 4, 8, 9 of a 16-bit number in memory.
3. Add the contents of a table of size 27 containing 16-bit numbers using the scaled indexed mode of addressing.
4. Find the difference between the processors 80386DX and 80386SX.
5. In virtual 8086 mode, can the BIOS and DOS instructions be used by 8086 based programs? How?

16 THE PENTIUM PROCESSOR



In this chapter, you will learn

- The features of 80486 which are an enhanced set of the '386 features.
- The idea of burst mode data transfers for both 80486 and Pentium.
- Why data alignment is important for multibyte processors.
- The features of Pentium.
- The cache structure of Pentium.
- Superscalar architecture and branch prediction.
- Improved features of the sixth generation of the x86 family.
- How multiprocessing relates to multi-core processing.
- The concepts of ILP and TLP.

Introduction

This chapter has been titled 'Pentium', but here we will discuss the 80486 as well. In Chapter 15, the first 32-bit processor of Intel i.e., the 80386 was discussed with the aim of cracking the intricacies of a complex processor. Once the '386 is understood well, the '486 and Pentium have nothing more to it than a few 'enhancements'. This means that the internal architecture of the '486 and Pentium is more or less the same as that of the '386. They have the same memory management unit, protection mechanisms and multitasking units as we have discussed in Chapter 15. Thus, a discussion on these features is redundant and we will concentrate on the additional features of the 80486 and Pentium, in addition to that of 80386.

16.1 | The Enhanced Features of 80486

- i) The 80486 processor was introduced by Intel in 1989. The first important difference that the '486 has in comparison with its predecessor is that it has an integrated floating point unit on the chip itself. Previous processors had the arithmetic unit as a separate unit – the 8086 had 8087, the 80186 had the 80187, the 80286 had the 80287, the '386 had the '387 and so on. Logically, the 80486 should have had an arithmetic co-processor 80487 external to the chip, but instead, Intel placed the arithmetic co-processor inside the chip and called the whole chip as 80486DX. To improve its market segment, Intel also sold a '486 processor without an FPU (floating point unit, as the arithmetic co-processor is called). It was named as 80486SX, but in fact it was just the 80486 processor with its FPU turned off.

- ii) The second important enhancement is that of adding an 8 KB of cache **on the chip**, for data as well as code. All predecessor processors had cache outside the processor. With this enhancement, on-chip cache came to be designated as L1 (level 1) cache while off chip cache is called L2 (level 2) cache. Control register CR0 is used to control the internal cache with two new control bits which were not present in the 80386 microprocessor.
- iii) The internal cache improved the memory access speed substantially, but later versions had something called ‘clock doubling’. New editions were released with higher clock frequencies, as they hit on the idea of *doubling* the *internal* clock frequency in relation to the external clock. These double-clocked processors were given the name, 80486DX2. A very popular model in this series had an external clock frequency of 33 MHz while working at 66 MHz internally. This principle (*over-clocking*) has been allowed in one way or another in all later generations of CPU’s, though there are problems associated with it. It is just an option allowed for a processor and is specified in the motherboard spec.
- iv) Compared to the 80386, the ’486 is a heavily pipelined processor. It has a 5 stage pipeline as shown in Fig 16.1. Each stage takes one clock cycle, but once the pipeline is full, each instruction will execute in a single clock.

Referring to Figure 16.1, the stages in the pipeline are as PF-Pre-fetch, D1-decode1, D2-decode2, EX-Execute and WB-Write back. I₁ to I₅ correspond to five instructions in the pipeline. As per this figure, there are two decoding stages. This is because of the varied addressing modes of 80486 and the necessity for protection checks before any access is allowed. This processor has a 5 stage pipeline – if the pipelining stages are split into smaller and smaller steps, it is called ‘super pipelining’, but such an idea is not used in the 808486.

- v) ’486 processors and above support burst mode transfers. This feature is helpful in saving time when accessing data from consecutive addresses. Let us try to understand this important feature. Figure 16.2 shows the read cycle timing for an ordinary memory read cycle. Two clock cycles are required to transfer data (which is 32 bits long). The first clock cycle T1 provides the memory address and control signals, while in T2, data is transferred between the memory and the microprocessor. The figure shows the case of two 32-bit words being transferred. $\overline{ADS}$ is address strobe. Thus two clock cycles are required to transfer 32 bits of data. If 128 bytes of data is to be transferred, obviously 8 clock cycles will be required. Now, let us contrast this with a burst cycle. A burst cycle will require only 5 clock cycles to transfer 128 bytes of data. How do we get such a number? See Fig 16.3.

A burst cycle transfers multiple bytes of data across the bus during one long memory cycle. For example, a transfer of a 128-bit data item across a 32-bit bus would normally occur in four groups, each group containing one 32-bit data. But in the burst mode, the




Figure 16.1 | Five stage pipeline of the 80486




Figure 16.2 | Non-burst bus cycle




Figure 16.3 | Burst mode data transfer

initial address (of the first byte) is used by the processor to compute the remaining addresses for the subsequent data bytes. There are various ways of accomplishing this, but to understand it, think of it this way – a sequential burst order from address zero would first transfer the 32-bit data residing at address zero. The next transfers would be to the data at addresses 4, 8 and 12, in order.

Figure 16.3 shows a burst cycle which reads four 32-bit words (a total of 128 bytes) in five clocking periods. To start the process, the processor sends out the first address and asserts the BLAST signal (not shown in the figure) high. The $\overline{BRDY}$ (Burst ready) signal is asserted in the first T_2 . The processor reads the data word in two cycles and outputs the next address. Since data are at successive addresses, only the lower address bits need to be changed. When the processor has read the required number of data words, it asserts the BLAST signal low to terminate the burst mode. Such a burst read cycle is referred to as a 2-1-1-1 read where four data reads are accomplished in 5 clock cycles.

- vi) The processor has a major addition in terms of four pins (one for each memory bank) for parity, which is generated by it, during each write cycle. Parity is generated as even parity




Figure 16.4 | Memory banks of 80486

and a parity bit is provided for each byte of memory. The parity check bits appear on pins DP0–DP3, which are parity inputs as well as parity outputs. These are typically stored in memory during each write cycle and read from memory during each read cycle. On a read, the microprocessor checks parity and generates a parity check error, if it occurs on the $\overline{PCHK}$ pin. A parity error causes no change in processing unless the user applies the $\overline{PCHK}$ signal to an interrupt input. Figure 16.4 shows the memory banks of 80486 with the bank enable signals ($\overline{BE}$ s) and the parity check bit for each channel.

- vii) Another enhancement which needs mention, is the addition of a few new instructions which are listed as below:

Menemonic	Function performed
BSWAP	Byte Swap
XADD	Exchange and add
CMPXCHG	Compare and Exchange
INVD	Invalidate cache
INVLPG	Invalidate TLB entry
WBINVD	Write back and invalidate cache

Of these, the last three instructions relate to cache operations. The first instruction BSWAP is very useful for systems which use big endian format, instead of the little endian format used by Intel. Changing from one format to the other just involves the switching of the bytes of a word.

Now, before going ahead to the next topic, let us try to understand the meaning and implications of an issue that needs careful consideration.

16.2 | Data Alignment

In Section 7.3 the idea of memory banks was presented, and the importance of using an even address when accessing 16-bit data was explained. It was made clear that when the memory address is even, memory access takes only one bus cycle, while an odd addressed word incurs a penalty of one more bus cycle. That was for the case of the 8086 where the maximum data

width was 16 bits. However, for '386 onwards the bus width is 32 bits and data is divided into 4 memory banks. For Pentium, the external data bus width is 64 bits. In all these cases, a bus penalty will occur if the memory address is not 'aligned'. A performance penalty occurs due to misaligned data. What is meant by data alignment?

2-byte Data A 16-bit data item is aligned if it is stored at an address that is a multiple of two. This implies that the least significant bit of the address must be 0. This is the case that we have discussed for the 8086, where the LSB of the address being 0, means that the address is an even number.

4-byte Date Let us think of it this way – if all the 4 bytes are in one row, they can be accessed in one cycle, otherwise one more cycle is required. This means that a 32-bit data item is aligned if it is stored at an address that is a multiple of 4. This implies that the least significant two bits of the address must be zero.

8-byte Data A 64-bit data item is aligned if it is stored at an address that is a multiple of 8. That means that the least significant three bits of the address must be 0. This alignment is important for processors such as Pentium that have a 64-bit data bus.

16.3 | The Pentium Processor

Now, we will come to Intel's most popular name in processors i.e., Pentium. In 1993, Intel developed its fifth generation x86 processor and decided to call it 80586 (according to the convention followed till then), but to keep away competitors such as AMD and a few others, Intel wanted to trade mark this processor title. However, the lawsuit that followed, declared that numbers cannot be trademarked, and hence Intel had to find a new name for its new processor, and that is how the word Pentium came into being. It is also called P5. The name Pentium was derived from the Greek word *pente*, meaning 'five', and the Latin ending -ium. Vinod Dham, an engineer from Pune who worked in Intel, is frequently hailed as 'the father of Pentium'.

The Pentium family of processors, which has its roots in the Intel 486 processor, uses the Intel 486 instruction set (with a few additional instructions). The term 'Pentium processor' refers to a family of microprocessors that share a common architecture and instruction set, but each new version of Pentium has newer and newer features. The first Pentium processor was released in the year 1993 and it had a lot of new features compared to its predecessor i.e., 80486. The foremost enhancement is that it has a 64-bit data bus, while continuing to be a 32-bit processor internally. See the list below.

Features of the Pentium Processor family:

- 32-bit Microprocessor (internal registers of 32 bits)
 - 32-bit address bus
 - 64-bit data bus
- Superscalar architecture
 - Two pipelined integer units
 - Capable of 'less than' one clock per instruction
 - Pipelined floating point unit
- Separate code and data caches
 - 8 K code cache, 8 K write back data

- Advanced design features
 - Branch prediction

Now, let us go into the details of each of these enhancements.

16.3.1 | Data Bus of 64 Bits

The data bus width has been doubled. What is the advantage of that? It is that data can be received as 64 bits in one go. What can be done with 64-bit data when the processor registers are only 32 bits in width? The answer is that since the bus bandwidth has doubled, double the amount of data can be sent/received from/to the memory.

See Fig 16.5. There are 8 memory banks activated by eight $\overline{BE}$ signals. From this memory bank array, data can be accessed for internal processing as 8, 16 or 32 bits. However, sending or receiving data outside the chip can be at the rate of 64 bits per memory cycle. This is a major enhancement.

16.3.2 | Superscalar Architecture

Pentium is a ‘superscalar processor’ which is pipelined. Pipelining is still only of 5 stages. Being superscalar means it has more than one executing unit – here, there are two integer execution units (each of them is pipelined). They are called the U and V pipes – both these units are slightly different in their processing power, which means that one unit can handle more complex operations than the other. It is the U pipe which can handle more complex operations. See Fig 16.6 which shows the execution units of the Pentium and the ‘486. Note the differences.




Figure 16.5 | Memory banks of Pentium




Figure 16.6 | Execution units of pentium and the ‘486 compared

16.3.3 | Faster Floating Point Unit

The floating-point unit has been completely redesigned over the Intel '486 CPU. Faster operations provide up to 10x speed for common operations including add, multiply, and load. The FPU has an eight stage pipeline to speed up operations. Figure 16.7 shows that the units responsible for fetch and decode are common – once decoding is over, it goes to the execution phase. In this phase, the execution occurs in either the U pipe, V pipe or the floating point unit, depending on the data type and complexity of the required operation. Once this is done, the result is written into the data cache.

16.3.4 | Separate Data and Instruction Cache

Figure 16.7 also shows that there are separate instruction and data caches. Pentium has two separate on-chip caches, one for data, and one for code. Each cache is 8 KB in size, with a 32-byte line




Figure 16.7 | Execution unit of Pentium

size and is 2-way set associative. Each cache has a dedicated Translation Lookaside Buffer (TLB) to translate linear addresses to physical addresses. The code (instruction) cache is an inherently write-protected cache. The caches can be enabled or disabled by software or hardware.

16.3.5 | Data Dependency Problem

Increasing the number of execution units brings with it a new set of problems. Programs are written by a programmer for serial execution. However, when attempting to run this program on two parallel execution units, the problem of 'data dependency' can arise for certain program sequences. For example, if the first program line asks for two numbers to be added, and the second line needs the result of the first instruction for its completion, obviously these two instructions cannot be executed in parallel. In such cases, re-ordering or re-scheduling of instructions has to be done, and at present this problem has been handed over to the compiler to resolve.

16.3.6 | Branch Prediction Unit

Another new feature that Pentium has, is the branch prediction unit. This is useful in preventing 'pipeline stalling' in the case of a branch instruction occurring in the sequence.

Let us take a closer look at what is called the 'branch unit'. This unit includes the 'branch execution unit' and the 'branch prediction unit'. Think of what should happen if a conditional branch instruction is encountered. First, the condition is to be evaluated, then the branch address calculated and then the code should be fetched from the new branch address and instruction execution should proceed. All these steps involve a delay. To avoid the delay, a technique called 'speculative execution' is implemented. In this, a calculated risk is taken and one particular branch is taken for continuing the instruction execution in sequence. In the mean while, the branch condition is evaluated. If the 'speculative execution' had taken the correct branch, all is well and good. Otherwise, the pipeline will have to be flushed and everything re-done. This is indeed a minor catastrophe considering that the delay so incurred causes a performance penalty. Thus it is imperative that the branch prediction unit does a good job.

There are two methods of branch prediction-static and dynamic. In the first case, the assumption is that most branch conditions pertain to whether to repeat a loop or not. In most cases, a loop's exit condition is false; hence repeating the loop is something that is normally assumed to be necessary. This type of static prediction is simple and fast, but its success rate is good only in cases of loop intensive programs.

The other technique i.e., dynamic prediction relies on using the past history of branching, for the program. There is a 'Branch History Table', as well as a 'Branch Target Buffer'. The latter stores the target addresses of previously executed branches, so that speculative execution can use any of these target addresses, when branching has to occur. This assumes that previously used target addresses are likely to be needed again.

Pentium uses both static and dynamic branch prediction and it is found that a success rate of around 75 to 85% is possible.

16.3.7 | Page Size

Remember that the '386 paging unit supported pages with a maximum size of 4K only. This is so for 80486 as well. The Pentium processors' Memory Management Unit contains optional extensions to the architecture which allow 2MB and 4MB page sizes.

16.3.8 | System Management Mode

The first Pentium was released in 1993, and a newer version was released in 1994, which had a new mode of operation, besides the real, virtual and protected modes. This mode is the System Management mode, which is an additional feature for efficient power management, and to run specific proprietary code.

In this mode, the power dissipation is reduced by switching off all peripherals and even the system when it is not in use. This is not the only feature of this mode, which operates as a high level manager and has also extra features for enhanced security.

16.4 | Pentium Pro

In 1995, Intel released its sixth generation processor P6 and called it Pentium Pro. This was followed by a number of new processors named Pentium II and Pentium III. Intel's P6 architecture, first instantiated in the Pentium Pro, is by any calculation, a super success story. Its performance was significantly better than that of the Pentium, it became a hot favorite in the desktop and workstation market, and paved the way for Intel to compete in the RISC markets as well. See Table 17.2, which compares the processors of this generation.

	Pentium Pro	Pentium-II	Pentium-III
Year of release	November 1, 1995	May 7, 1997	February 26, 1999
Clock speed at introduction	200 MHz	300 MHz	500 MHz
Transistor count	5.5 million	7.5 million	9.5 million
L1 cache size	8 K instruction, 8 K data	16 K instruction, 16 K data	16 K instruction, 16 K data
L2 cache size	256 K or 512 K (on chip)	512 K (off chip)	512 K (on chip)
Features	No MMX	MMX	MMX, SSE

We will have a brief review of the enhancements of these processors over their predecessor, the first Pentium. It introduced several unique architectural features that had never been seen in a PC processor before. The Pentium Pro was the first mainstream CPU to radically change how it executes instructions, by translating them into RISC-like microinstructions and executing these on a highly advanced internal core. The Pentium Pro achieves performance approximately 50% higher than a Pentium of the same clock speed. In addition to its new way of processing instructions, the Pentium Pro incorporates several other technical advancements that contribute to this increased performance.

- **Superpipelining** The Pentium Pro dramatically increases the number of execution steps, to 14, from the Pentium's 5.
- **Integrated Level 2 Cache** The Pentium Pro features a dramatically higher performance secondary cache compared to all earlier processors. Instead of using motherboard-based cache

running at the speed of the memory bus, it uses an integrated level 2 cache with its own bus, running at full processor speed, typically three times the speed that the cache runs at, on the Pentium. The Pentium Pro's cache is also non-blocking, which allows the processor to continue without waiting on a cache miss.

- **32-bit Optimization** The Pentium Pro is optimized for running 32-bit code (which most modern operating systems and applications use) and so gives a greater performance improvement over the Pentium when using the latest software.
- **Wider Address Bus** The address bus on the Pentium Pro is widened to 36 bits, giving it a maximum addressability of 64 GB of memory.
- **Greater Multiprocessing** Quad 'multiprocessor' configurations are supported with the Pentium Pro compared to only dual with the Pentium. This means that four Pentium Pro processors can be interconnected, on the Pentium Pro processor bus.
- **Out of Order Completion** Instructions flowing down the execution pipelines can be executed 'out of order'.
- **Superior Branch Prediction Unit** The branch target buffer is double the size of that of Pentium and its accuracy is higher, as much as 94%.
- **Register Renaming** This is one of the solutions for the pipeline stalling that might occur when data dependency prevents two sequential instructions to be executed in parallel in the U and V pipes. This feature improves parallel performance of the pipelines.
- **Speculative Execution** The processor uses speculative execution to reduce pipeline stall time in its RISC core.

16.5 | Pentium-II and Pentium-III

These two processors have the Pentium Pro as its architectural core, while providing a few new instructions and enhancements.

16.5.1 | MMX

In 1997, a new set of instructions catering to multimedia and called MMX (for Multimedia Extension) were introduced in Pentium processors. Such P5 processors were designated as Pentium MMX. MMX is also a feature in P-II and P-III as Table 17.1 shows. MMX technology is designed to accelerate multimedia and communications applications by including new instructions and data types that allow applications to achieve a new level of performance. It exploits the parallelism inherent in many multimedia and communications algorithms, yet maintains full compatibility with existing operating systems and applications.

16.5.2 | SSE

This stands for 'Streaming SIMD Extension' where SIMD means 'Single Instruction, Multiple Data'. SSE includes 50 new instructions, which enable simultaneous, advanced calculations of more floating-point numbers with a single instruction. It accelerates performance on a wide variety of applications like video, audio, 3D graphics and image processing.

16.6 | Pentium-IV

In November 2000, Intel introduced its seventh generation processor Pentium-IV, which had high clock speeds (starting from 1.5 GHz) as an important feature, and a totally new architectural feature called the ‘netburst’ architecture which includes the following:

- Hyper Pipelined Technology
- 400MHz System Bus
- Execution Trace Cache
- Rapid Execution Engine
- Advanced Transfer Cache
- Advanced Dynamic Execution
- Enhanced Floating Point and Multimedia Unit
- Streaming SIMD Extensions 2

We will not discuss all these in detail here – instead, we will talk about the new trends in processor design.

16.7 | Latest Trends in Microprocessor Design

With clock speeds reaching very high values, which leads to more and more power dissipation, newer ways of enhancing performances are being considered. Thus, the latest trends in microprocessor design are focused towards two ideas – multithreading and multi core processing. Let us now try to get a feel of what exactly these mean.

16.7.1 | Multi Core Processors

In Chapter 13, we discussed ‘multiprocessors’ – the idea of having two or more processors tied together with common resources – we discussed the associated problems and the possible solutions. In this, each processor is physically separate and self contained – but basically we can consider each of them as a ‘processing element’, and when many such processing elements work in synchronism, the resultant system is more powerful and gives a higher throughput.

In recent times, the same idea is being used with a new name – multicore technology – the only difference being that the processing elements are not separate chips, but rather two or more processing elements in the same chip – manufactured on the same die. Each of these processing elements is now called a ‘core’. Thus, now ‘multicore processing’ has come of age and is the new industry trend and is growing rapidly. Since the cores are on the same die, they are physically very close and communication between them is fast. Ideally a dual core processor (one with two cores) should be twice as powerful as a single core processor. However, practically, performance gains are said to be about fifty percent which means that a dual core processor is likely to be about one and a half times as powerful as a single core processor.

Multiple cores can be heterogeneous or homogeneous – in the former; two different types of cores are on the same die – for example, like having a general-purpose processor and a math processor in the same chip. This is already present in the 80486 and Pentium – but when we talk of multicore processors, we usually mean ‘homogenous’ cores.

Dual core and quad core processors have been manufactured by Intel, AMD and ARM for example. The cores are typically integrated onto a single integrated circuit die (known as




Figure 16.8 | **a** A single core microprocessor




Figure 16.8 | **b** A multi core microprocessor

a chip multiprocessor or CMP, or they may be integrated onto multiple dies in a single chip package. Figures 16.8a and b show the differences between a single core processor and a multi core processor. The figures should clarify what exactly a core should mean. In a single core processor, there is only one ALU and register file, which is what is meant by the ‘core’. In a multi core processor, there are many such cores. The cores in Fig 16.8b share the same interconnect to the rest of the system. Each ‘core’ independently implements optimizations such as superscalar execution, pipelining, and multithreading – the last term needs more explanation and that will be presented soon.

The benefits of multicore technology should be supported by software – then only the performance benefits will be obvious. Research is ongoing in this field and being able to parallelize a program effectively is the key to success. Now, to understand the idea of parallelism, let us get to know a few new terms.

16.7.1.1 | Thread

The idea of a thread is not very new, only it has become more relevant and important these days. A thread of execution results from a fork of a computer program into two or more concurrently running tasks. Each thread

- has its own instructions and data,
- may be part of a parallel program or independent programs,
- each thread has all the states (instructions, data, program counter, register state, and so on) needed to execute.

If a thread is considered as a 'task', a single processor will have to do context switching to execute multiple threads. However, a thread can also be considered as a subset of a task, in which case, the execution of concurrent threads can be managed by software itself. Thus, as we see below, the OS can manage multithreading with just one processing unit.

16.7.1.2 | Multithreading

This is the ability of an operating system to execute different parts of a program, called threads, simultaneously. The programmer must carefully design the program in such a way that all the threads can run at the same time without interfering with each other. Multithreading as a wide-spread programming and execution model allows multiple threads to exist within the context of a single process. These threads share the process resources but are able to execute independently. When there is hardware support for multithreading, the system becomes more effective

16.7.1.3 | TLP and ILP

TLP (Thread Level Parallelism) It is explicitly represented by multiple threads of execution that are inherently parallel. In this, multiple instruction streams are used to improve the throughput of computers that run many programs and reduce the execution time of multi threaded programs

ILP (Instruction Level Parallelism) We have seen a number of steps used in the design of Pentium processors for getting more instructions executed in unit time. They include instruction pipelining, superscalar architecture, out of order execution, branch prediction and speculative execution. All these are categorized under the title of 'instruction level parallelism'. Thus, superscalar processors exploit ILP by executing multiple instructions from a single program in a single cycle.

16.7.1.4 | Simultaneous multithreading (SMT)

This exploits ILP as well as TLP. Simultaneous multithreading, is a technique for improving the overall efficiency of superscalar CPUs with hardware multithreading. SMT permits multiple independent threads of execution to better utilize the hardware resources.

SMT permits multiple independent threads to execute concurrently on the same core. In simultaneous multithreading, instructions from more than one thread can be executing in any given pipeline stage at a time. This is done without great changes to the basic processor architecture – the main additions needed are the ability to fetch instructions from multiple threads in a cycle, and a larger register file to hold data from multiple threads. For instance, if one thread is waiting for a floating point operation to complete, another thread can use the integer units.

16.7.2 | Multi Core Processing

Multicore processors use TLP by executing different threads in parallel on the different cores. With insufficient TLP, the different cores will be idle.




Figure 16.9 | A dual core processor with separate L1 cache and common L2 cache

Dual-core Intel Xeon processors Most modern operating systems, such as Linux and Microsoft Windows, support TLP. In addition to operating system support, adjustments to existing software are required to maximize utilization of the computing resources provided by multi-core processors. Figure 16.9 shows a dual core xeon processor with dual cores and a separate L1 cache for each core and running multiple threads on each core.

16.8 | Multi-Core Technology and Intel

Pentium-D processor at 3.2 GHz was the first server processor from Intel to integrate multi-core technology.

Intel® Core™2 processor family

- 1. Intel® Core™ i7
- 2. Intel® Core™ 2 Extreme
- 3. Intel® Core™ 2 Quad
- 4. Intel® Core™ 2 Duo

Here, 1 and 3 are four-core processors and 2 and 4 are dual core processors.

DualCore Intel processors enable two threads to be fully processed in their own processor cores. Hyper threading technology enables dual core processors to execute up to four simultaneous threads in unison.

Quad Core processor allows four threads to be processed in parallel. Hyper threading technology enables quad core processors to execute up to eight simultaneous threads in unison.

A few of Intel’s processor names have been mentioned throughout the book. However, there are various other Intel processor names which you are likely to see and hear of. Xeon is one such name. Intel Xeon is a high-performance version of Intel desktop processors intended for use in servers and high-end workstations. Xeon CPUs have the same features as Pentium 4/D, Core 2 Duo/Quad desktop microprocessors. Additionally, Xeon CPUs can work in dual processor systems. Xeon CPUs often have larger size of level 2 cache and may include a large level 3 cache. All these extra features are necessary for its application base (servers), which is more challenging than a desktop PC.

16.9 | Mobile Processors

A number of mobile processors have been manufactured by various companies, where the word 'mobile' implies its use in laptops which are portable i.e., mobile. Intel has a number of mobile processors. Pentium M (M for mobile) is one of its processors being used for mobile computing. More about this and Intel's Centrino technology is covered in Section 17. Another new and promising processor for mobile computing is Intel's Atom processor hailed as the 'smallest' of processors with comparable performance.

16.10 | Legacy Support

We started our discussions of x86, from the 8086 processor onwards. We have covered a lot of distance in coming to the Pentium and then to Intel's multi core processors. All along, whenever Intel developed the next generation processor of the x86 family, one aspect that had to be ensured was the backward compatibility issue, which means that the newest x86 processor should be able to run the instructions of the 16-bit 8086. We have seen the Pentium processor which has a 64-bit data bus, but which does all computations in 32 bits. The next generation of processors will definitely be 64-bit; however, these processors will also have to provide the 'legacy support' to run 16/32 bit applications. Such issues cause the newer processors to expend a certain part of its transistor budget for legacy support. Figures as high as 30% have been suggested for this budget, which is pretty high, but this may continue to be necessary if the x86 family has to remain in use.

KEY POINTS OF THIS CHAPTER

- The 80486 processor has a number of enhanced features compared to its predecessor 80386.
- It has the burst mode of data transfer which increases data transfer speed significantly.
- It has level one cache placed on the chip which was a big step in microprocessor history.
- Intel's first Pentium processor was released in 1993.
- One important enhancement is that it has a 64-bit data bus, while the processor is internally 32-bit.
- It has a superscalar architecture which means that it has two integer execution units.
- It has a faster floating point unit.
- Its instruction and data caches are separate.
- It can have a page size of 2 MB/4 MB.
- Pentium Pro is the first of Intel's sixth generation processors of the x86 family.
- Pentium II and III are enhanced versions of the Pentium Pro family.
- The new trend in microprocessor design is in adding more cores on the same die.
- Concurrency can be achieved by thread level and instruction level parallelism.
- Multithreading is a software method of parallel execution of instructions.

QUESTIONS

1. What is the difference between 80386 DX and SX?
2. What is L1 cache and how is it different from L2 cache?
3. What is overclocking? What are the possible pitfalls in this idea?
4. How does a burst access cycle of memory contribute to memory speed?
5. What is the penalty incurred in data misalignment?
6. What advantage is obtained by having a 64-bit data bus for a 32-bit processor?
7. What do the following words mean in the context of a Pentium processor?
 - a) superscalar architecture
 - b) branch prediction
 - c) code cache and instruction cache
 - d) pipeline stalling
8. Which is the first processor in the sixth generation of the x86 family?
9. What are the uses of the MMX and SSE group of instructions?
10. In what way do multi core processors improve performance?
11. Compare and contrast ILP and TLP.

EXERCISE

1. Find the names and frequency of operation of other processors of the Pentium-IV class of processors.
2. Find out how 'register renaming' helps to solve the problems of 'pipeline stalling'.
3. In Pentium Pro, there is an 'integrated Level 2 cache'. What exactly is this?
4. How does 32-bit optimization of Pentium Pro improve program execution performance?
5. Find out the features of Intel's Atom processor.
6. Find the names and features of non-Intel multi core processors.
7. Find out the names and details of Intel's processors, which do not have the 'x86 architecture'.

17 THE x86 BASED PERSONAL COMPUTER



In this chapter, you will learn

- The hardware features of the x86 based personal computer.
- The components that make up the motherboard.
- The meaning of the word 'chipset' and its significance.
- The devices connected to the North and South bridges.
- Why high speed is necessary for the AGP.
- The I/O buses such as ISA and PCI.
- The relevance and features of the PCI Express expansion bus.
- The features that have made USB very popular.
- The transition from ATA to SATA.
- What SIMM and DIMM mean.
- How dual channel RAMs are supposed to double memory speeds.
- The features of laptop computers.

Introduction

All along the chapters of this book, it was mentioned that the most important application of the x86 processor of all the different generations, was in the personal computer. The more powerful of these processors are used as servers, but the popularity of the x86 series stems from its use as the processing unit of the PC, which has become a household as well as an office appliance – a piece of equipment without which survival is very difficult, as it has blended into our lifestyle.

In this context, this last chapter is dedicated to creating an understanding of the internals of a PC. This chapter has not been written from a layman's point of view. It is written for a reader who has a moderately good understanding of the content of the previous chapters of this book. The aim of this chapter is to impart a certain amount of knowledge of the meaning of the terminologies and techniques used in the design and use of a PC. This chapter is aimed at 'understanding the PC' – no claim is made that the reader will be able to set up his own PC, upgrade his PC or even repair his PC. The point is that, once a basic understanding has been developed, other references can be sought by which the aforementioned activities can be tried out. In short, this chapter is aimed at bringing our study of the x86 processors, to a successful culmination.

Note In many of the figures, new terms will be seen. It may be necessary to read through the whole chapter, and then come back to the figures to understand these terms. The way to grasp the contents of this chapter is to read it once, get a feel of the terms used, and then read it again from the beginning.

17.1 | The Modern PC

We have, in previous chapters, seen the evolution of the PC and how things have changed over the years. We will go into the details of the PC as it is now, with the latest features, and in so doing, we will also have a peep into how and why changes have occurred.

What are the parts associated with a PC? They are:

- i) Processor, RAM and ROM
- ii) Video monitor, keyboard, mouse
- iii) Hard disk
- iv) Printer, speaker, modem
- v) CD/DVD drives, floppy drives (obsolete, more or less)
- vi) Connectors to optional devices like external hard disk, flash memory, infra red devices, blue tooth devices, digital camera and so on.

In this list, the first three sets are mandatory for any system. The fourth and fifth sets are necessary to get the full utility of a computer i.e., we would like to print, listen to music, read or write CDs/DVDs, access the internet and so on. The last set is for specialized uses. All this shows that the application realm of computers is large, and we get a lot of service from it – but such a large system with incompatible components have to work in synchronism. For all services, the CPU chip will be called on for action and this implies that each of these components should be able to communicate with the processor. This is where buses come into the picture. We know that buses carry address data and control information. Fast buses are used for communication between main memory and the processor, and relatively slow buses are used for peripherals. Since peripherals are of different types and specifications, there are different types of buses inside the PC – this will be one of the main areas of our discussion – which bus, at what speed connects to the hard disk, to the video processing unit and so on. Over the years, the data bandwidth of the buses, and speed has increased many fold – we will trace the evolution over the years, and converge on what is in vogue currently. We will also try to guess the direction in which changes are happening. Corresponding to each of the buses, there are connectors and we should know their features also.

The x86 PC started with 8088 as its CPU but the current popular processor is Pentium – it is internally 32 bits, but has an external data bus width of 64 bits – processors which are 64-bit ‘internally’ also, are coming up – but have not yet reached the general desktop/laptop PC market. So we will assume that the processor in our discussion is some version of Pentium, the different types of which has already been covered in Chapter 16. With this background, let us try to understand the working of the modern PC.

17.2 | The Motherboard

We will start with what is considered the most important part of the PC – the motherboard, the system board, or simply called the ‘board’. This is the circuit board which contains all the essential elements to make a computer work – this means a lot of things. The first and foremost item on this board is the CPU i.e., the processor. Besides holding the CPU, the motherboard is the centre through which all the communications in the computer system occur. It houses the RAM, ROM, other hardware and all the peripherals are plugged on to it. It also contains all the control circuitry essential for the working of the system.

The motherboard has slots or sockets for connecting various peripherals or support system/hardware. There is an Accelerated Graphics Port, which is used exclusively for video cards;

ATA adapter which provides the interfaces for the hard disk drives; slots for RAM cards; and Peripheral Component Interconnect (PCI) connectors, which provides electronic connections for sound cards, network cards and so on.

Figure 17.1 shows the picture of a motherboard released in 2003 (ASUS make). Now, in the year 2009, this is considered to be an old board. We will see a newer board later, but having a look at the older board will let us see the gradual changes that have happened in motherboard approaches, and also is helpful in understanding the functions of the various components that make up the motherboard.

In the figure, the processor is not in the socket. On the board, a number of electronic components and tracks can be seen. All these work together to facilitate the use of the motherboard as a hub which connects various components. The most important parts of the motherboard have been marked and labeled. Many of them are slots/sockets for cards that are plugged on to the motherboard and which reside inside the cabinet. As you read through the next sections, the functions of these parts will become clear.

Figure 17.2 shows the typical sockets for connecting devices external to the computer. They are placed on the side of the motherboard and are accessible at the back of the cabinet. In addition to these sockets, connectors and ports, the motherboard contains a number of other contacts. These include:

- i) *Jumpers*, which are used on some motherboards to configure voltage and various operating speeds.
- ii) A number of pins used to connect the reset button, like the LED for hard disk activity, built-in speaker.

These are not shown in the figures.




Figure 17.1 | A motherboard with important parts marked and labeled




Figure 17.2 | Connectors of the motherboard, accessible outside the PC

17.3 | Chipset

A word commonly used when talking about motherboards is ‘chipset’. What exactly is it? A chipset is a group of ICs working together to control the flow of data from/to and within the motherboard. We can call them ‘controllers’ of the communication done therein. In Chapter 6, we found that the 8086 needs various chips for buffering, de-multiplexing, generating control signals and so on. In Chapters 9, 10, and 11, more interfacing chips were discussed for handling parallel data, serial data, DMA and so on. Each of these functions call for the use of separate chips which have to be programmed individually. If all these functions could be done by a smaller set of dedicated controller chips, it would obviously be great – and that is the role taken up by the motherboard chipset.

To dig into a bit of history – in the IBM PC and PC-XT, a number of additional chips were used to realize the functions of the bus controller, clock generator, system timer, interrupt controllers, DMA controllers and so on. However, in 1986, after the 80286 based AT was launched, a company named Chips and Technology introduced a revolutionary concept by designing 82C206. This was a single chip that integrated into it all the functions of the main motherboard chips in an AT-compatible system. This chip included the functions of the 82284 Clock Generator, 82288 Bus Controller, 8254 System Timer, dual 8259 Interrupt Controllers, dual 8237 DMA Controllers, and even the MC146818 CMOS/Clock chip.

This meant that virtually many of the chips on a motherboard (except the RAM and ROM) could be replaced by a single chip. Four other chips augmented the 82C206 acting as buffers and memory controllers, thus the entire motherboard circuit was realized with a total of five chips. This first chipset was called the CS8220 chipset by Chips and Technologies.

This idea soon became very popular and many manufacturers brought out many different types of chipsets. An important point is that chipsets are to be designed ‘for’ the CPU being used. Thus, once a CPU is released, chipset manufacturers were to step in and design and sell the chipset – then only could the new chip be used in a motherboard. As Intel kept on releasing newer versions of its processors, it found that the delay in getting the processor on to a motherboard was caused by the delay in chipset manufacturing by other companies. Intel then

decided to manufacture chips and compatible chipsets in parallel, thus gaining a big margin in terms of time.

Now, Intel is the leading manufacturer of processors, and motherboards with compatible chipsets. Some other manufacturers of chipsets, for example, are VIA, NVIDIA, SiS and ASUS. Intel is a major manufacturer and there are other manufacturers in Taiwan, and slowly China is on its way to becoming a major hub for motherboard manufacturing.

17.3.1 | North Bridge and South Bridge

A chipset consists of two major microchips. These are known as the North Bridge and the South Bridge. The North Bridge handles data for the AGP (Accelerated Graphics Port) i.e., the adapter of the video card, and the main memory which includes the FSB (Front Side Bus-another name for the system bus) Although both chips are required for the PC to work, the North Bridge handles most of the very important tasks such as the connection between the CPU and main memory.

The North Bridge also handles all the faster traffic like the data transfer to/from and the AGP. On the motherboard, the North Bridge is situated close to the processor and is covered by a heat sink to take care of the large amount of heat it generates. In Fig 17.1, the placement of the North Bridge close to the processor slot is marked, but it is not seen as a heat sink covers it.

The South Bridge handles data from the PCI, USB, ATA and ISA (now obsolete) buses. The south bridge is a controller for I/O which can afford to be slower than the main system bus. In this set up, there are ports which connect the I/O to the motherboard by adapters which are plugged to the slots (sockets) on the motherboard. Fig 17.3 shows a schematic diagram of the bridges, and the peripherals each of them cater to.

The South Bridge is normally supplemented by a small Super I/O controller which takes care of a number of less critical functions that also used to be allotted to separate controller chips in the past. It used to be connected to the south bridge via the ISA bus; in modern architectures the LPC (Low Pin Count) interface is used. For example, most Super I/O chips contain controllers for the serial ports, parallel port, floppy drive, and keyboard/mouse interface. See Fig 17.4 and also refer Fig 17.1 that shows the super I/O chip, which is a controller for the serial port, parallel port, PS/2 mouse and key boards and a floppy drive. The speeds of these devices are relatively low.

Developments in recent years have led chipset manufacturers to attempt to place more and more functions in the chipset. These extra functions are typically:

- Video card (integrated into the north bridge)
- Sound card (in the south bridge)




Figure 17.3 | The north bridge and the south bridge




Figure 17.4 | South bridge supplemented by the super I/O chip

- Modem (in the south bridge)
- Network and Firewire (in the south bridge)

All these functions had traditionally been managed by separate devices, usually plug-in cards, which connect to the PC. However, it has been found that these functions can definitely be incorporated into the chipset itself. For example, one does not need to buy an extra sound card or even a separate video card, because they are already on the motherboard – we buy extra cards only to get better sound and video quality than that ensured in the chipset. Whether we purchase those extra cards or not depends on our need and also our budget.

17.3.2 | Intel Hub Architecture

The Intel Hub Architecture (IHA) has replaced the Northbridge/Southbridge chipset. The IHA chipset also has two parts, the Graphics and Memory Controller Hub (GMCH) and the I/O Controller Hub (ICH).

The GMCH interfaces between the high-speed processor bus and the hub interface and AGP bus, whereas the ICH interfaces between the hub interface and the ATA (IDE) ports, the SATA ports and the PCI bus. The ICH also includes a new low-pin-count (LPC) bus, consisting basically of a version of PCI designed primarily to support the motherboard ROM BIOS and Super I/O chips.

The IHA architecture started being used in Intel's 800 series chipsets, which is the first chipset architecture to move away from the Northbridge/Southbridge design. Figure 17.5 shows the IHA for Intel's 8xx chipset series.

17.4 | Transfer Speed

It should be convincing that certain components need to, and can transfer data at higher speeds than others. For data transfer from/to the RAM, the best situation will be when the RAM is as fast as the processor – that depends on the SDRAM available. For a 133MHz processor, a PC133 SDRAM corresponds to the same speed. However, now processor speeds have increased dynamically and this RAM will obviously not be good enough for a new and faster processor.




Figure 17.5 | Components of the IHA for a typical Intel motherboard (8xx series)

However, as discussed in Chapter 12, memory speed limitations are taken care of by the use of caches. We will see later that dual channel motherboard designs also attempt to improve data transfer speeds.

17.4.1 | Processor Bus

Also called the Front-Side Bus (FSB), this is the highest-speed bus in the system and is at the core of the chipset and motherboard. This bus is used primarily by the processor to pass information to and from cache or main memory and the North Bridge of the chipset. The processor bus in a modern system runs at 800MHz or higher and is normally 64 bits wide.

17.4.2 | AGP

Another data transfer that occurs through the North Bridge is to the AGP. What is AGP? It stands for Accelerated (Advanced) Graphics Port – it is a dedicated port that links the graphics controller card on a personal computer directly to the computer's memory.

The AGP bus is a high-speed 32-bit bus specifically designed for a video card. Since video and 3D graphics need high speed access to the system RAM, special video cards have been designed, with extra on-board RAM, a specialized graphics processor and other features. However, this card will be useless unless it has high speed access to the CPU and system RAM. This need made Intel design a special port for it, placed close to the processor connected to the North Bridge or Memory Controller Hub of the chipset and is manifested as a single AGP slot in systems that support it.

AGP is a high speed bus with a clock frequency of 66.66 MHz. In the basic mode, called AGP 1x, there is only one channel of transfer of 32 bits (4 bytes) making its bandwidth to be $66.66 \times 4 = 266$ MBps. There are AGP cards numbered as 2x, 4x and 8x – 2, 4 and 8 are the number of data transfer channels possible, where the data rate for 8x should obviously be 266×8 MBps. In Fig 17.1, an AGP slot is seen which supports x4 transfer rate. Note that it is the same size as the PCI socket, but is of a different color and not in the same line as the PCI sockets.

The reason for having a special AGP port in the north bridge is that graphics processing especially 3D graphics need high speed processing. To connect a dedicated graphics card to the PCI (we will see what it is, soon) meant that graphics processing would be subject to the speed limitation of the PCI bus. It was Intel's idea to have an AGP port in the north bridge, instead of data having to be fetched via the PCI expansion bus. AGP can handle at least twice the throughput of a standard PCI bus. AGP was developed by Intel, and AGP support was added to later versions of Microsoft Windows. However, this has also lost its relevance with the advent of the new PCI Express. In motherboards made after 2006 the newer and faster PCI Express (PCIe) graphics interface has more or less replaced AGP.

17.5 | Expansion Buses

The word 'expansion bus' means the I/O buses through which I/O devices are connected to the processor (at the South Bridge or ICH). This is used for slower peripherals as seen in the Figures 17.3–17.5. Let us find out what the features of these buses are and what changes have occurred over the years.

17.5.1 | ISA Bus

ISA stands for 'Industry Standards Architecture' and was around for a very long time, during which time there were modifications. It was followed by EISA (Extended ISA) – then there was the VESA bus, VL bus and also the Microchannel Architecture (MCA).

ISA which was an 8 MHz, 16-bit bus has disappeared from recent systems after first appearing in the original PC in 8-bit, 5 MHz form and in the 1984 IBM AT in the 16-bit, 8 MHz form. This bus was directly related to the processors for which it was created – 8088 and 80286. In current standards, it is a very slow bus – but was enough in the early days for the slow peripherals of the time.

In the search of a faster bus, Intel created the MCA (Microchannel Architecture) which was a 32-bit bus but could be used for 16-bit applications as well. It was an improvement over ISA, but did not quite make it in the desktop PC market.

Around this time, attempts to improve on ISA were made, and buses like EISA (32 bits at 8 MHz), VL-Bus (developed for video by VESA, the Video Electronics Standards Association) were tried, but none of them really made a great difference. The VL bus was okay for just one

device, but when more devices came into the picture, there was a chance of interfering with the CPU activities – this was its drawback.

17.5.2 | PCI Bus

PCI stands for Peripheral Component Interconnect and it has replaced the ISA bus on the motherboard. During the early 1990s, Intel introduced a new bus standard – the Peripheral Component Interconnect (PCI) bus. PCI presents a hybrid of sorts between ISA and VL-Bus. Since this is the currently used I/O bus (but is rapidly being replaced by the PCI Express), let us delve a bit deeper. Figure 17.1 shows three PCI sockets.

Systems that integrated the PCI bus became available from 1993 and have remained as such. Information typically is transferred through the PCI bus at 33MHz and 32 bits at a time. The bandwidth is 133MBps, as the following formula shows:

$$33.33 \text{ MHz} \times 4 \text{ bytes (32 bits)} = 133 \text{ MBps}$$

Although 32-bit, 33MHz PCI is the standard found in most PCs, there are now several variations. There is the 66MHz, 32-bit PCI and the 64-bit, 66MHz versions and even the 64-bit 133MHz PCI. The last version is used mostly in servers.




Figure 17.6 | Four PCI slots on a motherboard

What are the important and distinct features of the PCI bus?

- i) It supports 3.3 and 5 V signaling.
- ii) It is processor independent.
- iii) Plug and Play.

It can dynamically configure a system, avoiding resource conflicts. The specification for Plug and Play was developed by Microsoft and Intel among others, and the idea was to provide a system where one can simply install an adapter for a component and it should work. It is not quite as simple in practice; a software driver has to be installed before an adapter will work. However, the actual cooperation between adapter, motherboard and operating system – happens automatically. During startup, communication takes place between the PC's startup programs, the PCI controller and each PCI device.

Earlier, configuration issues were handled by jumpers on each add-on card. These jumpers would select memory or I/O space, interrupt vectors and DMA channels. All this required expertise and a fairly good knowledge of the associated hardware and software. If wrongly configured, the result would be a system with faults that are difficult to diagnose. The advent of the PCI, with its plug 'n play feature, has sorted out such problems.

- iv) PCI bus allows what is called 'bus mastering', which means that it allows any peripheral to take charge of the bus – this is most relevant for DMA between peripherals and main memory.

Fig 17.6 shows four PCI slots on a motherboard.

17.5.3 | PCI Express

All buses we have talked of till now were parallel buses – data of 8 bits, 16 bits, 32 bits or 64 bits moved in parallel, along with address and control signals – that made the number of pins on a bus connector quite large. Over the years, the trend of data transfer has moved from parallel to serial because of the problems associated with parallel data transfer, the most significant being what is called 'skew'. For example, when data is sent over 32 parallel lines, all the bit information do not arrive at the destination at the same time, and thus some lines are 'skewed' with respect to the others. This problem had always been there, but now that transfer speeds have had to be increased, the effect caused by this has compounded. Moreover, with more number of data lines, the problem of cross talk and electromagnetic interference between them has also intensified.

The solution for all this has been the conversion gradually from parallel to serial data standards – the evolution of the PCI Express and the USB are the highly successful examples of this change. PCI Express specifications were finalized in July 2002, and were initially called 'Third Generation I/O bus specifications', where ISA was the first generation and PCI was the second generation. One feature of this new specification was the backward compatibility with the existing PCI bus, because one cannot expect the existing standard to be wiped out overnight.

Note There is no PCIe (PCI-Express) slot in the motherboard shown in Fig 17.1.

In PCI-Express, data is sent full duplex (two separate one way paths) over two pairs of differentially signaled wires called a lane. Each lane caters to 250 MBps throughput in each direction. A multiplication factor of 1 to 2, 4, 8, 16, or 32 is possible, where these numbers mean the number of such lanes (channels). With 8 channels, the bandwidth is $250 \times 8 = 2000$ MBps one way. PCIe 16x has a data rate of 4000 MBps. Compare this with PCI which has a bandwidth of 133 MBps (one way) and needs many more pins by virtue of being a parallel port.

One obvious change in the motherboard is the smaller sized PCIe connectors, compared to the PCI connectors. Over the past few years, motherboards have been designed with PCI

as well as PCIe connectors existing side by side. One direct consequence of the advent of this new standard has been the disappearance of the AGP connectors from the motherboard. The current design of AGP cards is based on PCIe.

PCIe slots started appearing on motherboards from 2004 onwards. See Fig 17.7, which shows three PCIe slots side by side. Why are they of different lengths? Obviously, because of catering to a different number of lanes. The longest one in Fig 17.7 is a 16x connector for an advanced graphics card, requiring high speed transfer and thus more lanes. PCIe connectors are available in 1x, 4x, 8x and 16x configurations.

The sizes of connectors with different scaling factors of PCI Express, are different. The PCIe 1x connector has 36 signal pins, the 4x connector has 64 signal pins, the 8x connector has 98 signal pins, and the 16x connector has 164 signal pins. A PCI express card is upward compatible, so a 1x card will fit in any card slot, a 4x card will fit into an 8 or 16x port and so on. An adaptor card using 16x lanes will only fit in an x16 size connector, obviously.

One important feature of this new bus is that it is 'hot pluggable' and hot swappable. What do these words mean? Hot swapping describes changing components without significant interruption to the system, while hot plugging describes changing or adding components which interact with the operating system. One does not need to switch off the computer while doing this – this is similar to the USB port, for example.

17.5.4 | USB

This stands for Universal Serial Bus and is a very important interface for users – all of us have experienced the ease of plugging in different types of devices to the PC through the USB port. The USB port is available on the cabinet (it can be placed on the front or back) – and is hot pluggable and hot swappable – these features make it very useful for us users – now the trend is to shift a lot of the peripherals to the USB port – thus we have printers, mice, keyboards, scanners, cameras, external hard disk and so on, all of which are interfaced to the PC through the USB port.

Let us start with its history. The USB 1.0 specification was introduced in 1996. The original USB 1.0 specification had a data transfer rate of 12 Mbit/s. USB was created by a core group of companies that consisted of *Compaq*, *Digital*, *IBM*, *Intel*, *Northern Telecom*, and *Microsoft*. One of the co-inventors of USB was *Ajay Bhatt*, who was later given credit by Intel. The USB 2.0 specification was released in April 2000 and was standardized at the end of 2001, with a data rate of 480 Mbit/s.




Figure 17.7 | Three PCIe sockets on a motherboard

The USB 3.0 specification was released on November 12, 2008 by the USB 3.0 Promoter Group. Its maximum transfer rate is up to 10 times faster than the USB 2.0 release. It has been dubbed the 'Super Speed USB'. Intel says that the speed boost is made possible by using fiber-optics in addition to the regular wires. The USB 3.0 port will be backward compatible with current USB ones. Which are the equipments that need this high signaling rate? Current hard-drives, mice and printers are probably too speed limited to make use of the new 4.8 Gigabit/sec bandwidth, but we could use the bandwidth to transport high definition video signals. Very soon, we can expect innovative applications for this new standard which is just beginning to spread its wings.

17.5.4.1 | Features of USB

The Universal Serial Bus is host controlled and there can be only one host per bus. A USB system consists of a host controller and multiple devices connected in a tree-like fashion using special hub devices. Hubs may be cascaded, up to 5 levels. Up to 127 devices may be connected to a single host controller. To the user, this means that up to 127 devices can be connected to any one USB bus at any one given time. What if more devices need to be used? Simply add another port/host.

USB, as its name suggests is a serial bus. It uses 4 shielded wires of which two are power (+ 5 V and GND). The remaining two are twisted pair differential data signals. It uses a NRZI (Non Return to Zero Invert) encoding scheme to send data with a sync field to synchronize the host and receiver clocks. USB supports plug and play with dynamically loadable and unloadable drivers. This means that the user simply needs to plug the device on the bus. The host will detect this addition, interrogate the newly inserted device and load the appropriate driver, provided a driver is installed for the plugged-in device. The end user need not worry about terms such as IRQs and port addresses, or rebooting the computer. Once the use is over, the user can simply plug the cable out – the host will detect its absence and automatically unload the driver.

All USB peripherals are slaves that obey a defined protocol. They must react to request transactions sent from the host PC. The peripheral responds to control transactions that, for example, request detailed information about the device and its configuration. The peripheral sends and receives data to/from the host using a standard USB data format. This standardized data movement to/from the PC host and interpretation by the peripheral, gives USB its enormous flexibility with little PC host software changes. Now, the USB port is made available in various other devices like dedicated music players – thus it has made itself ubiquitous (meaning that it is like being omnipresent).

17.5.4.2 | USB – Power Issues

The USB connector provides a single 5 volt wire from which connected USB devices may power them. A given segment of the bus is specified to deliver up to 500 mA. This is often enough to power several devices, although this budget must be shared among all devices downstream of an unpowered hub. A bus-powered device may use as much of that power as allowed by the port it is plugged into. When USB devices (including hubs) are first connected they are interrogated by the host controller, which enquires of each, their maximum power requirement. Devices that need more than 500 mA current can use an external power source – thus we see printers and certain external hard disks with external power supply, while USB based mice and keyboards do not need any extra power.

17.6 | ATA

Now, let us go back and look at the South Bridge/IO controller hub in Fig 17.1. There is a word ATA – what does it stand for and what does it connect to? The word ATA stands for AT Attachment where AT stands for the AT bus (the ISA bus) which was found in the PC-AT. This designation refers to the fact that this interface was originally designed to connect a combined drive and controller directly to the 16-bit bus found in the PC-AT type computers. This type of interface is also called IDE or Integrated Drive Electronics which is a marketing term used by some companies to imply that the controllers of the hard disk drive have been integrated onto the drive itself, instead of having it as a separate add-on card plugged on to the ISA bus (as was done earlier). ATA is mostly thought of as the interface to the hard disk, though it also refers to the drive of the optical drives (CD/DVD) as well.

17.6.1 | Hard Disk

A hard disk is a sealed unit containing a number of platters in a stack. Hard disks may be mounted in a horizontal or a vertical position. Let us assume a horizontal hard disk. Electro-magnetic read/write heads are positioned above and below each platter. As the platters spin, the drive heads move in toward the center surface and out toward the edge. In this way, the drive heads can reach the entire surface of each platter.

On a hard disk, data is stored in thin, concentric bands. A drive head, while in one position can read or write a circular ring or band called a track. There can be more than a thousand tracks on a 3.5-inch hard disk. Sections within each track are called sectors. A sector is the smallest physical storage unit on a disk, and is almost always 512 bytes in size (Fig 17.8).

Each hard disk plate is divided into tracks. Each track is subdivided into a number of sectors, the disk's smallest unit. A sector normally holds 512 bytes of data. The individual files are written across a number of disk sectors (at least one), and this task is handled by the file system which is part of the operating system. Over the years, the capacity of hard disk has increased, due to two factors – the magnetic plates have become denser and read/write heads have become faster. Together this has led to faster hard disks with very high capacity.

17.6.2 | Hard Disk Interface

The interface to the hard disk has two sets of controllers and one interconnecting cable. See the PCB under the hard disk in Fig 17.9. Part of the drive electronics is here, and the other part is integrated into the South bridge as the ATA controller (Fig 17.3).

The ATA interface can be seen as a bus, which is managed by a host controller. Up to four devices can be connected per host controller. The unusual thing about the ATA interface is that there are two channels, which can each have two devices connected. They are called the primary and secondary channels. If two devices are connected to one channel, they have to be configured as one master and one slave device. These four channels are standard on today's motherboards using parallel ATA. Motherboards typically have connectors for two ATA cables, which can each connect two devices (master and slave). With all four channels being used, it could be two hard disks, one CD/DVD reader and one R/W DVD drive.

Many versions of ATA have appeared over the years. The most recent data rates are ATA/66 at 66 MB/second, ATA/100 at 100 MB/second and ATA/133 at 133 MB/second. Up to the ATA/66 standard, the connecting cable had only 40 conductors, but newer ones have 80 conductors,




Figure 17.8 | Tracks and sectors of a hard disk



Figure 17.9 | PCB under the hard disk with a controller and associated drive electronics




Figure 17.10 | The ATA system's four channels.




Figure 17.11 | Two ATA connectors on a motherboard




Figure 17.12 | Four SATA connectors on a motherboard

with each line having an individual ground connection, thus eliminating noise and cross talk and improving bandwidth. Fig 17.11 shows two ATA connectors on a motherboard.

After ATA/133, serial ATA (SATA) stepped in – this was designed to be backward compatible with parallel ATA (also called PATA), which only means that Serial ATA is software compatible with parallel ATA – it fully emulates all the commands, registers, and controls, so that existing software will run on the new architecture without any changes. In other words, the existing BIOS, operating systems, and utilities that work on parallel ATA also work on Serial ATA. However, SATA is a serial standard; it is an interface using cables with only 7 wires instead of the 40/80-wired cables used in the Parallel ATA-interface.

Optical drives are also becoming more readily available with SATA connections. SATA data rates are 150MBps for SATA-150, 300MBps for SATA-300 and 600MBps for SATA-600. See the smaller sized SATA connectors (four) in Fig 17.12.

17.7 | Memory – SIMM and DIMM

The theoretical aspect of semiconductor memory was discussed in Chapter 12. Here, we will see how these theoretical aspects matter, in the world of personal computers. The amount of physical memory that can be used for a Pentium Pro chip is 2^{36} bytes, where 36 is the number of address lines it possesses. This comes to 64 GB of physical memory space, which is a very big number. Having this much of RAM is prohibitively expensive, and takes a number of DDR chips and a number of sockets to put them in. The maximum capacity of current DDRs is only 2/4 GB. In practice, the RAM that most motherboards support is less than the physically usable amount. We know that the concept of virtual memory enables the user to feel that he has memory in excess of what is physically available.

RAM in the early days was manufactured as a DIP chip. Now, RAMs are in the form of modules, which are cards with the RAM chips embedded on it and which can be plugged into sockets made for this purpose. These are printed circuit boards with one side devoted to placing the memory chips. The earlier most popular module was called SIMM (Single In Line Memory Module). There used to be SIMMs for 8-bit and 16-bit memories for the earlier generation PCs. There are two types of SIMM modules, according to the number of connectors: SIMM modules with 30 connectors are 8-bit memories with which first-generation PCs were equipped ('286, '386), SIMM modules with 72 connectors are memories able to store 32 bits of data simultaneously. These memories are found on PCs from the 386DX to the first Pentiums. On Pentium which has a 64-bit data bus, two such SIMM modules are needed.

The current popular memory module is the DIMM (Dual In Line Memory Module) and its most important difference is that it is a 64-bit memory. Since Pentium and many other processors have a 64-bit bus width, two SIMMs installed in matched pairs, can be replaced by just one DIMM module. Physically, DIMMs differ from SIMMs in an important way. SIMMs have contacts on both side of the circuit board but they are tied together (they are called redundant connections). DIMMs also have contacts on both sides – however they have separate connections on each side of the circuit board. So a 168-pin DIMM has 84 pads on each side but they are not redundant. This allows the packaging to be made smaller, but makes DIMMs a bit more sensitive to correct insertion and good electrical contact.

DIMMs are available with ranges from 72 pins to 240 pins, which cater to different applications. For example, a small outline (SO) DIMM with 72 pins is used in laptops to support miniaturization. See the pictures of DIMM DDR module and DDR2 module in Figures 17.13 and 17.14 respectively.

RAM is rated with a number of the form x-y-y-y, which is a measure of the number of clock cycles required for access in the burst mode (Section 12.3). Standard 60ns asynchronous DRAM normally runs with 5-3-3-3 burst mode timing (Section 12.3). This means the first access takes a total of five cycles and the consecutive cycles take three cycles each. Without the burst technique, memory access would be 5-5-5-5 because the full latency is necessary for each memory transfer. This would take 20 cycles, while the burst modes need only 14 cycles.




Figure 17.13 | DIMM for DDR – with three notches on the sides and one notch along the contacts



Figure 17.14 | DIMM for DDR-2 with similar notches




Figure 17.15 | Four dual channel DIMM memory slots

SDRAM performance is still better. A typical SDRAM timing for a burst access would be 5-1-1-1, meaning that four memory reads would complete in only eight system bus cycles.

17.7.1 | Dual Channel Memory

This is a technique advertised to say that it gives double the speed for memory. The difference is in the motherboard, or rather the memory controller design, and not in the memory technology. In Fig 17.5 the word 'dual channel' SDRAM socket is found to be used. The diagram shows an Intel chipset where the **memory controller** supports dual channels for RAMs. In AMD processors and in some processors of Intel (Nehalem, for example), the memory controller is inside the chip.

In memory controllers that support dual channels, the data bus connection to the RAM is $64 \times 2 = 128$ lines, which are to be connected to two memory modules. The CPU can then access these 128 bits in parallel, which theoretically doubles the memory bandwidth and brings memory access speeds to be close to processor speeds. For example, many of the Intel 8xx family chipsets use dual-channel DDR memory. They also support the 800 MHz Pentium-4 processor system bus, which can transfer 8 bytes (64 bits) at a time for a bandwidth of $800 \times 8 = 6,400$ MBps. This is the speed at which the processor 'can' access memory. But, how to get memory to match this speed? With an 800 MHz FSB processor installed, these boards have memory controllers with 128 data lines which connect to two standard PC3200 memory modules, which operate in parallel. Thus the total bandwidth is now $3200 \times 2 = 6400$ Bps which thus can

match the speed of the processor. Practical results do not match with these calculated figures, but there is definitely an improvement. (There is also a triple channel architecture for some of Intel's latest processor chipsets).

To ensure the benefits of dual channeling, this facility must be enabled in the BIOS and, it is mandatory to have two memory modules in two sockets, installed correctly (refer the chipset manual). Thus, if 2 GB memory is to be used, it should be installed as two 1 GB modules (in sockets of the same color), rather than as a single 2 GB module. For help in installing correctly, the DIMM sockets are color coded in Intel chipsets. Fig 17.15 shows four such dual channel DIMM sockets. Other motherboards do not follow the same identification method.

17.8 | System BIOS

You have seen the term BIOS used many times throughout this book, and have also used a few of the BIOS routines for accessing PC hardware. So, an idea of what it is and what can be done with it, must be apparent. Here, let us see the relevance of BIOS for PCs and the changes that have taken place in the use of BIOS, over the years.

BIOS stands for **Basic Input Output System** and it is a collection of functions which can be called upon to access hardware directly. (The software which interacts with a particular hardware directly is called it 'driver'). These functions are activated by using software interrupt calls, and they are stored in system ROM which is an area in the physical address space of the PC. Table 7.2 shows the address allocation for BIOS to be 64K from F0000H to FFFFFH.

In Section 6.1.9, you have seen it mentioned that on startup, the first instruction will be taken from the location FFFF0H. This is for the 8086 processor. However, for all advanced processors beyond 8086, on startup, the processor wakes up in the real mode where only the lowest 1 MB of memory is effective. At the wakeup location there is a jump instruction (Section 8.4) to another part of memory which contains the bootstrap loader, which is a program for loading the OS from disk to system RAM. The ROM also contains the program for Power on Self Test (POST), which is a program which tests all hardware and memory before the system switches on.

Thus, we see that the drivers for all hardware are in system ROM as 'BIOS' routines. This idea of putting all drivers in system ROM was thought of and implemented very long back, during the initial period of PC development, and surprisingly this idea is still in vogue – but a few things have changed, obviously. We know that many years back, the number of external devices that could be interfaced to the PC were only a few – so, naturally, the device drivers of those could fit into the available ROM.

Now, things have changed – besides the lots of new hardware devices which need to be interfaced to the PC, there are various kinds and brands of each of these devices – like video powered by extra hardware which entered into the PC as an adapter card – there are many types of video cards with wide ranges of specifications made by different manufactures. Naturally, the drivers of all these cannot be pre-programmed into the system BIOS. The method of using these new adapter cards is that of having on that card, a ROM containing the necessary device drivers. The system ROM is preprogrammed to scan a predetermined area of memory looking for any adapter card ROMs and, if any are found, their code is tested and subsequently executed, essentially linking them into and adding their functionality to the existing BIOS. Thus, the system ROM acquired the drivers of these adapter cards, thus expanding the functionality of the system BIOS. This idea is important, especially for hardware like video adapters, because the display will not work unless the display adapter is first installed. For the not so essential

(for starting up) devices, the drivers are placed in the OS files which can be loaded into the system RAM, and also linked to the drivers in the system BIOS ROM.

Over time, things have kept changing. The trend has turned to manufacturers of new hardware and adapters writing their own device drivers and loading it to RAM during startup. Since the newer operating systems are either 32-bit/64-bit, the device drivers are also similar. These new device drivers are loaded into RAM just after startup. Thus, the functionality of the original 16-bit drivers in system BIOS is replaced by the new drivers, and the ROM BIOS exists only to get the system started, to initialize specific hardware, to offer security and to perform some basic initial configuration. However, after the OS is loaded, a whole new set of drivers which are now in RAM takes over. However, we know that we can use the ROM BIOS in the DOS mode, because DOS is a 16-bit operating system.

17.8.1 | ROM BIOS and Motherboards

BIOS routines are those which interact directly with the hardware – so they have to be designed ‘for the motherboard’. The motherboard BIOS usually includes drivers for all the basic system components, including the keyboard, floppy drive, hard drive, serial and parallel ports, and more. Because the drivers in the BIOS communicate with the hardware, the BIOS must be specific to the hardware and match it completely.

As mentioned earlier, many devices can have their drivers in the hard disk along with the OS, and they can be loaded into RAM during boot time. However, this cannot apply to essential devices like the display for example – thus their minimum drivers should be available in the system ROM as BIOS functions. Since the BIOS has to be compatible with the motherboard hardware, it is designed for specific hardware. The practice is that BIOS designers supply the BIOS according to the specifications put up by the motherboard manufacturers. The most popular BIOS sellers are AMI (American Megatrends Inc) and Phoenix. Earlier the ‘AWARD’ BIOS chip was seen commonly on PCs – later Phoenix acquired Award and still later, they stopped selling BIOS in the name of Award. These are the software vendors for Intel motherboards, but there are other players in the market also.

17.8.2 | Flash BIOS

From the last ten years or so, BIOS is in flash ROM and the advantage is that since flash ROM is re-writable while on the circuit board itself, its content can be changed. This makes sense if BIOS upgrades are available. The term ‘flashing the BIOS’ implies the changing of the contents of the flash ROM and is done to replace the original contents with a upgraded version of the same BIOS, if the user thinks the new version has some useful features for him. See the flash BIOS chip in Fig 17.1.

17.8.3 | CMOS NVRAM

We see that the system time is correct, once it is set – in spite of the fact that we may switch off the computer frequently. This implies that there is some ROM which keeps count of time continuously. In reality, there is no such ROM, but a real time clock (RTC) is kept continually running inside an NVRAM (non-volatile RAM). This is a RAM made of CMOS which has a battery back up, making its contents non-volatile.

In the original IBM AT system, a Motorola 146818 chip was used as the real-time clock (RTC) and CMOS RAM chip. This special chip had a simple digital clock that used 14 bytes of RAM and 50 bytes remained in which some system configuration settings were stored.

Modern PC systems do not use the Motorola chip anymore – instead, they incorporate the functions of this chip into the motherboard chipset (South Bridge) or Super I/O chip, or they use a special battery and NVRAM module. Because of the low power capability of CMOS, the battery life lasts up to ten years. The memory and *real-time clock* are generally powered by a CR2032 *lithium coin cell* (See Fig 17.1).

Many newer systems have CMOS RAM with 2KB to 4KB capacity or more. The extra room is used to store the Plug and Play information detailing the configuration of adapter cards and other options in the system. As such, no 100% compatible standards exists for how CMOS information is stored in all systems, but all the information stored herein is considered as a part of ROM BIOS itself, and thus this again means that it is directly related to the motherboard hardware.

17.9 | New Motherboards

Thus we have discussed most of the important parts of a motherboard. Now, let us look at a motherboard that is more recent than that shown in Fig 17.1. This is an Intel motherboard, released in the year 2006, and was quite an advanced board of that year. It has four PCI sockets and three PCIe sockets – the graphics card is to be plugged into the long PCIe slot. It also has two parallel ATA connectors and four SATA connectors. There is also a connector for the floppy drive, which is like ATA but with a different connector. The processor is usually covered by a fan and heat sink – this has been removed to show the processor (Fig 17.16). The ATX connector is the connection to the power supply for all the motherboard components.

In the figure, the processor is seen as a square chip placed in its socket. In the working condition, the processor is covered with a heat sink and a fan to dissipate the heat produced




Figure 17.16 | Intel motherboard showing all the important sockets and connectors




Figure 17.17 | Another view of the same motherboard with the processor covered, and a DIMM (memory module) inserted




Figure 17.18 | Connectors seen at the back of a new PC

by it. These are seen in Fig 17.17. In this figure, a DDR module is shown in one of the DIMM sockets.

Now, it is up to the interested reader to find out how a very new motherboard looks like. The most observable change will be that the number of PCI slots will be less in number—there may be one PCI slot and many PCIe slots. Another important change will be the connectors at the back of the cabinet – see Fig 17.18. There is no parallel port or serial port; as mentioned in previous chapters, they are on the way to becoming obsolete – but they are sometimes retained as legacy ports or for special appliances. For instance, in Chapter 9, we used a trainer kit to which the binary code of our program was downloaded from the PC using a serial port. This means that a serial port is to be specified when ordering the PC for using along with the trainer kit. Another option would be to use a USB to serial port converter.

The other major change is that there are a number of USB ports – why? The USB has become a very popular bus and many different devices can be connected through this – thus mice, keyboards, printers, modems and memory sticks can be connected to this very versatile port called the USB port. This board, shown in Fig 17.18, has the PS/2 mouse and keyboard connectors, but many new PCs are equipped only with USB ports for these two peripherals too.

There are two video connectors – one is the ordinary VGA port, to connect to an analog monitor – the other is the DVI (Digital Video In) port which gives a digital output to connect to a digital monitor. There is also a Firewire port. What is a firewire port? It is a serial port similar to the USB. FireWire is a method of transferring information between digital devices, especially audio and video equipment. It is also known as IEEE 1394 port. The FireWire port has been in common use since 1995, when Apple Inc. first began to include the port on a number of digital camcorders. IEEE 1394 was adopted as the High-Definition Audio-Video Network Alliance (HANA) standard connection interface for A/V (audio/visual) component communication and control. Many types of equipment like digital cameras and camcorders used to have firewire ports, which have now been replaced by USB ports. However, premium and professional equipment for audio and video still have these ports. FireWire is fast – the latest version achieves speeds up to 800Mbps. At some time in the future, that number is expected to jump to an unbelievable 3.2 Gbps.

17.10 | Other I/O Devices

We have discussed most of the important components of a general-purpose desk top PC. What is left are devices like keyboards and mice. We will not go onto them in detail, but make a general statement that each of them requires controller hardware which is taken care of by the south bridge. Recently, their connectors have also become USB. However, still PS/2 versions of these input devices exist.

Many items like network cards, optical drives, modems and Ethernet connections have not been discussed here. The interested reader is advised to read and find out about these and also the relevance of terms like SCSI, RAID and CNR (it is shown in Fig 17.5).

17.11 | PS/2

What is PS/2? This stands for IBM Personal System/2, a version of the PC released by IBM in 1987. There is a bit of history associated with this. It was IBM who manufactured the first PC and popularized it as well. However, soon, many other manufacturers came forward to make IBM-PC clones. Feeling that it should establish its grip over the PC market, IBM released a proprietary model calling it IBM PS/2. It had a lot of new and good features, like the MCA bus, a new version of SIMM and so on. However, because of the higher cost of the proprietary architecture, it could not hold the market and was soon considered a commercial failure. However, many of the standards used in this model were taken up and remained in the PC world. It was in the PS/2 model that the small 3.5" floppy disk was introduced and also the 72 pin SIMM which became a standard by itself. Also the PS/2 connectors of the keyboard and mouse are still around, although they are gradually being replaced by USB versions. The VGA standard for video was introduced in the PS/2 and also became a standard. To conclude, let us say that the PS/2 was a failure for IBM but did make contributions to the PC industry.

17.12 | Form Factors

There is a term called form factor frequently used in the PC world. This refers to the physical dimension of the components in a PC, which dictate the physical size of the cabinet, the

positioning of connectors, screw holes and so on. The biggest component (and also the most important) in a PC is the motherboard – its size, position, angle at which it is placed are important. It is based on this that the cooling system, the plug in cards and the position of the power supply can be decided. Because of all these, the form factor is a standard and over the years, standards have changed – the older motherboards were made for the AT cabinet. From 1996, the ATX (AT Extended) cabinet is the most popular standard and there are versions of it like mini and micro-ATX. The ATX standard has changed a bit from the year it was introduced to now. See the **ATX connectors** in Fig 17.1 and Fig 17.15. They are the connectors for the power supply required for the motherboard components (there are other power supply connectors for the hard disk and optical drives). Earlier, ATX connectors had 20 pins, but the newer connectors have 24 connections to account for some extra voltage dimensions.

An ATX power supply does not directly connect to the system power button, thereby allowing the computer to be turned off via software. We all have seen the sequence of how a computer shuts down with the OS initiating and controlling it. However, all ATX power supplies have a manual switch on the back to ensure the computer is truly OFF and no power is being sent to the components. If this switch is ON, power still flows to the components even when the computer appears to be off. In 2003, Intel proposed the BTX (Balanced Technology Extended) as a successor to ATX. Such cabinets are getting to be available – actually, the change from ATX to BTX has been rather slow. The BTX is designed to eventually replace the ATX form factor. One of the main focus points is on cooling requirements within the cabinet, because of the ever-increasing power requirements of the motherboard components.

With processors exceeding 1000W in thermal output, as well as voltage regulators, motherboard chipsets, and video cards adding to the thermal load in a system, BTX has been designed to allow all the high-heat-producing core components to be mounted inline from front to back, so that a single high efficiency thermal module (heatsink) can cool the system. So the idea of having more and more fans need not be adhered to. The motherboard is placed in such a way that air flow through the cabinet is easier. The case itself is modified so that front and rear vents are available to send the required airflow across the motherboard and out through the rear of the computer. The rest of the motherboard components have also been shifted to accommodate all this positioning.

Effectively, the BTX form-factor attempts to increase cabinet cooling so that it can provide a steady flow of fresh air that modern high power components need, while minimizing the amount of noise it creates. The motherboard redesign has to allow this to happen with a minimum of active cooling (like, using fans), with a single ‘thermal module’ both cooling the processor and creating the stream of air that will cool the other components. Besides ATX and BTX, there are smaller sized cabinets with these standards – they are the mini ATX, micro ATX, mini BTX and so on.

These are standards for general motherboard and cabinet design – but PC manufacturing companies do not generally adhere to these standards. They have their own proprietary designs – think of Dell, HP, Compaq and their computers. The problem of such proprietary designs is that one has to go back to them for upgrading, repairing and so on. One cannot just take an ‘off the shelf’ component and use it in these proprietary designs.

17.13 | Laptops

The PCs that we have discussed so far are called desktops, as they are stationed on desks and designed to be stationary – so weight and dimensions are not much of an issue. The use of

laptops (also called notebooks), rather than desktops is on the upswing and possessing a laptop is getting to be a fad and a craze. So, let us have a look at a laptop and its features.

A laptop is a portable PC – which means that two of its important features should be compactness and low weight. Thus, in laptops, many sacrifices have to be made resulting in less performance – but to overcome these limitations, special low power processors, small heat sinks, special boards, small cabinets and such things are being thought of, researched on, and are being developed. The probability is that laptops have all their functions integrated into the motherboard, including the controllers and processors for video, sound, and so on. The possibility of adding extra cards is almost zero. The hard disk capacity and RAM are also likely to be less. When buying a laptop, the specifications are to be decided then and there – because the possibility of upgrading is unlikely. At the most, an extra memory slot is available – but that may be all.

There is also a mini laptop that is now available, which is called a ‘netbook’. It is not very different from a laptop, but its specifications are lower – which means that its computing power is much less, but it is sufficient for ordinary applications like web browsing and office computing – its main features are low weight and small size – screens are small (9" to 10"). It can be carried comfortably while traveling, but while at home (or office), external screens and keyboards can be connected to it for ease of using. Optical drives may also be external, connectable through the USB port.

At the other end of the spectrum, there are high performance laptops with special models designed with extra and excellent graphic features added. For example ‘Alienware’ is a high end gaming laptop which has extraordinary graphic capabilities, which is also useful for video editing and similar applications. The company was established in 1996, and was acquired by Dell in 2006. However, it has the status of a subsidiary of Dell, maintaining its autonomy while using Dell’s market tactics and economies of scale, serving to reduce its cost factor – but it is still quite an expensive laptop. On the whole, a laptop is a good idea, but for the same price, a higher end model of a desktop can be bought.

17.13.1 | Intel’s Centrino Technology

Many of you when confronted with the prospect of buying a laptop, would have heard the word ‘Centrino’ referring to a processor of Intel which is used in laptops. However, actually, the word Centrino does not refer to an Intel processor – it refers to a system which Intel has manufactured for mobile applications – this means that Intel has designed a motherboard with a specific processor and chipsets which when used in a laptop makes it a ‘Centrino’ system. Because of the high amount of advertising by Intel, and reasonably good performance, the Centrino label has definitely made an impact in the laptop market.

To be considered a Centrino board, it should have the following components: the Intel Pentium – M processor, an Intel 855 chipset and an Intel 802.11 wireless network module. It should also have an Intel PRO/wireless mini PCI card. The mini PCI card is about the size of a credit card, and is plugged on the motherboard through a mini PCI connector.

However, things are also moving in the direction in which newer versions of Pentium-M are being used in other vendors’ motherboards which give better or comparable performance. Another new development promised by Intel is a system on chip (SOC) for a laptop, based on the Atom processor, which is a ‘small’ and relatively cheap processor.

Laptops and Desktops While Fig 17.19 shows two laptops with screen sizes of 12" and 9" – understand that screen sizes are measured ‘diagonally’ – Fig 17.20 shows a desktop with a tall, proprietary cabinet of Dell along with a keyboard and monitor.



Figure 17.19 | Laptops with screen sizes of 12" and 9"



Figure 17.20 | A desktop computer

KEY POINTS OF THIS CHAPTER

- The most important part of a PC is its motherboard.
- The controllers on the motherboard are called the chipsets.
- The chipsets consist of two important chips which are the North Bridge and the South Bridge.
- The South Bridge can have an additional chip called the Super I/O chip.
- Intel has changed its chipset architecture to what is called the 'Intel Hub Architecture'.
- There has been a number of expansion buses used in PCs and one such bus is the ISA bus which is obsolete now.
- The more recent bus is the PCI bus which is also gradually giving way to the PCI-Express.

- PCI Express and USB are serial buses unlike the parallel buses of earlier times.
- The controller for the hard disk drives and optical drives is integrated onto the drive itself.
- The adapter for the above drives is called ATA.
- Parallel ATA is gradually being replaced by serial ATA (SATA).
- Memory comes in the form of modules and the current popular module is called DIMM.
- PS/2 version of personal computers was a commercial disaster for IBM but the standards and technology it introduced have enriched the PC world.
- The word 'form factor' refers to the physical dimensions of the important components of a system. The form factor corresponding to BTX is the current market trend.
- Laptops are a portable form of the desktop PC, and many compromises are made in its design, for stressing on portability.

QUESTIONS

1. What are the functions of a PC motherboard?
2. What does the word 'chipset' mean?
3. Why is it that there are many chipset manufacturers for a particular processor?
4. Why should the Northbridge have a very high speed factor?
5. Which accessories are controlled by the LPC chip?
6. What is meant by the word 'memory wall'?
7. What limits the speed of the front side bus?
8. There used to be a term called 'back side bus'. What did it refer to?
9. Which are the buses considered to be the forerunners of the PCI bus?
10. In what way is the PCI Express different from the PCI bus?
11. What features have made the USB a very versatile bus?
12. What does the word IDE refer to?
13. What are the ways in which a DIMM module differs from a SIMM?
14. Why is it that dual channel memory is not as fast as it is advertised to be?
15. What is the role of system BIOS now?
16. What information does the NVRAM store?
17. What is the relevance of the word 'form factor'?

EXERCISE

1. Find out the components and features of the following cards:
 - a) network card
 - b) modem
 - c) sound card
 - d) graphics card
2. What is the state of the firewire port now?
3. Find out the history of the RDRAM memory, including how it has lost its relevance, in spite of having been a very promising technology.

APPENDIX A

8086

16-BIT HMOS MICROPROCESSOR

8086/8086-2/8086-1

- Direct Addressing Capability 1 MByte of Memory
- Architecture Designed for Powerful Assembly Language and Efficient High Level Languages
- 14 Word, by 16-Bit Register Set with Symmetrical Operations
- 24 Operand Addressing Modes
- Bit, Byte, Word, and Block Operations
- 8 and 16-Bit Signed and Unsigned Arithmetic in Binary or Decimal Including Multiply and Divide
- Range of Clock Rates:
5 MHz for 8086,
8 MHz for 8086-2,
10 MHz for 8086-1
- MULTIBUS System Compatible Interface
- Available in EXPRESS
— Standard Temperature Range
— Extended Temperature Range
- Available in 40-Lead Cerdip and Plastic Package (See Packaging Spec. Order #231369)

The Intel 8086 high performance 16-bit CPU is available in three clock rates: 5, 8 and 10 MHz. The CPU is implemented in N-Channel, depletion load, silicon gate technology (HMOS-III), and packaged in a 40-pin Cerdip or plastic package. The 8086 operates in both single processor and multiple processor configurations to achieve high performance levels.




Figure 1 | 8086 CPU block diagram





Figure 2 | 8086 pin configuration

MINIMUM MODE



(Continued)

READ CYCLE




(Continued)

WRITE CYCLE


8288 OUTPUTS
SEE NOTES 5,6

SOFTWARE HALT- (DEN = V_{OL} ; $\overline{RD}$, $\overline{MRDC}$, $\overline{IORC}$, $\overline{MWTC}$, $\overline{AMWC}$, $\overline{IOWC}$, $\overline{AIOWC}$, $\overline{INTA}$, $=V_{OH}$)
AD₁₅-AD₀




1. All signals switch between VOH and VOL unless otherwise specified.
2. RDY is sampled near the end of T_2 , T_3 , T_W to determine if TW machines states are to be inserted.
3. Cascade address is valid between first and second INTA cycle.
4. Two INTA cycles run back-to-back. The 8086 LOCAL ADDR/DATA BUS is floating during both INTA cycles. Control for pointer address is shown for second INTA cycle.
5. Signals at 8284A or 8288 are shown for reference only.
6. The issuance of the 8288 command and control signals (MRDC, MWTC, AMWC, IORC, IOWC, AIOWC, INTA and DEN) lags the active high 8288 CEN.
7. All timing measurements are made at 1.5V unless otherwise noted.
8. Status inactive in state just prior to T_4 .

ASYNCHRONOUS SIGNAL RECOGNITION




Note: Setup requirements for asynchronous signals only to guarantee recognition at next CLK.


BUS LOCK SIGNAL TIMING (MAXIMUM MODE ONLY)



RESET TIMING



REQUEST / GRANT SEQUENCE TIMING (MAXIMUM MODE ONLY)



Note: The coprocessor may not drive the buses outside the region shown without risking contention.

HOLD/HOLD ACKNOWLEDGE TIMING (MINIMUM MODE ONLY)




Table 1 | Instruction Set Summary

Mnemonic and Description	Instruction Code			
	76543210	76543210	76543210	76543210
DATA TRANSFER				
MOV = Move:				
Register/Memory to/from Register	100010 dw	mod reg r/m		
Immediate to Register/Memory	1100011 w	mod 000 r/m	data	data if w=1
Immediate to Register	1011 w reg	data	data if w=1	
Memory to Accumulator	1010000 w	addr-low	addr-high	
Accumulator to Memory	1010001 w	addr-low	addr-high	
Register/Memory to Segment Register	10001110	mod 0 reg r/m		
Segment Register to Register/Memory	10001100	mod 0 reg r/m		
PUSH = Push:				
Register/Memory	11111111	mod 110 r/m		
Register	01010 reg			
Segment Register	000 reg 110			
POP = Pop:				
Register/Memory	10001111	mod 000 r/m		
Register	01011 reg			
Segment Register	000 reg 111			
XCHG = Exchange:				
Register/Memory with Register	1000011 w	mod regr/m		
Register with Accumulator	10010 reg			
IN = Input from:				
Fixed Port	1110010 w	port		
Variable Port	1110110 w			
OUT = Output to:				
Fixed Port	1110011 w	port		
Variable Port	1110111 w			
XLAT = Translate Byte to AL	11010111			
LEA = Load EA to Register	10001101	mod reg r/m		
LDS = Load Pointer to DS	11000101	mod reg r/m		
LES = Load Pointer to ES	11000100	mod reg r/m		
LAHF = Load AH with Flags	10011111			
SAHF = Store AH into Flags	10011110			
PUSHIF = Push Flags	10011100			
POPF = Pop Flags	10011101			

(Continued)

Table 1 (Continued)

Mnemonic and Description	Instruction Code			
	76543210	76543210	76543210	76543210
ARITHMETIC				
ADD = Add:				
Reg./Memory with Register to Either	000000 dw	mod reg r/m		
Immediate to Register/Memory	100000 sw	mod 000 r/m	data	data if s: w=01
Immediate to Accumulator	0000010 w	data	data if w=1	
ADC = Add with Carry:				
Reg./Memory with Register to Either	000100 dw	mod reg r/m		
Immediate to Register/Memory	100000 sw	mod 010 r/m	data	data if s: w=01
Immediate to Accumulator	0001010 w	data	data if w=1	
INC = Increment:				
Register/Memory	1111111 w	mod 000 r/m		
Register	01000 reg			
AAA = ASCII Adjust for Add	00110111			
BAA = Decimal Adjust for Add	00100111			
SUB = Subtract:				
Reg./Memory and Register to either	001010 dw	mod reg r/m		
Immediate from Register/Memory	100000 sw	mod 101 r/m	data	data if sw=01
Immediate from Accumulator	0010110 w	data	data if w=1	
SSB = Subtract with Borrow				
Reg./Memory and Register to Either	000110 dw	mod reg r/m		
Immediate from Register/Memory	100000 sw	mod 011 r/m	data	data if sw=01
Immediate from Accumulator	000111 w	data	data if w=1	
DEC = Decrement:				
Register/Memory	1111111 w	mod 001 r/m		
Register	01001 reg			
NEG = Change sign	1111011 w	mod 011 r/m		
CMP = Compare:				
Register/Memory and Register	001110 dw	mod reg r/m		
Immediate with Register/Memory	100000 sw	mod 111 r/m	data	data if sw=01
Immediate with Accumulator	0011110 w	data	data if w=1	
AAS = ASCII Adjust for Subtract	00111111			
DAS = Decimal Adjust for Subtract	00100111			

(Continued)

Table 1 (Continued)

Mnemonic and Description	Instruction Code			
	76543210	76543210	76543210	76543210
ARITHMETIC				
ADD = Add:				
MUL = Multiply (Unsigned)	1111011 w	mod 100 r/m		
IMUL = Integer Multiply (Signed)	1111011 w	mod 101 r/m		
AAM = ASCII Adjust for Multiply	11010100	00001010		
DIV = Divide (Unsigned)	1111011 w	mod 110 r/m		
IDIV = Integer Divide (Signed)	1111011 w	mod 111 r/m		
AAD = ASCII Adjust for Divide	11010101	00001010		
CBW = Convert Byte to Word	10011000			
CWD = Convert Word to Double Word	10011001			
LOGIC				
NOT = Invert	1111011 w	mod 010 r/m		
SHL/SAL = Shift Logical/Arithmetic Left	110100 vw	mod 100 r/m		
SHR = Shift Logical Right	110100 vw	mod 101 r/m		
SAR = Shift Arithmetic Right	110100 vw	mod 111 r/m		
ROL = Rotate Left	110100 vw	mod 000 r/m		
ROR = Rotate Right	110100 vw	mod 001 r/m		
RCL = Rotate Through Carry Flag Left	110100 vw	mod 010 r/m		
RCR = Rotate Through Carry Right	110100 vw	mod 011 r/m		
AND = And:				
Reg./Memory and Register to Either	001000 dw	mod reg r/m		
Immediate to Register/Memory	1000000 w	mod 100 r/m	data	data if w=1
Immediate to Accumulator	0010010 w	data	data if w=1	
TEST = And Function to Flags, No Result:				
Register/Memory and Register	1000010 w	mod reg r/m		
Immediate Data and Register/Memory	1111011 w	mod 000 r/m	data	data if w=1
Immediate Data and Accumulator	1010100 w	data	data if w=1	
OR = Or:				
Reg./Memory and Register to Either	000010 dw	mod reg r/m		
Immediate to Register/Memory	1000000 w	mod 001 r/m	data	data if w=1
Immediate to Accumulator	0000110 w	data	data if w=1	

(Continued)

Table 1 (Continued)

Mnemonic and Description	Instruction Code			
	76543210	76543210	76543210	76543210
XOR = Exclusive or:				
Reg./Memory and Register to Either	001100 dw	mod reg r/m		
Immediate to Register/Memory	1000000 w	mod 110 r/m	data	data if w=1
Immediate to Accumulator	0011010 w	data	data if w=1	
STRING MANIPULATION				
REP = Repeat	1111001 z			
MOVS = Move Byte/Word	1010010 w			
CMPS = Compare Byte/Word	1010011 w			
SCAS = Scan Byte/Word	1010111 w			
LODS = Local Byte/Wd to AL/AX	1010110 w			
STOS = Store Byte/Wd from AL/A	1010101 w			
CONTROL TRANSFER				
CALL = Call:				
Direct within Segment	11101000	dep-low	dap-high	
Indirect within Segment	11111111	mod 010 r/m		
Direct Intersegment	10011010	offset-low	offset-high	
		seg-low	seg-high	
Indirect Intersegment	11111111	mod 011 r/m		
JMP = Unconditional Jump:				
Direct within Segment	11101001	disp-low	disp-high	
Direct within Segment-Short	11101011	disp		
Indirect within Segment	11111111	mod 100 r/m		
Direct Intersegment	11101010	offset-low	offset-high	
		seg-low	seg-high	
Indirect Intersegment	11111111	mod 101 r/m		
RET = Return from CALL:				
Within Segment	11000011			
Within Segment Adding Immediate to SP	11000010	data-low	data-high	
Intersegment	11001011			
Intersegment Adding Immediate to SP	11001010	data-low	data-high	

(Continued)

Table 1 (Continued)

Mnemonic and Description	Instruction Code		
	76543210	76543210	76543210
JE/JZ = Jump on Equal/Zero	01110100	disp	
JL/JNGE = Jump on Less/Not Greater of Equal	01111100	disp	
JL/JNG = Jump on Less or Equal/ Nor Greater	01111110	disp	
JB/JNAE = Jump on Below/Not Above or Equal	01110010	disp	
JBE/JNA = Jump on Below or Equal/ Not Above	01110110	disp	
JP/JPE = Jump on Parity/Parity Even	01111010	disp	
JO = Jump on Overflow	01110000	disp	
JS = Jump on Sign	01111000	disp	
JNE/JNZ = Jump on Not Equal/Not Zero	01110101	disp	
JNL/JGE = Jump on Not Less/Greater of Equal	01111101	disp	
JNLE/JG = Jump on Not Less or Equal/ Greater	01111111	disp	
JNB/JAE = Jump on Not Below/Above of Equal	01110011	disp	
JNBE/JA = Jump on Not Below or Equal/ Above	01110111	disp	
JNP/JPO = Jump on Not Par/Par Odd	01111011	disp	
JNO = Jump on Not Overflow	01110001	disp	
JNS = Jump on Not Sign	01111001	disp	
LOOP = Lop CX Times	11100010	disp	
LOOPZ/LOOPE = Loop While Zero/Equal	11100001	disp	
LOOPNZ/LOOPNE = Loop While Not Zero/Equal	11100000	disp	
JCXZ = Jump on CX Zero	11100011	disp	
INT = Interrupt			
Type Specified	11001101	type	
Type 3	11001100		
INTO = Interrupt on Overflow	11001110		
IRET = Interrupt Return	11001111		

(Continued)

Table 1 (Continued)

Mnemonic and Description	Instruction Code	
	76543210	76543210
CLC = Clear Carry	11111000	
CMC = Complement Carry	11110101	
STC = Set Carry	11111001	
CLD = Clear Direction	11111100	
STD = Set Direction	11111101	
CLI = Clear Interrupt	11111010	
STI = Set Interrupt	11111011	
HLT = Halt	11110100	
WAIT = Wait	10011011	
ESC = Escape (to External Device)	11011 xxx	mod xxx r/m
LOCK = Bus Lock Prefix	11110000	

NOTES:

AL=8-bit accumulator
AX=16-bit accumulator
CX=Count register
DS=Data segment
ES=Extra segment
Above/below refers to unsigned value
Greater=more positive;
Less=less positive (more negative) signed values
if d = 1 then "to" reg; if d=0 then "from" reg
if w = 1 then word instruction; if w=0 then byte instruction
if mod=11 then r/m is treated as a REG field
if mod=00 then DISP=0*, disp-low and disp-high are absent
if mod=01 then DISP=disp-low sign-extended to 16 bits, disp-high is absent
if mod=10 then DISP=disp-high; disp-low
if r/m=000 then EA=(BX)+(SI)+DISP
if r/m=001 then EA=(BX)+(DI)+DISP
if r/m=010 then EA=(BP)+(SI)+DISP
if r/m=011 then EA=(BP)+(DI)+DISP
if r/m=100 then EA=(SI)+DISP
if r/m=101 then EA=(DI)+DISP
if r/m=110 then EA=(BP)+DISP*
if r/m=111 then EA=(BX)+DISP
DISP follows 2nd byte of instruction (before data if required)
*except if mod=00 and r/m=110 then EA=disp-high; disp-low.

if s w=01 then 16 bits of immediate data from the operand
if s w=11 then an immediate data byte is sign extended to from the 16-bit operand
if v=0 then "count" = 1; if v=1 then "count" in (CL)
x= don't care
z is used for string primitives for comparison with ZF FLAG

SEGMENT OVERRIDE PREFIX

001 reg 110

REG is assigned according to the following table:

16-Bit (W=1)	8-Bit (W=0)	Segment
000 AX	000 AL	00 ES
001 CX	001 CL	01 CS
010 DX	010 DL	10 SS
011 BX	011 BL	11 DS
100 SP	100 AH	
101 BP	101 CH	
110 SI	110 DH	
111 DI	111 BH	

Instructions which reference the flag register file as a 16-bit object use the symbol FLAGS to represent the file:
FLAGS = X:X:X:X:(OF):(DF):(IF):(TF):(SF):(ZF):X:(AF):X:(PF):X:(CF)

APPENDIX B

USE OF DOS AND DEBUG COMMANDS

MASM 6.14, is an assembler (basically for x86 architecture) that runs under DOS, which is not a GUI (graphical user interface) operating system. It is a command line interface – we have to give commands to run various tasks. Nowadays, all computers have some version of Windows running (unless we are ‘linux only’ users). In Windows, there is a ‘command line’ shell:

- i) Go to the Start menu and click on ‘Run’
- ii) A small window opens. Type cmd in the space and click ‘OK’
- iii) A small black window opens. (This can be enlarged by pressing ‘Alt’ and ‘Enter’ keys simultaneously).

In this window, various commands can be typed. Only a few of them will be mentioned here:

- i) cd – change directory .

To go to the root directory, the command is cd\

```
C:\Documents and Settings\admin>cd\  
C:\>
```

The current drive is the C drive now.

To go to another drive, say type ‘D’:

```
C:\>D:  
D:\>
```

The current drive is D.

Assume that the MASM6.14 assembler is in this drive. To go to this directory, type
cd MASM6.14

The result is

```
D:\>MASM6.14
```

To go to the BIN directory, use the ‘cd’ command again. Type **cd BIN**. Now, we see

```
D:\>MASM6.14\BIN
```

All commands to run the assembler and the debugger have to be used when in the BIN directory.

List of Some Important DOS Commands

Some of the commands that may be useful at some time or other:

edit – opens an editor in which the program can be written.

dir – lists the subdirectories and files in the current directory

dir/p – same as ‘dir’, however, lists one page at a time

dir directory name – lists the files and subdirectories of the named directory.

cls – clears the screen

del filename – deletes the file specified

exit – closes the command window

Debug Commands

The commands for using the MASM assembler have been discussed in Chapter 2. But some of the commands of the debugger were left out. We will discuss them here.

On typing 'debug' an underscore is obtained at the prompt. Type '?'.
 The list of the debug commands are displayed and are as shown below. (This list can be obtained in the C directory itself as it is the native debugger of x86. It can be obtained in any other directory as well.)

C:\>debug

-?

assemble	A [address]
compare	C range address
dump	D [range]
enter	E address [list]
fill	F range list
go	G [= address] [addresses]
hex	H value1 value2
input	I port
load	L [address] [drive] [firstsector] [number]
move	M range address
name	N [pathname] [arglist]
output	O port byte
proceed	P [= address] [number]
quit	Q
register	R [register]
search	S range list
trace	T [= address] [value]
unassemble	U [range]
write	W [address] [drive] [firstsector] [number]

Some of these commands will be necessary to debug the assembly programs that we write. We will discuss them for a small assembly program named `tinym.asm`.

After the program is assembled and linked, a com file named `tinym.com` is obtained as the executable file (the com file is obtained only because the tiny model has been used in the program. If the small model had been used, an exe file would be the result).

To debug this program, type **debug tinym.com**. The debug prompt (hyphen) is obtained. Now, the various debugger commands can be used.

Note All numbers used in the debugger are by default, in the hex format.

i) The Unassemble command 'u'

Format: u

C:\masm6.14\BIN>debug tinym.com

-u

```

13C6:0100 B067  MOV AL,67
13C6:0102 B345  MOV BL,45
13C6:0104 02C3  ADD AL,BL
13C6:0106 8AD0  MOV DL,AL
13C6:0108 B44C  MOV AH,4C

```

The unassemble command converts machine code in the executable file to assembly code, as shown above. The number 13C6 is the content of the CS register and 0100 is the offset. The number B067 corresponds to the machine code in hex format for the instruction MOV AL,67. Note that any com program starts at an offset of 0100.

ii) The Go command 'g'

Format: G [start address] break address

The start address is optional. Normally, only the break address is required.

For a com file, the program to be executed starts at location 0100.

Typing g0100 causes the execution of instructions up to that address.

Below it, the instruction MOV AL,67 is scheduled to be executed.

The content of all the registers at that time is shown.

-g0100

```
AX=0000 BX=0000 CX=000C DX=0000 SP=FFFE BP=0000 SI=0000 DI=0000
DS=13C6 ES=13C6 SS=13C6 CS=13C6 IP=0100 NV UP EI PL NZ NA PO NC
13C6:0100 B067  MOV AL,67
```

Now to execute from address 0100 to 0106

-g0100 0106

The result is as follows. All instructions up to the break address 0106 have been executed.

```
AX=00AC BX=0045 CX=000C DX=0000 SP=FFFE BP=0000 SI=0000 DI=0000
DS=13C6 ES=13C6 SS=13C6 CS=13C6 IP=0106 OV UP EI NG NZ NA PE NC
13C6:0106 8AD0  MOV DL,AL
```

iii) The Trace command 't'

This command causes the next instruction to be executed. The contents of registers and flag status are displayed.

-t

```
AX=0067 BX=0000 CX=000C DX=0000 SP=FFFE BP=0000 SI=0000 DI=0000
DS=13C6 ES=13C6 SS=13C6 CS=13C6 IP=0102 NV UP EI PL NZ NA PO NC
13C6:0102 B345  MOV BL,45
```

The part

NV UP EI PL NZ NA PO NC is the flag status currently. The interpretation of the flag status is explained in the last part of this appendix.

To execute more than one instruction use the command 't' followed by the number of instructions to be executed. For example, using t2 shows the execution of two instructions with corresponding register contents and flag status.

-t2

```
AX=0067 BX=0045 CX=000C DX=0000 SP=FFFE BP=0000 SI=0000 DI=0000
DS=13C6 ES=13C6 SS=13C6 CS=13C6 IP=0104 NV UP EI PL NZ NA PO NC
13C6:0104 02C3  ADD AL,BL
```

```
AX=00AC BX=0045 CX=000C DX=0000 SP=FFFE BP=0000 SI=0000 DI=0000
DS=13C6 ES=13C6 SS=13C6 CS=13C6 IP=0106 OV UP EI NG NZ NA PE NC
13C6:0106 8AD0  MOV DL,AL
```

iv) The Register command 'r'

This gives the contents of the registers at the particular time.

-r

```
AX=0067 BX=0000 CX=000C DX=0000 SP=FFFE BP=0000 SI=0000 DI=0000
DS=13C6 ES=13C6 SS=13C6 CS=13C6 IP=0102 NV UP EI PL NZ NA PO NC
13C6:0102 B345  MOV BL,45
```

v) The Dump or Display command 'd'

This command displays the contents of the memory portion specified.

By default, the content of the data segment is displayed. But the contents of other segments can be displayed by specifying the address in the form

base address: offset

The different formats of the 'd' command are as follows.

a) d start address [length] .

It is optional to specify the length. If not specified, contents of 128 byte locations will be displayed.

b) d start address end address

Examples

a)

-d0000

```
13C6:0000 CD 20 FF 9F 00 9A F0 FE-1D F0 4F 03 DA 0D 8A 03  ....O....
13C6:0010 DA 0D 17 03 DA 0D C9 0D-01 01 01 00 02 FF FF FF  ....
13C6:0020 FF FF FF FF FF FF FF FF-FF FF FF FF FF 87 13 F1 49  ....I
13C6:0030 DA 0D 14 00 18 00 C6 13-FF FF FF FF FF 00 00 00 00  ....
13C6:0040 05 00 00 00 00 00 00 00-00 00 00 00 00 00 00 00  ....
13C6:0050 CD 21 CB 00 00 00 00 00-00 00 00 00 00 20 20 20 20  !.....
13C6:0060 20 20 20 20 20 20 20 20-00 00 00 00 00 20 20 20  ....
13C6:0070 20 20 20 20 20 20 20 20-00 00 00 00 00 00 00 00  ....
```

-

In this case, the data in 128 locations from DS:0000 are displayed, 16 bytes on a line.

b) **-d**

In this case, the 128 bytes will be the contents of memory, beginning at the end of the last address displayed.

c) **-d 0000 0002**

In this case, the contents of only 3 bytes (of addresses 0000 to 0002) will be displayed.

d) **-d 0000 3**

In this case, the contents of 4 locations will be displayed, since any count starts from 0.

e) To display the contents of other segments, find the content of their segment registers and use it along with the d command in the **base address: offset** format.

vi) The Assemble command 'a'

This command is used to write program lines in the debugger itself, execute it, and view the result. Once the program is typed in, it can be assembled and run.

The format of this command is

-a[starting address]

If the starting address is not specified, the default address is 0100.

Let us type in a few instructions after specifying the 'a' command with an address .

Then the CS value (here it is 1374) is supplied by the debugger. See the code below.

- a0220

1374:0220 MOV AL,09

1374:0222 MOV AH,76

1374:0224 ADD AL,AH

1374:0226 MOV [0124],AL

1374:0229 <Enter>

After entering the whole program, it can be unassembled using the 'u' command along with the starting address.

-u0220

1374:0220 B009 MOV AL,09

1374:0222 B476 MOV AH,76

1374:0224 00E0 ADD AL,AH

1374:0226 A22401 MOV [0124],AL

The op codes of the program are now in the unassembled form.

It can be run using the 't' or 'g' commands .The value [0124] is an address in the data segment, and the content of this location can be checked after program execution.

vii) The Compare command 'c'

This command compares two areas in memory. The default register is the DS register.

There are two formats for this.

a) C start of 'from' address Length of byte area start of 'to' address

Here, the starting address of the first area, the starting address of the second area and number of bytes to be compared are to be specified.

-c0200 L24 0400

This command compares the 24(H) bytes from addresses 0200 and 0400.

-c0200 L24 0400

1374:0220 B0 00 1374:0420

1374:0221 09 00 1374:0421

1374:0222 B4 00 1374:0422

1374:0223 76 00 1374:0423

The contents of locations that are different, is displayed.

b) C start of 'from' address end of 'from' address start of 'to' address

In this format, the range of addresses of the 'from' category is specified.

-c0200 0224 0400.

This format for comparison gives the same result as the previous format.

Instead of specifying a length of 24 locations by L24, the range of addresses from 0200 to 0224 is mentioned.

viii) The Enter command 'e'

This command is useful for keying in data into locations. By default, it is the data segment that is referred to. But other segments can be used in the base address: offset format.

a) In the following case, the three bytes 56, 45 and 56 (all in hex) are entered into locations starting from 0200.

-e 0200 56 45 56

To verify if this data has been entered, check using the 'd' command.

-d 0200 0202

1374:0200 56 45 56

b) Here, a string is entered into the addresses from 0200 onwards. Then a verification is done using the 'd' command.

-e 0203 "notme"

-d 0203 0210

1374:0200 6E 6F 74 6D 65-00 00 00 00 00 00 00 00 notme.....

1374:0210 00 .

The ASCII values of the characters are seen in the specified locations.

ix) The Fill command 'f'

This command fills up an area in memory with a particular data. Like the enter command, this area may be specified by the starting address, the length of the area in memory and the data to be replicated in these locations. It may also be specified by the starting address, the ending address and the data.

Thus, there are two formats for this command.

a) -f starting address Length data

-f 0110 L10 "HI"

This command causes the data corresponding to HI to be replicated 10(H) times starting from the address 0110.

b) -f starting address ending address data

-f 0110 0120 "hi"

In this case, the data 'hi' is stored in whole range of locations from 0110 to 0120.

-f 0110 0120 "hi"

-d 0110 0120

1381: 0110 68 69 68 69 68 69 68 69-68 69 68 69 68 69 ...hihihihihihihihi

1381: 0120 68h

-

x) The Proceed command 'p'

This command is useful for proceeding through a number of instructions. In programs where there are software interrupts, this command goes right over the interrupt commands, and executes the number of instructions mentioned in the program. The format of the proceed command is proceed P [=address] [number]

Either the address up to which line to proceed, or the number of instructions to execute, while proceeding may be specified. See the following program.

```

MOV CX,10
LEA SI,NUMBER
REPEA: MOV AH,01
        INT 21H
        SUB AL,30H
        CALL SQUARE

```

```

        MOV [SI],AL
        INC SI
        LOOP REPEA
        .EXIT
SQUARE  PROC NEAR
        MOV BL,AL
        MUL BL
        RET
        SQUARE ENDP
        END

```

If this is to be executed in single step, using the trace command, on reaching the INT instruction, control will go to bizarre locations within the interrupt routine. This can be avoided using the proceed command. Use it with a count – and that many instructions get executed. Or the address beyond that of the INT instruction can be used with the ‘p’ command.

(The same thing can be got done, by using the ‘g’ command with an address beyond that of the address of the INT instruction.)

xi) The Search command ‘s’

This command searches memory for characters in a list. The default register is the DS register. If the characters are found, the addresses are shown, otherwise, there is no response.

There are two formats for this.

a) s start-address length ‘data’

Here, the starting address of the first area, the starting address of the second area and number of bytes to be compared are to be specified. The following command searches for ‘C’ in 24H locations starting from 0100.

-s 0100 L24 ‘C’

1376:0110

1376:0116

1376:0118

The result of searching shows that there are three locations in which the character ‘C’ is present.

b) s start-from-address end-from-address start-to-address

This is the second format for the use of the string instruction.

xii) The Hex command ‘h’

The hexarithmic command is for adding and subtracting hexadecimal numbers. It has just two parameters – the two numbers to be added and subtracted. The response is the sum and difference of the numbers. The numbers can be one to four hexadecimal digits long. The addition and subtraction are unsigned, and no carry or borrow is shown beyond the fourth (high order) digit.

-

-h 6 2

0008 0100

-h 5678 1234

68AC 4444

-

In the first example, we are adding 0006 and 0002. The sum and difference is obtained.

xiii) The Input command ‘I’

The Input command can be used to read a byte of data from any of the I/O ports in the PC. The port address can be either one or two bytes long. The command reads the ports and displays the result.

```
-i 3fe
-23
```

In the above, a port with address 3FEH was read, to get a data of 23H.

xiv) The Output command ‘O’

The output command is the reverse of the input command. This can be used to send a single byte of data to a port.

```
-o 3fc 5
-
```

Advanced DOS Commands

The load, name and write commands are commands that need a very good and deep understanding of DOS, and hence will not be discussed here.

Flag Status

The status of the flags of the processor are seen in the debugger along with the contents of the registers. The set/reset status of the flags is shown below.

	Overflow	Direction	Interrupt	Sign	Zero	Aux.carry	Parity	Carry
Set	OV	DN	EI	NG	ZR	AC	PE	CY
Clear	NV	UP	OI	PL	NZ	NA	PO	NC

APPENDIX C

INSTRUCTION SET AND INSTRUCTION TIMING OF 8086

i) Instruction Set of 8086

In this table, the complete set of instructions of the 8086 has been listed, with their mnemonic and function. For knowing the details (format and examples) of each of the instruction, the section in the text book, where they are discussed is given in the third column of the table.

Data Transfer Instructions

Mnemonic	Function	Section	Table
MOV	Move byte or word to register or memory	3.2.1	3.1
IN, OUT	Input byte or word from port, output word to port	5.2	3.1
LEA	Load effective address	3.2.2	3.1
LDS, LES	Load pointer using data segment, extra segment	—	3.1
PUSH, POP	Push word onto stack, pop word off stack	3.2.4	3.1
XCHG	Exchange byte or word	—	3.1
XLAT	Translate byte using look-up table	3.2.3	3.1

Logical Instructions

Mnemonic	Function	Section	Table
NOT	Logical NOT of byte or word (one's complement)	—	3.7
AND	Logical AND of byte or word	—	3.7
OR	Logical OR of byte or word	—	3.7
XOR	Logical exclusive – OR of byte or word	—	3.7
TEST	Test byte or word (AND without storing)	—	3.7

Shift and Rotate Instructions

Mnemonic	Function	Section	Table
SHL, SHR	Logical shift left, right byte or word by 1 or CL	3.6.1	3.8
SAL, SAR	Arithmetic shift left, right byte or word by 1 or CL	4.7.3	3.8
ROL, ROR	Rotate left, right, byte or word by 1 or CL	3.6.2	3.8
RCL, RCR	Rotate left, right, through carry, byte or word by 1 or CL	3.6.2	3.8

Arithmetic Instructions

Mnemonic	Function	Section	Table
ADD	Add byte or word	3.4.2	3.4
SUB	Subtract byte or word	3.4.3	3.4
ADC	Add byte or word and carry	3.4.2	3.4
SBB	Subtract byte or word and carry (borrow)	3.4.3	3.4
INC	Increment byte or word	3.4.2	3.4
DEC	Decrement byte or word	3.4.3	3.4
NEG	Negate byte or word (two's complement)	—	3.4
CMP	Compare byte or word (subtract without storing)	3.4.4	3.4
MUL	Multiply byte or word (unsigned)	3.4.5	3.4
DIV	Divide byte or word (unsigned)	3.4.6	3.4
IMUL	Integer multiply byte or word (signed)	4.7.2	3.4
IDIV	Integer divide byte or word (signed)	4.7.2	3.4
CBW, CWD	Convert byte to word, word to double word	4.7	3.4
AAA	ASCII adjust for addition	4.5.1	3.4
AAS	ASCII adjust for subtraction	4.5.2	3.4
AAM	ASCII adjust for multiplication	4.5.3	3.4
AAD	ASCII adjust for division	4.5.3	3.4
DAA, DAS	Decimal adjust for addition, subtraction (binary coded decimal numbers)	4.4.2	3.4

Jump and Loop Instructions

Mnemonic	Function	Section	Table
JMP	Unconditional jump	3.3.1 and 3.3.2	—
JA (JNBE)	Jump if above (not below or equal)	—	3.3
JAE (JNB)	Jump if above or equal (not below)	—	3.3
JB (JNAE)	Jump if below (not above or equal)	—	3.3
JBE (JNA)	Jump if below or equal (not above)	—	3.3
JE (JZ)	Jump if equal (zero)	—	3.3

Mnemonic	Function	Section	Table
JG (JNLE)	Jump if greater (not less or equal)	—	4.2
JGE (JNL)	Jump if greater or equal (not less)	—	4.2
JL (JNGE)	Jump if less (not greater nor equal)	—	4.2
JLE (JNG)	Jump if less or equal (not greater)	—	4.2
JC, JNC	Jump if carry set, carry not set	—	3.3
JO, JNO	Jump if overflow, no overflow	—	4.2
JS, JNS	Jump if sign, no sign	—	4.2
JNP (JPO)	Jump if no parity (parity odd)	—	3.3
JP (JPE)	Jump if parity (parity even)	—	3.3
LOOP	Loop unconditional, count in CX	3.3.3	—
LOOPE (LOOPZ)	Loop if equal (zero), count in CX	3.3.3	—
LOOPNE (LOOPNZ)	Loop if not equal (not zero),count in CX	3.3.3	—
JCXZ	Jump if CX equals zero	—	3.3

Call and Return Instructions

Mnemonic	Function	Section	Table
CALL, RET	Call, return from procedure	4.2.2 and 4.2.3	—
INT, INTO	Software interrupt, interrupt if overflow	8.3 and 8.2.5	—
IRET	Return from interrupt	8.1.1	—

String Instructions

Mnemonic	Function	Section	Table
MOVS	Move byte or word string	4.1.1	4.1
CMPS	Compare byte or word string	4.1.2	4.1
SCAS	Scan byte or word string	4.1.3	4.1
LODS, STOS	Load, store byte or word string	4.1.4	4.1
REP	Repeat	—	4.1
REPE, REPZ	Repeat while equal, zero	—	4.1
REPNE, REPNZ	Repeat while not equal (not zero)	—	4.1

Processor and Flag Control Instructions

Mnemonic	Function	Section	Table
STC, CLC, CMC	Set, clear, complement carry flag	3.4.1	—
STD, CLD	Set, clear direction flag	—	4.1
STI, CLI	Set, clear interrupt enable flag	—	—
LAHF, SAHF	Load AH from flags, store AH into flags	—	3.1

(Continued)

Mnemonic	Function	Section	Table
PUSHF, POPF	Push flags onto stack, pop flags off stack	—	3.1
ESC	Escape to external processor interface	13.2.1	—
LOCK	Lock bus during next instruction	6.4.4	—
NOP	No operation (do nothing)	6.5.1	—
WAIT	Wait for signal on TEST input	13.2.1	—
HLT	Halt processor	6.3.5	—

ii) Instruction Timing of 8086

a) Number of cycles expended in calculating the 'Effective Address'

No.	Addressing mode	No. of clocks for calculation of EA
1	Direct	6
2	Register indirect	5
3	Register relative	9
4	Based indexed with BP as the base register	8
5	Based indexed with BX as the base register	7
6	Relative based indexed with BP as the base register	12
7	Relative based indexed with BX as the base register	11

b) Number of clock cycles expended for each instruction of 8086

Code	Description	8086
AAA	ASCII adjust for addition	8
AAD	ASCII adjust for division	60
AAM	ASCII adjust for multiplication	83
AAS	ASCII adjust for subtraction	8
ADC	Add with carry	
	reg to reg	3
	mem to reg	9+EA
	reg to mem	16+EA
	immed to reg	4
	immed to mem	17+EA
	immed to acc	4
ADD	Addition	
	reg to reg	3
	mem to reg	9+EA
	reg to mem	16+EA
	immed to reg	4
	immed to mem	17+EA
	immed to acc	4
AND	Logical AND	

Code	Description	8086
	reg to reg	3
	mem to reg	9+EA
	reg to mem	16+EA
	immed to reg	4
	immed to mem	17+EA
	immed to acc	4
CALL	Call a procedure	
	intra segment direct	19
	intra segment direct through register	16
	intra segment direct through memory	21+EA
	intra segment direct	28
CBW	Convert byte to word	2
CLC	Clear carry flag	2
CLD	Clear direction flag	2
CLI	Clear interrupt flag	2
CMC	Complement carry flag	2
CMP	Compare	
	reg to reg	3
	mem to mem	9+EA
	reg to mem	9+EA
	immed to reg	4
	immed to mem	10+EA
	immed to acc	4
CMPS/ CMPSB/ CMPSW	Compare string/ Compare byte string/ Compare word string	
	Not repeated	22
	REPE/REPNE CMPS/CMPSB/CMPSW	9+22/rep
CWD	Convert word to doubleword	5
DAA	Decimal adjust for addition	4
DAS	Decimal adjust for subtraction	4
DEC	Decrement by 1	
	16-bit reg	3
	8-bit reg	3
	memory	15+EA
DIV	Unsigned division	
	8-bit reg	80–90
	16-bit reg	144–162
	8-bit mem	(86–96)+EA
	16-bit mem	(150–168)+EA
ESC	Escape	
	reg	2
	mem	8+EA

(Continued)

Code	Description	8086
HLT	Halt	2
IDIV	Integer division	
	8-bit reg	101–112
	16-bit reg	165–184
	8-bit mem	(107–118)+EA
	16-bit mem	(171–190)+EA
IMUL	Integer multiplication	
	8-bit reg	80–98
	16-bit reg	128–154
	8-bit mem	(86–104)+EA
	16-bit mem	(134–160)+EA
IN	Input from I/O port	
	Fixed port	10
	Variable port through DX	8
INC	Increment by 1	
	16-bit reg	3
	8-bit reg	3
	mem	15+EA
INT	Interrupt	
	type=3	52
	type3	51
INTO	Interrupt if overflow	
	interrupt taken	53
	interrupt not taken	4
IRET	Return from interrupt	32
JA/	Jump if above/	16, noj 4
JNBE	Jump if not below or equal	
JAE/	Jump if above or equal	16, noj 4
JNB	Jump if not below/	
JNA	Jump if not above	
JCXZ	Jump if CX is Zero	18, noj 6
JE/	Jump if equal/	16, noj 4
JZ	Jump if Zero	
JG/	Jump if greater	16, noj 4
JNLE	Jump if not less, or equal	
JGE/	Jump if greater or equal/	16, noj 4
JNL	Jump if not less	
JL/	Jump if less/	16, noj 4
JNGE	Jump if not greater, or equal	
JLE/	Jump if less or equal/	16, noj 4
JNG	Jump if not greater	
JMP	Jump	
	intrasegment direct short	15
	intrasegment direct	15
	intersegment direct	15

Code	Description	8086
	intrasegment indirect through memory	18+EA
	intrasegment indirect through register	11
	intrasegment indirect	24+EA
JNE/	Jump if not equal/	16, noj 4
JNZ	Jump if not Zero	
JNO	Jump if not overflow	16, noj 4
JNP/	Jump if not parity/	16, noj 4
JPO	Jump if parity odd	
JNS	Jump if not sign	16, noj 4
JO	Jump if overflow	16, noj 4
JP/	Jump if parity/	16, noj 4
JPE	Jump if parity even	
JS	Jump if sign	16, noj 4
LAHF	Load AH from flags	4
LDS	Load pointer using DS/	
LES	Load pointer using ES	16+EA
LEA	Load effective address	2+EA
LOCK	Lock bus	2
LODS/	Load string/	
LODSB	Load byte string	
LODSW	Load word string	
	not repeated	12
	repeated	9+13/rep
LOOP	Loop	17, noj 5
LOOPE/	Loop if equal/	
LOOPZ	Loop if zero	18, noj 6
LOOPNE/	Loop if not equal/	
LOOPNZ	Loop if not zero	19, noj 5
MOV	Move	
	acc to mem	10
	mem to acc	10
	reg to reg	2
	mem to reg	8+EA
	reg to mem	9+EA
	immed to reg	4
	immed to mem	10+EA
	reg to SS/DS/ES	2
	mem to SS/DS/ES	8+EA
	segment reg to reg	2
	segment reg to mem	9+EA
MOVS/	Move string/	
MOVSB/	Move byte string/	
MOVSW	Move word string	

(Continued)

Code	Description	8086
MUL	Not repeated	18
	REP MOVSB/MOVSX/MOVSW	9+17/rep
	Unsigned multiplication	
	8-bit reg	70–77
	16-bit reg	118–133
NEG	8-bit mem	(76–83)+EA
	16-bit mem	(124–139)+EA
	Negate	
	reg	3
	mem	16+EA
NOP	No operation	3
NOT	Logical NOT	
	reg	3
OR	mem	16+EA
	Logical OR	
	reg to reg	3
	mem to reg	9+EA
	reg to mem	16+EA
OUT	immed to acc	4
	immed to reg	4
	immed to mem	17+EA
	Output to I/O port	
	fixed port	10
POP	variable port	8
	Pop word off stack	
	reg	8
	segment reg	8
POPF	memory	17+EA
	Pop flags off stack	8
PUSH	Push word onto stack	
	reg	11
	segment reg:ES/SS/CS	10
	memory	16+EA
PUSHD	Push double word onto stack	10
RCL	Rotate left through carry /	
	Rotate right through carry /	
	reg with single-shift	2
	reg with variable-shift	8+4/bit
	mem with single-shift	15+EA
RET	mem with variable-shift	20+EA+4/bnit
	Return from procedure /	
	Return far /	
	Return near	
	intra-segment	16
RETN	intra-segment with constant	20
	intra-segment	26

Code	Description	8086
	intra-segment with constant	25
ROL/	Rotate left	
ROR	Rotate right	
	reg with single-shift	2
	reg with variable-shift	8+4/bit
	mem with single-shift	15+EA
	mem with variable-shift	20+EA+4/bit
SAHF	Store AH into flags	4
SAL/	Shift arithmetic left/	
SAR/	Shift arithmetic right/	
SHL/	Shift logical left/	
SHR	Shift logical right	
	reg with single-shift	2
	reg with variable-shift	8+4/bit
	mem with single-shift	15+EA
	mem with variable shift	20+EA+4/bit
SBB	Subtract with borrow	
	reg from reg	3
	mem from reg	9+EA
	reg from mem	16+EA
	immed from acc	4
	immed from reg	4
	immed from mem	17+EA
SCAS/	Scan string/	
SCASB	Scan byte string	
SCASW	Scan word string	
	not repeated	15
	REPE/REPNE SCAS/SCASB/SCASW	9+15/rep
STC	Set carry flag	2
STD	Set direction flag	2
STI	Set interrupt flag	2
STOS/	Store string/	
STOSB/	Store byte string/	
STOSW	Store word string	
	Not repeated	11
	REP STOS/STOSB/STOSW	9+10/rep
STR	Store task register	
SUB	Subtraction	
	reg from reg	3
	mem from reg	9+EA
	reg from mem	16+EA
	immed from acc	4
	immed from reg	4
	immed from mem	17+EA
TEST	Test	
	reg with reg	3

(Continued)

Code	Description	8086
	mem with reg	9+EA
	immed with acc	4
	immed with reg	5
	immed with mem	11+EA
WAIT	Wait while TEST pin not asserted	4
XADD	Exchange and add	
XCHG	Exchange	
	reg with acc	3
	reg with mem	17+EA
	reg with reg	4
XLAT/	Translate	11
XOR	Logical exclusive OR	
	reg with reg	3
	mem with reg	9+EA
	reg with mem	16+EA
	immed with acc	4
	immed with reg	4
	immed with mem	17+EA

APPENDIX D

LIST OF DOS AND BIOS FUNCTIONS

In this, some of the DOS 21H functions and the BIOS 10H and 16H functions are listed.

i) List of important DOS INT 21H services

DOS system calls use the 21H vector. Some of the important INT 21H services are listed here.

Function No.	Action	Returns/Expects
AH=01	Read character from standard input device with echo	Returns: AL=ASCII of the input key
AH=02	Write character to standard output device	Expects: DL=ASCII code of output data
AH=05	Write character to printer	Expects: DL=ASCII of the output
AH=08	Read character input from standard input device, without echo	Returns: AL=ASCII of the input key
AH=09	Display string	Expects: DS:DX=segment: offset of string, terminated by '\$'
AH=0AH	Buffered string output	Expects: DS:DX=segment: offset of buffer
AH=25H	Set interrupt vector	Expects: AL=machine interrupt number DS:DX=segment offset of interrupt service routine
AH=2AH	Get DOS system date	Returns: DL=Day (1 to 31) AL=Day (0=Sunday, 1=Monday etc.) DH=Month (1 to 12), CX=Year (1980 to 2099)
AH=2BH	Set DOS system date	Expects: DL=Day (1 to 31) DH=Month (1 to 12), CX=Year (1980 to 2099)
AH=2CH	Get DOS system time	Returns: CH=Hour (0 to 23) CL=Minutes (0 to 59) DH=seconds (0 to 59) DL=100 th of seconds (0 to 99)

Function No.	Action	Returns/Expects
AH=2DH	Set DOS system time	Expects: CH=Hour (0 to 23) CL=Minutes (0 to 59) DH=seconds (0 to 59) DL=100 th of seconds (0 to 99)
AH=30H	Get DOS version number	Returns: AL=Major version number (DOS 6.2=6 etc) AH=Minor version number (DOS 6.2=2 etc)
AH=31H	Terminate and stay resident (TSR)	Expects: AL=return code DX=number of paragraphs to make resident in memory
AH=35H	Get interrupt vector	Expects: AH=35H AL=interrupt type number Returns: ES:BX of interrupt service routine
AH=4CH	Terminate with return code	Expects: AL=return code

ii) BIOS 10H functions for video

1. INT 10H, Function 00: Set Video Mode

Expects: AH=00
AL=display mode

Returns: None

Description: The function 00 is used to set the video mode. This function clears the screen, sets the BIOS variables, and initializes the video mode.

2. INT 10H, Function 01: Set Cursor Shape

Expects: AH=01
CH bits 0–4=starting line for cursor (20H=no cursor)
CL bits 0–4=ending line for cursor

Returns: None

Description: The function 01 is used to set the shape of the cursor. This is performed by selecting the starting and ending lines for the blinking hardware cursor (works in text mode only).

3. INT 10H, Function 02: Set Cursor Position

Expects: AH=02
BH=video page (must be zero in graphics mode)

DH=row (y coordinate)
DL=column (x coordinate)

Returns: None

Description: The function 02 is used to set the position of a cursor on the display device (monitor) using text coordinates (row and column).

4. INT 10H, Function 03 : Read Cursor Position

Expects: AH=03
BH=video page

Returns: DH=current row (y coordinate)
DL=current column (x coordinate)
CH=starting line for cursor
CL=ending line for cursor

Description: The function 03 is used to read the current position of the cursor on the display in text coordinates.

5. INT 10H, Function 04 : Read Light Pen Position

Expects: AH=04

Returns: AH=0 if light pen not down or not triggered
1 if light pen down or triggered
CH=pixel row (y coordinate in graphics mode 04-06)
CX=pixel row (y coordinate in graphics mode 0DH-13H)
BX=pixel column (x coordinate in graphics mode)
DH=character row(y coordinate 0-24 in text mode)
DL=character column (x coordinate 0-79 or 0-39 in text mode)

Description: The function 04 is used to read the current status and position of the light pen.

6. INT 10H, Function 05 : Set Active Video Page

Expects: AH=05
AL=page number

Returns: None for standard PC

Description: The function 05 is used to select the active display page for the video display.

7. INT 10H, Function 06 : Scroll (Initialize) Rectangle Window Up

Expects: AH=06
AL=Number of lines to scroll up
(if AL=Zero, entire window is cleared or blanked)
BH=blanked area attributes
CH=y coordinate, upper left corner of window
CL=x coordinate, upper left corner of window
DH=y coordinate, lower right corner of window
DL=x coordinate, lower right corner of window

Returns: None

Description: The function 06 is used to initialize a specified rectangular window of the display to ASCII characters with a given attribute, or scrolls the contents of a window up by a specified number of lines.

8. INT 10H, Function 07 : Scroll (Initialize) Rectangle Window Down

Expects: AH=07

AL=Number of lines to scroll down

(if AL=Zero, entire window is cleared or blanked)

BH=blanked area attributes

CH=y coordinate, upper left corner of window

CL=x coordinate, upper left corner of window

DH=y coordinate, lower right corner of window

DL=x coordinate, lower right corner of window

Returns: None

Description: The function 07 is used to initialize a specified rectangular window of the display to ASCII character with a given attribute, or scrolls the contents of a window down by a specified number of lines.

9. INT 10H, Function 08 : Read Character and Attribute at Cursor

Expects: AH=08

BH=video page (on CGA, must be 0 in graphics mode)

Returns: AH=attribute byte

AL=ASCII character code

Description: The function 08 used to read the ASCII character and its attribute at the current cursor position for the specified video page.

10. INT 10H, Function 09 : Write Character and Attribute at Cursor

Expects: AH=09

AL=ASCII character code

BH=video page

BL=attribute (in text mode) or color (in graphics mode)

CX=count of character to write (replication factor)

Returns: None

Description: The function 09 is used to write a specified ASCII character and its attribute to the video display at the current cursor position.

11. INT 10H, Function 0AH : Write Character Only at Cursor

Expects: AH=0AH

AL=ASCII character code

BH=video page

BL=color (graphics mode)

CX=count of characters to write (replication factor)

Returns: None

Description: The function 0AH is used to write an ASCII character to the video display at the current cursor position. The character uses the attribute of the previous character displayed at the same position.

12. INT 10H, Function 0BH : Set Color Palette

Expects: AH=0BH
In text mode
BH=00 selects border color
BL=border color (0-1FH: 10H to 1FH for high intensity)
In graphics mode
BH=01 select palette combination
BL=0 to select Red Green Blue combination
BL=1 to select Cyan Magenta White combination

Returns: None

Description: The function 0BH is used to set the contents of a color palette.

13. INT 10H, Function 0CH : Set Pixel

Expects: AH=0CH
AL=pixel color value
BH=video page
CX=column number (x coordinate in graphics mode)
DX=row number (y coordinate in graphics mode)

Returns: None

Description: The function 0CH is used to set the pixel point on the video display at the specified graphics coordinates. The range of possible valid (x, y) coordinate depends on the current video mode.

14. INT 10H, Function 0DH : Get Pixel

Expects: AH=0DH
BH=video page
CX=column number (x coordinate)
DX=row number (y coordinate)

Returns: AL=pixel value (pixel attribute-color)

Description: The function 0DH is used to read the attributes of pixel point on the video at the specified graphics coordinates. The range of possible valid (x, y) coordinate depends on the current video mode.

15. INT 10H, Function 0EH : Write Text in Teletype Mode

Expects: AH=0EH
AL=ASCII character code

BH=video page (in text mode)
 BL=foreground color (in graphics mode)

Returns: None

Description: The function 0EH is used to write an ASCII character to the video display at the current cursor position, using the specified color both in text and graphics mode.

16. INT 10H, Function 0FH : Get Video Mode

Expects: AH=0FH

Returns: AH=number of character columns on screen
 AL=display mode
 BH=active video page

Description: The function 0FH is used to read the current video mode of the active video controller.

17. INT 10H, Function 10H : Set Color Palette Registers

Expects: AH=10H

AL=00 for setting palette register
 01 for setting border color register
 02 for setting all palette register and border register
 03 for toggling blink or intensity bit (only on EGA)
 BH=color value
 BL=palette register to set (00 to 02) if AL=00
 blink/intensity bit AL=03
 0 enable intensity
 1 enable blinking
 ES:DX=segment:offset of color list (if AL=02)

Returns: None

Description: The function 10H is used to set palette register to the new color combination.

18. INT 10H, Function 13H : Display String

Expects: AH=13H

AL=write mode
 0 use attribute in BL cursor position is not updated after write
 1 use attribute in BL cursor position is updated after write
 2 string format: char, attr, ..., char, attr cursor position is not updated after write
 3 string format: char, attr, ..., char, attr cursor position is updated after write
 BH=video page
 BL=attribute (write modes 0 and 1)
 CX=length of string
 DH, DL=row, column to start display of string
 ES:BP=segment:offset of source string

Returns: None

Description: The function 13H is used to transfer a string to video buffer for current active display.

19. INT 10H, Function 0FEH : Get Video Buffer Pointer

Expects: AH=0FEH

ES:DI=segment:offset of assumed video buffer

B000:000H for monochrome adapter

B800:0000 for standard or color graphic adapter

Returns: ES:DI segment:offset of actual video buffer for current process

Description: The function 0FEH is used to obtain the memory pointer of the video buffer for the currently executing task under top view.

20. INT 10H, Function 0FFH : Update Video Buffer

Expects: AH=0FFH

CX=number of sequential characters that have been modified

DI=offset of first character that has been modified within shadow video buffer

ES=segment of shadow video buffer

Returns: None

Description: The function 0FFH is used to copy the contents of the application's shadow video buffer to the true video refresh buffer under Top View.

iii) BIOS 16H functions for keyboard control

Here is presented a few BIOS interrupts for keyboard control.

Some of the early BIOS functions catered to the old 83-key keyboard only. Later, some more BIOS functions were added for the enhanced keyboard.

- i) AH=0 or AH=10H. The first function number is for the older keyboard. The second one is the equivalent one for the enhanced keyboard. This function checks for a character in the keyboard buffer. If available, the scan code is returned in AH, and ASCII value in AL. For function keys (F0 to F12) which have no ASCII values, AL=0.
If no character is available in the keyboard buffer, the function waits for a key press.
- ii) AH=01 or AH=11H. The first function number is for the older keyboard. The second one is the equivalent one for the enhanced keyboard. This function is similar to the previous one, except that if no character is available in the keyboard buffer, it does not wait for a key press. It simply sets ZF (ZF=1) and returns.
- iii) AH=02 or AH=12H. The first function number is for the older keyboard. The second one is the equivalent one for the enhanced keyboard. This function returns the first keyboard status byte in the AL register. This status byte is also available in the BIOS data area 0040 : 0017.

APPENDIX E

80x87 INSTRUCTION SET (x87 – PENTIUM)

Legend

Generals

- reg** = floating point register, st(0), st(1) ... st(7)
- mem** = memory address
- mem32** = memory address of 32-bit item
- mem64** = memory address of 64-bit item
- mem80** = memory address of 80-bit item

FPU Instruction Timings

- FX** = pairs with FXCH
- NP** = no pairing
- Timings with a **hyphen** indicate a range of possible timings
- Timings with a **slash** (unless otherwise noted) are latency and throughput.
- Latency is the time between instructions dependent on the result.
- Throughput is the pipeline throughput between non conflicting instructions.
- EA** = cycles to calculate the effective address

FPU Instruction Sizing

- All FPU instructions that do not access memory are two bytes in length (except FWAIT which is one byte).
- FPU instructions that access memory are four bytes for 16-bit addressing and six bytes for 32-bit addressing.
- (end of legend)

Instruction Formats, Clock Cycles and Pentium® Pairing Info

F2XM1		Compute 2 ^x -1				
8087	287	387	486	Pentium		
310-630	310-630	211-476	140-279	13-57	NP	
FABS		Absolute value				
8087	287	387	486	Pentium		
10-17	10-17	22	3	1	FX	

FADD Floating point add
FADDP Floating point add and pop

variations/						
operand	8087	287	387	486	Pentium	
fadd	70-100	70-100	23-34	8-20	3/1	FX
fadd mem32	90-120+EA	90-120	24-32	8-20	3/1	FX
fadd mem64	95-125+EA	95-125	29-37	8-20	3/1	FX
faddp	75-105	75-105	23-31	8-20	3/1	FX

FBLD Load BCD

operand	8087	287	387	486	Pentium	
mem	(290-310)+EA	290-310	266-275	70-103	48-58	NP

FBSTP Store BCD and pop

	8087	287	387	486	Pentium	
	(520-540)+EA	520-540	512-534	172-176	148-154	NP

FCBS Change sign

	8087	287	387	486	Pentium	
	10-17	10-17	24-25	6	1	FX

FCLEX Clear exceptions

FNCLEX Clear exceptions, no wait

variations	8087	287	387	486	Pentium	
fclex	2-8	2-8	11	7	9	NP
fnclex	2-8	2-8	11	7	9	NP

The wait version may take additional cycles

FCOM Floating point compare

FCOMP Floating point compare and pop

FCOMPP Floating point compare and pop twice

variations/						
operand	8087	287	387	486	Pentium	
fcom reg	40-50	40-50	24	4	4/1	FX
fcom mem32	(60-70)+EA	60-70	26	4	4/1	FX
fcom mem64	(65-75)+EA	65-75	31	4	4/1	FX
fcomp	42-52	42-52	26	4	4/1	FX
fcompp	45-55	45-55	26	5	4/1	FX

FCOS Floating point cosine (387+)

8087	287	387	486	Pentium	
-	-	123-772	257-354	18-124	NP

Additional cycles required if operand > pi/4 (~3.141/4 = ~.785)

FDECSTP Decrement floating point stack pointer

8087	287	387	486	Pentium	
6-12	6-12	22	3	1	NP

FDISI Disable interrupts (8087 only, others do fnop)**FNDISI** Disable interrupts, no wait (8087 only, others do fnop)

Variations	8087	287	387	486	Pentium	
fdisi	2-8	2	2	3	1	NP
fndisi	2-8	2	2	3	1	NP

The wait version may take additional cycles

FDIV Floating divide**FDIVP** Floating divide and pop

variations/

operand	8087	287	387	486	Pentium	
fdiv reg	193-203	193-203	88-91	73	39	FX
fdiv mem32	(215-225)+EA	215-225	89	73	39	FX
fdiv mem64	(220-230)+EA	220-230	94	73	39	FX
fdivp	197-207	197-207	91	73	39	FX

FDIVR Floating divide reversed**FDIVRP** Floating divide reversed and pop

variations/

operand	8087	287	387	486	Pentium	
fdiv reg	194-204	194-204	88-91	73	39	FX
fdiv mem32	(216-226)+EA	216-226	89	73	39	FX
fdiv mem64	(221-231)+EA	221-231	94	73	39	FX
fdivp	198-208	198-208	91	73	39	FX

FENI Enable interrupts (8087 only, others do fnop)**FNENI** Enable interrupts, nowait (8087 only, others do fnop)

variations	8087	287	387	486	Pentium	
feni	2-8	2	2	3	1	NP
fneni	2-8	2	2	3	1	NP

FFREE Free register

8087	287	387	486	Pentium	
9-16	9-16	18	3	1	NP

FIADD Integer add

operand	8087	287	387	486	Pentium	
mem16	(102-137)+EA	102-137	71-85	20-35	7/4	NP
mem32	(108-143)+EA	108-143	57-72	19-32	7/4	NP

FICOM Integer compare**FICOMP** Integer compare and pop

variations/

operand	8087	287	387	486	Pentium	
ficom mem16	(72-86)+EA	72-86	71-75	16-20	8/4	NP
ficom mem32	(78-91)+EA	78-91	56-63	15-17	8/4	NP
ficomp mem16	(74-88)+EA	74-88	71-75	16-20	8/4	NP
ficomp mem32	(80-93)+EA	80-93	56-63	15-17	8/4	NP

FIDIV Integer divide**FIDIVR** Integer divide reversed

variations/

operand	8087	287	387	486	Pentium	
fidiv mem16	(224-238)+EA	224-238	136-140	85-89	42	NP
fidiv mem32	(230-243)+EA	230-243	120-127	84-86	42	NP
fidivr mem16	(225-239)+EA	225-239	135-141	85-89	42	NP
fidivr mem32	(231-245)+EA	231-245	121-128	84-86	42	NP

FILD Load integer

operand	8087	287	387	486	Pentium	
mem16	(46-54)+EA	46-54	61-65	13-16	3/1	NP
mem32	(52-60)+EA	52-60	45-52	9-12	3/1	NP
mem64	(60-68)+EA	60-68	56-67	10-18	3/1	NP

FIMUL Integer multiply

operand	8087	287	387	486	Pentium	
mem16	(124-138)+EA	124-138	76-87	23-27	7/4	NP
mem32	(130-144)+EA	130-144	61-82	22-24	7/4	NP

FINCSTP Increment floating point stack pointer

	8087	287	387	486	Pentium	
6-12	6-12	21	3		1	NP

FINIT Initialize floating point processor**FNINIT** Initialize floating point processor, no wait

variations	8087	287	387	486	Pentium	
finit	2-8	2-8	33	17	16	NP
fninit	2-8	2-8	33	17	12	NP

The wait version may take additional cycles

FIST Store integer**FISTP** Store integer and pop

variations/

operand	8087	287	387	486	Pentium	
fist mem16	(80-90)+EA	80-90	82-95	29-34	6	NP

fist	mem32	(82-92)+EA	82-92	79-93	28-34	6	NP
fistp	mem16	(82-92)+EA	82-92	82-95	29-34	6	NP
fistp	mem32	(84-94)+EA	84-94	79-93	28-34	6	NP
fistp	mem64	(94-105)+EA	94-105	80-97	28-34	6	NP

FISUB Integer subtract

FISUBR Integer subtract reversed

variations/

operand	8087	287	387	486	Pentium	
fisub	mem16	(102-137)+EA	102-137	71-85	20-35	7/4 NP
fisubr	mem32	(108-143)+EA	108-143	57-82	19-32	7/4 NP

FLD Floating point load

operand	8087	287	387	486	Pentium	
reg	17-22	17-22	14	4	1	FX
mem32	(38-56)+EA	38-56	20	3	1	FX
mem64	(40-60)+EA	40-60	25	3	1	FX
mem80	(53-65)+EA	53-65	44	6	3	NP

Load floating point constants

FLDZ Load constant onto stack, 0.0

FLD1 Load constant onto stack, 1.0

FLDL2E Load constant onto stack, logarithm base 2 (e)

FLDL2T Load constant onto stack, logarithm base 2 (10)

FLDLG2 Load constant onto stack, logarithm base 10 (2)

FLDLN2 Load constant onto stack, natural logarithm (2)

FLDPI Load constant onto stack, pi (3.14159...)

Variations	8087	287	387	486	Pentium	
fldz	11-17	11-17	20	4	2	NP
fld1	15-21	15-21	24	4	2	NP
fldl2e	15-21	15-21	40	8	5/3	NP
fldl2t	16-22	16-22	40	8	5/3	NP
fldlg2	18-24	18-24	41	8	5/3	NP
fldln2	17-23	17-23	41	8	5/3	NP
fldpi	16-22	16-22	40	8	5/3	NP

FLDCW Load control word

operand	8087	287	387	486	Pentium	
mem16	(7-14)+EA	7-14	19	4	7	NP

FLDENV Load environment state

operand	8087	287	387	486	Pentium	
mem	(35-45)+EA	35-45	71	44/34	37/32-33	NP

cycles for real mode/protected mode

FMUL Floating point multiply
FMULP Floating point multiply and pop

variations/

operand	8087	287	387	486	Pentium	
fmul reg s	90-105	90-105	29-52	16	3/1	FX
fmul reg	130-145	130-145	46-57	16	3/1	FX
fmul mem32	(110-125)+EA	110-125	27-35	11	3/1	FX
fmul mem64	(154-168)+EA	154-168	32-57	14	3/1	FX
fmlp reg s	94-108	94-108	29-52	16	3/1	FX
fmlp reg	134-148	134-148	29-57	16	3/1	FX

s = register with 40 trailing zeros in fraction

FNOP no operation

	8087	287	387	486	Pentium	
	10-16	10-16	12	3	1	NP

FPATAN Partial arctangent

	8087	287	387	486	Pentium	
	250-800	250-800	314-487	218-303	17-173	

FPREM Partial remainder

FPREM1 Partial remainder (IEEE compatible, 387+)

variations	8087	287	387	486	Pentium	
fprem	15-190	15-190	74-155	70-138	16-64	NP
fprem1	-	-	95-185	72-167	20-70	NP

FPTAN Partial tangent

	8087	287	387	486	Pentium	
	30-540	30-540	191-497	200-273	17-173	NP

Additional cycles required if operand > pi/4 (~3.141/4 =~.785)

FRNDINT Round to integer

	8087	287	387	486	Pentium	
	16-50	16-50	66-80	21-30	9-20	NP

FRSTOR Restore saved state

variations/

operand	8087	287	387	486	Pentium	
frstor mem	(197-207)+EA	197-207	308	131/120	75-95/70	NP
frstorw mem	-	-	308	131/120	75-95/70	NP
frstord mem	-	-	308	131/120	75-95/70	NP

cycles for real mode/protected mode

Save FPU State

FSAVE Save FPU state
FSAVEW Save FPU state, 16-bit format (387+)
FSAVED Save FPU state, 32-bit format (387+)
FSAVE Save FPU state, no wait
FSAVEW Save FPU state, no wait, 16-bit format (387+)
FSAVED Save FPU state, no wait, 32-bit format (387+)

variations	8087	287	387	486	Pentium	
fsave	(197-207)+EA	197-207	375-376	154/143	127-151/124	NP
fsavew			375-376	154/143	127-151/124	NP
fsaved			375-376	154/143	127-151/124	NP
fnsave	(197-207)+EA	197-207	375-376	154/143	127-151/124	NP
fnsavew			375-376	154/143	127-151/124	NP
fnsaved			375-376	154/143	127-151/124	NP

Cycles for real mode/protected mode
 The wait version may take additional cycles

FSCALE Scale by factor of 2

	8087	287	387	486	Pentium	
32-38	32-38	67-86	30-32	20-31	NP	

FSETPM Set protected mode (287 only, 387+ = fnop)

	8087	287	387	486	Pentium	
-	2-8	12	3	1	NP	

FSIN Sine (387+)

FSINCOS Sine and cosine (387+)

Variations	8087	287	387	486	Pentium	
Fsin	-	-	122-771	257-354	16-126	NP
Fsincos	-	-	194-809	292-365	17-137	NP

Additional cycles required if operand > pi/4 (~3.141/4 = ~.785)

FSQRT Square root

	8087	287	387	486	Pentium	
180-186	180-186	122-129	83-87	70	NP	

FST Floating point store

FSTP Floating point store and pop

variations/ operand	8087	287	387	486	Pentium	
fst reg	15-22	15-22	11	3	1	NP
fst mem32	(84-90)+EA	84-90	44	7	2	NP
fst mem64	(96-104)+EA	96-104	45	8	2	NP
fstp reg	17-24	17-24	12	3	1	NP

fstp mem32	(86-92)+EA	86-92	44	7	2	NP
fstp mem64	(98-106)+EA	98-106	45	8	2	NP
fstp mem80	(52-58)+EA	52-58	53	6	3	NP

FSTCW Store control word

FNSTCW Store control word, no wait

variations/

operand	8087	287	387	486	Pentium	
fstcw mem	12-18	12-18	15	3	2	NP
fnstcw mem	12-18	12-18	15	3	2	NP

The wait version may take additional cycles

Store FPU environment

FSTENV Store FPU environment

FSTENVW Store FPU environment, 16-bit format (387+)

FSTENVVD Store FPU environment, 32-bit format (387+)

FNSTENV Store FPU environment, no wait

FNSTENVW Store FPU environment, no wait, 16-bit format (387+)

FNSTENVVD Store FPU environment, no wait, 32-bit format (387+)

variations/

operand	8087	287	387	486	Pentium	
fstenv mem (40-50)+EA	40-50	103-104	67/56	48-50	NP	
fstenvw mem		103-104	67/56	48-50	NP	
fstenvd mem		103-104	67/56	48-50	NP	
fnstenv mem (40-50)+EA	40-50	103-104	67/56	48-50	NP	
fnstenvw mem		103-104	67/56	48-50	NP	
fnstenvd mem		103-104	67/56	48-50	NP	

Cycles for real mode/protected mode

The wait version may take additional cycles

FSTSW Store status word

FNSTSW Store status word, no wait

variations/

operand	8087	287	387	486	Pentium	
fstsw mem	12-18	12-18	15	3	2	NP
fstsw ax	-	10-16	13	3	2	NP
fnstsw mem	12-18	12-18	15	3	2	NP
fnstsw ax	-	10-16	13	3	2	NP

The wait version may take additional cycles

FSUB Floating point subtract**FSUBP** Floating point subtract and pop

variations/

operand	8087	287	387	486	Pentium	
fsub reg	70-100	70-100	26-37	8-20	3/1	FX
fsub mem32	(90-120)+EA	90-120	24-32	8-20	3/1	FX
fsub mem64	(95-125)+EA	95-125	28-36	8-20	3/1	FX
fsubp reg	75-105	75-105	26-34	8-20	3/1	FX

FSUBr Floating point subtract**FSUBR** Floating point subtract and pop

variations/

operand	8087	287	387	486	Pentium	
fsubr reg	70-100	70-100	26-37	8-20	3/1	FX
fsubr mem32	(90-120)+EA	90-120	24-32	8-20	3/1	FX
fsubr mem64	(95-125)+EA	95-125	28-36	8-20	3/1	FX
fsubrp reg	75-105	75-105	26-34	8-20	3/1	FX

FTST Floating point test for zero

8087	287	387	486	Pentium	
38-48	38-48	28	4	4/1	FX

FUCOM Unordered floating point compare (387+)**FUCOMP** Unordered floating point compare and pop (387+)**FUCOMPP** Unordered floating point compare and pop twice (387+)

Variations	8087	287	387	486	Pentium	
Fucom	-	-	24	4	4/1	FX
Fucomp	-	-	26	4	4/1	FX
Fucompp	-	-	26	5	4/1	FX

FWAIT Wait while FPU is executing

8087	287	387	486	Pentium	
4	3	6	1-3	1-3	NP

FXAM Examine condition flags

8087	287	387	486	Pentium	
12-23	12-23	30-38	8	21	NP

FXCH Exchange floating point registers

8087	287	387	486	Pentium
10-15	10-15	18	4	0-1 *

* FCXH is pairable in the V pipe with all FX pairable instructions

FXTRACT Extract exponent and significand

8087	287	387	486	Pentium
27-55	27-55	70-76	16-20	13 NP

FYL2X Compute $Y * \log_2(x)$

FYL2XP1 Compute $Y * \log_2(x+1)$

variations	8087	287	387	486	Pentium
fyl2x	900-1100	900-1100	120-538	196-329	22-111 NP
fyl2xp1	700-1000	700-1000	257-547	171-326	22-103 NP

BIBLIOGRAPHY

- Abel, Peter, *IBM PC Assembly Language Programming*, Pearson Education, 2002, pp. 156–163.
- Antonakos, James L., *An Introduction to the Intel Family of Microprocessors*, 3rd ed., Pearson Education, 2006.
- Ayala, Kenneth J., *The 8086 Microprocessor, Programming and Interfacing the PC*, Thomson Education, 2004, pp. 240–243.
- Brey, Barry B., *The Intel Microprocessors – Architecture, Programming and Interfacing*, 8th ed., Pearson Education, 2009.
- Cypress Manual for 6264 Static RAM, October 1994, Revised June 1996.
- Dandamudi, Sivarama P., *Introduction to Assembly Language Programming*, Springer International Edition, 2004.
- Hall, Douglas V., *Microprocessors and Interfacing – Programming and Hardware*, Tata McGraw Hill Publishers, 2009, p. 315.
- Intel Manual for 16-bit HMOS Microprocessor 8086/8086-2/8086-1, September 1990.
- Intel Manual for 8253/8254 Programmable Interval Timer, November 1986.
- Intel Manual for 8259A Programmable Interrupt Controller, December 1988.
- Intel Manual for 8279A Keyboard Display Interface, September 1987.
- Intel Manual for 8251A Programmable Communications Interface, November 1988.
- Intel 80386 Programmer's Reference Manual, 1986.
- Intel Manual for 80486 Family of Microprocessors, December 1992.
- Intel Manual for Pentium Processor at iCOMP® Index 610\75 MHz, June 1997.
- Intel Manual for 82C55A CHMOS Programmable Interface, October 1995.
- Intel Manual for 8087 Math Coprocessor, October 1989.
- Intersil Manual for 8289 Bus Arbiter, March 1997.
- Lecture Notes on the Status of IEEE Standard 754 for Binary Floating-Point Arithmetic, Prof. W. Kahan, Department of Electrical Engineering and Computer Science, University of California, October 1997.
- Liu, Yu Cheng and Gibson, Glenn A., *Microcomputer Systems: The 8086/8088 Family*, Prentice Hall of India, 1991, pp. 463–465.
- Mazidi, Muhammed Ali and Mazidi, Janice Gillispie, *The 80x86 IBM PC and Compatible Computers – Assembly Language, Design and Interfacing*, 4th ed., Prentice Hall, 2002, pp. 124–127, 484–489, 542–550, 731–737.
- Mazidi, Muhammed Ali, Mazidi, Janice Gillispie and McKinlay, Rolin D., *The 8051 Microcontroller and Embedded Systems*, 2nd ed. (IE), Pearson Education, 2008, pp. 435–438.
- Mueller, Scott, *Upgrading and Repairing PCs*, 16th ed., Pearson Education, 2004, pp. 239–242, 398–399, 990–992, 1161–1162.
- National Semiconductor Manual for DAC0800/DAC0802 8-bit Digital-to-Analog Converter, September 2006.

Rafiquzzaman, Mohammed, *Fundamentals of Digital Logic and Microcomputer Design*, 5th ed., Wiley Publishers, 2005, pp. 419–420.

Ray, A. K. and Bhurchandi, K. M., *Advanced Microprocessors and Peripherals*, Tata McGraw Hill Publishers, 2002.

Texas Instruments Manual for MAX232, MAX232I Dual EIA-232 Drivers/Receivers, October 2002.

Triebel, Walter A. and Singh, Avatar, *The 8088 and 8086 Microprocessors, Programming, Interfacing, Software, Hardware and Applications*, 4th ed., Pearson Education, 2002.

Users' Guide to ASUS Motherboard A7V266–MX, October 2003.

http://en.wikibooks.org/wiki/X86_Assembly/Print_Version, accessed in December 2008.

[http://msdn.microsoft.com/en-us/library/c9tff918\(VS.80\).aspx](http://msdn.microsoft.com/en-us/library/c9tff918(VS.80).aspx), accessed in January 2009.

http://webster.cs.ucr.edu/Page_TechDocs/MASMDoc/ProgrammersGuide/Appendix_A.htm, accessed in November 2008.

<http://www.beyondlogic.org/serial/serial.htm>, accessed in April 2009.

<http://www.ecs.umass.edu/ece/koren/architecture/Cache/tutorial.html>, accessed in April 2009.

http://www.cs.iastate.edu/~prabhu/Tutorial/CACHE/bl_place.html, accessed in April 2009.

<http://www.laynetworks.com/Direct%20Mapped%20Cache.htm>, accessed in April 2009.

http://arstechnica.com/paedia/r/ram_guide/ram_guide.part1-1.html, accessed in July 2009.

<http://www.karbosguide.com/books/pcarchitecture/start.htm>, accessed in August 2009.

<http://www.pcguides.com/>, accessed in August 2009.

INDEX

4004, 1
 64-bit, 3
 data bus, 2–3, 6–7, 9–10, 32, 42–45
 8008, 1–2
 80186, 487–88, 492–96, 501
 80188, 487
 80286, 2–4, 503–05, 514, 518, 524, 540
 80386 SX, 504
 80386, 503–05, 507–08, 510–16, 524–28, 530–31, 535, 538–42
 enhancements, 503–06, 513
 hardware features, 505, 510, 528, 535
 internal architecture, 505
 interrupts of, 538
 privilege levels of, 530
 programming enhancements, 503, 506
 real mode operations, 513
 80486, 3–4, 6, 545–46, 548–49, 552, 555
 enhanced features of, 545
 8080, 1–2
 8085 microprocessor, 1, 31
 clock, 1, 5–6, 8–9, 12, 31, 42–4
 hardware aspects of, 41
 power supply, 43
 programming model, 1, 31–2
 read and write, 8, 43
 reset and ready pins, 43
 8086, 51–5, 58, 60–1, 63
 in a coprocessor configuration, 455, 458, 461
 internal block diagram of, 51–3
 register set of, 52
 8086 and 8089, 462
 8086 based system, 380
 interrupt sequence for, 380
 8087, 455, 458–61, 463–64, 471–72, 477–79, 481–82
 features of, 461, 472, 477, 561, 568, 570, 572
 registers of, 478
 8088, 503
 8253/8254, 349–350
 8255, 291–99, 340, 308–10, 313, 315–18, 323–24, 332–33, 337–39, 342–44
 control word, 351–54, 356–58, 360–61, 363, 368, 371–73, 375, 381, 384–85, 389
 internal block diagram, 292–93

 pin configuration, 292, 303–04
 8259, 349, 377–81, 383–85, 389–94
 block diagram, 350–51, 353, 363, 368, 373, 378–79, 404–05, 419–22
 features, 349, 363, 377, 385, 389
 using the, 363–64, 367, 381, 389
 8279, 349, 363–66, 368–69, 371

A

Accelerated Graphics Port (AGP), 562, 565
 access time, 431–33, 437–38, 440, 446
 accumulator, 31–2, 35–7, 46
 active low signal, 226
 ADD, 10, 17, 19, 21, 23, 30, 35–9
 address bus, 2, 6–7, 9–10, 32, 42–5
 address, 51–2, 55, 58–66
 base, 55, 58–62, 65–7
 branch, 58
 effective, 61, 64, 66–7
 logical, 59–64
 physical, 58–9, 62, 64
 addressing, 2, 34–7, 39, 51, 55, 61, 63–67
 base indexed, 66–7
 direct, 10–2, 35, 46, 148–49, 160, 166, 173, 443, 447–50
 fixed port, 181
 immediate, 35–7
 indirect, 35–7, 39, 148–49
 modes, 34–5
 register, 6, 10–11, 32, 34–6, 39, 489–90, 492–501
 relative base indexed, 66–7
 register indirect, 64–5
 register relative, 66–7
 variable port, 181
 ALE, 257
 Alienware, 584
 ALU, 7, 31–2, 219
 AMD, 2, 3
 Apple Mac, 5
 architecture, 1–2, 5, 11–2
 internal, 2–3, 6, 43, 51–3
 arithmetic, 139, 151, 158–59, 165–66, 172
 signed number, 139, 148, 166–69, 172
 arithmetic co-processing, 461
 arithmetic co-processor, 471

ascending order, 153
 ASCII, 19–20, 39
 addition, 10, 17, 22–5, 27–30, 36–7, 39
 code, 10, 11, 17–20, 34–5
 division, 15, 37
 multiplication, 10, 37–8
 subtraction, 10, 25–8, 29, 40
 assembler, 11
 assemblers, 71–4, 86
 features of, 72, 84
 for x86, 71, 73
 macro, 72, 74, 76
 MASM, 73–9, 84–6
 TASM, 73–4
 two-pass, 73
 assembly error, 79
 assembly language, 71–4, 76, 78–9, 81, 83, 85–6
 general rules for writing, 85
 process, 71–4, 76, 86
 program development tools, 71, 86
 asymmetric square wave, 499
 asynchronous read, 438
 ATA, 561, 563, 565–66, 573–75, 580

B

backward references, 73
 basic input output system, 97
 baud rate, 397–98, 405–06, 408, 410, 412
 BCD numbers, 17–8, 23, 25
 bias, 473–74, 476
 BIN directory, 74–5
 binary numbers, 13, 17–8, 22, 25
 addition, 17, 22–5, 27–30
 fractional, 13–4
 BIOS 10H, 259, 264–65, 267
 functions, 258–59, 264–65, 267, 272, 274–78, 280, 283
 BIOS INT 09, 274
 Bit Set Reset mode, 300
 BIU, 219
 blanking command, 363
 blast signal, 547
 bottom up approach, 96
 bound, 490
 address, 487, 489–91, 493–94, 500
 register, 490
 bps, 397–98, 408
 branch prediction unit, 552, 554

- branch address, 552
- branch history table, 552
- branch target buffer, 552, 554
- dynamic prediction, 552
- speculative execution, 552, 554, 557
- static prediction, 552
- BRDY* signal, 547
- buffer, 257, 274–77
- burst mode, 441
- burst mode data transfer, 547
- bus arbitration, 455, 457, 464–66, 468
 - between different processing modules, 468
 - using the 8289 bus arbiter IC, 466
- bus contention, 455, 459
- bus controller, 213–16, 218
- Bus High Enable (BHE), 237
- bus interface unit, 51, 58, 60
- Bus Low Enable (BLE), 237

C

- C, 32, 35, 38–9, 41
- cache, 431–32, 435, 444–54
 - and the x86 family, 451
 - hit, 445, 450
 - memory, 431–53
 - organization, 434, 446–47
 - replacement policies, 451
 - write policies, 451
- Call gates, 533–34, 538
 - access control policy for, 533
- CAPSLOCK, 276
- carry, 23–5, 29–30, 33–4, 37–8, 39–40, 42
- carry flag, 34
 - auxiliary, 34
- cascade mode, 382–83, 390
 - programmable interrupt controllers in the PC, 390
- CBW, 168
- centrino board, 584
- centronics, 310–12
 - printer interface, 310–11
- CGA, 4
- character display, 262, 270
- chip, 289–96, 298, 308, 310, 313–18, 343
 - programmable, 289–91, 322–23, 349–50, 355, 363, 371, 377–79, 390–91
- chip select logic, 487, 494
- chips, 1, 8, 10
 - 1KB, 8, 31
 - 1MB, 8, 31
 - 256 bytes, 8, 31
- chipset, 561, 564–68, 577–78, 580, 584
- CISC architecture, 12
- Clock, 1, 5–6, 8–9, 12, 31, 43–4, 51, 197, 199, 203, 206
 - generation, 203–04, 207
 - period, 8
 - system, 1, 5–11, 14–5, 26
- clock doubling, 546
- CMOS NVRAM, 579
 - Digital Video In (DVI), 582
 - Real Time Clock (RTC), 579–80
- code segment, 58, 60, 62
- code segments, 530, 532–34
 - conforming, 518, 530, 532–33
 - non-conforming, 532–33
- code, 7, 10–11, 17–20, 34–5, 71–92
 - designing, 12, 88–9
 - machine, 3–5, 10–11, 32
 - object, 11, 72–6, 78
 - source, 11, 35–6, 40, 71–3, 78, 88–9
 - symbolic, 10–11
- color configuration, 263
- COM Files, 78
- Command Word, 407, 410–15, 425
- common anode, 341–42
- communication, 395–99, 402–07, 411, 415
 - asynchronous, 396, 403, 405–08, 410–12, 415
 - full duplex, 395–96
 - half duplex, 395–96, 405
 - simplex, 395
 - synchronous, 396, 403, 405–08, 412, 438, 440, 442–43
- compare register, 497–99
- compatibility, 3, 5
 - backward, 3
- compiler, 189, 191
 - Borland's Turbo C++, 189
- computer arithmetic, 1, 22, 29, 472
- control bus, 6–7, 9–10, 32
- control transfers, 532
 - privilege checks for, 532
- controller, 9
- conversion ports, 244
 - analog to digital, 244
- conversions, 1, 12, 15, 158–59, 165
 - binary to hexadecimal, 16
 - decimal to binary, 15
 - decimal to hexadecimal, 16
 - from two's complement form 21
 - logical address to a physical address, 522–23
- coprocessor, 213, 217
- counters, 350, 352–54, 357
 - programming the chip, 352, 368
- CP/M, 96
- CPU, 5–6, 8
- CWD, 168
- cycle, 7–9, 42, 195–97, 199, 200–201, 203, 205–15, 218–21
 - memory read, 6, 9
 - memory write, 6, 9

D

- D and W bits, 87
- DAA, 160–61
- DAC, 289, 317–19, 321
- daisy chaining, 464–65, 469, 471
- Darlington pair, 332, 335
- data alignment, 545, 548–49
- data bus, 2–3, 6–7, 9–10, 32, 42–5
 - external, 2–3, 5, 43–4
- data byte, 79
- data dependency problem, 552
- data segments, 505, 529, 531–32
- data transfer, 289, 291–92, 296, 302, 304, 310, 312–13
 - parallel, 289, 291, 298, 310, 312–13, 328, 332
- data types, 74, 79, 86
- decoders, 229, 231
 - address range, 225–29, 233, 240
 - block, 229, 231–32
- decoding, 225, 227–35, 237, 239, 244
 - exhaustive, 234
 - memory address, 225, 227, 234, 239
 - partial address, 225, 234–35
- delay loops, 195, 220, 222–23
- descending order, 152
- descriptor tables, 518–19, 523, 528
 - global, 519
 - local, 519–20, 528
- design of 8086, 87
 - instruction set, 86–7, 487–88, 490
- desktops, 583–84
- devices, 2, 5–6, 9–10, 43–4
 - input, 5, 9, 35–6, 301–10, 313–18, 328, 337–38
 - output, 5, 9, 39–40, 41–2, 291–92, 294–98, 300–10, 313, 315–21, 334, 337, 344–48
- DIP, 31
- direct memory access, 395, 415
- directive, 76–7, 79, 81–5
 - assume, 84–5
 - DUP, 82
 - ends, 84
 - EQU, 82
 - ORG, 83–4, 92
 - segment, 71, 75–82, 84–8, 92
- directives, 71–2, 74, 79, 84
- disk, 4
 - floppy, 4
- disk operating systems, 96
- display interface, 363, 365, 367–68
- disassembling, 77
- distributed arbitration, 466
- division, 114, 120, 126–28
 - by zero, 127
 - dividend, 120, 126–28
 - quotient, 120, 126–28
 - signed, 108, 111, 113–14, 131, 139, 148, 166–72

unsigned, 95, 111, 114, 121, 124, 126, 131

DMA, 6, 45–6

features, 419

fetch and deposit, 417–18

FLYBY, 416

initialization and programming, 425

transfer modes, 416–18

DMA controller, 212, 395, 416–20, 423–24, 426–27

base address register, 420–23

base word count register, 420

block diagram of, 492

DMA and IBM-PC, 426

DMA and PC-AT, 428

double precision, 472–75, 477

DRAM, 431, 433, 435–40, 442–43, 445–46

chip, 431–41, 444, 452

fast page mode dram, 439

read cycle of, 433, 435–37

synchronous, 396, 403, 405–08, 412, 438, 440, 442–43

drivers, 11

device, 1, 7, 9–11, 30–1, 39, 45–6

dual channel memory, 567, 577

Dual In Line Memory Module (DIMM), 576

dynamic multiplexed display, 289, 344

E

echo, 97, 123

edge triggered mode, 378

EDORAM, 439

EFLAG, 505

embedded microprocessor, 488

history and development, 488

enhancements, 487, 490

instruction set, 487–88, 490

EPROM, 226, 230–31, 234

execution unit, 51–2

EXE Files, 78

expansion buses, 568

ISA Bus, 565, 568–69, 573

PCI Bus, 566, 568–70

PCI Express, 561, 565, 568–71

USB, see universal serial bus

EXTRN, 184–85, 187–88

F

factorial, 125–26

fetching, 58

FIFO RAM, 367–68, 370, 374–75

firewire port, 582

Flag, 33–4, 40, 55–7

busy, 323, 348

carry, 23–5, 29–30, 33–4, 37–40, 42

conditional, 55–7

control, 51, 55, 57

interrupt, 57

parity, 33–4

sign, 19, 21–2, 27, 29, 30, 33–4

trap, 57

Flash ROM, 444

floating point number, 473

floating point unit, 551

form factors, 582

ATX, 580, 583

BTX, 583

cooling system, 583

format, 56, 59, 61, 63

big endian, 63

little endian, 63, 65

forward references, 73

full segment, 71, 75, 84–5

function calls, 97, 101

BIOS, 96–7, 249, 253, 255–56, 258–61, 264–65, 267, 270, 272, 274–77, 283

DOS, 95–7, 99–102, 105–06, 109, 122–23, 128, 249, 255–56, 259–60, 269, 272, 274, 277–78, 280–81, 283–87

function keys, 272, 274, 276

functional block, 95–6

G

gallium arsenide, 336

Global Descriptor Table Register (GDTR), 519, 522

Graphic User Interface (GUI), 504

H

handshaking, 296, 302–05, 308, 310, 312

in the RS232 protocol, 402

signals, 289–90, 292, 296, 302–06, 308, 310, 311–13, 322, 325, 348

hard disk, 562–63, 571, 573–74, 579, 583–84

interface, 565–66, 568, 571, 573, 575, 582

platters, 573

sectors, 573–74

track, 573

hardware, 5, 12, 31, 41, 44–6, 195–97, 217

structure of 8086, 196

hardware retriggerable one shot, 359

hardware triggered mode, 362

hex keyboard, 289, 337–39

hex representation, 17–8, 30

high level language constructs, 172–73, 175–76

hit rate, 445

hooking, 249, 277, 280

hardware interrupts, 249–50, 256, 258, 280

interrupts, 249–50, 255–59, 261, 272, 274–75, 277–78, 283

hot key, 280–83, 285

I

I/O Read, 6

I/O, 1–2, 6–7, 9–10, 32, 36, 42–3, 46

isolated, 180, 201, 212, 225, 239, 243–44, 336, 415, 423, 528

ports, 1–2, 45

IBM, 2, 4–5, 503

PC-AT, 4

PC-XT, 4

personal computer, 2, 4–5

IBM PC, 225, 239–40

memory map of, 239

IBM personal system/2, 582

IDIV, 170–71

IMUL, 170–71

inequality, 142–43

Initialization Command Words (ICWs), 381

ICW1, 381–85

ICW2, 381–85, 387

ICW3, 381–82, 384–86

ICW4, 381–86

inline assembly, 189

Input/Output (I/O), 179, 183

address decoding, 225, 227–35, 237, 239, 244, 293–94, 354, 384, 422, 462

decoding 16-bit addresses, 244

memory mapped, 179–80

peripheral or isolated, 180

input/output, 5

instruction design, 71, 86

manual coding, 86

Instruction Level Parallelism (ILP), 557

instruction queue, 58

FIFO, 58

instructions, 3, 5, 7–8, 11, 12, 27, 31–8, 40–1, 44, 71–2, 74, 77–9, 81, 84–9, 91–2

arithmetic, 1, 5–6, 19, 22, 29, 31–2, 36–8

branch, 7, 34, 36, 38–9, 58

call, 5, 7, 12, 38

compare, 40–1

complex, 12

data transfer, 6, 36, 38, 46

decimal adjust, 36–7

decode, 7–8

execute, 7–8

fetch, 7–8

flag control, 114, 609

LODS, 144–45, 596, 609, 613

logical, 36, 129–130

RET, 38
 return, 19, 38
 rotate, 40
 SHIFT, 364, 367–68
 source, 11, 35–6, 40, 71–3, 78, 88–9
 STOS, 144–45, 596, 609, 615
 STRING, 157, 174–76, 596
 target, 71
 INT 10H, 259, 265–71, 279–84
 intel, 1–5, 11–2, 31, 545, 548–49, 551, 558
 centrino technology, 584
 intel hub architecture, 566
 interfacing, 289–92, 298–99, 313, 316–17, 322, 327, 336–37
 hex keyboard, 289, 337–39
 peripheral, 289–91, 302–04, 310
 interfacing chip, 10
 keyboard display, 10
 interfacing to the 8086, 289
 analog to digital converter, 313
 digital to analog converter, 317–18
 LED Displays, 341
 liquid crystal displays, 322
 stepper motor, 289, 327–37
 interleaved memory architecture, 440
 internal reset, 407, 411, 414–15
 inter processor communication, 214, 455, 456
 interrupt, 6–7, 44–5, 46
 acknowledge, 6, 46
 breakpoint, 254, 256
 execution, 7–8, 12, 32, 34–5, 38, 43, 45
 flag, 33–4, 40, 55–7, 250
 handler, 251, 278, 377, 390, 533, 537
 hooking, 277, 280
 instruction, 1, 3, 7–8, 11–2, 31–2, 34–40, 45, 52, 54, 56, 58, 60–1, 67, 195, 199, 206, 213, 215, 217–21
 service routine, 45
 vector, 45–46
 Interrupt Descriptor Table (IDT), 503, 537–38
 exceptions, 503, 515, 538–40
 interrupt hooks, 278
 Interrupt Service Routine (ISR), 250–51
 interrupt types, 253, 255
 breakpoint interrupt, 254, 256
 divide by zero error, 253, 255–56
 non maskable interrupt, 254–56, 260
 overflow interrupt, 255
 single stepping, 254
 interrupt vector table, 249, 252–53, 255–57, 274, 278–79
 interrupts, 249–50, 252, 255–59, 261, 272, 274–75, 277–78, 280, 283
 BIOS, 249, 253, 255–56, 258–61, 264–65, 267, 270, 272, 274–77, 283

DOS, 249, 255–56, 259–60, 269, 272, 274, 277–78, 280–81, 283–85
 hardware, 249–50, 254–56, 258–61, 272–74, 280, 285
 of 8086, 249–50, 252, 256
 priority of, 259
 software, 249–50, 255–56, 259–61, 274
 intrasegment call, 148
 direct, 10–2, 35, 46, 148–49, 160, 166, 173, 443, 447–50
 indirect, 35–7, 39, 148–49
 intersegment call, 149
 direct, 10–2, 35, 46, 148–49, 160, 166, 173, 443, 447–50
 indirect, 35–7, 39, 148–49
 IRQ line, 368, 390–91

J

jump instruction, 107–12
 far, 99, 107, 111–12, 121
 inter-segment, 107
 near, 107–10
 SHORT, 108–09
 unconditional, 95, 107–10, 112
 jumpers, 563, 570

K

key debounce, 338–40
 key release, 273–74
 keyboard display interface, 363
 programmable, 289–91, 322–23, 349–50, 355, 363, 371, 377–79, 390–91
 keyboard scans, 367
 keypress, 367–68, 374
 keys, 272–76, 281–82
 ALT, 272, 274, 276, 280–81, 283–86
 application, 255, 258, 273, 277
 Caps Lock, 273–74
 character, 260–64, 266–78, 283
 CTRL, 272, 274, 282–85
 Enter (Return), 272
 Escape, 272
 Function, 256, 258–61, 264–70, 272, 274–81, 283–84
 Num Lock, 273
 windows, 266, 273

L

label, 72–3, 82, 85–6
 language, 1, 3, 5, 10–1, 34
 assembly, 1, 5, 10–1, 34

computer, 1–2, 4–12, 17, 19, 22, 29, 38, 272–73
 high level, 1, 10–1
 lower level, 11
 machine, 3–5, 10–1, 32
 opcodes, 10
 laptops, 576, 583–86
 latches, 195–96, 200–01, 209, 216
 LCD, 292, 322–26, 341
 pins of, 289, 297, 300, 304, 306, 308, 310, 313, 315–16, 322–23, 332
 LED, 289, 291, 297–98, 322–23, 336, 341–45
 seven segment, 341–44
 left entry, 369–70, 373
 legacy support, 559
 level triggered mode, 378
 lexical nesting, 489
 limit checking, 528–29
 line print terminal, 312
 linking, 72, 74–5
 list file, 76
 Load Effective Address (LEA), 101
 Local Descriptor Table Register (LDTR), 520
 locality of reference, 446
 location counter, 73, 79
 lock, 196, 215, 217–18
 logical address space, 523
 calculating the size of, 523
 loosely coupled configurations, 457–58, 464
 LPC, 565–66

M

machine codes, 72, 77, 86
 macros, 139, 155–58
 local directive, 157
 magnitude, 19, 22, 29, 40
 mapping techniques, 431, 447, 450
 direct, 447–50
 fully associative, 447–50
 set associative, 449–51
 MASM 6, 14, 96
 master 8, 259, 390, 391
 Maximum Count (MC), 496–97, 499, 501
 maximum count register, 496–97, 499
 maximum mode, 195, 213–18
 memory, 2, 4–10, 30–2, 34–7, 39–46
 address, 463, 489, 505, 525, 546, 548
 bank, 200, 213, 235–39, 244–45, 298, 440–41, 512, 547–50
 primary, 8, 431, 514
 secondary, 8, 431, 445–46, 514, 524
 virtual, 2, 4, 431, 503, 514, 516, 524, 576
 memory access, 63
 memory banks, 235–39, 244–47

memory capacity, 30–31
 units of, 6, 30
 memory controller, 566, 568, 577
 memory management unit, 515, 524, 545, 552
 memory models, 71, 75, 85
 full segment, 71, 75, 84
 tiny, 75, 77–80, 85
 memory organization, 62
 memory read, 6–9
 memory write, 6–9
 microchip, 1, 565
 microcontroller, 487, 495
 microprocessor design, 555
 latest trends in, 555
 multi core processors, 555, 558
 multi core technology, 558
 microprocessors, 1–2, 12, 23–4, 31, 46
 embedded, 2, 12
 history, 1, 4, 12, 31
 minimum mode pin, 196–98
 functions, 195–98, 203, 214, 216, 256, 258–59, 264–65, 267, 272, 274–78, 280, 283
 MMX, 553, 554
 mnemonics, 10
 mobile processors, 559
 mode control word, 405, 407–08, 410, 412, 414
 mode instruction format, 408, 411
 mode, 2, 13, 35
 asynchronous, 396, 403, 405–08, 410–12, 415
 bidirectional, 6, 7
 hardware triggered, 355, 362
 protected, 2
 software triggered, 355, 362
 modem, 397–98, 403, 562, 566
 modules, 179, 183–86, 188–90
 motherboard, 561–67, 569–71, 575, 577, 579–81, 583–84
 accelerated graphics port, 562, 565
 sound card, 565–66
 video card, 565–66, 568
 peripheral component interconnect, 563, 569
 move (MOV), 98
 direct addressing, 100
 MSB, 19, 22, 25, 28–30, 34
 MS-DOS, 4
 MUL, 10
 multi core processing, 545, 555, 557
 multi tasking, 504, 534
 issues in, 532, 535
 systems, 503, 514–15, 526, 533–35
 multibus, 467, 470–71
 multiplexed dynamic display, 364
 multiplication, 114, 124–26, 131
 signed, 108, 111, 113–14, 131, 133, 139, 148, 166–72
 unsigned, 95, 111, 114, 121, 124, 126, 131
 multiprocessor, 455–59, 462, 466–67, 469–71

 using 8086, 457, 472
 multiprocessor systems, 456–57
 tightly coupled systems, 456, 464
 multithreading, 555–57

N

NAND, 294
 negative numbers, 1, 19, 21, 24, 29
 addition of, 22–4, 28–9, 36
 representation of, 17–9, 21–2
 netbook, 584
 network and firewire, 566
 nibble, 18, 23, 27
 NMI pin, 256
 normalization, 473–74
 north bridge, 563, 565, 567–68
 null modem cables, 400
 number systems, 1, 12–4
 binary, 1–2, 10–11, 13–9, 21–6, 28
 decimal, 12–8, 21–30, 36, 37, 49
 hexadecimal, 1, 14–7, 21–2, 24–6
 numbers of different lengths, 29
 addition, 10, 17, 22–5, 27–30, 36–7, 39
 sign extended, 29–30
 NUMSLOCK, 276
 NVRAM, 444, 579–80

O

odd addresses, 236, 238
 opamp, 319
 operating system, 11
 kernel, 11
 operational command words, 385
 OCW1, 385, 387, 390
 OCW2, 385, 387, 389
 OCW3, 385, 389–90
 optocouplers, 336
 output ports, 240
 over-clocking, 546
 overflow flag, 166–68

P

packed bcd, 18, 23, 26–7
 addition of, 22–4, 28–9, 36
 subtraction of, 27–9
 packing density, 438, 445
 page size, 552
 page table entry, 525–26
 page translation, 515, 524
 paging, 505, 515–16, 524–29, 535
 parallel arbitration, 465, 469–70
 implementation of, 469, 472
 parameters, 151–53, 156–57
 passing through memory, 152

 passing through registers, 151
 passing To and From procedures, 151
 parasitic capacitance, 438
 PC-AT, 503
 PC-XT, 503
 PC clones, 503
 PC, 2, 4, 5, 12–3, 30–1, 34, 44
 desktops, 3
 laptops, 4
 servers, 4
 supercomputers, 4
 PCI based computers, 428
 PCI, 561, 563, 565–71, 580–81, 584
 express, 561, 565, 568–71
 socket, 563, 568, 577, 580
 Pentium II, 553–54
 Pentium III, 553–54
 Pentium IV, 555
 Pentium M, 559
 Pentium PRO, 3, 12, 553–54
 Pentium processor, 549
 cache structure of, 545
 father of Pentium, 549
 Pentium, 3–5, 12, 29
 peripheral blocks, 487
 peripheral control block, 493, 497
 peripherals, 2, 9, 10, 44, 46
 personal computer, 561, 567
 modern, 562, 565, 567, 580, 583
 pins, 31, 41–4, 46
 configuration, 195, 205, 292, 303
 HLDA, 41, 45–6
 HOLD, 31, 34, 39, 41, 45–6
 INTR, 41, 45–6
 minimum mode, 195–200, 203, 213–14
 queue status, 217, 460
 request/grant, 457
 serial input data see SIID, 46
 serial output data see SOD, 46
 SID, 41, 45–6
 SOD, 41, 45–6
 pipelined processor, 546
 pipelining, 58
 pointer, 55, 60–1, 64–5, 99, 101, 103, 107, 112–13, 115–17, 122–23
 instruction, 52, 54, 107, 146, 478–79, 489, 505–06, 535–36
 stack, 55, 58, 60–2, 67–9, 573
 port address, 179–83
 port, 289–310, 312–13, 315–17, 319–21, 324–26, 332–34, 337–39, 343–46
 extended capabilities, 312
 extended parallel, 312
 legacy, 310
 standard parallel, 312
 USB, see universal serial bus
 power dissipation, 553, 555
 power on reset, 206
 power on, 253, 258
 pre-fetching, 58
 prefix, 140, 142–45
 REP, 140–42, 144–45

pre-scaler, 371
 privilege levels, 528, 530–35, 537
 protection, 503–04, 515–17, 520, 526, 528–30, 532, 540
 trust, 530
 privileged instructions, 539–40
 procedure calls, 150
 procedures, 139, 145, 148–49, 151, 179
 processing power, 5
 processor bus, 566–67
 processor, 1–12, 29, 31, 33–5, 41, 43–6
 program counter, 2, 33–4, 43, 45
 Program Segment Prefix (PSP), 106
 program, 2–3, 5, 7, 10, 11, 33–4, 37–9, 43, 45, 141–54, 156–57, 163–69, 171–72, 174–78
 called, 139–40, 145, 147–49, 151, 155–56, 158
 calling, 145, 148, 150–51, 171
 counter, 2, 33–4, 38–9, 43, 45
 execution, 7–8, 12, 32, 34–5, 38, 43, 45
 main, 145–47, 149–51, 153–54
 source, 11, 35–6, 40, 71–3, 78, 88–9
 programmable interval timer, 349, 391
 programmable one shot, 359
 programing using, 172
 programming in C, 189
 with assembly modules, 189
 Programming Peripheral Interface (PPI), 291
 programming, 289, 290–92, 295, 395, 407, 412, 420, 425–26, 489, 495, 497
 programming the 8087, 479
 instruction format, 480–81
 programming the 8251, 407, 412
 loop back mode, 412–13
 programming, 95–7, 105, 289–92, 295, 395, 407, 412, 420, 425–26
 approaches to, 95
 input/output, 179, 183
 interactive, 96
 modular, 179, 183, 189
 top down approach, 95
 Protected Virtual Addressing Mode (PVAM), 504, 515
 protection mechanisms, 545
 protection, 503–04, 515–17, 520, 526, 528–30, 532, 540
 pseudo-instructions, 72
 PTR directive, 115
 public, 184–89

R

RAM, 1, 4, 8, 34, 52
 on-chip, 52, 58
 rate generator, 360

read and write, 201, 216
 bus timing, 195, 201
 control signals for, 201, 211, 216
 real mode, 504–06, 513–14, 516, 540
 reduced instruction set computer architecture see RISC architecture
 re-entrant, 280
 procedure, 255–56, 279–80
 program, 249–51, 254–55, 258, 260–61, 265–69, 271–73, 276–85
 refreshing rate, 439
 register, 31–7, 39, 51–2, 54–61, 63–7
 flag, 33–4, 40, 55–7, 250
 general purpose, 32
 scratchpad, 32, 52
 registers, 51–2, 54–5, 58–67
 internal, 6, 52, 352, 419–20, 423, 459
 program visible, 52
 requested privilege level, 520, 531
 reset logic, 493
 resistors, 242
 right entry, 369, 370–71, 374, 376
 RISC architecture, 12
 ROM, 1, 4
 ports, 1–2, 45
 RS-232, 395, 398–400, 402, 413
 connectors, 399–400
 DCE and DTE devices, 399
 level converters, 399
 standards, 395, 398, 402
 RS232c, 290

S

scan code, 273–76, 281–82, 284–86
 Schmitt trigger, 206–07
 segment override, 87, 92
 prefix, 67
 segment registers, 52, 58–62, 64
 segment, 52, 58–62, 64–67
 selectable CAS latency, 441
 selector, 515, 520–23, 527, 529–31, 533–38
 segment, 52, 58–62, 64–67, 505–06, 515–36, 538–40
 semiconductor memory, 431–32
 dynamic RAM, 435, 439, 446
 static RAM, 432, 446
 serial communication, 10
 programmable, 403
 servers, 561, 569
 seven segment displays, 341–42
 dynamic, 289, 342–44
 multiplexed, 289, 313–14, 342, 344
 static, 342
 shadow memory, 259
 sign extension, 21, 30
 signed numbers, 166, 168–70
 comparison of, 168–69
 significant, 479
 Simultaneous Multithreading (SMT), 557
 Single In Line Memory Module (SIMM), 576
 single precision, 472–73, 475–76
 small model, 75, 80–1
 software, 5, 11–2, 45
 software triggered mode, 362
 south bridge, 563, 565–66, 568, 573, 580, 582
 special values, 474, 476
 denormalized, 477
 infinity, 477
 not a number, 477
 special operations, 477
 zero, 473, 476–77, 479–80
 square wave generator, 362
 SRAM chip, 432–34
 internal architecture of, 431, 434, 437
 SSE, 553, 554
 stack pointer, 2, 33–4
 stack, 55, 58, 60–2, 99, 104–06, 119, 128, 146–47, 149–51, 153–55, 573
 defining a, 105–06
 operation, 98, 104–06, 112, 114, 118, 120–21, 128–33, 135
 overflow, 155, 166–69, 171
 underflow, 155
 status byte, 274–76, 281, 284
 status register, 352, 368, 374–76
 status word, 292, 305, 410–11
 step angle, 331–33
 stepper motor, 289, 327–37
 bipolar, 329–30
 driving a, 330
 two-phase, 329
 types, 328
 unipolar, 329–30
 string instructions, 139, 140, 144–45
 CMPS, 142
 MOVS, 140
 pre-requisites, 139–40
 SCAS, 143
 strobe, 355
 hardware triggered, 355, 362
 software triggered, 355, 362
 strobed mode, 302, 304
 input, 302, 305
 output, 305, 308
 subtraction, 10, 25–9, 40, 114, 119–21
 signed numbers, 22, 25, 27–8, 29–30
 unsigned numbers, 22, 25, 29
 superscalar architecture, 545, 549–50, 557
 switching to protected mode, 540
 symbols, 72
 synchronous vs asynchronous, 442
 syntax errors, 74

system BIOS, 578–79
 flash BIOS, 563, 567, 579
 ROM BIOS and motherboards, 579
 system bus, 6–7, 10
 system descriptors, 518
 system management mode, 553
 system on chip, 584

T

task linking, 538
 task register, 536–37, 540
 task switching, 503, 534–35, 537–38
 terminal count, 355, 358
 terminate and stay resident, 277–78
 Thread Level Parallelism (TLP), 557
 thread, 557
 time, 1, 4, 6, 8, 12, 31, 34, 39, 42, 44, 46
 access, 6, 8, 11, 31, 43–4
 timer unit, 488, 494–95
 timers, 289, 350–51
 timing, 495–97, 501
 events, 495–96
 sequence, 494–95, 499
 trainer kit, 290, 298
 transfer speed, 566
 translate, 99, 102–03

Translation Lookaside Buffer (TLB), 526–27
 translator, 71–2
 transmission rate, 397
 trigger pin, 351
 tri-state buffer, 226, 242–44
 TSB, 257–58
 TSR, 249, 278, 280–81, 283, 285
 TTL, 226
 type checking, 528–29
 type number, 252–60, 264, 274, 277–78, 282, 284
 generation of, 257–58
 typematic, 274

U

universal serial bus, 310, 571–72
 features of, 72, 78, 84, 461, 472, 477, 561, 568, 570, 572
 power issues, 572
 unsigned numbers, 22, 25, 29
 addition of, 22–4, 28, 29–36
 USART 8251, 395, 403
 block diagram, 350, 351, 353, 363, 368, 373, 378, 379, 404, 405, 419–22
 modem control block, 406
 pin diagram, 400, 404
 programming, 95–7, 105
 receive buffer, 405
 USES construct, 150–51, 160

V

video adaptor, 261
 video card, 4
 monochrome, 4
 virtual 8086 mode, 513
 virtual memory, 503–04, 514, 516, 524

W

wait cycles, 210
 wait state generator, 487
 waveform, 289, 300, 303, 305, 319–21
 saw-tooth, 321
 staircase, 320–21
 triangular, 319
 write machine cycle, 211
 write strobes, 239

X

x86, 2–5, 11–2, 31

Z

Z-80, 2
 ZF flag, 509

The x86 Microprocessors

Architecture, Programming and Interfacing (8086 to Pentium)

Lyla B. Das

The book is designed for an undergraduate course on the 16-bit microprocessor and Pentium processor. The text comprehensively covers both the hardware and software aspects of the subject with equal emphasis on architecture, programming and interfacing.

FEATURES



A thorough description of the binary math required for proficiency in assembly programming



Comprehensive analysis of programming and interfacing, with practical examples



Discusses the features and enhancements of the 80386, 80486 and Pentium processors



All concepts are presented with worked-out examples and programs



A chapter devoted to the internal details of PCs and the current trends in computer design

Lyla B. Das is Associate Professor, Department of Electronics and Communication Engineering, National Institute of Technology Calicut, Kozhikode, Kerala.



Online resources available at
www.pearsoned.co.in/lylabdas

ISBN 978-81-317-3246-5



9 788131 732465

www.pearsoned.co.in